
MongoDB Reference Manual

Release 2.6.1

MongoDB Documentation Project

June 05, 2014

1	About MongoDB Documentation	3
1.1	License	3
1.2	Editions	3
1.3	Version and Revisions	4
1.4	Report an Issue or Make a Change Request	4
1.5	Contribute to the Documentation	4
2	Interfaces Reference	21
2.1	mongo Shell Methods	21
2.2	Database Commands	198
2.3	Operators	373
2.4	Aggregation Reference	484
3	MongoDB and SQL Interface Comparisons	495
3.1	SQL to MongoDB Mapping Chart	495
3.2	SQL to Aggregation Mapping Chart	500
4	Program and Tool Reference Pages	503
4.1	MongoDB Package Components	503
5	Internal Metadata	593
5.1	Config Database	593
5.2	The local Database	598
5.3	System Collections	600
6	General System Reference	603
6.1	Exit Codes and Statuses	603
6.2	MongoDB Limits and Thresholds	604
6.3	Glossary	609
7	Release Notes	619
7.1	Current Stable Release	619
7.2	Previous Stable Releases	647
7.3	Other MongoDB Release Notes	692
	Index	693

This document contains all of the reference material from the `MongoDB Manual`, reflecting the 2.6.1 release. See the full manual, for complete documentation of MongoDB, it's operation, and use.

About MongoDB Documentation

The [MongoDB Manual](#)¹ contains comprehensive documentation on the MongoDB *document*-oriented database management system. This page describes the manual's licensing, editions, and versions, and describes how to make a change request and how to contribute to the manual.

For more information on MongoDB, see [MongoDB: A Document Oriented Database](#)². To download MongoDB, see the [downloads page](#)³.

1.1 License

This manual is licensed under a Creative Commons “[Attribution-NonCommercial-ShareAlike 3.0 Unported](#)”⁴ (i.e. “CC-BY-NC-SA”) license.

The MongoDB Manual is copyright © 2011-2014 MongoDB, Inc.

1.2 Editions

In addition to the [MongoDB Manual](#)⁵, you can also access this content in the following editions:

- [ePub Format](#)⁶
- [Single HTML Page](#)⁷
- [PDF Format](#)⁸ (without reference.)
- [HTML tar.gz](#)⁹

You also can access PDF files that contain subsets of the MongoDB Manual:

- [MongoDB Reference Manual](#)¹⁰
- [MongoDB CRUD Operations](#)¹¹

¹<http://docs.mongodb.org/manual/#>

²<http://www.mongodb.org/about/>

³<http://www.mongodb.org/downloads>

⁴<http://creativecommons.org/licenses/by-nc-sa/3.0/>

⁵<http://docs.mongodb.org/manual/#>

⁶<http://docs.mongodb.org/master/MongoDB-manual.epub>

⁷<http://docs.mongodb.org/master/single/>

⁸<http://docs.mongodb.org/master/MongoDB-manual.pdf>

⁹<http://docs.mongodb.org/master/manual.tar.gz>

¹⁰<http://docs.mongodb.org/master/MongoDB-reference-manual.pdf>

¹¹<http://docs.mongodb.org/master/MongoDB-crud-guide.pdf>

- [Data Models for MongoDB](#)¹²
- [MongoDB Data Aggregation](#)¹³
- [Replication and MongoDB](#)¹⁴
- [Sharding and MongoDB](#)¹⁵
- [MongoDB Administration](#)¹⁶
- [MongoDB Security](#)¹⁷

MongoDB Reference documentation is also available as part of [dash](#)¹⁸. You can also access the [MongoDB Man Pages](#)¹⁹ which are also distributed with the official MongoDB Packages.

1.3 Version and Revisions

This version of the manual reflects version 2.6 of MongoDB.

See the [MongoDB Documentation Project Page](#)²⁰ for an overview of all editions and output formats of the MongoDB Manual. You can see the full revision history and track ongoing improvements and additions for all versions of the manual from its [GitHub repository](#)²¹.

This edition reflects “master” branch of the documentation as of the “12da929a91193eefa8bcc6bb55382301a7546232” revision. This branch is explicitly accessible via “<http://docs.mongodb.org/master>” and you can always reference the commit of the current manual in the [release.txt](#)²² file.

The most up-to-date, current, and stable version of the manual is always available at “<http://docs.mongodb.org/manual/>”.

1.4 Report an Issue or Make a Change Request

To report an issue with this manual or to make a change request, file a ticket at the [MongoDB DOCS Project on Jira](#)²³.

1.5 Contribute to the Documentation

1.5.1 MongoDB Manual Translation

The original authorship language for all MongoDB documentation is American English. However, ensuring that speakers of other languages can read and understand the documentation is of critical importance to the documentation project.

¹²<http://docs.mongodb.org/master/MongoDB-data-models-guide.pdf>

¹³<http://docs.mongodb.org/master/MongoDB-aggregation-guide.pdf>

¹⁴<http://docs.mongodb.org/master/MongoDB-replication-guide.pdf>

¹⁵<http://docs.mongodb.org/master/MongoDB-sharding-guide.pdf>

¹⁶<http://docs.mongodb.org/master/MongoDB-administration-guide.pdf>

¹⁷<http://docs.mongodb.org/master/MongoDB-security-guide.pdf>

¹⁸<http://kapeli.com/dash>

¹⁹<http://docs.mongodb.org/master/manpages.tar.gz>

²⁰<http://docs.mongodb.org>

²¹<https://github.com/mongodb/docs>

²²<http://docs.mongodb.org/master/release.txt>

²³<https://jira.mongodb.org/browse/DOCS>

In this direction, the MongoDB Documentation project uses the service provided by [Smartling](#)²⁴ to translate the MongoDB documentation into additional non-English languages. This translation project is largely supported by the work of volunteer translators from the MongoDB community who contribute to the translation effort.

If you would like to volunteer to help translate the MongoDB documentation, please:

- complete the [MongoDB Contributor Agreement](#)²⁵, and
- create an account on Smartling at translate.docs.mongodb.org²⁶.

Please use the same email address you use to sign the contributor as you use to create your Smartling account.

The [mongodb-translators](#)²⁷ user group exists to facilitate collaboration between translators and the documentation team at large. You can join the Google Group without signing the contributor's agreement.

We currently have the following languages configured:

- [Arabic](#)²⁸
- [Chinese](#)²⁹
- [Czech](#)³⁰
- [French](#)³¹
- [German](#)³²
- [Hungarian](#)³³
- [Indonesian](#)³⁴
- [Italian](#)³⁵
- [Japanese](#)³⁶
- [Korean](#)³⁷
- [Lithuanian](#)³⁸
- [Polish](#)³⁹
- [Portuguese](#)⁴⁰
- [Romanian](#)⁴¹
- [Russian](#)⁴²
- [Spanish](#)⁴³

²⁴<http://smartling.com/>

²⁵<http://www.mongodb.com/legal/contributor-agreement>

²⁶<http://translate.docs.mongodb.org/>

²⁷<http://groups.google.com/group/mongodb-translators>

²⁸<http://ar.docs.mongodb.org>

²⁹<http://cn.docs.mongodb.org>

³⁰<http://cs.docs.mongodb.org>

³¹<http://fr.docs.mongodb.org>

³²<http://de.docs.mongodb.org>

³³<http://hu.docs.mongodb.org>

³⁴<http://id.docs.mongodb.org>

³⁵<http://it.docs.mongodb.org>

³⁶<http://jp.docs.mongodb.org>

³⁷<http://ko.docs.mongodb.org>

³⁸<http://lt.docs.mongodb.org>

³⁹<http://pl.docs.mongodb.org>

⁴⁰<http://pt.docs.mongodb.org>

⁴¹<http://ro.docs.mongodb.org>

⁴²<http://ru.docs.mongodb.org>

⁴³<http://es.docs.mongodb.org>

- [Turkish](#)⁴⁴
- [Ukrainian](#)⁴⁵

If you would like to initiate a translation project to an additional language, please report this issue using the “*Report a Problem*” link above or by posting to the [mongodb-translators](#)⁴⁶ list.

Currently the translation project only publishes rendered translation. While the translation effort is currently focused on the web site we are evaluating how to retrieve the translated phrases for use in other media.

See also:

- [Contribute to the Documentation](#) (page 4)
- [Style Guide and Documentation Conventions](#) (page 6)
- [MongoDB Manual Organization](#) (page 15)
- [MongoDB Documentation Practices and Processes](#) (page 12)
- [MongoDB Documentation Build System](#) (page 16)

The entire documentation source for this manual is available in the [mongodb/docs repository](#)⁴⁷, which is one of the [MongoDB project repositories on GitHub](#)⁴⁸.

To contribute to the documentation, you can open a [GitHub account](#)⁴⁹, fork the [mongodb/docs repository](#)⁵⁰, make a change, and issue a pull request.

In order for the documentation team to accept your change, you must complete the [MongoDB Contributor Agreement](#)⁵¹.

You can clone the repository by issuing the following command at your system shell:

```
git clone git://github.com/mongodb/docs.git
```

1.5.2 About the Documentation Process

The MongoDB Manual uses [Sphinx](#)⁵², a sophisticated documentation engine built upon [Python Docutils](#)⁵³. The original [reStructured Text](#)⁵⁴ files, as well as all necessary Sphinx extensions and build tools, are available in the same repository as the documentation.

For more information on the MongoDB documentation process, see:

Style Guide and Documentation Conventions

This document provides an overview of the style for the MongoDB documentation stored in this repository. The overarching goal of this style guide is to provide an accessible base style to ensure that our documentation is easy to read, simple to use, and straightforward to maintain.

For information regarding the MongoDB Manual organization, see [MongoDB Manual Organization](#) (page 15).

⁴⁴<http://tr.docs.mongodb.org>

⁴⁵<http://uk.docs.mongodb.org>

⁴⁶<http://groups.google.com/group/mongodb-translators>

⁴⁷<https://github.com/mongodb/docs>

⁴⁸<http://github.com/mongodb>

⁴⁹<https://github.com/>

⁵⁰<https://github.com/mongodb/docs>

⁵¹<http://www.mongodb.com/contributor>

⁵²<http://sphinx-doc.org/>

⁵³<http://docutils.sourceforge.net/>

⁵⁴<http://docutils.sourceforge.net/rst.html>

Document History

2011-09-27: Document created with a (very) rough list of style guidelines, conventions, and questions.

2012-01-12: Document revised based on slight shifts in practice, and as part of an effort of making it easier for people outside of the documentation team to contribute to documentation.

2012-03-21: Merged in content from the Jargon, and cleaned up style in light of recent experiences.

2012-08-10: Addition to the “Referencing” section.

2013-02-07: Migrated this document to the manual. Added “map-reduce” terminology convention. Other edits.

2013-11-15: Added new table of preferred terms.

Naming Conventions

This section contains guidelines on naming files, sections, documents and other document elements.

- File naming Convention:
 - For Sphinx, all files should have a `.txt` extension.
 - Separate words in file names with hyphens (i.e. `-`.)
 - For most documents, file names should have a terse one or two word name that describes the material covered in the document. Allow the path of the file within the document tree to add some of the required context/categorization. For example it's acceptable to have `http://docs.mongodb.org/manualcore/sharding.rst` and `http://docs.mongodb.org/manualadministration/sharding.rst`.
 - For tutorials, the full title of the document should be in the file name. For example, `http://docs.mongodb.org/manualtutorial/replace-one-configuration-server-in-a-shard-`
- Phrase headlines and titles so users can determine what questions the text will answer, and material that will be addressed, without needing them to read the content. This shortens the amount of time that people spend looking for answers, and improvise search/scanning, and possibly “SEO.”
- Prefer titles and headers in the form of “Using foo” over “How to Foo.”
- When using target references (i.e. `:ref:` references in documents), use names that include enough context to be intelligible through all documentation. For example, use “`replica-set-secondary-only-node`” as opposed to “`secondary-only-node`”. This makes the source more usable and easier to maintain.

Style Guide

This includes the local typesetting, English, grammatical, conventions and preferences that all documents in the manual should use. The goal here is to choose good standards, that are clear, and have a stylistic minimalism that does not interfere with or distract from the content. A uniform style will improve user experience and minimize the effect of a multi-authored document.

Punctuation

- Use the Oxford comma.

Oxford commas are the commas in a list of things (e.g. “something, something else, and another thing”) before the conjunction (e.g. “and” or “or”).
- Do not add two spaces after terminal punctuation, such as periods.

- Place commas and periods inside quotation marks.

Headings Use title case for headings and document titles. Title case capitalizes the first letter of the first, last, and all significant words.

Verbs Verb tense and mood preferences, with examples:

- **Avoid** the first person. For example do not say, “We will begin the backup process by locking the database,” or “I begin the backup process by locking my database instance.”
- **Use** the second person. “If you need to back up your database, start by locking the database first.” In practice, however, it’s more concise to imply second person using the imperative, as in “Before initiating a backup, lock the database.”
- When indicated, use the imperative mood. For example: “Backup your databases often” and “To prevent data loss, back up your databases.”
- The future perfect is also useful in some cases. For example, “Creating disk snapshots without locking the database will lead to an invalid state.”
- Avoid helper verbs, as possible, to increase clarity and concision. For example, attempt to avoid “this does foo” and “this will do foo” when possible. Use “does foo” over “will do foo” in situations where “this foos” is unacceptable.

Referencing

- To refer to future or planned functionality in MongoDB or a driver, *always* link to the Jira case. The Manual’s `conf.py` provides an `:issue:` role that links directly to a Jira case (e.g. `:issue:\`SERVER-9001\``).
- For non-object references (i.e. functions, operators, methods, database commands, settings) always reference only the first occurrence of the reference in a section. You should *always* reference objects, except in section headings.
- Structure references with the *why* first; the link second.

For example, instead of this:

Use the [http://docs.mongodb.org/manual/tutorial/convert-replica-set-to-replicated-shard-cl](http://docs.mongodb.org/manual/tutorial/convert-replica-set-to-replicated-shard-cluster/) procedure if you have an existing replica set.

Type this:

To deploy a sharded cluster for an existing replica set, see [http://docs.mongodb.org/manual/tutorial/convert-re](http://docs.mongodb.org/manual/tutorial/convert-replica-set-to-replicated-shard-cluster/)

General Formulations

- Contractions are acceptable insofar as they are necessary to increase readability and flow. Avoid otherwise.
- Make lists grammatically correct.
 - Do not use a period after every item unless the list item completes the unfinished sentence before the list.
 - Use appropriate commas and conjunctions in the list items.
 - Typically begin a bulleted list with an introductory sentence or clause, with a colon or comma.
- The following terms are one word:
 - standalone
 - workflow

- Use “unavailable,” “offline,” or “unreachable” to refer to a `mongod` instance that cannot be accessed. Do not use the colloquialism “down.”
- Always write out units (e.g. “megabytes”) rather than using abbreviations (e.g. “MB”).

Structural Formulations

- There should be at least two headings at every nesting level. Within an “h2” block, there should be either: no “h3” blocks, 2 “h3” blocks, or more than 2 “h3” blocks.
- Section headers are in title case (capitalize first, last, and all important words) and should effectively describe the contents of the section. In a single document you should strive to have section titles that are not redundant and grammatically consistent with each other.
- Use paragraphs and paragraph breaks to increase clarity and flow. Avoid burying critical information in the middle of long paragraphs. Err on the side of shorter paragraphs.
- Prefer shorter sentences to longer sentences. Use complex formations only as a last resort, if at all (e.g. compound complex structures that require semi-colons).
- Avoid paragraphs that consist of single sentences as they often represent a sentence that has unintentionally become too complex or incomplete. However, sometimes such paragraphs are useful for emphasis, summary, or introductions.

As a corollary, most sections should have multiple paragraphs.

- For longer lists and more complex lists, use bulleted items rather than integrating them inline into a sentence.
- Do not expect that the content of any example (inline or blocked) will be self explanatory. Even when it feels redundant, make sure that the function and use of every example is clearly described.

ReStructured Text and Typesetting

- Place spaces between nested parentheticals and elements in JavaScript examples. For example, prefer `{ [a, a, a] }` over `{[a,a,a]}`.
- For underlines associated with headers in RST, use:
 - `=` for heading level 1 or h1s. Use underlines and overlines for document titles.
 - `--` for heading level 2 or h2s.
 - `~` for heading level 3 or h3s.
 - ``` for heading level 4 or h4s.

- Use hyphens (`-`) to indicate items of an ordered list.
- Place footnotes and other references, if you use them, at the end of a section rather than the end of a file.

Use the footnote format that includes automatic numbering and a target name for ease of use. For instance a footnote tag may look like: `[#note]_` with the corresponding directive holding the body of the footnote that resembles the following: `.. [#note]`.

Do **not** include `.. code-block:: [language]` in footnotes.

- As it makes sense, use the `.. code-block:: [language]` form to insert literal blocks into the text. While the double colon, `::`, is functional, the `.. code-block:: [language]` form makes the source easier to read and understand.
- For all mentions of referenced types (i.e. commands, operators, expressions, functions, statuses, etc.) use the reference types to ensure uniform formatting and cross-referencing.

Jargon and Common Terms

Preferred Term	Concept	Dispreferred Alternatives	Notes
<i>document</i>	A single, top-level object/record in a MongoDB collection.	record, object, row	Prefer document over object because of concerns about cross-driver language handling of objects. Reserve record for “allocation” of storage. Avoid “row,” as possible.
<i>database</i>	A group of collections. Refers to a group of data files. This is the “logical” sense of the term “database.”		Avoid genericizing “database.” Avoid using database to refer to a server process or a data set. This applies both to the datastoring contexts as well as other (related) operational contexts (command context, authentication/authorization context.)
instance	A daemon process. (e.g. mongos or mongod)	process (acceptable sometimes), node (never acceptable), server.	Avoid using instance, unless it modifies something specifically. Having a descriptor for a process/instance makes it possible to avoid needing to make mongod or mongos plural. Server and node are both vague and contextually difficult to disambiguate with regards to application servers, and underlying hardware.
<i>field name</i>	The identifier of a value in a document.	key, column	Avoid introducing unrelated terms for a single field. In the documentation we’ve rarely had to discuss the identifier of a field, so the extra word here isn’t burdensome.
<i>field/value</i>	The name/value pair that describes a unit of data in MongoDB.	key, slot, attribute	Use to emphasize the difference between the name of a field and its value. For example, “_id” is the field and the default value is an ObjectId.
value	The data content of a field.	data	
Mon- goDB	A group of processes, or deployment that implement the MongoDB interface.	mongo, mongodb, cluster	Stylistic preference, mostly. In some cases it’s useful to be able to refer generically to instances (that may be either mongod or mongos .)
sub- document	An embedded or nested document within a document or an array.	embedded document, nested document	
<i>map- reduce</i>	An operation performed by the mapReduce command.	mapReduce, map reduce, map/reduce	Avoid confusion with the command, shell helper, and driver interfaces. Makes it possible to discuss the operation generally.
clus- ter	A sharded cluster.	grid, shard cluster, set, deployment	Cluster is a great word for a group of processes; however, it’s important to avoid letting the term become generic. Do not use for any group of MongoDB processes or deployments.
sharded clus- ter	A <i>sharded cluster</i> .	shard cluster, cluster, sharded system	
<i>replica set</i>	A deployment of replicating mongod programs that provide redundancy and automatic failover.	set, replication deployment	
de- ploy- ment	A group of MongoDB processes, or a standalone mongod instance.	cluster, system	Typically in the form MongoDB deployment.
data set	The collection of physical databases provided by a MongoDB deployment.	database, data	Includes standalones, replica sets and sharded clusters. Important to keep the distinction between the data provided by a mongod or a sharded cluster as distinct from each “database” (i.e. a logical

Database Systems and Processes

- To indicate the entire database system, use “MongoDB,” not `mongo` or `Mongo`.
- To indicate the database process or a server instance, use `mongod` or `mongos`. Refer to these as “processes” or “instances.” Reserve “database” for referring to a database structure, i.e., the structure that holds collections and refers to a group of files on disk.

Distributed System Terms

- Refer to partitioned systems as “sharded clusters.” Do not use `shard clusters` or `sharded systems`.
- Refer to configurations that run with replication as “replica sets” (or “master/slave deployments”) rather than “clusters” or other variants.

Data Structure Terms

- “document” refers to “rows” or “records” in a MongoDB database. Potential confusion with “JSON Documents.”

Do not refer to documents as “objects,” because drivers (and MongoDB) do not preserve the order of fields when fetching data. If the order of objects matter, use an array.
- “field” refers to a “key” or “identifier” of data within a MongoDB document.
- “value” refers to the contents of a “field”.
- “sub-document” describes a nested document.

Other Terms

- Use `example.net` (and `.org` or `.com` if needed) for all examples and samples.
- Hyphenate “map-reduce” in order to avoid ambiguous reference to the command name. Do not camel-case.

Notes on Specific Features

- Geo-Location
 1. While MongoDB *is capable* of storing coordinates in sub-documents, in practice, users should only store coordinates in arrays. (See: [DOCS-41](#)⁵⁵.)

MongoDB Documentation Practices and Processes

This document provides an overview of the practices and processes.

Commits

When relevant, include a Jira case identifier in a commit message. Reference documentation cases when applicable, but feel free to reference other cases from jira.mongodb.org⁵⁶.

Err on the side of creating a larger number of discrete commits rather than bundling large set of changes into one commit.

⁵⁵<https://jira.mongodb.org/browse/DOCS-41>

⁵⁶<http://jira.mongodb.org/>

For the sake of consistency, remove trailing whitespaces in the source file.

“Hard wrap” files to between 72 and 80 characters per-line.

Standards and Practices

- At least two people should vet all non-trivial changes to the documentation before publication. One of the reviewers should have significant technical experience with the material covered in the documentation.
- All development and editorial work should transpire on GitHub branches or forks that editors can then merge into the publication branches.

Collaboration

To propose a change to the documentation, do either of the following:

- Open a ticket in the [documentation project](#)⁵⁷ proposing the change. Someone on the documentation team will make the change and be in contact with you so that you can review the change.
- Using [GitHub](#)⁵⁸, fork the [mongodb/docs repository](#)⁵⁹, commit your changes, and issue a pull request. Someone on the documentation team will review and incorporate your change into the documentation.

Builds

Building the documentation is useful because [Sphinx](#)⁶⁰ and docutils can catch numerous errors in the format and syntax of the documentation. Additionally, having access to an example documentation as it *will* appear to the users is useful for providing more effective basis for the review process. Besides Sphinx, Pygments, and Python-Docutils, the documentation repository contains all requirements for building the documentation resource.

Talk to someone on the documentation team if you are having problems running builds yourself.

Publication

The makefile for this repository contains targets that automate the publication process. Use `make html` to publish a test build of the documentation in the `build/` directory of your repository. Use `make publish` to build the full contents of the manual from the current branch in the `../public-docs/` directory relative the docs repository.

Other targets include:

- `man` - builds UNIX Manual pages for all MongoDB utilities.
- `push` - builds and deploys the contents of the `../public-docs/`.
- `pdfs` - builds a PDF version of the manual (requires LaTeX dependencies.)

Branches

This section provides an overview of the git branches in the MongoDB documentation repository and their use.

⁵⁷<https://jira.mongodb.org/browse/DOCS>

⁵⁸<https://github.com/>

⁵⁹<https://github.com/mongodb/docs>

⁶⁰<http://sphinx.pocoo.org/>

At the present time, future work transpires in the `master`, with the main publication being `current`. As the documentation stabilizes, the documentation team will begin to maintain branches of the documentation for specific MongoDB releases.

Migration from Legacy Documentation

The MongoDB.org Wiki contains a wealth of information. As the transition to the Manual (i.e. this project and resource) continues, it's *critical* that no information disappears or goes missing. The following process outlines *how* to migrate a wiki page to the manual:

1. Read the relevant sections of the Manual, and see what the new documentation has to offer on a specific topic.
In this process you should follow cross references and gain an understanding of both the underlying information and how the parts of the new content relates its constituent parts.
2. Read the wiki page you wish to redirect, and take note of all of the factual assertions, examples presented by the wiki page.
3. Test the factual assertions of the wiki page to the greatest extent possible. Ensure that example output is accurate. In the case of commands and reference material, make sure that documented options are accurate.
4. Make corrections to the manual page or pages to reflect any missing pieces of information.
The target of the redirect need *not* contain every piece of information on the wiki page, **if** the manual as a whole does, and relevant section(s) with the information from the wiki page are accessible from the target of the redirection.
5. As necessary, get these changes reviewed by another writer and/or someone familiar with the area of the information in question.
At this point, update the relevant Jira case with the target that you've chosen for the redirect, and make the ticket unassigned.
6. When someone has reviewed the changes and published those changes to Manual, you, or preferably someone else on the team, should make a final pass at both pages with fresh eyes and then make the redirect.
Steps 1-5 should ensure that no information is lost in the migration, and that the final review in step 6 should be trivial to complete.

Review Process

Types of Review The content in the Manual undergoes many types of review, including the following:

Initial Technical Review Review by an engineer familiar with MongoDB and the topic area of the documentation. This review focuses on technical content, and correctness of the procedures and facts presented, but can improve any aspect of the documentation that may still be lacking. When both the initial technical review and the content review are complete, the piece may be “published.”

Content Review Textual review by another writer to ensure stylistic consistency with the rest of the manual. Depending on the content, this may precede or follow the initial technical review. When both the initial technical review and the content review are complete, the piece may be “published.”

Consistency Review This occurs post-publication and is content focused. The goals of consistency reviews are to increase the internal consistency of the documentation as a whole. Insert relevant cross-references, update the style as needed, and provide background fact-checking.

When possible, consistency reviews should be as systematic as possible and we should avoid encouraging stylistic and information drift by editing only small sections at a time.

Subsequent Technical Review If the documentation needs to be updated following a change in functionality of the server or following the resolution of a user issue, changes may be significant enough to warrant additional technical review. These reviews follow the same form as the “initial technical review,” but is often less involved and covers a smaller area.

Review Methods If you’re not a usual contributor to the documentation and would like to review something, you can submit reviews in any of the following methods:

- If you’re reviewing an open pull request in GitHub, the best way to comment is on the “overview diff,” which you can find by clicking on the “diff” button in the upper left portion of the screen. You can also use the following URL to reach this interface:

```
https://github.com/mongodb/docs/pull/[pull-request-id]/files
```

Replace `[pull-request-id]` with the identifier of the pull request. Make all comments inline, using GitHub’s comment system.

You may also provide comments directly on commits, or on the pull request itself but these commit-comments are archived in less coherent ways and generate less useful emails, while comments on the pull request lead to less specific changes to the document.

- Leave feedback on Jira cases in the [DOCS⁶¹](#) project. These are better for more general changes that aren’t necessarily tied to a specific line, or affect multiple files.
- Create a fork of the repository in your GitHub account, make any required changes and then create a pull request with your changes.

If you insert lines that begin with any of the following annotations:

```
.. TODO:
TODO:
.. TODO
TODO
```

followed by your comments, it will be easier for the original writer to locate your comments. The two dots `..` format is a comment in reStructured Text, which will hide your comments from Sphinx and publication if you’re worried about that.

This format is often easier for reviewers with larger portions of content to review.

MongoDB Manual Organization

This document provides an overview of the global organization of the documentation resource. Refer to the notes below if you are having trouble understanding the reasoning behind a file’s current location, or if you want to add new documentation but aren’t sure how to integrate it into the existing resource.

If you have questions, don’t hesitate to open a ticket in the [Documentation Jira Project⁶²](#) or contact the [documentation team⁶³](#).

⁶¹<http://jira.mongodb.org/browse/DOCS>

⁶²<https://jira.mongodb.org/browse/DOCS>

⁶³docs@mongodb.com

Global Organization

Indexes and Experience The documentation project has two “index files”: `http://docs.mongodb.org/manualcontents.txt` and `http://docs.mongodb.org/manualindex.txt`. The “contents” file provides the documentation’s tree structure, which Sphinx uses to create the left-pane navigational structure, to power the “Next” and “Previous” page functionality, and to provide all overarching outlines of the resource. The “index” file is not included in the “contents” file (and thus builds will produce a warning here) and is the page that users first land on when visiting the resource.

Having separate “contents” and “index” files provides a bit more flexibility with the organization of the resource while also making it possible to customize the primary user experience.

Topical Organization The placement of files in the repository depends on the *type* of documentation rather than the *topic* of the content. Like the difference between `contents.txt` and `index.txt`, by decoupling the organization of the files from the organization of the information the documentation can be more flexible and can more adequately address changes in the product and in users’ needs.

Files in the `source/` directory represent the tip of a logical tree of documents, while *directories* are containers of types of content. The `administration` and `applications` directories, however, are legacy artifacts and with a few exceptions contain sub-navigation pages.

With several exceptions in the `reference/` directory, there is only one level of sub-directories in the `source/` directory.

Tools

The organization of the site, like all Sphinx sites derives from the `toctree`⁶⁴ structure. However, in order to annotate the table of contents and provide additional flexibility, the MongoDB documentation generates `toctree`⁶⁵ structures using data from YAML files stored in the `source/includes/` directory. These files start with `ref-toc` or `toc` and generate output in the `source/includes/toc/` directory. Briefly this system has the following behavior:

- files that start with `ref-toc` refer to the documentation of API objects (i.e. commands, operators and methods), and the build system generates files that hold `toctree`⁶⁶ directives as well as files that hold *tables* that list objects and a brief description.
- files that start with `toc` refer to all other documentation and the build system generates files that hold `toctree`⁶⁷ directives as well as files that hold *definition lists* that contain links to the documents and short descriptions the content.
- file names that have `spec` following `toc` or `ref-toc` will generate aggregated tables or definition lists and allow ad-hoc combinations of documents for landing pages and quick reference guides.

MongoDB Documentation Build System

This document contains more direct instructions for building the MongoDB documentation.

Getting Started

Install Dependencies The MongoDB Documentation project depends on the following tools:

⁶⁴<http://sphinx-doc.org/markup/toctree.html#directive-toctree>

⁶⁵<http://sphinx-doc.org/markup/toctree.html#directive-toctree>

⁶⁶<http://sphinx-doc.org/markup/toctree.html#directive-toctree>

⁶⁷<http://sphinx-doc.org/markup/toctree.html#directive-toctree>

- GNU Make
- GNU Tar
- Python
- Git
- Sphinx (documentation management toolchain)
- Pygments (syntax highlighting)
- PyYAML (for the generated tables)
- Droopy (Python package for static text analysis)
- Fabric (Python package for scripting and orchestration)
- Inkscape (Image generation.)
- python-argparse (For Python 2.6.)
- LaTeX/PDF LaTeX (typically texlive; for building PDFs)
- Common Utilities (rsync, tar, gzip, sed)

OS X Install Sphinx, Docutils, and their dependencies with `easy_install` the following command:

```
easy_install Sphinx Jinja2 Pygments docutils PyYAML droopy fabric
```

Feel free to use `pip` rather than `easy_install` to install python packages.

To generate the images used in the documentation, [download and install Inkscape](#)⁶⁸.

Optional

To generate PDFs for the full production build, install a TeX distribution (for building the PDF.) If you do not have a LaTeX installation, use [MacTeX](#)⁶⁹. This is **only** required to build PDFs.

Arch Linux Install packages from the system repositories with the following command:

```
pacman -S python2-sphinx python2-yaml inkscape python2-pip
```

Then install the following Python packages:

```
pip install droopy fabric
```

Optional

To generate PDFs for the full production build, install the following packages from the system repository:

```
pacman -S texlive-bin texlive-core texlive-latexextra
```

Debian/Ubuntu Install the required system packages with the following command:

```
apt-get install python-sphinx python-yaml python-argparse inkscape python-pip
```

Then install the following Python packages:

⁶⁸<http://inkscape.org/download/>

⁶⁹<http://www.tug.org/mactex/2011/>

```
pip install droopy fabric
```

Optional

To generate PDFs for the full production build, install the following packages from the system repository:

```
apt-get install texlive-latex-recommended texlive-latex-recommended
```

Setup and Configuration Clone the repository:

```
git clone git://github.com/mongodb/docs.git
```

Then run the `bootstrap.py` script in the `docs/` repository, to configure the build dependencies:

```
python bootstrap.py
```

This downloads and configures the [mongodb/docs-tools](http://github.com/mongodb/docs-tools)⁷⁰ repository, which contains the authoritative build system shared between branches of the MongoDB Manual and other MongoDB documentation projects.

You can run `bootstrap.py` regularly to update build system.

Building the Documentation

The MongoDB documentation build system is entirely accessible via `make` targets. For example, to build an HTML version of the documentation issue the following command:

```
make html
```

You can find the build output in `build/<branch>/html`, where `<branch>` is the name of the current branch.

In addition to the `html` target, the build system provides the following targets:

publish Builds and integrates all output for the production build. Build output is in `build/public/<branch>/`. When you run `publish` in the master, the build will generate some output in `build/public/`.

push; stage Uploads the production build to the production or staging web servers. Depends on `publish`. Requires access production or staging environment.

push-all; stage-all Uploads the entire content of `build/public/` to the web servers. Depends on `publish`. Not used in common practice.

push-with-delete; stage-with-delete Modifies the action of `push` and `stage` to remove remote file that don't exist in the local build. Use with caution.

html; latex; dirhtml; epub; texinfo; man; json These are standard targets derived from the default Sphinx Makefile, with adjusted dependencies. Additionally, for all of these targets you can append `-nitpick` to increase Sphinx's verbosity, or `-clean` to remove all Sphinx build artifacts.

`latex` performs several additional post-processing steps on `.tex` output generated by Sphinx. This target will also compile PDFs using `pdflatex`.

`html` and `man` also generates a `.tar.gz` file of the build outputs for inclusion in the final releases.

⁷⁰<http://github.com/mongodb/docs-tools/>

Build Mechanics and Tools

Internally the build system has a number of components and processes. See the [docs-tools README](#)⁷¹ for more information on the internals. This section documents a few of these components from a very high level and lists useful operations for contributors to the documentation.

Fabric Fabric is an orchestration and scripting package for Python. The documentation uses Fabric to handle the deployment of the build products to the web servers and also unifies a number of independent build operations. Fabric commands have the following form:

```
fab <module>.<task>[:<argument>]
```

The `<argument>` is optional in most cases. Additionally some tasks are available at the root level, without a module. To see a full list of fabric tasks, use the following command:

```
fab -l
```

You can chain fabric tasks on a single command line, although this doesn't always make sense.

Important fabric tasks include:

tools.bootstrap Runs the `bootstrap.py` script. Useful for re-initializing the repository without needing to be in root of the repository.

tools.dev; tools.reset `tools.dev` switches the `origin` remote of the `docs-tools` checkout in `build` directory, to `../docs-tools` to facilitate build system testing and development. `tools.reset` resets the `origin` remote for normal operation.

tools.conf `tools.conf` returns the content of the configuration object for the current project. These data are useful during development.

stats.report:<filename> Returns, a collection of readability statistics. Specify file names relative to `source/` tree.

make Provides a thin wrapper around Make calls. Allows you to start make builds from different locations in the project repository.

process.refresh_dependencies Updates the time stamp of `.txt` source files with changed include files, to facilitate Sphinx's incremental rebuild process. This task runs internally as part of the build process.

Buildcloth [Buildcloth](#)⁷² is a meta-build tool, used to generate Makefiles programmatically. This makes the build system easier to maintain, and makes it easier to use the same fundamental code to generate various branches of the Manual as well as related documentation projects. See [makecloth/ in the docs-tools repository](#)⁷³ for the relevant code.

Running `make` with no arguments will regenerate these parts of the build system automatically.

Rstcloth [Rstcloth](#)⁷⁴ is a library for generating reStructuredText programmatically. This makes it possible to generate content for the documentation, such as tables, tables of contents, and API reference material programmatically and transparently. See [rstcloth/ in the docs-tools repository](#)⁷⁵ for the relevant code.

If you have any questions, please feel free to open a [Jira Case](#)⁷⁶.

⁷¹<https://github.com/mongodb/docs-tools/blob/master/README.rst>

⁷²<https://pypi.python.org/pypi/buildcloth/>

⁷³<https://github.com/mongodb/docs-tools/tree/master/makecloth>

⁷⁴<https://pypi.python.org/pypi/rstcloth>

⁷⁵<https://github.com/mongodb/docs-tools/tree/master/rstcloth>

⁷⁶<https://jira.mongodb.org/browse/DOCS>

Interfaces Reference

2.1 mongo Shell Methods

JavaScript in MongoDB

Although these methods use JavaScript, most interactions with MongoDB do not use JavaScript but use an idiomatic driver in the language of the interacting application.

2.1.1 Collection

Collection Methods

Name	Description
<code>db.collection.aggregate()</code> (page 22)	Provides access to the aggregation pipeline.
<code>db.collection.count()</code> (page 25)	Wraps <code>count</code> (page 201) to return a count of the number of documents in a collection.
<code>db.collection.copyTo()</code> (page 26)	Wraps <code>eval</code> (page 238) to copy data between collections in the same database.
<code>db.collection.createIndex()</code> (page 27)	Builds an index on a collection. Use <code>db.collection.ensureIndex()</code> to create an index if it does not currently exist.
<code>db.collection.getIndexStats()</code> (page 27)	Renders a human-readable view of the data collected by index.
<code>db.collection.indexStats()</code> (page 28)	Renders a human-readable view of the data collected by index.
<code>db.collection.dataSize()</code> (page 28)	Returns the size of the collection. Wraps the <code>size</code> (page 301) method.
<code>db.collection.distinct()</code> (page 29)	Returns an array of documents that have distinct values for the specified field.
<code>db.collection.drop()</code> (page 29)	Removes the specified collection from the database.
<code>db.collection.dropIndex()</code> (page 29)	Removes a specified index on a collection.
<code>db.collection.dropIndexes()</code> (page 30)	Removes all indexes on a collection.
<code>db.collection.ensureIndex()</code> (page 30)	Creates an index if it does not currently exist. If the index exists, it does nothing.
<code>db.collection.find()</code> (page 34)	Performs a query on a collection and returns a cursor object.
<code>db.collection.findAndModify()</code> (page 39)	Atomically modifies and returns a single document.
<code>db.collection.findOne()</code> (page 43)	Performs a query and returns a single document.
<code>db.collection.getIndexes()</code> (page 45)	Returns an array of documents that describe the existing indexes on a collection.
<code>db.collection.getShardDistribution()</code> (page 45)	For collections in sharded clusters, <code>db.collection.getShardDistribution()</code> returns an array of documents that describe the distribution of data across shards.
<code>db.collection.getShardVersion()</code> (page 47)	Internal diagnostic method for shard cluster.
<code>db.collection.group()</code> (page 47)	Provides simple data aggregation function. Groups documents by a key and performs an aggregation on the results.
<code>db.collection.initializeOrderedBulkOp()</code> (page 51)	Initializes a <code>Bulk()</code> (page 128) operations builder for an ordered bulk write.
<code>db.collection.initializeUnorderedBulkOp()</code> (page 51)	Initializes a <code>Bulk()</code> (page 128) operations builder for an unordered bulk write.
<code>db.collection.insert()</code> (page 52)	Creates a new document in a collection.
<code>db.collection.isCapped()</code> (page 55)	Reports if a collection is a <i>capped collection</i> .

Table 2.1 – continued from

Name	Description
<code>db.collection.mapReduce()</code> (page 55)	Performs map-reduce style data aggregation.
<code>db.collection.reIndex()</code> (page 62)	Rebuilds all existing indexes on a collection.
<code>db.collection.remove()</code> (page 62)	Deletes documents from a collection.
<code>db.collection.renameCollection()</code> (page 65)	Changes the name of a collection.
<code>db.collection.save()</code> (page 66)	Provides a wrapper around an <code>insert()</code> (page 52) and <code>update()</code> (page 69).
<code>db.collection.stats()</code> (page 68)	Reports on the state of a collection. Provides a wrapper around the <code>stats()</code> command.
<code>db.collection.storageSize()</code> (page 68)	Reports the total size used by the collection in bytes. Provides a wrapper around the <code>storageSize()</code> command.
<code>db.collection.totalSize()</code> (page 68)	Reports the total size of a collection, including the size of the indexes.
<code>db.collection.totalIndexSize()</code> (page 69)	Reports the total size used by the indexes on a collection. Provides a wrapper around the <code>totalIndexSize()</code> command.
<code>db.collection.update()</code> (page 69)	Modifies a document in a collection.
<code>db.collection.validate()</code> (page 76)	Performs diagnostic operations on a collection.

db.collection.aggregate()

New in version 2.2.

Definition

`db.collection.aggregate(pipeline, options)`

Calculates aggregate values for the data in a collection.

param array pipeline A sequence of data aggregation operations or stages. See the [aggregation pipeline operators](#) (page 438) for details.

Changed in version 2.6: The method can still accept the pipeline stages as separate arguments instead of as elements in an array; however, if you do not specify the `pipeline` as an array, you cannot specify the `options` parameter.

param document options Additional options that `aggregate()` (page 22) passes to the `aggregate` (page 198) command.

New in version 2.6: Available only if you specify the `pipeline` as an array.

The `options` document can contain the following fields and values:

field boolean explain Specifies to return the information on the processing of the pipeline. See [Return Information on Aggregation Pipeline Operation](#) (page 24) for an example.

New in version 2.6.

field boolean allowDiskUse Enables writing to temporary files. When set to `true`, aggregation operations can write data to the `_tmp` subdirectory in the `dbPath` directory. See [Perform Large Sort Operation with External Sort](#) (page 24) for an example.

New in version 2.6.

field document cursor Specifies the *initial* batch size for the cursor. The value of the `cursor` field is a document with the field `batchSize`. See [Specify an Initial Batch Size](#) (page 24) for syntax and example.

New in version 2.6.

Returns

A *cursor* to the documents produced by the final stage of the aggregation pipeline operation, or if you include the `explain` option, the document that provides details on the processing of the aggregation operation.

If the pipeline includes the `$out` (page 453) operator, `aggregate()` (page 22) returns an empty cursor. See `$out` (page 453) for more information.

Changed in version 2.6: The `db.collection.aggregate()` (page 22) method returns a cursor and can return result sets of any size. Previous versions returned all results in a single document, and the result set was subject to a size limit of 16 megabytes.

Changed in version 2.4: If an error occurs, the `aggregate()` (page 22) helper throws an exception. In previous versions, the helper returned a document with the error message and code, and `ok` status field not equal to 1, same as the `aggregate` (page 198) command.

See also:

For more information, see <http://docs.mongodb.org/manualcore/aggregation-pipeline>, [Aggregation Reference](#) (page 484), <http://docs.mongodb.org/manualcore/aggregation-pipeline-limits>, and `aggregate` (page 198).

Cursor Behavior In the `mongo` (page 527) shell, if the cursor returned from the `db.collection.aggregate()` (page 22) is not assigned to a variable using the `var` keyword, then the `mongo` (page 527) shell automatically iterates the cursor up to 20 times. See <http://docs.mongodb.org/manualcore/cursors> for cursor behavior in the `mongo` (page 527) shell and <http://docs.mongodb.org/manualtutorial/iterate-a-cursor> for handling cursors in the `mongo` (page 527) shell.

Cursors returned from aggregation only supports cursor methods that operate on evaluated cursors (i.e. cursors whose first batch has been retrieved), such as the following methods:

- `cursor.hasNext()` (page 85)
- `cursor.next()` (page 91)
- `cursor.toArray()` (page 96)
- `cursor.forEach()` (page 85)
- `cursor.map()` (page 87)
- `cursor.objsLeftInBatch()` (page 91)
- `cursor.itcount()`
- `cursor.pretty()`

Examples The examples in this section use the `db.collection.aggregate()` (page 22) helper provided in the 2.6 version of the `mongo` (page 527) shell.

The following examples use the collection `orders` that contains the following documents:

```
{ _id: 1, cust_id: "abc1", ord_date: ISODate("2012-11-02T17:04:11.102Z"), status: "A", amount: 50 }
{ _id: 2, cust_id: "xyz1", ord_date: ISODate("2013-10-01T17:04:11.102Z"), status: "A", amount: 100 }
{ _id: 3, cust_id: "xyz1", ord_date: ISODate("2013-10-12T17:04:11.102Z"), status: "D", amount: 25 }
{ _id: 4, cust_id: "xyz1", ord_date: ISODate("2013-10-11T17:04:11.102Z"), status: "D", amount: 125 }
{ _id: 5, cust_id: "abc1", ord_date: ISODate("2013-11-12T17:04:11.102Z"), status: "A", amount: 25 }
```

Group by and Calculate a Sum The following aggregation operation selects documents with status equal to "A", groups the matching documents by the `cust_id` field and calculates the total for each `cust_id` field from the sum of the amount field, and sorts the results by the `total` field in descending order:

```
db.orders.aggregate([
  { $match: { status: "A" } },
  { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },
  { $sort: { total: -1 } }
])
```

The operation returns a cursor with the following documents:

```
{ "_id" : "xyz1", "total" : 100 }
{ "_id" : "abc1", "total" : 75 }
```

The `mongo` (page 527) shell iterates the returned cursor automatically to print the results. See <http://docs.mongodb.org/manual/tutorial/iterate-a-cursor> for handling cursors manually in the `mongo` (page 527) shell.

Return Information on Aggregation Pipeline Operation The following aggregation operation sets the option `explain` to `true` to return information about the aggregation operation.

```
db.orders.aggregate(
  [
    { $match: { status: "A" } },
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },
    { $sort: { total: -1 } }
  ],
  {
    explain: true
  }
)
```

The operation returns a cursor with the document that contains detailed information regarding the processing of the aggregation pipeline. For example, the document may show, among other details, which index, if any, the operation used.¹ If the `orders` collection is a sharded collection, the document would also show the division of labor between the shards and the merge operation, and for targeted queries, the targeted shards.

Note: The intended readers of the `explain` output document are humans, and not machines, and the output format is subject to change between releases.

The `mongo` (page 527) shell iterates the returned cursor automatically to print the results. See <http://docs.mongodb.org/manual/tutorial/iterate-a-cursor> for handling cursors manually in the `mongo` (page 527) shell.

Perform Large Sort Operation with External Sort Aggregation pipeline stages have *maximum memory use limit*. To handle large datasets, set `allowDiskUse` option to `true` to enable writing data to temporary files, as in the following example:

```
var results = db.stocks.aggregate(
  [
    { $project : { cusip: 1, date: 1, price: 1, _id: 0 } },
    { $sort : { cusip : 1, date: 1 } }
  ],
  {
    allowDiskUse: true
  }
)
```

Specify an Initial Batch Size To specify an initial batch size for the cursor, use the following syntax for the `cursor` option:

¹ *index-filters* can affect the choice of index used. See *index-filters* for details.

```
cursor: { batchSize: <int> }
```

For example, the following aggregation operation specifies the *initial* batch size of 0 for the cursor:

```
db.orders.aggregate(
  [
    { $match: { status: "A" } },
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },
    { $sort: { total: -1 } },
    { $limit: 2 }
  ],
  {
    cursor: { batchSize: 0 }
  }
)
```

A `batchSize` of 0 means an empty first batch and is useful for quickly returning a cursor or failure message without doing significant server-side work. Specify subsequent batch sizes to *OP_GET_MORE*² operations as with other MongoDB cursors.

The `mongo` (page 527) shell iterates the returned cursor automatically to print the results. See <http://docs.mongodb.org/manual/tutorial/iterate-a-cursor> for handling cursors manually in the `mongo` (page 527) shell.

`db.collection.count()`

Definition

`db.collection.count(<query>)`

Returns the count of documents that would match a `find()` (page 34) query. The `db.collection.count()` (page 25) method does not perform the `find()` (page 34) operation but instead counts and returns the number of results that match a query.

The `db.collection.count()` (page 25) method has the following parameter:

param document query The query selection criteria.

See also:

`cursor.count()` (page 79)

Behavior On a sharded cluster, `db.collection.count()` (page 25) can result in an *inaccurate* count if *orphaned documents* exist or if a chunk migration is in progress.

To avoid these situations, on a sharded cluster, use the `$group` (page 447) stage of the `db.collection.aggregate()` (page 22) method to `$sum` (page 460) the documents. For example, the following operation counts the documents in a collection:

```
db.collection.aggregate(
  [
    { $group: { _id: null, count: { $sum: 1 } } }
  ]
)
```

To get a count of documents that match a query condition, include the `$match` (page 440) stage as well:

²<http://docs.mongodb.org/meta-driver/latest/legacy/mongodb-wire-protocol/#wire-op-get-more>

```
db.collection.aggregate([
  { $match: <query condition> },
  { $group: { _id: null, count: { $sum: 1 } } }
])
```

See [Perform a Count](#) (page 440) for an example.

Examples

Count all Documents in a Collection To count the number of all documents in the `orders` collection, use the following operation:

```
db.orders.count()
```

This operation is equivalent to the following:

```
db.orders.find().count()
```

Count all Documents that Match a Query Count the number of the documents in the `orders` collection with the field `ord_dt` greater than `new Date('01/01/2012')`:

```
db.orders.count( { ord_dt: { $gt: new Date('01/01/2012') } } )
```

The query is equivalent to the following:

```
db.orders.find( { ord_dt: { $gt: new Date('01/01/2012') } } ).count()
```

`db.collection.copyTo()`

Definition

`db.collection.copyTo(newCollection)`

Copies all documents from `collection` into `newCollection` using server-side JavaScript. If `newCollection` does not exist, MongoDB creates it.

If authentication is enabled, you must have access to all actions on all resources in order to run `db.collection.copyTo()` (page 26). Providing such access is not recommended, but if your organization requires a user to run `db.collection.copyTo()` (page 26), create a role that grants `anyAction` on `resource-anyresource`. Do not assign this role to any other user.

param string newCollection The name of the collection to write data to.

Warning: When using `db.collection.copyTo()` (page 26) check field types to ensure that the operation does not remove type information from documents during the translation from *BSON* to *JSON*. Consider using `cloneCollection()` (page 99) to maintain type fidelity.

`copyTo()` (page 26) returns the number of documents copied. If the copy fails, it throws an exception.

Behavior Because `copyTo()` (page 26) uses `eval` (page 238) internally, the copy operations will block all other operations on the `mongod` (page 503) instance.

Example The following operation copies all documents from the `source` collection into the `target` collection.

```
db.source.copyTo(target)
```

db.collection.createIndex()

Definition

`db.collection.createIndex(keys, options)`

Deprecated since version 1.8.

Creates indexes on collections.

param document keys For each field to index, a key-value pair with the field and the index order: 1 for ascending or -1 for descending.

param document options One or more key-value pairs that specify index options. For a list of options, see `db.collection.ensureIndex()` (page 30).

See also:

<http://docs.mongodb.org/manualindexes>, `db.collection.createIndex()` (page 27), `db.collection.dropIndex()` (page 29), `db.collection.dropIndexes()` (page 30), `db.collection.getIndexes()` (page 45), `db.collection.reIndex()` (page 62), and `db.collection.totalIndexSize()` (page 69)

db.collection.getIndexStats()

Definition

`db.collection.getIndexStats(index)`

Displays a human-readable summary of aggregated statistics about an index's B-tree data structure. The information summarizes the output returned by the `indexStats` (page 339) command and `indexStats()` (page 28) method. The `getIndexStats()` (page 27) method displays the information on the screen and does not return an object.

The `getIndexStats()` (page 27) method has the following form:

```
db.<collection>.getIndexStats( { index : "<index name>" } )
```

param document index The *index name*.

The `getIndexStats()` (page 27) method is available only when connected to a `mongod` (page 503) instance that uses the `--enableExperimentalIndexStatsCmd` option.

To view *index names* for a collection, use the `getIndexes()` (page 45) method.

Warning: Do not use `getIndexStats()` (page 27) or `indexStats` (page 339) with production deployments.

Example The following command returns information for an index named `type_1_traits_1`:

```
db.animals.getIndexStats({index:"type_1_traits_1"})
```

The command returns the following summary. For more information on the B-tree statistics, see `indexStats` (page 339).

```
-- index "undefined" --
  version 1 | key pattern {  "type" : 1,  "traits" : 1 } | storage namespace "test.animals.$type_1_t
  2 deep, bucket body is 8154 bytes

bucket count      45513    on average 99.401 % (±0.463 %) full      49.581 % (±4.135 %) bson keys,

-- depth 0 --
  bucket count      1        on average 71.511 % (±0.000 %) full      36.191 % (±0.000 %) bson keys,

-- depth 1 --
  bucket count      180      on average 98.954 % (±5.874 %) full      49.732 % (±5.072 %) bson keys,

-- depth 2 --
  bucket count      45332    on average 99.403 % (±0.245 %) full      49.580 % (±4.130 %) bson keys,
```

db.collection.indexStats()

Definition

`db.collection.indexStats(index)`

Aggregates statistics for the B-tree data structure that stores data for a MongoDB index. The `indexStats()` (page 28) method is a thin wrapper around the `indexStats` (page 339) command. The `indexStats()` (page 28) method is available only on `mongod` (page 503) instances running with the `--enableExperimentalIndexStatsCmd` option.

Important: The `indexStats()` (page 28) method is not intended for production deployments.

The `indexStats()` (page 28) method has the following form:

```
db.<collection>.indexStats( { index: "<index name>" } )
```

The `indexStats()` (page 28) method has the following parameter:

param document index *The index name.*

The method takes a read lock and pages into memory all the extents, or B-tree buckets, encountered. The method might be slow for large indexes if the underlying extents are not already in physical memory. Do not run `indexStats()` (page 28) on a *replica set primary*. When run on a *secondary*, the command causes the secondary to fall behind on replication.

The method aggregates statistics for the entire B-tree and for each individual level of the B-tree. For a description of the command's output, see `indexStats` (page 339).

For more information about running `indexStats()` (page 28), see <https://github.com/mongodb-labs/storage-viz#readme>.

db.collection.dataSize()

`db.collection.dataSize()`

Returns The size of the collection. This method provides a wrapper around the `size` (page 326) output of the `collStats` (page 325) (i.e. `db.collection.stats()` (page 68)) command.

db.collection.distinct()

Definition

`db.collection.distinct (field, query)`

Finds the distinct values for a specified field across a single collection and returns the results in an array.

param string field The field for which to return distinct values.

param document query A query that specifies the documents from which to retrieve the distinct values.

The `db.collection.distinct()` (page 29) method provides a wrapper around the `distinct` (page 203) command. Results must not be larger than the maximum *BSON size* (page 604).

When possible to use covered indexes, the `db.collection.distinct()` (page 29) method will use an index to find the documents in the query as well as to return the data.

Examples The following are examples of the `db.collection.distinct()` (page 29) method:

- Return an array of the distinct values of the field `ord_dt` from all documents in the `orders` collection:

```
db.orders.distinct( 'ord_dt' )
```

- Return an array of the distinct values of the field `sku` in the subdocument `item` from all documents in the `orders` collection:

```
db.orders.distinct( 'item.sku' )
```

- Return an array of the distinct values of the field `ord_dt` from the documents in the `orders` collection where the `price` is greater than 10:

```
db.orders.distinct( 'ord_dt', { price: { $gt: 10 } } )
```

db.collection.drop()

`db.collection.drop()`

Call the `db.collection.drop()` (page 29) method on a collection to drop it from the database. The method provides a wrapper around the `drop` (page 304) command.

`db.collection.drop()` (page 29) takes no arguments and will produce an error if called with any arguments.

This method also removes any indexes associated with the dropped collection.

Warning: This method obtains a write lock on the affected database and will block other operations until it has completed.

db.collection.dropIndex()**Definition**

`db.collection.dropIndex (index)`

Drops or removes the specified index from a collection. The `db.collection.dropIndex()` (page 29) method provides a wrapper around the `dropIndexes` (page 311) command.

Note: You cannot drop the default index on the `_id` field.

The `db.collection.dropIndex()` (page 29) method takes the following parameter:

param string,document index Specifies the index to drop. You can specify the index either by the index name or by the index specification document.³

To drop a `text` index, specify the index name.

To get the index name or the index specification document for the `db.collection.dropIndex()` (page 29) method, use the `db.collection.getIndexes()` (page 45) method.

Example Consider a `pets` collection. Calling the `getIndexes()` (page 45) method on the `pets` collection returns the following indexes:

```
[
  {
    "v" : 1,
    "key" : { "_id" : 1 },
    "ns" : "test.pets",
    "name" : "_id_"
  },
  {
    "v" : 1,
    "key" : { "cat" : -1 },
    "ns" : "test.pets",
    "name" : "catIdx"
  },
  {
    "v" : 1,
    "key" : { "cat" : 1, "dog" : -1 },
    "ns" : "test.pets",
    "name" : "cat_1_dog_-1"
  }
]
```

The single field index on the field `cat` has the user-specified name of `catIdx`⁴ and the index specification document of `{ "cat" : -1 }`.

To drop the index `catIdx`, you can use either the index name:

```
db.pets.dropIndex( "catIdx" )
```

Or you can use the index specification document `{ "cat" : -1 }`:

```
db.pets.dropIndex( { "cat" : -1 } )
```

db.collection.dropIndexes()

`db.collection.dropIndexes()`

Drops all indexes other than the required index on the `_id` field. Only call `dropIndexes()` (page 30) as a method on a collection object.

db.collection.ensureIndex()

Definition

³ When using a `mongo` (page 527) shell version earlier than 2.2.2, if you specified a name during the index creation, you must use the name to drop the index.

⁴ During index creation, if the user does **not** specify an index name, the system generates the name by concatenating the index key field and value with an underscore, e.g. `cat_1`.

`db.collection.ensureIndex(keys, options)`

Creates an index on the specified field if the index does not already exist.

The `ensureIndex()` (page 30) method has the following fields:

param document keys A document that contains the field and value pairs where the field is the index key and the value describes the type of index for that field. For an ascending index on a field, specify a value of 1; for descending index, specify a value of -1.

MongoDB supports several different index types including *text*, *geospatial*, and *hashed* indexes. See *index-type-list* for more information.

param document options A document that contains a set of options that controls the creation of the index. See *Options* (page 31) for details.

Options The `options` document contains a set of options that controls the creation of the index. Different index types can have additional options specific for that type.

Options for All Index Types The following options are available for all index types unless otherwise specified:

param Boolean background Builds the index in the background so that building an index does *not* block other database activities. Specify `true` to build in the background. The default value is `false`.

param Boolean unique Creates a unique index so that the collection will not accept insertion of documents where the index key or keys match an existing value in the index. Specify `true` to create a unique index. The default value is `false`.

The option is *unavailable* for hashed indexes.

param string name The name of the index. If unspecified, MongoDB generates an index name by concatenating the names of the indexed fields and the sort order.

Whether user specified or MongoDB generated, index names including their full namespace (i.e. `database.collection`) cannot be longer than the *Index Name Limit* (page 605).

param Boolean dropDups Creates a unique index on a field that *may* have duplicates. MongoDB indexes only the first occurrence of a key and **removes** all documents from the collection that contain subsequent occurrences of that key. Specify `true` to create unique index. The default value is `false`.

The option is *unavailable* for hashed indexes.

Deprecated since version 2.6.

Warning: `dropDups` will delete data from your collection when building the index.

param Boolean sparse If `true`, the index only references documents with the specified field. These indexes use less space but behave differently in some situations (particularly sorts). The default value is `false`. See <http://docs.mongodb.org/manualcore/index-sparse> for more information.

Changed in version 2.6: `2dsphere` indexes are sparse by default and ignore this option. For a compound index that includes `2dsphere` index key(s) along with keys of other types, only the `2dsphere` index fields determine whether the index references a document.

`2d`, `geoHaystack`, and `text` indexes behave similarly to the `2dsphere` indexes.

param integer expireAfterSeconds Specifies a value, in seconds, as a *TTL* to control how long MongoDB retains documents in this collection. See

<http://docs.mongodb.org/manual/tutorial/expire-data> for more information on this functionality. This applies only to *TTL* indexes.

param index version v The index version number. The default index version depends on the version of *mongod* (page 503) running when creating the index. Before version 2.0, the this value was 0; versions 2.0 and later use version 1, which provides a smaller and faster index format. Specify a different index version *only* in unusual situations.

Options for text Indexes The following options are available for *text* indexes only:

param document weights For *text* indexes, a document that contains field and weight pairs. The weight is an integer ranging from 1 to 99,999 and denotes the significance of the field relative to the other indexed fields in terms of the score. You can specify weights for some or all the indexed fields. See <http://docs.mongodb.org/manual/tutorial/control-results-of-text-search> to adjust the scores. The default value is 1.

param string default_language For *text* indexes, the language that determines the list of stop words and the rules for the stemmer and tokenizer. See *text-search-languages* for the available languages and <http://docs.mongodb.org/manual/tutorial/specify-language-for-text-index> for more information and examples. The default value is *english*.

param string language_override For *text* indexes, the name of the field, in the collection's documents, that contains the override language for the document. The default value is *language*. See *specify-language-field-text-index-example* for an example.

param integer textIndexVersion For *text* indexes, the *text* index version number. Version can be either 1 or 2.

In MongoDB 2.6, the default version is 2. MongoDB 2.4 can only support version 1.

New in version 2.6.

Options for 2dsphere Indexes The following option is available for *2dsphere* indexes only:

param integer 2dsphereIndexVersion For *2dsphere* indexes, the *2dsphere* index version number. Version can be either 1 or 2.

In MongoDB 2.6, the default version is 2. MongoDB 2.4 can only support version 1.

New in version 2.6.

Options for 2d Indexes The following options are available for *2d* indexes only:

param integer bits For *2d* indexes, the number of precision of the stored *geohash* value of the location data.

The *bits* value ranges from 1 to 32 inclusive. The default value is 26.

param number min For *2d* indexes, the lower inclusive boundary for the longitude and latitude values. The default value is -180.0 .

param number max For *2d* indexes, the upper inclusive boundary for the longitude and latitude values. The default value is 180.0 .

Options for geoHaystack Indexes The following option is available for geoHaystack indexes only:

param number bucketSize For geoHaystack indexes, specify the number of units within which to group the location values; i.e. group in the same bucket those location values that are within the specified number of units to each other.

The value must be greater than 0.

Behaviors The `ensureIndex()` (page 30) method has the behaviors described here.

- To add or change index options you must drop the index using the `dropIndex()` (page 29) method and issue another `ensureIndex()` (page 30) operation with the new options.

If you create an index with one set of options, and then issue the `ensureIndex()` (page 30) method with the same index fields and different options without first dropping the index, `ensureIndex()` (page 30) will *not* rebuild the existing index with the new options.

- If you call multiple `ensureIndex()` (page 30) methods with the same index specification at the same time, only the first operation will succeed, all other operations will have no effect.
- Non-background indexing operations will block all other operations on a database.
- MongoDB will **not** `create an index` (page 30) on a collection if the index entry for an existing document exceeds the `Maximum Index Key Length`. Previous versions of MongoDB would create the index but not index such documents.

Changed in version 2.6.

Examples

Create an Ascending Index on a Single Field The following example creates an ascending index on the field `orderDate`.

```
db.collection.ensureIndex( { orderDate: 1 } )
```

If the `keys` document specifies more than one field, then `ensureIndex()` (page 30) creates a *compound index*.

Create an Index on a Multiple Fields The following example creates a compound index on the `orderDate` field (in ascending order) and the `zipcode` field (in descending order.)

```
db.collection.ensureIndex( { orderDate: 1, zipcode: -1 } )
```

A compound index cannot include a *hashed index* component.

Note: The order of an index is important for supporting `sort()` (page 93) operations using the index.

See also:

- The <http://docs.mongodb.org/manualindexes> section of this manual for full documentation of indexes and indexing in MongoDB.
- <http://docs.mongodb.org/manualcore/index-text> for details on creating text indexes.
- *index-feature-geospatial* and *index-geohaystack-index* for geospatial queries.
- *index-feature-ttl* for expiration of data.
- `db.collection.getIndexes()` (page 45) to view the specifications of existing indexes for a collection.

db.collection.find()**Definition**

`db.collection.find(<criteria>, <projection>)`

Selects documents in a collection and returns a *cursor* to the selected documents.⁵

param document criteria Specifies selection criteria using *query operators* (page 373). To return all documents in a collection, omit this parameter or pass an empty document (`{ }`).

param document projection Specifies the fields to return using *projection operators* (page 406). To return all fields in the matching document, omit this parameter.

Returns

A *cursor* to the documents that match the query criteria. When the `find()` (page 34) method “returns documents,” the method is actually returning a cursor to the documents.

If the `projection` argument is specified, the matching documents contain only the `projection` fields and the `_id` field. You can optionally exclude the `_id` field.

Executing `find()` (page 34) directly in the `mongo` (page 527) shell automatically iterates the cursor to display up to the first 20 documents. Type `it` to continue iteration.

To access the returned documents with a driver, use the appropriate cursor handling mechanism for the driver language.

The `projection` parameter takes a document of the following form:

```
{ field1: <boolean>, field2: <boolean> ... }
```

The `<boolean>` value can be any of the following:

- 1 or `true` to include the field. The `find()` (page 34) method always includes the `_id` field even if the field is not explicitly stated to return in the *projection* parameter.
- 0 or `false` to exclude the field.

A *projection cannot contain both include and exclude specifications*, except for the exclusion of the `_id` field. In projections that *explicitly include* fields, the `_id` field is the only field that you can *explicitly exclude*.

Examples

Find All Documents in a Collection The `find()` (page 34) method with no parameters returns all documents from a collection and returns all fields for the documents. For example, the following operation returns all documents in the `bios` collection:

```
db.bios.find()
```

Find Documents that Match Query Criteria To find documents that match a set of selection criteria, call `find()` with the `<criteria>` parameter. The following operation returns all the documents from the collection `products` where `qty` is greater than 25:

```
db.products.find( { qty: { $gt: 25 } } )
```

⁵ `db.collection.find()` (page 34) is a wrapper for the more formal query structure that uses the `$query` (page 483) operator.

Query for Equality The following operation returns documents in the `bios` collection where `_id` equals 5:

```
db.bios.find( { _id: 5 } )
```

Query Using Operators The following operation returns documents in the `bios` collection where `_id` equals either 5 or `ObjectId("507c35dd8fada716c89d0013")`:

```
db.bios.find(
  {
    _id: { $in: [ 5, ObjectId("507c35dd8fada716c89d0013") ] }
  }
)
```

Query for Ranges Combine comparison operators to specify ranges. The following operation returns documents with `field` between `value1` and `value2`:

```
db.collection.find( { field: { $gt: value1, $lt: value2 } } );
```

Query a Field that Contains an Array If a field contains an array and your query has multiple conditional operators, the field as a whole will match if either a single array element meets the conditions or a combination of array elements meet the conditions.

Given a collection `students` that contains the following documents:

```
{ "_id" : 1, "score" : [ -1, 3 ] }
{ "_id" : 2, "score" : [ 1, 5 ] }
{ "_id" : 3, "score" : [ 5, 5 ] }
```

The following query:

```
db.students.find( { score: { $gt: 0, $lt: 2 } } )
```

Matches the following documents:

```
{ "_id" : 1, "score" : [ -1, 3 ] }
{ "_id" : 2, "score" : [ 1, 5 ] }
```

In the document with `_id` equal to 1, the `score: [-1, 3]` meets the conditions because the element `-1` meets the `$lt: 2` condition and the element `3` meets the `$gt: 0` condition.

In the document with `_id` equal to 2, the `score: [1, 5]` meets the conditions because the element `1` meets both the `$lt: 2` condition and the `$gt: 0` condition.

Query Arrays

Query for an Array Element The following operation returns documents in the `bios` collection where the array field `contribs` contains the element `"UNIX"`:

```
db.bios.find( { contribs: "UNIX" } )
```

Query an Array of Documents The following operation returns documents in the `bios` collection where `awards` array contains a subdocument element that contains the `award` field equal to `"Turing Award"` and the `year` field greater than 1980:

```
db.bios.find(
  {
    awards: {
      $elemMatch: {
        award: "Turing Award",
        year: { $gt: 1980 }
      }
    }
  }
)
```

Query Subdocuments

Query Exact Matches on Subdocuments The following operation returns documents in the `bios` collection where the subdocument name is *exactly* `{ first: "Yukihiro", last: "Matsumoto" }`, including the order:

```
db.bios.find(
  {
    name: {
      first: "Yukihiro",
      last: "Matsumoto"
    }
  }
)
```

The `name` field must match the sub-document exactly. The query does **not** match documents with the following `name` fields:

```
{
  first: "Yukihiro",
  aka: "Matz",
  last: "Matsumoto"
}

{
  last: "Matsumoto",
  first: "Yukihiro"
}
```

Query Fields of a Subdocument The following operation returns documents in the `bios` collection where the subdocument name contains a field `first` with the value `"Yukihiro"` and a field `last` with the value `"Matsumoto"`. The query uses *dot notation* to access fields in a subdocument:

```
db.bios.find(
  {
    "name.first": "Yukihiro",
    "name.last": "Matsumoto"
  }
)
```

The query matches the document where the `name` field contains a subdocument with the field `first` with the value `"Yukihiro"` and a field `last` with the value `"Matsumoto"`. For instance, the query would match documents with `name` fields that held either of the following values:


```
{
  first: "Yukihiro",
  aka: "Matz",
  last: "Matsumoto"
}

{
  last: "Matsumoto",
  first: "Yukihiro"
}
```

Projections The `projection` parameter specifies which fields to return. The parameter contains either include or exclude specifications, not both, unless the exclude is for the `_id` field.

Specify the Fields to Return The following operation returns all the documents from the `products` collection where `qty` is greater than 25 and returns only the `_id`, `item` and `qty` fields:

```
db.products.find( { qty: { $gt: 25 } }, { item: 1, qty: 1 } )
```

The operation returns the following:

```
{ "_id" : 11, "item" : "pencil", "qty" : 50 }
{ "_id" : ObjectId("50634d86be4617f17bb159cd"), "item" : "bottle", "qty" : 30 }
{ "_id" : ObjectId("50634dbcbe4617f17bb159d0"), "item" : "paper", "qty" : 100 }
```

The following operation finds all documents in the `bios` collection and returns only the `name` field, `contribs` field and `_id` field:

```
db.bios.find( { }, { name: 1, contribs: 1 } )
```

Explicitly Excluded Fields The following operation queries the `bios` collection and returns all fields *except* the `first` field in the `name` subdocument and the `birth` field:

```
db.bios.find(
  { contribs: 'OOP' },
  { 'name.first': 0, birth: 0 }
)
```

Explicitly Exclude the `_id` Field The following operation excludes the `_id` and `qty` fields from the result set:

```
db.products.find( { qty: { $gt: 25 } }, { _id: 0, qty: 0 } )
```

The documents in the result set contain all fields *except* the `_id` and `qty` fields:

```
{ "item" : "pencil", "type" : "no.2" }
{ "item" : "bottle", "type" : "blue" }
{ "item" : "paper" }
```

The following operation finds documents in the `bios` collection and returns only the `name` field and the `contribs` field:

```
db.bios.find(
  { },
  { name: 1, contribs: 1, _id: 0 }
)
```

On Arrays and Subdocuments The following operation queries the `bios` collection and returns the `last` field in the `name` subdocument and the first two elements in the `contribs` array:

```
db.bios.find(
  { },
  {
    _id: 0,
    'name.last': 1,
    contribs: { $slice: 2 }
  }
)
```

Iterate the Returned Cursor The `find()` (page 34) method returns a *cursor* to the results. In the `mongo` (page 527) shell, if the returned cursor is not assigned to a variable using the `var` keyword, the cursor is automatically iterated up to 20 times to access up to the first 20 documents that match the query. You can use the `DBQuery.shellBatchSize` to change the number of iterations. See *Flags* (page 78) and *cursor-behaviors*. To iterate manually, assign the returned cursor to a variable using the `var` keyword.

With Variable Name The following example uses the variable `myCursor` to iterate over the cursor and print the matching documents:

```
var myCursor = db.bios.find( );

myCursor
```

With `next()` Method The following example uses the cursor method `next()` (page 91) to access the documents:

```
var myCursor = db.bios.find( );

var myDocument = myCursor.hasNext() ? myCursor.next() : null;

if (myDocument) {
  var myName = myDocument.name;
  print (toJson(myName));
}
```

To print, you can also use the `printjson()` method instead of `print (toJson())`:

```
if (myDocument) {
  var myName = myDocument.name;
  printjson(myName);
}
```

With `forEach()` Method The following example uses the cursor method `forEach()` (page 85) to iterate the cursor and access the documents:

```
var myCursor = db.bios.find( );

myCursor.forEach(printjson);
```

Modify the Cursor Behavior The `mongo` (page 527) shell and the drivers provide several cursor methods that call on the *cursor* returned by the `find()` (page 34) method to modify its behavior.

Order Documents in the Result Set The `sort()` (page 93) method orders the documents in the result set. The following operation returns documents in the `bios` collection sorted in ascending order by the `name` field:

```
db.bios.find().sort( { name: 1 } )
```

`sort()` (page 93) corresponds to the `ORDER BY` statement in SQL.

Limit the Number of Documents to Return The `limit()` (page 86) method limits the number of documents in the result set. The following operation returns at most 5 documents in the `bios` collection:

```
db.bios.find().limit( 5 )
```

`limit()` (page 86) corresponds to the `LIMIT` statement in SQL.

Set the Starting Point of the Result Set The `skip()` (page 92) method controls the starting point of the results set. The following operation skips the first 5 documents in the `bios` collection and returns all remaining documents:

```
db.bios.find().skip( 5 )
```

Combine Cursor Methods The following example chains cursor methods:

```
db.bios.find().sort( { name: 1 } ).limit( 5 )
db.bios.find().limit( 5 ).sort( { name: 1 } )
```

Regardless of the order you chain the `limit()` (page 86) and the `sort()` (page 93), the request to the server has the structure that treats the query and the `sort()` (page 93) modifier as a single object. Therefore, the `limit()` (page 86) operation method is always applied after the `sort()` (page 93) regardless of the specified order of the operations in the chain. See the *meta query operators* (page 477).

`db.collection.findAndModify()`

Definition

`db.collection.findAndModify(<document>)`

Modifies and returns a single document. By default, the returned document does not include the modifications made on the update. To return the document with the modifications made on the update, use the `new` option. The `findAndModify()` (page 39) method is a shell helper around the `findAndModify` (page 229) command.

The `findAndModify()` (page 39) method has the following form:

```
db.collection.findAndModify({
  query: <document>,
  sort: <document>,
  remove: <boolean>,
  update: <document>,
  new: <boolean>,
  fields: <document>,
  upsert: <boolean>
});
```

The `db.collection.findAndModify()` (page 39) method takes a document parameter with the following subdocument fields:

param document query The selection criteria for the modification. The `query` field employs the same *query selectors* (page 373) as used in the `db.collection.find()` (page 34) method. Although the query may match multiple documents, `findAndModify()` (page 39) will select only one document to modify.

param document sort Determines which document the operation modifies if the query selects multiple documents. `findAndModify()` (page 39) modifies the first document in the sort order specified by this argument.

param Boolean remove Must specify either the `remove` or the `update` field. Removes the document specified in the `update` field. Set this to `true` to remove the selected document. The default is `false`.

param document update Must specify either the `remove` or the `update` field. Performs an update of the selected document. The `update` field employs the same *update operators* (page 412) or `field: value` specifications to modify the selected document.

param Boolean new When `true`, returns the modified document rather than the original. The `findAndModify()` (page 39) method ignores the `new` option for `remove` operations. The default is `false`.

param document fields A subset of fields to return. The `fields` document specifies an inclusion of a field with `1`, as in: `fields: { <field1>: 1, <field2>: 1, ... }`. See *projection*.

param Boolean upsert Used in conjunction with the `update` field.

When `true`, `findAndModify()` (page 39) creates a new document if no document matches the query, or if documents match the query, `findAndModify()` (page 39) performs an update.

The default is `false`.

Return Data The `findAndModify()` (page 39) method returns either: the pre-modification document or, if `new: true` is set, the modified document.

Note:

- If the query finds no document for update or remove operations, `findAndModify()` (page 39) returns `null`.
 - If the query finds no document for an update with an upsert operation, `findAndModify()` (page 39) creates a new document. If `new` is `false`, **and** the `sort` option is **NOT** specified, the method returns `null`.
 - If the query finds no document for an update with an upsert operation, `findAndModify()` (page 39) creates a new document. If `new` is `false`, **and** a `sort` option, the method returns an empty document `{}`.
-

Behaviors

Upsert and Unique Index When `findAndModify()` (page 39) includes the `upsert: true` option **and** the query field(s) is not uniquely indexed, the method could insert a document multiple times in certain circumstances. For instance, if multiple clients each invoke the method with the same query condition and these methods complete the `find` phase before any of methods perform the `modify` phase, these methods could result in the insertion of the same document.

In the following example, no document with the name `Andy` exists, and multiple clients issue the following command:

```
db.people.findAndModify({
  query: { name: "Andy" },
  sort: { rating: 1 },
  update: { $inc: { score: 1 } },
  upsert: true
})
```

Then, if these clients' `findAndModify()` (page 39) methods finish the `query` phase before any command starts the `modify` phase, **and** there is no unique index on the `name` field, the commands may all perform an upsert. To prevent this condition, create a *unique index* on the `name` field. With the unique index in place, the multiple methods would observe one of the following behaviors:

- Exactly one `findAndModify()` (page 39) would successfully insert a new document.
- Zero or more `findAndModify()` (page 39) methods would update the newly inserted document.
- Zero or more `findAndModify()` (page 39) methods would fail when they attempted to insert a duplicate. If the method fails due to a unique index constraint violation, you can retry the method. Absent a delete of the document, the retry should not fail.

Sharded Collections When using `findAndModify` (page 229) in a *sharded* environment, the `query` **must** contain the *shard key* for all operations against the shard cluster for the *sharded* collections.

`findAndModify` (page 229) operations issued against `mongos` (page 518) instances for *non-sharded* collections function normally.

Comparisons with the `update` Method When updating a document, `findAndModify()` (page 39) and the `update()` (page 69) method operate differently:

- By default, both operations modify a single document. However, the `update()` (page 69) method with its `multi` option can modify more than one document.
- If multiple documents match the update criteria, for `findAndModify()` (page 39), you can specify a `sort` to provide some measure of control on which document to update.

With the default behavior of the `update()` (page 69) method, you cannot specify which single document to update when multiple documents match.

- By default, `findAndModify()` (page 39) method returns the pre-modified version of the document. To obtain the updated document, use the `new` option.

The `update()` (page 69) method returns a `WriteResult` (page 188) object that contains the status of the operation. To return the updated document, use the `find()` (page 34) method. However, other updates may have modified the document between your update and the document retrieval. Also, if the update modified only a single document but multiple documents matched, you will need to use additional logic to identify the updated document.

- You cannot specify a `write concern` to `findAndModify()` (page 39) to override the default write concern whereas, starting in MongoDB 2.6, you can specify a write concern to the `update()` (page 69) method.

When modifying a *single* document, both `findAndModify()` (page 39) and the `update()` (page 69) method *atomically* update the document. See <http://docs.mongodb.org/manual/tutorial/isolate-sequence-of-operations> for more details about interactions and order of operations of these methods.

Examples

Update and Return The following method updates and returns an existing document in the `people` collection where the document matches the query criteria:

```
db.people.findAndModify({
  query: { name: "Tom", state: "active", rating: { $gt: 10 } },
  sort: { rating: 1 },
```

```
    update: { $inc: { score: 1 } }
  })
```

This method performs the following actions:

1. The query finds a document in the `people` collection where the `name` field has the value `Tom`, the `state` field has the value `active` and the `rating` field has a value `greater than` (page 373) `10`.
2. The `sort` orders the results of the query in ascending order. If multiple documents meet the query condition, the method will select for modification the first document as ordered by this `sort`.
3. The update increments the value of the `score` field by `1`.
4. The method returns the original (i.e. pre-modification) document selected for this update:

```
{
  "_id" : ObjectId("50f1e2c99beb36a0f45c6453"),
  "name" : "Tom",
  "state" : "active",
  "rating" : 100,
  "score" : 5
}
```

To return the modified document, add the `new:true` option to the method.

If no document matched the query condition, the method returns `null`:

`null`

Upsert The following method includes the `upsert: true` option for the update operation to either update a matching document or, if no matching document exists, create a new document:

```
db.people.findAndModify({
  query: { name: "Gus", state: "active", rating: 100 },
  sort: { rating: 1 },
  update: { $inc: { score: 1 } },
  upsert: true
})
```

If the method finds a matching document, the method performs an update.

If the method does **not** find a matching document, the method creates a new document. Because the method included the `sort` option, it returns an empty document `{ }` as the original (pre-modification) document:

```
{ }
```

If the method did **not** include a `sort` option, the method returns `null`.

`null`

Return New Document The following method includes both the `upsert: true` option and the `new:true` option. The method either updates a matching document and returns the updated document or, if no matching document exists, inserts a document and returns the newly inserted document in the `value` field.

In the following example, no document in the `people` collection matches the query condition:

```
db.people.findAndModify({
  query: { name: "Pascal", state: "active", rating: 25 },
  sort: { rating: 1 },
  update: { $inc: { score: 1 } },
  upsert: true,
  new: true
})
```

```

    upsert: true,
    new: true
  })

```

The method returns the newly inserted document:

```

{
  "_id" : ObjectId("50f49ad6444c11ac2448a5d6"),
  "name" : "Pascal",
  "rating" : 25,
  "score" : 1,
  "state" : "active"
}

```

Sort and Remove By including a sort specification on the `rating` field, the following example removes from the `people` collection a single document with the `state` value of `active` and the lowest rating among the matching documents:

```

db.people.findAndModify(
  {
    query: { state: "active" },
    sort: { rating: 1 },
    remove: true
  }
)

```

The method returns the deleted document:

```

{
  "_id" : ObjectId("52fba867ab5fdca1299674ad"),
  "name" : "XYZ123",
  "score" : 1,
  "state" : "active",
  "rating" : 3
}

```

db.collection.findOne()

Definition

`db.collection.findOne(<criteria>, <projection>)`

Returns one document that satisfies the specified query criteria. If multiple documents satisfy the query, this method returns the first document according to the *natural order* which reflects the order of documents on the disk. In *capped collections*, natural order is the same as insertion order.

param document criteria Specifies query selection criteria using *query operators* (page 373).

param document projection Specifies the fields to return using *projection operators* (page 406).

Omit this parameter to return all fields in the matching document.

The `projection` parameter takes a document of the following form:

```

{ field1: <boolean>, field2: <boolean> ... }

```

The `<boolean>` can be one of the following include or exclude values:

- 1 or `true` to include. The `findOne()` (page 43) method always includes the `_id` field even if the field is not explicitly specified in the *projection* parameter.
- 0 or `false` to exclude.

The projection argument cannot mix include and exclude specifications, with the exception of excluding the `_id` field.

returns One document that satisfies the criteria specified as the first argument to this method.

If you specify a projection parameter, `findOne()` (page 43) returns a document that only contains the projection fields. The `_id` field is always included unless you explicitly exclude it.

Although similar to the `find()` (page 34) method, the `findOne()` (page 43) method returns a document rather than a cursor.

Examples

With Empty Query Specification The following operation returns a single document from the `bios` collection:

```
db.bios.findOne()
```

With a Query Specification The following operation returns the first matching document from the `bios` collection where either the field `first` in the subdocument `name` starts with the letter `G` or where the field `birth` is less than `new Date('01/01/1945')`:

```
db.bios.findOne(
  {
    $or: [
      { 'name.first' : /^G/ },
      { birth: { $lt: new Date('01/01/1945') } }
    ]
  }
)
```

With a Projection The projection parameter specifies which fields to return. The parameter contains either include or exclude specifications, not both, unless the exclude is for the `_id` field.

Specify the Fields to Return The following operation finds a document in the `bios` collection and returns only the `name`, `contribs` and `_id` fields:

```
db.bios.findOne(
  { },
  { name: 1, contribs: 1 }
)
```

Return All but the Excluded Fields The following operation returns a document in the `bios` collection where the `contribs` field contains the element `OOP` and returns all fields *except* the `_id` field, the first field in the `name` subdocument, and the `birth` field:

```
db.bios.findOne(
  { contribs: 'OOP' },
  { _id: 0, 'name.first': 0, birth: 0 }
)
```


The `findOne` Result Document You cannot apply cursor methods to the result of `findOne()` (page 43) because a single document is returned. You have access to the document directly:

```
var myDocument = db.bios.findOne();

if (myDocument) {
  var myName = myDocument.name;

  print (toJson(myName));
}
```

`db.collection.getIndexes()`

`db.collection.getIndexes()`

Returns an array that holds a list of documents that identify and describe the existing indexes on the collection. You must call the `db.collection.getIndexes()` (page 45) on a collection. For example:

```
db.collection.getIndexes()
```

Change `collection` to the name of the collection whose indexes you want to learn.

The `db.collection.getIndexes()` (page 45) items consist of the following fields:

`system.indexes.v`

Holds the version of the index.

The index version depends on the version of `mongod` (page 503) that created the index. Before version 2.0 of MongoDB, the this value was 0; versions 2.0 and later use version 1.

`system.indexes.key`

Contains a document holding the keys held in the index, and the order of the index. Indexes may be either descending or ascending order. A value of negative one (e.g. `-1`) indicates an index sorted in descending order while a positive value (e.g. `1`) indicates an index sorted in an ascending order.

`system.indexes.ns`

The namespace context for the index.

`system.indexes.name`

A unique name for the index comprised of the field names and orders of all keys.

`db.collection.getShardDistribution()`

Definition

`db.collection.getShardDistribution()`

Returns

Prints the data distribution statistics for a *sharded* collection. You must call the `getShardDistribution()` (page 45) method on a sharded collection, as in the following example:

```
db.myShardedCollection.getShardDistribution()
```

In the following example, the collection has two shards. The output displays both the individual shard distribution information as well the total shard distribution:

```
Shard <shard-a> at <host-a>
  data : <size-a> docs : <count-a> chunks : <number of chunks-a>
  estimated data per chunk : <size-a>/<number of chunks-a>
  estimated docs per chunk : <count-a>/<number of chunks-a>

Shard <shard-b> at <host-b>
  data : <size-b> docs : <count-b> chunks : <number of chunks-b>
  estimated data per chunk : <size-b>/<number of chunks-b>
  estimated docs per chunk : <count-b>/<number of chunks-b>

Totals
  data : <stats.size> docs : <stats.count> chunks : <calc total chunks>
  Shard <shard-a> contains <estDataPercent-a>% data, <estDocPercent-a>% docs in cluster, avg obj
  Shard <shard-b> contains <estDataPercent-b>% data, <estDocPercent-b>% docs in cluster, avg obj
```

See also:

<http://docs.mongodb.org/manual/sharding>

Output The output information displays:

- <shard-x> is a string that holds the shard name.
- <host-x> is a string that holds the host name(s).
- <size-x> is a number that includes the size of the data, including the unit of measure (e.g. b, Mb).
- <count-x> is a number that reports the number of documents in the shard.
- <number of chunks-x> is a number that reports the number of chunks in the shard.
- <size-x>/<number of chunks-x> is a calculated value that reflects the estimated data size per chunk for the shard, including the unit of measure (e.g. b, Mb).
- <count-x>/<number of chunks-x> is a calculated value that reflects the estimated number of documents per chunk for the shard.
- <stats.size> is a value that reports the total size of the data in the sharded collection, including the unit of measure.
- <stats.count> is a value that reports the total number of documents in the sharded collection.
- <calc total chunks> is a calculated number that reports the number of chunks from all shards, for example:
$$\text{<calc total chunks>} = \text{<number of chunks-a>} + \text{<number of chunks-b>}$$
- <estDataPercent-x> is a calculated value that reflects, for each shard, the data size as the percentage of the collection's total data size, for example:
$$\text{<estDataPercent-x>} = \text{<size-x>/<stats.size>}$$
- <estDocPercent-x> is a calculated value that reflects, for each shard, the number of documents as the percentage of the total number of documents for the collection, for example:
$$\text{<estDocPercent-x>} = \text{<count-x>/<stats.count>}$$
- stats.shards[<shard-x>].avgObjSize is a number that reflects the average object size, including the unit of measure, for the shard.

Example Output For example, the following is a sample output for the distribution of a sharded collection:

```
Shard shard-a at shard-a/MyMachine.local:30000,MyMachine.local:30001,MyMachine.local:30002
data : 38.14Mb docs : 1000003 chunks : 2
estimated data per chunk : 19.07Mb
estimated docs per chunk : 500001
```

```
Shard shard-b at shard-b/MyMachine.local:30100,MyMachine.local:30101,MyMachine.local:30102
data : 38.14Mb docs : 999999 chunks : 3
estimated data per chunk : 12.71Mb
estimated docs per chunk : 333333
```

Totals

```
data : 76.29Mb docs : 2000002 chunks : 5
Shard shard-a contains 50% data, 50% docs in cluster, avg obj size on shard : 40b
Shard shard-b contains 49.99% data, 49.99% docs in cluster, avg obj size on shard : 40b
```

db.collection.getShardVersion()

`db.collection.getShardVersion()`

This method returns information regarding the state of data in a *sharded cluster* that is useful when diagnosing underlying issues with a sharded cluster.

For internal and diagnostic use only.

db.collection.group()

Definition

`db.collection.group({ key, reduce, initial, [keyf,] [cond,] finalize })`

Groups documents in a collection by the specified keys and performs simple aggregation functions such as computing counts and sums. The method is analogous to a `SELECT <...> GROUP BY` statement in SQL. The `group()` (page 47) method returns an array.

The `db.collection.group()` (page 47) accepts a single *document* that contains the following:

field document key The field or fields to group. Returns a “key object” for use as the grouping key.

field function reduce An aggregation function that operates on the documents during the grouping operation. These functions may return a sum or a count. The function takes two arguments: the current document and an aggregation result document for that group.

field document initial Initializes the aggregation result document.

field function keyf Alternative to the `key` field. Specifies a function that creates a “key object” for use as the grouping key. Use `keyf` instead of `key` to group by calculated fields rather than existing document fields.

field document cond The selection criteria to determine which documents in the collection to process. If you omit the `cond` field, `db.collection.group()` (page 47) processes all the documents in the collection for the group operation.

field function finalize A function that runs each item in the result set before `db.collection.group()` (page 47) returns the final value. This function can either modify the result document or replace the result document as a whole.

The `db.collection.group()` (page 47) method is a shell wrapper for the `group` (page 204) command. However, the `db.collection.group()` (page 47) method takes the `keyf` field and the `reduce` field whereas the `group` (page 204) command takes the `$keyf` field and the `$reduce` field.

Behavior

Limits and Restrictions The `db.collection.group()` (page 47) method does not work with *sharded clusters*. Use the *aggregation framework* or *map-reduce* in *sharded environments*.

The result set must fit within the *maximum BSON document size* (page 604).

In version 2.2, the returned array can contain at most 20,000 elements; i.e. at most 20,000 unique groupings. For group by operations that results in more than 20,000 unique groupings, use `mapReduce` (page 208). Previous versions had a limit of 10,000 elements.

Prior to 2.4, the `db.collection.group()` (page 47) method took the `mongod` (page 503) instance's JavaScript lock, which blocked all other JavaScript execution.

mongo Shell JavaScript Functions/Properties Changed in version 2.4: In MongoDB 2.4, *map-reduce operations* (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions **cannot** access certain global functions or properties, such as `db`, that are available in the `mongo` (page 527) shell.

When upgrading to MongoDB 2.4, you will need to refactor your code if your *map-reduce operations* (page 208), `group` (page 204) commands, or `$where` (page 391) operator expressions include any global shell functions or properties that are no longer available, such as `db`.

The following JavaScript functions and properties **are available** to *map-reduce operations* (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions in MongoDB 2.4:

Available Properties	Available Functions	
args MaxKey MinKey	assert() BinData() DBPointer() DBRef() doassert() emit() gc() HexData() hex_md5() isNumber() isObject() ISODate() isString()	Map() MD5() NumberInt() NumberLong() ObjectId() print() printjson() printjsononeline() sleep() Timestamp() tojson() tojsononeline() tojsonObject() UUID() version()

Examples The following examples assume an `orders` collection with documents of the following prototype:

```
{
  _id: ObjectId("5085a95c8fada716c89d0021"),
  ord_dt: ISODate("2012-07-01T04:00:00Z"),
  ship_dt: ISODate("2012-07-02T04:00:00Z"),
  item: { sku: "abc123",
    price: 1.99,
    uom: "pcs",
```

```

    qty: 25 }
}

```

Group by Two Fields The following example groups by the `ord_dt` and `item.sku` fields those documents that have `ord_dt` greater than 01/01/2011:

```

db.orders.group(
  {
    key: { ord_dt: 1, 'item.sku': 1 },
    cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
    reduce: function ( curr, result ) { },
    initial: { }
  }
)

```

The result is an array of documents that contain the group by fields:

```

[
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc123"},
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc456"},
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "bcd123"},
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "efg456"},
  { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "abc123"},
  { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "efg456"},
  { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "ijk123"},
  { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc123"},
  { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc456"},
  { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc123"},
  { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc456"}
]

```

The method call is analogous to the SQL statement:

```

SELECT ord_dt, item_sku
FROM orders
WHERE ord_dt > '01/01/2012'
GROUP BY ord_dt, item_sku

```

Calculate the Sum The following example groups by the `ord_dt` and `item.sku` fields, those documents that have `ord_dt` greater than 01/01/2011 and calculates the sum of the `qty` field for each grouping:

```

db.orders.group(
  {
    key: { ord_dt: 1, 'item.sku': 1 },
    cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
    reduce: function( curr, result ) {
      result.total += curr.item.qty;
    },
    initial: { total : 0 }
  }
)

```

The result is an array of documents that contain the group by fields and the calculated aggregation field:

```

[ { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc123", "total" : 25 },
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc456", "total" : 25 },
  { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "bcd123", "total" : 10 },

```

```
{ "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "efg456", "total" : 10 },
{ "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "abc123", "total" : 25 },
{ "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "efg456", "total" : 15 },
{ "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "ijk123", "total" : 20 },
{ "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc123", "total" : 45 },
{ "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc456", "total" : 25 },
{ "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc123", "total" : 25 },
{ "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc456", "total" : 25 } ]
```

The method call is analogous to the SQL statement:

```
SELECT ord_dt, item_sku, SUM(item_qty) as total
FROM orders
WHERE ord_dt > '01/01/2012'
GROUP BY ord_dt, item_sku
```

Calculate Sum, Count, and Average The following example groups by the calculated `day_of_week` field, those documents that have `ord_dt` greater than 01/01/2011 and calculates the sum, count, and average of the `qty` field for each grouping:

```
db.orders.group(
  {
    keyf: function(doc) {
      return { day_of_week: doc.ord_dt.getDay() };
    },
    cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
    reduce: function( curr, result ) {
      result.total += curr.item.qty;
      result.count++;
    },
    initial: { total : 0, count: 0 },
    finalize: function(result) {
      var weekdays = [
        "Sunday", "Monday", "Tuesday",
        "Wednesday", "Thursday",
        "Friday", "Saturday"
      ];
      result.day_of_week = weekdays[result.day_of_week];
      result.avg = Math.round(result.total / result.count);
    }
  }
)
```

The result is an array of documents that contain the group by fields and the calculated aggregation field:

```
[
  { "day_of_week" : "Sunday", "total" : 70, "count" : 4, "avg" : 18 },
  { "day_of_week" : "Friday", "total" : 110, "count" : 6, "avg" : 18 },
  { "day_of_week" : "Tuesday", "total" : 70, "count" : 3, "avg" : 23 }
]
```

See also:

<http://docs.mongodb.org/manualcore/aggregation>

`db.collection.initializeOrderedBulkOp()`

Definition

`db.collection.initializeOrderedBulkOp()`

Initializes and returns a new `Bulk()` (page 128) operations builder for a collection. The builder constructs an ordered list of write operations that MongoDB executes in bulk.

Returns new `Bulk()` (page 128) operations builder object.

Behavior

Order of Operation With an *ordered* operations list, MongoDB executes the write operations in the list serially.

Stop on Error If an error occurs during the processing of one of the write operations, MongoDB will return without processing any remaining write operations in the list.

Examples The following initializes a `Bulk()` (page 128) operations builder on the `users` collection, adds a series of write operations, and executes the operations:

```
var bulk = db.users.initializeOrderedBulkOp();
bulk.insert( { user: "abc123", status: "A", points: 0 } );
bulk.insert( { user: "ijk123", status: "A", points: 0 } );
bulk.insert( { user: "mop123", status: "P", points: 0 } );
bulk.find( { status: "D" } ).remove();
bulk.find( { status: "P" } ).update( { $set: { comment: "Pending" } } );
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.find.removeOne()` (page 130)
- `Bulk.execute()` (page 137)

`db.collection.initializeUnorderedBulkOp()`**Definition**

`db.collection.initializeUnorderedBulkOp()`

New in version 2.6.

Initializes and returns a new `Bulk()` (page 128) operations builder for a collection. The builder constructs an *unordered* list of write operations that MongoDB executes in bulk.

Behavior

Order of Operation With an *unordered* operations list, MongoDB can execute in parallel the write operations in the list and in any order. If the order of operations matter, use `db.collection.initializeOrderedBulkOp()` (page 51) instead.

Continue on Error If an error occurs during the processing of one of the write operations, MongoDB will continue to process remaining write operations in the list.

Example The following initializes a `Bulk()` (page 128) operations builder and adds a series of insert operations to add multiple documents:

```
var bulk = db.users.initializeUnorderedBulkOp();
bulk.insert( { user: "abc123", status: "A", points: 0 } );
bulk.insert( { user: "ijk123", status: "A", points: 0 } );
bulk.insert( { user: "mop123", status: "P", points: 0 } );
bulk.execute();
```

See also:

- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk()` (page 128)
- `Bulk.insert()` (page 129)
- `Bulk.execute()` (page 137)

`db.collection.insert()`

Definition

`db.collection.insert()`

Inserts a document or documents into a collection.

The `insert()` (page 52) method has the following syntax:

Changed in version 2.6.

```
db.collection.insert(
  <document or array of documents>,
  {
    writeConcern: <document>,
    ordered: <boolean>
  }
)
```

param document,array document A document or array of documents to insert into the collection.

param document writeConcern A document expressing the write concern. Omit to use the default write concern. See *Safe Writes* (page 53).

New in version 2.6.

param boolean ordered If `true`, perform an ordered insert of the documents in the array, and if an error occurs with one of documents, MongoDB will return without processing the remaining documents in the array. Defaults to `false`.

New in version 2.6.

Changed in version 2.6: The `insert()` (page 52) returns an object that contains the status of the operation.

Returns

- A *WriteResult* (page 54) object for single inserts.
- A *BulkWriteResult* (page 55) object for bulk inserts.

Behaviors

Safe Writes Changed in version 2.6.

The `insert()` (page 52) method uses the `insert` (page 220) command, which uses the default write concern. To specify a different write concern, include the write concern in the options parameter.

Create Collection If the collection does not exist, then the `insert()` (page 52) method will create the collection.

`_id` Field If the document does not specify an `_id` field, then MongoDB will add the `_id` field and assign a unique <http://docs.mongodb.org/manualreference/object-id> for the document before inserting. Most drivers create an `ObjectId` and insert the `_id` field, but the `mongod` (page 503) will create and populate the `_id` if the driver or application does not.

If the document contains an `_id` field, the `_id` value must be unique within the collection to avoid duplicate key error.

Examples The following examples insert documents into the `products` collection. If the collection does not exist, the `insert()` (page 52) method creates the collection.

Insert a Document without Specifying an `_id` Field In the following example, the document passed to the `insert()` (page 52) method does not contain the `_id` field:

```
db.products.insert( { item: "card", qty: 15 } )
```

During the insert, `mongod` (page 503) will create the `_id` field and assign it a unique <http://docs.mongodb.org/manualreference/object-id> value, as verified by the inserted document:

```
{ "_id" : ObjectId("5063114bd386d8fadbd6b004"), "item" : "card", "qty" : 15 }
```

The `ObjectId` values are specific to the machine and time when the operation is run. As such, your values may differ from those in the example.

Insert a Document Specifying an `_id` Field In the following example, the document passed to the `insert()` (page 52) method includes the `_id` field. The value of `_id` must be unique within the collection to avoid duplicate key error.

```
db.products.insert( { _id: 10, item: "box", qty: 20 } )
```

The operation inserts the following document in the `products` collection:

```
{ "_id" : 10, "item" : "box", "qty" : 20 }
```

Insert Multiple Documents The following example performs a bulk insert of three documents by passing an array of documents to the `insert()` (page 52) method.

The documents in the array do not need to have the same fields. For instance, the first document in the array has an `_id` field and a `type` field. Because the second and third documents do not contain an `_id` field, `mongod` (page 503) will create the `_id` field for the second and third documents during the insert:

```
db.products.insert(
  [
    { _id: 11, item: "pencil", qty: 50, type: "no.2" },
    { item: "pen", qty: 20 },
    { item: "eraser", qty: 25 }
  ]
)
```

The operation inserted the following three documents:

```
{ "_id" : 11, "item" : "pencil", "qty" : 50, "type" : "no.2" }
{ "_id" : ObjectId("51e0373c6f35bd826f47e9a0"), "item" : "pen", "qty" : 20 }
{ "_id" : ObjectId("51e0373c6f35bd826f47e9a1"), "item" : "eraser", "qty" : 25 }
```

Perform an Ordered Insert The following example performs an *ordered* insert of four documents. Unlike *unordered* inserts which continue on error, *ordered* inserts return on error without processing the remaining documents in the array.

```
db.products.insert(
  [
    { _id: 20, item: "lamp", qty: 50, type: "desk" },
    { _id: 21, item: "lamp", qty: 20, type: "floor" },
    { _id: 22, item: "bulk", qty: 100 }
  ],
  { ordered: true }
)
```

Override Default Write Concern The following operation to a replica set specifies a write concern of "w: majority" with a wtimeout of 5000 milliseconds such that the method returns after the write propagates to a majority of the replica set members or the method times out after 5 seconds.

```
db.products.insert(
  { item: "envelopes", qty : 100, type: "Clasp" },
  { writeConcern: { w: "majority", wtimeout: 5000 } }
)
```

WriteResult Changed in version 2.6.

When passed a single document, `insert()` (page 52) returns a `WriteResult` object.

Successful Results The `insert()` (page 52) returns a `WriteResult` (page 188) object that contains the status of the operation. Upon success, the `WriteResult` (page 188) object contains information on the number of documents inserted:

```
WriteResult({ "nInserted" : 1 })
```

Write Concern Errors If the `insert()` (page 52) method encounters write concern errors, the results include the `WriteResult.writeConcernError` (page 188) field:

```
WriteResult({
  "nInserted" : 1,
  "writeConcernError" : {
    "code" : 64,
    "errmsg" : "waiting for replication timed out at shard-a"
  }
})
```

Errors Unrelated to Write Concern If the `insert()` (page 52) method encounters a non-write concern error, the results include the `WriteResult.writeError` (page 188) field:

```
WriteResult({
  "nInserted" : 0,
  "writeError" : {
    "code" : 11000,
    "errmsg" : "insertDocument :: caused by :: 11000 E11000 duplicate key error index: test.foo.$_id"
  }
})
```

BulkWriteResult Changed in version 2.6.

When passed an array of documents, `insert()` (page 52) returns a *bulk-write-result*. See *bulk-write-result* for details.

`db.collection.isCapped()`

```
db.collection.isCapped()
```

Returns Returns `true` if the collection is a *capped collection*, otherwise returns `false`.

See also:

<http://docs.mongodb.org/manualcore/capped-collections>

`db.collection.mapReduce()`

```
db.collection.mapReduce(map, reduce, {<out>, <query>, <sort>, <limit>, <finalize>, <scope>,
  <jsMode>, <verbose>})
```

The `db.collection.mapReduce()` (page 55) method provides a wrapper around the `mapReduce` (page 208) command.

```
db.collection.mapReduce(
  <map>,
  <reduce>,
  {
    out: <collection>,
    query: <document>,
    sort: <document>,
    limit: <number>,
    finalize: <function>,
    scope: <document>,
    jsMode: <boolean>,
    verbose: <boolean>
  }
)
```

`db.collection.mapReduce()` (page 55) takes the following parameters:

field Javascript function map A JavaScript function that associates or “maps” a value with a key and emits the key and value pair.

See *Requirements for the map Function* (page 57) for more information.

field JavaScript function reduce A JavaScript function that “reduces” to a single object all the values associated with a particular key.

See *Requirements for the reduce Function* (page 58) for more information.

field document options A document that specifies additional parameters to `db.collection.mapReduce()` (page 55).

The following table describes additional arguments that `db.collection.mapReduce()` (page 55) can accept.

field string or document out Specifies the location of the result of the map-reduce operation. You can output to a collection, output to a collection with an action, or output inline. You may output to a collection when performing map reduce operations on the primary members of the set; on *secondary* members you may only use the `inline` output.

See *out Options* (page 58) for more information.

field document query Specifies the selection criteria using *query operators* (page 373) for determining the documents input to the `map` function.

field document sort Sorts the *input* documents. This option is useful for optimization. For example, specify the sort key to be the same as the emit key so that there are fewer reduce operations. The sort key must be in an existing index for this collection.

field number limit Specifies a maximum number of documents to return from the collection.

field Javascript function finalize Follows the `reduce` method and modifies the output.

See *Requirements for the finalize Function* (page 59) for more information.

field document scope Specifies global variables that are accessible in the `map`, `reduce` and `finalize` functions.

field Boolean jsMode Specifies whether to convert intermediate data into BSON format between the execution of the `map` and `reduce` functions. Defaults to `false`.

If `false`:

- Internally, MongoDB converts the JavaScript objects emitted by the `map` function to BSON objects. These BSON objects are then converted back to JavaScript objects when calling the `reduce` function.
- The map-reduce operation places the intermediate BSON objects in temporary, on-disk storage. This allows the map-reduce operation to execute over arbitrarily large data sets.

If `true`:

- Internally, the JavaScript objects emitted during `map` function remain as JavaScript objects. There is no need to convert the objects for the `reduce` function, which can result in faster execution.
- You can only use `jsMode` for result sets with fewer than 500,000 distinct `key` arguments to the mapper's `emit()` function.

The `jsMode` defaults to `false`.

field Boolean verbose Specifies whether to include the `timing` information in the result information. The `verbose` defaults to `true` to include the `timing` information.

Note: Changed in version 2.4.

In MongoDB 2.4, *map-reduce operations* (page 208), the *group* (page 204) command, and *\$where* (page 391) operator expressions **cannot** access certain global functions or properties, such as `db`, that are available in the *mongo* (page 527) shell.

When upgrading to MongoDB 2.4, you will need to refactor your code if your *map-reduce operations* (page 208), *group* (page 204) commands, or *\$where* (page 391) operator expressions include any global shell functions or properties that are no longer available, such as `db`.

The following JavaScript functions and properties **are available** to `map-reduce` operations (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions in MongoDB 2.4:

Available Properties	Available Functions	
args MaxKey MinKey	assert() BinData() DBPointer() DBRef() doassert() emit() gc() HexData() hex_md5() isNumber() isObject() ISODate() isString()	Map() MD5() NumberInt() NumberLong() ObjectId() print() printjson() printjsononeline() sleep() Timestamp() tojson() tojsononeline() tojsonObject() UUID() version()

Requirements for the map Function The map function has the following prototype:

```
function() {
  ...
  emit(key, value);
}
```

The map function exhibits the following behaviors:

- In the map function, reference the current document as `this` within the function.
- The map function should *not* access the database for any reason.
- The map function should be pure, or have *no* impact outside of the function (i.e. side effects.)
- The `emit(key, value)` function associates the key with a value.
 - A single emit can only hold half of MongoDB's *maximum BSON document size* (page 604).
 - The map function can call `emit(key, value)` any number of times, including 0, per each input document.

The following map function may call `emit(key, value)` either 0 or 1 times depending on the value of the input document's `status` field:

```
function() {
  if (this.status == 'A')
    emit(this.cust_id, 1);
}
```

The following map function may call `emit(key, value)` multiple times depending on the number of elements in the input document's `items` field:

```
function() {  
  this.items.forEach(function(item){ emit(item.sku, 1); });  
}
```

- The map function can access the variables defined in the `scope` parameter.

Requirements for the `reduce` Function The `reduce` function has the following prototype:

```
function(key, values) {  
  ...  
  return result;  
}
```

The `reduce` function exhibits the following behaviors:

- The `reduce` function should *not* access the database, even to perform read operations.
- The `reduce` function should *not* affect the outside system.
- MongoDB will **not** call the `reduce` function for a key that has only a single value. The `values` argument is an array whose elements are the value objects that are “mapped” to the key.
- MongoDB can invoke the `reduce` function more than once for the same key. In this case, the previous output from the `reduce` function for that key will become one of the input values to the next `reduce` function invocation for that key.
- The `reduce` function can access the variables defined in the `scope` parameter.

Because it is possible to invoke the `reduce` function more than once for the same key, the following properties need to be true:

- the *type* of the return object must be **identical** to the type of the value emitted by the map function to ensure that the following operations is true:

```
reduce(key, [ C, reduce(key, [ A, B ]) ] ) == reduce( key, [ C, A, B ] )
```

- the `reduce` function must be *idempotent*. Ensure that the following statement is true:

```
reduce( key, [ reduce(key, valuesArray) ] ) == reduce( key, valuesArray )
```

- the order of the elements in the `valuesArray` should not affect the output of the `reduce` function, so that the following statement is true:

```
reduce( key, [ A, B ] ) == reduce( key, [ B, A ] )
```

out Options You can specify the following options for the `out` parameter:

Output to a Collection

```
out: <collectionName>
```

Output to a Collection with an Action This option is only available when passing `out` a collection that already exists. This option is not available on secondary members of replica sets.

```
out: { <action>: <collectionName>  
  [, db: <dbName>]  
  [, sharded: <boolean> ]  
  [, nonAtomic: <boolean> ] }
```

When you output to a collection with an action, the `out` has the following parameters:

- `<action>`: Specify one of the following actions:
 - `replace`
Replace the contents of the `<collectionName>` if the collection with the `<collectionName>` exists.
 - `merge`
Merge the new result with the existing result if the output collection already exists. If an existing document has the same key as the new result, *overwrite* that existing document.
 - `reduce`
Merge the new result with the existing result if the output collection already exists. If an existing document has the same key as the new result, apply the `reduce` function to both the new and the existing documents and overwrite the existing document with the result.
- `db`:
Optional. The name of the database that you want the map-reduce operation to write its output. By default this will be the same database as the input collection.
- `sharded`:
Optional. If `true` *and* you have enabled sharding on output database, the map-reduce operation will shard the output collection using the `_id` field as the shard key.
- `nonAtomic`:
New in version 2.2.
Optional. Specify output operation as non-atomic and is valid *only* for `merge` and `reduce` output modes which may take minutes to execute.

If `nonAtomic` is `true`, the post-processing step will prevent MongoDB from locking the database; however, other clients will be able to read intermediate states of the output collection. Otherwise the map reduce operation must lock the database during post-processing.

Output Inline Perform the map-reduce operation in memory and return the result. This option is the only available option for `out` on secondary members of replica sets.

```
out: { inline: 1 }
```

The result must fit within the *maximum size of a BSON document* (page 604).

Requirements for the `finalize` Function The `finalize` function has the following prototype:

```
function(key, reducedValue) {
  ...
  return modifiedObject;
}
```

The `finalize` function receives as its arguments a key value and the `reducedValue` from the `reduce` function. Be aware that:

- The `finalize` function should *not* access the database for any reason.
- The `finalize` function should be pure, or have *no* impact outside of the function (i.e. side effects.)
- The `finalize` function can access the variables defined in the `scope` parameter.

Map-Reduce Examples Consider the following map-reduce operations on a collection `orders` that contains documents of the following prototype:

```
{
  _id: ObjectId("50a8240b927d5d8b5891743c"),
  cust_id: "abc123",
  ord_date: new Date("Oct 04, 2012"),
  status: 'A',
  price: 25,
  items: [ { sku: "mmm", qty: 5, price: 2.5 },
            { sku: "nnn", qty: 5, price: 2.5 } ]
}
```

Return the Total Price Per Customer Perform the map-reduce operation on the `orders` collection to group by the `cust_id`, and calculate the sum of the `price` for each `cust_id`:

1. Define the map function to process each input document:

- In the function, `this` refers to the document that the map-reduce operation is processing.
- The function maps the `price` to the `cust_id` for each document and emits the `cust_id` and price pair.

```
var mapFunction1 = function() {
    emit(this.cust_id, this.price);
};
```

2. Define the corresponding reduce function with two arguments `keyCustId` and `valuesPrices`:

- The `valuesPrices` is an array whose elements are the price values emitted by the map function and grouped by `keyCustId`.
- The function reduces the `valuesPrice` array to the sum of its elements.

```
var reduceFunction1 = function(keyCustId, valuesPrices) {
    return Array.sum(valuesPrices);
};
```

3. Perform the map-reduce on all documents in the `orders` collection using the `mapFunction1` map function and the `reduceFunction1` reduce function.

```
db.orders.mapReduce(
    mapFunction1,
    reduceFunction1,
    { out: "map_reduce_example" }
)
```

This operation outputs the results to a collection named `map_reduce_example`. If the `map_reduce_example` collection already exists, the operation will replace the contents with the results of this map-reduce operation:

Calculate Order and Total Quantity with Average Quantity Per Item In this example, you will perform a map-reduce operation on the `orders` collection for all documents that have an `ord_date` value greater than 01/01/2012. The operation groups by the `item.sku` field, and calculates the number of orders and the total quantity ordered for each `sku`. The operation concludes by calculating the average quantity per order for each `sku` value:

1. Define the map function to process each input document:

- In the function, `this` refers to the document that the map-reduce operation is processing.

- For each item, the function associates the sku with a new object value that contains the count of 1 and the item qty for the order and emits the sku and value pair.

```
var mapFunction2 = function() {
    for (var idx = 0; idx < this.items.length; idx++) {
        var key = this.items[idx].sku;
        var value = {
            count: 1,
            qty: this.items[idx].qty
        };
        emit(key, value);
    }
};
```

2. Define the corresponding reduce function with two arguments keySKU and countObjVals:

- countObjVals is an array whose elements are the objects mapped to the grouped keySKU values passed by map function to the reducer function.
- The function reduces the countObjVals array to a single object reducedValue that contains the count and the qty fields.
- In reducedVal, the count field contains the sum of the count fields from the individual array elements, and the qty field contains the sum of the qty fields from the individual array elements.

```
var reduceFunction2 = function(keySKU, countObjVals) {
    reducedVal = { count: 0, qty: 0 };

    for (var idx = 0; idx < countObjVals.length; idx++) {
        reducedVal.count += countObjVals[idx].count;
        reducedVal.qty += countObjVals[idx].qty;
    }

    return reducedVal;
};
```

3. Define a finalize function with two arguments key and reducedVal. The function modifies the reducedVal object to add a computed field named avg and returns the modified object:

```
var finalizeFunction2 = function (key, reducedVal) {

    reducedVal.avg = reducedVal.qty/reducedVal.count;

    return reducedVal;

};
```

4. Perform the map-reduce operation on the orders collection using the mapFunction2, reduceFunction2, and finalizeFunction2 functions.

```
db.orders.mapReduce( mapFunction2,
                    reduceFunction2,
                    {
                        out: { merge: "map_reduce_example" },
                        query: { ord_date:
                                { $gt: new Date('01/01/2012') }
                            },
                        finalize: finalizeFunction2
                    }
                )
```

This operation uses the `query` field to select only those documents with `ord_date` greater than `new Date(01/01/2012)`. Then it output the results to a collection `map_reduce_example`. If the `map_reduce_example` collection already exists, the operation will merge the existing contents with the results of this map-reduce operation.

For more information and examples, see the Map-Reduce page and <http://docs.mongodb.org/manual/tutorial/perfor>

See also:

- <http://docs.mongodb.org/manual/tutorial/troubleshoot-map-function>
- <http://docs.mongodb.org/manual/tutorial/troubleshoot-reduce-function>
- [mapReduce](#) (page 208) command
- <http://docs.mongodb.org/manualcore/aggregation>

`db.collection.reIndex()`

`db.collection.reIndex()`

The `db.collection.reIndex()` (page 62) drops all indexes on a collection and recreates them. This operation may be expensive for collections that have a large amount of data and/or a large number of indexes.

Call this method, which takes no arguments, on a collection object. For example:

```
db.collection.reIndex()
```

Normally, MongoDB compacts indexes during routine updates. For most users, the `db.collection.reIndex()` (page 62) is unnecessary. However, it may be worth running if the collection size has changed significantly or if the indexes are consuming a disproportionate amount of disk space.

Behavior

Note: For replica sets, `db.collection.reIndex()` (page 62) will not propagate from the *primary* to *secondaries*. `db.collection.reIndex()` (page 62) will only affect a single `mongod` (page 503) instance.

Important: `db.collection.reIndex()` (page 62) will rebuild indexes in the *background* if the index was originally specified with this option. However, `db.collection.reIndex()` (page 62) will rebuild the `_id` index in the foreground, which takes the database's write lock.

Changed in version 2.6: Reindexing operations will error if the index entry for an indexed field exceeds the `Maximum Index Key Length`. Reindexing operations occur as part of `compact` (page 313) and `repairDatabase` (page 319) commands as well as the `db.collection.reIndex()` (page 62) method.

Because these operations drop *all* the indexes from a collection and then recreate them sequentially, the error from the `Maximum Index Key Length` prevents these operations from rebuilding any remaining indexes for the collection and, in the case of the `repairDatabase` (page 319) command, from continuing with the remainder of the process.

See

<http://docs.mongodb.org/manualcore/index-creation> for more information on the behavior of indexing operations in MongoDB.

`db.collection.remove()`

Definition

```
db.collection.remove()
```

Removes documents from a collection.

The `db.collection.remove()` (page 62) method can have one of two syntaxes. The `remove()` (page 62) method can take a query document and an optional `justOne` boolean:

```
db.collection.remove(
  <query>,
  <justOne>
)
```

Or the method can take a query document and an optional remove options document:

New in version 2.6.

```
db.collection.remove(
  <query>,
  {
    justOne: <boolean>,
    writeConcern: <document>
  }
)
```

param document query Specifies deletion criteria using *query operators* (page 373). To delete all documents in a collection, pass an empty document (`{}`).

Changed in version 2.6: In previous versions, the method invoked with no query parameter deleted all documents in a collection.

param boolean justOne To limit the deletion to just one document, set to `true`. Omit to use the default value of `false` and delete all documents matching the deletion criteria.

param document writeConcern A document expressing the write concern. Omit to use the default write concern. See *Safe Writes* (page 63).

New in version 2.6.

Changed in version 2.6: The `remove()` (page 62) returns an object that contains the status of the operation.

Returns A *WriteResult* (page 64) object that contains the status of the operation.

Behavior

Safe Writes Changed in version 2.6.

The `remove()` (page 62) method uses the `delete` (page 226) command, which uses the default write concern. To specify a different write concern, include the write concern in the options parameter.

Query Considerations By default, `remove()` (page 62) removes all documents that match the query expression. Specify the `justOne` option to limit the operation to removing a single document. To delete a single document sorted by a specified order, use the *findAndModify()* (page 43) method.

Capped Collections You cannot use the `remove()` (page 62) method with a *capped collection*.

Sharded Collections All `remove()` (page 62) operations for a sharded collection that specify the `justOne` option must include the *shard key* or the `_id` field in the query specification. `remove()` (page 62) operations specifying `justOne` in a sharded collection without the *shard key* or the `_id` field return an error.

Examples The following are examples of the `remove()` (page 62) method.

Remove All Documents from a Collection To remove all documents in a collection, call the `remove` (page 62) method with an empty query document `{}`. The following operation deletes all documents from the `bios` collection:

```
db.bios.remove( { } )
```

This operation is not equivalent to the `drop()` (page 29) method.

To remove all documents from a collection, it may be more efficient to use the `drop()` (page 29) method to drop the entire collection, including the indexes, and then recreate the collection and rebuild the indexes.

Remove All Documents that Match a Condition To remove the documents that match a deletion criteria, call the `remove()` (page 62) method with the `<query>` parameter:

The following operation removes all the documents from the collection `products` where `qty` is greater than 20:

```
db.products.remove( { qty: { $gt: 20 } } )
```

Override Default Write Concern The following operation to a replica set removes all the documents from the collection `products` where `qty` is greater than 20 and specifies a write concern of `"w: majority"` with a `wtimeout` of 5000 milliseconds such that the method returns after the write propagates to a majority of the replica set members or the method times out after 5 seconds.

```
db.products.remove(
  { qty: { $gt: 20 } },
  { writeConcern: { w: "majority", wtimeout: 5000 } }
)
```

Remove a Single Document that Matches a Condition To remove the first document that match a deletion criteria, call the `remove` (page 62) method with the query criteria and the `justOne` parameter set to `true` or `1`.

The following operation removes the first document from the collection `products` where `qty` is greater than 20:

```
db.products.remove( { qty: { $gt: 20 } }, true )
```

Isolate Removal Operations If the `<query>` argument to the `remove()` (page 62) method matches multiple documents in the collection, the delete operation may interleave with other write operations to that collection. For an unsharded collection, you have the option to override this behavior with the `$isolated` (page 436) isolation operator, effectively isolating the delete operation and blocking other write operations during the delete. To isolate the query, include `$isolated: 1` in the `<query>` parameter as in the following examples:

```
db.products.remove( { qty: { $gt: 20 }, $isolated: 1 } )
```

WriteResult Changed in version 2.6.

Successful Results The `remove()` (page 62) returns a `WriteResult` (page 188) object that contains the status of the operation. Upon success, the `WriteResult` (page 188) object contains information on the number of documents removed:

```
WriteResult({ "nRemoved" : 4 })
```

See also:

`WriteResult.nRemoved` (page 188)

Write Concern Errors If the `remove()` (page 62) method encounters write concern errors, the results include the `WriteResult.writeConcernError` (page 188) field:

```
WriteResult({
  "nRemoved" : 21,
  "writeConcernError" : {
    "code" : 64,
    "errInfo" : {
      "wtimeout" : true
    },
    "errmsg" : "waiting for replication timed out"
  }
})
```

See also:

`WriteResult.hasWriteConcernError()` (page 189)

Errors Unrelated to Write Concern If the `remove()` (page 62) method encounters a non-write concern error, the results include `WriteResult.writeError` (page 188) field:

```
WriteResult({
  "nRemoved" : 0,
  "writeError" : {
    "code" : 2,
    "errmsg" : "unknown top level operator: $invalidFieldName"
  }
})
```

See also:

`WriteResult.hasWriteError()` (page 189)

db.collection.renameCollection()**Definition**

`db.collection.renameCollection(target, dropTarget)`

Renames a collection. Provides a wrapper for the `renameCollection` (page 299) *database command*.

param string target The new name of the collection. Enclose the string in quotes.

param boolean dropTarget If `true`, `mongod` (page 503) drops the *target* of `renameCollection` (page 299) prior to renaming the collection.

Example Call the `db.collection.renameCollection()` (page 65) method on a collection object. For example:

```
db.rrecord.renameCollection("record")
```

This operation will rename the `rrecord` collection to `record`. If the target name (i.e. `record`) is the name of an existing collection, then the operation will fail.

Limitations The method has the following limitations:

- `db.collection.renameCollection()` (page 65) cannot move a collection between databases. Use `renameCollection` (page 299) for these rename operations.
- `db.collection.renameCollection()` (page 65) cannot operation on sharded collections.

The `db.collection.renameCollection()` (page 65) method operates within a collection by changing the metadata associated with a given collection.

Refer to the documentation `renameCollection` (page 299) for additional warnings and messages.

Warning: The `db.collection.renameCollection()` (page 65) method and `renameCollection` (page 299) command will invalidate open cursors which interrupts queries that are currently returning data.

`db.collection.save()`

Definition

`db.collection.save()`

Updates an existing document or inserts a new document, depending on its `document` parameter.

The `save()` (page 66) method has the following form:

Changed in version 2.6.

```
db.collection.save(  
  <document>,  
  {  
    writeConcern: <document>  
  }  
)
```

param document document A document to save to the collection.

param document writeConcern A document expressing the write concern. Omit to use the default write concern. See *Safe Writes* (page 66).

New in version 2.6.

Changed in version 2.6: The `save()` (page 66) returns an object that contains the status of the operation.

Returns A *WriteResult* (page 68) object that contains the status of the operation.

Behavior

Safe Writes Changed in version 2.6.

The `save()` (page 66) method uses either the `insert` (page 220) or the `update` (page 222) command, which use the default write concern. To specify a different write concern, include the write concern in the options parameter.

Insert If the document does **not** contain an `_id` field, then the `save()` (page 66) method calls the `insert()` (page 52) method. During the operation, the `mongo` (page 527) shell will create an `http://docs.mongodb.org/manualreference/object-id` and assign it to the `_id` field.

Note: Most MongoDB driver clients will include the `_id` field and generate an `ObjectId` before sending the insert

operation to MongoDB; however, if the client sends a document without an `_id` field, the `mongod` (page 503) will add the `_id` field and generate the `ObjectId`.

Upsert If the document contains an `_id` field, then the `save()` (page 66) method calls the `update()` (page 69) method with the *upsert option* (page 70) and a query on the `_id` field. If a document does not exist with the specified `_id` value, the `save()` (page 66) method, i.e. the *update() method with the upsert option* (page 70), results in an insertion of the document. If a document exists with the specified `_id` value, the `save()` (page 66) method performs an update that replaces **all** fields in the existing document with the fields from the document.

Examples

Save a New Document without Specifying an `_id` Field In the following example, `save()` (page 66) method performs an insert since the document passed to the method does not contain the `_id` field:

```
db.products.save( { item: "book", qty: 40 } )
```

During the insert, `mongod` (page 503) will create the `_id` field with a unique <http://docs.mongodb.org/manualreference/object-id> value, as verified by the inserted document:

```
{ "_id" : ObjectId("50691737d386d8fadbd6b01d"), "item" : "book", "qty" : 40 }
```

The `ObjectId` values are specific to the machine and time when the operation is run. As such, your values may differ from those in the example.

Save a New Document Specifying an `_id` Field In the following example, `save()` (page 66) performs an update with `upsert` since the document contains an `_id` field:

```
db.products.save( { _id: 100, item: "water", qty: 30 } )
```

Because the `_id` field holds a value that *does not* exist in the collection, the update operation results in an insertion of the document. The results of these operations are identical to an *update() method with the upsert flag* (page 70) set to `true` or `1`.

The operation results in the following new document in the `products` collection:

```
{ "_id" : 100, "item" : "water", "qty" : 30 }
```

Replace an Existing Document The `products` collection contains the following document:

```
{ "_id" : 100, "item" : "water", "qty" : 30 }
```

The `save()` (page 66) method performs an update with `upsert` since the document contains an `_id` field:

```
db.products.save( { _id : 100, item : "juice" } )
```

Because the `_id` field holds a value that exists in the collection, the operation performs an update to replace the document and results in the following document:

```
{ "_id" : 100, "item" : "juice" }
```

Override Default Write Concern The following operation to a replica set specifies a `writeConcern` of "w: majority" with a `wtimeout` of 5000 milliseconds such that the method returns after the write propagates to a majority of the replica set members or the method times out after 5 seconds.

```
db.products.save(  
  { item: "envelopes", qty : 100, type: "Clasp" },  
  { writeConcern: { w: "majority", wtimeout: 5000 } }  
)
```

WriteResult Changed in version 2.6.

The `save()` (page 66) returns a `WriteResult` (page 188) object that contains the status of the insert or update operation. See *WriteResult for insert* (page 54) and *WriteResult for update* (page 75) for details.

`db.collection.stats()`

Definition

`db.collection.stats(scale)`

Returns statistics about the collection. The method includes the following parameter:

param number scale The scale used in the output to display the sizes of items. By default, output displays sizes in bytes. To display kilobytes rather than bytes, specify a `scale` value of 1024.

Returns A *document* containing statistics that reflecting the state of the specified collection.

The `stats()` (page 68) method provides a wrapper around the database command `collStats` (page 325).

Example The following operation returns stats on the `people` collection:

```
db.people.stats()
```

See also:

collStats (page 325) for an overview of the output of this command.

`db.collection.storageSize()`

`db.collection.storageSize()`

Returns The total amount of storage allocated to this collection for document storage. Provides a wrapper around the `storageSize` (page 326) field of the `collStats` (page 325) (i.e. `db.collection.stats()` (page 68)) output.

`db.collection.totalSize()`

`db.collection.totalSize()`

Returns The total size of the data in the collection plus the size of every indexes on the collection.

db.collection.totalIndexSize()

```
db.collection.totalIndexSize()
```

Returns The total size of all indexes for the collection. This method provides a wrapper around the `totalIndexSize` (page 327) output of the `collStats` (page 325) (i.e. `db.collection.stats()` (page 68)) operation.

db.collection.update()**Definition**

```
db.collection.update(query, update, options)
```

Modifies an existing document or documents in a collection. The method can modify specific fields of an existing document or documents or replace an existing document entirely, depending on the *update parameter* (page 70).

By default, the `update()` (page 69) method updates a **single** document. Set the *Multi Parameter* (page 71) to update all documents that match the query criteria.

The `update()` (page 69) method has the following form:

Changed in version 2.6.

```
db.collection.update(
  <query>,
  <update>,
  {
    upsert: <boolean>,
    multi: <boolean>,
    writeConcern: <document>
  }
)
```

The `update()` (page 69) method takes the following parameters:

param document query The selection criteria for the update. Use the same *query selectors* (page 373) as used in the `find()` (page 34) method.

param document update The modifications to apply. For details see *Update Parameter* (page 70).

param boolean upsert If set to `true`, creates a new document when no document matches the query criteria. The default value is `false`, which does *not* insert a new document when no match is found.

param boolean multi If set to `true`, updates multiple documents that meet the query criteria. If set to `false`, updates one document. The default value is `false`. For additional information, see *Multi Parameter* (page 71).

param document writeConcern A document expressing the write concern. Omit to use the default write concern. See *Safe Writes* (page 70).

New in version 2.6.

Changed in version 2.6: The `update()` (page 69) method returns an object that contains the status of the operation.

Returns A *WriteResult* (page 75) object that contains the status of the operation.

Behavior

Safe Writes Changed in version 2.6.

The `update()` (page 69) method uses the `update` (page 222) command, which uses the default write concern. To specify a different write concern, include the `writeConcern` option in the options parameter. See *Override Default Write Concern* (page 73) for an example.

Update Parameter The `update()` (page 69) method either modifies specific fields in existing documents or replaces an existing document entirely.

Update Specific Fields If the `<update>` document contains *update operator* (page 412) expressions, such as those using the `$set` (page 416) operator, then:

- The `<update>` document must contain *only update operator* (page 412) expressions.
- The `update()` (page 69) method updates only the corresponding fields in the document. For an example, see *Update Specific Fields* (page 71).

Replace Document Entirely If the `<update>` document contains *only field:value* expressions, then:

- The `update()` (page 69) method *replaces* the matching document with the `<update>` document. The `update()` (page 69) method *does not* replace the `_id` value. For an example, see *Replace All Fields* (page 72).
- `update()` (page 69) *cannot update multiple documents* (page 71).

Upsert Parameter

Upsert Use If `upsert` is `true` and if no document matches the query criteria, `update()` (page 69) inserts a *single* document. The *upsert* creates the new document with either:

- The fields and values of the `<update>` parameter, or
- The fields and values of both the `<query>` and `<update>` parameters. The *upsert* creates a document with data from both `<query>` and `<update>` if the `<update>` parameter contains *only update operator* (page 412) expressions.

If `upsert` is `true` and there are documents that match the query criteria, `update()` (page 69) performs an update.

Use Unique Indexes

Warning: To avoid inserting the same document more than once, only use `upsert: true` if the query filter is uniquely indexed.

Given a collection named `people` where no documents have a `name` field that holds the value `Andy`. Consider when multiple clients issue the following `update` with an `upsert` parameter at the same time:

```
db.people.update(
  { name: "Andy" },
  {
    name: "Andy",
    rating: 1,
    score: 1
  },
  { upsert: true }
)
```

If all `update()` (page 69) operations complete the `query` portion before any client successfully inserts data, **and** there is no unique index on the `name` field, then each update operation may result in an insert.

To prevent MongoDB from inserting the same document more than once, create a *unique index* on the `name` field. With a unique index, if an applications issues a group of upsert operations, *exactly one* `update()` (page 69) would successfully insert a new document.

The remaining operations would either:

- update the newly inserted document, or
- fail when they attempted to insert a duplicate.

If the operation fails because of a duplicate index key error, applications may retry the operation which will succeed as an update operation.

Multi Parameter If `multi` is set to `true`, the `update()` (page 69) method updates all documents that meet the `<query>` criteria. The `multi` update operation may interleave with other write operations. For unsharded collections, you can override this behavior with the `$isolated` (page 436) operator, which isolates the update operation and blocks other write operations during the update.

If the `<update>` (page 70) document contains *only* `field:value` expressions, then `update()` (page 69) *cannot* update multiple documents.

For an example, see *Update Multiple Documents* (page 73).

Sharded Collections All `update()` (page 69) operations for a sharded collection that specify the `multi: false` option must include the *shard key* or the `_id` field in the query specification. `update()` (page 69) operations specifying `multi: false` in a sharded collection without the *shard key* or the `_id` field return an error.

See also:

`findAndModify()` (page 39)

Examples

Update Specific Fields To update specific fields in a document, use *update operators* (page 412) in the `<update>` parameter. If the `<update>` parameter refers to non-existent fields in the current document, the `update()` (page 69) method adds the fields to the document.

For example, given a `books` collection with the following document:

```
{ "_id" : 11, "item" : "Divine Comedy", "stock" : 2 }
```

The following operation adds a `price` field to the document and increments the `stock` field by 5.

```
db.books.update(
  { item: "Divine Comedy" },
  {
    $set: { price: 18 },
    $inc: { stock: 5 }
  }
)
```

The updated document is now the following:

```
{ "_id" : 11, "item" : "Divine Comedy", "price" : 18, "stock" : 7 }
```

See also:

`$set` (page 416), `$inc` (page 412), *Update Operators* (page 412)

Remove Fields The following operation uses the `$unset` (page 417) operator to remove the `stock` field:

```
db.books.update( { _id: 11 }, { $unset: { stock: 1 } } )
```

See also:

`$unset` (page 417), `$rename` (page 414), *Update Operators* (page 412)

Replace All Fields Given the following document in the `books` collection:

```
{
  "_id" : 22,
  "item" : "The Banquet",
  "author" : "Dante",
  "price" : 20,
  "stock" : 4
}
```

The following operation passes an `<update>` document that contains only field and value pairs, which means the document replaces all the fields in the original document. The operation *does not* replace the `_id` value. The operation contains the same value for the `item` field in both the `<query>` and `<update>` documents, which means the field does not change:

```
db.books.update(
  { item: "The Banquet" },
  { item: "The Banquet", price: 19 , stock: 3 }
)
```

The operation creates the following new document. The operation removed the `author` field and changed the values of the `price` and `stock` fields:

```
{
  "_id" : 22,
  "item" : "The Banquet",
  "price" : 19,
  "stock" : 3
}
```

Insert a New Document if No Match Exists (Upsert) The following command sets the `upsert` option to `true` so that `update()` (page 69) creates a new document in the `books` collection if no document matches the `<query>` parameter:

```
db.books.update(
  { item: "The New Life" },
  { item: "The New Life", author: "Dante", price: 15 },
  { upsert: true }
)
```

If no document matches the `<query>` parameter, the `upsert` inserts a document with the fields and values of the `<update>` parameter and a new unique `ObjectId` for the `_id` field:

```
{
  "_id" : ObjectId("51e5990c95098ed69d4a89f2"),

```

```

    "author" : "Dante",
    "item" : "The New Life",
    "price" : 15
  }

```

Update Multiple Documents To update multiple documents, set the `multi` option to `true`. For example, the following operation updates all documents where `stock` is less than 5:

```

db.books.update(
  { stock: { $lt: 5 } },
  { $set: { reorder: true } },
  { multi: true }
)

```

Override Default Write Concern The following operation on a replica set specifies a `write` concern of `"w: majority"` with a `wtimeout` of 5000 milliseconds such that the method returns after the write propagates to a majority of the replica set members or the method times out after 5 seconds.

```

db.books.update(
  { stock: { $lt: 5 } },
  { $set: { reorder: true } },
  {
    multi: true,
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)

```

Combine the Upsert and Multi Parameters Given a `books` collection that includes the following documents:

```

{ _id: 11, author: "Dante", item: "Divine Comedy", price: 18, translatedBy: "abc123" }
{ _id: 12, author: "Dante", item: "Divine Comedy", price: 21, translatedBy: "jkl123" }
{ _id: 13, author: "Dante", item: "Divine Comedy", price: 15, translatedBy: "xyz123" }

```

The following command specifies the `multi` parameter to update all documents where `item` is "Divine Comedy" and the `author` is "Dante" and specifies the `upsert` parameter to create a new document if no matching documents are found:

```

db.books.update(
  { item: "Divine Comedy", author: "Dante" },
  { $set: { reorder: false, price: 10 } },
  { upsert: true, multi: true }
)

```

The operation updates all three matching documents and results in the following:

```

{ _id: 11, author: "Dante", item: "Divine Comedy", price: 10, translatedBy: "abc123", reorder: false }
{ _id: 12, author: "Dante", item: "Divine Comedy", price: 10, translatedBy: "jkl123", reorder: false }
{ _id: 13, author: "Dante", item: "Divine Comedy", price: 10, translatedBy: "xyz123", reorder: false }

```

If the collection had *no* matching document, the operation would result in the insertion of a document:

```

{ _id: ObjectId("536aa66422363a21bc16bfd7"), author: "Dante", item: "Divine Comedy", reorder: false,

```

Update Arrays

Update an Element by Position If the update operation requires an update of an element in an array field, the `update()` (page 69) method can perform the update using the position of the element and *dot notation*. Arrays in MongoDB are zero-based.

The following operation queries the `bios` collection for the first document with the `_id` field equal to 1 and updates the second element in the `contribs` array:

```
db.bios.update(
  { _id: 1 },
  { $set: { "contribs.1": "ALGOL 58" } }
)
```

Update an Element if Position is Unknown If the position in the array is not known, the `update()` (page 69) method can perform the update using the `$` positional operator. The array field must appear in the `<query>` parameter in order to determine which array element to update.

The following operation queries the `bios` collection for the first document where the `_id` field equals 3 and the `contribs` array contains an element equal to `compiler`. If found, the `update()` (page 69) method updates the first matching element in the array to `A compiler` in the document:

```
db.bios.update(
  { _id: 3, "contribs": "compiler" },
  { $set: { "contribs.$": "A compiler" } }
)
```

Update a Document Element The `update()` (page 69) method can perform the update of an array that contains subdocuments by using the positional operator (i.e. `$`) and the *dot notation*.

The following operation queries the `bios` collection for the first document where the `_id` field equals 4 and the `awards` array contains a subdocument element with the `by` field equal to `ACM`. If found, the `update()` (page 69) method updates the `by` field in the first matching subdocument:

```
db.bios.update(
  { _id: 4, "awards.by": "ACM" } ,
  { $set: { "awards.$.by": "Association for Computing Machinery" } }
)
```

Add an Element The following operation queries the `bios` collection for the first document that has an `_id` field equal to 1 and adds a new element to the `awards` field:

```
db.bios.update(
  { _id: 1 },
  {
    $push: { awards: { award: "IBM Fellow", year: 1963, by: "IBM" } }
  }
)
```

In the next example, the `$set` (page 416) operator uses *dot notation* to access the middle field in the `name` subdocument. The `$push` (page 425) operator adds another document as an element to the field `awards`.

Consider the following operation:

```
db.bios.update(
  { _id: 1 },
  {
    $set: { "name.middle": "Warner" },
    $push: { awards: {
```

```

        award: "IBM Fellow",
        year: "1963",
        by: "IBM"
      }
    }
  }
)

```

This `update()` (page 69) operation:

- Modifies the field `name` whose value is another document. Specifically, the `$set` (page 416) operator updates the middle field in the `name` subdocument. The document uses *dot notation* to access a field in a subdocument.
- Adds an element to the field `awards`, whose value is an array. Specifically, the `$push` (page 425) operator adds another document as an element to the field `awards`.

WriteResult Changed in version 2.6.

Successful Results The `update()` (page 69) method returns a `WriteResult` (page 188) object that contains the status of the operation. Upon success, the `WriteResult` (page 188) object contains the number of documents that matched the query condition, the number of documents inserted via an `upsert`, and the number of documents modified:

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

See

`WriteResult.nMatched` (page 188), `WriteResult.nUpserted` (page 188), `WriteResult.nModified` (page 188)

Write Concern Errors If the `update()` (page 69) method encounters write concern errors, the results include the `WriteResult.writeConcernError` (page 188) field:

```

WriteResult({
  "nMatched" : 1,
  "nUpserted" : 0,
  "nModified" : 1,
  "writeConcernError" : {
    "code" : 64,
    "errmsg" : "waiting for replication timed out at shard-a"
  }
})

```

See also:

`WriteResult.hasWriteConcernError()` (page 189)

Errors Unrelated to Write Concern If the `update()` (page 69) method encounters a non-write concern error, the results include the `WriteResult.writeError` (page 188) field:

```

WriteResult({
  "nMatched" : 0,
  "nUpserted" : 0,
  "nModified" : 0,
  "writeError" : {

```

```
    "code" : 7,  
    "errmsg" : "could not contact primary for replica set shard-a"  
  }  
})
```

See also:

`WriteResult.hasWriteError()` (page 189)

db.collection.validate()

Description

`db.collection.validate()` (*full*)

Validates a collection. The method scans a collection's data structures for correctness and returns a single *document* that describes the relationship between the logical collection and the physical representation of the data.

The `validate()` (page 76) method has the following parameter:

param Boolean full Specify `true` to enable a full validation and to return full statistics. MongoDB disables full validation by default because it is a potentially resource-intensive operation.

The `validate()` (page 76) method output provides an in-depth view of how the collection uses storage. Be aware that this command is potentially resource intensive and may impact the performance of your MongoDB instance.

The `validate()` (page 76) method is a wrapper around the `validate` (page 335) *database command*.

See also:

validate (page 335)

2.1.2 Cursor

Cursor Methods

Name	Description
<code>cursor.addOption()</code> (page 77)	Adds special wire protocol flags that modify the behavior of the query.'
<code>cursor.batchSize()</code> (page 78)	Controls the number of documents MongoDB will return to the client in a single network message.
<code>cursor.count()</code> (page 79)	Returns a count of the documents in a cursor.
<code>cursor.explain()</code> (page 80)	Reports on the query execution plan, including index use, for a cursor.
<code>cursor.forEach()</code> (page 85)	Applies a JavaScript function for every document in a cursor.
<code>cursor.hasNext()</code> (page 85)	Returns true if the cursor has documents and can be iterated.
<code>cursor.hint()</code> (page 86)	Forces MongoDB to use a specific index for a query.
<code>cursor.limit()</code> (page 86)	Constrains the size of a cursor's result set.
<code>cursor.map()</code> (page 87)	Applies a function to each document in a cursor and collects the return values in an array.
<code>cursor.maxTimeMS()</code> (page 87)	Specifies a cumulative time limit in milliseconds for processing operations on a cursor.
<code>cursor.max()</code> (page 88)	Specifies an exclusive upper index bound for a cursor. For use with <code>cursor.hint()</code> (page 86)
<code>cursor.min()</code> (page 89)	Specifies an inclusive lower index bound for a cursor. For use with <code>cursor.hint()</code> (page 86)
<code>cursor.next()</code> (page 91)	Returns the next document in a cursor.
<code>cursor.objsLeftInBatch()</code> (page 91)	Returns the number of documents left in the current cursor batch.
<code>cursor.readPref()</code> (page 91)	Specifies a <i>read preference</i> to a cursor to control how the client directs queries to a <i>replica set</i> .
<code>cursor.showDiskLocation()</code> (page 91)	Returns a cursor with modified documents that include the on-disk location of the document.
<code>cursor.size()</code> (page 92)	Returns a count of the documents in the cursor after applying <code>skip()</code> (page 92) and <code>limit()</code> (page 86) methods.
<code>cursor.skip()</code> (page 92)	Returns a cursor that begins returning results only after passing or skipping a number of documents.
<code>cursor.snapshot()</code> (page 92)	Forces the cursor to use the index on the <code>_id</code> field. Ensures that the cursor returns each document, with regards to the value of the <code>_id</code> field, only once.
<code>cursor.sort()</code> (page 93)	Returns results ordered according to a sort specification.
<code>cursor.toArray()</code> (page 96)	Returns an array that contains all documents returned by the cursor.

`cursor.addOption()`

Definition

`cursor.addOption(flag)`

Adds OP_QUERY wire protocol flags, such as the `tailable` flag, to change the behavior of queries.

The `cursor.addOption()` (page 77) method has the following parameter:

param flag flag OP_QUERY wire protocol flag. See [MongoDB wire protocol](#)⁶ for more information on MongoDB Wire Protocols and the OP_QUERY flags. For the `mongo` (page 527) shell, you can use *cursor flags* (page 78). For the driver-specific list, see your driver documentation.

Flags The `mongo` (page 527) shell provides several additional cursor flags to modify the behavior of the cursor.

`DBQuery.Option.tailable`

`DBQuery.Option.slaveOk`

`DBQuery.Option.oplogReplay`

`DBQuery.Option.noTimeout`

`DBQuery.Option.awaitData`

`DBQuery.Option.exhaust`

`DBQuery.Option.partial`

For a description of the flags, see [MongoDB wire protocol](#)⁷.

Example The following example adds the `DBQuery.Option.tailable` flag and the `DBQuery.Option.awaitData` flag to ensure that the query returns a tailable cursor. The sequence creates a cursor that will wait for few seconds after returning the full result set so that it can capture and return additional data added during the query:

```
var t = db.myCappedCollection;
var cursor = t.find().addOption(DBQuery.Option.tailable).
                    addOption(DBQuery.Option.awaitData)
```

Warning: Adding incorrect wire protocol flags can cause problems and/or extra server load.

`cursor.batchSize()`

Definition

`cursor.batchSize(size)`

Specifies the number of documents to return in each batch of the response from the MongoDB instance. In most cases, modifying the batch size will not affect the user or the application, as the `mongo` (page 527) shell and most drivers return results as if MongoDB returned a single batch.

The `batchSize()` (page 78) method takes the following parameter:

param integer size The number of documents to return per batch. Do **not** use a batch size of 1.

Note: Specifying 1 or a negative number is analogous to using the `limit()` (page 86) method.

⁶<http://docs.mongodb.org/meta-driver/latest/legacy/mongodb-wire-protocol/?pageVersion=106#op-query>

⁷<http://docs.mongodb.org/meta-driver/latest/legacy/mongodb-wire-protocol/?pageVersion=106#op-query>

Example The following example sets the batch size for the results of a query (i.e. `find()` (page 34)) to 10. The `batchSize()` (page 78) method does not change the output in the `mongo` (page 527) shell, which, by default, iterates over the first 20 documents.

```
db.inventory.find().batchSize(10)
```

cursor.count()

Definition

`cursor.count()`

Counts the number of documents referenced by a cursor. Append the `count()` (page 79) method to a `find()` (page 34) query to return the number of matching documents. The operation does not perform the query but instead counts the results that would be returned by the query.

Changed in version 2.6: MongoDB supports the use of `hint()` (page 86) with `count()` (page 79). See *Specify the Index to Use* (page 80) for an example.

The `count()` (page 79) method has the following prototype form:

```
db.collection.find().count()
```

The `count()` (page 79) method has the following parameter:

param Boolean applySkipLimit Specifies whether to consider the effects of the `cursor.skip()` (page 92) and `cursor.limit()` (page 86) methods in the count. By default, the `count()` (page 79) method ignores the effects of the `cursor.skip()` (page 92) and `cursor.limit()` (page 86). Set `applySkipLimit` to `true` to consider the effect of these methods.

MongoDB also provides the shell wrapper `db.collection.count()` (page 25) for the `db.collection.find().count()` construct.

See also:

`cursor.size()` (page 92)

Behavior On a sharded cluster, `count()` (page 79) method can result in an *inaccurate* count if *orphaned documents* exist or if a chunk migration is in progress.

To avoid these situations, on a sharded cluster, use the `$group` (page 447) stage of the `db.collection.aggregate()` (page 22) method to `$sum` (page 460) the documents. For example, the following operation counts the documents in a collection:

```
db.collection.aggregate([
  { $group: { _id: null, count: { $sum: 1 } } }
])
```

To get a count of documents that match a query condition, include the `$match` (page 440) stage as well:

```
db.collection.aggregate([
  { $match: <query condition> },
  { $group: { _id: null, count: { $sum: 1 } } }
])
```

See *Perform a Count* (page 440) for an example.

Examples The following are examples of the `count()` (page 79) method.

Count All Documents The following operation counts the number of all documents in the `orders` collection:

```
db.orders.find().count()
```

Count Documents That Match a Query The following operation counts the number of the documents in the `orders` collection with the field `ord_dt` greater than `new Date('01/01/2012')`:

```
db.orders.find( { ord_dt: { $gt: new Date('01/01/2012') } } ).count()
```

Limit Documents in Count The following operation counts the number of the documents in the `orders` collection with the field `ord_dt` greater than `new Date('01/01/2012')` *taking into account* the effect of the `limit(5)`:

```
db.orders.find( { ord_dt: { $gt: new Date('01/01/2012') } } ).limit(5).count(true)
```

Specify the Index to Use The following operation uses the index `{ status: 1 }` to return a count of the documents in the `orders` collection with the field `ord_dt` greater than `new Date('01/01/2012')` and the `status` field is equal to `"D"`:

```
db.orders.find(
  { ord_dt: { $gt: new Date('01/01/2012') }, status: "D" }
).hint( { status: 1 } ).count()
```

`cursor.explain()`

Definition

`cursor.explain(verbose)`

Provides information on the query plan. The query plan is the plan the server uses to find the matches for a query. This information may be useful when optimizing a query. The `explain()` (page 80) method returns a document that describes the process used to return the query results.

The `explain()` (page 80) method has the following form:

```
db.collection.find().explain()
```

The `explain()` (page 80) method has the following parameter:

param Boolean verbose Specifies the level of detail to include in the output. If `true` or `1`, includes the `allPlans` and `oldPlan` fields in the output.

For an explanation of output, see *Explain on Queries on Sharded Collections* (page 82) and *Core Explain Output Fields* (page 83).

Behavior The `explain()` (page 80) method runs the actual query to determine the result. Although there are some differences between running the query with `explain()` (page 80) and running without, generally, the performance will be similar between the two. So, if the query is slow, the `explain()` (page 80) operation is also slow.

Additionally, the `explain()` (page 80) operation reevaluates a set of candidate query plans, which may cause the `explain()` (page 80) operation to perform differently than a normal query. As a result, these operations generally provide an accurate account of *how* MongoDB would perform the query, but do not reflect the length of these queries.

See also:

- `$explain` (page 478)
- <http://docs.mongodb.org/manualadministration/optimization> page for information regarding optimization strategies.
- <http://docs.mongodb.org/manualtutorial/manage-the-database-profiler> tutorial for information regarding the database profile.
- *Current Operation Reporting* (page 103)

Explain Results

Explain on Queries on Unsharded Collections For queries on unsharded collections, `explain()` (page 80) returns the following core information.

```
{
  "cursor" : "<Cursor Type and Index>",
  "isMultiKey" : <boolean>,
  "n" : <num>,
  "nscannedObjects" : <num>,
  "nscanned" : <num>,
  "nscannedObjectsAllPlans" : <num>,
  "nscannedAllPlans" : <num>,
  "scanAndOrder" : <boolean>,
  "indexOnly" : <boolean>,
  "nYields" : <num>,
  "nChunkSkips" : <num>,
  "millis" : <num>,
  "indexBounds" : { <index bounds> },
  "allPlans" : [
    {
      "cursor" : "<Cursor Type and Index>",
      "n" : <num>,
      "nscannedObjects" : <num>,
      "nscanned" : <num>,
      "indexBounds" : { <index bounds> }
    },
    ...
  ],
  "oldPlan" : {
    "cursor" : "<Cursor Type and Index>",
    "indexBounds" : { <index bounds> }
  },
  "server" : "<host:port>",
  "filterSet" : <boolean>
}
```

For details on the fields, see *Core Explain Output Fields* (page 83).

Explain on \$or Queries Queries with `$or` (page 377) operator execute each clause of the `$or` (page 377) expression in parallel and can use separate indexes on the individual clauses. If the query uses indexes on any or all of the query's clause, `explain()` (page 80) contains *output* (page 83) for each clause as well as the cumulative data for the entire query:

```
{
  "clauses" : [
    {
      <core explain output>
    }
  ]
}
```

```
        },
        {
            <core explain output>
        },
        ...
    ],
    "n" : <num>,
    "nscannedObjects" : <num>,
    "nscanned" : <num>,
    "nscannedObjectsAllPlans" : <num>,
    "nscannedAllPlans" : <num>,
    "millis" : <num>,
    "server" : "<host:port>"
}
```

For details on the fields, see *\$or Query Output Fields* (page 84) and *Core Explain Output Fields* (page 83).

Explain on Queries on Sharded Collections For queries on sharded collections, `explain()` (page 80) returns information for each shard the query accesses. For queries on unsharded collections, see *Core Explain Output Fields* (page 83).

For queries on a sharded collection, the output contains the *Core Explain Output Fields* (page 83) for each accessed shard and *cumulative shard information* (page 85):

```
{
  "clusteredType" : "<Shard Access Type>",
  "shards" : {
    "<shard1>" : [
      {
        <core explain output>
      }
    ],
    "<shard2>" : [
      {
        <core explain output>
      }
    ],
    ...
  },
  "millisShardTotal" : <num>,
  "millisShardAvg" : <num>,
  "numQueries" : <num>,
  "numShards" : <num>,
  "cursor" : "<Cursor Type and Index>",
  "n" : <num>,
  "nChunkSkips" : <num>,
  "nYields" : <num>,
  "nscanned" : <num>,
  "nscannedAllPlans" : <num>,
  "nscannedObjects" : <num>,
  "nscannedObjectsAllPlans" : <num>,
  "millis" : <num>
}
```

For details on these fields, see *Core Explain Output Fields* (page 83) for each accessed shard and *Sharded Collections Output Fields* (page 85).

Explain Output Fields

Core Explain Output Fields This section explains output for queries on collections that are *not sharded*. For queries on sharded collections, see [Explain on Queries on Sharded Collections](#) (page 82).

`explain.cursor`

`cursor` (page 83) is a string that reports the type of cursor used by the query operation:

- `BasicCursor` indicates a full collection scan.
- `BtreeCursor` indicates that the query used an index. The cursor includes name of the index. When a query uses an index, the output of `explain()` (page 80) includes `indexBounds` (page 84) details.
- `GeoSearchCursor` indicates that the query used a geospatial index.
- `Complex Plan` indicates that MongoDB used `index intersection`.

For `BtreeCursor` cursors, MongoDB will append the name of the index to the cursor string. Additionally, depending on how the query uses an index, MongoDB may append one or both of the following strings to the cursor string:

- `reverse` indicates that query transverses the index from the highest values to the lowest values (e.g. “right to left”).
- `multi` indicates that the query performed multiple look-ups. Otherwise, the query uses the index to determine a range of possible matches.

`explain.isMultiKey`

`isMultiKey` (page 83) is a boolean. When `true`, the query uses a *multikey index*, where one of the fields in the index holds an array.

`explain.n`

`n` (page 83) is a number that reflects the number of documents that match the query selection criteria.

`explain.nscannedObjects`

Specifies the total number of documents scanned during the query. The `nscannedObjects` (page 83) may be lower than `nscanned` (page 83), such as if the index *covers* a query. See `indexOnly` (page 84). Additionally, the `nscannedObjects` (page 83) may be lower than `nscanned` (page 83) in the case of multikey index on an array field with duplicate documents.

`explain.nscanned`

Specifies the total number of documents or index entries scanned during the database operation. You want `n` (page 83) and `nscanned` (page 83) to be close in value as possible. The `nscanned` (page 83) value may be higher than the `nscannedObjects` (page 83) value, such as if the index *covers* a query. See `indexOnly` (page 84).

`explain.nscannedObjectsAllPlans`

New in version 2.2.

`nscannedObjectsAllPlans` (page 83) is a number that reflects the total number of documents scanned for all query plans during the database operation.

`explain.nscannedAllPlans`

New in version 2.2.

`nscannedAllPlans` (page 83) is a number that reflects the total number of documents or index entries scanned for all query plans during the database operation.

`explain.scanAndOrder`

`scanAndOrder` (page 83) is a boolean that is `true` when the query **cannot** use the order of documents in the index for returning sorted results: MongoDB must sort the documents after it receives the documents from a cursor.

If `scanAndOrder` (page 83) is `false`, MongoDB *can* use the order of the documents in an index to return sorted results.

`explain.indexOnly`

`indexOnly` (page 84) is a boolean value that returns `true` when the query is *covered* by the index indicated in the `cursor` (page 83) field. When an index covers a query, MongoDB can both match the *query conditions* **and** return the results using only the index because:

- all the fields in the *query* are part of that index, **and**
- all the fields returned in the results set are in the same index.

`explain.nYields`

`nYields` (page 84) is a number that reflects the number of times this query yielded the read lock to allow waiting writes to execute.

`explain.nChunkSkips`

`nChunkSkips` (page 84) is a number that reflects the number of documents skipped because of active chunk migrations in a sharded system. Typically this will be zero. A number greater than zero is ok, but indicates a little bit of inefficiency.

`explain.millis`

`millis` (page 84) is a number that reflects the time in milliseconds to complete the query.

`explain.indexBounds`

`indexBounds` (page 84) is a document that contains the lower and upper index key bounds. This field resembles one of the following:

```
"indexBounds" : {
  "start" : { <index key1> : <value>, ... },
  "end"   : { <index key1> : <value>, ... }
},

"indexBounds" : { "<field>" : [ [ <lower bound>, <upper bound> ] ],
  ...
}
```

`explain.allPlans`

`allPlans` (page 84) is an array that holds the list of plans the query optimizer runs in order to select the index for the query. Displays only when the `<verbose>` parameter to `explain()` (page 80) is `true` or `1`.

`explain.oldPlan`

New in version 2.2.

`oldPlan` (page 84) is a document value that contains the previous plan selected by the query optimizer for the query. Displays only when the `<verbose>` parameter to `explain()` (page 80) is `true` or `1`.

`explain.server`

New in version 2.2.

`server` (page 84) is a string that reports the MongoDB server.

`explain.filterSet`

New in version 2.6.

`filterSet` (page 84) is a boolean that indicates whether MongoDB applied an *index filter* for the query.

\$or Query Output Fields

`explain.clauses`

`clauses` (page 84) is an array that holds the *Core Explain Output Fields* (page 83) information for each clause

of the `$or` (page 377) expression. `clauses` (page 84) is only included when the clauses in the `$or` (page 377) expression use indexes.

Sharded Collections Output Fields

`explain.clusteredType`

`clusteredType` (page 85) is a string that reports the access pattern for shards. The value is:

- `ParallelSort`, if the `mongos` (page 518) queries shards in parallel.
- `SerialServer`, if the `mongos` (page 518) queries shards sequentially.

`explain.shards`

`shards` (page 85) contains fields for each shard in the cluster accessed during the query. Each field holds the *Core Explain Output Fields* (page 83) for that shard.

`explain.millisShardTotal`

`millisShardTotal` (page 85) is a number that reports the total time in milliseconds for the query to run on the shards.

`explain.millisShardAvg`

`millisShardAvg` (page 85) is a number that reports the average time in millisecond for the query to run on each shard.

`explain.numQueries`

`numQueries` (page 85) is a number that reports the total number of queries executed.

`explain.numShards`

`numShards` (page 85) is a number that reports the total number of shards queried.

`cursor.forEach()`

Description

`cursor.forEach()` (function)

Iterates the cursor to apply a JavaScript function to each document from the cursor.

The `forEach()` (page 85) method has the following prototype form:

```
db.collection.find().forEach(<function>)
```

The `forEach()` (page 85) method has the following parameter:

param JavaScript function A JavaScript function to apply to each document from the cursor. The `<function>` signature includes a single argument that is passed the current document to process.

Example The following example invokes the `forEach()` (page 85) method on the cursor returned by `find()` (page 34) to print the name of each user in the collection:

```
db.users.find().forEach( function(myDoc) { print( "user: " + myDoc.name ); } );
```

See also:

`cursor.map()` (page 87) for similar functionality.

`cursor.hasNext()`

```
cursor.hasNext()
```

Returns Boolean.

`cursor.hasNext()` (page 85) returns `true` if the cursor returned by the `db.collection.find()` (page 34) query can iterate further to return more documents.

`cursor.hint()`

Definition

`cursor.hint(index)`

Call this method on a query to override MongoDB's default index selection and query optimization process. Use `db.collection.getIndexes()` (page 45) to return the list of current indexes on a collection.

The `cursor.hint()` (page 86) method has the following parameter:

param string,document index The index to “hint” or force MongoDB to use when performing the query. Specify the index either by the index name or by the index specification document.

Behavior When an *index filter* exists for the query shape, MongoDB ignores the `hint()` (page 86). The `explain.filterSet` (page 84) field of the `explain()` (page 80) output indicates whether MongoDB applied an index filter for the query.

You cannot use `hint()` (page 86) if the query includes a `$text` (page 387) query expression.

Example The following example returns all documents in the collection named `users` using the index on the `age` field.

```
db.users.find().hint( { age: 1 } )
```

You can also specify the index using the index name:

```
db.users.find().hint( "age_1" )
```

See also:

- <http://docs.mongodb.org/manualcore/indexes-introduction>
- <http://docs.mongodb.org/manualadministration/indexes>
- <http://docs.mongodb.org/manualcore/query-plans>
- *index-filters*
- `$hint` (page 479)

`cursor.limit()`

Definition

`cursor.limit()`

Use the `limit()` (page 86) method on a cursor to specify the maximum number of documents the cursor will return. `limit()` (page 86) is analogous to the `LIMIT` statement in a SQL database.

Note: You must apply `limit()` (page 86) to the cursor before retrieving any documents from the database.

Use `limit()` (page 86) to maximize performance and prevent MongoDB from returning more results than required for processing.

Behavior

Supported Values The behavior of `limit()` (page 86) is undefined for values less than -2^{31} and greater than 2^{31} .

Negative Values A `limit()` (page 86) value of 0 (i.e. `.limit(0)` (page 86)) is equivalent to setting no limit. A negative limit is similar to a positive limit, but a negative limit prevents the creation of a cursor such that only a single batch of results is returned. As such, with a negative limit, if the limited result set does not fit into a single batch, the number of documents received will be less than the limit.

`cursor.map()`

`cursor.map(function)`

Applies function to each document visited by the cursor and collects the return values from successive application into an array.

The `cursor.map()` (page 87) method has the following parameter:

param function function A function to apply to each document visited by the cursor.

Example

```
db.users.find().map( function(u) { return u.name; } );
```

See also:

`cursor.forEach()` (page 85) for similar functionality.

`cursor.maxTimeMS()`

Definition New in version 2.6.

`cursor.maxTimeMS(<time limit>)`

Specifies a cumulative time limit in milliseconds for processing operations on a cursor.

The `maxTimeMS()` (page 87) method has the following parameter:

param integer milliseconds Specifies a cumulative time limit in milliseconds for processing operations on the cursor.

Important: `maxTimeMS()` (page 87) is not related to the `NoCursorTimeout` query flag. `maxTimeMS()` (page 87) relates to processing time, while `NoCursorTimeout` relates to idle time. A cursor's idle time does not contribute towards its processing time.

Behaviors MongoDB targets operations for termination if the associated cursor exceeds its allotted time limit. MongoDB terminates operations that exceed their allotted time limit, using the same mechanism as `db.killOp()` (page 114). MongoDB only terminates an operation at one of its designated interrupt points.

MongoDB does not count network latency towards a cursor's time limit.

Queries that generate multiple batches of results continue to return batches until the cursor exceeds its allotted time limit.

Examples

Example

The following query specifies a time limit of 50 milliseconds:

```
db.collection.find({description: /August [0-9]+, 1969/}).maxTimeMS(50)
```

`cursor.max()`

Definition

`cursor.max()`

Specifies the *exclusive* upper bound for a specific index in order to constrain the results of `find()` (page 34). `max()` (page 88) provides a way to specify an upper bound on compound key indexes.

The `max()` (page 88) method has the following parameter:

param document indexBounds The exclusive upper bound for the index keys.

The `indexBounds` parameter has the following prototype form:

```
{ field1: <max value>, field2: <max value2> ... fieldN:<max valueN> }
```

The fields correspond to *all* the keys of a particular index *in order*. You can explicitly specify the particular index with the `hint()` (page 86) method. Otherwise, `mongodb` (page 503) selects the index using the fields in the `indexBounds`; however, if multiple indexes exist on same fields with different sort orders, the selection of the index may be ambiguous.

See also:

`min()` (page 89).

Note: `max()` (page 88) is a shell wrapper around the query modifier `$max` (page 480).

Behavior

- Because `max()` (page 88) requires an index on a field, and forces the query to use this index, you may prefer the `$lt` (page 375) operator for the query if possible. Consider the following example:

```
db.products.find( { _id: 7 } ).max( { price: 1.39 } )
```

The query will use the index on the `price` field, even if the index on `_id` may be better.

- `max()` (page 88) exists primarily to support the `mongos` (page 518) (sharding) process.
- If you use `max()` (page 88) with `min()` (page 89) to specify a range, the index bounds specified in `min()` (page 89) and `max()` (page 88) must both refer to the keys of the same index.

Example This example assumes a collection named `products` that holds the following documents:

```
{ "_id" : 6, "item" : "apple", "type" : "cortland", "price" : 1.29 }
{ "_id" : 2, "item" : "apple", "type" : "fuji", "price" : 1.99 }
{ "_id" : 1, "item" : "apple", "type" : "honey crisp", "price" : 1.99 }
{ "_id" : 3, "item" : "apple", "type" : "jonagold", "price" : 1.29 }
{ "_id" : 4, "item" : "apple", "type" : "jonathan", "price" : 1.29 }
{ "_id" : 5, "item" : "apple", "type" : "mcintosh", "price" : 1.29 }
{ "_id" : 7, "item" : "orange", "type" : "cara cara", "price" : 2.99 }
{ "_id" : 10, "item" : "orange", "type" : "navel", "price" : 1.39 }
```

```
{ "_id" : 9, "item" : "orange", "type" : "satsuma", "price" : 1.99 }
{ "_id" : 8, "item" : "orange", "type" : "valencia", "price" : 0.99 }
```

The collection has the following indexes:

```
{ "_id" : 1 }
{ "item" : 1, "type" : 1 }
{ "item" : 1, "type" : -1 }
{ "price" : 1 }
```

- Using the ordering of { item: 1, type: 1 } index, `max()` (page 88) limits the query to the documents that are below the bound of item equal to apple and type equal to jonagold:

```
db.products.find().max( { item: 'apple', type: 'jonagold' } ).hint( { item: 1, type: 1 } )
```

The query returns the following documents:

```
{ "_id" : 6, "item" : "apple", "type" : "cortland", "price" : 1.29 }
{ "_id" : 2, "item" : "apple", "type" : "fuji", "price" : 1.99 }
{ "_id" : 1, "item" : "apple", "type" : "honey crisp", "price" : 1.99 }
```

If the query did not explicitly specify the index with the `hint()` (page 86) method, it is ambiguous as to whether `mongod` (page 503) would select the { item: 1, type: 1 } index ordering or the { item: 1, type: -1 } index ordering.

- Using the ordering of the index { price: 1 }, `max()` (page 88) limits the query to the documents that are below the index key bound of price equal to 1.99 and `min()` (page 89) limits the query to the documents that are at or above the index key bound of price equal to 1.39:

```
db.products.find().min( { price: 1.39 } ).max( { price: 1.99 } ).hint( { price: 1 } )
```

The query returns the following documents:

```
{ "_id" : 6, "item" : "apple", "type" : "cortland", "price" : 1.29 }
{ "_id" : 4, "item" : "apple", "type" : "jonathan", "price" : 1.29 }
{ "_id" : 5, "item" : "apple", "type" : "mcintosh", "price" : 1.29 }
{ "_id" : 3, "item" : "apple", "type" : "jonagold", "price" : 1.29 }
{ "_id" : 10, "item" : "orange", "type" : "navel", "price" : 1.39 }
```

cursor.min()

Definition

`cursor.min()`

Specifies the *inclusive* lower bound for a specific index in order to constrain the results of `find()` (page 34). `min()` (page 89) provides a way to specify lower bounds on compound key indexes.

The `min()` (page 89) has the following parameter:

param document indexBounds The inclusive lower bound for the index keys.

The `indexBounds` parameter has the following prototype form:

```
{ field1: <min value>, field2: <min value2>, fieldN:<min valueN> }
```

The fields correspond to *all* the keys of a particular index *in order*. You can explicitly specify the particular index with the `hint()` (page 86) method. Otherwise, MongoDB selects the index using the fields in the `indexBounds`; however, if multiple indexes exist on same fields with different sort orders, the selection of the index may be ambiguous.

See also:

`max()` (page 88).

Note: `min()` (page 89) is a shell wrapper around the query modifier `$min` (page 481).

Behaviors

- Because `min()` (page 89) requires an index on a field, and forces the query to use this index, you may prefer the `$gte` (page 374) operator for the query if possible. Consider the following example:

```
db.products.find( { _id: 7 } ).min( { price: 1.39 } )
```

The query will use the index on the `price` field, even if the index on `_id` may be better.

- `min()` (page 89) exists primarily to support the `mongos` (page 518) process.
- If you use `min()` (page 89) with `max()` (page 88) to specify a range, the index bounds specified in `min()` (page 89) and `max()` (page 88) must both refer to the keys of the same index.

Example This example assumes a collection named `products` that holds the following documents:

```
{ "_id" : 6, "item" : "apple", "type" : "cortland", "price" : 1.29 }
{ "_id" : 2, "item" : "apple", "type" : "fuji", "price" : 1.99 }
{ "_id" : 1, "item" : "apple", "type" : "honey crisp", "price" : 1.99 }
{ "_id" : 3, "item" : "apple", "type" : "jonagold", "price" : 1.29 }
{ "_id" : 4, "item" : "apple", "type" : "jonathan", "price" : 1.29 }
{ "_id" : 5, "item" : "apple", "type" : "mcintosh", "price" : 1.29 }
{ "_id" : 7, "item" : "orange", "type" : "cara cara", "price" : 2.99 }
{ "_id" : 10, "item" : "orange", "type" : "navel", "price" : 1.39 }
{ "_id" : 9, "item" : "orange", "type" : "satsuma", "price" : 1.99 }
{ "_id" : 8, "item" : "orange", "type" : "valencia", "price" : 0.99 }
```

The collection has the following indexes:

```
{ "_id" : 1 }
{ "item" : 1, "type" : 1 }
{ "item" : 1, "type" : -1 }
{ "price" : 1 }
```

- Using the ordering of the `{ item: 1, type: 1 }` index, `min()` (page 89) limits the query to the documents that are at or above the index key bound of `item` equal to `apple` and `type` equal to `jonagold`, as in the following:

```
db.products.find().min( { item: 'apple', type: 'jonagold' } ).hint( { item: 1, type: 1 } )
```

The query returns the following documents:

```
{ "_id" : 3, "item" : "apple", "type" : "jonagold", "price" : 1.29 }
{ "_id" : 4, "item" : "apple", "type" : "jonathan", "price" : 1.29 }
{ "_id" : 5, "item" : "apple", "type" : "mcintosh", "price" : 1.29 }
{ "_id" : 7, "item" : "orange", "type" : "cara cara", "price" : 2.99 }
{ "_id" : 10, "item" : "orange", "type" : "navel", "price" : 1.39 }
{ "_id" : 9, "item" : "orange", "type" : "satsuma", "price" : 1.99 }
{ "_id" : 8, "item" : "orange", "type" : "valencia", "price" : 0.99 }
```

If the query did not explicitly specify the index with the `hint()` (page 86) method, it is ambiguous as to whether `mongod` (page 503) would select the `{ item: 1, type: 1 }` index ordering or the `{ item: 1, type: -1 }` index ordering.

- Using the ordering of the index `{ price: 1 }`, `min()` (page 89) limits the query to the documents that are at or above the index key bound of `price` equal to 1.39 and `max()` (page 88) limits the query to the documents that are below the index key bound of `price` equal to 1.99:

```
db.products.find().min( { price: 1.39 } ).max( { price: 1.99 } ).hint( { price: 1 } )
```

The query returns the following documents:

```
{ "_id" : 6, "item" : "apple", "type" : "cortland", "price" : 1.29 }
{ "_id" : 4, "item" : "apple", "type" : "jonathan", "price" : 1.29 }
{ "_id" : 5, "item" : "apple", "type" : "mcintosh", "price" : 1.29 }
{ "_id" : 3, "item" : "apple", "type" : "jonagold", "price" : 1.29 }
{ "_id" : 10, "item" : "orange", "type" : "navel", "price" : 1.39 }
```

`cursor.next()`

```
cursor.next()
```

Returns The next document in the cursor returned by the `db.collection.find()` (page 34) method. See `cursor.hasNext()` (page 85) related functionality.

`cursor.objsLeftInBatch()`

```
cursor.objsLeftInBatch()
```

`cursor.objsLeftInBatch()` (page 91) returns the number of documents remaining in the current batch.

The MongoDB instance returns response in batches. To retrieve all the documents from a cursor may require multiple batch responses from the MongoDB instance. When there are no more documents remaining in the current batch, the cursor will retrieve another batch to get more documents until the cursor exhausts.

`cursor.readPref()`

Definition

```
cursor.readPref(mode, tagSet)
```

Append `readPref()` (page 91) to a cursor to control how the client routes the query to members of the replica set.

param string mode One of the following *read preference* modes: `primary`, `primaryPreferred`, `secondary`, `secondaryPreferred`, or `nearest`

param array tagSet A tag set used to specify custom read preference modes. For details, see *replica-set-read-preference-tag-sets*.

Note: You must apply `readPref()` (page 91) to the cursor before retrieving any documents from the database.

`cursor.showDiskLoc()`

```
cursor.showDiskLoc()
```

Returns A modified cursor object that contains documents with appended information that describes the on-disk location of the document.

See also:

`$showDiskLoc` (page 482) for related functionality.

`cursor.size()`

`cursor.size()`

Returns A count of the number of documents that match the `db.collection.find()` (page 34) query after applying any `cursor.skip()` (page 92) and `cursor.limit()` (page 86) methods.

`cursor.skip()`

`cursor.skip()`

Call the `cursor.skip()` (page 92) method on a cursor to control where MongoDB begins returning results. This approach may be useful in implementing “paged” results.

Note: You must apply `cursor.skip()` (page 92) to the cursor before retrieving any documents from the database.

Consider the following JavaScript function as an example of the skip function:

```
function printStudents(pageNumber, nPerPage) {
  print("Page: " + pageNumber);
  db.students.find().skip(pageNumber > 0 ? ((pageNumber-1)*nPerPage) : 0).limit(nPerPage).forEach(
}
```

The `cursor.skip()` (page 92) method is often expensive because it requires the server to walk from the beginning of the collection or index to get the offset or skip position before beginning to return result. As offset (e.g. `pageNumber` above) increases, `cursor.skip()` (page 92) will become slower and more CPU intensive. With larger collections, `cursor.skip()` (page 92) may become IO bound.

Consider using range-based pagination for these kinds of tasks. That is, query for a range of objects, using logic within the application to determine the pagination rather than the database itself. This approach features better index utilization, if you do not need to easily jump to a specific page.

`cursor.snapshot()`

`cursor.snapshot()`

Append the `snapshot()` (page 92) method to a cursor to toggle the “snapshot” mode. This ensures that the query will not return a document multiple times, even if intervening write operations result in a move of the document due to the growth in document size.

Warning:

- You must apply `snapshot()` (page 92) to the cursor before retrieving any documents from the database.
- You can only use `snapshot()` (page 92) with **unsharded** collections.

The `snapshot()` (page 92) does not guarantee isolation from insertion or deletions.

The `snapshot()` (page 92) traverses the index on the `_id` field. As such, `snapshot()` (page 92) **cannot** be used with `sort()` (page 93) or `hint()` (page 86).

Queries with results of less than 1 megabyte are effectively implicitly snapshotted.

`cursor.sort()`

Definition

`cursor.sort (sort)`

Specifies the order in which the query returns matching documents. You must apply `sort()` (page 93) to the cursor before retrieving any documents from the database.

The `sort()` (page 93) method has the following parameter:

param document sort A document that defines the sort order of the result set.

The `sort` parameter contains field and value pairs, in the following form:

```
{ field: value }
```

The sort document can specify *ascending or descending sort on existing fields* (page 93) or *sort on computed metadata* (page 94).

Behaviors

Ascending/Descending Sort Specify in the sort parameter the field or fields to sort by and a value of 1 or -1 to specify an ascending or descending sort respectively.

The following sample document specifies a descending sort by the `age` field and then an ascending sort by the `posts` field:

```
{ age : -1, posts: 1 }
```

When comparing values of different *BSON* types, MongoDB uses the following comparison order, from lowest to highest:

1. MinKey (internal type)
2. Null
3. Numbers (ints, longs, doubles)
4. Symbol, String
5. Object
6. Array
7. BinData
8. ObjectId
9. Boolean
10. Date, Timestamp
11. Regular Expression
12. MaxKey (internal type)

MongoDB treats some types as equivalent for comparison purposes. For instance, numeric types undergo conversion before comparison.

The comparison treats a non-existent field as it would an empty BSON Object. As such, a sort on the `a` field in documents `{ }` and `{ a: null }` would treat the documents as equivalent in sort order.

With arrays, a less-than comparison or an ascending sort compares the smallest element of arrays, and a greater-than comparison or a descending sort compares the largest element of the arrays. As such, when comparing a field whose value is a single-element array (e.g. [1]) with non-array fields (e.g. 2), the comparison is between 1 and 2. A comparison of an empty array (e.g. []) treats the empty array as less than `null` or a missing field.

Metadata Sort Specify in the sort parameter a new field name for the computed metadata and specify the `$meta` (page 410) expression as its value.

The following sample document specifies a descending sort by the "textScore" metadata:

```
{ score: { $meta: "textScore" } }
```

The specified metadata determines the sort order. For example, the "textScore" metadata sorts in descending order. See `$meta` (page 410) for details.

Limit Results The sort operation requires that the entire sort be able to complete within 32 megabytes.

When the sort operation consumes more than 32 megabytes, MongoDB returns an error. To avoid this error, either create an index to support the sort operation or use `sort()` (page 93) in conjunction with `limit()` (page 86). The specified limit must result in a number of documents that fall within the 32 megabyte limit.

For example, if the following sort operation `stocks_quotes` exceeds the 32 megabyte limit:

```
db.stocks.find().sort( { ticker: 1, date: -1 } )
```

Either create an index to support the sort operation:

```
db.stocks.ensureIndex( { ticker: 1, date: -1 } )
```

Or use `sort()` (page 93) in conjunction with `limit()` (page 86):

```
db.stocks.find().sort( { ticker: 1, date: -1 } ).limit(100)
```

Examples A collection `orders` contain the following documents:

```
{ _id: 1, item: { category: "cake", type: "chiffon" }, amount: 10 }
{ _id: 2, item: { category: "cookies", type: "chocolate chip" }, amount: 50 }
{ _id: 3, item: { category: "cookies", type: "chocolate chip" }, amount: 15 }
{ _id: 4, item: { category: "cake", type: "lemon" }, amount: 30 }
{ _id: 5, item: { category: "cake", type: "carrot" }, amount: 20 }
{ _id: 6, item: { category: "brownies", type: "blondie" }, amount: 10 }
```

The following query, which returns all documents from the `orders` collection, does not specify a sort order:

```
db.orders.find()
```

The query returns the documents in indeterminate order:

```
{ "_id" : 1, "item" : { "category" : "cake", "type" : "chiffon" }, "amount" : 10 }
{ "_id" : 2, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 50 }
{ "_id" : 3, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 15 }
{ "_id" : 4, "item" : { "category" : "cake", "type" : "lemon" }, "amount" : 30 }
{ "_id" : 5, "item" : { "category" : "cake", "type" : "carrot" }, "amount" : 20 }
{ "_id" : 6, "item" : { "category" : "brownies", "type" : "blondie" }, "amount" : 10 }
```

The following query specifies a sort on the `amount` field in descending order.

```
db.orders.find().sort( { amount: -1 } )
```

The query returns the following documents, in descending order of amount:

```
{ "_id" : 2, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 50 }
{ "_id" : 4, "item" : { "category" : "cake", "type" : "lemon" }, "amount" : 30 }
{ "_id" : 5, "item" : { "category" : "cake", "type" : "carrot" }, "amount" : 20 }
{ "_id" : 3, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 15 }
{ "_id" : 1, "item" : { "category" : "cake", "type" : "chiffon" }, "amount" : 10 }
{ "_id" : 6, "item" : { "category" : "brownies", "type" : "blondie" }, "amount" : 10 }
```

The following query specifies the sort order using the fields from a sub-document `item`. The query sorts first by the `category` field in ascending order, and then within each category, by the `type` field in ascending order.

```
db.orders.find().sort( { "item.category": 1, "item.type": 1 } )
```

The query returns the following documents, ordered first by the `category` field, and within each category, by the `type` field:

```
{ "_id" : 6, "item" : { "category" : "brownies", "type" : "blondie" }, "amount" : 10 }
{ "_id" : 5, "item" : { "category" : "cake", "type" : "carrot" }, "amount" : 20 }
{ "_id" : 1, "item" : { "category" : "cake", "type" : "chiffon" }, "amount" : 10 }
{ "_id" : 4, "item" : { "category" : "cake", "type" : "lemon" }, "amount" : 30 }
{ "_id" : 2, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 50 }
{ "_id" : 3, "item" : { "category" : "cookies", "type" : "chocolate chip" }, "amount" : 15 }
```

Return in Storage Order The `$natural` (page 483) parameter returns items according to their storage order within the collection level extents.

Typically, the storage order reflects insertion order, *except* when documents relocate because of *document growth due to updates* or remove operations free up space which are then taken up by newly inserted documents.

Consider the sequence of insert operations to the `trees` collection:

```
db.trees.insert( { _id: 1, common_name: "oak", genus: "quercus" } )
db.trees.insert( { _id: 2, common_name: "chestnut", genus: "castanea" } )
db.trees.insert( { _id: 3, common_name: "maple", genus: "aceraceae" } )
db.trees.insert( { _id: 4, common_name: "birch", genus: "betula" } )
```

The following query returns the documents in the storage order:

```
db.trees.find().sort( { $natural: 1 } )
```

The documents return in the following order:

```
{ "_id" : 1, "common_name" : "oak", "genus" : "quercus" }
{ "_id" : 2, "common_name" : "chestnut", "genus" : "castanea" }
{ "_id" : 3, "common_name" : "maple", "genus" : "aceraceae" }
{ "_id" : 4, "common_name" : "birch", "genus" : "betula" }
```

Update a document such that the document outgrows its current allotted space:

```
db.trees.update(
  { _id: 1 },
  { $set: { famous_oaks: [ "Emancipation Oak", "Goethe Oak" ] } }
)
```

Rerun the query to returns the documents in the storage order:

```
db.trees.find().sort( { $natural: 1 } )
```

The documents return in the following storage order:

```
{ "_id" : 2, "common_name" : "chestnut", "genus" : "castanea" }
{ "_id" : 3, "common_name" : "maple", "genus" : "aceraceae" }
{ "_id" : 4, "common_name" : "birch", "genus" : "betula" }
{ "_id" : 1, "common_name" : "oak", "genus" : "quercus", "famous_oaks" : [ "Emancipation Oak", "Goeth"
```

See also:

`$natural` (page 483)

cursor.toArray()

`cursor.toArray()`

The `toArray()` (page 96) method returns an array that contains all the documents from a cursor. The method iterates completely the cursor, loading all the documents into RAM and exhausting the cursor.

Returns An array of documents.

Consider the following example that applies `toArray()` (page 96) to the cursor returned from the `find()` (page 34) method:

```
var allProductsArray = db.products.find().toArray();

if (allProductsArray.length > 0) { printjson (allProductsArray[0]); }
```

The variable `allProductsArray` holds the array of documents returned by `toArray()` (page 96).

2.1.3 Database

Database Methods

Name	Description
<code>db.addUser()</code> (page 143)	Adds a user to a database, and allows administrators to configure the user's
<code>db.auth()</code> (page 99)	Authenticates a user to a database.
<code>db.changeUserPassword()</code> (page 147)	Changes an existing user's password.
<code>db.cloneCollection()</code> (page 99)	Copies data directly between MongoDB instances. Wraps <code>cloneCollection()</code>
<code>db.cloneDatabase()</code> (page 100)	Copies a database from a remote host to the current host. Wraps <code>clone()</code> (p
<code>db.commandHelp()</code> (page 100)	Returns help information for a <i>database command</i> .
<code>db.copyDatabase()</code> (page 100)	Copies a database to another database on the current host. Wraps <code>copyDatabase()</code>
<code>db.createCollection()</code> (page 102)	Creates a new collection. Commonly used to create a capped collection.
<code>db.currentOp()</code> (page 103)	Reports the current in-progress operations.
<code>db.dropDatabase()</code> (page 108)	Removes the current database.
<code>db.eval()</code> (page 108)	Passes a JavaScript function to the <code>mongod</code> (page 503) instance for server
<code>db.fsyncLock()</code> (page 109)	Flushes writes to disk and locks the database to prevent write operations and
<code>db.fsyncUnlock()</code> (page 110)	Allows writes to continue on a database locked with <code>db.fsyncLock()</code>
<code>db.getCollection()</code> (page 110)	Returns a collection object. Used to access collections with names that are
<code>db.getCollectionNames()</code> (page 110)	Lists all collections in the current database.
<code>db.getLastError()</code> (page 110)	Checks and returns the status of the last operation. Wraps <code>getLastError()</code>
<code>db.getLastErrorMessage()</code> (page 111)	Returns the status document for the last operation. Wraps <code>getLastError()</code>
<code>db.getMongo()</code> (page 111)	Returns the <code>Mongo()</code> (page 193) connection object for the current connection

Table 2.2 – continued from previous page

Name	Description
<code>db.getName()</code> (page 111)	Returns the name of the current database.
<code>db.getPrevError()</code> (page 111)	Returns a status document containing all errors since the last error reset. W
<code>db.getProfilingLevel()</code> (page 111)	Returns the current profiling level for database operations.
<code>db.getProfilingStatus()</code> (page 112)	Returns a document that reflects the current profiling level and the profiling
<code>db.getReplicationInfo()</code> (page 112)	Returns a document with replication statistics.
<code>db.getSiblingDB()</code> (page 113)	Provides access to the specified database.
<code>db.help()</code> (page 113)	Displays descriptions of common <code>db</code> object methods.
<code>db.hostInfo()</code> (page 113)	Returns a document with information about the system MongoDB runs on
<code>db.isMaster()</code> (page 114)	Returns a document that reports the state of the replica set.
<code>db.killOp()</code> (page 114)	Terminates a specified operation.
<code>db.listCommands()</code> (page 115)	Displays a list of common database commands.
<code>db.loadServerScripts()</code> (page 115)	Loads all scripts in the <code>system.js</code> collection for the current database int
<code>db.logout()</code> (page 115)	Ends an authenticated session.
<code>db.printCollectionStats()</code> (page 115)	Prints statistics from every collection. Wraps <code>db.collection.stats</code>
<code>db.printReplicationInfo()</code> (page 116)	Prints a report of the status of the replica set from the perspective of the pr
<code>db.printShardingStatus()</code> (page 116)	Prints a report of the sharding configuration and the chunk ranges.
<code>db.printSlaveReplicationInfo()</code> (page 116)	Prints a report of the status of the replica set from the perspective of the se
<code>db.removeUser()</code> (page 147)	Removes a user from a database.
<code>db.repairDatabase()</code> (page 117)	Runs a repair routine on the current database.
<code>db.resetError()</code> (page 117)	Resets the error message returned by <code>db.getPrevError()</code> (page 111)
<code>db.runCommand()</code> (page 118)	Runs a <i>database command</i> (page 198).
<code>db.serverBuildInfo()</code> (page 118)	Returns a document that displays the compilation parameters for the <code>mong</code>
<code>db.serverStatus()</code> (page 118)	Returns a document that provides an overview of the state of the database p
<code>db.setProfilingLevel()</code> (page 119)	Modifies the current level of database profiling.
<code>db.shutdownServer()</code> (page 119)	Shuts down the current <code>mongod</code> (page 503) or <code>mongos</code> (page 518) proces
<code>db.stats()</code> (page 119)	Returns a document that reports on the state of the current database.
<code>db.version()</code> (page 120)	Returns the version of the <code>mongod</code> (page 503) instance.
<code>db.upgradeCheck()</code> (page 120)	Performs a preliminary check for upgrade preparedness for a specific datab
<code>db.upgradeCheckAllDBs()</code> (page 121)	Performs a preliminary check for upgrade preparedness for all databases a

db.addUser()

Deprecated since version 2.6: Use `db.createUser()` (page 141) and `db.updateUser()` (page 145) instead of `db.addUser()` (page 143) to add users to MongoDB.

In 2.6, MongoDB introduced a new model for user credentials and privileges, as described in <http://docs.mongodb.org/manualcore/security-introduction>. To use `db.addUser()` (page 143) on MongoDB 2.4, see [db.addUser\(\) in the version 2.4 of the MongoDB Manual](#)⁸.

Definition

`db.addUser` (*document*)

Adds a new user on the database where you run the method. The `db.addUser()` (page 143) method takes a user document as its argument:

```
db.addUser(<user document>)
```

Specify a document that resembles the following as an argument to `db.addUser()` (page 143):

```
{ user: "<name>",
  pwd: "<cleartext password>",
```

⁸<http://docs.mongodb.org/v2.4/reference/method/db.addUser>

```
customData: { <any information> },
roles: [
  { role: "<role>", db: "<database>" } | "<role>",
  ...
],
writeConcern: { <write concern> }
}
```

The `db.addUser()` (page 143) user document has the following fields:

field string user The name of the new user.

field string pwd The user's password. The `pwd` field is not required if you run `db.addUser()` (page 143) on the `$external` database to create users who have credentials stored externally to MongoDB.

any document customData Any arbitrary information.

field array roles The roles granted to the user.

field document writeConcern The level of `write concern` for the creation operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Users created on the `$external` database should have credentials stored externally to MongoDB, as, for example, with MongoDB Enterprise installations that use Kerberos.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.addUser()` (page 143) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Considerations The `db.addUser()` (page 143) method returns a *duplicate user* error if the user exists.

When interacting with 2.6 and later MongoDB instances, `db.addUser()` (page 143) sends unencrypted passwords. To encrypt the password in transit use SSL.

In the 2.6 version of the shell, `db.addUser()` (page 143) is backwards compatible with both the 2.4 version of `db.addUser()`⁹ and the 2.2 version of `db.addUser()`¹⁰. In 2.6, for backwards compatibility `db.addUser()` (page 143) creates users that approximate the privileges available in previous versions of MongoDB.

Example The following `db.addUser()` (page 143) method creates a user Carlos on the database where the command runs. The command gives Carlos the `clusterAdmin` and `readAnyDatabase` roles on the `admin` database and the `readWrite` role on the current database:

```
{ user: "Carlos",
  pwd: "cleartext password",
  customData: { employeeId: 12345 },
  roles: [
```

⁹<http://docs.mongodb.org/v2.4/reference/method/db.addUser>

¹⁰<http://docs.mongodb.org/v2.2/reference/method/db.addUser>

```
{ role: "clusterAdmin", db: "admin" },
{ role: "readAnyDatabase", db: "admin" },
"readWrite"
],
writeConcern: { w: "majority" , wtimeout: 5000 }
}
```

db.auth()**Definition**

db.auth (*username*, *password*)

Allows a user to authenticate to the database from within the shell.

param string username Specifies an existing username with access privileges for this database.

param string password Specifies the corresponding password.

Alternatively, you can use *mongo --username* and *--password* to specify authentication credentials.

Note: The *mongo* (page 527) shell excludes all *db.auth()* (page 99) operations from the saved history.

Returns *db.auth()* (page 99) returns 0 when authentication is **not** successful, and 1 when the operation is successful.

db.changeUserPassword()**Definition**

db.changeUserPassword (*username*, *password*)

Updates a user's password.

param string username Specifies an existing username with access privileges for this database.

param string password Specifies the corresponding password.

Example The following operation changes the *reporting* user's password to *SOh3TbYhx8ypJPxmt1oOfL*:

```
db.changeUserPassword("reporting", "SOh3TbYhx8ypJPxmt1oOfL")
```

db.cloneCollection()**Definition**

db.cloneCollection (*from*, *collection*, *query*)

Copies data directly between MongoDB instances. The *db.cloneCollection()* (page 99) wraps the *cloneCollection* (page 306) database command and accepts the following arguments:

param string from Host name of the MongoDB instance that holds the collection to copy.

param string collection The collection in the MongoDB instance that you want to copy. *db.cloneCollection()* (page 99) will only copy the collection with this name from *database* of the same name as the current database the remote MongoDB instance. If you want to copy a collection from a different database name you must use the *cloneCollection* (page 306) directly.

param document query A standard query document that limits the documents copied as part of the `db.cloneCollection()` (page 99) operation. All *query selectors* (page 373) available to the `find()` (page 34) are available here.

`db.cloneCollection()` (page 99) does not allow you to clone a collection through a `mongos` (page 518). You must connect directly to the `mongod` (page 503) instance.

`db.cloneDatabase()`

Definition

`db.cloneDatabase("hostname")`

Copies a remote database to the current database. The command assumes that the remote database has the same name as the current database.

param string hostname The hostname of the database to copy.

This method provides a wrapper around the MongoDB *database command* “`clone` (page 305).” The `copydb` (page 300) database command provides related functionality.

Example To clone a database named `importdb` on a host named `hostname`, issue the following:

```
use importdb
db.cloneDatabase("hostname")
```

New databases are implicitly created, so the current host does not need to have a database named `importdb` for this command to succeed.

`db.commandHelp()`

Description

`db.commandHelp(command)`

Displays help text for the specified *database command*. See the *Database Commands* (page 198).

The `db.commandHelp()` (page 100) method has the following parameter:

param string command The name of a *database command*.

`db.copyDatabase()`

Definition

`db.copyDatabase(fromdb, todb, fromhost, username, password)`

Copies a database from a remote host to the current host or copies a database to another database within the current host. `db.copyDatabase()` (page 100) wraps the `copydb` (page 300) command and takes the following arguments:

param string fromdb The name of the source database.

param string todb The name of the destination database.

param string fromhost The name of the source database host. Omit the hostname to copy from one database to another on the same server.

field string username The username credentials on the `fromhost` for authentication and authorization.

field string password The password on the `fromhost` for authentication and authorization. The method does not transmit the password in plaintext.

Behavior Be aware of the following properties of `db.copyDatabase()` (page 100):

- `db.copyDatabase()` (page 100) runs on the destination `mongod` (page 503) instance, i.e. the host receiving the copied data.
- `db.copyDatabase()` (page 100) creates the target database if it does not exist.
- `db.copyDatabase()` (page 100) requires enough free disk space on the host instance for the copied database. Use the `db.stats()` (page 119) operation to check the size of the database on the source `mongod` (page 503) instance.
- `db.copyDatabase()` (page 100) and `clone` (page 305) do not produce point-in-time snapshots of the source database. Write traffic to the source or destination database during the copy process will result in divergent data sets.
- `db.copyDatabase()` (page 100) does not lock the destination server during its operation, so the copy will occasionally yield to allow other operations to complete.

Required Access Changed in version 2.6.

The `copydb` (page 300) command requires the following authorization on the target and source databases.

Source Database (`fromdb`)

Source is non-admin Database If the source database is a non-admin database, you must have privileges that specify `find` action on the source database, and `find` action on the `system.js` collection in the source database. For example:

```
{ resource: { db: "mySourceDB", collection: "" }, actions: [ "find" ] }
{ resource: { db: "mySourceDB", collection: "system.js" }, actions: [ "find" ] }
```

If the source database is on a remote server, you also need the `find` action on the `system.indexes` and `system.namespaces` collections in the source database; e.g.

```
{ resource: { db: "mySourceDB", collection: "system.indexes" }, actions: [ "find" ] }
{ resource: { db: "mySourceDB", collection: "system.namespaces" }, actions: [ "find" ] }
```

Source is admin Database If the source database is the admin database, you must have privileges that specify `find` action on the admin database, and `find` action on the `system.js`, `system.users`, `system.roles`, and `system.version` collections in the admin database. For example:

```
{ resource: { db: "admin", collection: "" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.js" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.users" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.roles" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.version" }, actions: [ "find" ] }
```

If the source database is on a remote server, the you also need the `find` action on the `system.indexes` and `system.namespaces` collections in the admin database; e.g.

```
{ resource: { db: "admin", collection: "system.indexes" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.namespaces" }, actions: [ "find" ] }
```

Source Database is on a Remote Server If copying from a remote server and the remote server has authentication enabled, you must authenticate to the remote host as a user with the proper authorization. See [Authentication](#) (page 102).

Target Database (todb)

Copy from non-admin Database If the source database is not the admin database, you must have privileges that specify insert and createIndex actions on the target database, and insert action on the system.js collection in the target database. For example:

```
{ resource: { db: "myTargetDB", collection: "" }, actions: [ "insert", "createIndex" ] }
{ resource: { db: "myTargetDB", collection: "system.js" }, actions: [ "insert" ] }
```

Copy from admin Database If the source database is the admin database, you must have privileges that specify insert and createIndex actions on the target database, and insert action on the system.js, system.users, system.roles, and system.version collections in the target database. For example:

```
{ resource: { db: "myTargetDB", collection: "" }, actions: [ "insert", "createIndex" ] },
{ resource: { db: "myTargetDB", collection: "system.js" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.users" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.roles" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.version" }, actions: [ "insert" ] }
```

Authentication If copying from a remote server and the remote server has authentication enabled, then you must include the <username> and <password>. The method does not transmit the password in plaintext.

Example To copy a database named records into a database named archive_records, use the following invocation of db.copyDatabase() (page 100):

```
db.copyDatabase('records', 'archive_records')
```

See also:

clone (page 305)

db.createCollection()

Definition

db.createCollection(*name*, *options*)

Creates a new collection explicitly.

Because MongoDB creates a collection implicitly when the collection is first referenced in a command, this method is used primarily for creating new *capped collections*. This is also used to pre-allocate space for an ordinary collection.

The db.createCollection() (page 102) method has the following prototype form:

```
db.createCollection(name, {capped: <boolean>, autoIndexId: <boolean>, size: <number>, max: <number>})
```

The db.createCollection() (page 102) method has the following parameters:

param string name The name of the collection to create.

param document options Configuration options for creating a capped collection or for preallocating space in a new collection.

The `options` document creates a capped collection or preallocates space in a new ordinary collection. The `options` document contains the following fields:

field Boolean capped Enables a *capped collection*. To create a capped collection, specify `true`. If you specify `true`, you must also set a maximum size in the `size` field.

field Boolean autoIndexId Specify `false` to disable the automatic creation of an index on the `_id` field. Before 2.2, the default value for `autoIndexId` was `false`. See *[_id Fields and Indexes on Capped Collections](#)* (page 672) for more information.

field number size Specifies a maximum size in bytes for a capped collection. The `size` field is required for capped collections. If `capped` is `false`, you can use this field to preallocate space for an ordinary collection.

field number max The maximum number of documents allowed in the capped collection. The `size` limit takes precedence over this limit. If a capped collection reaches its maximum `size` before it reaches the maximum number of documents, MongoDB removes old documents. If you prefer to use this limit, ensure that the `size` limit, which is required, is sufficient to contain the documents limit.

field boolean usePowerOf2Sizes New in version 2.6: `usePowerOf2Sizes` (page 316) became an option to `db.createCollection()` (page 102) when `usePowerOf2Sizes` (page 316) became the default allocation strategy for all new collections by default.

Set to `false` to disable the `usePowerOf2Sizes` (page 316) allocation strategy for this collection. Defaults to `true` unless the `newCollectionsUsePowerOf2Sizes` parameter is set to `false`.

Example The following example creates a capped collection. Capped collections have maximum size or document counts that prevent them from growing beyond maximum thresholds. All capped collections must specify a maximum size and may also specify a maximum document count. MongoDB removes older documents if a collection reaches the maximum size limit before it reaches the maximum document count. Consider the following example:

```
db.createCollection("log", { capped : true, size : 5242880, max : 5000 } )
```

This command creates a collection named `log` with a maximum size of 5 megabytes and a maximum of 5000 documents.

The following command simply pre-allocates a 2-gigabyte, uncapped collection named `people`:

```
db.createCollection("people", { size: 2147483648 } )
```

This command provides a wrapper around the database command `create` (page 304). See <http://docs.mongodb.org/manualcore/capped-collections> for more information about capped collections.

db.currentOp()

Definition

`db.currentOp()`

Returns A *document* that reports in-progress operations for the database instance.

The `db.currentOp()` (page 103) method can take no arguments or take the `true` argument, which returns a more verbose output, including idle connections and system operations. The following example uses the `true` argument:

```
db.currentOp(true)
```

`db.currentOp()` (page 103) is available only for users with administrative privileges.

You can use `db.killOp()` (page 114) in conjunction with the `opid` (page 105) field to terminate a currently running operation. The following JavaScript operations for the `mongo` (page 527) shell filter the output of specific types of operations:

- Return all pending write operations:

```
db.currentOp().inprog.forEach(  
  function(d){  
    if(d.waitingForLock && d.lockType != "read")  
      printjson(d)  
  })
```

- Return the active write operation:

```
db.currentOp().inprog.forEach(  
  function(d){  
    if(d.active && d.lockType == "write")  
      printjson(d)  
  })
```

- Return all active read operations:

```
db.currentOp().inprog.forEach(  
  function(d){  
    if(d.active && d.lockType == "read")  
      printjson(d)  
  })
```

Warning: Terminate running operations with extreme caution. Only use `db.killOp()` (page 114) to terminate operations initiated by clients and *do not* terminate internal database operations.

Example The following is an example of `db.currentOp()` (page 103) output. If you specify the `true` argument, `db.currentOp()` (page 103) returns more verbose output.

```
{  
  "inprog": [  
    {  
      "opid" : 3434473,  
      "active" : <boolean>,  
      "secs_running" : 0,  
      "op" : "<operation>",  
      "ns" : "<database>.<collection>",  
      "query" : {  
      },  
      "client" : "<host>:<outgoing>",  
      "desc" : "conn57683",  
      "threadId" : "0x7f04a637b700",  
      "connectionId" : 57683,  
      "locks" : {  
        "^" : "w",  
        "^local" : "W",  
        "^<database>" : "W"  
      },  
      "waitingForLock" : false,  
    }  
  ]  
}
```

```

    "msg": "<string>"
    "numYields" : 0,
    "progress" : {
      "done" : <number>,
      "total" : <number>
    }
    "lockStats" : {
      "timeLockedMicros" : {
        "R" : NumberLong(),
        "W" : NumberLong(),
        "r" : NumberLong(),
        "w" : NumberLong()
      },
      "timeAcquiringMicros" : {
        "R" : NumberLong(),
        "W" : NumberLong(),
        "r" : NumberLong(),
        "w" : NumberLong()
      }
    }
  },
]
}

```

Output Changed in version 2.2.

The `db.currentOp()` (page 103) returns a document with an array named `inprog`. The `inprog` array contains a document for each in-progress operation. The fields that appear for a given operation depend on the kind of operation and its state.

`currentOp.opid`

Holds an identifier for the operation. You can pass this value to `db.killOp()` (page 114) in the `mongo` (page 527) shell to terminate the operation.

`currentOp.active`

A boolean value, that is `true` if the operation has started or `false` if the operation is queued and waiting for a lock to run. `active` (page 105) may be `true` even if the operation has yielded to another operation.

`currentOp.secs_running`

The duration of the operation in seconds. MongoDB calculates this value by subtracting the current time from the start time of the operation.

If the operation is not running, (i.e. if `active` (page 105) is `false`), this field may not appear in the output of `db.currentOp()` (page 103).

`currentOp.op`

A string that identifies the type of operation. The possible values are:

- insert
- query
- update
- remove
- getmore
- command

currentOp.ns

The *namespace* the operation targets. MongoDB forms namespaces using the name of the *database* and the name of the *collection*.

currentOp.query

A document containing the current operation's query. The document is empty for operations that do not have queries: *getmore*, *insert*, and *command*.

currentOp.client

The IP address (or hostname) and the ephemeral port of the client connection where the operation originates. If your *inprog* array has operations from many different clients, use this string to relate operations to clients.

For some commands, including *findAndModify* (page 229) and *db.eval()* (page 108), the client will be *0.0.0.0:0*, rather than an actual client.

currentOp.desc

A description of the client. This string includes the *connectionId* (page 106).

currentOp.threadId

An identifier for the thread that services the operation and its connection.

currentOp.connectionId

An identifier for the connection where the operation originated.

currentOp.locks

New in version 2.2.

The *locks* (page 106) document reports on the kinds of locks the operation currently holds. The following kinds of locks are possible:

currentOp.locks.^

^ (page 106) reports on the use of the global lock for the *mongod* (page 503) instance. All operations must hold the global lock for some phases of operation.

currentOp.locks.^local

^local (page 106) reports on the lock for the *local* database. MongoDB uses the *local* database for a number of operations, but the most frequent use of the *local* database is for the *oplog* used in replication.

currentOp.locks.^<database>

locks.^<database> (page 106) reports on the lock state for the database that this operation targets.

currentOp.waitingForLock

Returns a boolean value. *waitingForLock* (page 106) is *true* if the operation is waiting for a lock and *false* if the operation has the required lock.

currentOp.msg

The *msg* (page 106) provides a message that describes the status and progress of the operation. In the case of indexing or mapReduce operations, the field reports the completion percentage.

currentOp.progress

Reports on the progress of mapReduce or indexing operations. The *progress* (page 106) fields corresponds to the completion percentage in the *msg* (page 106) field. The *progress* (page 106) specifies the following information:

currentOp.progress.done

Reports the number completed.

currentOp.progress.total

Reports the total number.

currentOp.killed

Returns true if `mongod` (page 503) instance is in the process of killing the operation.

currentOp.numYields

`numYields` (page 107) is a counter that reports the number of times the operation has yielded to allow other operations to complete.

Typically, operations yield when they need access to data that MongoDB has not yet fully read into memory. This allows other operations that have data in memory to complete quickly while MongoDB reads in data for the yielding operation.

currentOp.lockStats

New in version 2.2.

The `lockStats` (page 107) document reflects the amount of time the operation has spent both acquiring and holding locks. `lockStats` (page 107) reports data on a per-lock type, with the following possible lock types:

- R represents the global read lock,
- W represents the global write lock,
- r represents the database specific read lock, and
- w represents the database specific write lock.

currentOp.timeLockedMicros

The `timeLockedMicros` (page 107) document reports the amount of time the operation has spent holding a specific lock.

For operations that require more than one lock, like those that lock the `local` database to update the *oplog*, then the values in this document can be longer than this value may be longer than the total length of the operation (i.e. `secs_running` (page 105).)

currentOp.timeLockedMicros.R

Reports the amount of time in microseconds the operation has held the global read lock.

currentOp.timeLockedMicros.W

Reports the amount of time in microseconds the operation has held the global write lock.

currentOp.timeLockedMicros.r

Reports the amount of time in microseconds the operation has held the database specific read lock.

currentOp.timeLockedMicros.w

Reports the amount of time in microseconds the operation has held the database specific write lock.

currentOp.timeAcquiringMicros

The `timeAcquiringMicros` (page 107) document reports the amount of time the operation has spent *waiting* to acquire a specific lock.

currentOp.timeAcquiringMicros.R

Reports the mount of time in microseconds the operation has waited for the global read lock.

currentOp.timeAcquiringMicros.W

Reports the mount of time in microseconds the operation has waited for the global write lock.

currentOp.timeAcquiringMicros.r

Reports the mount of time in microseconds the operation has waited for the database specific read lock.

```
currentOp.timeAcquiringMicros.w
```

Reports the mount of time in microseconds the operation has waited for the database specific write lock.

db.dropDatabase()

db.dropDatabase()

Removes the current database. Does not change the current database, so the insertion of any documents in this database will allocate a fresh set of data files.

db.eval()

Definition

db.eval() (*function, arguments*)

Provides the ability to run JavaScript code on the MongoDB server.

If authentication is enabled, you must have access to all actions on all resources in order to run `db.eval()` (page 108). Providing such access is not recommended, but if your organization requires a user to run `db.eval()` (page 108), create a role that grants `anyAction` on *resource-anyresource*. Do not assign this role to any other user.

The helper `db.eval()` (page 108) in the `mongo` (page 527) shell wraps the `eval` (page 238) command. Therefore, the helper method shares the characteristics and behavior of the underlying command with *one exception*: `db.eval()` (page 108) method does not support the `nolock` option.

The method accepts the following parameters:

param JavaScript function function A JavaScript function to execute.

param list arguments A list of arguments to pass to the JavaScript function. Omit if the function does not take arguments.

The JavaScript function need not take any arguments, as in the first example, or may optionally take arguments as in the second:

```
function () {  
    // ...  
}  
  
function (arg1, arg2) {  
    // ...  
}
```

Examples The following is an example of the `db.eval()` (page 108) method:

```
db.eval( function(name, incAmount) {  
    var doc = db.myCollection.findOne( { name : name } );  
  
    doc = doc || { name : name , num : 0 , total : 0 , avg : 0 };  
  
    doc.num++;  
    doc.total += incAmount;  
    doc.avg = doc.total / doc.num;  
  
    db.myCollection.save( doc );  
    return doc;  
})
```



```
    },
    "eliot", 5 );
```

- The `db` in the function refers to the current database.
- `"eliot"` is the argument passed to the function, and corresponds to the `name` argument.
- `5` is an argument to the function and corresponds to the `incAmount` field.

If you want to use the server's interpreter, you must run `db.eval()` (page 108). Otherwise, the `mongo` (page 527) shell's JavaScript interpreter evaluates functions entered directly into the shell.

If an error occurs, `db.eval()` (page 108) throws an exception. The following is an example of an invalid function that uses the variable `x` without declaring it as an argument:

```
db.eval( function() { return x + x; }, 3 );
```

The statement results in the following exception:

```
{
  "errmsg" : "exception: JavaScript execution failed: ReferenceError: x is not defined near '{ return",
  "code" : 16722,
  "ok" : 0
}
```

Warning:

- By default, `db.eval()` (page 108) takes a global write lock before evaluating the JavaScript function. As a result, `db.eval()` (page 108) blocks all other read and write operations to the database while the `db.eval()` (page 108) operation runs. Set `noLock` to `true` on the `eval` (page 238) *command* to prevent the `eval` (page 238) *command* from taking the global write lock before evaluating the JavaScript. `noLock` does not impact whether operations within the JavaScript code itself takes a write lock.
 - Do not use `db.eval()` (page 108) for long running operations as `db.eval()` (page 108) blocks all other operations. Consider using other server side code execution options.
 - You can not use `db.eval()` (page 108) with *sharded* data. In general, you should avoid using `db.eval()` (page 108) in *sharded cluster*; nevertheless, it is possible to use `db.eval()` (page 108) with non-sharded collections and databases stored in a *sharded cluster*.
 - With authentication enabled, `db.eval()` (page 108) will fail during the operation if you do not have the permission to perform a specified task.
- Changed in version 2.4: You must have full admin access to run.

Changed in version 2.4: The V8 JavaScript engine, which became the default in 2.4, allows multiple JavaScript operations to execute at the same time. Prior to 2.4, `db.eval()` (page 108) executed in a single thread.

See also:

<http://docs.mongodb.org/manualcore/server-side-javascript>

db.fsyncLock()

`db.fsyncLock()`

Forces the `mongod` (page 503) to flush all pending write operations to the disk and locks the *entire* `mongod` (page 503) instance to prevent additional writes until the user releases the lock with the `db.fsyncUnlock()` (page 110) command. `db.fsyncLock()` (page 109) is an administrative command.

This command provides a simple wrapper around a `fsync` (page 312) database command with the following syntax:

```
{ fsync: 1, lock: true }
```

This function locks the database and create a window for backup operations.

The database cannot be locked with `db.fsyncLock()` (page 109) while profiling is enabled. You must disable profiling before locking the database with `db.fsyncLock()` (page 109). Disable profiling using `db.setProfilingLevel()` (page 119) as follows in the `mongo` (page 527) shell:

```
db.setProfilingLevel(0)
```

db.fsyncUnlock()

`db.fsyncUnlock()`

Unlocks a `mongod` (page 503) instance to allow writes and reverses the operation of a `db.fsyncLock()` (page 109) operation. Typically you will use `db.fsyncUnlock()` (page 110) following a database backup operation.

`db.fsyncUnlock()` (page 110) is an administrative command.

db.getCollection()

Description

`db.getCollection()` (*name*)

Returns a collection name. This is useful for a collection whose name might interact with the shell itself, such names that begin with `_` or that mirror the *database commands* (page 198).

The `db.getCollection()` (page 110) method has the following parameter:

param string name The name of the collection.

db.getCollectionNames()

`db.getCollectionNames()`

Returns An array containing all collections in the existing database.

db.getLastError()

`db.getLastError()`

Changed in version 2.6: A new protocol for *write operations* (page 623) integrates write concerns with the write operations, eliminating the need for a separate `db.getLastError()` (page 110) method. Write methods now return the status of the write operation, including error information.

In previous versions, clients typically used the `db.getLastError()` (page 110) method in combination with the write operations to ensure that the write succeeds.

Returns The last error message string.

Sets the level of *write concern* for confirming the success of write operations.

See

`getLastError` (page 235) and <http://docs.mongodb.org/manualreference/write-concern> for all options, *Write Concern* for a conceptual overview, <http://docs.mongodb.org/manualcore/write-operati> for information about all write operations in MongoDB.

db.getLastErrorObj()**db.getLastErrorObj()**

Changed in version 2.6: A new protocol for *write operations* (page 623) integrates write concerns with the write operations, eliminating the need for a separate `db.getLastError()` (page 110) method. Write methods now return the status of the write operation, including error information.

In previous versions, clients typically used the `db.getLastError()` (page 110) method in combination with the write operations to ensure that the write succeeds.

Returns A full *document* with status information.

See also:

Write Concern, <http://docs.mongodb.org/manualreference/write-concern>, and *replica-set-write-concern*.

db.getMongo()**db.getMongo()**

Returns The current database connection.

`db.getMongo()` (page 111) runs when the shell initiates. Use this command to test that the `mongo` (page 527) shell has a connection to the proper database instance.

db.getName()**db.getName()**

Returns the current database name.

db.getPrevError()**db.getPrevError()**

Returns A status document, containing the errors.

Deprecated since version 1.6.

This output reports all errors since the last time the database received a `resetError` (page 237) (also `db.resetError()` (page 117)) command.

This method provides a wrapper around the `getPrevError` (page 237) command.

db.getProfilingLevel()**db.getProfilingLevel()**

This method provides a wrapper around the database command “`profile` (page 334)” and returns the current profiling level.

Deprecated since version 1.8.4: Use `db.getProfilingStatus()` (page 112) for related functionality.

db.getProfilingStatus()

`db.getProfilingStatus()`

Returns The current `profile` (page 334) level and `slowOpThresholdMs` setting.

db.getReplicationInfo()

Definition

`db.getReplicationInfo()`

Returns A document with the status of the replica status, using data polled from the “*oplog*”. Use this output when diagnosing issues with replication.

Output

`db.getReplicationInfo.logSizeMB`

Returns the total size of the *oplog* in megabytes. This refers to the total amount of space allocated to the oplog rather than the current size of operations stored in the oplog.

`db.getReplicationInfo.usedMB`

Returns the total amount of space used by the *oplog* in megabytes. This refers to the total amount of space currently used by operations stored in the oplog rather than the total amount of space allocated.

`db.getReplicationInfo.errmsg`

Returns an error message if there are no entries in the oplog.

`db.getReplicationInfo.oplogMainRowCount`

Only present when there are no entries in the oplog. Reports a the number of items or rows in the *oplog* (e.g. 0).

`db.getReplicationInfo.timeDiff`

Returns the difference between the first and last operation in the *oplog*, represented in seconds.

Only present if there are entries in the oplog.

`db.getReplicationInfo.timeDiffHours`

Returns the difference between the first and last operation in the *oplog*, rounded and represented in hours.

Only present if there are entries in the oplog.

`db.getReplicationInfo.tFirst`

Returns a time stamp for the first (i.e. earliest) operation in the *oplog*. Compare this value to the last write operation issued against the server.

Only present if there are entries in the oplog.

`db.getReplicationInfo.tLast`

Returns a time stamp for the last (i.e. latest) operation in the *oplog*. Compare this value to the last write operation issued against the server.

Only present if there are entries in the oplog.

`db.getReplicationInfo.now`

Returns a time stamp that reflects reflecting the current time. The shell process generates this value, and the datum may differ slightly from the server time if you’re connecting from a remote host as a result. Equivalent to `Date()` (page 186).

Only present if there are entries in the oplog.

db.getSiblingDB()**Definition**

`db.getSiblingDB(<database>)`

param string database The name of a MongoDB database.

Returns A database object.

Used to return another database without modifying the `db` variable in the shell environment.

Example You can use `db.getSiblingDB()` (page 113) as an alternative to the `use <database>` helper. This is particularly useful when writing scripts using the `mongo` (page 527) shell where the `use` helper is not available. Consider the following sequence of operations:

```
db = db.getSiblingDB('users')
db.active.count()
```

This operation sets the `db` object to point to the database named `users`, and then returns a *count* (page 25) of the collection named `active`. You can create multiple `db` objects, that refer to different databases, as in the following sequence of operations:

```
users = db.getSiblingDB('users')
records = db.getSiblingDB('records')
```

```
users.active.count()
users.active.findOne()
```

```
records.requests.count()
records.requests.findOne()
```

This operation creates two `db` objects referring to different databases (i.e. `users` and `records`) and then returns a *count* (page 25) and an *example document* (page 43) from one collection in that database (i.e. `active` and `requests` respectively.)

db.help()

`db.help()`

Returns Text output listing common methods on the `db` object.

db.hostInfo()

`db.hostInfo()`

New in version 2.2.

Returns A document with information about the underlying system that the `mongod` (page 503) or `mongos` (page 518) runs on. Some of the returned fields are only included on some platforms.

`db.hostInfo()` (page 113) provides a helper in the `mongo` (page 527) shell around the `hostInfo` (page 344) The output of `db.hostInfo()` (page 113) on a Linux system will resemble the following:

```
{
  "system" : {
    "currentTime" : ISODate("<timestamp>"),
```

```
    "hostname" : "<hostname>",
    "cpuAddrSize" : <number>,
    "memSizeMB" : <number>,
    "numCores" : <number>,
    "cpuArch" : "<identifier>",
    "numaEnabled" : <boolean>
  },
  "os" : {
    "type" : "<string>",
    "name" : "<string>",
    "version" : "<string>"
  },
  "extra" : {
    "versionString" : "<string>",
    "libcVersion" : "<string>",
    "kernelVersion" : "<string>",
    "cpuFrequencyMHz" : "<string>",
    "cpuFeatures" : "<string>",
    "pageSize" : <number>,
    "numPages" : <number>,
    "maxOpenFiles" : <number>
  },
  "ok" : <return>
}
```

See `hostInfo` (page 345) for full documentation of the output of `db.hostInfo()` (page 113).

db.isMaster()

`db.isMaster()`

Returns A document that describes the role of the `mongod` (page 503) instance.

If the `mongod` (page 503) is a member of a *replica set*, then the `ismaster` (page 280) and `secondary` (page 281) fields report if the instance is the *primary* or if it is a *secondary* member of the replica set.

See

`isMaster` (page 280) for the complete documentation of the output of `db.isMaster()` (page 114).

db.killOp()

Description

`db.killOp` (*opid*)

Terminates an operation as specified by the operation ID. To find operations and their corresponding IDs, see `db.currentOp()` (page 103).

The `db.killOp()` (page 114) method has the following parameter:

param number opid An operation ID.

Warning: Terminate running operations with extreme caution. Only use `db.killOp()` (page 114) to terminate operations initiated by clients and *do not* terminate internal database operations.

db.listCommands()**db.listCommands()**

Provides a list of all database commands. See the *Database Commands* (page 198) document for a more extensive index of these options.

db.loadServerScripts()**db.loadServerScripts()**

`db.loadServerScripts()` (page 115) loads all scripts in the `system.js` collection for the current database into the `mongo` (page 527) shell session.

Documents in the `system.js` collection have the following prototype form:

```
{ _id : "<name>" , value : <function> }
```

The documents in the `system.js` collection provide functions that your applications can use in any JavaScript context with MongoDB in this database. These contexts include `$where` (page 391) clauses and `mapReduce` (page 208) operations.

db.logout()**db.logout()**

Ends the current authentication session. This function has no effect if the current session is not authenticated.

Note: If you're not logged in and using authentication, `db.logout()` (page 115) has no effect.

Changed in version 2.4: Because MongoDB now allows users defined in one database to have privileges on another database, you must call `db.logout()` (page 115) while using the same database context that you authenticated to.

If you authenticated to a database such as `users` or `$external`, you must issue `db.logout()` (page 115) against this database in order to successfully log out.

Example

Use the `use <database-name>` helper in the interactive `mongo` (page 527) shell, or the following `db.getSiblingDB()` (page 113) in the interactive shell or in `mongo` (page 527) shell scripts to change the `db` object:

```
db = db.getSiblingDB('<database-name>')
```

When you have set the database context and `db` object, you can use the `db.logout()` (page 115) to log out of database as in the following operation:

```
db.logout()
```

`db.logout()` (page 115) function provides a wrapper around the database command `logout` (page 249).

db.printCollectionStats()**db.printCollectionStats()**

Provides a wrapper around the `db.collection.stats()` (page 68) method. Returns statistics from every collection separated by three hyphen characters.

Note: The `db.printCollectionStats()` (page 115) in the `mongo` (page 527) shell does **not** return *JSON*. Use `db.printCollectionStats()` (page 115) for manual inspection, and `db.collection.stats()` (page 68) in scripts.

See also:

`collStats` (page 325)

`db.printReplicationInfo()`

`db.printReplicationInfo()`

Provides a formatted report of the status of a *replica set* from the perspective of the *primary* set member. See the *replSetGetStatus* (page 273) for more information regarding the contents of this output.

Note: The `db.printReplicationInfo()` (page 116) in the `mongo` (page 527) shell does **not** return *JSON*. Use `db.printReplicationInfo()` (page 116) for manual inspection, and `rs.status()` (page 168) in scripts.

`db.printShardingStatus()`

Definition

`db.printShardingStatus()`

Prints a formatted report of the sharding configuration and the information regarding existing chunks in a *sharded cluster*.

Only use `db.printShardingStatus()` (page 116) when connected to a `mongos` (page 518) instance.

The `db.printShardingStatus()` (page 116) method has the following parameter:

param Boolean verbose If `true`, the method displays details of the document distribution across chunks when you have 20 or more chunks.

See *sh.status()* (page 180) for details of the output.

Note: The `db.printShardingStatus()` (page 116) in the `mongo` (page 527) shell does **not** return *JSON*. Use `db.printShardingStatus()` (page 116) for manual inspection, and *Config Database* (page 593) in scripts.

See also:

`sh.status()` (page 180)

`db.printSlaveReplicationInfo()`

Definition

`db.printSlaveReplicationInfo()`

Returns a formatted report of the status of a *replica set* from the perspective of the *secondary* member of the set. The output is identical to that of `rs.printSlaveReplicationInfo()` (page 166).

Output The following is example output from the `rs.printSlaveReplicationInfo()` (page 166) method issued on a replica set with two secondary members:

```
source: m1.example.net:27017
  syncedTo: Thu Apr 10 2014 10:27:47 GMT-0400 (EDT)
  0 secs (0 hrs) behind the primary
source: m2.example.net:27017
  syncedTo: Thu Apr 10 2014 10:27:47 GMT-0400 (EDT)
  0 secs (0 hrs) behind the primary
```

Note: The `db.printSlaveReplicationInfo()` (page 116) in the `mongo` (page 527) shell does **not** return *JSON*. Use `db.printSlaveReplicationInfo()` (page 116) for manual inspection, and `rs.status()` (page 168) in scripts.

A *delayed member* may show as 0 seconds behind the primary when the inactivity period on the primary is greater than the `slaveDelay` value.

`db.removeUser()`

Deprecated since version 2.6: Use `db.dropUser()` (page 148) instead of `db.removeUser()` (page 147)

Definition

`db.removeUser(username)`

Removes the specified username from the database.

The `db.removeUser()` (page 147) method has the following parameter:

param string username The database username.

`db.repairDatabase()`

`db.repairDatabase()`

Warning: During normal operations, only use the `repairDatabase` (page 319) command and wrappers including `db.repairDatabase()` (page 117) in the `mongo` (page 527) shell and `mongod --repair`, to compact database files and/or reclaim disk space. Be aware that these operations remove and do not save any corrupt data during the repair process.

If you are trying to repair a *replica set* member, and you have access to an intact copy of your data (e.g. a recent backup or an intact member of the *replica set*), you should restore from that intact copy, and **not** use `repairDatabase` (page 319).

When using *journaling*, there is almost never any need to run `repairDatabase` (page 319). In the event of an unclean shutdown, the server will be able restore the data files to a pristine state automatically.

`db.repairDatabase()` (page 117) provides a wrapper around the database command `repairDatabase` (page 319), and has the same effect as the run-time option `mongod --repair` option, limited to *only* the current database. See `repairDatabase` (page 319) for full documentation.

`db.resetError()`

`db.resetError()`

Deprecated since version 1.6.

Resets the error message returned by `db.getPrevError` (page 111) or `getPrevError` (page 237). Provides a wrapper around the `resetError` (page 237) command.

db.runCommand()

Definition

`db.runCommand (command)`

Provides a helper to run specified *database commands* (page 198). This is the preferred method to issue database commands, as it provides a consistent interface between the shell and drivers.

param document, string command “A *database command*, specified either in *document* form or as a string. If specified as a string, `db.runCommand()` (page 118) transforms the string into a document.”

New in version 2.6: To specify a time limit in milliseconds, see <http://docs.mongodb.org/manual/tutorial/terminate-running-operations>.

Behavior `db.runCommand()` (page 118) runs the command in the context of the current database. Some commands are only applicable in the context of the `admin` database, and you must change your `db` object to before running these commands.

db.serverBuildInfo()

`db.serverBuildInfo()`

Provides a wrapper around the `buildInfo` (page 324) *database command*. `buildInfo` (page 324) returns a document that contains an overview of parameters used to compile this `mongod` (page 503) instance.

db.serverStatus()

`db.serverStatus()`

Returns a *document* that provides an overview of the database process’s state.

This command provides a wrapper around the database command `serverStatus` (page 347).

Changed in version 2.4: In 2.4 you can dynamically suppress portions of the `db.serverStatus()` (page 118) output, or include suppressed sections in a document passed to the `db.serverStatus()` (page 118) method, as in the following example:

```
db.serverStatus( { repl: 0, indexCounters: 0, locks: 0 } )
db.serverStatus( { workingSet: 1, metrics: 0, locks: 0 } )
```

`db.serverStatus()` (page 118) includes all fields by default, except `workingSet` (page 359), by default.

Note: You may only dynamically include top-level fields from the *serverStatus* (page 346) document that are not included by default. You can exclude any field that `db.serverStatus()` (page 118) includes by default.

See also:

serverStatus (page 346) for complete documentation of the output of this function.

db.setProfilingLevel()**Definition**

db.setProfilingLevel (*level*, *slowms*)

Modifies the current *database profiler* level used by the database profiling system to capture data about performance. The method provides a wrapper around the *database command profile* (page 334).

param integer level Specifies a profiling level, which is either 0 for no profiling, 1 for only slow operations, or 2 for all operations.

param integer slowms Sets the threshold in milliseconds for the profile to consider a query or operation to be slow.

The level chosen can affect performance. It also can allow the server to write the contents of queries to the log, which might have information security implications for your deployment.

Configure the `slowOpThresholdMs` option to set the threshold for the profiler to consider a query “slow.” Specify this value in milliseconds to override the default, 100ms.

`mongod` (page 503) writes the output of the database profiler to the `system.profile` collection.

`mongod` (page 503) prints information about queries that take longer than the `slowOpThresholdMs` to the log even when the database profiler is not active.

The database cannot be locked with `db.fsyncLock()` (page 109) while profiling is enabled. You must disable profiling before locking the database with `db.fsyncLock()` (page 109). Disable profiling using `db.setProfilingLevel()` (page 119) as follows in the `mongo` (page 527) shell:

```
db.setProfilingLevel(0)
```

db.shutdownServer()

db.shutdownServer ()

Shuts down the current `mongod` (page 503) or `mongos` (page 518) process cleanly and safely.

This operation fails when the current database *is not* the *admin database*.

This command provides a wrapper around the `shutdown` (page 321).

db.stats()**Description**

db.stats (*scale*)

Returns statistics that reflect the use state of a single *database*.

The `db.stats()` (page 119) method has the following parameter:

param number scale The scale at which to deliver results. Unless specified, this command returns all data in bytes.

Returns A *document* with statistics reflecting the database system’s state. For an explanation of the output, see *dbStats* (page 331).

The `db.stats()` (page 119) method is a wrapper around the `dbStats` (page 331) database command.

Example The following example converts the returned values to kilobytes:

```
db.stats(1024)
```

Note: The scale factor rounds values to whole numbers. This can produce unpredictable and unexpected results in some situations.

db.version()

```
db.version()
```

Returns The version of the [mongod](#) (page 503) or [mongos](#) (page 518) instance.

db.upgradeCheck()

Definition

`db.upgradeCheck (<document>)`

New in version 2.6.

Performs a preliminary check for upgrade preparedness to 2.6. The helper, available in the 2.6 [mongo](#) (page 527) shell, can run connected to either a 2.4 or a 2.6 server.

The method checks for:

- documents with index keys *longer than the index key limit* (page 627),
- documents with `illegal field names` (page 609),
- collections without an `_id` index, and
- indexes with invalid specifications, such as an index key with an empty or illegal field name.

The method can accept a document parameter which determine the scope of the check:

param document scope Document to limit the scope of the check to the specified collection in the database.

Omit to perform the check on all collections in the database.

The optional scope document has the following form:

```
{
  collection: <string>
}
```

Additional 2.6 changes that affect compatibility with older versions require manual checks and intervention. See *Compatibility Changes in MongoDB 2.6* (page 627) for details.

See also:

[db.upgradeCheckAllDBs\(\)](#) (page 121)

Behavior `db.upgradeCheck()` (page 120) performs collection scans and has an impact on performance. To mitigate the performance impact:

- For sharded clusters, configure to read from secondaries and run the command on the [mongos](#) (page 518).
- For replica sets, run the command on the secondary members.

`db.upgradeCheck()` (page 120) can miss new data during the check when run on a live system with active write operations.

Required Access On systems running with `authorization`, a user must have access that includes the `find` action on all collections, including the *system collections* (page 600).

Example The following example connects to a secondary running on `localhost` and runs `db.upgradeCheck()` (page 120) against the `employees` collection in the `records` database. Because the output from the method can be quite large, the example pipes the output to a file.

```
./mongo --eval "db.getMongo().setSlaveOk(); db.upgradeCheck( { collection: 'employees' } )" localhost
```

Error Output The upgrade check can return the following errors when it encounters incompatibilities in your data:

Index Key Exceed Limit

Document Error: key for index '<indexName>' (<indexSpec>) too long on document: <doc>

To resolve, remove the document. Ensure that the query to remove the document does not specify a condition on the invalid field or field.

Documents with Illegal Field Names

Document Error: document is no longer valid in 2.6 because <errmsg>: <doc>

To resolve, remove the document and re-insert with the appropriate corrections.

Index Specification Invalid

Index Error: invalid index spec for index '<indexName>': <indexSpec>

To resolve, remove the invalid index and recreate with a valid index specification.

Missing `_id` Index

Collection Error: lack of `_id` index on collection: <collectionName>

To resolve, create a unique index on `_id`.

`db.upgradeCheckAllDBs()`

Definition

`db.upgradeCheckAllDBs()`

New in version 2.6.

Performs a preliminary check for upgrade preparedness to 2.6. The helper, available in the 2.6 `mongo` (page 527) shell, can run connected to either a 2.4 or a 2.6 server in the `admin` database.

The method cycles through all the databases and checks for:

- documents with index keys *longer than the index key limit* (page 627),
- documents with `illegal field names` (page 609),
- collections without an `_id` index, and

- indexes with invalid specifications, such as an index key with an empty or illegal field name.

Additional 2.6 changes that affect compatibility with older versions require manual checks and intervention. See *Compatibility Changes in MongoDB 2.6* (page 627) for details.

See also:

`db.upgradeCheck()` (page 120)

Behavior `db.upgradeCheckAllDBs()` (page 121) performs collection scans and has an impact on performance. To mitigate the performance impact:

- For sharded clusters, configure to read from secondaries and run the command on the `mongos` (page 518).
- For replica sets, run the command on the secondary members.

`db.upgradeCheckAllDBs()` (page 121) can miss new data during the check when run on a live system with active write operations.

Required Access On systems running with authorization, a user must have access that includes the `listDatabases` action on all databases and the `find` action on all collections, including the *system collections* (page 600).

Example The following example connects to a secondary running on `localhost` and runs `db.upgradeCheckAllDBs()` (page 121) against the `admin` database. Because the output from the method can be quite large, the example pipes the output to a file.

```
./mongo --eval "db.getMongo().setSlaveOk(); db.upgradeCheckAllDBs();" localhost/admin | tee /tmp/upg
```

Error Output The upgrade check can return the following errors when it encounters incompatibilities in your data:

Index Key Exceed Limit

```
Document Error: key for index '<indexName>' (<indexSpec>) too long on document: <doc>
```

To resolve, remove the document. Ensure that the query to remove the document does not specify a condition on the invalid field or field.

Documents with Illegal Field Names

```
Document Error: document is no longer valid in 2.6 because <errmsg>: <doc>
```

To resolve, remove the document and re-insert with the appropriate corrections.

Index Specification Invalid

```
Index Error: invalid index spec for index '<indexName>': <indexSpec>
```

To resolve, remove the invalid index and recreate with a valid index specification.

Missing `_id` Index

```
Collection Error: lack of _id index on collection: <collectionName>
```

To resolve, create a unique index on `_id`.

2.1.4 Query Plan Cache

Query Plan Cache Methods

The PlanCache methods are only accessible from a collection's plan cache object. To retrieve the plan cache object, use the `db.collection.getPlanCache()` (page 123) method.

Name	Description
<code>db.collection.getPlanCache()</code> (page 123)	Returns an interface to access the query plan cache object and associated PlanCache methods for a collection."
<code>PlanCache.help()</code> (page 124)	Displays the methods available for a collection's query plan cache. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().help()</code> .
<code>PlanCache.listQueryShapes()</code> (page 124)	Displays the query shapes for which cached query plans exist. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().listQueryShapes()</code> .
<code>PlanCache.getPlansByQuery()</code> (page 125)	Displays the cached query plans for the specified query shape. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().getPlansByQuery()</code> .
<code>PlanCache.clearPlansByQuery()</code> (page 126)	Clears the cached query plans for the specified query shape. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().clearPlansByQuery()</code> .
<code>PlanCache.clear()</code> (page 127)	Clears all the cached query plans for a collection. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().clear()</code> .

`db.collection.getPlanCache()`

Definition

`db.collection.getPlanCache()`

Returns an interface to access the query plan cache for a collection. The interface provides methods to view and clear the query plan cache.

Returns Interface to access the query plan cache.

The query optimizer only caches the plans for those query shapes that can have more than one viable plan.

Methods The following methods are available through the interface:

Name	Description
<code>PlanCache.help()</code> (page 124)	Displays the methods available for a collection's query plan cache. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().help()</code> .
<code>PlanCache.listQueryShapes()</code> (page 124)	Displays the query shapes for which cached query plans exist. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().listQueryShapes()</code> .
<code>PlanCache.getPlansByQueryShape()</code> (page 125)	Displays the cached query plans for the specified query shape. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().getPlansByQueryShape()</code> .
<code>PlanCache.clearPlansByQueryShape()</code> (page 126)	Clears the cached query plans for the specified query shape. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().clearPlansByQueryShape()</code> .
<code>PlanCache.clear()</code> (page 127)	Clears all the cached query plans for a collection. Accessible through the plan cache object of a specific collection, i.e. <code>db.collection.getPlanCache().clear()</code> .

PlanCache.help()

Definition

`PlanCache.help()`

Displays the methods available to view and modify a collection's query plan cache.

The method is only available from the `plan cache object` (page 123) of a specific collection; i.e.

```
db.collection.getPlanCache().help()
```

See also:

`db.collection.getPlanCache()` (page 123)

PlanCache.listQueryShapes()

Definition

`PlanCache.listQueryShapes()`

Displays the *query shapes* for which cached query plans exist.

The query optimizer only caches the plans for those query shapes that can have more than one viable plan.

The method is only available from the `plan cache object` (page 123) of a specific collection; i.e.

```
db.collection.getPlanCache().listQueryShapes()
```

Returns Array of *query shape* documents.

The method wraps the `planCacheListQueryShapes` (page 245) command.

Required Access On systems running with authorization, a user must have access that includes the `planCacheRead` action.

Example The following returns the *query shapes* that have cached plans for the `orders` collection:

```
db.orders.getPlanCache().listQueryShapes()
```


The method returns an array of the query shapes currently in the cache. In the example, the `orders` collection had cached query plans associated with the following shapes:

```
[
  {
    "query" : { "qty" : { "$gt" : 10 } },
    "sort" : { "ord_date" : 1 },
    "projection" : { }
  },
  {
    "query" : { "$or" :
      [
        { "qty" : { "$gt" : 15 }, "item" : "xyz123" },
        { "status" : "A" }
      ]
    },
    "sort" : { },
    "projection" : { }
  },
  {
    "query" : { "$or" : [ { "qty" : { "$gt" : 15 } }, { "status" : "A" } ] },
    "sort" : { },
    "projection" : { }
  }
]
```

See also:

- `db.collection.getPlanCache()` (page 123)
- `PlanCache.getPlansByQuery()` (page 125)
- `PlanCache.help()` (page 124)
- `planCacheListQueryShapes` (page 245)

PlanCache.getPlansByQuery()

Definition

`PlanCache.getPlansByQuery(<query>, <projection>, <sort>)`

Displays the cached query plans for the specified *query shape*.

The query optimizer only caches the plans for those query shapes that can have more than one viable plan.

The method is only available from the `plan cache object` (page 123) of a specific collection; i.e.

```
db.collection.getPlanCache().getPlansByQuery( <query>, <projection>, <sort> )
```

The `PlanCache.getPlansByQuery()` (page 125) method accepts the following parameters:

param document query The query predicate of the *query shape*. Only the structure of the predicate, including the field names, are significant to the shape; the values in the query predicate are insignificant.

param document projection The projection associated with the *query shape*. Required if specifying the `sort` parameter.

param document sort The sort associated with the *query shape*.

Returns Array of cached query plans for a query shape.

To see the query shapes for which cached query plans exist, use the `PlanCache.listQueryShapes()` (page 124) method.

Required Access On systems running with authorization, a user must have access that includes the `planCacheRead` action.

Example If a collection `orders` has the following query shape:

```
{
  "query" : { "qty" : { "$gt" : 10 } },
  "sort" : { "ord_date" : 1 },
  "projection" : { }
}
```

The following operation displays the query plan cached for the shape:

```
db.orders.getPlanCache().getPlansByQuery(
  { "qty" : { "$gt" : 10 } },
  { },
  { "ord_date" : 1 }
)
```

See also:

- `db.collection.getPlanCache()` (page 123)
- `PlanCache.listQueryShapes()` (page 124)
- `PlanCache.help()` (page 124)

PlanCache.clearPlansByQuery()

Definition

`PlanCache.clearPlansByQuery(<query>, <projection>, <sort>)`

Clears the cached query plans for the specified *query shape*.

The method is only available from the `plan cache object` (page 123) of a specific collection; i.e.

```
db.collection.getPlanCache().clearPlansByQuery( <query>, <projection>, <sort> )
```

The `PlanCache.clearPlansByQuery()` (page 126) method accepts the following parameters:

param document query The query predicate of the *query shape*. Only the structure of the predicate, including the field names, are significant to the shape; the values in the query predicate are insignificant.

param document projection The projection associated with the *query shape*. Required if specifying the `sort` parameter.

param document sort The sort associated with the *query shape*.

To see the query shapes for which cached query plans exist, use the `PlanCache.listQueryShapes()` (page 124) method.

Required Access On systems running with authorization, a user must have access that includes the `planCacheWrite` action.

Example If a collection `orders` has the following query shape:

```
{
  "query" : { "qty" : { "$gt" : 10 } },
  "sort" : { "ord_date" : 1 },
  "projection" : { }
}
```

The following operation removes the query plan cached for the shape:

```
db.orders.getPlanCache().clearPlansByQuery(
  { "qty" : { "$gt" : 10 } },
  { },
  { "ord_date" : 1 }
)
```

See also:

- `db.collection.getPlanCache()` (page 123)
- `PlanCache.listQueryShapes()` (page 124)
- `PlanCache.clear()` (page 127)

PlanCache.clear()

Definition

`PlanCache.clear()`

Removes all cached query plans for a collection.

The method is only available from the `plan cache object` (page 123) of a specific collection; i.e.

```
db.collection.getPlanCache().clear()
```

For example, to clear the cache for the `orders` collection:

```
db.orders.getPlanCache().clear()
```

Required Access On systems running with authorization, a user must have access that includes the `planCacheWrite` action.

See also:

- `db.collection.getPlanCache()` (page 123)
- `PlanCache.clearPlansByQuery()` (page 126)

2.1.5 Bulk Write Operation

Bulk Operation Methods

New in version 2.6.

Name	Description
<code>Bulk()</code> (page 128)	Bulk operations builder.
<code>Bulk.insert()</code> (page 129)	Adds an insert operation to a list of operations.
<code>Bulk.find()</code> (page 130)	Specifies the query condition for an update or a remove operation.
<code>Bulk.find.removeOne()</code> (page 130)	Adds a single document remove operation to a list of operations.
<code>Bulk.find.remove()</code> (page 131)	Adds a multiple document remove operation to a list of operations.
<code>Bulk.find.replaceOne()</code> (page 131)	Adds a single document replacement operation to a list of operations.
<code>Bulk.find.updateOne()</code> (page 132)	Adds a single document update operation to a list of operations.
<code>Bulk.find.update()</code> (page 134)	Adds a multi update operation to a list of operations.
<code>Bulk.find.upsert()</code> (page 134)	Specifies the <i>upsert</i> flag for an update operation.
<code>Bulk.execute()</code> (page 137)	Executes a list of operations in bulk.
<code>Bulk.getOperations()</code> (page 138)	Returns an array of write operations executed in the <code>Bulk()</code> (page 128) operations object.
<code>Bulk.toJson()</code> (page 140)	Returns a JSON document that contains the number of operations and batches in the <code>Bulk()</code> (page 128) operations object.
<code>Bulk.toString()</code> (page 140)	Returns the <code>Bulk.toJson()</code> (page 140) results as a string.

Bulk()

Description

Bulk()

New in version 2.6.

Bulk operations builder used to construct a list of write operations to perform in bulk for a single collection. To instantiate the builder, use either the `db.collection.initializeOrderedBulkOp()` (page 51) or the `db.collection.initializeUnorderedBulkOp()` (page 51) method.

Ordered and Unordered Bulk Operations The builder can construct the list of operations as *ordered* or *unordered*.

Ordered Operations With an *ordered* operations list, MongoDB executes the write operations in the list serially. If an error occurs during the processing of one of the write operations, MongoDB will return without processing any remaining write operations in the list.

Use `db.collection.initializeOrderedBulkOp()` (page 51) to create a builder for an ordered list of write commands.

Unordered Operations With an *unordered* operations list, MongoDB can execute in parallel, as well as in a non-deterministic order, the write operations in the list. If an error occurs during the processing of one of the write operations, MongoDB will continue to process remaining write operations in the list.

Use `db.collection.initializeUnorderedBulkOp()` (page 51) to create a builder for an unordered list of write commands.

Methods The `Bulk()` (page 128) builder has the following methods:

Name	Description
<code>Bulk.insert()</code> (page 129)	Adds an insert operation to a list of operations.
<code>Bulk.find()</code> (page 130)	Specifies the query condition for an update or a remove operation.
<code>Bulk.find.removeOne()</code> (page 130)	Adds a single document remove operation to a list of operations.
<code>Bulk.find.remove()</code> (page 131)	Adds a multiple document remove operation to a list of operations.
<code>Bulk.find.replaceOne()</code> (page 131)	Adds a single document replacement operation to a list of operations.
<code>Bulk.find.updateOne()</code> (page 132)	Adds a single document update operation to a list of operations.
<code>Bulk.find.update()</code> (page 134)	Adds a multi update operation to a list of operations.
<code>Bulk.find.upsert()</code> (page 134)	Specifies the <i>upsert</i> flag for an update operation.
<code>Bulk.execute()</code> (page 137)	Executes a list of operations in bulk.
<code>Bulk.getOperations()</code> (page 138)	Returns an array of write operations executed in the <code>Bulk()</code> (page 128) operations object.
<code>Bulk.toJson()</code> (page 140)	Returns a JSON document that contains the number of operations and batches in the <code>Bulk()</code> (page 128) operations object.
<code>Bulk.toString()</code> (page 140)	Returns the <code>Bulk.toJson()</code> (page 140) results as a string.

Bulk.insert()

Description

`Bulk.insert(<document>)`

New in version 2.6.

Adds an insert operation to a bulk operations list.

`Bulk.insert()` (page 129) accepts the following parameter:

param document doc Document to insert. The size of the document must be less than or equal to the maximum BSON document size (page 604).

Example The following initializes a `Bulk()` (page 128) operations builder for the `items` collection and adds a series of insert operations to add multiple documents:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.insert( { item: "abc123", defaultQty: 100, status: "A", points: 100 } );
bulk.insert( { item: "ijk123", defaultQty: 200, status: "A", points: 200 } );
bulk.insert( { item: "mop123", defaultQty: 0, status: "P", points: 0 } );
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.execute()` (page 137)

Bulk.find()

Description

`Bulk.find(<query>)`

New in version 2.6.

Specifies a query condition for an update or a remove operation.

`Bulk.find()` (page 130) accepts the following parameter:

param document query Specifies a query condition using *Query Selectors* (page 373) to select documents for an update or a remove operation.

With update operations, the sum of the query document and the update document must be less than or equal to the `maximum BSON document size` (page 604).

With remove operations, the query document must be less than or equal to the `maximum BSON document size` (page 604).

Use `Bulk.find()` (page 130) with the following write operations:

- `Bulk.find.removeOne()` (page 130)
- `Bulk.find.remove()` (page 131)
- `Bulk.find.replaceOne()` (page 131)
- `Bulk.find.updateOne()` (page 132)
- `Bulk.find.update()` (page 134)

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection and adds a remove operation and an update operation to the list of operations. The remove operation and the update operation use the `Bulk.find()` (page 130) method to specify a condition for their respective actions:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "D" } ).remove();
bulk.find( { status: "P" } ).update( { $set: { points: 0 } } )
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.execute()` (page 137)

Bulk.find.removeOne()

Description

`Bulk.find.removeOne()`

New in version 2.6.

Adds a single document remove operation to a bulk operations list. Use the `Bulk.find()` (page 130) method to specify the condition that determines which document to remove. The `Bulk.find.removeOne()` (page 130) limits the removal to one document. To remove multiple documents, see `Bulk.find.remove()` (page 131).

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection and adds two `Bulk.find.removeOne()` (page 130) operations to the list of operations.

Each remove operation removes just one document: one document with the `status` equal to "D" and another document with the `status` equal to "P".

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "D" } ).removeOne();
bulk.find( { status: "P" } ).removeOne();
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.find.remove()` (page 131)
- `Bulk.execute()` (page 137)
- *All Bulk Methods* (page 129)

Bulk.find.remove()

Description

`Bulk.find.remove()`

New in version 2.6.

Adds a remove operation to a bulk operations list. Use the `Bulk.find()` (page 130) method to specify the condition that determines which documents to remove. The `Bulk.find.remove()` (page 131) method removes all matching documents in the collection. To limit the remove to a single document, see `Bulk.find.removeOne()` (page 130).

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection and adds a remove operation to the list of operations. The remove operation removes all documents in the collection where the `status` equals "D":

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "D" } ).remove();
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.find.removeOne()` (page 130)
- `Bulk.execute()` (page 137)

Bulk.find.replaceOne()

Description

`Bulk.find.replaceOne(<document>)`

New in version 2.6.

Adds a single document replacement operation to a bulk operations list. Use the `Bulk.find()` (page 130) method to specify the condition that determines which document to replace. The `Bulk.find.replaceOne()` (page 131) method limits the replacement to a single document.

`Bulk.find.replaceOne()` (page 131) accepts the following parameter:

param document replacement A replacement document that completely replaces the existing document. Contains only field and value pairs.

The sum of the associated `<query>` document from the `Bulk.find()` (page 130) and the replacement document must be less than or equal to the `maximum BSON document size` (page 604).

To specify an *upsert* for this operation, see `Bulk.find.upsert()` (page 134).

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection, and adds various `replaceOne` (page 131) operations to the list of operations.

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { item: "abc123" } ).replaceOne( { item: "abc123", status: "P", points: 100 } );
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.execute()` (page 137)
- *All Bulk Methods* (page 129)

`Bulk.find.updateOne()`

Description

`Bulk.find.updateOne(<update>)`

New in version 2.6.

Adds a single document update operation to a bulk operations list. The operation can either replace an existing document or update specific fields in an existing document.

Use the `Bulk.find()` (page 130) method to specify the condition that determines which document to update. The `Bulk.find.updateOne()` (page 132) method limits the update or replacement to a single document. To update multiple documents, see `Bulk.find.update()` (page 134).

`Bulk.find.updateOne()` (page 132) accepts the following parameter:

param document update An update document that updates specific fields or a replacement document that completely replaces the existing document.

An update document only contains *update operator* (page 412) expressions. A replacement document contains only field and value pairs.

The sum of the associated `<query>` document from the `Bulk.find()` (page 130) and the update/replacement document must be less than or equal to the `maximum BSON document size`.

To specify an *upsert* for this operation, see `Bulk.find.upsert()` (page 134).

Behavior

Update Specific Fields If the <update> document contains only *update operator* (page 412) expressions, as in:

```
{
  $set: { status: "D" },
  points: { $inc: 2 }
}
```

Then, `Bulk.find.updateOne()` (page 132) updates only the corresponding fields, `status` and `points`, in the document.

Replace a Document If the <update> document contains only `field:value` expressions, as in:

```
{
  item: "TBD",
  points: 0,
  inStock: true,
  status: "I"
}
```

Then, `Bulk.find.updateOne()` (page 132) *replaces* the matching document with the <update> document with the exception of the `_id` field. The `Bulk.find.updateOne()` (page 132) method *does not* replace the `_id` value.

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection, and adds various `updateOne` (page 132) operations to the list of operations.

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "D" } ).updateOne( { $set: { status: "I", points: "0" } } );
bulk.find( { item: null } ).updateOne(
  {
    item: "TBD",
    points: 0,
    inStock: true,
    status: "I"
  }
);
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.find.update()` (page 134)
- `Bulk.execute()` (page 137)
- *All Bulk Methods* (page 129)

Bulk.find.update()

Description

`Bulk.find.update(<update>)`

New in version 2.6.

Adds a multi update operation to a bulk operations list. The method updates specific fields in existing documents.

Use the `Bulk.find()` (page 130) method to specify the condition that determines which documents to update. The `Bulk.find.update()` (page 134) method updates all matching documents. To specify a single document update, see `Bulk.find.updateOne()` (page 132).

`Bulk.find.update()` (page 134) accepts the following parameter:

param document update Specifies the fields to update. Only contains *update operator* (page 412) expressions.

The sum of the associated <query> document from the `Bulk.find()` (page 130) and the update document must be less than or equal to the `maximum BSON document size` (page 604).

To specify an *upsert* for this operation, see `Bulk.find.upsert()` (page 134). With `Bulk.find.upsert()` (page 134), if no documents match the `Bulk.find()` (page 130) query condition, the update operation inserts only a single document.

Example The following example initializes a `Bulk()` (page 128) operations builder for the `items` collection, and adds various multi update operations to the list of operations.

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "D" } ).update( { $set: { status: "I", points: "0" } } );
bulk.find( { item: null } ).update( { $set: { item: "TBD" } } );
bulk.execute();
```

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.find.updateOne()` (page 132)
- `Bulk.execute()` (page 137)
- *All Bulk Methods* (page 129)

Bulk.find.upsert()

Description

`Bulk.find.upsert()`

New in version 2.6.

Sets the optional *upsert* flag for an update or a replacement operation and has the following syntax:

```
Bulk.find(<query>).upsert().update(<update>);
Bulk.find(<query>).upsert().updateOne(<update>);
Bulk.find(<query>).upsert().replaceOne(<replacement>);
```

With the *upsert* flag, if no matching documents exist for the `Bulk.find()` (page 130) condition, then the update or the replacement operation performs an insert. If a matching document does exist, then the update or replacement operation performs the specified update or replacement.

Use `Bulk.find.upsert()` (page 134) with the following write operations:

- `Bulk.find.replaceOne()` (page 131)
- `Bulk.find.updateOne()` (page 132)
- `Bulk.find.update()` (page 134)

Behavior The following describe the insert behavior of various write operations when used in conjunction with `Bulk.find.upsert()` (page 134).

Insert for `Bulk.find.replaceOne()` The `Bulk.find.replaceOne()` (page 131) method accepts, as its parameter, a replacement document that only contains field and value pairs:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { item: "abc123" } ).upsert().replaceOne(
  {
    item: "abc123",
    status: "P",
    points: 100,
  }
);
bulk.execute();
```

If the replacement operation with the `Bulk.find.upsert()` (page 134) option performs an insert, the inserted document is the replacement document. If the replacement document does not specify an `_id` field, MongoDB adds the `_id` field:

```
{
  "_id" : ObjectId("52ded3b398ca567f5c97ac9e"),
  "item" : "abc123",
  "status" : "P",
  "points" : 100
}
```

Insert for `Bulk.find.updateOne()` The `Bulk.find.updateOne()` (page 132) method accepts, as its parameter, an <update> document that contains only field and value pairs or only *update operator* (page 412) expressions.

Field and Value Pairs If the <update> document contains only field and value pairs:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "P" } ).upsert().updateOne(
  {
    item: "TBD",
    points: 0,
    inStock: true,
    status: "I"
  }
);
bulk.execute();
```

Then, if the update operation with the `Bulk.find.upsert()` (page 134) option performs an insert, the inserted document is the <update> document. If the update document does not specify an `_id` field, MongoDB adds the `_id` field:

```
{
  "_id" : ObjectId("52ded5a898ca567f5c97ac9f"),
  "item" : "TBD",
  "points" : 0,
  "inStock" : true,
  "status" : "I"
}
```

Update Operator Expressions If the <update> document contains only *update operator* (page 412) expressions:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "P", item: null } ).upsert().updateOne(
  {
    $setOnInsert: { defaultQty: 0, inStock: true },
    $currentDate: { lastModified: true },
    $set: { points: "0" }
  }
);
bulk.execute();
```

Then, if the update operation with the `Bulk.find.upsert()` (page 134) option performs an insert, the update operation inserts a document with field and values from the <query> document of the `Bulk.find()` (page 130) method and then applies the specified update from the <update> document:

```
{
  "_id" : ObjectId("52ded68c98ca567f5c97aca0"),
  "item" : null,
  "status" : "P",
  "defaultQty" : 0,
  "inStock" : true,
  "lastModified" : ISODate("2014-01-21T20:20:28.786Z"),
  "points" : "0"
}
```

If neither the <query> document nor the <update> document specifies an `_id` field, MongoDB adds the `_id` field.

Insert for `Bulk.find.update()` When using `upsert()` (page 134) with the multiple document update method `Bulk.find.update()` (page 134), if no documents match the query condition, the update operation inserts a *single* document.

The `Bulk.find.update()` (page 134) method accepts, as its parameter, an <update> document that contains only *update operator* (page 412) expressions:

```
var bulk = db.items.initializeUnorderedBulkOp();
bulk.find( { status: "P" } ).upsert().update(
  {
    $setOnInsert: { defaultQty: 0, inStock: true },
    $currentDate: { lastModified: true },
    $set: { status: "I", points: "0" }
  }
);
```

```
);
bulk.execute();
```

Then, if the update operation with the `Bulk.find.upsert()` (page 134) option performs an insert, the update operation inserts a single document with the fields and values from the `<query>` document of the `Bulk.find()` (page 130) method and then applies the specified update from the `<update>` document:

```
{
  "_id": ObjectId("52ded81a98ca567f5c97aca1"),
  "status": "I",
  "defaultQty": 0,
  "inStock": true,
  "lastModified": ISODate("2014-01-21T20:27:06.691Z"),
  "points": "0"
}
```

If neither the `<query>` document nor the `<update>` document specifies an `_id` field, MongoDB adds the `_id` field.

See also:

- `db.collection.initializeUnorderedBulkOp()` (page 51)
- `db.collection.initializeOrderedBulkOp()` (page 51)
- `Bulk.find()` (page 130)
- `Bulk.execute()` (page 137)
- *All Bulk Methods* (page 129)

Bulk.execute()

Description

`Bulk.execute()`

New in version 2.6.

Executes the list of operations built by the `Bulk()` (page 128) operations builder.

`Bulk.execute()` (page 137) accepts the following parameter:

param document writeConcern Write concern document for the bulk operation as a whole. Omit to use default. For a standalone `mongod` (page 503) server, the write concern defaults to `{ w: 1 }`. With a replica set, the default write concern for a `mongod` (page 503) server is set as a replica set configuration option.

Returns A `BulkWriteResult` (page 189) object that contains the status of the operation.

After execution, you cannot re-execute the `Bulk()` (page 128) object without reinitializing. See `db.collection.initializeUnorderedBulkOp()` (page 51) and `db.collection.initializeOrderedBulkOp()` (page 51).

Example The following initializes a `Bulk()` (page 128) operations builder on the `items` collection, adds a series of write operations, and executes the operations:

```
var bulk = db.items.initializeOrderedBulkOp();
bulk.insert( { item: "abc123", status: "A", defaultQty: 500, points: 5 } );
bulk.insert( { item: "ijk123", status: "A", defaultQty: 100, points: 10 } );
```

```
bulk.find( { status: "D" } ).removeOne();
bulk.find( { status: "D" } ).update( { $set: { points: 0 } } );
bulk.execute();
```

The operation returns the following `BulkWriteResult()` (page 189) object:

```
BulkWriteResult({
  "writeErrors" : [ ],
  "writeConcernErrors" : [ ],
  "nInserted" : 2,
  "nUpserted" : 0,
  "nMatched" : 3,
  "nModified" : 3,
  "nRemoved" : 1,
  "upserted" : [ ]
})
```

For details on the return object, see `BulkWriteResult()` (page 189).

See

`Bulk()` (page 128) for a listing of methods available for bulk operations.

Bulk.getOperations()

`Bulk.getOperations()`

New in version 2.6.

Returns an array of write operations executed through `Bulk.execute()` (page 137). Only use after a `Bulk.execute()` (page 137). Calling `Bulk.getOperations()` (page 138) before you execute will result in an *incomplete* list.

Example The following initializes a `Bulk()` (page 128) operations builder on the `items` collection, adds a series of write operations, executes the operations, and then calls `getOperations()` (page 138) on the bulk builder object:

```
var bulk = db.items.initializeOrderedBulkOp();
bulk.insert( { item: "abc123", status: "A", defaultQty: 500, points: 5 } );
bulk.insert( { item: "ijk123", status: "A", defaultQty: 100, points: 10 } );
bulk.find( { status: "D" } ).removeOne();
bulk.find( { status: "D" } ).update( { $set: { points: 0 } } );
bulk.execute();
bulk.getOperations();
```

The `getOperations()` (page 138) returns the following array:

```
[
  {
    "originalZeroIndex" : 0,
    "batchType" : 1,
    "operations" : [
      {
        "_id" : ObjectId("5345aea9ac681b0cddcd0a87"),
        "item" : "abc123",
        "status" : "A",
        "defaultQty" : 500,
```

```

        "points" : 5
      },
      {
        "_id" : ObjectId("5345aea9ac681b0cddcd0a88"),
        "item" : "ijk123",
        "status" : "A",
        "defaultQty" : 100,
        "points" : 10
      }
    ]
  },
  {
    "originalZeroIndex" : 2,
    "batchType" : 3,
    "operations" : [
      {
        "q" : {
          "status" : "D"
        },
        "limit" : 1
      }
    ]
  },
  {
    "originalZeroIndex" : 3,
    "batchType" : 2,
    "operations" : [
      {
        "q" : {
          "status" : "D"
        },
        "u" : {
          "$set" : {
            "points" : 0
          }
        },
        "multi" : true,
        "upsert" : false
      }
    ]
  }
]

```

Returned Fields The array contains documents with the following fields:

originalZeroIndex

Specifies the order in which the operation was added to the bulk operations builder, based on a zero index; e.g. first operation added to the bulk operations builder will have `originalZeroIndex` (page 139) value of 0.

batchType

Specifies the write operations type.

batchType	Operation
1	Insert
2	Update
3	Remove

operations

Array of documents that contain the details of the operation.

See also:

`Bulk()` (page 128) and `Bulk.execute()` (page 137).

Bulk.toJson()

`Bulk.toJson()`

New in version 2.6.

Returns a JSON document that contains the number of operations and batches in the `Bulk()` (page 128) object.

Example The following initializes a `Bulk()` (page 128) operations builder on the `items` collection, adds a series of write operations, and calls `Bulk.toJson()` (page 140) on the bulk builder object.

```
var bulk = db.items.initializeOrderedBulkOp();
bulk.insert( { item: "abc123", status: "A", defaultQty: 500, points: 5 } );
bulk.insert( { item: "ijk123", status: "A", defaultQty: 100, points: 10 } );
bulk.find( { status: "D" } ).removeOne();
bulk.toJson();
```

The `Bulk.toJson()` (page 140) returns the following JSON document

```
{ "nInsertOps": 2, "nUpdateOps": 0, "nRemoveOps": 1, "nBatches": 2 }
```

See also:

`Bulk()` (page 128)

Bulk.toString()

`Bulk.toString()`

New in version 2.6.

Returns as a string a JSON document that contains the number of operations and batches in the `Bulk()` (page 128) object.

Example The following initializes a `Bulk()` (page 128) operations builder on the `items` collection, adds a series of write operations, and calls `Bulk.toString()` (page 140) on the bulk builder object.

```
var bulk = db.items.initializeOrderedBulkOp();
bulk.insert( { item: "abc123", status: "A", defaultQty: 500, points: 5 } );
bulk.insert( { item: "ijk123", status: "A", defaultQty: 100, points: 10 } );
bulk.find( { status: "D" } ).removeOne();
bulk.toString();
```

The `Bulk.toString()` (page 140) returns the following JSON document

```
{ "nInsertOps": 2, "nUpdateOps": 0, "nRemoveOps": 1, "nBatches": 2 }
```

See also:

`Bulk()` (page 128)

2.1.6 User Management

User Management Methods

Name	Description
<code>db.createUser()</code> (page 141)	Creates a new user.
<code>db.addUser()</code> (page 143)	Deprecated. Adds a user to a database, and allows administrators to configure the user's privileges.
<code>db.updateUser()</code> (page 145)	Updates user data.
<code>db.changeUserPassword()</code> (page 147)	Changes an existing user's password.
<code>db.removeUser()</code> (page 147)	Deprecated. Removes a user from a database.
<code>db.dropAllUsers()</code> (page 147)	Deletes all users associated with a database.
<code>db.dropUser()</code> (page 148)	Removes a single user.
<code>db.grantRolesToUser()</code> (page 148)	Grants a role and its privileges to a user.
<code>db.revokeRolesFromUser()</code> (page 150)	Removes a role from a user.
<code>db.getUser()</code> (page 151)	Returns information about the specified user.
<code>db.getUsers()</code> (page 151)	Returns information about all users associated with a database.

`db.createUser()`

Definition

`db.createUser` (*user*, *writeConcern*)

Creates a new user for the database where the method runs. `db.createUser()` (page 141) returns a *duplicate user* error if the user already exists on the database.

The `db.createUser()` (page 141) method has the following syntax:

field document user The document with authentication and access information about the user to create.

field document writeConcern The level of `write concern` for the creation operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The user document defines the user and has the following form:

```
{ user: "<name>",
  pwd: "<cleartext password>",
  customData: { <any information> },
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ]
}
```

The user document has the following fields:

field string user The name of the new user.

field string pwd The user's password. The `pwd` field is not required if you run `db.createUser()` (page 141) on the `$external` database to create users who have credentials stored externally to MongoDB.

any document customData Any arbitrary information.

field array roles The roles granted to the user.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.createUser()` (page 141) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.createUser()` (page 141) method wraps the `createUser` (page 250) command.

Behavior

Encryption `db.createUser()` (page 141) sends password to the MongoDB instance *without* encryption. To encrypt the password during transmission, use SSL.

External Credentials Users created on the `$external` database should have credentials stored externally to MongoDB, as, for example, with MongoDB Enterprise installations that use Kerberos.

Required Access You must have the `createUser` *action* on a database to create a new user on that database.

You must have the `grantRole` *action* on a role's database to grant the role to another user.

If you have the `userAdmin` or `userAdminAnyDatabase` role, or if you are authenticated using the *localhost exception*, you have those actions.

Examples The following `db.createUser()` (page 141) operation creates the `accountAdmin01` user on the `products` database.

```
use products
db.createUser( { "user" : "accountAdmin01",
  "pwd": "cleartext password",
  "customData" : { employeeId: 12345 },
  "roles" : [ { role: "clusterAdmin", db: "admin" },
    { role: "readAnyDatabase", db: "admin" },
    "readWrite"
  ] },
  { w: "majority" , wtimeout: 5000 } )
```

The operation gives `accountAdmin01` the following roles:

- the `clusterAdmin` and `readAnyDatabase` roles on the `admin` database
- the `readWrite` role on the `products` database

Create User with Roles The following operation creates `accountUser` in the `products` database and gives the user the `readWrite` and `dbAdmin` roles.

```
use products
db.createUser(
  {
    user: "accountUser",
    pwd: "password",
    roles: [ "readWrite", "dbAdmin" ]
  }
)
```

Create User Without Roles The following operation creates a user named `reportsUser` in the `admin` database but does not yet assign roles:

```
use admin
db.createUser(
  {
    user: "reportsUser",
    pwd: "password",
    roles: [ ]
  }
)
```

Create Administrative User with Roles The following operation creates a user named `appAdmin` in the `admin` database and gives the user `readWrite` access to the `config` database, which lets the user change certain settings for sharded clusters, such as to the balancer setting.

```
use admin
db.createUser(
  {
    user: "appAdmin",
    pwd: "password",
    roles:
      [
        "clusterAdmin",
        { db: "config", role: "readWrite" }
      ]
  }
)
```

db.addUser()

Deprecated since version 2.6: Use `db.createUser()` (page 141) and `db.updateUser()` (page 145) instead of `db.addUser()` (page 143) to add users to MongoDB.

In 2.6, MongoDB introduced a new model for user credentials and privileges, as described in <http://docs.mongodb.org/manualcore/security-introduction>. To use `db.addUser()` (page 143) on MongoDB 2.4, see `db.addUser()` in the version 2.4 of the [MongoDB Manual](#)¹¹.

Definition

¹¹<http://docs.mongodb.org/v2.4/reference/method/db.addUser>

`db.addUser (document)`

Adds a new user on the database where you run the method. The `db.addUser ()` (page 143) method takes a user document as its argument:

```
db.addUser(<user document>)
```

Specify a document that resembles the following as an argument to `db.addUser ()` (page 143):

```
{ user: "<name>",
  pwd: "<cleartext password>",
  customData: { <any information> },
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ],
  writeConcern: { <write concern> }
}
```

The `db.addUser ()` (page 143) user document has the following fields:

field string user The name of the new user.

field string pwd The user's password. The `pwd` field is not required if you run `db.addUser ()` (page 143) on the `$external` database to create users who have credentials stored externally to MongoDB.

any document customData Any arbitrary information.

field array roles The roles granted to the user.

field document writeConcern The level of `write concern` for the creation operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Users created on the `$external` database should have credentials stored externally to MongoDB, as, for example, with MongoDB Enterprise installations that use Kerberos.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.addUser ()` (page 143) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Considerations The `db.addUser ()` (page 143) method returns a *duplicate user* error if the user exists.

When interacting with 2.6 and later MongoDB instances, `db.addUser ()` (page 143) sends unencrypted passwords. To encrypt the password in transit use SSL.

In the 2.6 version of the shell, `db.addUser ()` (page 143) is backwards compatible with both the 2.4 version of `db.addUser()`¹² and the 2.2 version of `db.addUser()`¹³. In 2.6, for backwards compatibility `db.addUser ()` (page 143) creates users that approximate the privileges available in previous versions of MongoDB.

¹²<http://docs.mongodb.org/v2.4/reference/method/db.addUser>

¹³<http://docs.mongodb.org/v2.2/reference/method/db.addUser>

Example The following `db.addUser()` (page 143) method creates a user Carlos on the database where the command runs. The command gives Carlos the `clusterAdmin` and `readAnyDatabase` roles on the `admin` database and the `readWrite` role on the current database:

```
{ user: "Carlos",
  pwd: "cleartext password",
  customData: { employeeId: 12345 },
  roles: [
    { role: "clusterAdmin", db: "admin" },
    { role: "readAnyDatabase", db: "admin" },
    "readWrite"
  ],
  writeConcern: { w: "majority" , wtimeout: 5000 }
}
```

db.updateUser()

Definition

`db.updateUser` (*username*, *update*, *writeConcern*)

Updates the user's profile on the database on which you run the method. An update to a field **completely replaces** the previous field's values. This includes updates to the user's `roles` array.

Warning: When you update the `roles` array, you completely replace the previous array's values. To add or remove roles without replacing all the user's existing roles, use the `db.grantRolesToUser()` (page 148) or `db.revokeRolesFromUser()` (page 150) methods.

The `db.updateUser()` (page 145) method uses the following syntax:

```
db.updateUser(
  "<username>",
  {
    customData : { <any information> },
    roles : [
      { role: "<role>", db: "<database>" } | "<role>",
      ...
    ],
    pwd: "<cleartext password>"
  },
  writeConcern: { <write concern> }
)
```

The `db.updateUser()` (page 145) method has the following arguments.

param string username The name of the user to update.

param document update A document containing the replacement data for the user. This data completely replaces the corresponding data for the user.

field document writeConcern The level of write concern for the update operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The update document specifies the fields to update and their new values. All fields in the update document are optional, but *must* include at least one field.

The update document has the following fields:

field document customData Any arbitrary information.

field array roles The roles granted to the user. An update to the `roles` array overrides the previous array's values.

field string pwd The user's password.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.updateUser()` (page 145) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.updateUser()` (page 145) method wraps the `updateUser` (page 252) command.

Behavior `db.updateUser()` (page 145) sends password to the MongoDB instance *without* encryption. To encrypt the password during transmission, use SSL.

Required Access You must have access that includes the `revokeRole` *action* on all databases in order to update a user's `roles` array.

You must have the `grantRole` *action* on a role's database to add a role to a user.

To change another user's `pwd` or `customData` field, you must have the `changeAnyPassword` and `changeAnyCustomData` *actions* respectively on that user's database.

To modify your own password or custom data, you must have the `changeOwnPassword` and `changeOwnCustomData` *actions* respectively on the `cluster` resource.

Example Given a user `appClient01` in the `products` database with the following user info:

```
{
  "_id" : "products.appClient01",
  "user" : "appClient01",
  "db" : "products",
  "customData" : { "empID" : "12345", "badge" : "9156" },
  "roles" : [
    { "role" : "readWrite",
      "db" : "products"
    },
    { "role" : "read",
      "db" : "inventory"
    }
  ]
}
```

The following `db.updateUser()` (page 145) method **completely** replaces the user's `customData` and `roles` data:

```
use products
db.updateUser( "appClient01",
  {
    customData : { employeeId : "0x3039" },
    roles : [
```

```

        { role : "read", db : "assets" }
      ]
    }
  )
}

```

The user `appClient01` in the `products` database now has the following user information:

```

{
  "_id" : "products.appClient01",
  "user" : "appClient01",
  "db" : "products",
  "customData" : { "employeeId" : "0x3039" },
  "roles" : [
    { "role" : "read",
      "db" : "assets"
    }
  ]
}

```

`db.changeUserPassword()`

Definition

`db.changeUserPassword (username, password)`

Updates a user's password.

param string username Specifies an existing username with access privileges for this database.

param string password Specifies the corresponding password.

Example The following operation changes the reporting user's password to `SOh3TbYhx8ypJPxmt1oOfL`:

```
db.changeUserPassword("reporting", "SOh3TbYhx8ypJPxmt1oOfL")
```

`db.removeUser()`

Deprecated since version 2.6: Use `db.dropUser()` (page 148) instead of `db.removeUser()` (page 147)

Definition

`db.removeUser (username)`

Removes the specified username from the database.

The `db.removeUser()` (page 147) method has the following parameter:

param string username The database username.

`db.dropAllUsers()`

Definition

`db.dropAllUsers (writeConcern)`

Removes all users from the current database.

Warning: The `dropAllUsers` method removes all users from the database.

The `dropAllUsers` method takes the following arguments:

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `db.dropAllUsers()` (page 147) method wraps the `dropAllUsersFromDatabase` (page 254) command.

Required Access You must have the `dropUser` *action* on a database to drop a user from that database.

Example The following `db.dropAllUsers()` (page 147) operation drops every user from the `products` database.

```
use products
db.dropAllUsers( {w: "majority", wtimeout: 5000} )
```

The `n` field in the results document shows the number of users removed:

```
{ "n" : 12, "ok" : 1 }
```

db.dropUser()

Definition

`db.dropUser` (*username*, *writeConcern*)

Removes the user from the current database.

The `db.dropUser()` (page 148) method takes the following arguments:

param string username The name of the user to remove from the database.

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `db.dropUser()` (page 148) method wraps the `dropUser` (page 253) command.

Required Access You must have the `dropUser` *action* on a database to drop a user from that database.

Example The following `db.dropUser()` (page 148) operation drops the `accountAdmin01` user on the `products` database.

```
use products
db.dropUser("accountAdmin01", {w: "majority", wtimeout: 5000})
```

db.grantRolesToUser()

Definition

`db.grantRolesToUser` (*username*, *roles*, *writeConcern*)

Grants additional roles to a user.

The `grantRolesToUser` method uses the following syntax:


```
db.grantRolesToUser( "<username>", [ <roles> ], { <writeConcern> } )
```

The `grantRolesToUser` method takes the following arguments:

param string user The name of the user to whom to grant roles.

field array roles An array of additional roles to grant to the user.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.grantRolesToUser()` (page 148) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.grantRolesToUser()` (page 148) method wraps the `grantRolesToUser` (page 255) command.

Required Access You must have the `grantRole` *action* on a database to grant a role on that database.

Example Given a user `accountUser01` in the `products` database with the following roles:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  }
]
```

The following `grantRolesToUser()` operation gives `accountUser01` the `readWrite` role on the `products` database and the `read` role on the `stock` database.

```
use products
db.grantRolesToUser(
  "accountUser01",
  [ "readWrite" , { role: "read", db: "stock" } ],
  { w: "majority" , wtimeout: 4000 }
)
```

The user `accountUser01` in the `products` database now has the following roles:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  },
  { "role" : "read",
    "db" : "stock"
  },
  { "role" : "readWrite",
    "db" : "products"
  }
]
```

```
    }  
  ]
```

db.revokeRolesFromUser()

Definition

db.revokeRolesFromUser()

Removes a one or more roles from a user on the current database. The `db.revokeRolesFromUser()` (page 150) method uses the following syntax:

```
db.revokeRolesFromUser( "<username>", [ <roles> ], { <writeConcern> } )
```

The `revokeRolesFromUser` method takes the following arguments:

param string user The name of the user from whom to revoke roles.

field array roles The roles to remove from the user.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.revokeRolesFromUser()` (page 150) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.revokeRolesFromUser()` (page 150) method wraps the `revokeRolesFromUser` (page 256) command.

Required Access You must have the `revokeRole` *action* on a database to revoke a role on that database.

Example The `accountUser01` user in the `products` database has the following roles:

```
"roles" : [  
  { "role" : "assetsReader",  
    "db" : "assets"  
  },  
  { "role" : "read",  
    "db" : "stock"  
  },  
  { "role" : "readWrite",  
    "db" : "products"  
  }  
]
```

The following `db.revokeRolesFromUser()` (page 150) method removes the two of the user's roles: the `read` role on the `stock` database and the `readWrite` role on the `products` database, which is also the database on which the method runs:

```

use products
db.revokeRolesFromUser( "accountUser01",
                        [ { role: "read", db: "stock" }, "readWrite" ],
                        { w: "majority" }
                      )

```

The user `accountUser01` user in the `products` database now has only one remaining role:

```

"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  }
]

```

db.getUser()

Definition

db.getUser (*username*)

Returns user information for a specified user. Run this method on the user's database. The user must exist on the database on which the method runs.

The `db.getUser()` (page 151) method has the following parameter:

param string username The name of the user for which to retrieve information.

`db.getUser()` (page 151) wraps the `usersInfo` (page 257) command.

Required Access You must have the `viewUser` *action* on another user's database to view the other user's credentials.

You can view your own information.

Example The following sequence of operations returns information about the `appClient` user on the `accounts` database:

```

use accounts
db.getUser("appClient")

```

db.getUsers()

Definition

db.getUsers ()

Returns information for all the users in the database.

`db.getUsers()` (page 151) wraps the `usersInfo` (page 257) command.

Required Access You must have the `viewUser` *action* on another user's database to view the other user's credentials.

You can view your own information.

2.1.7 Role Management

Role Management Methods

Name	Description
<code>db.createRole()</code> (page 152)	Creates a role and specifies its privileges.
<code>db.updateRole()</code> (page 153)	Updates a user-defined role.
<code>db.dropRole()</code> (page 155)	Deletes a user-defined role.
<code>db.dropAllRoles()</code> (page 156)	Deletes all user-defined roles associated with a database.
<code>db.grantPrivilegesToRole()</code> (page 156)	Assigns privileges to a user-defined role.
<code>db.revokePrivilegesFromRole()</code> (page 158)	Removes the specified privileges from a user-defined role.
<code>db.grantRolesToRole()</code> (page 160)	Specifies roles from which a user-defined role inherits privileges.
<code>db.revokeRolesFromRole()</code> (page 161)	Removes a role from a user.
<code>db.getRole()</code> (page 162)	Returns information for the specified role.
<code>db.getRoles()</code> (page 163)	Returns information for all the user-defined roles in a database.

`db.createRole()`

Definition

`db.createRole` (*role*, *writeConcern*)

Creates a role and specifies its *privileges*. The role applies to the database on which you run the method. The `db.createRole()` (page 152) method returns a *duplicate role* error if the role already exists in the database.

The `db.createRole()` (page 152) method takes the following arguments:

param document role A document containing the name of the role and the role definition.

field document writeConcern The level of `write concern` to apply to this operation. The `writeConcern` document uses the same fields as the `getLastError` (page 235) command.

The role document has the following form:

```
{ role: "<name>",
  privileges: [
    { resource: { <resource> }, actions: [ "<action>", ... ] },
    ...
  ],
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ]
}
```

The role document has the following fields:

field string role The name of the new role.

field array privileges The privileges to grant the role. A privilege consists of a resource and permitted actions. You must specify the `privileges` field. Use an empty array to specify *no* privileges. For the syntax of a privilege, see the `privileges` array.

field array roles An array of roles from which this role inherits privileges. You must specify the `roles` field. Use an empty array to specify *no* roles.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.createRole()` (page 152) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.createRole()` (page 152) method wraps the `createRole` (page 259) command.

Behavior A role's privileges apply to the database where the role is created. The role can inherit privileges from other roles in its database. A role created on the `admin` database can include privileges that apply to all databases or to the *cluster* and can inherit privileges from roles in other databases.

Required Access You must have the `createRole` *action* on a database to create a role on that database.

You must have the `grantRole` *action* on the database that a privilege targets in order to grant that privilege to a role. If the privilege targets multiple databases or the `cluster` resource, you must have the `grantRole` action on the `admin` database.

You must have the `grantRole` *action* on a role's database to grant the role to another role.

Example The following `db.createRole()` (page 152) method creates the `myClusterwideAdmin` role on the `admin` database:

```
use admin
db.createRole({ role: "myClusterwideAdmin",
  privileges: [
    { resource: { cluster: true }, actions: [ "addShard" ] },
    { resource: { db: "config", collection: "" }, actions: [ "find", "update", "insert", "remove" ] },
    { resource: { db: "users", collection: "usersCollection" }, actions: [ "update", "insert", "remove" ] },
    { resource: { db: "", collection: "" }, actions: [ "find" ] }
  ],
  roles: [
    { role: "read", db: "admin" }
  ],
  writeConcern: { w: "majority", wtimeout: 5000 }
})
```

db.updateRole()

Definition

`db.updateRole(rolename, update, writeConcern)`

Updates a *user-defined role*. The `db.updateRole()` (page 153) method must run on the role's database.

An update to a field **completely replaces** the previous field's values. To grant or remove roles or *privileges* without replacing all values, use one or more of the following methods:

- `db.grantRolesToRole()` (page 160)
- `db.grantPrivilegesToRole()` (page 156)
- `db.revokeRolesFromRole()` (page 161)

•`db.revokePrivilegesFromRole()` (page 158)

Warning: An update to the `privileges` or `roles` array completely replaces the previous array's values.

The `updateRole()` method uses the following syntax:

```
db.updateRole(  
  "<rolename>",  
  {  
    privileges:  
      [  
        { resource: { <resource> }, actions: [ "<action>", ... ] },  
        ...  
      ],  
    roles:  
      [  
        { role: "<role>", db: "<database>" } | "<role>",  
        ...  
      ]  
  },  
  { <writeConcern> }  
)
```

The `db.updateRole()` (page 153) method takes the following arguments.

param string rolename The name of the *user-defined role* to update.

param document update A document containing the replacement data for the role. This data completely replaces the corresponding data for the role.

field document writeConcern The level of write concern for the update operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `update` document specifies the fields to update and the new values. Each field in the `update` document is optional, but the document must include at least one field. The `update` document has the following fields:

field array privileges Required if you do not specify `roles` array. The privileges to grant the role. An update to the `privileges` array overrides the previous array's values. For the syntax for specifying a privilege, see the `privileges` array.

field array roles Required if you do not specify `privileges` array. The roles from which this role inherits privileges. An update to the `roles` array overrides the previous array's values.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.updateRole()` (page 153) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.updateRole()` (page 153) method wraps the `updateRole` (page 260) command.

Behavior A role's privileges apply to the database where the role is created. The role can inherit privileges from other roles in its database. A role created on the `admin` database can include privileges that apply to all databases or to the `cluster` and can inherit privileges from roles in other databases.

Required Access You must have the `revokeRole` *action* on all databases in order to update a role.

You must have the `grantRole` *action* on the database of each role in the `roles` array to update the array.

You must have the `grantRole` *action* on the database of each privilege in the `privileges` array to update the array. If a privilege's resource spans databases, you must have `grantRole` on the `admin` database. A privilege spans databases if the privilege is any of the following:

- a collection in all databases
- all collections and all database
- the `cluster` resource

Example The following `db.updateRole()` (page 153) method replaces the `privileges` and the `roles` for the `inventoryControl` role that exists in the `products` database. The method runs on the database that contains `inventoryControl`:

```
use products
db.updateRole(
  "inventoryControl",
  {
    privileges:
      [
        {
          resource: { db:"products", collection:"clothing" },
          actions: [ "update", "createCollection", "createIndex" ]
        }
      ],
    roles:
      [
        {
          role: "read",
          db: "products"
        }
      ]
  },
  { w:"majority" }
)
```

To view a role's privileges, use the `rolesInfo` (page 270) command.

db.dropRole()

Definition

`db.dropRole (rolename, writeConcern)`

Deletes a *user-defined* role from the database on which you run the method.

The `db.dropRole()` (page 155) method takes the following arguments:

param string rolename The name of the *user-defined role* to remove from the database.

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `db.dropRole()` (page 155) method wraps the `dropRole` (page 262) command.

Required Access You must have the `dropRole` *action* on a database to drop a role from that database.

Example The following operations remove the `readPrices` role from the `products` database:

```
use products
db.dropRole( "readPrices", { w: "majority" } )
```

`db.dropAllRoles()`

Definition

`db.dropAllRoles` (*writeConcern*)

Deletes all *user-defined* roles on the database where you run the method.

Warning: The `dropAllRoles` method removes *all user-defined* roles from the database.

The `dropAllRoles` method takes the following argument:

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Returns The number of *user-defined* roles dropped.

The `db.dropAllRoles()` (page 156) method wraps the `dropAllRolesFromDatabase` (page 263) command.

Required Access You must have the `dropRole` *action* on a database to drop a role from that database.

Example The following operations drop all *user-defined* roles from the `products` database and uses a *write concern* of `majority`.

```
use products
db.dropAllRoles( { w: "majority" } )
```

The method returns the number of roles dropped:

4

`db.grantPrivilegesToRole()`

Definition

`db.grantPrivilegesToRole` (*rolename, privileges, writeConcern*)

Grants additional *privileges* to a *user-defined* role.

The `grantPrivilegesToRole()` method uses the following syntax:


```

db.grantPrivilegesToRole(
  "<rolename>",
  [
    { resource: { <resource> }, actions: [ "<action>", ... ] },
    ...
  ],
  { < writeConcern > }
)

```

The `grantPrivilegesToRole()` method takes the following arguments:

param string rolename The name of the role to grant privileges to.

field array privileges The privileges to add to the role. For the format of a privilege, see [privileges](#).

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `grantPrivilegesToRole()` method can grant one or more privileges. Each `<privilege>` has the following syntax:

```
{ resource: { <resource> }, actions: [ "<action>", ... ] }
```

The `db.grantPrivilegesToRole()` (page 156) method wraps the `grantPrivilegesToRole` (page 264) command.

Behavior A role's privileges apply to the database where the role is created. A role created on the `admin` database can include privileges that apply to all databases or to the *cluster*.

Required Access You must have the `grantRole` action on the database a privilege targets in order to grant the privilege. To grant a privilege on multiple databases or on the `cluster` resource, you must have the `grantRole` action on the `admin` database.

Example The following `db.grantPrivilegesToRole()` (page 156) operation grants two additional privileges to the role `inventoryCntrl01`, which exists on the `products` database. The operation is run on that database:

```

use products
db.grantPrivilegesToRole(
  "inventoryCntrl01",
  [
    {
      resource: { db: "products", collection: "" },
      actions: [ "insert" ]
    },
    {
      resource: { db: "products", collection: "system.indexes" },
      actions: [ "find" ]
    }
  ],
  { w: "majority" }
)

```

The first privilege permits users with this role to perform the `insert` *action* on all collections of the `products` database, except the *system collections* (page 600). To access a system collection, a privilege must explicitly specify the system collection in the resource document, as in the second privilege.

The second privilege permits users with this role to perform the `find` *action* on the `product` database's system collection named `system.indexes` (page 601).

`db.revokePrivilegesFromRole()`

Definition

`db.revokePrivilegesFromRole` (*rolename*, *privileges*, *writeConcern*)

Removes the specified privileges from the *user-defined* role on the database where the method runs. The `revokePrivilegesFromRole` method has the following syntax:

```
db.revokePrivilegesFromRole(  
  "<rolename>",  
  [  
    { resource: { <resource> }, actions: [ "<action>", ... ] },  
    ...  
  ],  
  { <writeConcern> }  
)
```

The `revokePrivilegesFromRole` method takes the following arguments:

param string rolename The name of the *user-defined* role from which to revoke privileges.

field array privileges An array of privileges to remove from the role. See `privileges` for more information on the format of the privileges.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

The `db.revokePrivilegesFromRole()` (page 158) method wraps the `revokePrivilegesFromRole` (page 265) command.

Behavior To revoke a privilege, the `resource` document pattern must match **exactly** the `resource` field of that privilege. The `actions` field can be a subset or match exactly.

For example, given the role `accountRole` in the `products` database with the following privilege that specifies the `products` database as the resource:

```
{  
  "resource" : {  
    "db" : "products",  
    "collection" : ""  
  },  
  "actions" : [  
    "find",  
    "update"  
  ]  
}
```

You *cannot* revoke `find` and/or `update` from just *one* collection in the `products` database. The following operations result in no change to the role:

```

use products
db.revokePrivilegesFromRole(
  "accountRole",
  [
    {
      resource : {
        db : "products",
        collection : "gadgets"
      },
      actions : [
        "find",
        "update"
      ]
    }
  ]
)

db.revokePrivilegesFromRole(
  "accountRole",
  [
    {
      resource : {
        db : "products",
        collection : "gadgets"
      },
      actions : [
        "find"
      ]
    }
  ]
)

```

To revoke the "find" and/or the "update" action from the role `accountRole`, you must match the resource document exactly. For example, the following operation revokes just the "find" action from the existing privilege.

```

use products
db.revokePrivilegesFromRole(
  "accountRole",
  [
    {
      resource : {
        db : "products",
        collection : ""
      },
      actions : [
        "find"
      ]
    }
  ]
)

```

Required Access You must have the `revokeRole` *action* on the database a privilege targets in order to revoke that privilege. If the privilege targets multiple databases or the `cluster` resource, you must have the `revokeRole` action on the `admin` database.

Example The following operation removes multiple privileges from the `associates` role:

```
db.revokePrivilegesFromRole(
  "associate",
  [
    {
      resource: { db: "products", collection: "" },
      actions: [ "createCollection", "createIndex", "find" ]
    },
    {
      resource: { db: "products", collection: "orders" },
      actions: [ "insert" ]
    }
  ],
  { w: "majority" }
)
```

db.grantRolesToRole()

Definition

db.grantRolesToRole (*rolename*, *roles*, *writeConcern*)

Grants roles to a *user-defined* role.

The `grantRolesToRole` method uses the following syntax:

```
db.grantRolesToRole( "<rolename>", [ <roles> ], { <writeConcern> } )
```

The `grantRolesToRole` method takes the following arguments:

param string rolename The name of the role to which to grant sub roles.

field array roles An array of roles from which to inherit.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.grantRolesToRole()` (page 160) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.grantRolesToRole()` (page 160) method wraps the `grantRolesToRole` (page 267) command.

Behavior A role can inherit privileges from other roles in its database. A role created on the `admin` database can inherit privileges from roles in any database.

Required Access You must have the `grantRole` *action* on a database to grant a role on that database.

Example The following `grantRolesToRole()` operation updates the `productsReaderWriter` role in the `products` database to *inherit* the *privileges* of `productsReader` role:

```
use products
db.grantRolesToRole(
  "productsReaderWriter",
  [ "productsReader" ],
  { w: "majority" , wtimeout: 5000 }
)
```

db.revokeRolesFromRole()

Definition

`db.revokeRolesFromRole (rolename, roles, writeConcern)`

Removes the specified inherited roles from a role.

The `revokeRolesFromRole` method uses the following syntax:

```
db.revokeRolesFromRole( "<rolename>", [ <roles> ], { <writeConcern> } )
```

The `revokeRolesFromRole` method takes the following arguments:

param string rolename The name of the role from which to revoke roles.

field array roles The inherited roles to remove.

field document writeConcern The level of `write concern` to apply to this operation. The `writeConcern` document uses the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `db.revokeRolesFromRole()` (page 161) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

The `db.revokeRolesFromRole()` (page 161) method wraps the `revokeRolesFromRole` (page 268) command.

Required Access You must have the `revokeRole` *action* on a database to revoke a role on that database.

Example The `purchaseAgents` role in the `emea` database inherits privileges from several other roles, as listed in the `roles` array:

```
{
  "_id" : "emea.purchaseAgents",
  "role" : "purchaseAgents",
  "db" : "emea",
  "privileges" : [],
  "roles" : [
    {
      "role" : "readOrdersCollection",
      "db" : "emea"
    }
  ]
}
```

```
    },
    {
      "role" : "readAccountsCollection",
      "db" : "emea"
    },
    {
      "role" : "writeOrdersCollection",
      "db" : "emea"
    }
  ]
}
```

The following `db.revokeRolesFromRole()` (page 161) operation on the `emea` database removes two roles from the `purchaseAgents` role:

```
use emea
db.revokeRolesFromRole( "purchaseAgents",
  [
    "writeOrdersCollection",
    "readOrdersCollection"
  ],
  { w: "majority" , wtimeout: 5000 }
)
```

The `purchaseAgents` role now contains just one role:

```
{
  "_id" : "emea.purchaseAgents",
  "role" : "purchaseAgents",
  "db" : "emea",
  "privileges" : [],
  "roles" : [
    {
      "role" : "readAccountsCollection",
      "db" : "emea"
    }
  ]
}
```

db.getRole()

Definition

`db.getRole(rolename, showPrivileges)`

Returns the roles from which this role inherits privileges. Optionally, the method can also return all the role's privileges.

Run `db.getRole()` (page 162) from the database that contains the role. The command can retrieve information for both *user-defined roles* and *built-in roles*.

The `db.getRole()` (page 162) method takes the following arguments:

param string rolename The name of the role.

param document showPrivileges If `true`, returns the role's privileges. Pass this argument as a document: `{showPrivileges: true}`.

`db.getRole()` (page 162) wraps the `rolesInfo` (page 270) command.

Required Access To view a role's information, you must be explicitly granted the role or must have the `viewRole` *action* on the role's database.

Examples The following operation returns role inheritance information for the role `associate` defined on the `products` database:

```
use products
db.getRole( "associate" )
```

The following operation returns role inheritance information *and privileges* for the role `associate` defined on the `products` database:

```
use products
db.getRole( "associate", { showPrivileges: true } )
```

`db.getRoles()`

Definition

`db.getRoles()`

Returns information for all the roles in the database on which the command runs. The method can be run with or without an argument.

If run without an argument, `db.getRoles()` (page 163) returns inheritance information for the database's *user-defined* roles.

To return more information, pass the `db.getRoles()` (page 163) a document with the following fields:

field integer rolesInfo Set this field to 1 to retrieve all user-defined roles.

field Boolean showBuiltinRoles Set to true to display *built-in roles* as well as user-defined roles.

field Boolean showPrivileges Set the field to `true` to show role privileges, including both privileges inherited from other roles and privileges defined directly. By default, the command returns only the roles from which this role inherits privileges and does not return specific privileges.

`db.getRoles()` (page 163) wraps the `rolesInfo` (page 270) command.

Required Access To view a role's information, you must be explicitly granted the role or must have the `viewRole` *action* on the role's database.

Example The following operations return documents for all the roles on the `products` database, including role privileges and built-in roles:

```
db.getRoles(
  {
    rolesInfo: 1,
    showPrivileges: true,
    showBuiltinRoles: true
  }
)
```

2.1.8 Replication

Replication Methods

Name	Description
<code>rs.add()</code> (page 164)	Adds a member to a replica set.
<code>rs.addArb()</code> (page 165)	Adds an <i>arbiter</i> to a replica set.
<code>rs.conf()</code> (page 165)	Returns the replica set configuration document.
<code>rs.freeze()</code> (page 165)	Prevents the current member from seeking election as primary for a period of time.
<code>rs.help()</code> (page 166)	Returns basic help text for <i>replica set</i> functions.
<code>rs.initiate()</code> (page 166)	Initializes a new replica set.
<code>rs.printReplicationInfo()</code> (page 166)	Prints a report of the status of the replica set from the perspective of the primary.
<code>rs.printSlaveReplicationInfo()</code> (page 166)	Prints a report of the status of the replica set from the perspective of the secondaries.
<code>rs.reconfig()</code> (page 167)	Re-configures a replica set by applying a new replica set configuration object.
<code>rs.remove()</code> (page 168)	Remove a member from a replica set.
<code>rs.slaveOk()</code> (page 168)	Sets the <code>slaveOk</code> property for the current connection. Deprecated. Use <code>readPref()</code> (page 91) and <code>Mongo.setReadPref()</code> (page 192) to set <i>read preference</i> .
<code>rs.status()</code> (page 168)	Returns a document with information about the state of the replica set.
<code>rs.stepDown()</code> (page 168)	Causes the current <i>primary</i> to become a secondary which forces an <i>election</i> .
<code>rs.syncFrom()</code> (page 169)	Sets the member that this replica set member will sync from, overriding the default sync target selection logic.

`rs.add()`

Definition

`rs.add(host, arbiterOnly)`

Adds a member to a *replica set*.

param string,document host The new member to add to the replica set. If a string, specify the hostname and optionally the port number for the new member. If a document, specify a replica-set members document, as found in the `members` array. To view a replica set's members array, run `rs.conf()` (page 165).

param boolean arbiterOnly Applies only if the `<host>` value is a string. If `true`, the added host is an arbiter."

You may specify new hosts in one of two ways:

- 1.as a "hostname" with an optional port number to use the default configuration as in the *replica-set-add-member* example.
- 2.as a configuration *document*, as in the *replica-set-add-member-alternate-procedure* example.

`rs.add()` (page 164) provides a wrapper around some of the functionality of the "`replSetReconfig` (page 276)" *database command* and the corresponding shell helper `rs.reconfig()` (page 167). See the <http://docs.mongodb.org/manualreference/replica-configuration> document for full documentation of all replica set configuration options.

Behavior `rs.add()` (page 164) can in some cases force an election for primary which will disconnect the shell. In such cases, the shell displays an error even if the operation succeeds.

Example To add a `mongod` (page 503) accessible on the default port 27017 running on the host `mongodb3.example.net`, use the following `rs.add()` (page 164) invocation:

```
rs.add('mongodb3.example.net:27017')
```

If `mongodb3.example.net` is an arbiter, use the following form:

```
rs.add('mongodb3.example.net:27017', true)
```

To add `mongodb3.example.net` as a *secondary-only* member of set, use the following form of `rs.add()` (page 164):

```
rs.add( { "_id": 3, "host": "mongodbd3.example.net:27017", "priority": 0 } )
```

Replace, 3 with the next unused `_id` value in the replica set. See `rs.conf()` (page 165) to see the existing `_id` values in the replica set configuration document.

See the <http://docs.mongodb.org/manualreference/replica-configuration> and <http://docs.mongodb.org/manualadministration/replica-sets> documents for more information.

`rs.addArb()`

Description

`rs.addArb(host)`

Adds a new *arbiter* to an existing replica set.

The `rs.addArb()` (page 165) method takes the following parameter:

param string host Specifies the hostname and optionally the port number of the arbiter member to add to replica set.

This function briefly disconnects the shell and forces a reconnection as the replica set renegotiates which member will be *primary*. As a result, the shell displays an error even if this command succeeds.

`rs.conf()`

`rs.conf()`

Returns a *document* that contains the current *replica set* configuration document.

See <http://docs.mongodb.org/manualreference/replica-configuration> for more information on the replica set configuration document.

`rs.config()`

`rs.config()` (page 165) is an alias of `rs.conf()` (page 165).

`rs.freeze()`

Description

`rs.freeze(seconds)`

Makes the current *replica set* member ineligible to become *primary* for the period specified.

The `rs.freeze()` (page 165) method has the following parameter:

param number seconds The duration the member is ineligible to become primary.

`rs.freeze()` (page 165) provides a wrapper around the *database command* `replSetFreeze` (page 273).

rs.help()

`rs.help()`

Returns a basic help text for all of the *replication* related shell functions.

rs.initiate()

Description

`rs.initiate()` (*configuration*)

Initiates a *replica set*. Optionally takes a configuration argument in the form of a *document* that holds the configuration of a replica set.

The `rs.initiate()` (page 166) method has the following parameter:

param document configuration A *document* that specifies configuration settings for the new replica set. If a configuration is not specified, MongoDB uses a default configuration.

The `rs.initiate()` (page 166) method provides a wrapper around the “`replSetInitiate` (page 275)” *database command*.

Replica Set Configuration See <http://docs.mongodb.org/manualadministration/replica-set-member-conf> and <http://docs.mongodb.org/manualreference/replica-configuration> for examples of replica set configuration and invitation objects.

rs.printReplicationInfo()

`rs.printReplicationInfo()`

New in version 2.6.

Returns a formatted report of the status of a *replica set* from the perspective of the *primary* member of the set. The output is identical to that of `db.printReplicationInfo()` (page 116). See the *replSetGetStatus* (page 273) for more information regarding the contents of this output.

Note: The `rs.printReplicationInfo()` (page 166) in the *mongo* (page 527) shell does **not** return *JSON*. Use `rs.printReplicationInfo()` (page 166) for manual inspection, and `rs.status()` (page 168) in scripts.

rs.printSlaveReplicationInfo()

Definition

`rs.printSlaveReplicationInfo()`

Returns a formatted report of the status of a *replica set* from the perspective of the *secondary* member of the set. The output is identical to that of `db.printSlaveReplicationInfo()` (page 116).

Output The following is example output from the `rs.printSlaveReplicationInfo()` (page 166) method issued on a replica set with two secondary members:

```
source: m1.example.net:27017
  syncedTo: Thu Apr 10 2014 10:27:47 GMT-0400 (EDT)
  0 secs (0 hrs) behind the primary
source: m2.example.net:27017
  syncedTo: Thu Apr 10 2014 10:27:47 GMT-0400 (EDT)
  0 secs (0 hrs) behind the primary
```

A *delayed member* may show as 0 seconds behind the primary when the inactivity period on the primary is greater than the `slaveDelay` value.

rs.reconfig()

Definition

`rs.reconfig(configuration, force)`

Initializes a new *replica set* configuration. Disconnects the shell briefly and forces a reconnection as the replica set renegotiates which member will be *primary*. As a result, the shell will display an error even if this command succeeds.

param document configuration A *document* that specifies the configuration of a replica set.

param document force “If set as `{ force: true }`, this forces the replica set to accept the new configuration even if a majority of the members are not accessible. Use with caution, as this can lead to *term:rollback* situations.”

`rs.reconfig()` (page 167) overwrites the existing replica set configuration. Retrieve the current configuration object with `rs.conf()` (page 165), modify the configuration as needed and then use `rs.reconfig()` (page 167) to submit the modified configuration object.

`rs.reconfig()` (page 167) provides a wrapper around the “*replSetReconfig* (page 276)” *database command*.

Examples To reconfigure a replica set, use the following sequence of operations:

```
conf = rs.conf()

// modify conf to change configuration

rs.reconfig(conf)
```

If you want to force the reconfiguration if a majority of the set is not connected to the current member, or you are issuing the command against a secondary, use the following form:

```
conf = rs.conf()

// modify conf to change configuration

rs.reconfig(conf, { force: true } )
```

Warning: Forcing a `rs.reconfig()` (page 167) can lead to *rollback* situations and other difficult to recover from situations. Exercise caution when using this option.

See also:

<http://docs.mongodb.org/manualreference/replica-configuration>
<http://docs.mongodb.org/manualadministration/replica-sets>.

and

rs.remove()

Definition

rs.remove (*hostname*)

Removes the member described by the *hostname* parameter from the current *replica set*. This function will disconnect the shell briefly and forces a reconnection as the *replica set* renegotiates which member will be *primary*. As a result, the shell will display an error even if this command succeeds.

The **rs.remove()** (page 168) method has the following parameter:

param string hostname The hostname of a system in the replica set.

Note: Before running the **rs.remove()** (page 168) operation, you must *shut down* the replica set member that you're removing.

Changed in version 2.2: This procedure is no longer required when using **rs.remove()** (page 168), but it remains good practice.

rs.slaveOk()

rs.slaveOk()

Provides a shorthand for the following operation:

```
db.getMongo().setSlaveOk()
```

This allows the current connection to allow read operations to run on *secondary* members. See the **readPref()** (page 91) method for more fine-grained control over read preference in the **mongo** (page 527) shell.

rs.status()

rs.status()

Returns A *document* with status information.

This output reflects the current status of the replica set, using data derived from the heartbeat packets sent by the other members of the replica set.

This method provides a wrapper around the **replSetGetStatus** (page 273) *database command*. See the documentation of the command for a complete description of the *output* (page 274).

rs.stepDown()

Description

rs.stepDown (*seconds*)

Forces the current *replica set* member to step down as *primary* and then attempt to avoid election as primary for the designated number of seconds. Produces an error if the current member is not the primary.

The **rs.stepDown()** (page 168) method has the following parameter:

param number seconds The duration of time that the stepped-down member attempts to avoid re-election as primary. If this parameter is not specified, the method uses the default value of 60 seconds.

This function disconnects the shell briefly and forces a reconnection as the replica set renegotiates which member will be primary. As a result, the shell will display an error even if this command succeeds.

`rs.stepDown()` (page 168) provides a wrapper around the *database command* `replSetStepDown` (page 277).

rs.syncFrom()

`rs.syncFrom()`

New in version 2.2.

Provides a wrapper around the `replSetSyncFrom` (page 278), which allows administrators to configure the member of a replica set that the current member will pull data from. Specify the name of the member you want to replicate from in the form of `[hostname]:[port]`.

See `replSetSyncFrom` (page 278) for more details.

2.1.9 Sharding

Sharding Methods

Name	Description
<code>sh._adminCommand</code> (page 172)	Runs a <i>database command</i> against the admin database, like <code>db.runCommand()</code> (page 118), but can confirm that it is issued against a <i>mongos</i> (page 518).
<code>sh._checkFullName()</code> (page 172)	Tests a namespace to determine if its well formed.
<code>sh._checkMongos()</code> (page 172)	Tests to see if the <i>mongo</i> (page 527) shell is connected to a <i>mongos</i> (page 518) instance.
<code>sh._lastMigration()</code> (page 172)	Reports on the last <i>chunk</i> migration.
<code>sh.addShard()</code> (page 173)	Adds a <i>shard</i> to a sharded cluster.
<code>sh.addShardTag()</code> (page 174)	Associates a shard with a tag, to support tag aware sharding.
<code>sh.addTagRange()</code> (page 174)	Associates range of shard keys with a shard tag, to support tag aware sharding.
<code>sh.disableBalancing()</code> (page 175)	Disable balancing on a single collection in a sharded database. Does not affect balancing of other collections in a sharded cluster.
<code>sh.enableBalancing()</code> (page 175)	Activates the sharded collection balancer process if previously disabled using <code>sh.disableBalancing()</code> (page 175).
<code>sh.enableSharding()</code> (page 175)	Enables sharding on a specific database.
<code>sh.getBalancerHost()</code> (page 176)	Returns the name of a <i>mongos</i> (page 518) that's responsible for the balancer process.
<code>sh.getBalancerState()</code> (page 176)	Returns a boolean to report if the <i>balancer</i> is currently enabled.
<code>sh.help()</code> (page 176)	Returns help text for the <i>sh</i> methods.
<code>sh.isBalancerRunning()</code> (page 177)	Returns a boolean to report if the balancer process is currently migrating chunks.
<code>sh.moveChunk()</code> (page 177)	Migrates a <i>chunk</i> in a <i>sharded cluster</i> .
<code>sh.removeShardTag()</code> (page 178)	Removes the association between a shard and a shard tag <i>shard tag</i> .
<code>sh.setBalancerState()</code> (page 178)	Enables or disables the <i>balancer</i> which migrates <i>chunks</i> between <i>shards</i> .
<code>sh.shardCollection()</code> (page 178)	Enables sharding for a collection.
<code>sh.splitAt()</code> (page 179)	Divides an existing <i>chunk</i> into two chunks using a specific value of the <i>shard key</i> as the dividing point.
<code>sh.splitFind()</code> (page 179)	Divides an existing <i>chunk</i> that contains a document matching a query into two approximately equal chunks.
<code>sh.startBalancer()</code> (page 180)	Enables the <i>balancer</i> and waits for balancing to start.
<code>sh.status()</code> (page 180)	Reports on the status of a <i>sharded cluster</i> , as <code>db.printShardingStatus()</code> (page 116).
<code>sh.stopBalancer()</code> (page 182)	Disables the <i>balancer</i> and waits for any in progress balancing rounds to complete.
<code>sh.waitForBalancer()</code> (page 183)	Internal. Waits for the balancer state to change.
<code>sh.waitForBalancerOn()</code> (page 183)	Internal. Waits until the balancer stops running.
<code>sh.waitForDLock()</code> (page 184)	Internal. Waits for a specified distributed <i>sharded cluster</i> lock.
2.1.10 mongo Shell Methods	
<code>sh.waitForPingChange()</code> (page 184)	Internal. Waits for a change in ping state from one of the <i>mongos</i> (page 518) in the <i>sharded cluster</i> .

sh._adminCommand()

Definition

sh._adminCommand (command, checkMongos)

Runs a database command against the admin database of a [mongos](#) (page 518) instance.

param string command A database command to run against the admin database.

param boolean checkMongos Require verification that the shell is connected to a [mongos](#) (page 518) instance.

See also:

[db.runCommand\(\)](#) (page 118)

sh._checkFullName()

Definition

sh._checkFullName (namespace)

Verifies that a *namespace* name is well formed. If the namespace is well formed, the [sh._checkFullName\(\)](#) (page 172) method exits *with no message*.

Throws If the namespace is not well formed, [sh._checkFullName\(\)](#) (page 172) throws: “name needs to be fully qualified <db>.<collection>”

The [sh._checkFullName\(\)](#) (page 172) method has the following parameter:

param string namespace The *namespace* of a collection. The namespace is the combination of the database name and the collection name. Enclose the namespace in quotation marks.

sh._checkMongos()

sh._checkMongos ()

Returns nothing

Throws “not connected to a mongos”

The [sh._checkMongos\(\)](#) (page 172) method throws an error message if the [mongo](#) (page 527) shell is not connected to a [mongos](#) (page 518) instance. Otherwise it exits (no return document or return code).

sh._lastMigration()

Definition

sh._lastMigration (namespace)

Returns information on the last migration performed on the specified database or collection.

The [sh._lastMigration\(\)](#) (page 172) method has the following parameter:

param string namespace The *namespace* of a database or collection within the current database.

Output The [sh._lastMigration\(\)](#) (page 172) method returns a document with details about the last migration performed on the database or collection. The document contains the following output:

sh._lastMigration._id

The id of the migration task.

`sh.__lastMigration.server`

The name of the server.

`sh.__lastMigration.clientAddr`

The IP address and port number of the server.

`sh.__lastMigration.time`

The time of the last migration, formatted as *ISODate*.

`sh.__lastMigration.what`

The specific type of migration.

`sh.__lastMigration.ns`

The complete *namespace* of the collection affected by the migration.

`sh.__lastMigration.details`

A document containing details about the migrated chunk. The document includes `min` and `max` sub-documents with the bounds of the migrated chunk.

`sh.addShard()`

Definition

`sh.addShard(host)`

Adds a database instance or replica set to a *sharded cluster*. The optimal configuration is to deploy shards across *replica sets*. This method must be run on a `mongos` (page 518) instance.

The `sh.addShard()` (page 173) method has the following parameter:

param string host The hostname of either a standalone database instance or of a replica set. Include the port number if the instance is running on a non-standard port. Include the replica set name if the instance is a replica set, as explained below.

The `sh.addShard()` (page 173) method has the following prototype form:

```
sh.addShard("<host>")
```

The `host` parameter can be in any of the following forms:

```
[hostname]
[hostname]:[port]
[replica-set-name]/[hostname]
[replica-set-name]/[hostname]:port
```

Warning: Do not use `localhost` for the hostname unless your *configuration server* is also running on `localhost`.

New in version 2.6: `mongos` (page 518) installed from official `.deb` and `.rpm` packages have the `bind_ip` configuration set to `127.0.0.1` by default.

The `sh.addShard()` (page 173) method is a helper for the `addShard` (page 284) command. The `addShard` (page 284) command has additional options which are not available with this helper.

Important: You cannot include a `hidden` member in the seed list provided to `sh.addShard()` (page 173).

Example To add a shard on a replica set, specify the name of the replica set and the hostname of at least one member of the replica set, as a seed. If you specify additional hostnames, all must be members of the same replica set.

The following example adds a replica set named `rep10` and specifies one member of the replica set:

```
sh.addShard("repl0/mongodb3.example.net:27327")
```

sh.addShardTag()

Definition

`sh.addShardTag(shard, tag)`

New in version 2.2.

Associates a shard with a tag or identifier. MongoDB uses these identifiers to direct *chunks* that fall within a tagged range to specific shards. `sh.addTagRange()` (page 174) associates chunk ranges with tag ranges.

param string shard The name of the shard to which to give a specific tag.

param string tag The name of the tag to add to the shard.

Only issue `sh.addShardTag()` (page 174) when connected to a *mongos* (page 518) instance.

Example The following example adds three tags, NYC, LAX, and NRT, to three shards:

```
sh.addShardTag("shard0000", "NYC")
sh.addShardTag("shard0001", "LAX")
sh.addShardTag("shard0002", "NRT")
```

See also:

`sh.addTagRange()` (page 174) and `sh.removeShardTag()` (page 178).

sh.addTagRange()

Definition

`sh.addTagRange(namespace, minimum, maximum, tag)`

New in version 2.2.

Attaches a range of shard key values to a shard tag created using the `sh.addShardTag()` (page 174) method. `sh.addTagRange()` (page 174) takes the following arguments:

param string namespace The *namespace* of the sharded collection to tag.

param document minimum The minimum value of the *shard key* range to include in the tag. Specify the minimum value in the form of `<fieldname>:<value>`. This value must be of the same BSON type or types as the shard key.

param document maximum The maximum value of the shard key range to include in the tag. Specify the maximum value in the form of `<fieldname>:<value>`. This value must be of the same BSON type or types as the shard key.

param string tag The name of the tag to attach the range specified by the *minimum* and *maximum* arguments to.

Use `sh.addShardTag()` (page 174) to ensure that the balancer migrates documents that exist within the specified range to a specific shard or set of shards.

Only issue `sh.addTagRange()` (page 174) when connected to a *mongos* (page 518) instance.

Note: If you add a tag range to a collection using `sh.addTagRange()` (page 174) and then later drop the collection or its database, MongoDB does not remove the tag association. If you later create a new collection with the same name, the old tag association will apply to the new collection.

Example Given a shard key of `{state: 1, zip: 1}`, the following operation creates a tag range covering zip codes in New York State:

```
sh.addTagRange( "exampledb.collection",
               { state: "NY", zip: MinKey },
               { state: "NY", zip: MaxKey },
               "NY"
             )
```

sh.disableBalancing()

Description

sh.disableBalancing (*namespace*)

Disables the balancer for the specified sharded collection. This does not affect the balancing of *chunks* for other sharded collections in the same cluster.

The `sh.disableBalancing()` (page 175) method has the following parameter:

param string namespace The *namespace* of the collection.

For more information on the balancing process, see <http://docs.mongodb.org/manual/tutorial/manage-sharded-and-sharding-balancing>.

sh.enableBalancing()

Description

sh.enableBalancing (*collection*)

Enables the balancer for the specified sharded collection.

The `sh.enableBalancing()` (page 175) method has the following parameter:

param string collection The *namespace* of the collection.

Important: `sh.enableBalancing()` (page 175) does not *start* balancing. Rather, it allows balancing of this collection the next time the balancer runs.

For more information on the balancing process, see <http://docs.mongodb.org/manual/tutorial/manage-sharded-and-sharding-balancing>.

sh.enableSharding()

Definition

sh.enableSharding (*database*)

Enables sharding on the specified database. This does not automatically shard any collections but makes it possible to begin sharding collections using `sh.shardCollection()` (page 178).

The `sh.enableSharding()` (page 175) method has the following parameter:

param string database The name of the database shard. Enclose the name in quotation marks.

See also:

`sh.shardCollection()` (page 178)

sh.getBalancerHost()

sh.getBalancerHost()

Returns String in form *hostname:port*

sh.getBalancerHost() (page 176) returns the name of the server that is running the balancer.

See also:

- sh.enableBalancing() (page 175)
- sh.disableBalancing() (page 175)
- sh.getBalancerState() (page 176)
- sh.isBalancerRunning() (page 177)
- sh.setBalancerState() (page 178)
- sh.startBalancer() (page 180)
- sh.stopBalancer() (page 182)
- sh.waitForBalancer() (page 183)
- sh.waitForBalancerOff() (page 183)

sh.getBalancerState()

sh.getBalancerState()

Returns boolean

sh.getBalancerState() (page 176) returns `true` when the *balancer* is enabled and `false` if the balancer is disabled. This does not reflect the current state of balancing operations: use `sh.isBalancerRunning()` (page 177) to check the balancer's current state.

See also:

- sh.enableBalancing() (page 175)
- sh.disableBalancing() (page 175)
- sh.getBalancerHost() (page 176)
- sh.isBalancerRunning() (page 177)
- sh.setBalancerState() (page 178)
- sh.startBalancer() (page 180)
- sh.stopBalancer() (page 182)
- sh.waitForBalancer() (page 183)
- sh.waitForBalancerOff() (page 183)

sh.help()

sh.help()

Returns a basic help text for all sharding related shell functions.

sh.isBalancerRunning()**sh.isBalancerRunning()****Returns** boolean

Returns true if the *balancer* process is currently running and migrating chunks and false if the balancer process is not running. Use `sh.getBalancerState()` (page 176) to determine if the balancer is enabled or disabled.

See also:

- `sh.enableBalancing()` (page 175)
- `sh.disableBalancing()` (page 175)
- `sh.getBalancerHost()` (page 176)
- `sh.getBalancerState()` (page 176)
- `sh.setBalancerState()` (page 178)
- `sh.startBalancer()` (page 180)
- `sh.stopBalancer()` (page 182)
- `sh.waitForBalancer()` (page 183)
- `sh.waitForBalancerOff()` (page 183)

sh.moveChunk()**Definition****sh.moveChunk** (*namespace, query, destination*)

Moves the *chunk* that contains the document specified by the *query* to the destination shard. `sh.moveChunk()` (page 177) provides a wrapper around the `moveChunk` (page 296) database command and takes the following arguments:

- param string namespace** The *namespace* of the sharded collection that contains the chunk to migrate.
- param document query** An equality match on the shard key that selects the chunk to move.
- param string destination** The name of the shard to move.

Important: In most circumstances, allow the *balancer* to automatically migrate *chunks*, and avoid calling `sh.moveChunk()` (page 177) directly.

See also:

`moveChunk` (page 296), `sh.splitAt()` (page 179), `sh.splitFind()` (page 179), <http://docs.mongodb.org/manual/sharding>, and *chunk migration*.

Example Given the `people` collection in the `records` database, the following operation finds the chunk that contains the documents with the `zipcode` field set to 53187 and then moves that chunk to the shard named `shard0019`:

```
sh.moveChunk("records.people", { zipcode: "53187" }, "shard0019")
```

sh.removeShardTag()

Definition

sh.**removeShardTag** (*shard*, *tag*)

New in version 2.2.

Removes the association between a tag and a shard. Only issue `sh.removeShardTag()` (page 178) when connected to a `mongos` (page 518) instance.

param string shard The name of the shard from which to remove a tag.

param string tag The name of the tag to remove from the shard.

See also:

`sh.addShardTag()` (page 174), `sh.addTagRange()` (page 174)

sh.setBalancerState()

Description

sh.**setBalancerState** (*state*)

Enables or disables the *balancer*. Use `sh.getBalancerState()` (page 176) to determine if the balancer is currently enabled or disabled and `sh.isBalancerRunning()` (page 177) to check its current state.

The `sh.getBalancerState()` (page 176) method has the following parameter:

param Boolean state Set this to `true` to enable the balancer and `false` to disable it.

See also:

- `sh.enableBalancing()` (page 175)
- `sh.disableBalancing()` (page 175)
- `sh.getBalancerHost()` (page 176)
- `sh.getBalancerState()` (page 176)
- `sh.isBalancerRunning()` (page 177)
- `sh.startBalancer()` (page 180)
- `sh.stopBalancer()` (page 182)
- `sh.waitForBalancer()` (page 183)
- `sh.waitForBalancerOff()` (page 183)

sh.shardCollection()

Definition

sh.**shardCollection** (*namespace*, *key*, *unique*)

Shards a collection using the key as a the *shard key*. `sh.shardCollection()` (page 178) takes the following arguments:

param string namespace The *namespace* of the collection to shard.

param document key A *document* that specifies the *shard key* to use to *partition* and distribute objects among the shards. A shard key may be one field or multiple fields. A shard key with multiple fields is called a “compound shard key.”

param Boolean unique When true, ensures that the underlying index enforces a unique constraint. *Hashed shard keys* do not support unique constraints.

New in version 2.4: Use the form `{field: "hashed"}` to create a *hashed shard key*. Hashed shard keys may not be compound indexes.

Considerations MongoDB provides no method to deactivate sharding for a collection after calling `shardCollection` (page 291). Additionally, after `shardCollection` (page 291), you cannot change shard keys or modify the value of any field used in your shard key index.

Example Given the `people` collection in the `records` database, the following command shards the collection by the `zipcode` field:

```
sh.shardCollection("records.people", { zipcode: 1 } )
```

Additional Information `shardCollection` (page 291) for additional options, <http://docs.mongodb.org/manualsharding> and <http://docs.mongodb.org/manualcore/sharding-intro> for an overview of sharding, <http://docs.mongodb.org/manualtutorial/deploy-shard-cluster> for a tutorial, and *sharding-shard-key* for choosing a shard key.

sh.splitAt()

Definition

`sh.splitAt(namespace, query)`

Splits the chunk containing the document specified by the query as if that document were at the “middle” of the collection, even if the specified document is not the actual median of the collection.

param string namespace The namespace (i.e. `<database>.<collection>`) of the sharded collection that contains the chunk to split.

param document query A query to identify a document in a specific chunk. Typically specify the *shard key* for a document as the query.

Use this command to manually split chunks unevenly. Use the “`sh.splitFind()` (page 179)” function to split a chunk at the actual median.

In most circumstances, you should leave chunk splitting to the automated processes within MongoDB. However, when initially deploying a *sharded cluster* it is necessary to perform some measure of *pre-splitting* using manual methods including `sh.splitAt()` (page 179).

sh.splitFind()

Definition

`sh.splitFind(namespace, query)`

Splits the chunk containing the document specified by the query at its median point, creating two roughly equal chunks. Use `sh.splitAt()` (page 179) to split a collection in a specific point.

In most circumstances, you should leave chunk splitting to the automated processes. However, when initially deploying a *sharded cluster* it is necessary to perform some measure of *pre-splitting* using manual methods including `sh.splitFind()` (page 179).

param string namespace The namespace (i.e. `<database>.<collection>`) of the sharded collection that contains the chunk to split.

param document query A query to identify a document in a specific chunk. Typically specify the *shard key* for a document as the query.

sh.startBalancer()

Definition

sh.startBalancer (timeout, interval)

Enables the balancer in a sharded cluster and waits for balancing to initiate.

param integer timeout Milliseconds to wait.

param integer interval Milliseconds to sleep each cycle of waiting.

See also:

- sh.enableBalancing() (page 175)
- sh.disableBalancing() (page 175)
- sh.getBalancerHost() (page 176)
- sh.getBalancerState() (page 176)
- sh.isBalancerRunning() (page 177)
- sh.setBalancerState() (page 178)
- sh.stopBalancer() (page 182)
- sh.waitForBalancer() (page 183)
- sh.waitForBalancerOff() (page 183)

sh.status()

Definition

sh.status()

Prints a formatted report of the sharding configuration and the information regarding existing chunks in a *sharded cluster*. The default behavior suppresses the detailed chunk information if the total number of chunks is greater than or equal to 20.

The sh.status() (page 180) method has the following parameter:

param Boolean verbose If true, the method displays details of the document distribution across chunks when you have 20 or more chunks.

See also:

db.printShardingStatus() (page 116)

Output Examples The *Sharding Version* (page 181) section displays information on the *config database*:

```
--- Sharding Status ---
sharding version: {
  "_id" : <num>,
  "version" : <num>,
  "minCompatibleVersion" : <num>,
  "currentVersion" : <num>,
  "clusterId" : <ObjectId>
}
```


The *Shards* (page 182) section lists information on the shard(s). For each shard, the section displays the name, host, and the associated tags, if any.

```
shards:
  { "_id" : <shard name1>,
    "host" : <string>,
    "tags" : [ <string> ... ]
  }
  { "_id" : <shard name2>,
    "host" : <string>,
    "tags" : [ <string> ... ]
  }
  ...
```

The *Databases* (page 182) section lists information on the database(s). For each database, the section displays the name, whether the database has sharding enabled, and the *primary shard* for the database.

```
databases:
  { "_id" : <dbname1>,
    "partitioned" : <boolean>,
    "primary" : <string>
  }
  { "_id" : <dbname2>,
    "partitioned" : <boolean>,
    "primary" : <string>
  }
  ...
```

The *Sharded Collection* (page 182) section provides information on the sharding details for sharded collection(s). For each sharded collection, the section displays the shard key, the number of chunks per shard(s), the distribution of documents across chunks ¹⁴, and the tag information, if any, for shard key range(s).

```
<dbname>.<collection>
  shard key: { <shard key> : <1 or hashed> }
  chunks:
    <shard name1> <number of chunks>
    <shard name2> <number of chunks>
    ...
  { <shard key>: <min range1> } --> { <shard key> : <max range1> } on : <shard name> <last modified>
  { <shard key>: <min range2> } --> { <shard key> : <max range2> } on : <shard name> <last modified>
  ...
  tag: <tag1> { <shard key> : <min range1> } --> { <shard key> : <max range1> }
  ...
```

Output Fields

Sharding Version

`sh.status.sharding-version._id`

The `_id` (page 181) is an identifier for the version details.

`sh.status.sharding-version.version`

The `version` (page 181) is the version of the *config server* for the sharded cluster.

`sh.status.sharding-version.minCompatibleVersion`

The `minCompatibleVersion` (page 181) is the minimum compatible version of the config server.

¹⁴ The sharded collection section, by default, displays the chunk information if the total number of chunks is less than 20. To display the information when you have 20 or more chunks, call the `sh.status()` (page 180) methods with the `verbose` parameter set to `true`, i.e. `sh.status(true)`.

`sh.status.sharding-version.currentVersion`

The `currentVersion` (page 181) is the current version of the config server.

`sh.status.sharding-version.clusterId`

The `clusterId` (page 182) is the identification for the sharded cluster.

Shards

`sh.status.shards._id`

The `_id` (page 182) displays the name of the shard.

`sh.status.shards.host`

The `host` (page 182) displays the host location of the shard.

`sh.status.shards.tags`

The `tags` (page 182) displays all the tags for the shard. The field only displays if the shard has tags.

Databases

`sh.status.databases._id`

The `_id` (page 182) displays the name of the database.

`sh.status.databases.partitioned`

The `partitioned` (page 182) displays whether the database has sharding enabled. If `true`, the database has sharding enabled.

`sh.status.databases.primary`

The `primary` (page 182) displays the *primary shard* for the database.

Sharded Collection

`sh.status.databases.shard-key`

The `shard-key` (page 182) displays the shard key specification document.

`sh.status.databases.chunks`

The `chunks` (page 182) lists all the shards and the number of chunks that reside on each shard.

`sh.status.databases.chunk-details`

The `chunk-details` (page 182) lists the details of the chunks ¹:

- The range of shard key values that define the chunk,
- The shard where the chunk resides, and
- The last modified timestamp for the chunk.

`sh.status.databases.tag`

The `tag` (page 182) lists the details of the tags associated with a range of shard key values.

`sh.stopBalancer()`

Definition

`sh.stopBalancer` (*timeout, interval*)

Disables the balancer in a sharded cluster and waits for balancing to complete.

param integer timeout Milliseconds to wait.

param integer interval Milliseconds to sleep each cycle of waiting.

See also:

- `sh.enableBalancing()` (page 175)
- `sh.disableBalancing()` (page 175)

- `sh.getBalancerHost()` (page 176)
- `sh.getBalancerState()` (page 176)
- `sh.isBalancerRunning()` (page 177)
- `sh.setBalancerState()` (page 178)
- `sh.startBalancer()` (page 180)
- `sh.waitForBalancer()` (page 183)
- `sh.waitForBalancerOff()` (page 183)

sh.waitForBalancer()**Definition**

`sh.waitForBalancer` (*wait, timeout, interval*)

Waits for a change in the state of the balancer. `sh.waitForBalancer()` (page 183) is an internal method, which takes the following arguments:

- param Boolean wait** Set to `true` to ensure the balancer is now active. The default is `false`, which waits until balancing stops and becomes inactive.
- param integer timeout** Milliseconds to wait.
- param integer interval** Milliseconds to sleep.

sh.waitForBalancerOff()**Definition**

`sh.waitForBalancerOff` (*timeout, interval*)

Internal method that waits until the balancer is not running.

- param integer timeout** Milliseconds to wait.
- param integer interval** Milliseconds to sleep.

See also:

- `sh.enableBalancing()` (page 175)
- `sh.disableBalancing()` (page 175)
- `sh.getBalancerHost()` (page 176)
- `sh.getBalancerState()` (page 176)
- `sh.isBalancerRunning()` (page 177)
- `sh.setBalancerState()` (page 178)
- `sh.startBalancer()` (page 180)
- `sh.stopBalancer()` (page 182)
- `sh.waitForBalancer()` (page 183)

sh.waitForDLock()**Definition**

sh.waitForDLock (*lockname, wait, timeout, interval*)

Waits until the specified distributed lock changes state. `sh.waitForDLock()` (page 184) is an internal method that takes the following arguments:

param string lockname The name of the distributed lock.

param Boolean wait Set to `true` to ensure the balancer is now active. Set to `false` to wait until balancing stops and becomes inactive.

param integer timeout Milliseconds to wait.

param integer interval Milliseconds to sleep in each waiting cycle.

sh.waitForPingChange()**Definition**

sh.waitForPingChange (*activePings, timeout, interval*)

`sh.waitForPingChange()` (page 184) waits for a change in ping state of one of the `activePings`, and only returns when the specified ping changes state.

param array activePings An array of active pings from the `mongos` (page 597) collection.

param integer timeout Number of milliseconds to wait for a change in ping state.

param integer interval Number of milliseconds to sleep in each waiting cycle.

2.1.10 Subprocess

Subprocess Methods

Name	Description
<code>clearRawMongoProgramOutput()</code> (page 184)	For internal use.
<code>rawMongoProgramOutput()</code> (page 185)	For internal use.
<code>run()</code>	For internal use.
<code>runMongoProgram()</code> (page 185)	For internal use.
<code>runProgram()</code> (page 185)	For internal use.
<code>startMongoProgram()</code>	For internal use.
<code>stopMongoProgram()</code> (page 185)	For internal use.
<code>stopMongoProgramByPid()</code> (page 185)	For internal use.
<code>stopMongod()</code> (page 185)	For internal use.
<code>waitMongoProgramOnPort()</code> (page 185)	For internal use.
<code>waitProgram()</code> (page 186)	For internal use.

clearRawMongoProgramOutput()

clearRawMongoProgramOutput ()

For internal use.

rawMongoProgramOutput()**rawMongoProgramOutput ()**

For internal use.

run()**run ()**

For internal use.

runMongoProgram()**runMongoProgram ()**

For internal use.

runProgram()**runProgram ()**

For internal use.

startMongoProgram()**_startMongoProgram ()**

For internal use.

stopMongoProgram()**stopMongoProgram ()**

For internal use.

stopMongoProgramByPid()**stopMongoProgramByPid ()**

For internal use.

stopMongod()**stopMongod ()**

For internal use.

waitMongoProgramOnPort()**waitMongoProgramOnPort ()**

For internal use.

waitProgram()

waitProgram()
For internal use.

2.1.11 Constructors**Object Constructors and Methods**

Name	Description
<code>Date()</code> (page 186)	Creates a date object. By default creates a date object including the current date.
<code>UUID()</code> (page 186)	Converts a 32-byte hexadecimal string to the UUID BSON subtype.
<code>ObjectId.getTimestamp()</code> (page 187)	Returns the timestamp portion of an <i>ObjectId</i> .
<code>ObjectId.toString()</code> (page 187)	Displays the string representation of an <i>ObjectId</i> .
<code>ObjectId.valueOf()</code> (page 187)	Displays the <code>str</code> attribute of an <i>ObjectId</i> as a hexadecimal string.
<code>WriteResult()</code> (page 188)	Wrapper around the result set from write methods.
<code>WriteResult.hasWriteError()</code> (page 189)	Returns a boolean specifying whether the results include <code>WriteResult.writeError</code> (page 188).
<code>WriteResult.hasWriteConcernError()</code> (page 189)	Returns a boolean specifying whether whether the results include <code>WriteResult.writeConcernError</code> (page 188).
<code>BulkWriteResult()</code> (page 189)	Wrapper around the result set from <code>Bulk.execute()</code> (page 137).

Date()**Date()**

Returns Current date, as a string.

UUID()**Definition**

UUID (<string>)

Generates a BSON UUID object.

param string hex Specify a 32-byte hexadecimal string to convert to the UUID BSON subtype.

Returns A BSON UUID object.

Example Create a 32 byte hexadecimal string:

```
var myuuid = '0123456789abcdeffedcba9876543210'
```

Convert it to the BSON UUID subtype:

```
UUID(myuuid)
BinData(3, "ASNfZ4mrze/+3LqYdlQyEA==")
```

ObjectId.getTimestamp()`ObjectId.getTimestamp()`

Returns The timestamp portion of the *ObjectId()* object as a Date.

In the following example, call the `getTimestamp()` (page 187) method on an *ObjectId* (e.g. `ObjectId("507c7f79bcf86cd7994f6c0e")`):

```
ObjectId("507c7f79bcf86cd7994f6c0e").getTimestamp()
```

This will return the following output:

```
ISODate("2012-10-15T21:26:17Z")
```

ObjectId.toString()`ObjectId.toString()`

Returns The string representation of the *ObjectId()* object. This value has the format of `ObjectId(...)`.

Changed in version 2.2: In previous versions `ObjectId.toString()` (page 187) returns the value of the *ObjectId* as a hexadecimal string.

In the following example, call the `toString()` (page 187) method on an *ObjectId* (e.g. `ObjectId("507c7f79bcf86cd7994f6c0e")`):

```
ObjectId("507c7f79bcf86cd7994f6c0e").toString()
```

This will return the following string:

```
ObjectId("507c7f79bcf86cd7994f6c0e")
```

You can confirm the type of this object using the following operation:

```
typeof ObjectId("507c7f79bcf86cd7994f6c0e").toString()
```

ObjectId.valueOf()`ObjectId.valueOf()`

Returns The value of the *ObjectId()* object as a lowercase hexadecimal string. This value is the `str` attribute of the *ObjectId()* object.

Changed in version 2.2: In previous versions `ObjectId.valueOf()` (page 187) returns the *ObjectId()* object.

In the following example, call the `valueOf()` (page 187) method on an *ObjectId* (e.g. `ObjectId("507c7f79bcf86cd7994f6c0e")`):

```
ObjectId("507c7f79bcf86cd7994f6c0e").valueOf()
```

This will return the following string:

```
507c7f79bcf86cd7994f6c0e
```

You can confirm the type of this object using the following operation:

```
typeof ObjectId("507c7f79bcf86cd7994f6c0e").valueOf()
```

WriteResult()

Definition

WriteResult()

A wrapper that contains the result status of the `mongo` (page 527) shell write methods.

See

`db.collection.insert()` (page 52), `db.collection.update()` (page 69), `db.collection.remove()` (page 62), and `db.collection.save()` (page 66).

Properties The `WriteResult` (page 188) has the following properties:

`WriteResult.nInserted`

The number of documents inserted, excluding upserted documents. See `WriteResult.nUpserted` (page 188) for the number of documents inserted through an upsert.

`WriteResult.nMatched`

The number of documents selected for update. If the update operation results in no change to the document, e.g. `$set` (page 416) expression updates the value to the current value, `nMatched` (page 188) can be greater than `nModified` (page 188).

`WriteResult.nModified`

The number of existing documents updated. If the update/replacement operation results in no change to the document, such as setting the value of the field to its current value, `nModified` (page 188) can be less than `nMatched` (page 188).

`WriteResult.nUpserted`

The number of documents inserted by an *upsert* (page 70).

`WriteResult._id`

The `_id` of the document inserted by an upsert. Returned only if an upsert results in an insert.

`WriteResult.nRemoved`

The number of documents removed.

`WriteResult.writeError`

A document that contains information regarding any error, excluding write concern errors, encountered during the write operation.

`WriteResult.writeError.code`

An integer value identifying the error.

`WriteResult.writeError.errmsg`

A description of the error.

`WriteResult.writeConcernError`

A document that contains information regarding any write concern errors encountered during the write operation.

`WriteResult.writeConcernError.code`

An integer value identifying the write concern error.

`WriteResult.writeConcernError.errInfo`

A document identifying the write concern setting related to the error.

`WriteResult.writeError.errmsg`

A description of the error.

See also:

`WriteResult.hasWriteError()` (page 189), `WriteResult.hasWriteConcernError()` (page 189)

WriteResult.hasWriteError()

Definition

`WriteResult.hasWriteError()`

Returns true if the result of a `mongo` (page 527) shell write method has `WriteResult.writeError` (page 188). Otherwise, the method returns false.

See also:

`WriteResult()` (page 188)

WriteResult.hasWriteConcernError()

Definition

`WriteResult.hasWriteConcernError()`

Returns true if the result of a `mongo` (page 527) shell write method has `WriteResult.writeConcernError` (page 188). Otherwise, the method returns false.

See also:

`WriteResult()` (page 188)

BulkWriteResult()

BulkWriteResult()

New in version 2.6.

A wrapper that contains the results of the `Bulk.execute()` (page 137) method.

Properties The `BulkWriteResult` (page 189) has the following properties:

`BulkWriteResult.nInserted`

The number of documents inserted using the `Bulk.insert()` (page 129) method. For documents inserted through an upsert, see the `nUpserted` (page 189) field instead.

`BulkWriteResult.nMatched`

The number of existing documents selected for update or replacement. If the update/replacement operation results in no change to an existing document, e.g. `$set` (page 416) expression updates the value to the current value, `nMatched` (page 189) can be greater than `nModified` (page 189).

`BulkWriteResult.nModified`

The number of existing documents updated or replaced. If the update/replacement operation results in no change to an existing document, such as setting the value of the field to its current value, `nModified` (page 189) can be less than `nMatched` (page 189). Inserted documents do not affect the number of `nModified` (page 189); refer to the `nInserted` (page 189) and `nUpserted` (page 189) fields instead.

`BulkWriteResult.nRemoved`

The number of documents removed.

`BulkWriteResult.nUpserted`

The number of documents inserted through operations with the `Bulk.find.upsert()` (page 134) option.

`BulkWriteResult.upserted`

An array of documents that contains information for each upserted documents.

Each document contains the following information:

`BulkWriteResult.upserted.index`

An integer that identifies the operation in the bulk operations list, which uses a zero-based index.

`BulkWriteResult.upserted._id`

The `_id` value of the upserted document.

`BulkWriteResult.writeErrors`

An array of documents that contains information regarding any error, unrelated to write concerns, encountered during the update operation. The `writeErrors` (page 190) array contains an error document for each write operation that errors.

Each error document contains the following fields:

`BulkWriteResult.writeErrors.index`

An integer that identifies the write operation in the bulk operations list, which uses a zero-based index. See also `Bulk.getOperations()` (page 138).

`BulkWriteResult.writeErrors.code`

An integer value identifying the error.

`BulkWriteResult.writeErrors.errmsg`

A description of the error.

`BulkWriteResult.writeErrors.op`

A document identifying the operation that failed. For instance, an update/replace operation error will return a document specifying the query, the update, the `multi` and the `upsert` options; an insert operation will return the document the operation tried to insert.

`BulkWriteResult.writeConcernError`

Document that describe error related to write concern and contains the field:

`BulkWriteResult.writeConcernError.code`

An integer value identifying the cause of the write concern error.

`BulkWriteResult.writeConcernError.errInfo`

A document identifying the write concern setting related to the error.

`BulkWriteResult.writeConcernError.errmsg`

A description of the cause of the write concern error.

2.1.12 Connection

Connection Methods

Name	Description
<code>Mongo.getDB()</code> (page 191)	Returns a database object.
<code>Mongo.getReadPrefMode()</code> (page 191)	Returns the current read preference mode for the MongoDB connection.
<code>Mongo.getReadPrefTagSet()</code> (page 192)	Returns the read preference tag set for the MongoDB connection.
<code>Mongo.setReadPref()</code> (page 192)	Sets the <i>read preference</i> for the MongoDB connection.
<code>Mongo.setSlaveOk()</code> (page 193)	Allows operations on the current connection to read from <i>secondary</i> members.
<code>Mongo()</code> (page 193)	Creates a new connection object.
<code>connect()</code>	Connects to a MongoDB instance and to a specified database on that instance.

Mongo.getDB()

Description

`Mongo.getDB(<database>)`

Provides access to database objects from the `mongo` (page 527) shell or from a JavaScript file.

The `Mongo.getDB()` (page 191) method has the following parameter:

param string database The name of the database to access.

Example The following example instantiates a new connection to the MongoDB instance running on the localhost interface and returns a reference to "myDatabase":

```
db = new Mongo().getDB("myDatabase");
```

See also:

`Mongo()` (page 193) and `connect()` (page 193)

Mongo.getReadPrefMode()

`Mongo.getReadPrefMode()`

Returns The current *read preference* mode for the `Mongo()` (page 111) connection object.

See <http://docs.mongodb.org/manualcore/read-preference> for an introduction to read preferences in MongoDB. Use `getReadPrefMode()` (page 191) to return the current read preference mode, as in the following example:

```
db.getMongo().getReadPrefMode()
```

Use the following operation to return and print the current read preference mode:

```
print(db.getMongo().getReadPrefMode());
```

This operation will return one of the following read preference modes:

- primary

- primaryPreferred
- secondary
- secondaryPreferred
- nearest

See also:

<http://docs.mongodb.org/manualcore/read-preference>, [readPref\(\)](#) (page 91), [setReadPref\(\)](#) (page 192), and [getReadPrefTagSet\(\)](#) (page 192).

Mongo.getReadPrefTagSet()

`Mongo.getReadPrefTagSet()`

Returns The current *read preference* tag set for the `Mongo()` (page 111) connection object.

See <http://docs.mongodb.org/manualcore/read-preference> for an introduction to read preferences and tag sets in MongoDB. Use [getReadPrefTagSet\(\)](#) (page 192) to return the current read preference tag set, as in the following example:

```
db.getMongo().getReadPrefTagSet()
```

Use the following operation to return and print the current read preference tag set:

```
printjson(db.getMongo().getReadPrefTagSet());
```

See also:

<http://docs.mongodb.org/manualcore/read-preference>, [readPref\(\)](#) (page 91), [setReadPref\(\)](#) (page 192), and [getReadPrefTagSet\(\)](#) (page 192).

Mongo.setReadPref()**Definition**

`Mongo.setReadPref(mode, tagSet)`

Call the [setReadPref\(\)](#) (page 192) method on a `Mongo` (page 111) connection object to control how the client will route all queries to members of the replica set.

param string mode One of the following *read preference* modes: `primary`, `primaryPreferred`, `secondary`, `secondaryPreferred`, or `nearest`.

param array tagSet A tag set used to specify custom read preference modes. For details, see *replica-set-read-preference-tag-sets*.

Examples To set a read preference mode in the `mongo` (page 527) shell, use the following operation:

```
db.getMongo().setReadPref('primaryPreferred')
```

To set a read preference that uses a tag set, specify an array of tag sets as the second argument to `Mongo.setReadPref()` (page 192), as in the following:

```
db.getMongo().setReadPref('primaryPreferred', [ { "dc": "east" } ] )
```

You can specify multiple tag sets, in order of preference, as in the following:

```
db.getMongo().setReadPref('secondaryPreferred',
    [ { "dc": "east", "use": "production" },
      { "dc": "east", "use": "reporting" },
      { "dc": "east" },
      {}
    ] )
```

If the replica set cannot satisfy the first tag set, the client will attempt to use the second read preference. Each tag set can contain zero or more field/value tag pairs, with an “empty” document acting as a wildcard which matches a replica set member with any tag set or no tag set.

Note: You must call `Mongo.setReadPref()` (page 192) on the connection object before retrieving documents using that connection to use that read preference.

`mongo.setSlaveOk()`

`Mongo.setSlaveOk()`

For the current session, this command permits read operations from non-master (i.e. *slave* or *secondary*) instances. Practically, use this method in the following form:

```
db.getMongo().setSlaveOk()
```

Indicates that “*eventually consistent*” read operations are acceptable for the current application. This function provides the same functionality as `rs.slaveOk()` (page 168).

See the `readPref()` (page 91) method for more fine-grained control over read preference in the `mongo` (page 527) shell.

`Mongo()`

Description

Mongo (*host*)

JavaScript constructor to instantiate a database connection from the `mongo` (page 527) shell or from a JavaScript file.

The `Mongo()` (page 193) method has the following parameter:

param string host The host, either in the form of `<host>` or `<host>:<port>`.

Instantiation Options Use the constructor without a parameter to instantiate a connection to the localhost interface on the default port.

Pass the `<host>` parameter to the constructor to instantiate a connection to the `<host>` and the default port.

Pass the `<host>:<port>` parameter to the constructor to instantiate a connection to the `<host>` and the `<port>`.

See also:

`Mongo.getDB()` (page 191) and `db.getMongo()` (page 111).

`connect()`

connect (`<hostname>:<port>/<database>`)

The `connect()` method creates a connection to a MongoDB instance. However, use the `Mongo()` (page 193) object and its `getDB()` (page 191) method in most cases.

`connect()` accepts a string `<hostname><:port>/<database>` parameter to connect to the MongoDB instance on the `<hostname><:port>` and return the reference to the database `<database>`.

The following example instantiates a new connection to the MongoDB instance running on the localhost interface and returns a reference to `myDatabase`:

```
db = connect("localhost:27017/myDatabase")
```

See also:

`Mongo.getDB()` (page 191)

2.1.13 Native

Native Methods

Name	Description
<code>cat()</code>	Returns the contents of the specified file.
<code>version()</code>	Returns the current version of the <code>mongo</code> (page 527) shell instance.
<code>cd()</code>	Changes the current working directory to the specified path.
<code>copyDbpath()</code> (page 195)	Copies a local <code>dbPath</code> . For internal use.
<code>resetDbpath()</code> (page 195)	Removes a local <code>dbPath</code> . For internal use.
<code>fuzzFile()</code> (page 195)	For internal use to support testing.
<code>getHostName()</code> (page 195)	Returns the hostname of the system running the <code>mongo</code> (page 527) shell.
<code>getMemInfo()</code> (page 195)	Returns a document that reports the amount of memory used by the shell.
<code>hostname()</code>	Returns the hostname of the system running the shell.
<code>_isWindows()</code> (page 196)	Returns <code>true</code> if the shell runs on a Windows system; <code>false</code> if a Unix or Linux system.
<code>listFiles()</code> (page 196)	Returns an array of documents that give the name and size of each object in the directory.
<code>load()</code>	Loads and runs a JavaScript file in the shell.
<code>ls()</code>	Returns a list of the files in the current directory.
<code>md5sumFile()</code> (page 197)	The <code>md5</code> hash of the specified file.
<code>mkdir()</code>	Creates a directory at the specified path.
<code>pwd()</code>	Returns the current directory.
<code>quit()</code>	Exits the current shell session.
<code>_rand()</code> (page 197)	Returns a random number between 0 and 1.
<code>removeFile()</code> (page 198)	Removes the specified file from the local file system.
<code>_srand()</code> (page 198)	For internal use.

`cat()`

Definition

cat (*filename*)

Returns the contents of the specified file. The method returns with output relative to the current shell session and does not impact the server.

param string filename Specify a path and file name on the local file system.

version()**version()**

Returns The version of the `mongo` (page 527) shell as a string.

Changed in version 2.4: In previous versions of the shell, `version()` would print the version instead of returning a string.

cd()**Definition****cd(*path*)**

param string path A path on the file system local to the `mongo` (page 527) shell context.

`cd()` changes the directory context of the `mongo` (page 527) shell and has no effect on the MongoDB server.

copyDbpath()**copyDbpath()**

For internal use.

resetDbpath()**resetDbpath()**

For internal use.

fuzzFile()**Description****fuzzFile(*filename*)**

For internal use.

param string filename A filename or path to a local file.

getHostName()**getHostName()**

Returns The hostname of the system running the `mongo` (page 527) shell process.

getMemInfo()**getMemInfo()**

Returns a document with two fields that report the amount of memory used by the JavaScript shell process. The fields returned are *resident* and *virtual*.

hostname()

hostname ()

Returns The hostname of the system running the `mongo` (page 527) shell process.

_isWindows()

_isWindows ()

Returns boolean.

Returns “true” if the `mongo` (page 527) shell is running on a system that is Windows, or “false” if the shell is running on a Unix or Linux systems.

listFiles()

listFiles ()

Returns an array, containing one document per object in the directory. This function operates in the context of the `mongo` (page 527) process. The included fields are:

name

Returns a string which contains the name of the object.

isDirectory

Returns true or false if the object is a directory.

size

Returns the size of the object in bytes. This field is only present for files.

load()

Definition

load (*file*)

Loads and runs a JavaScript file into the current shell environment.

The `load ()` method has the following parameter:

param string filename Specifies the path of a JavaScript file to execute.

Specify filenames with relative or absolute paths. When using relative path names, confirm the current directory using the `pwd ()` method.

After executing a file with `load ()`, you may reference any functions or variables defined the file from the `mongo` (page 527) shell environment.

Example Consider the following examples of the `load ()` method:

```
load("scripts/myjstest.js")
load("/data/db/scripts/myjstest.js")
```


ls()**ls()**

Returns a list of the files in the current directory.

This function returns with output relative to the current shell session, and does not impact the server.

md5sumFile()**Description****md5sumFile** (*filename*)

Returns a *md5* hash of the specified file.

The `md5sumFile()` (page 197) method has the following parameter:

param string filename A file name.

Note: The specified filename must refer to a file located on the system running the `mongo` (page 527) shell.

mkdir()**Description****mkdir** (*path*)

Creates a directory at the specified path. This method creates the entire path specified if the enclosing directory or directories do not already exist.

This method is equivalent to **mkdir -p** with BSD or GNU utilities.

The `mkdir()` method has the following parameter:

param string path A path on the local filesystem.

pwd()**pwd()**

Returns the current directory.

This function returns with output relative to the current shell session, and does not impact the server.

quit()**quit()**

Exits the current shell session.

rand()**__rand()**

Returns A random number between 0 and 1.

This function provides functionality similar to the `Math.rand()` function from the standard library.

`removeFile()`

Description

removeFile (*filename*)

Removes the specified file from the local file system.

The `removeFile()` (page 198) method has the following parameter:

param string filename A filename or path to a local file.

`_srand()`

`_srand()`

For internal use.

2.2 Database Commands

All command documentation outlined below describes a command and its available parameters and provides a document template or prototype for each command. Some command documentation also includes the relevant `mongo` (page 527) shell helpers.

2.2.1 User Commands

Aggregation Commands

Aggregation Commands

Name	Description
<code>aggregate</code> (page 198)	Performs aggregation tasks such as group using the aggregation framework.
<code>count</code> (page 201)	Counts the number of documents in a collection.
<code>distinct</code> (page 203)	Displays the distinct values found for a specified key in a collection.
<code>group</code> (page 204)	Groups documents in a collection by the specified key and performs simple aggregation.
<code>mapReduce</code> (page 208)	Performs map-reduce aggregation for large data sets.

aggregate

aggregate

New in version 2.2.

Performs aggregation operation using the *aggregation pipeline* (page 438). The pipeline allows users to process data from a collection with a sequence of stage-based manipulations.

Changed in version 2.6.

- The `aggregate` (page 198) command adds support for returning a cursor, supports the `explain` option, and enhances its sort operations with an external sort facility.
- aggregation pipeline* (page 438) introduces the `$out` (page 453) operator to allow `aggregate` (page 198) command to store results to a collection.

The command has following syntax:

Changed in version 2.6.

```
{
  aggregate: "<collection>",
  pipeline: [ <stage>, <...> ],
  explain: <boolean>,
  allowDiskUse: <boolean>,
  cursor: <document>
}
```

The `aggregate` (page 198) command takes the following fields as arguments:

field string aggregate The name of the collection to use as the input for the aggregation pipeline.

field array pipeline An array of *aggregation pipeline stages* (page 438) that process and transform the document stream as part of the aggregation pipeline.

field boolean explain Specifies to return the information on the processing of the pipeline.

New in version 2.6.

field boolean allowDiskUse Enables writing to temporary files. When set to `true`, aggregation stages can write data to the `_tmp` subdirectory in the `dbPath` directory.

New in version 2.6.

field document cursor Specify a document that contains options that control the creation of the cursor object.

New in version 2.6.

For more information about the aggregation pipeline <http://docs.mongodb.org/manualcore/aggregation-pipeline>, *Aggregation Reference* (page 484), and <http://docs.mongodb.org/manualcore/aggregation-pipeline-limits>.

Example

Aggregate Data with Multi-Stage Pipeline A collection `articles` contains documents such as the following:

```
{
  _id: ObjectId("52769ea0f3dc6ead47c9a1b2"),
  author: "abc123",
  title: "zzz",
  tags: [ "programming", "database", "mongodb" ]
}
```

The following example performs an `aggregate` (page 198) operation on the `articles` collection to calculate the count of each distinct element in the `tags` array that appears in the collection.

```
db.runCommand(
  { aggregate: "articles",
    pipeline: [
      { $project: { tags: 1 } },
      { $unwind: "$tags" },
      { $group: {
          _id: "$tags",
          count: { $sum : 1 }
        }
      }
    ]
  })
```

```
    }  
  )
```

In the `mongo` (page 527) shell, this operation can use the `aggregate()` (page 22) helper as in the following:

```
db.articles.aggregate(  
  [  
    { $project: { tags: 1 } },  
    { $unwind: "$tags" },  
    { $group: {  
      _id: "$tags",  
      count: { $sum : 1 }  
    }  
  ]  
)
```

Note: In 2.6 and later, the `aggregate()` (page 22) helper always returns a cursor.

Changed in version 2.4: If an error occurs, the `aggregate()` (page 22) helper throws an exception. In previous versions, the helper returned a document with the error message and code, and `ok` status field not equal to 1, same as the `aggregate` (page 198) command.

Return Information on the Aggregation Operation The following aggregation operation sets the optional field `explain` to `true` to return information about the aggregation operation.

```
db.runCommand( { aggregate: "orders",  
  pipeline: [  
    { $match: { status: "A" } },  
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },  
    { $sort: { total: -1 } }  
  ],  
  explain: true  
} )
```

Note: The intended readers of the `explain` output document are humans, and not machines, and the output format is subject to change between releases.

See also:

`db.collection.aggregate()` (page 22) method

Aggregate Data using External Sort Aggregation pipeline stages have *maximum memory use limit*. To handle large datasets, set `allowDiskUse` option to `true` to enable writing data to temporary files, as in the following example:

```
db.runCommand(  
  { aggregate: "stocks",  
    pipeline: [  
      { $project : { cusip: 1, date: 1, price: 1, _id: 0 } },  
      { $sort : { cusip : 1, date: 1 } }  
    ],  
    allowDiskUse: true  
  }  
)
```

See also:

`db.collection.aggregate()` (page 22)

Aggregate Command Returns a Cursor

Note: Using the `aggregate` (page 198) command to return a cursor is a low-level operation, intended for authors of drivers. Most users should use the `db.collection.aggregate()` (page 22) helper provided in the `mongo` (page 527) shell or in their driver. In 2.6 and later, the `aggregate()` (page 22) helper always returns a cursor.

The following command returns a document that contains results with which to instantiate a cursor object.

```
db.runCommand(
  { aggregate: "records",
    pipeline: [
      { $project: { name: 1, email: 1, _id: 0 } },
      { $sort: { name: 1 } }
    ],
    cursor: { }
  }
)
```

To specify an *initial* batch size, specify the `batchSize` in the `cursor` field, as in the following example:

```
db.runCommand(
  { aggregate: "records",
    pipeline: [
      { $project: { name: 1, email: 1, _id: 0 } },
      { $sort: { name: 1 } }
    ],
    cursor: { batchSize: 0 }
  }
)
```

The `{batchSize: 0}` document specifies the size of the *initial* batch size only. Specify subsequent batch sizes to *OP_GET_MORE*¹⁵ operations as with other MongoDB cursors. A `batchSize` of 0 means an empty first batch and is useful if you want to quickly get back a cursor or failure message, without doing significant server-side work.

See also:

`db.collection.aggregate()` (page 22)

count

Definition

count

Counts the number of documents in a collection. Returns a document that contains this count and as well as the command status. `count` (page 201) has the following form:

Changed in version 2.6: `count` (page 201) now accepts the `hint` option to specify an index.

```
{ count: <collection>, query: <query>, limit: <limit>, skip: <skip>, hint: <hint> }
```

`count` (page 201) has the following fields:

field string count The name of the collection to count.

field document query A query that selects which documents to count in a collection.

field integer limit The maximum number of matching documents to return.

¹⁵<http://docs.mongodb.org/meta-driver/latest/legacy/mongodb-wire-protocol/#wire-op-get-more>

field integer skip The number of matching documents to skip before returning results.

field String,document hint The index to use. Specify either the index name as a string or the index specification document.

New in version 2.6.

MongoDB also provides the `count()` (page 79) and `db.collection.count()` (page 25) wrapper methods in the `mongo` (page 527) shell.

Behavior On a sharded cluster, `count` (page 201) can result in an *inaccurate* count if *orphaned documents* exist or if a chunk migration is in progress.

To avoid these situations, on a sharded cluster, use the `$group` (page 447) stage of the `db.collection.aggregate()` (page 22) method to `$sum` (page 460) the documents. For example, the following operation counts the documents in a collection:

```
db.collection.aggregate([
  { $group: { _id: null, count: { $sum: 1 } } }
])
```

To get a count of documents that match a query condition, include the `$match` (page 440) stage as well:

```
db.collection.aggregate([
  { $match: <query condition> },
  { $group: { _id: null, count: { $sum: 1 } } }
])
```

See [Perform a Count](#) (page 440) for an example.

Examples The following sections provide examples of the `count` (page 201) command.

Count All Documents The following operation counts the number of all documents in the `orders` collection:

```
db.runCommand( { count: 'orders' } )
```

In the result, the `n`, which represents the count, is 26, and the command status `ok` is 1:

```
{ "n" : 26, "ok" : 1 }
```

Count Documents That Match a Query The following operation returns a count of the documents in the `orders` collection where the value of the `ord_dt` field is greater than `Date('01/01/2012')`:

```
db.runCommand( { count: 'orders',
  query: { ord_dt: { $gt: new Date('01/01/2012') } }
} )
```

In the result, the `n`, which represents the count, is 13 and the command status `ok` is 1:

```
{ "n" : 13, "ok" : 1 }
```

Skip Documents in Count The following operation returns a count of the documents in the `orders` collection where the value of the `ord_dt` field is greater than `Date('01/01/2012')` and skip the first 10 matching documents:

```
db.runCommand( { count:'orders',
                  query: { ord_dt: { $gt: new Date('01/01/2012') } },
                  skip: 10 } )
```

In the result, the `n`, which represents the count, is 3 and the command status `ok` is 1:

```
{ "n" : 3, "ok" : 1 }
```

Specify the Index to Use The following operation uses the index `{ status: 1 }` to return a count of the documents in the `orders` collection where the value of the `ord_dt` field is greater than `Date('01/01/2012')` and the `status` field is equal to `"D"`:

```
db.runCommand(
  {
    count:'orders',
    query: {
      ord_dt: { $gt: new Date('01/01/2012') },
      status: "D"
    },
    hint: { status: 1 }
  }
)
```

In the result, the `n`, which represents the count, is 1 and the command status `ok` is 1:

```
{ "n" : 1, "ok" : 1 }
```

distinct

Definition

distinct

Finds the distinct values for a specified field across a single collection. `distinct` (page 203) returns a document that contains an array of the distinct values. The return document also contains a subdocument with query statistics and the query plan.

When possible, the `distinct` (page 203) command uses an index to find documents and return values.

The command takes the following form:

```
{ distinct: "<collection>", key: "<field>", query: <query> }
```

The command contains the following fields:

field string distinct The name of the collection to query for distinct values.

field string key The field to collect distinct values from.

field document query A query specification to limit the input documents in the *distinct* analysis.

Examples Return an array of the distinct values of the field `ord_dt` from all documents in the `orders` collection:

```
db.runCommand ( { distinct: "orders", key: "ord_dt" } )
```

Return an array of the distinct values of the field `sku` in the subdocument `item` from all documents in the `orders` collection:

```
db.runCommand ( { distinct: "orders", key: "item.sku" } )
```

Return an array of the distinct values of the field `ord_dt` from the documents in the `orders` collection where the `price` is greater than 10:

```
db.runCommand ( { distinct: "orders",  
                  key: "ord_dt",  
                  query: { price: { $gt: 10 } }  
                } )
```

Note: MongoDB also provides the shell wrapper method `db.collection.distinct()` (page 29) for the `distinct` (page 203) command. Additionally, many MongoDB *drivers* also provide a wrapper method. Refer to the specific driver documentation.

group

Definition

group

Groups documents in a collection by the specified key and performs simple aggregation functions, such as computing counts and sums. The command is analogous to a `SELECT <...> GROUP BY` statement in SQL. The command returns a document with the grouped records as well as the command meta-data.

The `group` (page 204) command takes the following prototype form:

```
{  
  group:  
  {  
    ns: <namespace>,  
    key: <key>,  
    $reduce: <reduce function>,  
    $keyf: <key function>,  
    cond: <query>,  
    finalize: <finalize function>  
  }  
}
```

The command accepts a document with the following fields:

- field string ns** The collection from which to perform the group by operation.
- field document key** The field or fields to group. Returns a “key object” for use as the grouping key.
- field function \$reduce** An aggregation function that operates on the documents during the grouping operation. These functions may return a sum or a count. The function takes two arguments: the current document and an aggregation result document for that group.
- field document initial** Initializes the aggregation result document.
- field function \$keyf** Alternative to the `key` field. Specifies a function that creates a “key object” for use as the grouping key. Use `$keyf` instead of `key` to group by calculated fields rather than existing document fields.
- field document cond** The selection criteria to determine which documents in the collection to process. If you omit the `cond` field, `group` (page 204) processes all the documents in the collection for the group operation.

field function finalize A function that runs each item in the result set before `group` (page 204) returns the final value. This function can either modify the result document or replace the result document as a whole. Unlike the `$keyf` and `$reduce` fields that also specify a function, this field name is `finalize`, *not* `$finalize`.

For the shell, MongoDB provides a wrapper method `db.collection.group()` (page 47). However, the `db.collection.group()` (page 47) method takes the `keyf` field and the `reduce` field whereas the `group` (page 204) command takes the `$keyf` field and the `$reduce` field.

Behavior

Limits and Restrictions The `group` (page 204) command does not work with *sharded clusters*. Use the *aggregation framework* or *map-reduce* in *sharded environments*.

The result set must fit within the *maximum BSON document size* (page 604).

Additionally, in version 2.2, the returned array can contain at most 20,000 elements; i.e. at most 20,000 unique groupings. For group by operations that results in more than 20,000 unique groupings, use `mapReduce` (page 208). Previous versions had a limit of 10,000 elements.

Prior to 2.4, the `group` (page 204) command took the `mongod` (page 503) instance's JavaScript lock which blocked all other JavaScript execution.

mongo Shell JavaScript Functions/Properties Changed in version 2.4.

In MongoDB 2.4, *map-reduce operations* (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions **cannot** access certain global functions or properties, such as `db`, that are available in the `mongo` (page 527) shell.

When upgrading to MongoDB 2.4, you will need to refactor your code if your *map-reduce operations* (page 208), `group` (page 204) commands, or `$where` (page 391) operator expressions include any global shell functions or properties that are no longer available, such as `db`.

The following JavaScript functions and properties **are available** to *map-reduce operations* (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions in MongoDB 2.4:

Available Properties	Available Functions	
args MaxKey MinKey	assert() BinData() DBPointer() DBRef() doassert() emit() gc() HexData() hex_md5() isNumber() isObject() ISODate() isString()	Map() MD5() NumberInt() NumberLong() ObjectId() print() printjson() printjsononeline() sleep() Timestamp() tojson() tojsononeline() tojsonObject() UUID() version()

JavaScript in MongoDB

Although `group` (page 204) uses JavaScript, most interactions with MongoDB do not use JavaScript but use an idiomatic driver in the language of the interacting application.

Examples The following are examples of the `db.collection.group()` (page 47) method. The examples assume an `orders` collection with documents of the following prototype:

```
{
  _id: ObjectId("5085a95c8fada716c89d0021"),
  ord_dt: ISODate("2012-07-01T04:00:00Z"),
  ship_dt: ISODate("2012-07-02T04:00:00Z"),
  item:
    {
      sku: "abc123",
      price: 1.99,
      uom: "pcs",
      qty: 25
    }
}
```

Group by Two Fields The following example groups by the `ord_dt` and `item.sku` fields those documents that have `ord_dt` greater than 01/01/2012:

```
db.runCommand(
  {
    group:
      {
        ns: 'orders',
        key: { ord_dt: 1, 'item.sku': 1 },
        cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
        $reduce: function ( curr, result ) { },
        initial: { }
      }
  }
)
```

The result is a documents that contain the `retval` field which contains the group by records, the `count` field which contains the total number of documents grouped, the `keys` field which contains the number of unique groupings (i.e. number of elements in the `retval`), and the `ok` field which contains the command status:

```
{ "retval" :
  [ { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc123"},
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "abc456"},
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "bcd123"},
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item.sku" : "efg456"},
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "abc123"},
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "efg456"},
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item.sku" : "ijk123"},
    { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc123"},
    { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item.sku" : "abc456"},
    { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc123"},
    { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item.sku" : "abc456"}
  ],
  "count" : 13,
  "ok" : 1 }
```

```
"keys" : 11,
"ok" : 1 }
```

The method call is analogous to the SQL statement:

```
SELECT ord_dt, item_sku
FROM orders
WHERE ord_dt > '01/01/2012'
GROUP BY ord_dt, item_sku
```

Calculate the Sum The following example groups by the `ord_dt` and `item_sku` fields those documents that have `ord_dt` greater than 01/01/2012 and calculates the sum of the `qty` field for each grouping:

```
db.runCommand(
  { group:
    {
      ns: 'orders',
      key: { ord_dt: 1, 'item_sku': 1 },
      cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
      $reduce: function ( curr, result ) {
        result.total += curr.item.qty;
      },
      initial: { total : 0 }
    }
  }
)
```

The `retval` field of the returned document is an array of documents that contain the group by fields and the calculated aggregation field:

```
{ "retval" :
  [ { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item_sku" : "abc123", "total" : 25 },
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item_sku" : "abc456", "total" : 25 },
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item_sku" : "bcd123", "total" : 10 },
    { "ord_dt" : ISODate("2012-07-01T04:00:00Z"), "item_sku" : "efg456", "total" : 10 },
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item_sku" : "abc123", "total" : 25 },
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item_sku" : "efg456", "total" : 15 },
    { "ord_dt" : ISODate("2012-06-01T04:00:00Z"), "item_sku" : "ijk123", "total" : 20 },
    { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item_sku" : "abc123", "total" : 45 },
    { "ord_dt" : ISODate("2012-05-01T04:00:00Z"), "item_sku" : "abc456", "total" : 25 },
    { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item_sku" : "abc123", "total" : 25 },
    { "ord_dt" : ISODate("2012-06-08T04:00:00Z"), "item_sku" : "abc456", "total" : 25 }
  ],
  "count" : 13,
  "keys" : 11,
  "ok" : 1 }
```

The method call is analogous to the SQL statement:

```
SELECT ord_dt, item_sku, SUM(item_qty) as total
FROM orders
WHERE ord_dt > '01/01/2012'
GROUP BY ord_dt, item_sku
```

Calculate Sum, Count, and Average The following example groups by the calculated `day_of_week` field, those documents that have `ord_dt` greater than 01/01/2012 and calculates the sum, count, and average of the `qty` field for each grouping:

```
db.runCommand(
  {
    group:
      {
        ns: 'orders',
        $keyf: function(doc) {
          return { day_of_week: doc.ord_dt.getDay() };
        },
        cond: { ord_dt: { $gt: new Date( '01/01/2012' ) } },
        $reduce: function( curr, result ) {
          result.total += curr.item.qty;
          result.count++;
        },
        initial: { total : 0, count: 0 },
        finalize: function(result) {
          var weekdays = [
            "Sunday", "Monday", "Tuesday",
            "Wednesday", "Thursday",
            "Friday", "Saturday"
          ];
          result.day_of_week = weekdays[result.day_of_week];
          result.avg = Math.round(result.total / result.count);
        }
      }
  }
)
```

The `retval` field of the returned document is an array of documents that contain the group by fields and the calculated aggregation field:

```
{
  "retval" :
    [
      { "day_of_week" : "Sunday", "total" : 70, "count" : 4, "avg" : 18 },
      { "day_of_week" : "Friday", "total" : 110, "count" : 6, "avg" : 18 },
      { "day_of_week" : "Tuesday", "total" : 70, "count" : 3, "avg" : 23 }
    ],
  "count" : 13,
  "keys" : 3,
  "ok" : 1
}
```

See also:

<http://docs.mongodb.org/manualcore/aggregation>

mapReduce

mapReduce

The `mapReduce` (page 208) command allows you to run *map-reduce* aggregation operations over a collection. The `mapReduce` (page 208) command has the following prototype form:

```
db.runCommand(
  {
    mapReduce: <collection>,
    map: <function>,
    reduce: <function>,
    out: <output>,
    query: <document>,
```

```

    sort: <document>,
    limit: <number>,
    finalize: <function>,
    scope: <document>,
    jsMode: <boolean>,
    verbose: <boolean>
  }
)

```

Pass the name of the collection to the `mapReduce` command (i.e. `<collection>`) to use as the source documents to perform the map reduce operation. The command also accepts the following parameters:

field collection `mapReduce` The name of the collection on which you want to perform map-reduce.

field Javascript function `map` A JavaScript function that associates or “maps” a value with a key and emits the key and value pair.

See *Requirements for the map Function* (page 211) for more information.

field JavaScript function `reduce` A JavaScript function that “reduces” to a single object all the values associated with a particular key.

See *Requirements for the reduce Function* (page 212) for more information.

field string or document `out` Specifies the location of the result of the map-reduce operation. You can output to a collection, output to a collection with an action, or output inline. You may output to a collection when performing map reduce operations on the primary members of the set; on *secondary* members you may only use the `inline` output.

See *out Options* (page 212) for more information.

field document `query` Specifies the selection criteria using *query operators* (page 373) for determining the documents input to the `map` function.

field document `sort` Sorts the *input* documents. This option is useful for optimization. For example, specify the sort key to be the same as the emit key so that there are fewer reduce operations. The sort key must be in an existing index for this collection.

field number `limit` Specifies a maximum number of documents to return from the collection.

field Javascript function `finalize` Follows the `reduce` method and modifies the output.

See *Requirements for the finalize Function* (page 213) for more information.

field document `scope` Specifies global variables that are accessible in the `map`, `reduce` and `finalize` functions.

field Boolean `jsMode` Specifies whether to convert intermediate data into BSON format between the execution of the `map` and `reduce` functions. Defaults to `false`.

If `false`:

- Internally, MongoDB converts the JavaScript objects emitted by the `map` function to BSON objects. These BSON objects are then converted back to JavaScript objects when calling the `reduce` function.
- The map-reduce operation places the intermediate BSON objects in temporary, on-disk storage. This allows the map-reduce operation to execute over arbitrarily large data sets.

If `true`:

- Internally, the JavaScript objects emitted during `map` function remain as JavaScript objects. There is no need to convert the objects for the `reduce` function, which can result in faster execution.

- You can only use `jsMode` for result sets with fewer than 500,000 distinct key arguments to the mapper's `emit()` function.

The `jsMode` defaults to `false`.

field Boolean verbose Specifies whether to include the `timing` information in the result information. The `verbose` defaults to `true` to include the `timing` information.

The following is a prototype usage of the `mapReduce` (page 208) command:

```
var mapFunction = function() { ... };
var reduceFunction = function(key, values) { ... };

db.runCommand(
  {
    mapReduce: 'orders',
    map: mapFunction,
    reduce: reduceFunction,
    out: { merge: 'map_reduce_results', db: 'test' },
    query: { ord_date: { $gt: new Date('01/01/2012') } }
  }
)
```

JavaScript in MongoDB

Although `mapReduce` (page 208) uses JavaScript, most interactions with MongoDB do not use JavaScript but use an idiomatic driver in the language of the interacting application.

Note: Changed in version 2.4.

In MongoDB 2.4, `map-reduce operations` (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions **cannot** access certain global functions or properties, such as `db`, that are available in the `mongo` (page 527) shell.

When upgrading to MongoDB 2.4, you will need to refactor your code if your `map-reduce operations` (page 208), `group` (page 204) commands, or `$where` (page 391) operator expressions include any global shell functions or properties that are no longer available, such as `db`.

The following JavaScript functions and properties **are available** to `map-reduce operations` (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions in MongoDB 2.4:

Available Properties	Available Functions	
args MaxKey MinKey	assert() BinData() DBPointer() DBRef() doassert() emit() gc() HexData() hex_md5() isNumber() isObject() ISODate() isString()	Map() MD5() NumberInt() NumberLong() ObjectId() print() printjson() printjsononeline() sleep() Timestamp() tojson() tojsononeline() tojsonObject() UUID() version()

Requirements for the map Function The map function has the following prototype:

```
function() {
  ...
  emit(key, value);
}
```

The map function exhibits the following behaviors:

- In the map function, reference the current document as `this` within the function.
- The map function should *not* access the database for any reason.
- The map function should be pure, or have *no* impact outside of the function (i.e. side effects.)
- The `emit(key, value)` function associates the key with a value.
 - A single emit can only hold half of MongoDB's *maximum BSON document size* (page 604).
 - The map function can call `emit(key, value)` any number of times, including 0, per each input document.

The following map function may call `emit(key, value)` either 0 or 1 times depending on the value of the input document's `status` field:

```
function() {
  if (this.status == 'A')
    emit(this.cust_id, 1);
}
```

The following map function may call `emit(key, value)` multiple times depending on the number of elements in the input document's `items` field:

```
function() {
  this.items.forEach(function(item) { emit(item.sku, 1); });
}
```

- The `map` function can access the variables defined in the `scope` parameter.

Requirements for the `reduce` Function The `reduce` function has the following prototype:

```
function(key, values) {  
  ...  
  return result;  
}
```

The `reduce` function exhibits the following behaviors:

- The `reduce` function should *not* access the database, even to perform read operations.
- The `reduce` function should *not* affect the outside system.
- MongoDB will **not** call the `reduce` function for a key that has only a single value. The `values` argument is an array whose elements are the value objects that are “mapped” to the key.
- MongoDB can invoke the `reduce` function more than once for the same key. In this case, the previous output from the `reduce` function for that key will become one of the input values to the next `reduce` function invocation for that key.
- The `reduce` function can access the variables defined in the `scope` parameter.

Because it is possible to invoke the `reduce` function more than once for the same key, the following properties need to be true:

- the *type* of the return object must be **identical** to the type of the value emitted by the `map` function to ensure that the following operations is true:

```
reduce(key, [ C, reduce(key, [ A, B ]) ] ) == reduce( key, [ C, A, B ] )
```

- the `reduce` function must be *idempotent*. Ensure that the following statement is true:

```
reduce( key, [ reduce(key, valuesArray) ] ) == reduce( key, valuesArray )
```

- the order of the elements in the `valuesArray` should not affect the output of the `reduce` function, so that the following statement is true:

```
reduce( key, [ A, B ] ) == reduce( key, [ B, A ] )
```

out Options You can specify the following options for the `out` parameter:

Output to a Collection

```
out: <collectionName>
```

Output to a Collection with an Action This option is only available when passing `out` a collection that already exists. This option is not available on secondary members of replica sets.

```
out: { <action>: <collectionName>  
      [, db: <dbName>]  
      [, sharded: <boolean> ]  
      [, nonAtomic: <boolean> ] }
```

When you output to a collection with an action, the `out` has the following parameters:

- `<action>`: Specify one of the following actions:

- `replace`

Replace the contents of the `<collectionName>` if the collection with the `<collectionName>` exists.

- `merge`

Merge the new result with the existing result if the output collection already exists. If an existing document has the same key as the new result, *overwrite* that existing document.

- `reduce`

Merge the new result with the existing result if the output collection already exists. If an existing document has the same key as the new result, apply the `reduce` function to both the new and the existing documents and overwrite the existing document with the result.

- `db:`

Optional. The name of the database that you want the map-reduce operation to write its output. By default this will be the same database as the input collection.

- `sharded:`

Optional. If `true` and you have enabled sharding on output database, the map-reduce operation will shard the output collection using the `_id` field as the shard key.

- `nonAtomic:`

New in version 2.2.

Optional. Specify output operation as non-atomic and is valid *only* for `merge` and `reduce` output modes which may take minutes to execute.

If `nonAtomic` is `true`, the post-processing step will prevent MongoDB from locking the database; however, other clients will be able to read intermediate states of the output collection. Otherwise the map reduce operation must lock the database during post-processing.

Output Inline Perform the map-reduce operation in memory and return the result. This option is the only available option for `out` on secondary members of replica sets.

```
out: { inline: 1 }
```

The result must fit within the *maximum size of a BSON document* (page 604).

Requirements for the `finalize` Function The `finalize` function has the following prototype:

```
function(key, reducedValue) {
  ...
  return modifiedObject;
}
```

The `finalize` function receives as its arguments a key value and the `reducedValue` from the `reduce` function. Be aware that:

- The `finalize` function should *not* access the database for any reason.
- The `finalize` function should be pure, or have *no* impact outside of the function (i.e. side effects.)
- The `finalize` function can access the variables defined in the `scope` parameter.

Examples In the `mongo` (page 527) shell, the `db.collection.mapReduce()` (page 55) method is a wrapper around the `mapReduce` (page 208) command. The following examples use the `db.collection.mapReduce()` (page 55) method:

Consider the following map-reduce operations on a collection `orders` that contains documents of the following prototype:

```
{
  _id: ObjectId("50a8240b927d5d8b5891743c"),
  cust_id: "abc123",
  ord_date: new Date("Oct 04, 2012"),
  status: 'A',
  price: 25,
  items: [ { sku: "mmm", qty: 5, price: 2.5 },
            { sku: "nnn", qty: 5, price: 2.5 } ]
}
```

Return the Total Price Per Customer Perform the map-reduce operation on the `orders` collection to group by the `cust_id`, and calculate the sum of the `price` for each `cust_id`:

1. Define the map function to process each input document:

- In the function, `this` refers to the document that the map-reduce operation is processing.
- The function maps the `price` to the `cust_id` for each document and emits the `cust_id` and price pair.

```
var mapFunction1 = function() {
    emit(this.cust_id, this.price);
};
```

2. Define the corresponding reduce function with two arguments `keyCustId` and `valuesPrices`:

- The `valuesPrices` is an array whose elements are the price values emitted by the map function and grouped by `keyCustId`.
- The function reduces the `valuesPrice` array to the sum of its elements.

```
var reduceFunction1 = function(keyCustId, valuesPrices) {
    return Array.sum(valuesPrices);
};
```

3. Perform the map-reduce on all documents in the `orders` collection using the `mapFunction1` map function and the `reduceFunction1` reduce function.

```
db.orders.mapReduce(
    mapFunction1,
    reduceFunction1,
    { out: "map_reduce_example" }
)
```

This operation outputs the results to a collection named `map_reduce_example`. If the `map_reduce_example` collection already exists, the operation will replace the contents with the results of this map-reduce operation:

Calculate Order and Total Quantity with Average Quantity Per Item In this example, you will perform a map-reduce operation on the `orders` collection for all documents that have an `ord_date` value greater than 01/01/2012. The operation groups by the `item.sku` field, and calculates the number of orders and the total

quantity ordered for each sku. The operation concludes by calculating the average quantity per order for each sku value:

1. Define the map function to process each input document:

- In the function, `this` refers to the document that the map-reduce operation is processing.
- For each item, the function associates the sku with a new object value that contains the count of 1 and the item qty for the order and emits the sku and value pair.

```
var mapFunction2 = function() {
    for (var idx = 0; idx < this.items.length; idx++) {
        var key = this.items[idx].sku;
        var value = {
            count: 1,
            qty: this.items[idx].qty
        };
        emit(key, value);
    }
};
```

2. Define the corresponding reduce function with two arguments `keySKU` and `countObjVals`:

- `countObjVals` is an array whose elements are the objects mapped to the grouped `keySKU` values passed by map function to the reducer function.
- The function reduces the `countObjVals` array to a single object `reducedValue` that contains the count and the qty fields.
- In `reducedVal`, the `count` field contains the sum of the count fields from the individual array elements, and the `qty` field contains the sum of the qty fields from the individual array elements.

```
var reduceFunction2 = function(keySKU, countObjVals) {
    reducedVal = { count: 0, qty: 0 };

    for (var idx = 0; idx < countObjVals.length; idx++) {
        reducedVal.count += countObjVals[idx].count;
        reducedVal.qty += countObjVals[idx].qty;
    }

    return reducedVal;
};
```

3. Define a finalize function with two arguments `key` and `reducedVal`. The function modifies the `reducedVal` object to add a computed field named `avg` and returns the modified object:

```
var finalizeFunction2 = function (key, reducedVal) {

    reducedVal.avg = reducedVal.qty/reducedVal.count;

    return reducedVal;

};
```

4. Perform the map-reduce operation on the orders collection using the `mapFunction2`, `reduceFunction2`, and `finalizeFunction2` functions.

```
db.orders.mapReduce( mapFunction2,
                    reduceFunction2,
                    {
                        out: { merge: "map_reduce_example" },
                    }
```

```
    query: { ord_date:
              { $gt: new Date('01/01/2012') }
            },
    finalize: finalizeFunction2
  }
}
```

This operation uses the `query` field to select only those documents with `ord_date` greater than `new Date(01/01/2012)`. Then it output the results to a collection `map_reduce_example`. If the `map_reduce_example` collection already exists, the operation will merge the existing contents with the results of this map-reduce operation.

For more information and examples, see the [Map-Reduce](#) page and <http://docs.mongodb.org/manual/tutorial/perform-map-reduce>.

See also:

- <http://docs.mongodb.org/manual/tutorial/troubleshoot-map-function>
- <http://docs.mongodb.org/manual/tutorial/troubleshoot-reduce-function>
- `db.collection.mapReduce()` (page 55)
- <http://docs.mongodb.org/manual/core/aggregation>

For a detailed comparison of the different approaches, see [Aggregation Commands Comparison](#) (page 488).

Geospatial Commands

Geospatial Commands

Name	Description
geoNear (page 216)	Performs a geospatial query that returns the documents closest to a given point.
geoSearch (page 219)	Performs a geospatial query that uses MongoDB's <i>haystack index</i> functionality.
geoWalk (page 219)	An internal command to support geospatial queries.

geoNear

Definition

geoNear

Specifies a point for which a *geospatial* query returns the closest documents first. The query returns the documents from nearest to farthest. The [geoNear](#) (page 216) command provides an alternative to the `$near` (page 394) operator. In addition to the functionality of `$near` (page 394), [geoNear](#) (page 216) returns additional diagnostic information.

The [geoNear](#) (page 216) command accepts a *document* that contains the following fields. Specify all distances in the same units as the document coordinate system:

field string `geoNear` The collection to query.

field `GeoJSON point`, **term**: *legacy coordinate pairs* <*legacy coordinate pairs*> `near`:

The point for which to find the closest documents.

field number `limit` The maximum number of documents to return. The default value is 100. See also the `num` option.

field number num The `num` option provides the same function as the `limit` option. Both define the maximum number of documents to return. If both options are included, the `num` value overrides the `limit` value.

field number maxDistance A distance from the center point. Specify the distance in meters for *GeoJSON* data and in radians for *legacy coordinate pairs*. MongoDB limits the results to those documents that fall within the specified distance from the center point.

field document query Limits the results to the documents that match the query. The query syntax is the usual MongoDB *read operation query* syntax.

field Boolean spherical If `true`, MongoDB references points using a spherical surface. The default value is `false`.

field number distanceMultiplier The factor to multiply all distances returned by the query. For example, use the `distanceMultiplier` to convert radians, as returned by a spherical query, to kilometers by multiplying by the radius of the Earth.

field Boolean includeLocs If this is `true`, the query returns the location of the matching documents in the results. The default is `false`. This option is useful when a location field contains multiple locations. To specify a field within a subdocument, use *dot notation*.

field Boolean uniqueDocs If this value is `true`, the query returns a matching document once, even if more than one of the document's location fields match the query.

Deprecated since version 2.6: Geospatial queries no longer return duplicate results. The `$uniqueDocs` (page 400) operator has no impact on results.

Considerations The `geoNear` (page 216) command can use either a *GeoJSON* point or *legacy coordinate pairs*. Queries that use a 2d index return a limit of 100 documents.

The `geoNear` (page 216) command requires that a collection have *at most* only one 2d index and/or only one 2dsphere.

Behavior `geoNear` (page 216) always returns the documents sorted by distance. Any other sort order requires to sort the documents in memory, which can be inefficient. To return results in a different sort order, use the `$geoWithin` (page 392) operator in combination with `sort()`

Command Format To query a 2dsphere index, use the following syntax:

```
db.runCommand( { geoNear : <collection> ,
                  near : { type : "Point" ,
                          coordinates: [ <coordinates> ] } ,
                  spherical : true } )
```

To query a 2d index, use:

```
{ geoNear : <collection> , near : [ <coordinates> ] }
```

Example The following example runs the `geoNear` (page 216) command on the collection `places`:

```
db.runCommand( { geoNear : "places",
                  near : { type : "Point",
                          coordinates : [ -74.00, 40.00 ] },
                  spherical : true } )
```

The operation returns the following output:

```
{
  "results" : [
    {
      "dis" : 85753.24625705236,
      "obj" : {
        "_id" : ObjectId("536943463d2edb9288571e55"),
        "loc" : {
          "type" : "Point",
          "coordinates" : [
            -73.97,
            40.77
          ]
        },
        "name" : "Central Park",
        "category" : "Parks"
      }
    },
    {
      "dis" : 87422.73772813451,
      "obj" : {
        "_id" : ObjectId("536943603d2edb9288571e56"),
        "loc" : {
          "type" : "Point",
          "coordinates" : [
            -73.88,
            40.78
          ]
        },
        "name" : "La Guardia Airport",
        "category" : "Airport"
      }
    }
  ],
  "stats" : {
    "nscanned" : NumberLong(32),
    "objectsLoaded" : NumberLong(32),
    "avgDistance" : 86587.99199259344,
    "maxDistance" : 87422.73772813451,
    "time" : 2
  },
  "ok" : 1
}
```

Output The `geoNear` (page 216) command returns a document with the following fields:

geoNear.results

An array with the results of the `geoNear` (page 216) command, sorted by distance with the nearest result listed first and farthest last.

geoNear.results[n].dis

For each document in the results, the distance from the coordinates defined in the `geoNear` (page 216) command.

geoNear.results[n].obj

The document from the collection.

geoNear.stats

An object with statistics about the query used to return the results of the `geoNear` (page 216) search.

geoNear.stats.nscanned

The total number of index entries scanned during the database operation.

geoNear.stats.objectsLoaded

The total number of documents read from disk during the database operation.

geoNear.stats.avgDistance

The average distance between the coordinates defined in the [geoNear](#) (page 216) command and coordinates of the documents returned as results.

geoNear.stats.maxDistance

The maximum distance between the coordinates defined in the [geoNear](#) (page 216) command and coordinates of the documents returned as results.

geoNear.stats.time

The execution time of the database operation, in milliseconds.

geoNear.ok

A value of 1 indicates the [geoNear](#) (page 216) search succeeded. A value of 0 indicates an error.

geoSearch**geoSearch**

The [geoSearch](#) (page 219) command provides an interface to MongoDB's *haystack index* functionality. These indexes are useful for returning results based on location coordinates *after* collecting results based on some other query (i.e. a “haystack.”) Consider the following example:

```
{ geoSearch : "places", near : [33, 33], maxDistance : 6, search : { type : "restaurant" }, limit : 30 }
```

The above command returns all documents with a type of restaurant having a maximum distance of 6 units from the coordinates [30, 33] in the collection `places` up to a maximum of 30 results.

Unless specified otherwise, the [geoSearch](#) (page 219) command limits results to 50 documents.

Important: [geoSearch](#) (page 219) is not supported for sharded clusters.

geoWalk**geoWalk**

[geoWalk](#) (page 219) is an internal command.

Query and Write Operation Commands

Query and Write Operation Commands

Name	Description
<code>insert</code> (page 220)	Inserts one or more documents.
<code>update</code> (page 222)	Updates one or more documents.
<code>delete</code> (page 226)	Deletes one or more documents.
<code>findAndModify</code> (page 229)	Returns and modifies a single document.
<code>text</code> (page 235)	Performs a text search.
<code>getLastError</code> (page 235)	Returns the success status of the last operation.
<code>getPrevError</code> (page 237)	Returns status document containing all errors since the last <code>resetError</code> (page 237) command.
<code>resetError</code> (page 237)	Resets the last error status.
<code>eval</code> (page 238)	Runs a JavaScript function on the database server.
<code>parallelCollectionScan</code> (page 240)	Lets applications use multiple parallel cursors when reading documents from a collection.

insert

Definition

insert

New in version 2.6.

The `insert` (page 220) command inserts one or more documents and returns a document containing the status of all inserts. The insert methods provided by the MongoDB drivers use this command internally.

The command has the following syntax:

```
{
  insert: <collection>,
  documents: [ <document>, <document>, <document>, ... ],
  ordered: <boolean>,
  writeConcern: { <write concern> }
}
```

The `insert` (page 220) command takes the following fields:

field string insert The name of the target collection.

field array documents An array of one or more documents to insert into the named collection.

field boolean ordered If `true`, then when an insert of a document fails, return without inserting any remaining documents listed in the `inserts` array. If `false`, then when an insert of a document fails, continue to insert the remaining documents. Defaults to `true`.

field document writeConcern A document expressing the `write concern` of the `insert` (page 220) command. Omit to use the default write concern.

Returns A document that contains the status of the operation. See *Output* (page 221) for details.

Behavior The total size of all the `documents` array elements must be less than or equal to the `maximum BSON document size` (page 604).

The total number of documents in the `documents` array must be less than or equal to the `maximum bulk size` (page 608).

Examples

Insert a Single Document Insert a document into the `users` collection:

```
db.runCommand(
  {
    insert: "users",
    documents: [ { _id: 1, user: "abc123", status: "A" } ]
  }
)
```

The returned document shows that the command successfully inserted a document. See *Output* (page 221) for details.

```
{ "ok" : 1, "n" : 1 }
```

Bulk Insert Insert three documents into the `users` collection:

```
db.runCommand(
  {
    insert: "users",
    documents: [
      { _id: 2, user: "ijk123", status: "A" },
      { _id: 3, user: "xyz123", status: "P" },
      { _id: 4, user: "mop123", status: "P" }
    ],
    ordered: false,
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)
```

The returned document shows that the command successfully inserted the three documents. See *Output* (page 221) for details.

```
{ "ok" : 1, "n" : 3 }
```

Output The returned document contains a subset of the following fields:

`insert.ok`

The status of the command.

`insert.n`

The number of documents inserted.

`insert.writeErrors`

An array of documents that contains information regarding any error encountered during the insert operation. The `writeErrors` (page 221) array contains an error document for each insert that errors.

Each error document contains the following fields:

`insert.writeErrors.index`

An integer that identifies the document in the `documents` array, which uses a zero-based index.

`insert.writeErrors.code`

An integer value identifying the error.

`insert.writeErrors.errmsg`

A description of the error.

insert.writeConcernError

Document that describe error related to write concern and contains the field:

insert.writeConcernError.code

An integer value identifying the cause of the write concern error.

insert.writeConcernError.errmsg

A description of the cause of the write concern error.

The following is an example document returned for a successful `insert` (page 220) of a single document:

```
{ ok: 1, n: 1 }
```

The following is an example document returned for an `insert` (page 220) of two documents that successfully inserted one document but encountered an error with the other document:

```
{
  "ok" : 1,
  "n" : 1,
  "writeErrors" : [
    {
      "index" : 1,
      "code" : 11000,
      "errmsg" : "insertDocument :: caused by :: 11000 E11000 duplicate key error index: test.user"
    }
  ]
}
```

update

Definition

update

New in version 2.6.

The `update` (page 222) command modifies documents in a collection. A single `update` (page 222) command can contain multiple update statements. The update methods provided by the MongoDB drivers use this command internally.

The `update` (page 222) command has the following syntax:

```
{
  update: <collection>,
  updates:
    [
      { q: <query>, u: <update>, upsert: <boolean>, multi: <boolean> },
      { q: <query>, u: <update>, upsert: <boolean>, multi: <boolean> },
      { q: <query>, u: <update>, upsert: <boolean>, multi: <boolean> },
      ...
    ],
  ordered: <boolean>,
  writeConcern: { <write concern> }
}
```

The command takes the following fields:

field string update The name of the target collection.

field array updates An array of one or more update statements to perform in the named collection.

field boolean ordered If `true`, then when an update statement fails, return without performing the remaining update statements. If `false`, then when an update fails, continue with the remaining update statements, if any. Defaults to `true`.

field document writeConcern A document expressing the `write concern` of the `update` (page 222) command. Omit to use the default write concern.

Each element of the `updates` array contains the following sub-fields:

field string q The query that matches documents to update. Use the same *query selectors* (page 373) as used in the `find()` (page 34) method.

field document u The modifications to apply. For details, see *Behaviors* (page 223).

field boolean upsert If `true`, perform an insert if no documents match the query. If both `upsert` and `multi` are `true` and no documents match the query, the update operation inserts only a single document.

field boolean multi If `true`, updates all documents that meet the query criteria. If `false`, limit the update to one document that meet the query criteria. Defaults to `false`.

Returns A document that contains the status of the operation. See *Output* (page 225) for details.

Behaviors The `<update>` document can contain either all *update operator* (page 412) expressions or all `field:value` expressions.

Update Operator Expressions If the `<update>` document contains all *update operator* (page 412) expressions, as in:

```
{
  $set: { status: "D" },
  $inc: { quantity: 2 }
}
```

Then, the `update` (page 222) command updates only the corresponding fields in the document.

Field: Value Expressions If the `<update>` document contains *only* `field:value` expressions, as in:

```
{
  status: "D",
  quantity: 4
}
```

Then the `update` (page 222) command *replaces* the matching document with the update document. The `update` (page 222) command can only replace a *single* matching document; i.e. the `multi` field cannot be `true`. The `update` (page 222) command *does not* replace the `_id` value.

Limits For each update element in the `updates` array, the sum of the query and the update sizes (i.e. `q` and `u`) must be less than or equal to the *maximum BSON document size* (page 604).

The total number of update statements in the `updates` array must be less than or equal to the *maximum bulk size* (page 608).

Examples

Update Specific Fields of One Document Use *update operators* (page 412) to update only the specified fields of a document.

For example, given a `users` collection, the following command uses the `$set` (page 416) and `$inc` (page 412) operators to modify the `status` and the `points` fields respectively of a document where the `user` equals "abc123":

```
db.runCommand(
  {
    update: "users",
    updates: [
      {
        q: { user: "abc123" }, u: { $set: { status: "A" }, $inc: { points: 1 } }
      }
    ],
    ordered: false,
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)
```

Because `<update>` document does not specify the optional `multi` field, the update only modifies one document, even if more than one document matches the `q` match condition.

The returned document shows that the command found and updated a single document. See *Output* (page 225) for details.

```
{ "ok" : 1, "nModified" : 1, "n" : 1 }
```

Update Specific Fields of Multiple Documents Use *update operators* (page 412) to update only the specified fields of a document, and include the `multi` field set to `true` in the update statement.

For example, given a `users` collection, the following command uses the `$set` (page 416) and `$inc` (page 412) operators to modify the `status` and the `points` fields respectively of all documents in the collection:

```
db.runCommand(
  {
    update: "users",
    updates: [
      { q: { }, u: { $set: { status: "A" }, $inc: { points: 1 } }, multi: true }
    ],
    ordered: false,
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)
```

The update modifies all documents that match the query specified in the `q` field, namely the empty query which matches all documents in the collection.

The returned document shows that the command found and updated multiple documents. See *Output* (page 225) for details.

```
{ "ok" : 1, "nModified" : 100, "n" : 100 }
```

Bulk Update The following example performs multiple update operations on the `users` collection:

```
db.runCommand(
  {
    update: "users",
    updates: [
```

```

    { q: { status: "P" }, u: { $set: { status: "D" } }, multi: true },
    { q: { _id: 5 }, u: { _id: 5, name: "abc123", status: "A" }, upsert: true }
  ],
  ordered: false,
  writeConcern: { w: "majority", wtimeout: 5000 }
}
)

```

The returned document shows that the command modified 10 documents and upserted a document with the `_id` value 5. See *Output* (page 225) for details.

```

{
  "ok" : 1,
  "nModified" : 10,
  "n" : 11,
  "upserted" : [
    {
      "index" : 1,
      "_id" : 5
    }
  ]
}

```

Output The returned document contains a subset of the following fields:

`update.ok`

The status of the command.

`update.n`

The number of documents selected for update. If the update operation results in no change to the document, e.g. `$set` (page 416) expression updates the value to the current value, `n` (page 225) can be greater than `nModified` (page 225).

`update.nModified`

The number of documents updated. If the update operation results in no change to the document, such as setting the value of the field to its current value, `nModified` (page 225) can be less than `n` (page 225).

`update.upserted`

An array of documents that contains information for each upserted documents.

Each document contains the following information:

`update.upserted.index`

An integer that identifies the upsert statement in the `updates` array, which uses a zero-based index.

`update.upserted._id`

The `_id` value of the upserted document.

`update.writeErrors`

An array of documents that contains information regarding any error encountered during the update operation. The `writeErrors` (page 225) array contains an error document for each update statement that errors.

Each error document contains the following fields:

`update.writeErrors.index`

An integer that identifies the update statement in the `updates` array, which uses a zero-based index.

`update.writeErrors.code`

An integer value identifying the error.

`update.writeErrors.errmsg`

A description of the error.

`update.writeConcernError`

Document that describe error related to write concern and contains the field:

`update.writeConcernError.code`

An integer value identifying the cause of the write concern error.

`update.writeConcernError.errmsg`

A description of the cause of the write concern error.

The following is an example document returned for a successful `update` (page 222) command that performed an `upsert`:

```
{
  "ok" : 1,
  "nModified" : 0,
  "n" : 1,
  "upserted" : [
    {
      "index" : 0,
      "_id" : ObjectId("52ccb2118908ccd753d65882")
    }
  ]
}
```

The following is an example document returned for a bulk update involving three update statements, where one update statement was successful and two other update statements encountered errors:

```
{
  "ok" : 1,
  "nModified" : 1,
  "n" : 1,
  "writeErrors" : [
    {
      "index" : 1,
      "code" : 16837,
      "errmsg" : "The _id field cannot be changed from {_id: 1.0} to {_id: 5.0}."
    },
    {
      "index" : 2,
      "code" : 16837,
      "errmsg" : "The _id field cannot be changed from {_id: 2.0} to {_id: 6.0}."
    }
  ]
}
```

delete

Definition

delete

New in version 2.6.

The `delete` (page 226) command removes documents from a collection. A single `delete` (page 226) command can contain multiple delete specifications. The command cannot operate on capped collections. The remove methods provided by the MongoDB drivers use this command internally.

The `delete` (page 226) command has the following syntax:

```

{
  delete: <collection>,
  deletes: [
    { q : <query>, limit : <integer> },
    { q : <query>, limit : <integer> },
    { q : <query>, limit : <integer> },
    ...
  ],
  ordered: <boolean>,
  writeConcern: { <write concern> }
}

```

The command takes the following fields:

field string delete The name of the target collection.

field array deletes An array of one or more delete statements to perform in the named collection.

field boolean ordered If `true`, then when a delete statement fails, return without performing the remaining delete statements. If `false`, then when a delete statement fails, continue with the remaining delete statements, if any. Defaults to `true`.

field document writeConcern A document expressing the `write concern` of the `delete` (page 226) command. Omit to use the default write concern.

Each element of the `deletes` array contains the following sub-fields:

field string q The query that matches documents to delete.

field integer limit The number of matching documents to delete. Specify either a 0 to delete all matching documents or 1 to delete a single document.

Returns A document that contains the status of the operation. See *Output* (page 228) for details.

Behavior The total size of all the queries (i.e. the `q` field values) in the `deletes` array must be less than or equal to the `maximum BSON document size` (page 604).

The total number of delete documents in the `deletes` array must be less than or equal to the `maximum bulk size` (page 608).

Examples

Limit the Number of Documents Deleted The following example deletes from the `orders` collection one document that has the `status` equal to `D` by specifying the `limit` of 1:

```

db.runCommand(
  {
    delete: "orders",
    deletes: [ { q: { status: "D" }, limit: 1 } ]
  }
)

```

The returned document shows that the command deleted 1 document. See *Output* (page 228) for details.

```
{ "ok" : 1, "n" : 1 }
```

Delete All Documents That Match a Condition The following example deletes from the `orders` collection all documents that have the `status` equal to `D` by specifying the `limit` of 0:

```
db.runCommand(
  {
    delete: "orders",
    deletes: [ { q: { status: "D" }, limit: 0 } ],
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)
```

The returned document shows that the command found and deleted 13 documents. See [Output](#) (page 228) for details.

```
{ "ok" : 1, "n" : 13 }
```

Delete All Documents from a Collection Delete all documents in the `orders` collection by specifying an empty query condition *and* a `limit` of 0:

```
db.runCommand(
  {
    delete: "orders",
    deletes: [ { q: { }, limit: 0 } ],
    writeConcern: { w: "majority", wtimeout: 5000 }
  }
)
```

The returned document shows that the command found and deleted 35 documents in total. See [Output](#) (page 228) for details.

```
{ "ok" : 1, "n" : 35 }
```

Bulk Delete The following example performs multiple delete operations on the `orders` collection:

```
db.runCommand(
  {
    delete: "orders",
    deletes: [
      { q: { status: "D" }, limit: 0 },
      { q: { cust_num: 99999, item: "abc123", status: "A" }, limit: 1 }
    ],
    ordered: false,
    writeConcern: { w: 1 }
  }
)
```

The returned document shows that the command found and deleted 21 documents in total for the two delete statements. See [Output](#) (page 228) for details.

```
{ "ok" : 1, "n" : 21 }
```

Output The returned document contains a subset of the following fields:

`delete.ok`

The status of the command.

`delete.n`

The number of documents deleted.

delete.writeErrors

An array of documents that contains information regarding any error encountered during the delete operation. The `writeErrors` (page 228) array contains an error document for each delete statement that errors.

Each error document contains the following information:

delete.writeErrors.index

An integer that identifies the delete statement in the `deletes` array, which uses a zero-based index.

delete.writeErrors.code

An integer value identifying the error.

delete.writeErrors.errmsg

A description of the error.

delete.writeConcernError

Document that describe error related to write concern and contains the field:

delete.writeConcernError.code

An integer value identifying the cause of the write concern error.

delete.writeConcernError.errmsg

A description of the cause of the write concern error.

The following is an example document returned for a successful `delete` (page 226) command:

```
{ ok: 1, n: 1 }
```

The following is an example document returned for a `delete` (page 226) command that encountered an error:

```
{
  "ok" : 1,
  "n" : 0,
  "writeErrors" : [
    {
      "index" : 0,
      "code" : 10101,
      "errmsg" : "can't remove from a capped collection: test.cappedLog"
    }
  ]
}
```

findAndModify**Definition****findAndModify**

The `findAndModify` (page 229) command modifies and returns a single document. By default, the returned document does not include the modifications made on the update. To return the document with the modifications made on the update, use the `new` option.

The command has the following syntax:

```
{
  findAndModify: <string>,
  query: <document>,
  sort: <document>,
  remove: <boolean>,
  update: <document>,
  new: <boolean>,
}
```

```
fields: <document>,
upsert: <boolean>
}
```

The `findAndModify` (page 229) command takes the following fields:

field string findAndModify The collection against which to run the command.

param document query The selection criteria for the modification. The `query` field employs the same *query selectors* (page 373) as used in the `db.collection.find()` (page 34) method. Although the query may match multiple documents, `findAndModify` (page 229) will select only one document to modify.

param document sort Determines which document the operation modifies if the query selects multiple documents. `findAndModify` (page 229) modifies the first document in the sort order specified by this argument.

param Boolean remove Must specify either the `remove` or the `update` field. Removes the document specified in the `update` field. Set this to `true` to remove the selected document. The default is `false`.

param document update Must specify either the `remove` or the `update` field. Performs an update of the selected document. The `update` field employs the same *update operators* (page 412) or `field: value` specifications to modify the selected document.

param Boolean new When `true`, returns the modified document rather than the original. The `findAndModify` (page 229) method ignores the `new` option for `remove` operations. The default is `false`.

param document fields A subset of fields to return. The `fields` document specifies an inclusion of a field with `1`, as in: `fields: { <field1>: 1, <field2>: 1, ... }`. See *projection*.

param Boolean upsert Used in conjunction with the `update` field.

When `true`, `findAndModify` (page 229) creates a new document if no document matches the query, or if documents match the query, `findAndModify` (page 229) performs an update.

The default is `false`.

Output The return document contains the following fields:

- The `lastErrorObject` field that returns the details of the command:
 - The `updatedExisting` field **only** appears if the command specifies an `update` or an `update` with the `upsert` option; i.e. the field does not appear for a `remove`.
 - The `upserted` field **only** appears if the update with the `upsert` operation results in an insertion.
- The `value` field that returns either:
 - the original (i.e. pre-modification) document if `new` is `false`, or
 - the modified or inserted document if `new: true`.
- The `ok` field that returns the status of the command.

Note: If the `findAndModify` (page 229) finds no matching document, then:

- for `update` or `remove` operations, `lastErrorObject` does not appear in the return document and the `value` field holds a `null`.

```
{ "value" : null, "ok" : 1 }
```

- for update with an upsert operation that results in an insertion, if the command also specifies `new` is false **and** specifies a sort, the return document has a `lastErrorObject`, `value`, and `ok` fields, but the `value` field holds an empty document `{}`.
- for update with an upsert operation that results in an insertion, if the command also specifies `new` is false but does **not** specify a sort, the return document has a `lastErrorObject`, `value`, and `ok` fields, but the `value` field holds a null.

Behavior When the `findAndModify` (page 229) command includes the `upsert: true` option **and** the query field(s) is not uniquely indexed, the command could insert a document multiple times in certain circumstances. For instance, if multiple clients issue the `findAndModify` (page 229) command and these commands complete the `find` phase before any one starts the `modify` phase, these commands could insert the same document.

Consider an example where no document with the name `Andy` exists and multiple clients issue the following command:

```
db.runCommand(
  {
    findAndModify: "people",
    query: { name: "Andy" },
    sort: { rating: 1 },
    update: { $inc: { score: 1 } },
    upsert: true
  }
)
```

If all the commands finish the `query` phase before any command starts the `modify` phase, **and** there is no unique index on the `name` field, the commands may all perform an upsert. To prevent this condition, create a *unique index* on the `name` field. With the unique index in place, then the multiple `findAndModify` (page 229) commands would observe one of the following behaviors:

- Exactly one `findAndModify` (page 229) would successfully insert a new document.
- Zero or more `findAndModify` (page 229) commands would update the newly inserted document.
- Zero or more `findAndModify` (page 229) commands would fail when they attempted to insert a duplicate. If the command fails due to a unique index constraint violation, you can retry the command. Absent a delete of the document, the retry should not fail.

Sharded Collections When using `findAndModify` (page 229) in a *sharded* environment, the query must contain the *shard key* for all operations against the shard cluster. `findAndModify` (page 229) operations issued against `mongos` (page 518) instances for non-sharded collections function normally.

Concurrency This command obtains a write lock on the affected database and will block other operations until it has completed; however, typically the write lock is short lived and equivalent to other similar `update()` (page 69) operations.

Comparisons with the update Method When updating a document, `findAndModify` (page 229) and the `update()` (page 69) method operate differently:

- By default, both operations modify a single document. However, the `update()` (page 69) method with its `multi` option can modify more than one document.

- If multiple documents match the update criteria, for `findAndModify` (page 229), you can specify a `sort` to provide some measure of control on which document to update.

With the default behavior of the `update()` (page 69) method, you cannot specify which single document to update when multiple documents match.

- By default, `findAndModify` (page 229) method returns an object that contains the pre-modified version of the document, as well as the status of the operation. To obtain the updated document, use the `new` option.

The `update()` (page 69) method returns a `WriteResult` (page 188) object that contains the status of the operation. To return the updated document, use the `find()` (page 34) method. However, other updates may have modified the document between your update and the document retrieval. Also, if the update modified only a single document but multiple documents matched, you will need to use additional logic to identify the updated document.

- You cannot specify a `write concern` to `findAndModify` (page 229) to override the default write concern whereas, starting in MongoDB 2.6, you can specify a write concern to the `update()` (page 69) method.

When modifying a *single* document, both `findAndModify` (page 229) and the `update()` (page 69) method *atomically* update the document. See <http://docs.mongodb.org/manual/tutorial/isolate-sequence-of-operations> for more details about interactions and order of operations of these methods.

Examples

Update and Return The following command updates an existing document in the `people` collection where the document matches the `query` criteria:

```
db.runCommand(
  {
    findAndModify: "people",
    query: { name: "Tom", state: "active", rating: { $gt: 10 } },
    sort: { rating: 1 },
    update: { $inc: { score: 1 } }
  }
)
```

This command performs the following actions:

1. The query finds a document in the `people` collection where the `name` field has the value `Tom`, the `state` field has the value `active` and the `rating` field has a value *greater than* (page 373) 10.
2. The `sort` orders the results of the query in ascending order. If multiple documents meet the query condition, the command will select for modification the first document as ordered by this `sort`.
3. The update increments the value of the `score` field by 1.
4. The command returns a document with the following fields:
 - The `lastErrorObject` field that contains the details of the command, including the field `updatedExisting` which is `true`, and
 - The `value` field that contains the original (i.e. pre-modification) document selected for this update:

```
{
  "lastErrorObject" : {
    "updatedExisting" : true,
    "n" : 1,
    "connectionId" : 1,
    "err" : null,
  },
  "value" : {
    "name" : "Tom",
    "state" : "active",
    "rating" : 11,
    "score" : 1
  }
}
```

```

    "ok" : 1
  },
  "value" : {
    "_id" : ObjectId("50f1d54e9beb36a0f45c6452"),
    "name" : "Tom",
    "state" : "active",
    "rating" : 100,
    "score" : 5
  },
  "ok" : 1
}

```

To return the modified document in the `value` field, add the `new:true` option to the command.

If no document match the `query` condition, the command returns a document that contains `null` in the `value` field:

```
{ "value" : null, "ok" : 1 }
```

The `mongo` (page 527) shell and many *drivers* provide a `findAndModify()` (page 39) helper method. Using the shell helper, this previous operation can take the following form:

```

db.people.findAndModify( {
  query: { name: "Tom", state: "active", rating: { $gt: 10 } },
  sort: { rating: 1 },
  update: { $inc: { score: 1 } }
} );

```

However, the `findAndModify()` (page 39) shell helper method returns only the unmodified document, or if `new` is `true`, the modified document.

```

{
  "_id" : ObjectId("50f1d54e9beb36a0f45c6452"),
  "name" : "Tom",
  "state" : "active",
  "rating" : 100,
  "score" : 5
}

```

Upsert The following `findAndModify` (page 229) command includes the `upsert: true` option for the update operation to either update a matching document or, if no matching document exists, create a new document:

```

db.runCommand(
  {
    findAndModify: "people",
    query: { name: "Gus", state: "active", rating: 100 },
    sort: { rating: 1 },
    update: { $inc: { score: 1 } },
    upsert: true
  }
)

```

If the command finds a matching document, the command performs an update.

If the command does **not** find a matching document, the update with *upsert* operation results in an insertion and returns a document with the following fields:

- The `lastErrorObject` field that contains the details of the command, including the field `upserted` that contains the `ObjectId` of the newly inserted document, and

- The `value` field that contains an empty document `{}` as the original document because the command included the `sort` option:

```
{
  "lastErrorObject" : {
    "updatedExisting" : false,
    "upserted" : ObjectId("50f2329d0092b46dae1dc98e"),
    "n" : 1,
    "connectionId" : 1,
    "err" : null,
    "ok" : 1
  },
  "value" : {

  },
  "ok" : 1
}
```

If the command did **not** include the `sort` option, the `value` field would contain `null`:

```
{
  "value" : null,
  "lastErrorObject" : {
    "updatedExisting" : false,
    "n" : 1,
    "upserted" : ObjectId("5102f7540cb5c8be998c2e99")
  },
  "ok" : 1
}
```

Return New Document The following `findAndModify` (page 229) command includes both `upsert: true` option and the `new:true` option. The command either updates a matching document and returns the updated document or, if no matching document exists, inserts a document and returns the newly inserted document in the `value` field.

In the following example, no document in the `people` collection matches the query condition:

```
db.runCommand(
  {
    findAndModify: "people",
    query: { name: "Pascal", state: "active", rating: 25 },
    sort: { rating: 1 },
    update: { $inc: { score: 1 } },
    upsert: true,
    new: true
  }
)
```

The command returns the newly inserted document in the `value` field:

```
{
  "lastErrorObject" : {
    "updatedExisting" : false,
    "upserted" : ObjectId("50f47909444c11ac2448a5ce"),
    "n" : 1,
    "connectionId" : 1,
    "err" : null,
    "ok" : 1
  },
  "value" : {
    "name" : "Pascal",
    "state" : "active",
    "rating" : 25,
    "score" : 1
  },
  "ok" : 1
}
```

```

"value" : {
  "_id" : ObjectId("50f47909444c11ac2448a5ce"),
  "name" : "Pascal",
  "rating" : 25,
  "score" : 1,
  "state" : "active"
},
"ok" : 1
}

```

text

Definition

text

Deprecated since version 2.6: Use `$text` (page 387) query operator instead.

For document on the `text` (page 235), refer to the 2.4 Manual [2.4 Manual](http://docs.mongodb.org/v2.4/reference/command/text)¹⁶.

getLastError

Definition

getLastError

Changed in version 2.6: A new protocol for *write operations* (page 623) integrates write concerns with the write operations, eliminating the need for a separate `getLastError` (page 235) command. Write methods now return the status of the write operation, including error information.

In previous versions, clients typically used the `getLastError` (page 235) command in combination with the write operations to ensure that the write succeeds.

Returns the error status of the preceding write operation on the *current connection*.

`getLastError` (page 235) uses the following prototype form:

```
{ getLastError: 1 }
```

`getLastError` (page 235) uses the following fields:

field Boolean j If `true`, wait for the next journal commit before returning, rather than waiting for a full disk flush. If `mongod` (page 503) does not have journaling enabled, this option has no effect. If this option is enabled for a write operation, `mongod` (page 503) will wait *no more* than 1/3 of the current `commitIntervalMs` before writing data to the journal.

field integer,string w When running with replication, this is the number of servers to replicate to before returning. A `w` value of 1 indicates the primary only. A `w` value of 2 includes the primary and at least one secondary, etc. In place of a number, you may also set `w` to `majority` to indicate that the command should wait until the latest write propagates to a majority of replica set members. If using `w`, you should also use `wtimeout`. Specifying a value for `w` without also providing a `wtimeout` may cause `getLastError` (page 235) to block indefinitely.

field Boolean fsync If `true`, wait for `mongod` (page 503) to write this data to disk before returning. Defaults to `false`. In most cases, use the `j` option to ensure durability and consistency of the data set.

¹⁶<http://docs.mongodb.org/v2.4/reference/command/text>

field integer wtimeout Milliseconds. Specify a value in milliseconds to control how long to wait for write propagation to complete. If replication does not complete in the given timeframe, the `getLastError` (page 235) command will return with an error status.

See also:

Write Concern, <http://docs.mongodb.org/manualreference/write-concern>, and *replica-set-write-concern*.

Output Each `getLastError()` command returns a document containing a subset of the fields listed below.

getLastError.ok

`ok` (page 236) is `true` when the `getLastError` (page 235) command completes successfully.

Note: A value of `true` does *not* indicate that the preceding operation did not produce an error.

getLastError.err

`err` (page 236) is `null` unless an error occurs. When there was an error with the preceding operation, `err` contains a textual description of the error.

getLastError.code

`code` (page 236) reports the preceding operation's error code.

getLastError.connectionId

The identifier of the connection.

getLastError.lastOp

When issued against a replica set member and the preceding operation was a write or update, `lastOp` (page 236) is the *optime* timestamp in the *oplog* of the change.

getLastError.n

`n` (page 236) reports the number of documents updated or removed, if the preceding operation was an update or remove operation.

getLastError.shards

When issued against a sharded cluster after a write operation, `shards` (page 236) identifies the shards targeted in the write operation. `shards` (page 236) is present in the output only if the write operation targets multiple shards.

getLastError.singleShard

When issued against a sharded cluster after a write operation, identifies the shard targeted in the write operation. `singleShard` (page 236) is only present if the write operation targets exactly one shard.

getLastError.updatedExisting

`updatedExisting` (page 236) is `true` when an update affects at least one document and does not result in an *upsert*.

getLastError.upserted

If the update results in an insert, `upserted` (page 236) is the value of `_id` field of the document.

Changed in version 2.6: Earlier versions of MongoDB included `upserted` (page 236) only if `_id` was an *ObjectId*.

getLastError.wnote

If set, `wnote` indicates that the preceding operation's error relates to using the `w` parameter to `getLastError` (page 235).

See

<http://docs.mongodb.org/manualreference/write-concern> for more information about w values.

`getLastError.wtimeout`

`wtimeout` (page 237) is `true` if the `getLastError` (page 235) timed out because of the `wtimeout` setting to `getLastError` (page 235).

`getLastError.waited`

If the preceding operation specified a timeout using the `wtimeout` setting to `getLastError` (page 235), then `waited` (page 237) reports the number of milliseconds `getLastError` (page 235) waited before timing out.

`getLastError.wtime`

`getLastError.wtime` (page 237) is the number of milliseconds spent waiting for the preceding operation to complete. If `getLastError` (page 235) timed out, `wtime` (page 237) and `getLastError.waited` are equal.

Examples

Confirm Replication to Two Replica Set Members The following example ensures the operation has replicated to two members (the primary and one other member):

```
db.runCommand( { getLastError: 1, w: 2 } )
```

Confirm Replication to a Majority of a Replica Set The following example ensures the write operation has replicated to a majority of the configured members of the set.

```
db.runCommand( { getLastError: 1, w: "majority" } )
```

Changed in version 2.6: In Master/Slave deployments, MongoDB treats `w: "majority"` as equivalent to `w: 1`. In earlier versions of MongoDB, `w: "majority"` produces an error in master/slave deployments.

Set a Timeout for a `getLastError` Response Unless you specify a timeout, a `getLastError` (page 235) command may block forever if MongoDB cannot satisfy the requested write concern. To specify a timeout of 5000 milliseconds, use an invocation that resembles the following:

```
db.runCommand( { getLastError: 1, w: 2, wtimeout: 5000 } )
```

When `wtimeout` is 0, the `getLastError` (page 235) operation will never time out.

`getPrevError`

`getPrevError`

The `getPrevError` (page 237) command returns the errors since the last `resetError` (page 237) command.

See also:

`db.getPrevError()` (page 111)

`resetError`

resetError

The `resetError` (page 237) command resets the last error status.

See also:

`db.resetError()` (page 117)

eval**eval**

The `eval` (page 238) command evaluates JavaScript functions on the database server.

If authentication is enabled, you must have access to all actions on all resources in order to run `eval` (page 238). Providing such access is not recommended, but if your organization requires a user to run `eval` (page 238), create a role that grants `anyAction` on *resource-anyresource*. Do not assign this role to any other user.

The `eval` (page 238) command has the following form:

```
{
  eval: <function>,
  args: [ <arg1>, <arg2> ... ],
  nolog: <boolean>
}
```

The command contains the following fields:

field function eval A JavaScript function.

field array args An array of arguments to pass to the JavaScript function. Omit if the function does not take arguments.

field boolean nolog By default, `eval` (page 238) takes a global write lock before evaluating the JavaScript function. As a result, `eval` (page 238) blocks all other read and write operations to the database while the `eval` (page 238) operation runs. Set `nolog` to `true` on the `eval` (page 238) command to prevent the `eval` (page 238) command from taking the global write lock before evaluating the JavaScript. `nolog` does not impact whether operations within the JavaScript code itself takes a write lock.

JavaScript in MongoDB

Although `eval` (page 238) uses JavaScript, most interactions with MongoDB do not use JavaScript but use an idiomatic driver in the language of the interacting application.

Behavior The following example uses `eval` (page 238) to perform an increment and calculate the average on the server:

```
db.runCommand( {
  eval: function(name, incAmount) {
    var doc = db.myCollection.findOne( { name : name } );

    doc = doc || { name : name , num : 0 , total : 0 , avg : 0 };

    doc.num++;
    doc.total += incAmount;
    doc.avg = doc.total / doc.num;

    db.myCollection.save( doc );
    return doc;
  },
```

```

    args: [ "eliot", 5 ]
  }
);

```

The `db` in the function refers to the current database.

The `mongo` (page 527) shell provides a helper method `db.eval()` (page 108)¹⁷, so you can express the above as follows:

```

db.eval( function(name, incAmount) {
    var doc = db.myCollection.findOne( { name : name } );

    doc = doc || { name : name , num : 0 , total : 0 , avg : 0 };

    doc.num++;
    doc.total += incAmount;
    doc.avg = doc.total / doc.num;

    db.myCollection.save( doc );
    return doc;
},
"eliot", 5 );

```

If you want to use the server's interpreter, you must run `eval` (page 238). Otherwise, the `mongo` (page 527) shell's JavaScript interpreter evaluates functions entered directly into the shell.

If an error occurs, `eval` (page 238) throws an exception. The following invalid function uses the variable `x` without declaring it as an argument:

```

db.runCommand(
  {
    eval: function() { return x + x; },
    args: [ 3 ]
  }
)

```

The statement will result in the following exception:

```

{
  "errmsg" : "exception: JavaScript execution failed: ReferenceError: x is not defined near '{ return",
  "code" : 16722,
  "ok" : 0
}

```

¹⁷ The helper `db.eval()` (page 108) in the `mongo` (page 527) shell wraps the `eval` (page 238) command. Therefore, the helper method shares the characteristics and behavior of the underlying command with *one exception*: `db.eval()` (page 108) method does not support the `nolock` option.

Warning:

- By default, `eval` (page 238) takes a global write lock before evaluating the JavaScript function. As a result, `eval` (page 238) blocks all other read and write operations to the database while the `eval` (page 238) operation runs. Set `nolock` to `true` on the `eval` (page 238) command to prevent the `eval` (page 238) command from taking the global write lock before evaluating the JavaScript. `nolock` does not impact whether operations within the JavaScript code itself takes a write lock.
 - Do not use `eval` (page 238) for long running operations as `eval` (page 238) blocks all other operations. Consider using other server side code execution options.
 - You can not use `eval` (page 238) with *sharded* data. In general, you should avoid using `eval` (page 238) in *sharded cluster*; nevertheless, it is possible to use `eval` (page 238) with non-sharded collections and databases stored in a *sharded cluster*.
 - With authentication enabled, `eval` (page 238) will fail during the operation if you do not have the permission to perform a specified task.
- Changed in version 2.4: You must have full admin access to run.

Changed in version 2.4: The V8 JavaScript engine, which became the default in 2.4, allows multiple JavaScript operations to execute at the same time. Prior to 2.4, `eval` (page 238) executed in a single thread.

See also:

<http://docs.mongodb.org/manualcore/server-side-javascript>

parallelCollectionScan**parallelCollectionScan**

New in version 2.6.

Allows applications to use multiple parallel cursors when reading all the documents from a collection, thereby increasing throughput. The `parallelCollectionScan` (page 240) command returns a document that contains an array of cursor information.

The command has the following syntax:

```
{
  parallelCollectionScan: "<collection>",
  numCursors: <integer>
}
```

The `parallelCollectionScan` (page 240) command takes the following fields:

field string parallelCollectionScan The name of the collection.

field integer numCursors The maximum number of cursors to return. Must be between 1 and 10000, inclusive.

Query Plan Cache Commands**Query Plan Cache Commands**

Name	Description
<code>planCacheListFilters</code> (page 241)	Lists the index filters for a collection.
<code>planCacheSetFilter</code> (page 242)	Sets an index filter for a collection.
<code>planCacheClearFilters</code> (page 243)	Clears index filter(s) for a collection.
<code>planCacheListQueryShapes</code> (page 245)	Displays the query shapes for which cached query plans exist.
<code>planCacheListPlans</code> (page 246)	Displays the cached query plans for the specified query shape.
<code>planCacheClear</code> (page 247)	Removes cached query plan(s) for a collection.

planCacheListFilters

Definition

planCacheListFilters

New in version 2.6.

Lists the *index filters* associated with *query shapes* for a collection.

The command has the following syntax:

```
db.runCommand( { planCacheListFilters: <collection> } )
```

The `planCacheListFilters` (page 241) command has the following field:

field string planCacheListFilters The name of the collection.

Returns Document listing the index filters. See *Output* (page 241).

Required Access A user must have access that includes the `planCacheIndexFilter` action.

Output The `planCacheListFilters` (page 241) command returns the document with the following form:

```
{
  "filters" : [
    {
      "query" : <query>
      "sort" : <sort>,
      "projection" : <projection>,
      "indexes" : [
        <index1>,
        ...
      ]
    },
    ...
  ],
  "ok" : 1
}
```

planCacheListFilters.filters

The array of documents that contain the index filter information.

Each document contains the following fields:

planCacheListFilters.filters.query

The query predicate associated with this filter. Although the *query* (page 241) shows the specific values used to create the index filter, the values in the predicate are insignificant; i.e. query predicates cover similar queries that differ only in the values.

For instance, a *query* (page 241) predicate of { "type": "electronics", "status" : "A" } covers the following query predicates:

```
{ type: "food", status: "A" }
{ type: "utensil", status: "D" }
```

Together with the *sort* (page 241) and the *projection* (page 242), the *query* (page 241) make up the *query shape* for the specified index filter.

`planCacheListFilters.filters.sort`

The sort associated with this filter. Can be an empty document.

Together with the [query](#) (page 241) and the [projection](#) (page 242), the [sort](#) (page 241) make up the *query shape* for the specified index filter.

`planCacheListFilters.filters.projection`

The projection associated with this filter. Can be an empty document.

Together with the [query](#) (page 241) and the [sort](#) (page 241), the [projection](#) (page 242) make up the *query shape* for the specified index filter.

`planCacheListFilters.filters.indexes`

The array of indexes for this *query shape*. To choose the optimal query plan, the query optimizer evaluates only the listed *indexes* and the collection scan.

`planCacheListFilters.ok`

The status of the command.

See also:

[planCacheClearFilters](#) (page 243), [planCacheSetFilter](#) (page 242)

planCacheSetFilter

Definition

planCacheSetFilter

New in version 2.6.

Set an *index filter* for a collection. If an index filter already exists for the *query shape*, the command overrides the previous index filter.

The command has the following syntax:

```
db.runCommand(  
  {  
    planCacheSetFilter: <collection>,  
    query: <query>,  
    sort: <sort>,  
    projection: <projection>,  
    indexes: [ <index1>, <index2>, ...]  
  }  
)
```

The [planCacheSetFilter](#) (page 242) command has the following field:

field string planCacheSetFilter The name of the collection.

field document query The query predicate associated with the index filter. Together with the [sort](#) and the [projection](#), the query predicate make up the *query shape* for the specified index filter.

Only the structure of the predicate, including the field names, are significant; the values in the query predicate are insignificant. As such, query predicates cover similar queries that differ only in the values.

field document sort The sort associated with the filter. Together with the [query](#) and the [projection](#), the [sort](#) make up the *query shape* for the specified index filter.

field document projection The projection associated with the filter. Together with the [query](#) and the [sort](#), the [projection](#) make up the *query shape* for the specified index filter.

field array indexes An array of index specification documents that act as the index filter for the specified *query shape*. Because the query optimizer chooses among the collection scan and these indexes, if the indexes are non-existent, the optimizer will choose the collection scan.

Index filters only exist for the duration of the server process and do not persist after shutdown; however, you can also clear existing index filters using the `planCacheClearFilters` (page 243) command.

Required Access A user must have access that includes the `planCacheIndexFilter` action.

Examples The following example creates an index filter on the `orders` collection such that for queries that consists only of an equality match on the `status` field without any projection and sort, the query optimizer evaluates only the two specified indexes and the collection scan for the winning plan:

```
db.runCommand(
  {
    planCacheSetFilter: "orders",
    query: { status: "A" },
    indexes: [
      { cust_id: 1, status: 1 },
      { status: 1, order_date: -1 }
    ]
  }
)
```

In the query predicate, only the structure of the predicate, including the field names, are significant; the values are insignificant. As such, the created filter applies to the following operations:

```
db.orders.find( { status: "D" } )
db.orders.find( { status: "P" } )
```

To see whether MongoDB applied an index filter for a query, check the `explain.filterSet` (page 84) field of the `explain()` (page 80) output.

See also:

`planCacheClearFilters` (page 243), `planCacheListFilters` (page 241)

planCacheClearFilters

Definition

planCacheClearFilters

New in version 2.6.

Removes *index filters* on a collection. Although index filters only exist for the duration of the server process and do not persist after shutdown, you can also clear existing index filters with the `planCacheClearFilters` (page 243) command.

Specify the *query shape* to remove a specific index filter. Omit the query shape to clear all index filters on a collection.

The command has the following syntax:

```
db.runCommand(
  {
    planCacheClearFilters: <collection>,
    query: <query pattern>,
    sort: <sort specification>,
```

```
    projection: <projection specification>
  }
)
```

The `planCacheClearFilters` (page 243) command has the following field:

field string planCacheClearFilters The name of the collection.

field document query The query predicate associated with the filter to remove. If omitted, clears all filters from the collection.

The values in the `query` predicate are insignificant in determining the *query shape*, so the values used in the query need not match the values shown using `planCacheListFilters` (page 241).

field document sort The sort associated with the filter to remove, if any.

field document projection The projection associated with the filter to remove, if any.

Required Access A user must have access that includes the `planCacheIndexFilter` action.

Examples

Clear Specific Index Filter on Collection The `orders` collection contains the following two filters:

```
{
  "query" : { "status" : "A" },
  "sort"  : { "ord_date" : -1 },
  "projection" : { },
  "indexes" : [ { "status" : 1, "cust_id" : 1 } ]
}

{
  "query" : { "status" : "A" },
  "sort"  : { },
  "projection" : { },
  "indexes" : [ { "status" : 1, "cust_id" : 1 } ]
}
```

The following command removes the second index filter only:

```
db.runCommand(
  {
    planCacheClearFilters: "orders",
    query: { "status" : "A" }
  }
)
```

Because the values in the `query` predicate are insignificant in determining the *query shape*, the following command would also remove the second index filter:

```
db.runCommand(
  {
    planCacheClearFilters: "orders",
    query: { "status" : "P" }
  }
)
```


Clear all Index Filters on a Collection The following example clears all index filters on the `orders` collection:

```
db.runCommand(
  {
    planCacheClearFilters: "orders"
  }
)
```

See also:

`planCacheListFilters` (page 241), `planCacheSetFilter` (page 242)

`planCacheListQueryShapes`

Definition

`planCacheListQueryShapes`

New in version 2.6.

Displays the *query shapes* for which cached query plans exist for a collection.

The query optimizer only caches the plans for those query shapes that can have more than one viable plan.

The `mongo` (page 527) shell provides the wrapper `PlanCache.listQueryShapes()` (page 124) for this command.

The command has the following syntax:

```
db.runCommand(
  {
    planCacheListQueryShapes: <collection>
  }
)
```

The `planCacheListQueryShapes` (page 245) command has the following field:

field string `planCacheListQueryShapes` The name of the collection.

Returns A document that contains an array of *query shapes* for which cached query plans exist.

Required Access On systems running with authorization, a user must have access that includes the `planCacheRead` action.

Example The following returns the *query shapes* that have cached plans for the `orders` collection:

```
db.runCommand(
  {
    planCacheListQueryShapes: "orders"
  }
)
```

The command returns a document that contains the field `shapes` that contains an array of the *query shapes* currently in the cache. In the example, the `orders` collection had cached query plans associated with the following shapes:

```
{
  "shapes" : [
    {
      "query" : { "qty" : { "$gt" : 10 } },
      "sort" : { "ord_date" : 1 },
      "projection" : { }
```

```
    },
    {
      "query" : { "$or" : [ { "qty" : { "$gt" : 15 } }, { "status" : "A" } ] },
      "sort" : { },
      "projection" : { }
    },
    {
      "query" : { "$or" :
        [
          { "qty" : { "$gt" : 15 }, "item" : "xyz123" },
          { "status" : "A" }
        ]
      },
      "sort" : { },
      "projection" : { }
    }
  ],
  "ok" : 1
}
```

See also:

- `PlanCache.listQueryShapes()` (page 124)

planCacheListPlans**Definition****planCacheListPlans**

New in version 2.6.

Displays the cached query plans for the specified *query shape*.

The query optimizer only caches the plans for those query shapes that can have more than one viable plan.

The `mongo` (page 527) shell provides the wrapper `PlanCache.getPlansByQuery()` (page 125) for this command.

The `planCacheListPlans` (page 246) command has the following syntax:

```
db.runCommand(
  {
    planCacheListPlans: <collection>,
    query: <query>,
    sort: <sort>,
    projection: <projection>
  }
)
```

The `planCacheListPlans` (page 246) command has the following field:

field document query The query predicate of the *query shape*. Only the structure of the predicate, including the field names, are significant to the shape; the values in the query predicate are insignificant.

field document projection The projection associated with the *query shape*.

field document sort The sort associated with the *query shape*.

To see the query shapes for which cached query plans exist, use the `planCacheListQueryShapes` (page 245) command.

Required Access On systems running with authorization, a user must have access that includes the `planCacheRead` action.

Example If a collection `orders` has the following query shape:

```
{
  "query" : { "qty" : { "$gt" : 10 } },
  "sort" : { "ord_date" : 1 },
  "projection" : { }
}
```

The following operation displays the query plan cached for the shape:

```
db.runCommand(
  {
    planCacheListPlans: "orders",
    query: { "qty" : { "$gt" : 10 } },
    sort: { "ord_date" : 1 }
  }
)
```

See also:

- `planCacheListQueryShapes` (page 245)
- `PlanCache.getPlansByQuery()` (page 125)
- `PlanCache.listQueryShapes()` (page 124)

planCacheClear

Definition

planCacheClear

New in version 2.6.

Removes cached query plans for a collection. Specify a *query shape* to remove cached query plans for that shape. Omit the query shape to clear all cached query plans.

The command has the following syntax:

```
db.runCommand(
  {
    planCacheClear: <collection>,
    query: <query>,
    sort: <sort>,
    projection: <projection>
  }
)
```

The `planCacheClear` (page 247) command has the following field:

field document query The query predicate of the *query shape*. Only the structure of the predicate, including the field names, are significant to the shape; the values in the query predicate are insignificant.

field document projection The projection associated with the *query shape*.

field document sort The sort associated with the *query shape*.

To see the query shapes for which cached query plans exist, use the `planCacheListQueryShapes` (page 245) command.

Required Access On systems running with authorization, a user must have access that includes the `planCacheWrite` action.

Examples

Clear Cached Plans for a Query Shape If a collection `orders` has the following query shape:

```
{
  "query" : { "qty" : { "$gt" : 10 } },
  "sort" : { "ord_date" : 1 },
  "projection" : { }
}
```

The following operation clears the query plan cached for the shape:

```
db.runCommand(
  {
    planCacheClear: "orders",
    query: { "qty" : { "$gt" : 10 } },
    sort: { "ord_date" : 1 }
  }
)
```

Clear All Cached Plans for a Collection The following example clears all the cached query plans for the `orders` collection:

```
db.runCommand(
  {
    planCacheClear: "orders"
  }
)
```

See also:

- `PlanCache.clearPlansByQuery()` (page 126)
- `PlanCache.clear()` (page 127)

2.2.2 Database Operations

Authentication Commands

Authentication Commands

Name	Description
<code>logout</code> (page 249)	Terminates the current authenticated session.
<code>authenticate</code> (page 249)	Starts an authenticated session using a username and password.
<code>copydbgetnonce</code> (page 250)	This is an internal command to generate a one-time password for use with the <code>copydb</code> (page 300) command.
<code>getnonce</code> (page 250)	This is an internal command to generate a one-time password for authentication.
<code>authSchemaUpgrade</code> (page 250)	Supports the upgrade process for user data between version 2.4 and 2.6.

`logout`

`logout`

The `logout` (page 249) command terminates the current authenticated session:

```
{ logout: 1 }
```

Note: If you're not logged in and using authentication, `logout` (page 249) has no effect.

Changed in version 2.4: Because MongoDB now allows users defined in one database to have privileges on another database, you must call `logout` (page 249) while using the same database context that you authenticated to.

If you authenticated to a database such as `users` or `$external`, you must issue `logout` (page 249) against this database in order to successfully log out.

Example

Use the `use <database-name>` helper in the interactive `mongo` (page 527) shell, or the following `db.getSiblingDB()` (page 113) in the interactive shell or in `mongo` (page 527) shell scripts to change the `db` object:

```
db = db.getSiblingDB('<database-name>')
```

When you have set the database context and `db` object, you can use the `logout` (page 249) to log out of database as in the following operation:

```
db.runCommand( { logout: 1 } )
```

`authenticate`

`authenticate`

Clients use `authenticate` (page 249) to authenticate a connection. When using the shell, use the `db.auth()` (page 99) helper as follows:

```
db.auth( "username", "password" )
```

See

`db.auth()` (page 99) and <http://docs.mongodb.org/manualcore/security> for more information.

copydbgetnonce

copydbgetnonce

Client libraries use `copydbgetnonce` (page 250) to get a one-time password for use with the `copydb` (page 300) command.

Note: This command obtains a write lock on the affected database and will block other operations until it has completed; however, the write lock for this operation is short lived.

getnonce

getnonce

Client libraries use `getnonce` (page 250) to generate a one-time password for authentication.

authSchemaUpgrade New in version 2.6.

authSchemaUpgrade

`authSchemaUpgrade` supports the upgrade from 2.4 to 2.6 process for existing systems that use *authentication* and *authorization*. Between 2.4 and 2.6 the schema for user credential documents changed requiring the `authSchemaUpgrade` process.

See *Upgrade User Authorization Data to 2.6 Format* (page 640) for more information.

User Management Commands

User Management Commands

Name	Description
<code>createUser</code> (page 250)	Creates a new user.
<code>updateUser</code> (page 252)	Updates a user's data.
<code>dropUser</code> (page 253)	Removes a single user.
<code>dropAllUsersFromDatabase</code> (page 254)	Deletes all users associated with a database.
<code>grantRolesToUser</code> (page 255)	Grants a role and its privileges to a user.
<code>revokeRolesFromUser</code> (page 256)	Removes a role from a user.
<code>usersInfo</code> (page 257)	Returns information about the specified users.

createUser

Definition

createUser

Creates a new user on the database where you run the command. The `createUser` (page 250) command returns a *duplicate user* error if the user exists. The `createUser` (page 250) command uses the following syntax:

```
{ createUser: "<name>",
  pwd: "<cleartext password>",
  customData: { <any information> },
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
```

```

    ...
  ],
  writeConcern: { <write concern> }
}

```

`createUser` (page 250) has the following fields:

field string `createUser` The name of the new user.

field string `pwd` The user's password. The `pwd` field is not required if you run `createUser` (page 250) on the `$external` database to create users who have credentials stored externally to MongoDB.

any document `customData` Any arbitrary information.

field array `roles` The roles granted to the user.

field document `writeConcern` The level of `write concern` for the creation operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `createUser` (page 250) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Behavior `createUser` (page 250) sends password to the MongoDB instance in cleartext. To encrypt the password in transit, use SSL.

Users created on the `$external` database should have credentials stored externally to MongoDB, as, for example, with MongoDB Enterprise installations that use Kerberos.

Required Access You must have the `createUser` *action* on a database to create a new user on that database.

You must have the `grantRole` *action* on a role's database to grant the role to another user.

If you have the `userAdmin` or `userAdminAnyDatabase` role, or if you are authenticated using the *localhost exception*, you have those actions.

Example The following `createUser` (page 250) command creates a user `accountAdmin01` on the `products` database. The command gives `accountAdmin01` the `clusterAdmin` and `readAnyDatabase` roles on the `admin` database and the `readWrite` role on the `products` database:

```

db.getSiblingDB("products").runCommand( { createUser: "accountAdmin01",
  pwd: "cleartext password",
  customData: { employeeId: 12345 },
  roles: [
    { role: "clusterAdmin", db: "admin" },
    { role: "readAnyDatabase", db: "admin" },
    "readWrite"
  ],

```

```
    writeConcern: { w: "majority" , wtimeout: 5000 }  
  } )
```

updateUser

Definition

updateUser

Updates the user's profile on the database on which you run the command. An update to a field **completely replaces** the previous field's values, including updates to the user's `roles` array.

Warning: When you update the `roles` array, you completely replace the previous array's values. To add or remove roles without replacing all the user's existing roles, use the [grantRolesToUser](#) (page 255) or [revokeRolesFromUser](#) (page 256) commands.

The `updateUser` (page 252) command uses the following syntax. To update a user, you must specify the `updateUser` field and at least one other field, other than `writeConcern`:

```
{ updateUser: "<username>",  
  pwd: "<cleartext password>",  
  customData: { <any information> },  
  roles: [  
    { role: "<role>", db: "<database>" } | "<role>",  
    ...  
  ],  
  writeConcern: { <write concern> }  
}
```

The command has the following fields:

field string updateUser The name of the user to update.

field string pwd The user's password.

field document customData Any arbitrary information.

field array roles The roles granted to the user. An update to the `roles` array overrides the previous array's values.

field document writeConcern The level of write concern for the update operation. The `writeConcern` document takes the same fields as the [getLastError](#) (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `updateUser` (page 252) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Behavior `updateUser` (page 252) sends the password to the MongoDB instance in cleartext. To encrypt the password in transit, use SSL.

Required Access You must have access that includes the `revokeRole` *action* on all databases in order to update a user's roles array.

You must have the `grantRole` *action* on a role's database to add a role to a user.

To change another user's `pwd` or `customData` field, you must have the `changeAnyPassword` and `changeAnyCustomData` *actions* respectively on that user's database.

To modify your own password or custom data, you must have the `changeOwnPassword` and `changeOwnCustomData` *actions* respectively on the `cluster` resource.

Example Given a user `appClient01` in the `products` database with the following user info:

```
{
  "_id" : "products.appClient01",
  "user" : "appClient01",
  "db" : "products",
  "customData" : { "empID" : "12345", "badge" : "9156" },
  "roles" : [
    { "role" : "readWrite",
      "db" : "products"
    },
    { "role" : "read",
      "db" : "inventory"
    }
  ]
}
```

The following `updateUser` (page 252) command **completely** replaces the user's `customData` and `roles` data:

```
use products
db.runCommand( { updateUser : "appClient01",
                  customData : { employeeId : "0x3039" },
                  roles : [
                      { role : "read", db : "assets" }
                  ]
                } )
```

The user `appClient01` in the `products` database now has the following user information:

```
{
  "_id" : "products.appClient01",
  "user" : "appClient01",
  "db" : "products",
  "customData" : { "employeeId" : "0x3039" },
  "roles" : [
    { "role" : "read",
      "db" : "assets"
    }
  ]
}
```

dropUser

Definition

dropUser

Removes the user from the database on which you run the command. The `dropUser` (page 253) command has the following syntax:

```
{
  dropUser: "<user>",
  writeConcern: { <write concern> }
}
```

The `dropUser` (page 253) command document has the following fields:

field string `dropUser` The name of the user to delete. You must issue the `dropUser` (page 253) command while using the database where the user exists.

field document `writeConcern` The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Required Access You must have the `dropUser` *action* on a database to drop a user from that database.

Example The following sequence of operations in the `mongo` (page 527) shell removes `accountAdmin01` from the `products` database:

```
use products
db.runCommand( { dropUser: "accountAdmin01",
                 writeConcern: { w: "majority", wtimeout: 5000 }
               } )
```

dropAllUsersFromDatabase

Definition

dropAllUsersFromDatabase

Removes all users from the database on which you run the command.

Warning: The `dropAllUsersFromDatabase` (page 254) removes all users from the database.

The `dropAllUsersFromDatabase` (page 254) command has the following syntax:

```
{ dropAllUsersFromDatabase: 1,
  writeConcern: { <write concern> }
}
```

The `dropAllUsersFromDatabase` (page 254) document has the following fields:

field integer `dropAllUsersFromDatabase` Specify 1 to drop all the users from the current database.

field document `writeConcern` The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Required Access You must have the `dropUser` *action* on a database to drop a user from that database.

Example The following sequence of operations in the `mongo` (page 527) shell drops every user from the `products` database:

```
use products
db.runCommand( { dropAllUsersFromDatabase: 1, writeConcern: { w: "majority" } } )
```

The `n` field in the results document shows the number of users removed:

```
{ "n" : 12, "ok" : 1 }
```

grantRolesToUser

Definition

grantRolesToUser

Grants additional roles to a user.

The `grantRolesToUser` (page 255) command uses the following syntax:

```
{ grantRolesToUser: "<user>",
  roles: [ <roles> ],
  writeConcern: { <write concern> }
}
```

The command has the following fields:

field string grantRolesToUser The name of the user to give additional roles.

field array roles An array of additional roles to grant to the user.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `grantRolesToUser` (page 255) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Required Access You must have the `grantRole` *action* on a database to grant a role on that database.

Example Given a user `accountUser01` in the `products` database with the following roles:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  }
]
```

The following `grantRolesToUser` (page 255) operation gives `accountUser01` the `read` role on the `stock` database and the `readWrite` role on the `products` database.

```
use products
db.runCommand( { grantRolesToUser: "accountUser01",
  roles: [
    { role: "read", db: "stock"},
    "readWrite"
  ]
})
```

```
    ],
    writeConcern: { w: "majority" , wtimeout: 2000 }
  } )
```

The user `accountUser01` in the `products` database now has the following roles:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  },
  { "role" : "read",
    "db" : "stock"
  },
  { "role" : "readWrite",
    "db" : "products"
  }
]
```

revokeRolesFromUser

Definition

revokeRolesFromUser

Removes a one or more roles from a user on the database where the roles exist. The `revokeRolesFromUser` (page 256) command uses the following syntax:

```
{ revokeRolesFromUser: "<user>",
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ],
  writeConcern: { <write concern> }
}
```

The command has the following fields:

field string `revokeRolesFromUser` The user to remove roles from.

field array `roles` The roles to remove from the user.

field document `writeConcern` The level of `write concern` for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `revokeRolesFromUser` (page 256) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Required Access You must have the `revokeRole` *action* on a database to revoke a role on that database.

Example The `accountUser01` user in the `products` database has the following roles:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  },
  { "role" : "read",
    "db" : "stock"
  },
  { "role" : "readWrite",
    "db" : "products"
  }
]
```

The following `revokeRolesFromUser` (page 256) command removes the two of the user's roles: the `read` role on the `stock` database and the `readWrite` role on the `products` database, which is also the database on which the command runs:

```
use products
db.runCommand( { revokeRolesFromUser: "accountUser01",
                  roles: [
                      { role: "read", db: "stock" },
                      "readWrite"
                    ],
                  writeConcern: { w: "majority" }
                } )
```

The user `accountUser01` in the `products` database now has only one remaining role:

```
"roles" : [
  { "role" : "assetsReader",
    "db" : "assets"
  }
]
```

usersInfo

Definition

usersInfo

Returns information about one or more users. To match a single user on the database, use the following form:

```
{ usersInfo: { user: <name>, db: <db> },
  showCredentials: <Boolean>,
  showPrivileges: <Boolean>
}
```

The command has the following fields:

field various usersInfo The user(s) about whom to return information. See *Behavior* (page 258) for type and syntax.

field Boolean showCredentials Set the field to true to display the user's password hash. By default, this field is false.

field Boolean showPrivileges Set the field to true to show the user's full set of privileges, including expanded information for the inherited roles. By default, this field is false. If viewing all users, you cannot specify this field.

Required Access Users can always view their own information.

To view another user's information, the user running the command must have privileges that include the `viewUser` action on the other user's database.

Behavior The argument to the `usersInfo` (page 257) command has multiple forms depending on the requested information:

Specify a Single User In the `usersInfo` field, specify a document with the user's name and database:

```
{ usersInfo: { user: <name>, db: <db> } }
```

Alternatively, for a user that exists on the same database where the command runs, you can specify the user by its name only.

```
{ usersInfo: <name> }
```

Specify Multiple Users In the `usersInfo` field, specify an array of documents:

```
{ usersInfo: [ { user: <name>, db: <db> }, { user: <name>, db: <db> }, ... ] }
```

Specify All Users for a Database In the `usersInfo` field, specify 1:

```
{ usersInfo: 1 }
```

Examples

View Specific Users To see information and privileges, but not the credentials, for the user "Kari" defined in "home" database, run the following command:

```
db.runCommand(
  {
    usersInfo: { user: "Kari", db: "home" },
    showPrivileges: true
  }
)
```

To view a user that exists in the *current* database, you can specify the user by name only. For example, if you are in the home database and a user named "Kari" exists in the home database, you can run the following command:

```
db.getSiblingDB("home").runCommand(
  {
    usersInfo: "Kari",
    showPrivileges: true
  }
)
```

View Multiple Users To view info for several users, use an array, with or without the optional fields `showPrivileges` and `showCredentials`. For example:

```
db.runCommand( { usersInfo: [ { user: "Kari", db: "home" }, { user: "Li", db: "myApp" } ],
  showPrivileges: true
} )
```

View All Users for a Database To view all users on the database the command is run, use a command document that resembles the following:

```
db.runCommand( { usersInfo: 1 } )
```

When viewing all users, you can specify the `showCredentials` field but not the `showPrivileges` field.

Role Management Commands

Role Management Commands

Name	Description
<code>createRole</code> (page 259)	Creates a role and specifies its privileges.
<code>updateRole</code> (page 260)	Updates a user-defined role.
<code>dropRole</code> (page 262)	Deletes the user-defined role.
<code>dropAllRolesFromDatabase</code> (page 263)	Deletes all user-defined roles from a database.
<code>grantPrivilegesToRole</code> (page 264)	Assigns privileges to a user-defined role.
<code>revokePrivilegesFromRole</code> (page 265)	Removes the specified privileges from a user-defined role.
<code>grantRolesToRole</code> (page 267)	Specifies roles from which a user-defined role inherits privileges.
<code>revokeRolesFromRole</code> (page 268)	Removes specified inherited roles from a user-defined role.
<code>rolesInfo</code> (page 270)	Returns information for the specified role or roles.
<code>invalidateUserCache</code> (page 272)	Flushes the in-memory cache of user information, including credentials and roles.

`createRole`

Definition

`createRole`

Creates a role and specifies its *privileges*. The role applies to the database on which you run the command. The `createRole` (page 259) command returns a *duplicate role* error if the role already exists in the database.

The `createRole` (page 259) command uses the following syntax:

```
{ createRole: "<new role>",
  privileges: [
    { resource: { <resource> }, actions: [ "<action>", ... ] },
    ...
  ],
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ],
  writeConcern: <write concern document>
}
```

The `createRole` (page 259) command has the following fields:

field string `createRole` The name of the new role.

field array `privileges` The privileges to grant the role. A privilege consists of a resource and permitted actions. You must specify the `privileges` field. Use an empty array to specify *no* privileges. For the syntax of a privilege, see the `privileges` array.

field array roles An array of roles from which this role inherits privileges. You must specify the `roles` field. Use an empty array to specify *no* roles.

field document writeConcern The level of `write concern` to apply to this operation. The `writeConcern` document uses the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `createRole` (page 259) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Behavior A role's privileges apply to the database where the role is created. The role can inherit privileges from other roles in its database. A role created on the `admin` database can include privileges that apply to all databases or to the *cluster* and can inherit privileges from roles in other databases.

Required Access You must have the `createRole` *action* on a database to create a role on that database.

You must have the `grantRole` *action* on the database that a privilege targets in order to grant that privilege to a role. If the privilege targets multiple databases or the `cluster` resource, you must have the `grantRole` action on the `admin` database.

You must have the `grantRole` *action* on a role's database to grant the role to another role.

Example The following `createRole` (page 259) command creates the `myClusterwideAdmin` role on the `admin` database:

```
use admin
db.runCommand({ createRole: "myClusterwideAdmin",
  privileges: [
    { resource: { cluster: true }, actions: [ "addShard" ] },
    { resource: { db: "config", collection: "" }, actions: [ "find", "update", "insert", "remove" ] },
    { resource: { db: "users", collection: "usersCollection" }, actions: [ "update", "insert", "remove" ] },
    { resource: { db: "", collection: "" }, actions: [ "find" ] }
  ],
  roles: [
    { role: "read", db: "admin" }
  ],
  writeConcern: { w: "majority", wtimeout: 5000 }
})
```

updateRole

Definition

updateRole

Updates a *user-defined role*. The `updateRole` (page 260) command must run on the role's database.

An update to a field **completely replaces** the previous field's values. To grant or remove roles or *privileges* without replacing all values, use one or more of the following commands:

- [grantRolesToRole](#) (page 267)
- [grantPrivilegesToRole](#) (page 264)
- [revokeRolesFromRole](#) (page 268)
- [revokePrivilegesFromRole](#) (page 265)

Warning: An update to the `privileges` or `roles` array completely replaces the previous array's values.

The [updateRole](#) (page 260) command uses the following syntax. To update a role, you must provide the `privileges` array, `roles` array, or both:

```
{
  updateRole: "<role>",
  privileges:
    [
      { resource: { <resource> }, actions: [ "<action>", ... ] },
      ...
    ],
  roles:
    [
      { role: "<role>", db: "<database>" } | "<role>",
      ...
    ],
  writeConcern: <write concern document>
}
```

The [updateRole](#) (page 260) command has the following fields:

- field string updateRole** The name of the *user-defined role* to update.
- field array privileges** Required if you do not specify `roles` array. The privileges to grant the role. An update to the `privileges` array overrides the previous array's values. For the syntax for specifying a privilege, see the `privileges` array.
- field array roles** Required if you do not specify `privileges` array. The roles from which this role inherits privileges. An update to the `roles` array overrides the previous array's values.
- field document writeConcern** The level of `write concern` for the update operation. The `writeConcern` document takes the same fields as the [getLastError](#) (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where [updateRole](#) (page 260) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Behavior A role's privileges apply to the database where the role is created. The role can inherit privileges from other roles in its database. A role created on the `admin` database can include privileges that apply to all databases or to the *cluster* and can inherit privileges from roles in other databases.

Required Access You must have the `revokeRole` *action* on all databases in order to update a role.

You must have the `grantRole` *action* on the database of each role in the `roles` array to update the array.

You must have the `grantRole` action on the database of each privilege in the `privileges` array to update the array. If a privilege's resource spans databases, you must have `grantRole` on the `admin` database. A privilege spans databases if the privilege is any of the following:

- a collection in all databases
- all collections and all database
- the `cluster` resource

Example The following is an example of the `updateRole` (page 260) command that updates the `myClusterwideAdmin` role on the `admin` database. While the `privileges` and the `roles` arrays are both optional, at least one of the two is required:

```
use admin
db.runCommand(
  {
    updateRole: "myClusterwideAdmin",
    privileges:
      [
        {
          resource: { db: "", collection: "" },
          actions: [ "find", "update", "insert", "remove" ]
        }
      ],
    roles:
      [
        { role: "dbAdminAnyDatabase", db: "admin" }
      ],
    writeConcern: { w: "majority" }
  }
)
```

To view a role's privileges, use the `rolesInfo` (page 270) command.

dropRole

Definition

dropRole

Deletes a *user-defined* role from the database on which you run the command.

The `dropRole` (page 262) command uses the following syntax:

```
{
  dropRole: "<role>",
  writeConcern: { <write concern> }
}
```

The `dropRole` (page 262) command has the following fields:

field string dropRole The name of the *user-defined* role to remove from the database.

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Required Access You must have the `dropRole` *action* on a database to drop a role from that database.

Example The following operations remove the `readPrices` role from the `products` database:

```
use products
db.runCommand(
  {
    dropRole: "readPrices",
    writeConcern: { w: "majority" }
  }
)
```

dropAllRolesFromDatabase

Definition

dropAllRolesFromDatabase

Deletes all *user-defined* roles on the database where you run the command.

Warning: The `dropAllRolesFromDatabase` (page 263) removes *all user-defined* roles from the database.

The `dropAllRolesFromDatabase` (page 263) command takes the following form:

```
{
  dropAllRolesFromDatabase: 1,
  writeConcern: { <write concern> }
}
```

The command has the following fields:

field integer dropAllRolesFromDatabase Specify 1 to drop all *user-defined* roles from the database where the command is run.

field document writeConcern The level of `write concern` for the removal operation. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Required Access You must have the `dropRole` *action* on a database to drop a role from that database.

Example The following operations drop all *user-defined* roles from the `products` database:

```
use products
db.runCommand(
  {
    dropAllRolesFromDatabase: 1,
    writeConcern: { w: "majority" }
  }
)
```

The `n` field in the results document reports the number of roles dropped:

```
{ "n" : 4, "ok" : 1 }
```

grantPrivilegesToRole

Definition**grantPrivilegesToRole**

Assigns additional *privileges* to a *user-defined* role defined on the database on which the command is run. The `grantPrivilegesToRole` (page 264) command uses the following syntax:

```
{
  grantPrivilegesToRole: "<role>",
  privileges: [
    {
      resource: { <resource> }, actions: [ "<action>", ... ]
    },
    ...
  ],
  writeConcern: { <write concern> }
}
```

The `grantPrivilegesToRole` (page 264) command has the following fields:

- field string grantPrivilegesToRole** The name of the user-defined role to grant privileges to.
- field array privileges** The privileges to add to the role. For the format of a privilege, see `privileges`.
- field document writeConcern** The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Behavior A role's privileges apply to the database where the role is created. A role created on the `admin` database can include privileges that apply to all databases or to the *cluster*.

Required Access You must have the `grantRole` *action* on the database a privilege targets in order to grant the privilege. To grant a privilege on multiple databases or on the `cluster` resource, you must have the `grantRole` action on the `admin` database.

Example The following `grantPrivilegesToRole` (page 264) command grants two additional privileges to the `service` role that exists in the `products` database:

```
use products
db.runCommand(
  {
    grantPrivilegesToRole: "service",
    privileges: [
      {
        resource: { db: "products", collection: "" }, actions: [ "find" ]
      },
      {
        resource: { db: "products", collection: "system.indexes" }, actions: [ "find" ]
      }
    ],
    writeConcern: { w: "majority" , wtimeout: 5000 }
  }
)
```

The first privilege in the `privileges` array allows the user to search on all non-system collections in the `products` database. The privilege does not allow searches on *system collections* (page 600), such as the `system.indexes` (page 601) collection. To grant access to these system collections, explicitly provision access in the `privileges` array. See <http://docs.mongodb.org/manualreference/resource-document>.

The second privilege explicitly allows the `find` action on `system.indexes` (page 601) collections on all databases.

revokePrivilegesFromRole

Definition

revokePrivilegesFromRole

Removes the specified privileges from the *user-defined* role on the database where the command is run. The `revokePrivilegesFromRole` (page 265) command has the following syntax:

```
{
  revokePrivilegesFromRole: "<role>",
  privileges:
    [
      { resource: { <resource> }, actions: [ "<action>", ... ] },
      ...
    ],
  writeConcern: <write concern document>
}
```

The `revokePrivilegesFromRole` (page 265) command has the following fields:

field string revokePrivilegesFromRole The *user-defined* role to revoke privileges from.

field array privileges An array of privileges to remove from the role. See `privileges` for more information on the format of the privileges.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

Behavior To revoke a privilege, the `resource` document pattern must match **exactly** the `resource` field of that privilege. The `actions` field can be a subset or match exactly.

For example, consider the role `accountRole` in the `products` database with the following privilege that specifies the `products` database as the resource:

```
{
  "resource" : {
    "db" : "products",
    "collection" : ""
  },
  "actions" : [
    "find",
    "update"
  ]
}
```

You *cannot* revoke `find` and/or `update` from just *one* collection in the `products` database. The following operations result in no change to the role:

```
use products
db.runCommand(
  {
    revokePrivilegesFromRole: "accountRole",
    privileges:
      [
```

```
        {
          resource : {
            db : "products",
            collection : "gadgets"
          },
          actions : [
            "find",
            "update"
          ]
        }
      ]
    }
  )

db.runCommand(
  {
    revokePrivilegesFromRole: "accountRole",
    privileges:
      [
        {
          resource : {
            db : "products",
            collection : "gadgets"
          },
          actions : [
            "find"
          ]
        }
      ]
  }
)
```

To revoke the "find" and/or the "update" action from the role `accountRole`, you must match the resource document exactly. For example, the following operation revokes just the "find" action from the existing privilege.

```
use products
db.runCommand(
  {
    revokePrivilegesFromRole: "accountRole",
    privileges:
      [
        {
          resource : {
            db : "products",
            collection : ""
          },
          actions : [
            "find"
          ]
        }
      ]
  }
)
```

Required Access You must have the `revokeRole` *action* on the database a privilege targets in order to revoke that privilege. If the privilege targets multiple databases or the `cluster` resource, you must have the `revokeRole` action on the `admin` database.

Example The following operation removes multiple privileges from the `associates` role in the `products` database:

```
use products
db.runCommand(
  {
    revokePrivilegesFromRole: "associate",
    privileges:
      [
        {
          resource: { db: "products", collection: "" },
          actions: [ "createCollection", "createIndex", "find" ]
        },
        {
          resource: { db: "products", collection: "orders" },
          actions: [ "insert" ]
        }
      ],
    writeConcern: { w: "majority" }
  }
)
```

grantRolesToRole

Definition

grantRolesToRole

Grants roles to a *user-defined role*.

The `grantRolesToRole` (page 267) command affects roles on the database where the command runs. `grantRolesToRole` (page 267) has the following syntax:

```
{ grantRolesToRole: "<role>",
  roles: [
    { role: "<role>", db: "<database>" },
    ...
  ],
  writeConcern: { <write concern> }
}
```

The `grantRolesToRole` (page 267) command has the following fields:

field string grantRolesToRole The name of a role to add subsidiary roles.

field array roles An array of roles from which to inherit.

field document writeConcern The level of write concern for the modification. The `writeConcern` document takes the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `grantRolesToRole` (page 267) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:

```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Behavior A role can inherit privileges from other roles in its database. A role created on the `admin` database can inherit privileges from roles in any database.

Required Access You must have the `grantRole` *action* on a database to grant a role on that database.

Example The following `grantRolesToRole` (page 267) command updates the `productsReaderWriter` role in the `products` database to *inherit* the *privileges* of the `productsReader` role in the `products` database:

```
use products
db.runCommand(
  { grantRolesToRole: "productsReaderWriter",
    roles: [
      "productsReader"
    ],
    writeConcern: { w: "majority" , wtimeout: 5000 }
  }
)
```

revokeRolesFromRole

Definition

revokeRolesFromRole

Removes the specified inherited roles from a role. The `revokeRolesFromRole` (page 268) command has the following syntax:

```
{ revokeRolesFromRole: "<role>",
  roles: [
    { role: "<role>", db: "<database>" } | "<role>",
    ...
  ],
  writeConcern: { <write concern> }
}
```

The command has the following fields:

field string `revokeRolesFromRole` The role from which to remove inherited roles.

field array `roles` The inherited roles to remove.

field document `writeConcern` The level of write concern to apply to this operation. The `writeConcern` document uses the same fields as the `getLastError` (page 235) command.

In the `roles` field, you can specify both *built-in roles* and *user-defined role*.

To specify a role that exists in the same database where `revokeRolesFromRole` (page 268) runs, you can either specify the role with the name of the role:

```
"readWrite"
```

Or you can specify the role with a document, as in:


```
{ role: "<role>", db: "<database>" }
```

To specify a role that exists in a different database, specify the role with a document.

Required Access You must have the `revokeRole` *action* on a database to revoke a role on that database.

Example The `purchaseAgents` role in the `emea` database inherits privileges from several other roles, as listed in the `roles` array:

```
{
  "_id" : "emea.purchaseAgents",
  "role" : "purchaseAgents",
  "db" : "emea",
  "privileges" : [],
  "roles" : [
    {
      "role" : "readOrdersCollection",
      "db" : "emea"
    },
    {
      "role" : "readAccountsCollection",
      "db" : "emea"
    },
    {
      "role" : "writeOrdersCollection",
      "db" : "emea"
    }
  ]
}
```

The following `revokeRolesFromRole` (page 268) operation on the `emea` database removes two roles from the `purchaseAgents` role:

```
use emea
db.runCommand( { revokeRolesFromRole: "purchaseAgents",
  roles: [
    "writeOrdersCollection",
    "readOrdersCollection"
  ],
  writeConcern: { w: "majority" , wtimeout: 5000 }
} )
```

The `purchaseAgents` role now contains just one role:

```
{
  "_id" : "emea.purchaseAgents",
  "role" : "purchaseAgents",
  "db" : "emea",
  "privileges" : [],
  "roles" : [
    {
      "role" : "readAccountsCollection",
      "db" : "emea"
    }
  ]
}
```

rolesInfo

Definition

rolesInfo

Returns inheritance and privilege information for specified roles, including both *user-defined roles* and *built-in roles*.

The `rolesInfo` (page 270) command can also retrieve all roles scoped to a database.

The command has the following fields:

field string,document,array,integer rolesInfo The role(s) to return information about. For the syntax for specifying roles, see *Behavior* (page 270).

field Boolean showPrivileges Set the field to `true` to show role privileges, including both privileges inherited from other roles and privileges defined directly. By default, the command returns only the roles from which this role inherits privileges and does not return specific privileges.

field Boolean showBuiltinRoles When the `rolesInfo` field is set to `1`, set `showBuiltinRoles` to `true` to include *built-in roles* in the output. By default this field is set to `false`, and the output for `rolesInfo: 1` displays only *user-defined roles*.

Behavior When specifying roles, use the syntax described here.

To specify a role from the current database, specify the role by its name:

```
rolesInfo: "<rolename>"
```

To specify a role from another database, specify the role by a document that specifies the role and database:

```
rolesInfo: { role: "<rolename>", db: "<database>" }
```

To specify multiple roles, use an array. Specify each role in the array as a document or string. Use a string only if the role exists on the database on which the command runs:

```
rolesInfo:
[
  "<rolename>",
  { role: "<rolename>", db: "<database>" },
  ...
]
```

To specify all roles in the database on which the command runs, specify `rolesInfo: 1`. By default MongoDB displays all the *user-defined roles* in the database. To include *built-in roles* as well, include the parameter-value pair `showBuiltinRoles: true`:

```
rolesInfo: 1, showBuiltinRoles: true
```

Required Access To view a role's information, you must be explicitly granted the role or must have the `viewRole` *action* on the role's database.

Output

`rolesInfo.role`

The name of the role.

`rolesInfo.db`

The database on which the role is defined. Every database has *built-in roles*. A database might also have *user-defined roles*.

rolesInfo.roles

The roles that directly provide privileges to this role and the databases on which the roles are defined.

rolesInfo.indirectRoles

All roles from which this role inherits privileges. This includes the roles in the `rolesInfo.roles` (page 270) array as well as the roles from which the roles in the `rolesInfo.roles` (page 270) array inherit privileges. All privileges apply to the current role. The documents in this field list the roles and the databases on which they are defined.

rolesInfo.privileges

All the privileges granted by this role. By default the output does not include this array. To include it, specify `showPrivileges: true` when running the `rolesInfo` (page 270) command.

The array includes privileges defined directly in the role as well as privileges inherited from other roles.

Each set of privileges in the array is contained in its own document. Each document specifies the *resources* the privilege accesses and the *actions* allowed.

rolesInfo.isBuiltin

A value of `true` indicates the role is a *built-in role*. A value of `false` indicates the role is a *user-defined role*.

Examples

View Information for a Single Role The following command returns the role inheritance information for the role `associate` defined in the `products` database:

```
db.runCommand(
  {
    rolesInfo: { role: "associate", db: "products" }
  }
)
```

The following command returns the role inheritance information for the role `siteManager` on the database on which the command runs:

```
db.runCommand(
  {
    rolesInfo: "siteManager"
  }
)
```

The following command returns *both* the role inheritance and the privileges for the role `associate` defined on the `products` database:

```
db.runCommand(
  {
    rolesInfo:
      { role: "associate", db: "products" },
    showPrivileges: true
  }
)
```

View Information for Several Roles The following command returns information for two roles on two different databases:

```
db.runCommand(
  {
```

```
    rolesInfo:
      [
        { role: "associate", db: "products" },
        { role: "manager", db: "resources" },
      ]
    }
  )
```

The following returns *both* the role inheritance and the privileges:

```
db.runCommand(
  {
    rolesInfo:
      [
        { role: "associate", db: "products" },
        { role: "manager", db: "resources" },
      ],
    showPrivileges: true
  }
)
```

View All User-Defined Roles for a Database The following operation returns all *user-defined roles* on the database on which the command runs and includes privileges:

```
db.runCommand(
  {
    rolesInfo: 1,
    showPrivileges: true
  }
)
```

View All User-Defined and Built-In Roles for a Database The following operation returns all roles on the database on which the command runs, including both built-in and user-defined roles:

```
db.runCommand(
  {
    rolesInfo: 1,
    showBuiltinRoles: true
  }
)
```

invalidateUserCache

Definition

invalidateUserCache

New in version 2.6.

Flushes user information from in-memory cache, including removal of each user's credentials and roles. This allows you to purge the cache at any given moment, regardless of the interval set in the `userCacheInvalidationIntervalSecs` parameter.

`invalidateUserCache` (page 272) has the following syntax:

```
db.runCommand( { invalidateUserCache: 1 } )
```

Required Access You must have privileges that include the `invalidateUserCache` action on the cluster resource in order to use this command.

Replication Commands

Replication Commands

Name	Description
<code>replSetFreeze</code> (page 273)	Prevents the current member from seeking election as <i>primary</i> for a period of time.
<code>replSetGetStatus</code> (page 273)	Returns a document that reports on the status of the replica set.
<code>replSetInitiate</code> (page 275)	Initializes a new replica set.
<code>replSetMaintenance</code> (page 276)	Enables or disables a maintenance mode, which puts a <i>secondary</i> node in a RECOVERING state.
<code>replSetReconfig</code> (page 276)	Applies a new configuration to an existing replica set.
<code>replSetStepDown</code> (page 277)	Forces the current <i>primary</i> to <i>step down</i> and become a <i>secondary</i> , forcing an election.
<code>replSetSyncFrom</code> (page 278)	Explicitly override the default logic for selecting a member to replicate from.
<code>resync</code> (page 279)	Forces a <code>mongod</code> (page 503) to re-synchronize from the <i>master</i> . For master-slave replication only.
<code>applyOps</code> (page 279)	Internal command that applies <i>oplog</i> entries to the current data set.
<code>isMaster</code> (page 280)	Displays information about this member's role in the replica set, including whether it is the master.
<code>getoptime</code> (page 282)	Internal command to support replication, returns the optime.

`replSetFreeze`

`replSetFreeze`

The `replSetFreeze` (page 273) command prevents a replica set member from seeking election for the specified number of seconds. Use this command in conjunction with the `replSetStepDown` (page 277) command to make a different node in the replica set a primary.

The `replSetFreeze` (page 273) command uses the following syntax:

```
{ replSetFreeze: <seconds> }
```

If you want to unfreeze a replica set member before the specified number of seconds has elapsed, you can issue the command with a seconds value of 0:

```
{ replSetFreeze: 0 }
```

Restarting the `mongod` (page 503) process also unfreezes a replica set member.

`replSetFreeze` (page 273) is an administrative command, and you must issue it against the *admin database*.

`replSetGetStatus`

Definition

replSetGetStatus

The `replSetGetStatus` command returns the status of the replica set from the point of view of the current server. You must run the command against the *admin database*. The command has the following prototype format:

```
{ replSetGetStatus: 1 }
```

The value specified does not affect the output of the command. Data provided by this command derives from data included in heartbeats sent to the current instance by other members of the replica set. Because of the frequency of heartbeats, these data can be several seconds out of date.

You can also access this functionality through the `rs.status()` (page 168) helper in the *mongo* (page 527) shell.

The *mongod* (page 503) must have replication enabled and be a member of a replica set for the `replSetGetStatus` (page 273) to return successfully.

Output**replSetGetStatus.set**

The `set` value is the name of the replica set, configured in the `replSetName` setting. This is the same value as `_id` in `rs.conf()` (page 165).

replSetGetStatus.date

The value of the `date` field is an *ISODate* of the current time, according to the current server. Compare this to the value of the `lastHeartbeat` (page 275) to find the operational lag between the current host and the other hosts in the set.

replSetGetStatus.myState

The value of `myState` (page 274) is an integer between 0 and 10 that represents the `replica state` of the current member.

replSetGetStatus.members

The `members` field holds an array that contains a document for every member in the replica set.

replSetGetStatus.members.name

The `name` field holds the name of the server.

replSetGetStatus.members.self

The `self` field is only included in the document for the current *mongod* instance in the `members` array. Its value is `true`.

replSetGetStatus.members.health

The `health` value is only present for the other members of the replica set (i.e. not the member that returns `rs.status` (page 168).) This field conveys if the member is up (i.e. 1) or down (i.e. 0.)

replSetGetStatus.members.state

The value of `state` (page 274) is an integer between 0 and 10 that represents the `replica state` of the member.

replSetGetStatus.members.stateStr

A string that describes `state` (page 274).

replSetGetStatus.members.uptime

The `uptime` (page 274) field holds a value that reflects the number of seconds that this member has been online.

This value does not appear for the member that returns the `rs.status()` (page 168) data.

replSetGetStatus.members.optime

A document that contains information regarding the last operation from the operation log that this member has applied.

`replSetGetStatus.members.optime.t`

A 32-bit timestamp of the last operation applied to this member of the replica set from the *oplog*.

`replSetGetStatus.members.optime.i`

An incremented field, which reflects the number of operations in since the last time stamp. This value only increases if there is more than one operation per second.

`replSetGetStatus.members.optimeDate`

An *ISODate* formatted date string that reflects the last entry from the *oplog* that this member applied. If this differs significantly from `lastHeartbeat` (page 275) this member is either experiencing “replication lag” or there have not been any new operations since the last update. Compare `members.optimeDate` between all of the members of the set.

`replSetGetStatus.members.lastHeartbeat`

The `lastHeartbeat` value provides an *ISODate* formatted date and time of the transmission time of last heartbeat received from this member. Compare this value to the value of the `date` (page 274) and `lastHeartBeatRecv` field to track latency between these members.

This value does not appear for the member that returns the `rs.status()` (page 168) data.

`replSetGetStatus.members.lastHeartbeatRecv`

The `lastHeartbeatRecv` value provides an *ISODate* formatted date and time that the last heartbeat was received from this member. Compare this value to the value of the `date` (page 274) and `lastHeartBeat` field to track latency between these members.

`replSetGetStatus.members.lastHeartbeatMessage`

When the last heartbeat included an extra message, the `lastHeartbeatMessage` (page 275) contains a string representation of that message.

`replSetGetStatus.members.pingMs`

The `pingMs` represents the number of milliseconds (ms) that a round-trip packet takes to travel between the remote member and the local instance.

This value does not appear for the member that returns the `rs.status()` (page 168) data.

`replSetGetStatus.syncingTo`

The `syncingTo` field is only present on the output of `rs.status()` (page 168) on *secondary* and recovering members, and holds the hostname of the member from which this instance is syncing.

replSetInitiate

replSetInitiate

The `replSetInitiate` (page 275) command initializes a new replica set. Use the following syntax:

```
{ replSetInitiate : <config_document> }
```

The `<config_document>` is a *document* that specifies the replica set’s configuration. For instance, here’s a config document for creating a simple 3-member replica set:

```
{
  _id : <setname>,
  members : [
    { _id : 0, host : <host0> },
    { _id : 1, host : <host1> },
    { _id : 2, host : <host2> },
  ]
}
```

A typical way of running this command is to assign the config document to a variable and then to pass the document to the `rs.initiate()` (page 166) helper:

```
config = {
  _id : "my_replica_set",
  members : [
    { _id : 0, host : "rs1.example.net:27017"},
    { _id : 1, host : "rs2.example.net:27017"},
    { _id : 2, host : "rs3.example.net", arbiterOnly: true },
  ]
}

rs.initiate(config)
```

Notice that omitting the port cause the host to use the default port of 27017. Notice also that you can specify other options in the config documents such as the `arbiterOnly` setting in this example.

See also:

<http://docs.mongodb.org/manualreference/replica-configuration>,
<http://docs.mongodb.org/manualadministration/replica-sets>, and *Replica Set Re-configuration*.

replSetMaintenance

Definition

replSetMaintenance

The `replSetMaintenance` (page 276) admin command enables or disables the maintenance mode for a *secondary* member of a *replica set*.

The command has the following prototype form:

```
{ replSetMaintenance: <boolean> }
```

Behavior Consider the following behavior when running the `replSetMaintenance` (page 276) command:

- You cannot run the command on the Primary.
- You must run the command against the `admin` database.
- When enabled `replSetMaintenance: true`, the member enters the `RECOVERING` state. While the secondary is `RECOVERING`:
 - The member is not accessible for read operations.
 - The member continues to sync its *oplog* from the Primary.
- On secondaries, the `compact` (page 313) command forces the secondary to enter `RECOVERING` state. Read operations issued to an instance in the `RECOVERING` state will fail. This prevents clients from reading during the operation. When the operation completes, the secondary returns to: `replstate:SECONDARY` state.
- See <http://docs.mongodb.org/manualreference/replica-states/> for more information about replica set member states.

See <http://docs.mongodb.org/manualtutorial/perform-maintenance-on-replica-set-members> for an example replica set maintenance procedure to maximize availability during maintenance operations.

replSetReconfig

replSetReconfig

The `replSetReconfig` (page 276) command modifies the configuration of an existing replica set. You can use this command to add and remove members, and to alter the options set on existing members. Use the following syntax:

```
{ replSetReconfig: <new_config_document>, force: false }
```

You may also run `replSetReconfig` (page 276) with the shell's `rs.reconfig()` (page 167) method.

Behaviors Be aware of the following `replSetReconfig` (page 276) behaviors:

- You must issue this command against the *admin database* of the current primary member of the replica set.
- You can optionally force the replica set to accept the new configuration by specifying `force: true`. Use this option if the current member is not primary or if a majority of the members of the set are not accessible.

Warning: Forcing the `replSetReconfig` (page 276) command can lead to a *rollback* situation. Use with caution.

Use the force option to restore a replica set to new servers with different hostnames. This works even if the set members already have a copy of the data.

- A majority of the set's members must be operational for the changes to propagate properly.
- This command can cause downtime as the set renegotiates primary-status. Typically this is 10-20 seconds, but could be as long as a minute or more. Therefore, you should attempt to reconfigure only during scheduled maintenance periods.
- In some cases, `replSetReconfig` (page 276) forces the current primary to step down, initiating an election for primary among the members of the replica set. When this happens, the set will drop all current connections.

`replSetReconfig` (page 276) obtains a special mutually exclusive lock to prevent more than one `replSetReconfig` (page 276) operation from occurring at the same time.

Additional Information <http://docs.mongodb.org/manualreference/replica-configuration>, `rs.reconfig()` (page 167), and `rs.conf()` (page 165).

replSetStepDown**Description****replSetStepDown**

Forces the *primary* of the replica set to become a *secondary*. This initiates an *election for primary*.

`replSetStepDown` (page 277) has the following prototype form:

```
db.runCommand( { replSetStepDown: <seconds> , force: <true|false> } )
```

`replSetStepDown` (page 277) has the following fields:

field number replSetStepDown A number of seconds for the member to avoid election to primary.

If you do not specify a value for `<seconds>`, `replSetStepDown` (page 277) will attempt to avoid reelection to primary for 60 seconds.

field Boolean force New in version 2.0: Forces the *primary* to step down even if there are no secondary members within 10 seconds of the primary's latest optime.

Warning: `replSetStepDown` (page 277) forces all clients currently connected to the database to disconnect. This helps ensure that clients maintain an accurate view of the replica set.

New in version 2.0: If there is no *secondary* within 10 seconds of the primary, `replSetStepDown` (page 277) will not succeed to prevent long running elections.

Example The following example specifies that the former primary avoids reelection to primary for 120 seconds:

```
db.runCommand( { replSetStepDown: 120 } )
```

replSetSyncFrom

Description

replSetSyncFrom

New in version 2.2.

Explicitly configures which host the current `mongod` (page 503) pulls *oplog* entries from. This operation is useful for testing different patterns and in situations where a set member is not replicating from the desired host.

The `replSetSyncFrom` (page 278) command has the following form:

```
{ replSetSyncFrom: "hostname[:port]" }
```

The `replSetSyncFrom` (page 278) command has the following field:

field string replSetSyncFrom The name and port number of the replica set member that this member should replicate from. Use the `[hostname] : [port]` form.

The Target Member

The member to replicate from must be a valid source for data in the set. The member cannot be:

- The same as the `mongod` (page 503) on which you run `replSetSyncFrom` (page 278). In other words, a member cannot replicate from itself.
- An arbiter, because arbiters do not hold data.
- A member that does not build indexes.
- An unreachable member.
- A `mongod` (page 503) instance that is not a member of the same replica set.

If you attempt to replicate from a member that is more than 10 seconds behind the current member, `mongod` (page 503) will log a warning but will still replicate from the lagging member.

If you run `replSetSyncFrom` (page 278) during initial sync, MongoDB produces no error messages, but the sync target will not change until after the initial sync operation.

Run from the mongo Shell

To run the command in the `mongo` (page 527) shell, use the following invocation:

```
db.adminCommand( { replSetSyncFrom: "hostname[:port]" } )
```

You may also use the `rs.syncFrom()` (page 169) helper in the `mongo` (page 527) shell in an operation with the following form:

```
rs.syncFrom("hostname<:port>")
```

Note: `replSetSyncFrom` (page 278) and `rs.syncFrom()` (page 169) provide a temporary override of default behavior. `mongod` (page 503) will revert to the default sync behavior in the following situations:

- The `mongod` (page 503) instance restarts.
- The connection between the `mongod` (page 503) and the sync target closes.

Changed in version 2.4: The sync target falls more than 30 seconds behind another member of the replica set; the `mongod` (page 503) will revert to the default sync target.

resync

resync

The `resync` (page 279) command forces an out-of-date slave `mongod` (page 503) instance to re-synchronize itself. Note that this command is relevant to master-slave replication only. It does not apply to replica sets.

Warning: This command obtains a global write lock and will block other operations until it has completed.

applyOps

Definition

applyOps

Applies specified *oplog* entries to a `mongod` (page 503) instance. The `applyOps` (page 279) command is primarily an internal command to support *sharded clusters*.

If authentication is enabled, you must have access to all actions on all resources in order to run `applyOps` (page 279). Providing such access is not recommended, but if your organization requires a user to run `applyOps` (page 279), create a role that grants `anyAction` on `resource-anyresource`. Do not assign this role to any other user.

The `applyOps` (page 279) command has the following prototype form:

```
db.runCommand( { applyOps: [ <operations> ],
                  preCondition: [ { ns: <namespace>, q: <query>, res: <result> } ] } )
```

The `applyOps` (page 279) command takes a document with the following fields:

field array applyOps The oplog entries to apply.

field array preCondition An array of documents that contain the conditions that must be true in order to apply the oplog entry. Each document contains a set of conditions, as described in the next table.

The `preCondition` array takes one or more documents with the following fields:

field string ns A *namespace*. If you use this field, `applyOps` (page 279) applies oplog entries only for the *collection* described by this namespace.

param string q Specifies the *query* that produces the results specified in the `res` field.

param string res The results of the query in the `q` field that must match to apply the oplog entry.

Warning: This command obtains a global write lock and will block other operations until it has completed.

isMaster

Definition

isMaster

`isMaster` (page 280) returns a document that describes the role of the `mongod` (page 503) instance.

If the instance is a member of a replica set, then `isMaster` (page 280) returns a subset of the replica set configuration and status including whether or not the instance is the *primary* of the replica set.

When sent to a `mongod` (page 503) instance that is not a member of a replica set, `isMaster` (page 280) returns a subset of this information.

MongoDB *drivers* and *clients* use `isMaster` (page 280) to determine the state of the replica set members and to discover additional members of a *replica set*.

The `db.isMaster()` (page 114) method in the `mongo` (page 527) shell provides a wrapper around `isMaster` (page 280).

The command takes the following form:

```
{ isMaster: 1 }
```

See also:

`db.isMaster()` (page 114)

Output

All Instances The following `isMaster` (page 280) fields are common across all roles:

`isMaster.ismaster`

A boolean value that reports when this node is writable. If `true`, then this instance is a *primary* in a *replica set*, or a *master* in a master-slave configuration, or a `mongos` (page 518) instance, or a standalone `mongod` (page 503).

This field will be `false` if the instance is a *secondary* member of a replica set or if the member is an *arbiter* of a replica set.

`isMaster.maxBsonObjectSize`

The maximum permitted size of a *BSON* object in bytes for this `mongod` (page 503) process. If not provided, clients should assume a max size of “16 * 1024 * 1024”.

`isMaster.maxMessageSizeBytes`

New in version 2.4.

The maximum permitted size of a *BSON* wire protocol message. The default value is 48000000 bytes.

`isMaster.localTime`

New in version 2.2.

Returns the local server time in UTC. This value is an *ISO date*.

`isMaster.minWireVersion`

New in version 2.6.

The earliest version of the wire protocol that this `mongod` (page 503) or `mongos` (page 518) instance is capable of using to communicate with clients.

Clients may use `minWireVersion` (page 280) to help negotiate compatibility with MongoDB.

`isMaster.maxWireVersion`

New in version 2.6.

The latest version of the wire protocol that this `mongod` (page 503) or `mongos` (page 518) instance is capable of using to communicate with clients.

Clients may use `maxWireVersion` (page 281) to help negotiate compatibility with MongoDB.

Sharded Instances `mongos` (page 518) instances add the following field to the `isMaster` (page 280) response document:

`isMaster.msg`

Contains the value `isdbgrid` when `isMaster` (page 280) returns from a `mongos` (page 518) instance.

Replica Sets `isMaster` (page 280) contains these fields when returned by a member of a replica set:

`isMaster.setName`

The name of the current `:replica set`.

`isMaster.secondary`

A boolean value that, when `true`, indicates if the `mongod` (page 503) is a *secondary* member of a *replica set*.

`isMaster.hosts`

An array of strings in the format of "[hostname] : [port]" that lists all members of the *replica set* that are neither *hidden*, *passive*, nor *arbiters*.

Drivers use this array and the `isMaster.passives` (page 281) to determine which members to read from.

`isMaster.passives`

An array of strings in the format of "[hostname] : [port]" listing all members of the *replica set* which have a priority of 0.

This field only appears if there is at least one member with a priority of 0.

Drivers use this array and the `isMaster.hosts` (page 281) to determine which members to read from.

`isMaster.arbiters`

An array of strings in the format of "[hostname] : [port]" listing all members of the *replica set* that are *arbiters*.

This field only appears if there is at least one arbiter in the replica set.

`isMaster.primary`

A string in the format of "[hostname] : [port]" listing the current *primary* member of the replica set.

`isMaster.arbiterOnly`

A boolean value that, when `true`, indicates that the current instance is an *arbiter*. The `arbiterOnly` (page 281) field is only present, if the instance is an arbiter.

`isMaster.passive`

A boolean value that, when `true`, indicates that the current instance is *hidden*. The `passive` (page 281) field is only present for hidden members.

`isMaster.hidden`

A boolean value that, when `true`, indicates that the current instance is *hidden*. The `hidden` (page 281) field is only present for hidden members.

`isMaster.tags`

A document that lists any tags assigned to this member. This field is only present if there are tags assigned to the member. See

<http://docs.mongodb.org/manual/tutorial/configure-replica-set-tag-sets> for more information.

`isMaster.me`

The `[hostname] : [port]` of the member that returned `isMaster` (page 280).

getoptime

getoptime

`getoptime` (page 282) is an internal command.

See also:

<http://docs.mongodb.org/manual/replication> for more information regarding replication.

Sharding Commands

Sharding Commands

Name	Description
<code>flushRouterConfig</code> (page 283)	Forces an update to the cluster metadata cached by a <code>mongos</code> (page 518).
<code>addShard</code> (page 284)	Adds a <i>shard</i> to a <i>sharded cluster</i> .
<code>cleanupOrphaned</code> (page 285)	Removes orphaned data with shard key values outside of the ranges of the chunks owned by a shard.
<code>checkShardingIndex</code> (page 288)	Internal command that validates index on shard key.
<code>enableSharding</code> (page 288)	Enables sharding on a specific database.
<code>listShards</code> (page 288)	Returns a list of configured shards.
<code>removeShard</code> (page 288)	Starts the process of removing a shard from a sharded cluster.
<code>getShardMap</code> (page 289)	Internal command that reports on the state of a sharded cluster.
<code>getShardVersion</code> (page 289)	Internal command that returns the <i>config server</i> version.
<code>mergeChunks</code> (page 289)	Provides the ability to combine chunks on a single shard.
<code>setShardVersion</code> (page 291)	Internal command to sets the <i>config server</i> version.
<code>shardCollection</code> (page 291)	Enables the sharding functionality for a collection, allowing the collection to be sharded.
<code>shardingState</code> (page 292)	Reports whether the <code>mongod</code> (page 503) is a member of a sharded cluster.
<code>unsetSharding</code> (page 293)	Internal command that affects connections between instances in a MongoDB deployment.
<code>split</code> (page 293)	Creates a new <i>chunk</i> .
<code>splitChunk</code> (page 295)	Internal command to split chunk. Instead use the methods <code>sh.splitFind()</code> (page 179) and <code>sh.splitAt()</code> (page 179).
<code>splitVector</code> (page 296)	Internal command that determines split points.
<code>medianKey</code> (page 296)	Deprecated internal command. See <code>splitVector</code> (page 296).
<code>moveChunk</code> (page 296)	Internal command that migrates chunks between shards.
<code>movePrimary</code> (page 297)	Reassigns the <i>primary shard</i> when removing a shard from a sharded cluster.
<code>isdbgrid</code> (page 298)	Verifies that a process is a <code>mongos</code> (page 518).

flushRouterConfig

flushRouterConfig

`flushRouterConfig` (page 283) clears the current cluster information cached by a `mongos` (page 518) instance and reloads all *sharded cluster* metadata from the *config database*.

This forces an update when the configuration database holds data that is newer than the data cached in the `mongos` (page 518) process.

Warning: Do not modify the config data, except as explicitly documented. A config database cannot typically tolerate manual manipulation.

`flushRouterConfig` (page 283) is an administrative command that is only available for `mongos` (page 518) instances.

New in version 1.8.2.

addShard

Definition

addShard

Adds either a database instance or a *replica set* to a *sharded cluster*. The optimal configuration is to deploy shards across replica sets.

Run `addShard` (page 284) when connected to a `mongos` (page 518) instance. The command takes the following form when adding a single database instance as a shard:

```
{ addShard: "<hostname><:port>", maxSize: <size>, name: "<shard_name>" }
```

When adding a replica set as a shard, use the following form:

```
{ addShard: "<replica_set>/<hostname><:port>", maxSize: <size>, name: "<shard_name>" }
```

The command contains the following fields:

field string addShard The hostname and port of the `mongod` (page 503) instance to be added as a shard. To add a replica set as a shard, specify the name of the replica set and the hostname and port of a member of the replica set.

field integer maxSize The maximum size in megabytes of the shard. If you set `maxSize` to 0, MongoDB does not limit the size of the shard.

field string name A name for the shard. If this is not specified, MongoDB automatically provides a unique name.

The `addShard` (page 284) command stores shard configuration information in the *config database*. Always run `addShard` (page 284) when using the `admin` database.

Specify a `maxSize` when you have machines with different disk capacities, or if you want to limit the amount of data on some shards. The `maxSize` constraint prevents the *balancer* from migrating chunks to the shard when the value of `mem.mapped` (page 351) exceeds the value of `maxSize`.

Important: You cannot include a `hidden` member in the seed list provided to `addShard` (page 284).

Examples The following command adds the database instance running on port 27027 on the host `mongodb0.example.net` as a shard:

```
use admin
db.runCommand({addShard: "mongodb0.example.net:27027"})
```

Warning: Do not use `localhost` for the hostname unless your *configuration server* is also running on `localhost`.

The following command adds a replica set as a shard:

```
use admin
db.runCommand( { addShard: "repl0/mongodb3.example.net:27327" } )
```


You may specify all members in the replica set. All additional hostnames must be members of the same replica set.

cleanupOrphaned

Definition

cleanupOrphaned

New in version 2.6.

Deletes from a shard the *orphaned documents* whose shard key values fall into a single or a single contiguous range that do not belong to the shard. For example, if two contiguous ranges do not belong to the shard, the `cleanupOrphaned` (page 285) examines both ranges for orphaned documents.

`cleanupOrphaned` (page 285) has the following syntax:

```
db.runCommand( {
  cleanupOrphaned: "<database>.<collection>",
  startingAtKey: <minimumShardKeyValue>,
  secondaryThrottle: <boolean>
} )
```

`cleanupOrphaned` (page 285) has the following fields:

field string cleanupOrphaned The namespace, i.e. both the database and the collection name, of the sharded collection for which to clean the orphaned data.

field document startingAtKey The *shard key* value that determines the lower bound of the cleanup range. The default value is `MinKey`.

If the range that contains the specified `startingAtKey` value belongs to a chunk owned by the shard, `cleanupOrphaned` (page 285) continues to examine the next ranges until it finds a range not owned by the shard. See *Determine Range* (page 285) for details.

field boolean secondaryThrottle If `true`, each delete operation must be replicated to another secondary before the cleanup operation proceeds further. If `false`, do not wait for replication. Defaults to `false`.

Independent of the `secondaryThrottle` setting, after the final delete, `cleanupOrphaned` (page 285) waits for all deletes to replicate to a majority of replica set members before returning.

Behavior Run `cleanupOrphaned` (page 285) in the `admin` database directly on the `mongod` (page 503) instance that is the primary replica set member of the shard. Do not run `cleanupOrphaned` (page 285) on a `mongos` (page 518) instance.

You do not need to disable the balancer before running `cleanupOrphaned` (page 285).

Determine Range The `cleanupOrphaned` (page 285) command uses the `startingAtKey` value, if specified, to determine the start of the range to examine for orphaned document:

- If the `startingAtKey` value falls into a range for a chunk not owned by the shard, `cleanupOrphaned` (page 285) begins examining at the start of this range, which may not necessarily be the `startingAtKey`.
- If the `startingAtKey` value falls into a range for a chunk owned by the shard, `cleanupOrphaned` (page 285) moves onto the next range until it finds a range for a chunk not owned by the shard.

The `cleanupOrphaned` (page 285) deletes orphaned documents from the start of the determined range and ends at the start of the chunk range that belongs to the shard.

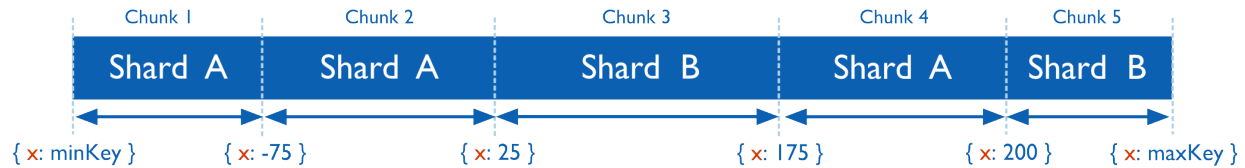


Figure 2.1: Diagram of shard key value space, showing chunk ranges and shards.

Consider the following key space with documents distributed across Shard A and Shard B.

Shard A owns:

- Chunk 1 with the range { x: minKey } --> { x: -75 },
- Chunk 2 with the range { x: -75 } --> { x: 25 }, and
- Chunk 4 with the range { x: 175 } --> { x: 200 }.

Shard B owns:

- Chunk 3 with the range { x: 25 } --> { x: 175 } and
- Chunk 5 with the range { x: 200 } --> { x: maxKey }.

If on Shard A, the `cleanupOrphaned` (page 285) command runs with `startingAtKey: { x: -70 }` or any other value belonging to range for Chunk 1 or Chunk 2, the `cleanupOrphaned` (page 285) command examines the Chunk 3 range of { x: 25 } --> { x: 175 } to delete orphaned data.

If on Shard B, the `cleanupOrphaned` (page 285) command runs with the `startingAtKey: { x: -70 }` or any other value belonging to range for Chunk 1, the `cleanupOrphaned` (page 285) command examines the combined contiguous range for Chunk 1 and Chunk 2, namely { x: minKey } --> { x: 25 } to delete orphaned data.

Required Access On systems running with authorization, you must have `clusterAdmin` privileges to run `cleanupOrphaned` (page 285).

Output

Return Document Each `cleanupOrphaned` (page 285) command returns a document containing a subset of the following fields:

`cleanupOrphaned.ok`

Equal to 1 on success.

A value of 1 indicates that `cleanupOrphaned` (page 285) scanned the specified shard key range, deleted any orphaned documents found in that range, and confirmed that all deletes replicated to a majority of the members of that shard's replica set. If confirmation does not arrive within 1 hour, `cleanupOrphaned` (page 285) times out.

A value of 0 could indicate either of two cases:

- `cleanupOrphaned` (page 285) found orphaned documents on the shard but could not delete them.
- `cleanupOrphaned` (page 285) found and deleted orphaned documents, but could not confirm replication before the 1 hour timeout. In this case, replication does occur, but only after `cleanupOrphaned` (page 285) returns.

cleanupOrphaned.stoppedAtKey

The upper bound of the cleanup range of shard keys. If present, the value corresponds to the lower bound of the next chunk on the shard. The absence of the field signifies that the cleanup range was the uppermost range for the shard.

Log Files The `cleanupOrphaned` (page 285) command prints the number of deleted documents to the `mongod` (page 503) log. For example:

```
m30000| 2013-10-31T15:17:28.972-0400 [conn1] Deleter starting delete for: foo.bar from { _id: -35.0 }
m30000| 2013-10-31T15:17:28.972-0400 [conn1] rangeDeleter deleted 0 documents for foo.bar from { _id:
```

Examples The following examples run the `cleanupOrphaned` (page 285) command directly on the primary of the shard.

Remove Orphaned Documents for a Specific Range For a sharded collection `info` in the `test` database, a shard owns a single chunk with the range: `{ x: MinKey } --> { x: 10 }`.

The shard also contains documents whose shard keys values fall in a range for a chunk *not* owned by the shard: `{ x: 10 } --> { x: MaxKey }`.

To remove orphaned documents within the `{ x: 10 } => { x: MaxKey }` range, you can specify a `startingAtKey` with a value that falls into this range, as in the following example:

```
use admin
db.runCommand( {
  "cleanupOrphaned": "test.info",
  "startingAtKey": { x: 10 },
  "secondaryThrottle": true
} )
```

Or you can specify a `startingAtKey` with a value that falls into the previous range, as in the following:

```
use admin
db.runCommand( {
  "cleanupOrphaned": "test.info",
  "startingAtKey": { x: 2 },
  "secondaryThrottle": true
} )
```

Since `{ x: 2 }` falls into a range that belongs to a chunk owned by the shard, `cleanupOrphaned` (page 285) examines the next range to find a range not owned by the shard, in this case `{ x: 10 } => { x: MaxKey }`.

Remove All Orphaned Documents from a Shard `cleanupOrphaned` (page 285) examines documents from a single contiguous range of shard keys. To remove all orphaned documents from the shard, you can run `cleanupOrphaned` (page 285) in a loop, using the returned `stoppedAtKey` as the next `startingFromKey`, as in the following:

```
use admin
var nextKey = { };

while ( nextKey = db.runCommand( {
  cleanupOrphaned: "test.user",
  startingFromKey: nextKey
} ).stoppedAtKey ) {
```

```
printjson(nextKey);  
}
```

checkShardingIndex

checkShardingIndex

`checkShardingIndex` (page 288) is an internal command that supports the sharding functionality.

enableSharding

enableSharding

The `enableSharding` (page 288) command enables sharding on a per-database level. Use the following command form:

```
{ enableSharding: "<database name>" }
```

Once you've enabled sharding in a database, you can use the `shardCollection` (page 291) command to begin the process of distributing data among the shards.

listShards

listShards

Use the `listShards` (page 288) command to return a list of configured shards. The command takes the following form:

```
{ listShards: 1 }
```

removeShard

removeShard

Removes a shard from a *sharded cluster*. When you run `removeShard` (page 288), MongoDB first moves the shard's chunks to other shards in the cluster. Then MongoDB removes the shard.

Behavior

Access Requirements You *must* run `removeShard` (page 288) while connected to a `mongos` (page 518). Issue the command against the `admin` database or use the `sh._adminCommand()` (page 172) helper.

If you have authorization enabled, you must have the `clusterManager` role or any role that includes the `removeShard` action.

Database Migration Requirements Each database in a sharded cluster has a primary shard. If the shard you want to remove is also the primary of one of the cluster's databases, then you must manually move the databases to a new shard after migrating all data from the shard. See the `movePrimary` (page 297) command and the <http://docs.mongodb.org/manual/tutorial/remove-shards-from-cluster> for more information.

Example From the `mongo` (page 527) shell, the `removeShard` (page 288) operation resembles the following:

```
use admin  
db.runCommand( { removeShard : "bristol01" } )
```

Replace `bristol01` with the name of the shard to remove. When you run `removeShard` (page 288), the command returns immediately, with the following message:

```
{
  "msg" : "draining started successfully",
  "state" : "started",
  "shard" : "bristol01",
  "ok" : 1
}
```

The balancer begins migrating chunks from the shard named `bristol01` to other shards in the cluster. These migrations happens slowly to avoid placing undue load on the overall cluster.

If you run the command again, `removeShard` (page 288) returns the following progress output:

```
{
  "msg" : "draining ongoing",
  "state" : "ongoing",
  "remaining" : {
    "chunks" : 23,
    "dbs" : 1
  },
  "ok" : 1
}
```

The remaining *document* specifies how many chunks and databases remain on the shard. Use `db.printShardingStatus()` (page 116) to list the databases that you must move from the shard. Use the `movePrimary` (page 297) to move databases.

After removing all chunks and databases from the shard, you can issue `removeShard` (page 288) again see the following:

```
{
  "msg" : "removeshard completed successfully",
  "state" : "completed",
  "shard" : "bristol01",
  "ok" : 1
}
```

getShardMap

getShardMap

`getShardMap` (page 289) is an internal command that supports the sharding functionality.

getShardVersion

getShardVersion

`getShardVersion` (page 289) is a command that supports sharding functionality and is not part of the stable client facing API.

mergeChunks

Definition

mergeChunks

For a sharded collection, `mergeChunks` (page 289) combines two contiguous *chunk* ranges the same shard into a single chunk. At least one of chunk must not have any documents. Issue the `mergeChunks` (page 289) command from a `mongos` (page 518) instance.

`mergeChunks` (page 289) has the following form:

```
db.runCommand( { mergeChunks : <namespace> ,
                  bounds : [ { <shardKeyField>: <minFieldValue> },
                             { <shardKeyField>: <maxFieldValue> } ] } )
```

For compound shard keys, you must include the full shard key in the `bounds` specification. If the shard key is `{ x: 1, y: 1 }`, `mergeChunks` (page 289) has the following form:

```
db.runCommand( { mergeChunks : <namespace> ,
                  bounds : [ { x: <minValue>, y: <minValue> },
                             { x: <maxValue>, y: <maxValue> } ] } )
```

The `mergeChunks` (page 289) command has the following fields:

field namespace mergeChunks The fully qualified *namespace* of the *collection* where both *chunks* exist. Namespaces take form of `<database>.<collection>`.

field array bounds An array that contains the minimum and maximum key values of the new chunk.

Behavior

Note: Use the `mergeChunks` (page 289) only in special circumstances such as cleaning up your *sharded cluster* after removing many documents.

In order to successfully merge chunks, the following *must* be true

- In the `bounds` field, `<minkey>` and `<maxkey>` must correspond to the lower and upper bounds of the *chunks* to merge.
- The two chunks must reside on the same shard.
- The two chunks must be contiguous.
- One or both chunks must be empty.

`mergeChunks` (page 289) returns an error if these conditions are not satisfied.

Return Messages On success, `mergeChunks` (page 289) returns to following document:

```
{ "ok" : 1 }
```

Another Operation in Progress `mergeChunks` (page 289) returns the following error message if another meta-data operation is in progress on the *chunks* (page 595) collection:

```
errmsg: "The collection's metadata lock is already taken."
```

If another process, such as balancer process, changes metadata while `mergeChunks` (page 289) is running, you may see this error. You can retry the `mergeChunks` (page 289) operation without side effects.

Chunks on Different Shards If the input *chunks* are not on the same *shard*, `mergeChunks` (page 289) returns an error similar to the following:

```
{
  "ok" : 0,
  "errmsg" : "could not merge chunks, collection test.users does not contain a chunk ending at { us
}
```

Noncontiguous Chunks If the input *chunks* are not contiguous, `mergeChunks` (page 289) returns an error similar to the following:

```
{
  "ok" : 0,
  "errmsg" : "could not merge chunks, collection test.users has more than 2 chunks between [{ usern
```

Documents in Both Chunks If neither input *chunk* is empty, `mergeChunks` (page 289) returns an error similar to the following:

```
{
  "ok" : 0,
  "errmsg" : "could not merge chunks, collection test.users has more than one non-empty chunk between
```

setShardVersion

setShardVersion

`setShardVersion` (page 291) is an internal command that supports sharding functionality.

shardCollection

Definition

shardCollection

Enables a collection for sharding and allows MongoDB to begin distributing data among shards. You must run `enableSharding` (page 288) on a database before running the `shardCollection` (page 291) command. `shardCollection` (page 291) has the following form:

```
{ shardCollection: "<database>.<collection>", key: <shardkey> }
```

`shardCollection` (page 291) has the following fields:

field string shardCollection The *namespace* of the collection to shard in the form `<database>.<collection>`.

field document key The index specification document to use as the shard key. The index must exist prior to the `shardCollection` (page 291) command, unless the collection is empty. If the collection is empty, in which case MongoDB creates the index prior to sharding the collection. New in version 2.4: The key may be in the form `{ field : "hashed" }`, which will use the specified field as a hashed shard key.

field Boolean unique When `true`, the `unique` option ensures that the underlying index enforces a unique constraint. Hashed shard keys do not support unique constraints.

field integer numInitialChunks To support *hashed sharding* added in MongoDB 2.4, `numInitialChunks` specifies the number of chunks to create when sharding an collection with a hashed shard key. MongoDB will then create and balance chunks across the cluster. The `numInitialChunks` must be less than 8192 per shard.

Considerations

Use Do **not** run more than one `shardCollection` (page 291) command on the same collection at the same time. MongoDB provides no method to deactivate sharding for a collection after calling `shardCollection` (page 291). Additionally, after `shardCollection` (page 291), you cannot change shard keys or modify the value of any field used in your shard key index.

Shard Keys Choosing the best shard key to effectively distribute load among your shards requires some planning. Review *sharding-shard-key* regarding choosing a shard key.

Hashed Shard Keys New in version 2.4.

Hashed shard keys use a hashed index of a single field as the shard key.

Note: If chunk migrations are in progress while creating a hashed shard key collection, the initial chunk distribution may be uneven until the balancer automatically balances the collection.

Example The following operation enables sharding for the `people` collection in the `records` database and uses the `zipcode` field as the *shard key*:

```
db.runCommand( { shardCollection: "records.people", key: { zipcode: 1 } } )
```

Additional

Information <http://docs.mongodb.org/manualsharding>, <http://docs.mongodb.org/manualcore/sharding>, and <http://docs.mongodb.org/manualtutorial/deplo>

shardingState

shardingState

`shardingState` (page 292) is an admin command that reports if `mongod` (page 503) is a member of a *sharded cluster*. `shardingState` (page 292) has the following prototype form:

```
{ shardingState: 1 }
```

For `shardingState` (page 292) to detect that a `mongod` (page 503) is a member of a sharded cluster, the `mongod` (page 503) must satisfy the following conditions:

- 1.the `mongod` (page 503) is a primary member of a replica set, and
- 2.the `mongod` (page 503) instance is a member of a sharded cluster.

If `shardingState` (page 292) detects that a `mongod` (page 503) is a member of a sharded cluster, `shardingState` (page 292) returns a document that resembles the following prototype:

```
{
  "enabled" : true,
  "configServer" : "<configdb-string>",
  "shardName" : "<string>",
  "shardHost" : "string:",
  "versions" : {
    "<database>.<collection>" : Timestamp(<...>),
    "<database>.<collection>" : Timestamp(<...>)
  },
  "ok" : 1
}
```

Otherwise, `shardingState` (page 292) will return the following document:


```
{ "note" : "from execCommand", "ok" : 0, "errmsg" : "not master" }
```

The response from `shardingState` (page 292) when used with a *config server* is:

```
{ "enabled": false, "ok": 1 }
```

Note: `mongos` (page 518) instances do not provide the `shardingState` (page 292).

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed; however, the operation is typically short lived.

unsetSharding

unsetSharding

`unsetSharding` (page 293) is an internal command that supports sharding functionality.

split

Definition

split

Splits a *chunk* in a *sharded cluster* into two chunks. The `mongos` (page 518) instance splits and manages chunks automatically, but for exceptional circumstances the `split` (page 293) command does allow administrators to manually create splits. See <http://docs.mongodb.org/manual/tutorial/split-chunks-in-sharded-cluster> for information on these circumstances, and on the MongoDB shell commands that wrap `split` (page 293).

The `split` (page 293) command uses the following form:

```
db.adminCommand( { split: <database>.<collection>,  
                  <find|middle|bounds> } )
```

The `split` (page 293) command takes a document with the following fields:

field string split The name of the *collection* where the *chunk* exists. Specify the collection's full *namespace*, including the database name.

field document find An query statement that specifies an equality match on the shard key. The match selects the chunk that contains the specified document. You must specify only one of the following: `find`, `bounds`, or `middle`.

You cannot use the `find` option on an empty collection.

field array bounds New in version 2.4: The bounds of a chunk to split. `bounds` applies to chunks in collections partitioned using a *hashed shard key*. The parameter's array must consist of two documents specifying the lower and upper shard-key values of the chunk. The values must match the minimum and maximum values of an existing chunk. Specify only one of the following: `find`, `bounds`, or `middle`.

You cannot use the `bounds` option on an empty collection.

field document middle The document to use as the split point to create two chunks. `split` (page 293) requires one of the following options: `find`, `bounds`, or `middle`.

Considerations When used with either the `find` or the `bounds` option, the `split` (page 293) command splits the chunk along the median. As such, the command cannot use the `find` or the `bounds` option to split an empty chunk since an empty chunk has no median.

To create splits in empty chunks, use either the `middle` option with the `split` (page 293) command or use the `splitAt` command.

Command Formats To create a chunk split, connect to a `mongos` (page 518) instance, and issue the following command to the `admin` database:

```
db.adminCommand( { split: <database>.<collection>,
                  find: <document> } )
```

Or:

```
db.adminCommand( { split: <database>.<collection>,
                  middle: <document> } )
```

Or:

```
db.adminCommand( { split: <database>.<collection>,
                  bounds: [ <lower>, <upper> ] } )
```

To create a split for a collection that uses a *hashed shard key*, use the `bounds` parameter. Do *not* use the `middle` parameter for this purpose.

Warning: Be careful when splitting data in a sharded collection to create new chunks. When you shard a collection that has existing data, MongoDB automatically creates chunks to evenly distribute the collection. To split data effectively in a sharded cluster you must consider the number of documents in a chunk and the average document size to create a uniform chunk size. When chunks have irregular sizes, shards may have an equal number of chunks but have very different data sizes. Avoid creating splits that lead to a collection with differently sized chunks.

See also:

`moveChunk` (page 296), `sh.moveChunk()` (page 177), `sh.splitAt()` (page 179), and `sh.splitFind()` (page 179), which wrap the functionality of `split` (page 293).

Examples The following sections provide examples of the `split` (page 293) command.

Split a Chunk in Half

```
db.runCommand( { split : "test.people", find : { _id : 99 } } )
```

The `split` (page 293) command identifies the chunk in the `people` collection of the `test` database, that holds documents that match `{ _id : 99 }`. `split` (page 293) does not require that a match exist, in order to identify the appropriate chunk. Then the command splits it into two chunks of equal size.

Note: `split` (page 293) creates two equal chunks by range as opposed to size, and does not use the selected point as a boundary for the new chunks

Define an Arbitrary Split Point To define an arbitrary split point, use the following form:

```
db.runCommand( { split : "test.people", middle : { _id : 99 } } )
```

The `split` (page 293) command identifies the chunk in the `people` collection of the `test` database, that would hold documents matching the query `{ _id : 99 }`. `split` (page 293) does not require that a match exist, in order to identify the appropriate chunk. Then the command splits it into two chunks, with the matching document as the lower bound of one of the split chunks.

This form is typically used when *pre-splitting* data in a collection.

Split a Chunk Using Values of a Hashed Shard Key This example uses the *hashed shard key* `userid` in a `people` collection of a `test` database. The following command uses an array holding two single-field documents to represent the minimum and maximum values of the hashed shard key to split the chunk:

```
db.runCommand( { split: "test.people",
                  bounds : [ { userid: NumberLong("-5838464104018346494") },
                             { userid: NumberLong("-5557153028469814163") }
                          ] } )
```

Note: MongoDB uses the 64-bit *NumberLong* type to represent the hashed value.

Use `sh.status()` (page 180) to see the existing bounds of the shard keys.

Metadata Lock Error If another process in the `mongos` (page 518), such as a balancer process, changes metadata while `split` (page 293) is running, you may see a metadata lock error.

```
errmsg: "The collection's metadata lock is already taken."
```

This message indicates that the split has failed with no side effects. Retry the `split` (page 293) command.

splitChunk

Definition

splitChunk

An internal administrative command. To split chunks, use the `sh.splitFind()` (page 179) and `sh.splitAt()` (page 179) functions in the `mongo` (page 527) shell.

Warning: Be careful when splitting data in a sharded collection to create new chunks. When you shard a collection that has existing data, MongoDB automatically creates chunks to evenly distribute the collection. To split data effectively in a sharded cluster you must consider the number of documents in a chunk and the average document size to create a uniform chunk size. When chunks have irregular sizes, shards may have an equal number of chunks but have very different data sizes. Avoid creating splits that lead to a collection with differently sized chunks.

See also:

`moveChunk` (page 296) and `sh.moveChunk()` (page 177).

The `splitChunk` (page 295) command takes a document with the following fields:

- field string ns** The complete *namespace* of the *chunk* to split.
- field document keyPattern** The *shard key*.
- field document min** The lower bound of the shard key for the chunk to split.
- field document max** The upper bound of the shard key for the chunk to split.
- field string from** The *shard* that owns the chunk to split.

field document splitKeys The split point for the chunk.

field document shardId The shard.

splitVector

splitVector

Is an internal command that supports meta-data operations in sharded clusters.

medianKey

medianKey

`medianKey` (page 296) is an internal command.

moveChunk

Definition

moveChunk

Internal administrative command. Moves *chunks* between *shards*. Issue the `moveChunk` (page 296) command via a `mongos` (page 518) instance while using the *admin database*. Use the following forms:

```
db.runCommand( { moveChunk : <namespace> ,
                  find : <query> ,
                  to : <string> ,
                  _secondaryThrottle : <boolean> ,
                  _waitForDelete : <boolean> } )
```

Alternately:

```
db.runCommand( { moveChunk : <namespace> ,
                  bounds : <array> ,
                  to : <string> ,
                  _secondaryThrottle : <boolean> ,
                  _waitForDelete : <boolean> } )
```

The `moveChunk` (page 296) command has the following fields:

field string moveChunk The *namespace* of the *collection* where the *chunk* exists. Specify the collection's full namespace, including the database name.

field document find An equality match on the shard key that specifies the shard-key value of the chunk to move. Specify either the `bounds` field or the `find` field but not both.

field array bounds The bounds of a specific chunk to move. The array must consist of two documents that specify the lower and upper shard key values of a chunk to move. Specify either the `bounds` field or the `find` field but not both. Use `bounds` to move chunks in collections partitioned using a *hashed shard key*.

field string to The name of the destination shard for the chunk.

field Boolean _secondaryThrottle Defaults to `true`. When `true`, the balancer waits for replication to *secondaries* when it copies and deletes data during chunk migrations. For details, see *sharded-cluster-config-secondary-throttle*.

field Boolean _waitForDelete Internal option for testing purposes. The default is `false`. If set to `true`, the delete phase of a `moveChunk` (page 296) operation blocks.

The value of `bounds` takes the form:

```
[ { hashedField : <minValue> } ,
  { hashedField : <maxValue> } ]
```

The *chunk migration* section describes how chunks move between shards on MongoDB.

See also:

`split` (page 293), `sh.moveChunk()` (page 177), `sh.splitAt()` (page 179), and `sh.splitFind()` (page 179).

Return Messages `moveChunk` (page 296) returns the following error message if another metadata operation is in progress on the `chunks` (page 595) collection:

```
errmsg: "The collection's metadata lock is already taken."
```

If another process, such as a balancer process, changes meta data while `moveChunk` (page 296) is running, you may see this error. You may retry the `moveChunk` (page 296) operation without side effects.

Note: Only use the `moveChunk` (page 296) in special circumstances such as preparing your *sharded cluster* for an initial ingestion of data, or a large bulk import operation. In most cases allow the balancer to create and balance chunks in sharded clusters. See <http://docs.mongodb.org/manual/tutorial/create-chunks-in-sharded-cluster> for more information.

movePrimary

movePrimary

In a *sharded cluster*, this command reassigns the database's *primary shard*, which holds all un-sharded collections in the database. `movePrimary` (page 297) is an administrative command that is only available for `mongos` (page 518) instances. Only use `movePrimary` (page 297) when removing a shard from a sharded cluster.

`movePrimary` (page 297) changes the primary shard for a database in the cluster metadata, and migrates all un-sharded collections to the specified shard. Use the command with the following form:

```
{ movePrimary : "test", to : "shard0001" }
```

When the command returns, the database's primary location will shift to the designated *shard*. To fully decommission a shard, use the `removeShard` (page 288) command.

Considerations

Limitations Only use `movePrimary` (page 297) when:

- the database does not contain any collections, *or*
- you have drained all sharded collections using the `removeShard` (page 288) command.

Use If you use the `movePrimary` (page 297) command to move un-sharded collections, you must either restart all `mongos` (page 518) instances, or use the `flushRouterConfig` (page 283) command on all `mongos` (page 518) instances before writing any data to the cluster. This action notifies the `mongos` (page 518) of the new shard for the database.

If you do not update the `mongos` (page 518) instances' metadata cache after using `movePrimary` (page 297), the `mongos` (page 518) may not write data to the correct shard, to recover you must manually intervene.

Additional Information See <http://docs.mongodb.org/manual/tutorial/remove-shards-from-cluster> for a complete procedure.

isdbgrid

isdbgrid

This command verifies that a process is a `mongos` (page 518).

If you issue the `isdbgrid` (page 298) command when connected to a `mongos` (page 518), the response document includes the `isdbgrid` field set to 1. The returned document is similar to the following:

```
{ "isdbgrid" : 1, "hostname" : "app.example.net", "ok" : 1 }
```

If you issue the `isdbgrid` (page 298) command when connected to a `mongod` (page 503), MongoDB returns an error document. The `isdbgrid` (page 298) command is not available to `mongod` (page 503). The error document, however, also includes a line that reads `"isdbgrid" : 1`, just as in the document returned for a `mongos` (page 518). The error document is similar to the following:

```
{
  "errmsg" : "no such cmd: isdbgrid",
  "bad cmd" : {
    "isdbgrid" : 1
  },
  "ok" : 0
}
```

You can instead use the `isMaster` (page 280) command to determine connection to a `mongos` (page 518). When connected to a `mongos` (page 518), the `isMaster` (page 280) command returns a document that contains the string `isdbgrid` in the `msg` field.

See also:

<http://docs.mongodb.org/manual/sharding> for more information about MongoDB's sharding functionality.

Instance Administration Commands

Administration Commands

Name	Description
<code>renameCollection</code> (page 299)	Changes the name of an existing collection.
<code>copydb</code> (page 300)	Copies a database from a remote host to the current host.
<code>dropDatabase</code> (page 304)	Removes the current database.
<code>drop</code> (page 304)	Removes the specified collection from the database.
<code>create</code> (page 304)	Creates a collection and sets collection parameters.
<code>clone</code> (page 305)	Copies a database from a remote host to the current host.
<code>cloneCollection</code> (page 306)	Copies a collection from a remote host to the current host.
<code>cloneCollectionAsCapped</code> (page 306)	Copies a non-capped collection as a new <i>capped collection</i> .
<code>closeAllDatabases</code> (page 307)	Internal command that invalidates all cursors and closes open database files.
<code>convertToCapped</code> (page 307)	Converts a non-capped collection to a capped collection.
<code>filemd5</code> (page 308)	Returns the <i>md5</i> hash for files stored using <i>GridFS</i> .
<code>createIndexes</code> (page 308)	Builds one or more indexes for a collection.
<code>dropIndexes</code> (page 311)	Removes indexes from a collection.
<code>fsync</code> (page 312)	Flushes pending writes to the storage layer and locks the database to allow backups.
<code>clean</code> (page 313)	Internal namespace administration command.
<code>connPoolSync</code> (page 313)	Internal command to flush connection pool.
<code>compact</code> (page 313)	Defragments a collection and rebuilds the indexes.
<code>collMod</code> (page 316)	Add flags to collection to modify the behavior of MongoDB.
<code>reIndex</code> (page 317)	Rebuilds all indexes on a collection.
<code>setParameter</code> (page 318)	Modifies configuration options.
<code>getParameter</code> (page 318)	Retrieves configuration options.
<code>repairDatabase</code> (page 319)	Repairs any errors and inconsistencies with the data storage.
<code>touch</code> (page 321)	Loads documents and indexes from data storage to memory.
<code>shutdown</code> (page 321)	Shuts down the <i>mongod</i> (page 503) or <i>mongos</i> (page 518) process.
<code>logRotate</code> (page 322)	Rotates the MongoDB logs to prevent a single file from taking too much space.

renameCollection

Definition

renameCollection

Changes the name of an existing collection. Specify collections to `renameCollection` (page 299) in the form of a complete *namespace*, which includes the database name. Issue the `renameCollection` (page 299) command against the *admin database*. The command takes the following form:

```
{ renameCollection: "<source_namespace>", to: "<target_namespace>", dropTarget: <true|false> }
```

The command contains the following fields:

field string renameCollection The *namespace* of the collection to rename. The namespace is a combination of the database name and the name of the collection.

field string to The new namespace of the collection. If the new namespace specifies a different database, the `renameCollection` (page 299) command copies the collection to the new database and drops the source collection.

field boolean dropTarget If `true`, `mongod` (page 503) will drop the target of `renameCollection` (page 299) prior to renaming the collection.

`renameCollection` (page 299) is suitable for production environments; *however*:

- `renameCollection` (page 299) blocks all database activity for the duration of the operation.
- `renameCollection` (page 299) is **not** compatible with sharded collections.

Warning: `renameCollection` (page 299) fails if `target` is the name of an existing collection *and* you do not specify `dropTarget: true`.
If the `renameCollection` (page 299) operation does not complete the `target` collection and indexes will not be usable and will require manual intervention to clean up.

Exceptions

exception 10026 Raised if the `source` namespace does not exist.

exception 10027 Raised if the `target` namespace exists and `dropTarget` is either `false` or unspecified.

exception 15967 Raised if the `target` namespace is an invalid collection name.

Shell Helper The shell helper `db.collection.renameCollection()` (page 65) provides a simpler interface to using this command within a database. The following is equivalent to the previous example:

```
db.source-namespace.renameCollection( "target" )
```

Warning: You cannot use `renameCollection` (page 299) with sharded collections.

Warning: This command obtains a global write lock and will block other operations until it has completed.

copydb

Definition

copydb

Copies a database from a remote host to the current host or copies a database to another database within the current host. Run `copydb` (page 300) in the `admin` database of the destination server with the following syntax:

```
{ copydb: 1,
  fromhost: <hostname>,
  fromdb: <database>,
  todb: <database>,
  slaveOk: <bool>,
  username: <username>,
  nonce: <nonce>,
  key: <key> }
```

`copydb` (page 300) accepts the following options:

field string fromhost Hostname of the remote source `mongod` (page 503) instance. Omit `fromhost` to copy from one database to another on the same server.

field string fromdb Name of the source database.

field string todb Name of the target databases.

field boolean slaveOk Set `slaveOk` to `true` to allow `copydb` (page 300) to copy data from secondary members as well as the primary. `fromhost` must also be set.

field string username The username credentials on the `fromhost` MongoDB deployment.

field string nonce A single use shared secret generated on the remote server, i.e. `fromhost`, using the `copydbgetnonce` (page 250) command. See [Authentication](#) (page 302) for details.

field string key A hash of the password used for authentication. See [Authentication](#) (page 302) for details.

The `mongo` (page 527) shell provides the `db.copyDatabase()` (page 100) wrapper for the `copydb` (page 300) command.

Behavior Be aware of the following properties of `copydb` (page 300):

- `copydb` (page 300) runs on the destination `mongod` (page 503) instance, i.e. the host receiving the copied data.
- `copydb` (page 300) creates the target database if it does not exist.
- `copydb` (page 300) requires enough free disk space on the host instance for the copied database. Use the `db.stats()` (page 119) operation to check the size of the database on the source `mongod` (page 503) instance.
- `copydb` (page 300) and `clone` (page 305) do not produce point-in-time snapshots of the source database. Write traffic to the source or destination database during the copy process will result in divergent data sets.
- `copydb` (page 300) does not lock the destination server during its operation, so the copy will occasionally yield to allow other operations to complete.

Required Access Changed in version 2.6.

On systems running with authorization, the `copydb` (page 300) command requires the following authorization on the target and source databases.

Source Database (`fromdb`)

Source is non-admin Database If the source database is a non-admin database, you must have privileges that specify `find` action on the source database, and `find` action on the `system.js` collection in the source database. For example:

```
{ resource: { db: "mySourceDB", collection: "" }, actions: [ "find" ] }
{ resource: { db: "mySourceDB", collection: "system.js" }, actions: [ "find" ] }
```

If the source database is on a remote server, you also need the `find` action on the `system.indexes` and `system.namespaces` collections in the source database; e.g.

```
{ resource: { db: "mySourceDB", collection: "system.indexes" }, actions: [ "find" ] }
{ resource: { db: "mySourceDB", collection: "system.namespaces" }, actions: [ "find" ] }
```

Source is admin Database If the source database is the admin database, you must have privileges that specify find action on the admin database, and find action on the system.js, system.users, system.roles, and system.version collections in the admin database. For example:

```
{ resource: { db: "admin", collection: "" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.js" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.users" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.roles" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.version" }, actions: [ "find" ] }
```

If the source database is on a remote server, then you also need the find action on the system.indexes and system.namespaces collections in the admin database; e.g.

```
{ resource: { db: "admin", collection: "system.indexes" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.namespaces" }, actions: [ "find" ] }
```

Source Database is on a Remote Server If copying from a remote server and the remote server has authentication enabled, you must authenticate to the remote host as a user with the proper authorization. See [Authentication](#) (page 302).

Target Database (todb)

Copy from non-admin Database If the source database is not the admin database, you must have privileges that specify insert and createIndex actions on the target database, and insert action on the system.js collection in the target database. For example:

```
{ resource: { db: "myTargetDB", collection: "" }, actions: [ "insert", "createIndex" ] }
{ resource: { db: "myTargetDB", collection: "system.js" }, actions: [ "insert" ] }
```

Copy from admin Database If the source database is the admin database, you must have privileges that specify insert and createIndex actions on the target database, and insert action on the system.js, system.users, system.roles, and system.version collections in the target database. For example:

```
{ resource: { db: "myTargetDB", collection: "" }, actions: [ "insert", "createIndex" ] },
{ resource: { db: "myTargetDB", collection: "system.js" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.users" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.roles" }, actions: [ "insert" ] },
{ resource: { db: "myTargetDB", collection: "system.version" }, actions: [ "insert" ] }
```

Authentication If copying from a remote server and the remote server has authentication enabled, then you must include a username, nonce, and key.

The nonce is a one-time password that you request from the remote server using the `copydbgetnonce` (page 250) command, as in the following:

```
use admin
mynonce = db.runCommand( { copydbgetnonce : 1, fromhost: <hostname> } ).nonce
```

If running the `copydbgetnonce` (page 250) command directly on the remote host, you can omit the fromhost field in the `copydbgetnonce` (page 250) command.

The key is a hash generated as follows:

```
hex_md5(mynonce + username + hex_md5(username + ":mongo:" + password))
```

Replica Sets With *read preference* configured to set the `slaveOk` option to `true`, you may run `copydb` (page 300) on a *secondary* member of a *replica set*.

Sharded Clusters

- Do not use `copydb` (page 300) from a `mongos` (page 518) instance.
- Do not use `copydb` (page 300) to copy databases that contain sharded collections.

Examples

Copy on the Same Host To copy from the same host, omit the `fromhost` field.

The following command copies the `test` database to a new `records` database on the current `mongod` (page 503) instance:

```
use admin
db.runCommand({
  copydb: 1,
  fromdb: "test",
  todb: "records"
})
```

Copy from a Remote Host to the Current Host To copy from a remote host, include the `fromhost` field.

The following command copies the `test` database from the remote host `example.net` to a new `records` database on the current `mongod` (page 503) instance:

```
use admin
db.runCommand({
  copydb: 1,
  fromdb: "test",
  todb: "records",
  fromhost: "example.net"
})
```

Copy Databases from Remote `mongod` Instances that Enforce Authentication To copy from a remote host that enforces authentication, include the `fromhost`, `username`, `nonce` and `key` fields.

The following command copies the `test` database from a remote host `example.net` that runs with authorization to a new `records` database on the local `mongod` (page 503) instance. Because the `example.net` has authorization enabled, the command includes the `username`, `nonce` and `key` fields:

```
use admin
db.runCommand({
  copydb: 1,
  fromdb: "test",
  todb: "records",
  fromhost: "example.net",
  username: "reportingAdmin",
  nonce: "<nonce>",
  key: "<key>"
})
```

```
    key: "<passwordhash>"
  })
```

See also:

- `db.copyDatabase()` (page 100)
- `clone` (page 305) and `db.cloneDatabase()` (page 100)
- <http://docs.mongodb.org/manualcore/backups> and <http://docs.mongodb.org/manualcore/import>

dropDatabase**dropDatabase**

The `dropDatabase` (page 304) command drops a database, deleting the associated data files. `dropDatabase` (page 304) operates on the current database.

In the shell issue the `use <database>` command, replacing `<database>` with the name of the database you wish to delete. Then use the following command form:

```
{ dropDatabase: 1 }
```

The `mongo` (page 527) shell also provides the following equivalent helper method:

```
db.dropDatabase();
```

Warning: This command obtains a global write lock and will block other operations until it has completed.

drop**drop**

The `drop` (page 304) command removes an entire collection from a database. The command has following syntax:

```
{ drop: <collection_name> }
```

The `mongo` (page 527) shell provides the equivalent helper method `db.collection.drop()` (page 29).

This command also removes any indexes associated with the dropped collection.

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed.

create**Definition****create**

Explicitly creates a collection. `create` (page 304) has the following form:

```
{ create: <collection_name>,
  capped: <true|false>,
  autoIndexId: <true|false>,
  size: <max_size>,
  max: <max_documents>,
  flags: <0|1>
}
```

`create` (page 304) has the following fields:

field string create The name of the new collection.

field Boolean capped To create a *capped collection*, specify `true`. If you specify `true`, you must also set a maximum size in the `size` field.

field Boolean autoIndexId Specify `false` to disable the automatic creation of an index on the `_id` field. Before 2.2, the default value for `autoIndexId` was `false`.

field integer size The maximum size for the capped collection. Once a capped collection reaches its maximum size, MongoDB overwrites older old documents with new documents. The `size` field is required for capped collections.

field integer max The maximum number of documents to keep in the capped collection. The `size` limit takes precedence over this limit. If a capped collection reaches its maximum size before it reaches the maximum number of documents, MongoDB removes old documents. If you use this limit, ensure that the `size` limit is sufficient to contain the documents limit.

field integer flags New in version 2.6.

Set to 0 to disable the `usePowerOf2Sizes` (page 316) allocation strategy for this collection, or 1 to enable `usePowerOf2Sizes` (page 316). Defaults to 1 unless the `newCollectionsUsePowerOf2Sizes` parameter is set to `false`.

For more information on the `autoIndexId` field in versions before 2.2, see *`_id` Fields and Indexes on Capped Collections* (page 672).

The `db.createCollection()` (page 102) method wraps the `create` (page 304) command.

Considerations The `create` (page 304) command obtains a write lock on the affected database and will block other operations until it has completed. The write lock for this operation is typically short lived. However, allocations for large capped collections may take longer.

Example To create a *capped collection* limited to 64 kilobytes, issue the command in the following form:

```
db.runCommand( { create: "collection", capped: true, size: 64 * 1024 } )
```

clone

clone

The `clone` (page 305) command clones a database from a remote MongoDB instance to the current host. `clone` (page 305) copies the database on the remote instance with the same name as the current database. The command takes the following form:

```
{ clone: "db1.example.net:27017" }
```

Replace `db1.example.net:27017` above with the resolvable hostname for the MongoDB instance you wish to copy from. Note the following behaviors:

- `clone` (page 305) can run against a *slave* or a non-*primary* member of a *replica set*.
- `clone` (page 305) does not snapshot the database. If any clients update the database you're copying at any point during the clone operation, the resulting database may be inconsistent.
- You must run `clone` (page 305) on the **destination server**.
- The destination server is not locked for the duration of the `clone` (page 305) operation. This means that `clone` (page 305) will occasionally yield to allow other operations to complete.

See `copydb` (page 300) for similar functionality with greater flexibility.

Warning: This command obtains an intermittent *write lock* on the destination server, which can block other operations until it completes.

cloneCollection

Definition

cloneCollection

Copies a collection from a remote `mongod` (page 503) instance to the current `mongod` (page 503) instance. `cloneCollection` (page 306) creates a collection in a database with the same name as the remote collection's database. `cloneCollection` (page 306) takes the following form:

```
{ cloneCollection: "<namespace>", from: "<hostname>", query: { <query> } }
```

Important: You cannot clone a collection through a `mongos` (page 518) but must connect directly to the `mongod` (page 503) instance.

`cloneCollection` (page 306) has the following fields:

field string cloneCollection The *namespace* of the collection to rename. The namespace is a combination of the database name and the name of the collection.

field string from Specify a resolvable hostname and optional port number of the remote server where the specified collection resides.

field document query A query that filters the documents in the remote collection that `cloneCollection` (page 306) will copy to the current database.

Example

```
{ cloneCollection: "users.profiles", from: "mongodb.example.net:27017", query: { active: true } }
```

This operation copies the `profiles` collection from the `users` database on the server at `mongodb.example.net`. The operation only copies documents that satisfy the query `{ active: true }` and does not copy indexes. `cloneCollection` (page 306) copies indexes by default. The query arguments is optional.

If, in the above example, the `profiles` collection exists in the `users` database, then MongoDB appends documents from the remote collection to the destination collection.

cloneCollectionAsCapped

cloneCollectionAsCapped

The `cloneCollectionAsCapped` (page 306) command creates a new *capped collection* from an existing, non-capped collection within the same database. The operation does not affect the original non-capped collection.

The command has the following syntax:

```
{ cloneCollectionAsCapped: <existing collection>, toCollection: <capped collection>, size: <capped size> }
```

The command copies an existing collection and creates a new capped collection with a maximum size specified by the `capped size` in bytes. The name of the new capped collection must be distinct and cannot be the same as that of the original existing collection. To replace the original non-capped collection with a capped collection, use the `convertToCapped` (page 307) command.

During the cloning, the `cloneCollectionAsCapped` (page 306) command exhibit the following behavior:

- MongoDB will transverse the documents in the original collection in *natural order* as they're loaded.
- If the `capped size` specified for the new collection is smaller than the size of the original uncapped collection, then MongoDB will begin overwriting earlier documents in insertion order, which is *first in, first out* (e.g “FIFO”).

`closeAllDatabases`

`closeAllDatabases`

`closeAllDatabases` (page 307) is an internal command that invalidates all cursors and closes the open database files. The next operation that uses the database will reopen the file.

Warning: This command obtains a global write lock and will block other operations until it has completed.

`convertToCapped`

`convertToCapped`

The `convertToCapped` (page 307) command converts an existing, non-capped collection to a *capped collection* within the same database.

The command has the following syntax:

```
{convertToCapped: <collection>, size: <capped size> }
```

`convertToCapped` (page 307) takes an existing collection (`<collection>`) and transforms it into a capped collection with a maximum size in bytes, specified to the `size` argument (`<capped size>`).

During the conversion process, the `convertToCapped` (page 307) command exhibit the following behavior:

- MongoDB transverses the documents in the original collection in *natural order* and loads the documents into a new capped collection.
- If the `capped size` specified for the capped collection is smaller than the size of the original uncapped collection, then MongoDB will overwrite documents in the capped collection based on insertion order, or *first in, first out* order.
- Internally, to convert the collection, MongoDB uses the following procedure
 - `cloneCollectionAsCapped` (page 306) command creates the capped collection and imports the data.
 - MongoDB drops the original collection.
 - `renameCollection` (page 299) renames the new capped collection to the name of the original collection.

Note: MongoDB does not support the `convertToCapped` (page 307) command in a sharded cluster.

Warning: The `convertToCapped` (page 307) will not recreate indexes from the original collection on the new collection, other than the index on the `_id` field. If you need indexes on this collection you will need to create these indexes after the conversion is complete.

See also:

`create` (page 304)

Warning: This command obtains a global write lock and will block other operations until it has completed.

filemd5

filemd5

The `filemd5` (page 308) command returns the *md5* hashes for a single file stored using the *GridFS* specification. Client libraries use this command to verify that files are correctly written to MongoDB. The command takes the `files_id` of the file in question and the name of the GridFS root collection as arguments. For example:

```
{ filemd5: ObjectId("4f1f10e37671b50e4ecd2776"), root: "fs" }
```

createIndexes New in version 2.6.

Definition

createIndexes

Builds one or more indexes on a collection. The `createIndexes` (page 308) command takes the following form:

```
db.runCommand(
  {
    createIndexes: <collection>,
    indexes: [
      {
        key: {
          <key-value_pair>,
          <key-value_pair>,
          ...
        },
        name: <index_name>,
        <option1>,
        <option2>,
        ...
      },
      { ... },
      { ... }
    ]
  }
)
```

The `createIndexes` (page 308) command takes the following fields:

field string createIndexes The collection for which to create indexes.

field array indexes Specifies the indexes to create. Each document in the array specifies a separate index.

Each document in the `indexes` array can take the following fields:

field document key Specifies the index's fields. For each field, specify a key-value pair in which the key is the name of the field to index and the value is either the index direction or `index` type. If specifying direction, specify 1 for ascending or -1 for descending.

field string name A name that uniquely identifies the index.

field string ns The *namespace* (i.e. `<database>.<collection>`) of the collection for which to create the index. If you omit `ns`, MongoDB generates the namespace.

param Boolean background Builds the index in the background so that building an index does *not* block other database activities. Specify `true` to build in the background. The default value is `false`.

param Boolean unique Creates a unique index so that the collection will not accept insertion of documents where the index key or keys match an existing value in the index. Specify `true` to create a unique index. The default value is `false`.

The option is *unavailable* for hashed indexes.

param Boolean dropDups Creates a unique index on a field that *may* have duplicates. MongoDB indexes only the first occurrence of a key and **removes** all documents from the collection that contain subsequent occurrences of that key. Specify `true` to create unique index. The default value is `false`.

The option is *unavailable* for hashed indexes.

Deprecated since version 2.6.

Warning: `dropDups` will delete data from your collection when building the index.

param Boolean sparse If `true`, the index only references documents with the specified field. These indexes use less space but behave differently in some situations (particularly sorts). The default value is `false`. See <http://docs.mongodb.org/manualcore/index-sparse> for more information.

Changed in version 2.6: 2dsphere indexes are sparse by default and ignore this option. For a compound index that includes 2dsphere index key(s) along with keys of other types, only the 2dsphere index fields determine whether the index references a document.

2d, geoHaystack, and text indexes behave similarly to the 2dsphere indexes.

param integer expireAfterSeconds Specifies a value, in seconds, as a *TTL* to control how long MongoDB retains documents in this collection. See <http://docs.mongodb.org/manualtutorial/expire-data> for more information on this functionality. This applies only to *TTL* indexes.

param index version v The index version number. The default index version depends on the version of `mongodb` (page 503) running when creating the index. Before version 2.0, the this value was 0; versions 2.0 and later use version 1, which provides a smaller and faster index format. Specify a different index version *only* in unusual situations.

param document weights For text indexes, a document that contains field and weight pairs. The weight is an integer ranging from 1 to 99,999 and denotes the significance of the field relative to the other indexed fields in terms of the score. You can specify weights for some or all the indexed fields. See <http://docs.mongodb.org/manualtutorial/control-results-of-text-search> to adjust the scores. The default value is 1.

param string default_language For text indexes, the language that determines the list of stop words and the rules for the stemmer and tokenizer. See *text-search-languages* for the available languages and <http://docs.mongodb.org/manualtutorial/specify-language-for-text-index> for more information and examples. The default value is `english`.

param string language_override For text indexes, the name of the field, in the collection's documents, that contains the override language for the document. The default value is `language`. See *specify-language-field-text-index-example* for an example.

param integer textIndexVersion For text indexes, the text index version number. Version can be either 1 or 2.

In MongoDB 2.6, the default version is 2. MongoDB 2.4 can only support version 1.

New in version 2.6.

param integer 2dsphereIndexVersion For 2dsphere indexes, the 2dsphere index version number. Version can be either 1 or 2.

In MongoDB 2.6, the default version is 2. MongoDB 2.4 can only support version 1.

New in version 2.6.

param integer bits For 2d indexes, the number of precision of the stored *geohash* value of the location data.

The bits value ranges from 1 to 32 inclusive. The default value is 26.

param number min For 2d indexes, the lower inclusive boundary for the longitude and latitude values. The default value is -180.0.

param number max For 2d indexes, the upper inclusive boundary for the longitude and latitude values. The default value is 180.0.

param number bucketSize For geoHaystack indexes, specify the number of units within which to group the location values; i.e. group in the same bucket those location values that are within the specified number of units to each other.

The value must be greater than 0.

Considerations An index name, including the *namespace*, cannot be longer than the *Index Name Length* (page 605) limit.

Behavior Non-background indexing operations block all other operations on a database. If you specify multiple indexes to *createIndexes* (page 308), MongoDB builds the indexes serially.

If you create an index with one set of options and then issue *createIndexes* (page 308) with the same index fields but different options, MongoDB will not change the options nor rebuild the index. To change index options, drop the existing index with *db.collection.dropIndex()* (page 29) before running the new *createIndexes* (page 308) with the new options.

Example The following command builds two indexes on the *inventory* collection of the *products* database:

```
db.getSiblingDB("products").runCommand({
  createIndexes: "inventory",
  indexes: [
    {
      key: {
        item: 1,
        manufacturer: 1,
        model: 1
      },
      name: "item_manufacturer_model",
      unique: true
    },
    {
      key: {
```

```

        item: 1,
        supplier: 1,
        model: 1
    },
    name: "item_supplier_model",
    unique: true
}
]
}
)

```

When the indexes successfully finish building, MongoDB returns a results document that includes a status of "ok" : 1.

Output The `createIndexes` (page 308) command returns a document that indicates the success of the operation. The document contains some but not all of the following fields, depending on outcome:

`createIndexes.createdCollectionAutomatically`

If `true`, then the collection didn't exist and was created in the process of creating the index.

`createIndexes.numIndexesBefore`

The number of indexes at the start of the command.

`createIndexes.numIndexesAfter`

The number of indexes at the end of the command.

`createIndexes.ok`

A value of 1 indicates the indexes are in place. A value of 0 indicates an error.

`createIndexes.note`

This `note` is returned if an existing index or indexes already exist. This indicates that the index was not created or changed.

`createIndexes.errmsg`

Returns information about any errors.

`createIndexes.code`

The error code representing the type of error.

dropIndexes

dropIndexes

The `dropIndexes` (page 311) command drops one or all indexes from the current collection. To drop all indexes, issue the command like so:

```
{ dropIndexes: "collection", index: "*" }
```

To drop a single, issue the command by specifying the name of the index you want to drop. For example, to drop the index named `age_1`, use the following command:

```
{ dropIndexes: "collection", index: "age_1" }
```

The shell provides a useful command helper. Here's the equivalent command:

```
db.collection.dropIndex("age_1");
```

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed.

fsync

Definition

fsync

Forces the `mongod` (page 503) process to flush all pending writes from the storage layer to disk. Optionally, you can use `fsync` (page 312) to lock the `mongod` (page 503) instance and block write operations for the purpose of capturing backups.

As applications write data, MongoDB records the data in the storage layer and then writes the data to disk within the `syncPeriodSecs` interval, which is 60 seconds by default. Run `fsync` (page 312) when you want to flush writes to disk ahead of that interval.

The `fsync` (page 312) command has the following syntax:

```
{ fsync: 1, async: <Boolean>, lock: <Boolean> }
```

The `fsync` (page 312) command has the following fields:

field integer fsync Enter “1” to apply `fsync` (page 312).

field Boolean async Runs `fsync` (page 312) asynchronously. By default, the `fsync` (page 312) operation is synchronous.

field Boolean lock Locks `mongod` (page 503) instance and blocks all write operations.

Behavior An `fsync` (page 312) lock is only possible on *individual* `mongod` (page 503) instances of a sharded cluster, not on the entire cluster. To backup an entire sharded cluster, please see <http://docs.mongodb.org/manualadministration/backup-sharded-clusters> for more information.

If your `mongod` (page 503) has *journaling* enabled, consider using *another method* to create a back up of the data set.

After `fsync` (page 312), with lock, runs on a `mongod` (page 503), all write operations will block until a subsequent unlock. Read operations *may* also block. As a result, `fsync` (page 312), with lock, is not a reliable mechanism for making a `mongod` (page 503) instance operate in a read-only mode.

The database cannot be locked with `db.fsyncLock()` (page 109) while profiling is enabled. You must disable profiling before locking the database with `db.fsyncLock()` (page 109). Disable profiling using `db.setProfilingLevel()` (page 119) as follows in the `mongo` (page 527) shell:

```
db.setProfilingLevel(0)
```

Examples

Run Asynchronously The `fsync` (page 312) operation is synchronous by default To run `fsync` (page 312) asynchronously, use the `async` field set to `true`:

```
{ fsync: 1, async: true }
```

The operation returns immediately. To view the status of the `fsync` (page 312) operation, check the output of `db.currentOp()` (page 103).

Lock mongod Instance The primary use of `fsync` (page 312) is to lock the `mongod` (page 503) instance in order to back up the files withing `mongod` (page 503)’s `dbPath`. The operation flushes all data to the storage layer and blocks all write operations until you unlock the `mongod` (page 503) instance.

To lock the database, use the `lock` field set to `true`:

```
{ fsync: 1, lock: true }
```

You may continue to perform read operations on a [mongod](#) (page 503) instance that has a [fsync](#) (page 312) lock. However, after the first write operation all subsequent read operations wait until you unlock the [mongod](#) (page 503) instance.

Check Lock Status To check the state of the `fsync` lock, use `db.currentOp()` (page 103). Use the following JavaScript function in the shell to test if [mongod](#) (page 503) instance is currently locked:

```
serverIsLocked = function () {
    var co = db.currentOp();
    if (co && co.fsyncLock) {
        return true;
    }
    return false;
}
```

After loading this function into your [mongo](#) (page 527) shell session call it, with the following syntax:

```
serverIsLocked()
```

This function will return `true` if the [mongod](#) (page 503) instance is currently locked and `false` if the [mongod](#) (page 503) is not locked. To unlock the [mongod](#) (page 503), make a request for an unlock using the following operation:

```
db.getSiblingDB("admin").$cmd.sys.unlock.findOne();
```

Unlock mongod Instance To unlock the [mongod](#) (page 503) instance, use `db.fsyncUnlock()` (page 110):

```
db.fsyncUnlock();
```

clean

clean

[clean](#) (page 313) is an internal command.

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed.

connPoolSync

connPoolSync

[connPoolSync](#) (page 313) is an internal command.

compact

Definition

compact

New in version 2.0.

Rewrites and defragments all data in a collection, as well as all of the indexes on that collection. [compact](#) (page 313) has the following form:

```
{ compact: <collection name> }
```

`compact` (page 313) has the following fields:

field string compact The name of the collection.

field boolean force If `true`, `compact` (page 313) can run on the *primary* in a *replica set*. If `false`, `compact` (page 313) returns an error when run on a primary, because the command blocks all other activity. Beginning with version 2.2, `compact` (page 313) blocks activity only for the database it is compacting.

field number paddingFactor Describes the *record size* allocated for each document as a factor of the document size for all records compacted during the `compact` (page 313) operation. The `paddingFactor` does not affect the padding of subsequent record allocations after `compact` (page 313) completes. For more information, see *paddingFactor* (page 314).

field integer paddingBytes Sets the padding as an absolute number of bytes for all records compacted during the `compact` (page 313) operation. After `compact` (page 313) completes, `paddingBytes` does not affect the padding of subsequent record allocations. For more information, see *paddingBytes* (page 315).

`compact` (page 313) is similar to `repairDatabase` (page 319); however, `repairDatabase` (page 319) operates on an entire database.

Warning: Always have an up-to-date backup before performing server maintenance such as the `compact` (page 313) operation.

paddingFactor New in version 2.2.

The `paddingFactor` field takes the following range of values:

- Default: 1.0
- Minimum: 1.0 (no padding)
- Maximum: 4.0

If your updates increase the size of the documents, padding will increase the amount of space allocated to each document and avoid expensive document relocation operations within the data files.

You can calculate the padding size by subtracting the document size from the record size or, in terms of the `paddingFactor`, by subtracting 1 from the `paddingFactor`:

```
padding size = (paddingFactor - 1) * <document size>.
```

For example, a `paddingFactor` of 1.0 specifies a padding size of 0 whereas a `paddingFactor` of 1.2 specifies a padding size of 0.2 or 20 percent (20%) of the document size.

With the following command, you can use the `paddingFactor` option of the `compact` (page 313) command to set the record size to 1.1 of the document size, or a padding factor of 10 percent (10%):

```
db.runCommand ( { compact: '<collection>', paddingFactor: 1.1 } )
```

`compact` (page 313) compacts existing documents but does not reset `paddingFactor` statistics for the collection. After the `compact` (page 313) MongoDB will use the existing `paddingFactor` when allocating new records for documents in this collection.

paddingBytes New in version 2.2.

Specifying `paddingBytes` can be useful if your documents start small but then increase in size significantly. For example, if your documents are initially 40 bytes long and you grow them by 1KB, using `paddingBytes: 1024` might be reasonable since using `paddingFactor: 4.0` would specify a record size of 160 bytes (4.0 times the initial document size), which would only provide a padding of 120 bytes (i.e. record size of 160 bytes minus the document size).

With the following command, you can use the `paddingBytes` option of the `compact` (page 313) command to set the padding size to 100 bytes on the collection named by `<collection>`:

```
db.runCommand ( { compact: '<collection>', paddingBytes: 100 } )
```

Behaviors The `compact` (page 313) has the behaviors described here.

Blocking In MongoDB 2.2, `compact` (page 313) blocks activities only for its database. Prior to 2.2, the command blocked all activities.

You may view the intermediate progress either by viewing the `mongod` (page 503) log file or by running the `db.currentOp()` (page 103) in another shell instance.

Operation Termination If you terminate the operation with the `db.killOp()` (page 114) method or restart the server before the `compact` (page 313) operation has finished:

- If you have journaling enabled, the data remains valid and usable, regardless of the state of the `compact` (page 313) operation. You may have to manually rebuild the indexes.
- If you do not have journaling enabled and the `mongod` (page 503) or `compact` (page 313) terminates during the operation, it is impossible to guarantee that the data is in a valid state.
- In either case, much of the existing free space in the collection may become un-reusable. In this scenario, you should rerun the compaction to completion to restore the use of this free space.

Disk Space `compact` (page 313) generally uses less disk space than `repairDatabase` (page 319) and is faster. However, the `compact` (page 313) command is still slow and blocks other database use. Only use `compact` (page 313) during scheduled maintenance periods.

`compact` (page 313) requires up to 2 gigabytes of additional disk space while running. Unlike `repairDatabase` (page 319), `compact` (page 313) does *not* free space on the file system.

To see how the storage space changes for the collection, run the `collStats` (page 325) command before and after compaction.

Size and Number of Data Files `compact` (page 313) may increase the total size and number of your data files, especially when run for the first time. However, this will not increase the total collection storage space since storage size is the amount of data allocated within the database files, and not the size/number of the files on the file system.

Replica Sets `compact` (page 313) commands do not replicate to secondaries in a *replica set*:

- Compact each member separately.
- Ideally run `compact` (page 313) on a secondary. See option `force:true` above for information regarding compacting the primary.

- On secondaries, the `compact` (page 313) command forces the secondary to enter `RECOVERING` state. Read operations issued to an instance in the `RECOVERING` state will fail. This prevents clients from reading during the operation. When the operation completes, the secondary returns to `replstate:SECONDARY` state.
- See <http://docs.mongodb.org/manualreference/replica-states/> for more information about replica set member states.

See <http://docs.mongodb.org/manualtutorial/perform-maintenance-on-replica-set-members> for an example replica set maintenance procedure to maximize availability during maintenance operations.

Sharded Clusters `compact` (page 313) is a command issued to a `mongod` (page 503). In a sharded environment, run `compact` (page 313) on each shard separately as a maintenance operation.

You cannot issue `compact` (page 313) against a `mongos` (page 518) instance.

Capped Collections It is not possible to compact *capped collections* because they don't have padding, and documents cannot grow in these collections. However, the documents of a *capped collection* are not subject to fragmentation.

collMod

Definition

collMod

New in version 2.2.

`collMod` (page 316) makes it possible to add flags to a collection to modify the behavior of MongoDB. Flags include `usePowerOf2Sizes` (page 316) and `index` (page 317). The command takes the following prototype form:

```
db.runCommand( { "collMod" : <collection> , "<flag>" : <value> } )
```

In this command substitute `<collection>` with the name of a collection in the current database, and `<flag>` and `<value>` with the flag and value you want to set.

Use the `userFlags` (page 327) field in the `db.collection.stats()` (page 68) output to check enabled collection flags.

Flags

Powers of Two Record Allocation

usePowerOf2Sizes

Changed in version 2.6: `usePowerOf2Sizes` (page 316) became the default allocation strategy for all new collections. Set `newCollectionsUsePowerOf2Sizes` to `false` to select the exact fit allocation strategy for new collections.

The `usePowerOf2Sizes` (page 316) flag changes the method that MongoDB uses to allocate space on disk for documents in this collection. By setting `usePowerOf2Sizes` (page 316), you ensure that MongoDB will allocate space for documents in sizes that are powers of 2 (e.g. 32, 64, 128, 256, 512...16777216.) The smallest allocation for a document is 32 bytes.

With `usePowerOf2Sizes` (page 316) MongoDB will be able to more effectively reuse space.

With `usePowerOf2Sizes` (page 316) MongoDB, allocates records that have power of 2 sizes, until record sizes equal 4 megabytes. For records larger than 4 megabytes with `usePowerOf2Sizes` (page 316) set, `mongod` (page 503) will allocate records in full megabytes by rounding up to the nearest megabyte.

Use `usePowerOf2Sizes` (page 316) for collections where applications insert and delete large numbers of documents to avoid storage fragmentation and ensure that MongoDB will effectively use space on disk.

TTL Collection Expiration Time

index

The `index` (page 317) flag changes the expiration time of a TTL Collection.

Specify the key and new expiration time with a document of the form:

```
{keyPattern: <index_spec>, expireAfterSeconds: <seconds> }
```

In this example, `<index_spec>` is an existing index in the collection and `seconds` is the number of seconds to subtract from the current time.

On success `collMod` (page 316) returns a document with fields `expireAfterSeconds_old` and `expireAfterSeconds_new` set to their respective values.

On failure, `collMod` (page 316) returns a document with no `expireAfterSeconds` field to update if there is no existing `expireAfterSeconds` field or cannot find index { ****key****: 1.0 } for ns ****namespace**** if the specified `keyPattern` does not exist.

Examples

Enable Powers of Two Allocation To enable `usePowerOf2Sizes` (page 316) on the collection named `products`, use the following operation:

```
db.runCommand( {collMod: "products", usePowerOf2Sizes : true } )
```

To disable `usePowerOf2Sizes` (page 316) on the collection `products`, use the following operation:

```
db.runCommand( { collMod: "products", usePowerOf2Sizes: false } )
```

`usePowerOf2Sizes` (page 316) only affects subsequent allocations caused by document insertion or record relocation as a result of document growth, and *does not* affect existing allocations.

Change Expiration Value for Indexes To update the expiration value for a collection named `sessions` indexed on a `lastAccess` field from 30 minutes to 60 minutes, use the following operation:

```
db.runCommand({collMod: "sessions",
               index: {keyPattern: {lastAccess:1},
                     expireAfterSeconds: 3600}})
```

Which will return the document:

```
{ "expireAfterSeconds_old" : 1800, "expireAfterSeconds_new" : 3600, "ok" : 1 }
```

reIndex

reIndex

The `reIndex` (page 317) command drops all indexes on a collection and recreates them. This operation may be expensive for collections that have a large amount of data and/or a large number of indexes. Use the following syntax:

```
{ reIndex: "collection" }
```

Normally, MongoDB compacts indexes during routine updates. For most users, the `reIndex` (page 317) command is unnecessary. However, it may be worth running if the collection size has changed significantly or if the indexes are consuming a disproportionate amount of disk space.

Call `reIndex` (page 317) using the following form:

```
db.collection.reIndex();
```

Note: For replica sets, `reIndex` (page 317) will not propagate from the *primary* to *secondaries*. `reIndex` (page 317) will only affect a single `mongod` (page 503) instance.

Important: `reIndex` (page 317) will rebuild indexes in the *background* if the index was originally specified with this option. However, `reIndex` (page 317) will rebuild the `_id` index in the foreground, which takes the database's write lock.

See

<http://docs.mongodb.org/manualcore/index-creation> for more information on the behavior of indexing operations in MongoDB.

setParameter

setParameter

`setParameter` (page 318) is an administrative command for modifying options normally set on the command line. You must issue the `setParameter` (page 318) command against the *admin database* in the form:

```
{ setParameter: 1, <option>: <value> }
```

Replace the `<option>` with one of the supported `setParameter` (page 318) options:

- journalCommitInterval
- logLevel
- logUserIds
- notablescan
- quiet
- replApplyBatchSize
- replIndexPrefetch
- syncdelay
- traceExceptions
- textSearchEnabled
- sslMode
- clusterAuthMode
- userCacheInvalidationIntervalSecs

getParameter

getParameter

`getParameter` (page 318) is an administrative command for retrieving the value of options normally set on the command line. Issue commands against the *admin database* as follows:

```
{ getParameter: 1, <option>: 1 }
```

The values specified for `getParameter` and `<option>` do not affect the output. The command works with the following options:

- `quiet`
- `notablescan`
- `logLevel`
- `syncdelay`

See also:

`setParameter` (page 318) for more about these parameters.

repairDatabase

- [Definition](#) (page 319)
- [Behavior](#) (page 320)
- [Example](#) (page 320)
- [Using repairDatabase to Reclaim Disk Space](#) (page 320)

Definition

repairDatabase

Checks and repairs errors and inconsistencies in data storage. `repairDatabase` (page 319) is analogous to a `fsck` command for file systems. Run the `repairDatabase` (page 319) command to ensure data integrity after the system experiences an unexpected system restart or crash, if:

1. The `mongod` (page 503) instance is not running with *journaling* enabled.

When using *journaling*, there is almost never any need to run `repairDatabase` (page 319). In the event of an unclean shutdown, the server will be able restore the data files to a pristine state automatically.

2. There are *no* other intact *replica set* members with a complete data set.

Warning: During normal operations, only use the `repairDatabase` (page 319) command and wrappers including `db.repairDatabase()` (page 117) in the `mongo` (page 527) shell and `mongod --repair`, to compact database files and/or reclaim disk space. Be aware that these operations remove and do not save any corrupt data during the repair process.

If you are trying to repair a *replica set* member, and you have access to an intact copy of your data (e.g. a recent backup or an intact member of the *replica set*), you should restore from that intact copy, and **not** use `repairDatabase` (page 319).

`repairDatabase` (page 319) takes the following form:

```
{ repairDatabase: 1 }
```

`repairDatabase` (page 319) has the following fields:

field boolean preserveClonedFilesOnFailure When `true`, `repairDatabase` will not delete temporary files in the backup directory on error, and all new files are created with the “backup” instead of “_tmp” directory prefix. By default `repairDatabase` does not delete temporary files, and uses the “_tmp” naming prefix for new files.

field boolean backupOriginalFiles When `true`, `repairDatabase` moves old database files to the backup directory instead of deleting them before moving new files into place. New files are

created with the “backup” instead of “_tmp” directory prefix. By default, `repairDatabase` leaves temporary files unchanged, and uses the “_tmp” naming prefix for new files.

You can explicitly set the options as follows:

```
{ repairDatabase: 1,
  preserveClonedFilesOnFailure: <boolean>,
  backupOriginalFiles: <boolean> }
```

Warning: This command obtains a global write lock and will block other operations until it has completed.

Note: `repairDatabase` (page 319) requires free disk space equal to the size of your current data set plus 2 gigabytes. If the volume that holds `dbpath` lacks sufficient space, you can mount a separate volume and use that for the repair. When mounting a separate volume for `repairDatabase` (page 319) you must run `repairDatabase` (page 319) from the command line and use the `--repairpath` switch to specify the folder in which to store temporary repair files.

See `mongod --repair` and `mongodump --repair` for information on these related options.

Behavior The `repairDatabase` (page 319) command compacts all collections in the database. It is identical to running the `compact` (page 313) command on each collection individually.

`repairDatabase` (page 319) reduces the total size of the data files on disk. It also recreates all indexes in the database.

The time requirement for `repairDatabase` (page 319) depends on the size of the data set.

You may invoke `repairDatabase` (page 319) from multiple contexts:

- Use the `mongo` (page 527) shell to run the command, as above.
- Use the `db.repairDatabase()` (page 117) in the `mongo` (page 527) shell.
- Run `mongod` (page 503) directly from your system’s shell. Make sure that `mongod` (page 503) isn’t already running, and that you invoke `mongod` (page 503) as a user that has access to MongoDB’s data files. Run as:

```
mongod --repair
```

To add a repair path:

```
mongod --repair --repairpath /opt/vol2/data
```

Note: `mongod --repair` will fail if your database is not a master or primary. In most cases, you should recover a corrupt secondary using the data from an existing intact node. To run repair on a secondary/slave restart the instance in standalone mode without the `--replSet` or `--slave` options.

Example

```
{ repairDatabase: 1 }
```

Using `repairDatabase` to Reclaim Disk Space You should not use `repairDatabase` (page 319) for data recovery unless you have no other option.

However, if you trust that there is no corruption and you have enough free space, then `repairDatabase` (page 319) is the appropriate and the only way to reclaim disk space.

touch
touch

New in version 2.2.

The `touch` (page 321) command loads data from the data storage layer into memory. `touch` (page 321) can load the data (i.e. documents) indexes or both documents and indexes. Use this command to ensure that a collection, and/or its indexes, are in memory before another operation. By loading the collection or indexes into memory, `mongod` (page 503) will ideally be able to perform subsequent operations more efficiently. The `touch` (page 321) command has the following prototypical form:

```
{ touch: [collection], data: [boolean], index: [boolean] }
```

By default, `data` and `index` are false, and `touch` (page 321) will perform no operation. For example, to load both the data and the index for a collection named `records`, you would use the following command in the `mongo` (page 527) shell:

```
db.runCommand({ touch: "records", data: true, index: true })
```

`touch` (page 321) will not block read and write operations on a `mongod` (page 503), and can run on *secondary* members of replica sets.

Note: Using `touch` (page 321) to control or tweak what a `mongod` (page 503) stores in memory may displace other records data in memory and hinder performance. Use with caution in production systems.

Warning: If you run `touch` (page 321) on a secondary, the secondary will enter a `RECOVERING` state to prevent clients from sending read operations during the `touch` (page 321) operation. When `touch` (page 321) finishes the secondary will automatically return to `SECONDARY` state. See `state` (page 274) for more information on replica set member states.

shutdown
shutdown

The `shutdown` (page 321) command cleans up all database resources and then terminates the process. You must issue the `shutdown` (page 321) command against the *admin database* in the form:

```
{ shutdown: 1 }
```

Note: Run the `shutdown` (page 321) against the *admin database*. When using `shutdown` (page 321), the connection must originate from `localhost` or use an authenticated connection.

If the node you're trying to shut down is a `replica set` primary, then the command will succeed only if there exists a secondary node whose oplog data is within 10 seconds of the primary. You can override this protection using the `force` option:

```
{ shutdown: 1, force: true }
```

Alternatively, the `shutdown` (page 321) command also supports a `timeoutSecs` argument which allows you to specify a number of seconds to wait for other members of the replica set to catch up:

```
{ shutdown: 1, timeoutSecs: 60 }
```

The equivalent `mongo` (page 527) shell helper syntax looks like this:

```
db.shutdownServer({timeoutSecs: 60});
```

logRotate

logRotate

The `logRotate` (page 322) command is an administrative command that allows you to rotate the MongoDB logs to prevent a single logfile from consuming too much disk space. You must issue the `logRotate` (page 322) command against the *admin database* in the form:

```
{ logRotate: 1 }
```

Note: Your `mongod` (page 503) instance needs to be running with the `--logpath [file]` option.

You may also rotate the logs by sending a `SIGUSR1` signal to the `mongod` (page 503) process. If your `mongod` (page 503) has a process ID of 2200, here's how to send the signal on Linux:

```
kill -SIGUSR1 2200
```

`logRotate` (page 322) renames the existing log file by appending the current timestamp to the filename. The appended timestamp has the following form:

```
<YYYY>--<mm>--<DD>T<HH>--<MM>--<SS>
```

Then `logRotate` (page 322) creates a new log file with the same name as originally specified by the `systemLog.path` setting to `mongod` (page 503) or `mongos` (page 518).

Note: New in version 2.0.3: The `logRotate` (page 322) command is available to `mongod` (page 503) instances running on Windows systems with MongoDB release 2.0.3 and higher.

Diagnostic Commands

Diagnostic Commands

Name	Description
<code>listDatabases</code> (page 323)	Returns a document that lists all databases and returns basic database statistics.
<code>dbHash</code> (page 324)	Internal command to support sharding.
<code>driverOIDTest</code> (page 324)	Internal command that converts an <code>ObjectId</code> to a string to support tests.
<code>listCommands</code> (page 324)	Lists all database commands provided by the current <code>mongod</code> (page 503) instance.
<code>availableQueryOptions</code> (page 324)	Internal command that reports on the capabilities of the current MongoDB instance.
<code>buildInfo</code> (page 324)	Displays statistics about the MongoDB build.
<code>collStats</code> (page 325)	Reports storage utilization statics for a specified collection.
<code>connPoolStats</code> (page 328)	Reports statistics on the outgoing connections from this MongoDB instance to other MongoDB instances in the deployment.
<code>shardConnPoolStats</code> (page 330)	Reports statistics on a <code>mongos</code> (page 518)'s connection pool for client operations against shards.
<code>dbStats</code> (page 331)	Reports storage utilization statistics for the specified database.
<code>cursorInfo</code> (page 333)	Deprecated. Reports statistics on active cursors.
<code>dataSize</code> (page 333)	Returns the data size for a range of data. For internal use.
<code>diagLogging</code> (page 333)	Provides a diagnostic logging. For internal use.
<code>getCmdLineOpts</code> (page 333)	Returns a document with the run-time arguments to the MongoDB instance and their parsed options.
<code>netstat</code> (page 334)	Internal command that reports on intra-deployment connectivity. Only available for <code>mongos</code> (page 518) instances.
<code>ping</code> (page 334)	Internal command that tests intra-deployment connectivity.
<code>profile</code> (page 334)	Interface for the <i>database profiler</i> .
<code>validate</code> (page 335)	Internal command that scans for a collection's data and indexes for correctness.
<code>top</code> (page 337)	Returns raw usage statistics for each database in the <code>mongod</code> (page 503) instance.
<code>indexStats</code> (page 339)	Experimental command that collects and aggregates statistics on all indexes.
<code>whatsmyuri</code> (page 344)	Internal command that returns information on the current client.
<code>getLog</code> (page 344)	Returns recent log messages.
<code>hostInfo</code> (page 344)	Returns data that reflects the underlying host system.
<code>serverStatus</code> (page 347)	Returns a collection metrics on instance-wide resource utilization and status.
<code>features</code> (page 364)	Reports on features available in the current MongoDB instance.
<code>isSelf</code>	Internal command to support testing.

listDatabases

listDatabases

The `listDatabases` (page 323) command provides a list of existing databases along with basic statistics about them:

```
{ listDatabases: 1 }
```

The value (e.g. 1) does not affect the output of the command. `listDatabases` (page 323) returns a document for each database. Each document contains a `name` field with the database name, a `sizeOnDisk` field with the total size of the database file on disk in bytes, and an `empty` field specifying whether the database has any data.

Example

The following operation returns a list of all databases:

```
db.runCommand( { listDatabases: 1 } )
```

See also:

<http://docs.mongodb.org/manual/tutorial/use-database-commands>.

dbHash

dbHash

`dbHash` (page 324) is a command that supports *config servers* and is not part of the stable client facing API.

driverOIDTest

driverOIDTest

`driverOIDTest` (page 324) is an internal command.

listCommands

listCommands

The `listCommands` (page 324) command generates a list of all database commands implemented for the current `mongod` (page 503) instance.

availableQueryOptions

availableQueryOptions

`availableQueryOptions` (page 324) is an internal command that is only available on `mongos` (page 518) instances.

buildInfo

buildInfo

The `buildInfo` (page 324) command is an administrative command which returns a build summary for the current `mongod` (page 503). `buildInfo` (page 324) has the following prototype form:

```
{ buildInfo: 1 }
```

In the `mongo` (page 527) shell, call `buildInfo` (page 324) in the following form:

```
db.runCommand( { buildInfo: 1 } )
```

Example

The output document of `buildInfo` (page 324) has the following form:

```
{
  "version" : "<string>",
  "gitVersion" : "<string>",
  "sysInfo" : "<string>",
  "loaderFlags" : "<string>",
  "compilerFlags" : "<string>",
  "allocator" : "<string>",
  "versionArray" : [ <num>, <num>, <...> ],
  "javascriptEngine" : "<string>",
  "bits" : <num>,
  "debug" : <boolean>,
  "maxBsonObjectSize" : <num>,
}
```



```
"ok" : <num>
}
```

Consider the following documentation of the output of `buildInfo` (page 324):

buildInfo

The document returned by the `buildInfo` (page 324) command.

buildInfo.gitVersion

The commit identifier that identifies the state of the code used to build the `mongod` (page 503).

buildInfo.sysInfo

A string that holds information about the operating system, hostname, kernel, date, and Boost version used to compile the `mongod` (page 503).

buildInfo.loaderFlags

The flags passed to the loader that loads the `mongod` (page 503).

buildInfo.compilerFlags

The flags passed to the compiler that builds the `mongod` (page 503) binary.

buildInfo allocator

Changed in version 2.2.

The memory allocator that `mongod` (page 503) uses. By default this is `tcmalloc` after version 2.2, and `system` before 2.2.

buildInfo.versionArray

An array that conveys version information about the `mongod` (page 503) instance. See `version` for a more readable version of this string.

buildInfo.javascriptEngine

Changed in version 2.4.

A string that reports the JavaScript engine used in the `mongod` (page 503) instance. By default, this is `V8` after version 2.4, and `SpiderMonkey` before 2.4.

buildInfo.bits

A number that reflects the target processor architecture of the `mongod` (page 503) binary.

buildInfo.debug

A boolean. `true` when built with debugging options.

buildInfo.maxBsonObjectSize

A number that reports the `Maximum BSON Document Size` (page 604).

collStats

Definition

collStats

The `collStats` (page 325) command returns a variety of storage statistics for a given collection. Use the following syntax:

```
{ collStats: "collection" , scale : 1024 }
```

Specify the `collection` you want statistics for, and use the `scale` argument to scale the output. The above example will display values in kilobytes.

Examine the following example output, which uses the `db.collection.stats()` (page 68) helper in the `mongo` (page 527) shell.

```
> db.users.stats()
{
  "ns" : "app.users",           // namespace
  "count" : 9,                  // number of documents
  "size" : 432,                 // collection size in bytes
  "avgObjSize" : 48,           // average object size in bytes
  "storageSize" : 3840,         // (pre)allocated space for the collection in bytes
  "numExtents" : 1,             // number of extents (contiguously allocated chunks of c
  "nindexes" : 2,               // number of indexes
  "lastExtentSize" : 3840,      // size of the most recently created extent in bytes
  "paddingFactor" : 1,          // padding can speed up updates if documents grow
  "flags" : 1,
  "totalIndexSize" : 16384,     // total index size in bytes
  "indexSizes" : {             // size of specific indexes in bytes
    "_id_" : 8192,
    "username" : 8192
  },
  "ok" : 1
}
```

Note: The scale factor rounds values to whole numbers. This can produce unpredictable and unexpected results in some situations.

Output

`collStats.ns`

The namespace of the current collection, which follows the format `[database].[collection]`.

`collStats.count`

The number of objects or documents in this collection.

`collStats.size`

The total size of all records in a collection. This value does not include the record header, which is 16 bytes per record, but *does* include the record's *padding*. Additionally *size* (page 326) does not include the size of any indexes associated with the collection, which the `totalIndexSize` (page 327) field reports.

The `scale` argument affects this value.

`collStats.avgObjSize`

The average size of an object in the collection. The `scale` argument affects this value.

`collStats.storageSize`

The total amount of storage allocated to this collection for *document* storage. The `scale` argument affects this value. The `storageSize` (page 326) does not decrease as you remove or shrink documents.

`collStats.numExtents`

The total number of contiguously allocated data file regions.

`collStats.nindexes`

The number of indexes on the collection. All collections have at least one index on the `_id` field.

Changed in version 2.2: Before 2.2, capped collections did not necessarily have an index on the `_id` field, and some capped collections created with pre-2.2 versions of `mongod` (page 503) may not have an `_id` index.

`collStats.lastExtentSize`

The size of the last extent allocated. The `scale` argument affects this value.

collStats.paddingFactor

The amount of space added to the end of each document at insert time. The document padding provides a small amount of extra space on disk to allow a document to grow slightly without needing to move the document. `mongod` (page 503) automatically calculates this padding factor

collStats.flags

Changed in version 2.2: Removed in version 2.2 and replaced with the `userFlags` (page 327) and `systemFlags` (page 327) fields.

Indicates the number of flags on the current collection. In version 2.0, the only flag notes the existence of an *index* on the `_id` field.

collStats.systemFlags

New in version 2.2.

Reports the flags on this collection that reflect internal server options. Typically this value is 1 and reflects the existence of an *index* on the `_id` field.

collStats.userFlags

New in version 2.2.

Reports the flags on this collection set by the user. In version 2.2 the only user flag is `usePowerOf2Sizes` (page 316). If `usePowerOf2Sizes` (page 316) is enabled, `userFlags` (page 327) will be set to 1, otherwise `userFlags` (page 327) will be 0.

See the `collMod` (page 316) command for more information on setting user flags and `usePowerOf2Sizes` (page 316).

collStats.totalIndexSize

The total size of all indexes. The `scale` argument affects this value.

collStats.indexSizes

This field specifies the key and size of every existing index on the collection. The `scale` argument affects this value.

Example The following is an example of `db.collection.stats()` (page 68) and `collStats` (page 325) output:

```
{
  "ns" : "<database>.<collection>",
  "count" : <number>,
  "size" : <number>,
  "avgObjSize" : <number>,
  "storageSize" : <number>,
  "numExtents" : <number>,
  "nindexes" : <number>,
  "lastExtentSize" : <number>,
  "paddingFactor" : <number>,
  "systemFlags" : <bit>,
  "userFlags" : <bit>,
  "totalIndexSize" : <number>,
  "indexSizes" : {
    "_id_" : <number>,
    "a_1" : <number>
  },
  "ok" : 1
}
```

connPoolStats

Definition**connPoolStats**

Note: `connPoolStats` (page 328) only returns meaningful results for `mongos` (page 518) instances and for `mongod` (page 503) instances in sharded clusters.

The command `connPoolStats` (page 328) returns information regarding the number of open connections to the current database instance, including client connections and server-to-server connections for replication and clustering. The command takes the following form:

```
{ connPoolStats: 1 }
```

The value of the argument (i.e. `1`) does not affect the output of the command.

Note: `connPoolStats` (page 328) only returns meaningful results for `mongos` (page 518) instances and for `mongod` (page 503) instances in sharded clusters.

Output**connPoolStats.hosts**

The sub-documents of the `hosts` (page 328) *document* report connections between the `mongos` (page 518) or `mongod` (page 503) instance and each component `mongod` (page 503) of the *sharded cluster*.

connPoolStats.hosts.[host].available

`available` (page 328) reports the total number of connections that the `mongos` (page 518) or `mongod` (page 503) could use to connect to this `mongod` (page 503).

connPoolStats.hosts.[host].created

`created` (page 328) reports the number of connections that this `mongos` (page 518) or `mongod` (page 503) has ever created for this host.

connPoolStats.replicaSets

`replicaSets` (page 328) is a *document* that contains *replica set* information for the *sharded cluster*.

connPoolStats.replicaSets.shard

The `shard` (page 328) *document* reports on each *shard* within the *sharded cluster*

connPoolStats.replicaSets.[shard].host

The `host` (page 328) field holds an array of *document* that reports on each host within the *shard* in the *replica set*.

These values derive from the *replica set status* (page 273) values.

connPoolStats.replicaSets.[shard].host[n].addr

`addr` (page 328) reports the address for the host in the *sharded cluster* in the format of “[hostname]:[port]”.

connPoolStats.replicaSets.[shard].host[n].ok

`ok` (page 328) reports false when:

- the `mongos` (page 518) or `mongod` (page 503) cannot connect to instance.
- the `mongos` (page 518) or `mongod` (page 503) received a connection exception or error.

This field is for internal use.

connPoolStats.replicaSets.[shard].host[n].ismaster

`ismaster` (page 328) reports true if this `host` (page 328) is the *primary* member of the *replica set*.

connPoolStats.replicaSets.[shard].host[n].hidden

`hidden` (page 328) reports true if this `host` (page 328) is a *hidden member* of the *replica set*.

`connPoolStats.replicaSets.[shard].host[n].secondary`
`secondary` (page 329) reports true if this `host` (page 328) is a *secondary* member of the *replica set*.

`connPoolStats.replicaSets.[shard].host[n].pingTimeMillis`
`pingTimeMillis` (page 329) reports the ping time in milliseconds from the `mongos` (page 518) or `mongod` (page 503) to this `host` (page 328).

`connPoolStats.replicaSets.[shard].host[n].tags`
 New in version 2.2.
`tags` (page 329) reports the tags, if this member of the set has tags configured.

`connPoolStats.replicaSets.[shard].master`
`master` (page 329) reports the ordinal identifier of the host in the `host` (page 328) array that is the *primary* of the *replica set*.

`connPoolStats.replicaSets.[shard].nextSlave`
 Deprecated since version 2.2.
`nextSlave` (page 329) reports the *secondary* member that the `mongos` (page 518) will use to service the next request for this *replica set*.

`connPoolStats.createdByType`
`createdByType` (page 329) *document* reports the number of each type of connection that `mongos` (page 518) or `mongod` (page 503) has created in all connection pools.

`mongos` (page 518) connect to `mongod` (page 503) instances using one of three types of connections. The following sub-document reports the total number of connections by type.

`connPoolStats.createdByType.master`
`master` (page 329) reports the total number of connections to the *primary* member in each *cluster*.

`connPoolStats.createdByType.set`
`set` (page 329) reports the total number of connections to a *replica set* member.

`connPoolStats.createdByType.sync`
`sync` (page 329) reports the total number of *config database* connections.

`connPoolStats.totalAvailable`
`totalAvailable` (page 329) reports the running total of connections from the `mongos` (page 518) or `mongod` (page 503) to all `mongod` (page 503) instances in the *sharded cluster* available for use.

`connPoolStats.totalCreated`
`totalCreated` (page 329) reports the total number of connections ever created from the `mongos` (page 518) or `mongod` (page 503) to all `mongod` (page 503) instances in the *sharded cluster*.

`connPoolStats.numDBClientConnection`
`numDBClientConnection` (page 329) reports the total number of connections from the `mongos` (page 518) or `mongod` (page 503) to all of the `mongod` (page 503) instances in the *sharded cluster*.

`connPoolStats.numAScopedConnection`
`numAScopedConnection` (page 329) reports the number of exception safe connections created from `mongos` (page 518) or `mongod` (page 503) to all `mongod` (page 503) in the *sharded cluster*. The `mongos` (page 518) or `mongod` (page 503) releases these connections after receiving a socket exception from the `mongod` (page 503).

shardConnPoolStats

Definition**shardConnPoolStats**

Returns information on the pooled and cached connections in the sharded connection pool. The command also returns information on the per-thread connection cache in the connection pool.

The `shardConnPoolStats` (page 330) command uses the following syntax:

```
{ shardConnPoolStats: 1 }
```

The sharded connection pool is specific to connections between members in a sharded cluster. The `mongos` (page 518) instances in a cluster use the connection pool to execute client reads and writes. The `mongod` (page 503) instances in a cluster use the pool when issuing `mapReduce` (page 208) to query temporary collections on other shards.

When the cluster requires a connection, MongoDB pulls a connection from the sharded connection pool into the per-thread connection cache. MongoDB returns the connection to the connection pool after every operation.

Output**shardConnPoolStats.hosts**

Displays connection status for each *config server*, *replica set*, and *standalone instance* in the cluster.

shardConnPoolStats.hosts.<host>.available

The number of connections available for this host to connect to the `mongos` (page 518).

shardConnPoolStats.hosts.<host>.created

The number of connections the host has ever created to connect to the `mongos` (page 518).

shardConnPoolStats.replicaSets

Displays information specific to replica sets.

shardConnPoolStats.replicaSets.<name>.host

Holds an array of documents that report on each replica set member. These values derive from the *replica set status* (page 273) values.

shardConnPoolStats.replicaSets.<name>.host[n].addr

The host address in the format `[hostname]:[port]`.

shardConnPoolStats.replicaSets.<name>.host[n].ok

This field is for internal use. Reports `false` when the `mongos` (page 518) either cannot connect to instance or received a connection exception or error.

shardConnPoolStats.replicaSets.<name>.host[n].ismaster

The host is the replica set's *primary* if this is set to `true`.

shardConnPoolStats.replicaSets.<name>.host[n].hidden

The host is a *hidden member* of the replica set if this is set to `true`.

shardConnPoolStats.replicaSets.<name>.host[n].secondary

The host is a *hidden member* of the replica set if this is set to `true`.

The host is a *secondary* member of the replica set if this is set to `true`.

shardConnPoolStats.replicaSets.<name>.host[n].pingTimeMillis

The latency, in milliseconds, from the `mongos` (page 518) to this member.

shardConnPoolStats.replicaSets.<name>.host[n].tags

The member has tags configured.

shardConnPoolStats.createdByType

The number connections in the cluster's connection pool.

`shardConnPoolStats.createdByType.master`

The number of connections to a shard.

`shardConnPoolStats.createdByType.set`

The number of connections to a replica set.

`shardConnPoolStats.createdByType.sync`

The number of connections to the config database.

`shardConnPoolStats.totalAvailable`

The number of connections available from the [mongos](#) (page 518) to the config servers, replica sets, and standalone [mongod](#) (page 503) instances in the cluster.

`shardConnPoolStats.totalCreated`

The number of connections the [mongos](#) (page 518) has ever created to other members of the cluster.

`shardConnPoolStats.threads`

Displays information on the per-thread connection cache.

`shardConnPoolStats.threads.hosts`

Displays each incoming client connection. For a [mongos](#) (page 518), this array field displays one document per incoming client thread. For a [mongod](#) (page 503), the array displays one entry per incoming sharded [mapReduce](#) (page 208) client thread.

`shardConnPoolStats.threads.hosts.host`

The host using the connection. The host can be a *config server*, *replica set*, or *standalone instance*.

`shardConnPoolStats.threads.hosts.created`

The number of times the host pulled a connection from the pool.

`shardConnPoolStats.threads.hosts.avail`

The thread's availability.

`shardConnPoolStats.threads.seenNS`

The namespaces used on this connection thus far.

dbStats

Definition

dbStats

The `dbStats` (page 331) command returns storage statistics for a given database. The command takes the following syntax:

```
{ dbStats: 1, scale: 1 }
```

The values of the options above do not affect the output of the command. The `scale` option allows you to specify how to scale byte values. For example, a `scale` value of 1024 will display the results in kilobytes rather than in bytes:

```
{ dbStats: 1, scale: 1024 }
```

Note: Because scaling rounds values to whole numbers, scaling may return unlikely or unexpected results.

The time required to run the command depends on the total size of the database. Because the command must touch all data files, the command may take several seconds to run.

In the [mongo](#) (page 527) shell, the `db.stats()` (page 119) function provides a wrapper around `dbStats` (page 331).

Output**dbStats.db**

Contains the name of the database.

dbStats.collections

Contains a count of the number of collections in that database.

dbStats.objects

Contains a count of the number of objects (i.e. *documents*) in the database across all collections.

dbStats.avgObjSize

The average size of each document in bytes. This is the *dataSize* (page 332) divided by the number of documents.

dbStats.dataSize

The total size in bytes of the data held in this database including the *padding factor*. The *scale* argument affects this value. The *dataSize* (page 332) will not decrease when *documents* shrink, but will decrease when you remove documents.

dbStats.storageSize

The total amount of space in bytes allocated to collections in this database for *document* storage. The *scale* argument affects this value. The *storageSize* (page 332) does not decrease as you remove or shrink documents.

dbStats.numExtents

Contains a count of the number of extents in the database across all collections.

dbStats.indexes

Contains a count of the total number of indexes across all collections in the database.

dbStats.indexSize

The total size in bytes of all indexes created on this database. The *scale* arguments affects this value.

dbStats.fileSize

The total size in bytes of the data files that hold the database. This value includes preallocated space and the *padding factor*. The value of *fileSize* (page 332) only reflects the size of the data files for the database and not the namespace file.

The *scale* argument affects this value.

dbStats.nsSizeMB

The total size of the *namespace* files (i.e. that end with *.ns*) for this database. You cannot change the size of the namespace file after creating a database, but you can change the default size for all new namespace files with the *nsSize* runtime option.

See also:

The *nsSize* option, and *Maximum Namespace File Size* (page 604)

dbStats.dataFileVersion

New in version 2.4.

Document that contains information about the on-disk format of the data files for the database.

dbStats.dataFileVersion.major

New in version 2.4.

The major version number for the on-disk format of the data files for the database.

dbStats.dataFileVersion.minor

New in version 2.4.

The minor version number for the on-disk format of the data files for the database.

cursorInfo Deprecated since version 2.6: Use the [serverStatus](#) (page 347) command to return the [serverStatus.metrics.cursor](#) (page 363) information instead.

cursorInfo

The `cursorInfo` (page 333) command returns information about current cursor allotment and use. Use the following form:

```
{ cursorInfo: 1 }
```

The value (e.g. 1 above) does not affect the output of the command.

`cursorInfo` (page 333) returns the total number of open cursors (`totalOpen`), the size of client cursors in current use (`clientCursors_size`), and the number of timed out cursors since the last server restart (`timedOut`).

dataSize

dataSize

The `dataSize` (page 333) command returns the data size for a set of data within a certain range:

```
{ dataSize: "database.collection", keyPattern: { field: 1 }, min: { field: 10 }, max: { field: 1
```

This will return a document that contains the size of all matching documents. Replace `database.collection` value with `database` and `collection` from your deployment. The `keyPattern`, `min`, and `max` parameters are options.

The amount of time required to return `dataSize` (page 333) depends on the amount of data in the collection.

diagLogging

diagLogging

`diagLogging` (page 333) is a command that captures additional data for diagnostic purposes and is not part of the stable client facing API.

diaglogging obtains a write lock on the affected database and will block other operations until it completes.

getCmdLineOpts

```
getCmdLineOpts
```

The `getCmdLineOpts` (page 333) command returns a document containing command line options used to start the given `mongod` (page 503) or `mongos` (page 518):

```
{ getCmdLineOpts: 1 }
```

This command returns a document with two fields, `argv` and `parsed`. The `argv` field contains an array with each item from the command string used to invoke `mongod` (page 503) or `mongos` (page 518). The document in the `parsed` field includes all runtime options, including those parsed from the command line and those specified in the configuration file, if specified.

Consider the following example output of `getCmdLineOpts` (page 333):

```
{
    "argv" : [
        "/usr/bin/mongod",
        "--config",
        "/etc/mongodb.conf",
        "--fork"
    ],
    "parsed" : {
```

```
    "bind_ip" : "127.0.0.1",
    "config"  : "/etc/mongodb/mongodb.conf",
    "dbpath"  : "/srv/mongodb",
    "fork"    : true,
    "logappend" : "true",
    "logpath"  : "/var/log/mongodb/mongod.log",
    "quiet"    : "true"
  },
  "ok" : 1
}
```

netstat

netstat

`netstat` (page 334) is an internal command that is only available on `mongos` (page 518) instances.

ping

ping

The `ping` (page 334) command is a no-op used to test whether a server is responding to commands. This command will return immediately even if the server is write-locked:

```
{ ping: 1 }
```

The value (e.g. 1 above) does not impact the behavior of the command.

profile

profile

Use the `profile` (page 334) command to enable, disable, or change the query profiling level. This allows administrators to capture data regarding performance. The database profiling system can impact performance and can allow the server to write the contents of queries to the log. Your deployment should carefully consider the security implications of this. Consider the following prototype syntax:

```
{ profile: <level> }
```

The following profiling levels are available:

Level	Setting
-1	No change. Returns the current profile level.
0	Off. No profiling.
1	On. Only includes slow operations.
2	On. Includes all operations.

You may optionally set a threshold in milliseconds for profiling using the `slowms` option, as follows:

```
{ profile: 1, slowms: 200 }
```

`mongod` (page 503) writes the output of the database profiler to the `system.profile` collection.

`mongod` (page 503) records queries that take longer than the `slowOpThresholdMs` to the server log even when the database profiler is not active.

See also:

Additional documentation regarding *Database Profiling*.

See also:

“`db.getProfilingStatus()` (page 112)” and “`db.setProfilingLevel()` (page 119)” provide wrappers around this functionality in the `mongo` (page 527) shell.

The database cannot be locked with `db.fsyncLock()` (page 109) while profiling is enabled. You must disable profiling before locking the database with `db.fsyncLock()` (page 109). Disable profiling using `db.setProfilingLevel()` (page 119) as follows in the `mongo` (page 527) shell:

```
db.setProfilingLevel(0)
```

Note: This command obtains a write lock on the affected database and will block other operations until it has completed. However, the write lock is only held while enabling or disabling the profiler. This is typically a short operation.

validate

Definition

validate

The `validate` (page 335) command checks the structures within a namespace for correctness by scanning the collection's data and indexes. The command returns information regarding the on-disk representation of the collection.

The `validate` command can be slow, particularly on larger data sets.

The following example validates the contents of the collection named `users`.

```
{ validate: "users" }
```

You may also specify one of the following options:

- **full:** `true` provides a more thorough scan of the data.
- **scandata:** `false` skips the scan of the base collection without skipping the scan of the index.

The `mongo` (page 527) shell also provides a wrapper:

```
db.collection.validate();
```

Use one of the following forms to perform the full collection validation:

```
db.collection.validate(true)
db.runCommand( { validate: "collection", full: true } )
```

Warning: This command is resource intensive and may have an impact on the performance of your MongoDB instance.

Output

validate.ns

The full namespace name of the collection. Namespaces include the database name and the collection name in the form `database.collection`.

validate.firstExtent

The disk location of the first extent in the collection. The value of this field also includes the namespace.

validate.lastExtent

The disk location of the last extent in the collection. The value of this field also includes the namespace.

validate.extentCount

The number of extents in the collection.

validate.extents

`validate` (page 335) returns one instance of this document for every extent in the collection. This sub-document is only returned when you specify the `full` option to the command.

validate.extents.loc

The disk location for the beginning of this extent.

validate.extents.xnext

The disk location for the extent following this one. “null” if this is the end of the linked list of extents.

validate.extents.xprev

The disk location for the extent preceding this one. “null” if this is the head of the linked list of extents.

validate.extents.nsdiag

The namespace this extent belongs to (should be the same as the namespace shown at the beginning of the `validate` listing).

validate.extents.size

The number of bytes in this extent.

validate.extents.firstRecord

The disk location of the first record in this extent.

validate.extents.lastRecord

The disk location of the last record in this extent.

validate.datasize

The number of bytes in all data records. This value does not include deleted records, nor does it include extent headers, nor record headers, nor space in a file unallocated to any extent. `datasize` (page 336) includes record *padding*.

validate.nrecords

The number of *documents* in the collection.

validate.lastExtentSize

The size of the last new extent created in this collection. This value determines the size of the *next* extent created.

validate.padding

A floating point value between 1 and 2.

When MongoDB creates a new record it uses the *padding factor* to determine how much additional space to add to the record. The padding factor is automatically adjusted by mongo when it notices that update operations are triggering record moves.

validate.firstExtentDetails

The size of the first extent created in this collection. This data is similar to the data provided by the `extents` (page 335) sub-document; however, the data reflects only the first extent in the collection and is always returned.

validate.firstExtentDetails.loc

The disk location for the beginning of this extent.

validate.firstExtentDetails.xnext

The disk location for the extent following this one. “null” if this is the end of the linked list of extents, which should only be the case if there is only one extent.

validate.firstExtentDetails.xprev

The disk location for the extent preceding this one. This should always be “null.”

validate.firstExtentDetails.nsdiag

The namespace this extent belongs to (should be the same as the namespace shown at the beginning of the `validate` listing).

`validate.firstExtentDetails.size`

The number of bytes in this extent.

`validate.firstExtentDetails.firstRecord`

The disk location of the first record in this extent.

`validate.firstExtentDetails.lastRecord`

The disk location of the last record in this extent.

`validate.objectsFound`

The number of records actually encountered in a scan of the collection. This field should have the same value as the `nrecords` (page 336) field.

`validate.invalidObjects`

The number of records containing BSON documents that do not pass a validation check.

Note: This field is only included in the validation output when you specify the `full` option.

`validate.bytesWithHeaders`

This is similar to `datasize`, except that `bytesWithHeaders` (page 337) includes the record headers. In version 2.0, record headers are 16 bytes per document.

Note: This field is only included in the validation output when you specify the `full` option.

`validate.bytesWithoutHeaders`

`bytesWithoutHeaders` (page 337) returns data collected from a scan of all records. The value should be the same as `datasize` (page 336).

Note: This field is only included in the validation output when you specify the `full` option.

`validate.deletedCount`

The number of deleted or “free” records in the collection.

`validate.deletedSize`

The size of all deleted or “free” records in the collection.

`validate.nIndexes`

The number of indexes on the data in the collection.

`validate.keysPerIndex`

A document containing a field for each index, named after the index’s name, that contains the number of keys, or documents referenced, included in the index.

`validate.valid`

Boolean. `true`, unless `validate` (page 335) determines that an aspect of the collection is not valid. When `false`, see the `errors` (page 337) field for more information.

`validate.errors`

Typically empty; however, if the collection is not valid (i.e `valid` (page 337) is `false`), this field will contain a message describing the validation error.

`validate.ok`

Set to 1 when the command succeeds. If the command fails the `ok` (page 337) field has a value of 0.

top

top

`top` (page 337) is an administrative command that returns usage statistics for each collection. `top` (page 337) provides amount of time, in microseconds, used and a count of operations for the following event types:

- total
- readLock
- writeLock
- queries
- getmore
- insert
- update
- remove
- commands

Issue the `top` (page 337) command against the *admin database* in the form:

```
{ top: 1 }
```

Example At the `mongo` (page 527) shell prompt, use `top` (page 337) with the following evocation:

```
db.adminCommand("top")
```

Alternately you can use `top` (page 337) as follows:

```
use admin
db.runCommand( { top: 1 } )
```

The output of the `top` command would resemble the following output:

```
{
  "totals" : {
    "records.users" : {
      "total" : {
        "time" : 305277,
        "count" : 2825
      },
      "readLock" : {
        "time" : 305264,
        "count" : 2824
      },
      "writeLock" : {
        "time" : 13,
        "count" : 1
      },
      "queries" : {
        "time" : 305264,
        "count" : 2824
      },
      "getmore" : {
        "time" : 0,
        "count" : 0
      },
      "insert" : {
        "time" : 0,
        "count" : 0
      },
      "update" : {
        "time" : 0,
```

```

        "count" : 0
      },
      "remove" : {
        "time" : 0,
        "count" : 0
      },
      "commands" : {
        "time" : 0,
        "count" : 0
      }
    }
  }
}

```

indexStats

Definition

indexStats

The `indexStats` (page 339) command aggregates statistics for the B-tree data structure that stores data for a MongoDB index.

Warning: This command is not intended for production deployments.

The command can be run *only* on a `mongod` (page 503) instance that uses the `--enableExperimentalIndexStatsCmd` option.

To aggregate statistics, issue the command like so:

```
db.runCommand( { indexStats: "<collection>", index: "<index name>" } )
```

Output The `db.collection.indexStats()` (page 28) method and equivalent `indexStats` (page 339) command aggregate statistics for the B-tree data structure that stores data for a MongoDB index. The commands aggregate statistics firstly for the entire B-tree and secondly for each individual level of the B-tree. The output displays the following values.

`indexStats.index`

The *index name*.

`indexStats.version`

The index version. For more information on index version numbers, see the `v` option in `db.collection.ensureIndex()` (page 30).

`indexStats.isIdIndex`

If `true`, the index is the default `_id` index for the collection.

`indexStats.keyPattern`

The indexed keys.

`indexStats.storageNs`

The namespace of the index's underlying storage.

`indexStats.bucketBodyBytes`

The fixed size, in bytes, of a B-tree bucket in the index, not including the record header. All indexes for a given version have the same value for this field. MongoDB allocates fixed size buckets on disk.

`indexStats.depth`

The number of levels in the B-tree, not including the root level.

indexStats.overall

This section of the output displays statistics for the entire B-tree.

indexStats.overall.numBuckets

The number of buckets in the entire B-tree, including all levels.

indexStats.overall.keyCount

Statistics about the number of keys in a bucket, evaluated on a per-bucket level.

indexStats.overall.usedKeyCount

Statistics about the number of used keys in a bucket, evaluated on a per-bucket level. Used keys are keys not marked as deleted.

indexStats.overall.bsonRatio

Statistics about the percentage of the bucket body that is occupied by the key objects themselves, excluding associated metadata.

For example, if you have the document { name: "Bob Smith" } and an index on { name: 1 }, the key object is the string Bob Smith.

indexStats.overall.keyNodeRatio

Statistics about the percentage of the bucket body that is occupied by the key node objects (the metadata and links pertaining to the keys). This does not include the key itself. In the current implementation, a key node's objects consist of: the pointer to the key data (in the same bucket), the pointer to the record the key is for, and the pointer to a child bucket.

indexStats.overall.fillRatio

The sum of the [bsonRatio](#) (page 340) and the [keyNodeRatio](#) (page 340). This shows how full the buckets are. This will be much higher for indexes with sequential inserts.

indexStats.perLevel

This section of the output displays statistics for each level of the B-tree separately, starting with the root level. This section displays a different document for each B-tree level.

indexStats.perLevel.numBuckets

The number of buckets at this level of the B-tree.

indexStats.perLevel.keyCount

Statistics about the number of keys in a bucket, evaluated on a per-bucket level.

indexStats.perLevel.usedKeyCount

Statistics about the number of used keys in a bucket, evaluated on a per-bucket level. Used keys are keys not marked as deleted.

indexStats.perLevel.bsonRatio

Statistics about the percentage of the bucket body that is occupied by the key objects themselves, excluding associated metadata.

indexStats.perLevel.keyNodeRatio

Statistics about the percentage of the bucket body that is occupied by the key node objects (the metadata and links pertaining to the keys).

indexStats.perLevel.fillRatio

The sum of the [bsonRatio](#) (page 340) and the [keyNodeRatio](#) (page 340). This shows how full the buckets are. This will be much higher in the following cases:

- For indexes with sequential inserts, such as the `_id` index when using ObjectId keys.
- For indexes that were recently built in the foreground with existing data.
- If you recently ran [compact](#) (page 313) or `--repair`.

Example The following is an example of `db.collection.indexStats()` (page 28) and `indexStats` (page 339) output.

```
{
  "index" : "type_1_traits_1",
  "version" : 1,
  "isIdIndex" : false,
  "keyPattern" : {
    "type" : 1,
    "traits" : 1
  },
  "storageNs" : "test.animals.$type_1_traits_1",
  "bucketBodyBytes" : 8154,
  "depth" : 2,
  "overall" : {
    "numBuckets" : 45513,
    "keyCount" : {
      "count" : NumberLong(45513),
      "mean" : 253.89602970579836,
      "stddev" : 21.784799875240708,
      "min" : 52,
      "max" : 290,
      "quantiles" : {
        "0.01" : 201.99785091648775,
        // ...
        "0.99" : 289.9999655156967
      }
    },
    "usedKeyCount" : {
      "count" : NumberLong(45513),
      // ...
      "quantiles" : {
        "0.01" : 201.99785091648775,
        // ...
        "0.99" : 289.9999655156967
      }
    },
    "bsonRatio" : {
      "count" : NumberLong(45513),
      // ...
      "quantiles" : {
        "0.01" : 0.4267797891997124,
        // ...
        "0.99" : 0.5945548174629648
      }
    },
    "keyNodeRatio" : {
      "count" : NumberLong(45513),
      // ...
      "quantiles" : {
        "0.01" : 0.3963656628236211,
        // ...
        "0.99" : 0.5690457993930765
      }
    },
    "fillRatio" : {
      "count" : NumberLong(45513),
      // ...
      "quantiles" : {
```

```
        "0.01" : 0.9909134214926929,
        // ...
        "0.99" : 0.9960755457453732
    }
}
},
"perLevel" : [
    {
        "numBuckets" : 1,
        "keyCount" : {
            "count" : NumberLong(1),
            "mean" : 180,
            "stddev" : 0,
            "min" : 180,
            "max" : 180
        },
        "usedKeyCount" : {
            "count" : NumberLong(1),
            // ...
            "max" : 180
        },
        "bsonRatio" : {
            "count" : NumberLong(1),
            // ...
            "max" : 0.3619082658817758
        },
        "keyNodeRatio" : {
            "count" : NumberLong(1),
            // ...
            "max" : 0.35320088300220753
        },
        "fillRatio" : {
            "count" : NumberLong(1),
            // ...
            "max" : 0.7151091488839834
        }
    },
    {
        "numBuckets" : 180,
        "keyCount" : {
            "count" : NumberLong(180),
            "mean" : 250.84444444444443,
            "stddev" : 26.30057503009355,
            "min" : 52,
            "max" : 290
        },
        "usedKeyCount" : {
            "count" : NumberLong(180),
            // ...
            "max" : 290
        },
        "bsonRatio" : {
            "count" : NumberLong(180),
            // ...
            "max" : 0.5945548197203826
        },
        "keyNodeRatio" : {
            "count" : NumberLong(180),
```

```

        // ...
        "max" : 0.5690458670591121
    },
    "fillRatio" : {
        "count" : NumberLong(180),
        // ...
        "max" : 0.9963208241353937
    }
},
{
    "numBuckets" : 45332,
    "keyCount" : {
        "count" : NumberLong(45332),
        "mean" : 253.90977675813994,
        "stddev" : 21.761620836279018,
        "min" : 167,
        "max" : 290,
        "quantiles" : {
            "0.01" : 202.0000012563603,
            // ...
            "0.99" : 289.99996486571894
        }
    },
    "usedKeyCount" : {
        "count" : NumberLong(45332),
        // ...
        "quantiles" : {
            "0.01" : 202.0000012563603,
            // ...
            "0.99" : 289.99996486571894
        }
    },
    "bsonRatio" : {
        "count" : NumberLong(45332),
        // ...
        "quantiles" : {
            "0.01" : 0.42678446958950583,
            // ...
            "0.99" : 0.5945548175411283
        }
    },
    "keyNodeRatio" : {
        "count" : NumberLong(45332),
        // ...
        "quantiles" : {
            "0.01" : 0.39636988227885306,
            // ...
            "0.99" : 0.5690457981176729
        }
    },
    "fillRatio" : {
        "count" : NumberLong(45332),
        // ...
        "quantiles" : {
            "0.01" : 0.9909246995605362,
            // ...
            "0.99" : 0.996075546919481
        }
    }
}

```

```
    }  
  }  
  ],  
  "ok" : 1  
}
```

Additional Resources For more information on the command's limits and output, see the following:

- The equivalent `db.collection.indexStats()` (page 28) method,
- *indexStats* (page 339), and
- <https://github.com/mongodb-labs/storage-viz#readme>.

whatsmyuri

whatsmyuri

`whatsmyuri` (page 344) is an internal command.

getLog

getLog

The `getLog` (page 344) command returns a document with a `log` array that contains recent messages from the `mongod` (page 503) process log. The `getLog` (page 344) command has the following syntax:

```
{ getLog: <log> }
```

Replace `<log>` with one of the following values:

- `global` - returns the combined output of all recent log entries.
- `rs` - if the `mongod` (page 503) is part of a *replica set*, `getLog` (page 344) will return recent notices related to replica set activity.
- `startupWarnings` - will return logs that *may* contain errors or warnings from MongoDB's log from when the current process started. If `mongod` (page 503) started without warnings, this filter may return an empty array.

You may also specify an asterisk (e.g. `*`) as the `<log>` value to return a list of available log filters. The following interaction from the `mongo` (page 527) shell connected to a replica set:

```
db.adminCommand({getLog: "*" })  
{ "names" : [ "global", "rs", "startupWarnings" ], "ok" : 1 }
```

`getLog` (page 344) returns events from a RAM cache of the `mongod` (page 503) events and *does not* read log data from the log file.

hostInfo

hostInfo

New in version 2.2.

Returns A document with information about the underlying system that the `mongod` (page 503) or `mongos` (page 518) runs on. Some of the returned fields are only included on some platforms.

You must run the `hostInfo` (page 344) command, which takes no arguments, against the `admin` database. Consider the following invocations of `hostInfo` (page 344):

```
db.hostInfo()  
db.adminCommand( { "hostInfo" : 1 } )
```

In the `mongo` (page 527) shell you can use `db.hostInfo()` (page 113) as a helper to access `hostInfo` (page 344). The output of `hostInfo` (page 344) on a Linux system will resemble the following:

```
{
  "system" : {
    "currentTime" : ISODate("<timestamp>"),
    "hostname" : "<hostname>",
    "cpuAddrSize" : <number>,
    "memSizeMB" : <number>,
    "numCores" : <number>,
    "cpuArch" : "<identifier>",
    "numaEnabled" : <boolean>
  },
  "os" : {
    "type" : "<string>",
    "name" : "<string>",
    "version" : "<string>"
  },
  "extra" : {
    "versionString" : "<string>",
    "libcVersion" : "<string>",
    "kernelVersion" : "<string>",
    "cpuFrequencyMHz" : "<string>",
    "cpuFeatures" : "<string>",
    "pageSize" : <number>,
    "numPages" : <number>,
    "maxOpenFiles" : <number>
  },
  "ok" : <return>
}
```

Consider the following documentation of these fields:

hostInfo

The document returned by the `hostInfo` (page 344).

hostInfo.system

A sub-document about the underlying environment of the system running the `mongod` (page 503) or `mongos` (page 518)

hostInfo.system.currentTime

A time stamp of the current system time.

hostInfo.system.hostname

The system name, which should correspond to the output of `hostname -f` on Linux systems.

hostInfo.system.cpuAddrSize

A number reflecting the architecture of the system. Either 32 or 64.

hostInfo.system.memSizeMB

The total amount of system memory (RAM) in megabytes.

hostInfo.system.numCores

The total number of available logical processor cores.

hostInfo.system.cpuArch

A string that represents the system architecture. Either `x86` or `x86_64`.

hostInfo.system.numaEnabled

A boolean value. `false` if NUMA is interleaved (i.e. disabled), otherwise `true`.

hostInfo.os

A sub-document that contains information about the operating system running the [mongod](#) (page 503) and [mongos](#) (page 518).

hostInfo.os.type

A string representing the type of operating system, such as Linux or Windows.

hostInfo.os.name

If available, returns a display name for the operating system.

hostInfo.os.version

If available, returns the name of the distribution or operating system.

hostInfo.extra

A sub-document with extra information about the operating system and the underlying hardware. The content of the [extra](#) (page 346) sub-document depends on the operating system.

hostInfo.extra.versionString

A complete string of the operating system version and identification. On Linux and OS X systems, this contains output similar to `uname -a`.

hostInfo.extra.libcVersion

The release of the system `libc`.

[libcVersion](#) (page 346) only appears on Linux systems.

hostInfo.extra.kernelVersion

The release of the Linux kernel in current use.

[kernelVersion](#) (page 346) only appears on Linux systems.

hostInfo.extra.alwaysFullSync

[alwaysFullSync](#) (page 346) only appears on OS X systems.

hostInfo.extra.nfsAsync

[nfsAsync](#) (page 346) only appears on OS X systems.

hostInfo.extra.cpuFrequencyMHz

Reports the clock speed of the system's processor in megahertz.

hostInfo.extra.cpuFeatures

Reports the processor feature flags. On Linux systems this the same information that <http://docs.mongodb.org/manualproc/cpuinfo> includes in the `flags` fields.

hostInfo.extra.pageSize

Reports the default system page size in bytes.

hostInfo.extra.numPages

[numPages](#) (page 346) only appears on Linux systems.

hostInfo.extra.maxOpenFiles

Reports the current system limits on open file handles. See <http://docs.mongodb.org/manualreference/ulimit> for more information.

[maxOpenFiles](#) (page 346) only appears on Linux systems.

hostInfo.extra.scheduler

Reports the active I/O scheduler. [scheduler](#) (page 346) only appears on OS X systems.

serverStatus

Definition**serverStatus**

The `serverStatus` (page 347) command returns a document that provides an overview of the database process's state. Most monitoring applications run this command at a regular interval to collection statistics about the instance:

```
{ serverStatus: 1 }
```

The value (i.e. 1 above), does not affect the operation of the command.

Changed in version 2.4: In 2.4 you can dynamically suppress portions of the `serverStatus` (page 347) output, or include suppressed sections by adding fields to the command document as in the following examples:

```
db.runCommand( { serverStatus: 1, repl: 0, indexCounters: 0 } )
db.runCommand( { serverStatus: 1, workingSet: 1, metrics: 0, locks: 0 } )
```

`serverStatus` (page 347) includes all fields by default, except `workingSet` (page 359), by default.

Note: You may only dynamically include top-level fields from the `serverStatus` (page 346) document that are not included by default. You can exclude any field that `serverStatus` (page 347) includes by default.

See also:

`db.serverStatus()` (page 118) and <http://docs.mongodb.org/manualreference/server-status>

Output The `serverStatus` (page 347) command returns a collection of information that reflects the database's status. These data are useful for diagnosing and assessing the performance of your MongoDB instance. This reference catalogs each datum included in the output of this command and provides context for using this data to more effectively administer your database.

See also:

Much of the output of `serverStatus` (page 347) is also displayed dynamically by `mongostat` (page 570). See the `mongostat` (page 569) command for more information.

For examples of the `serverStatus` (page 347) output, see <http://docs.mongodb.org/manualreference/server-status>

Instance Information For an example of the instance information, see the *Instance Information section* of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.host

The `host` (page 347) field contains the system's hostname. In Unix/Linux systems, this should be the same as the output of the `hostname` command.

serverStatus.version

The `version` (page 347) field contains the version of MongoDB running on the current `mongod` (page 503) or `mongos` (page 518) instance.

serverStatus.process

The `process` (page 347) field identifies which kind of MongoDB instance is running. Possible values are:

- `mongos` (page 518)
- `mongod` (page 503)

serverStatus.uptime

The value of the `uptime` (page 347) field corresponds to the number of seconds that the `mongos` (page 518) or `mongod` (page 503) process has been active.

serverStatus.uptimeEstimate

`uptimeEstimate` (page 347) provides the uptime as calculated from MongoDB's internal course-grained time keeping system.

serverStatus.localTime

The `localTime` (page 348) value is the current time, according to the server, in UTC specified in an ISODate format.

locks New in version 2.1.2: All `locks` (page 348) statuses first appeared in the 2.1.2 development release for the 2.2 series.

For an example of the `locks` output, see the *locks section* of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.locks

The `locks` (page 348) document contains sub-documents that provides a granular report on MongoDB database-level lock use. All values are of the `NumberLong()` type.

Generally, fields named:

- R refer to the global read lock,
- W refer to the global write lock,
- r refer to the database specific read lock, and
- w refer to the database specific write lock.

If a document does not have any fields, it means that no locks have existed with this context since the last time the `mongod` (page 503) started.

serverStatus.locks..

A field named `.` holds the first document in `locks` (page 348) that contains information about the global lock.

serverStatus.locks...timeLockedMicros

The `timeLockedMicros` (page 348) document reports the amount of time in microseconds that a lock has existed in all databases in this `mongod` (page 503) instance.

serverStatus.locks...timeLockedMicros.R

The `R` field reports the amount of time in microseconds that any database has held the global read lock.

serverStatus.locks...timeLockedMicros.W

The `W` field reports the amount of time in microseconds that any database has held the global write lock.

serverStatus.locks...timeLockedMicros.r

The `r` field reports the amount of time in microseconds that any database has held the local read lock.

serverStatus.locks...timeLockedMicros.w

The `w` field reports the amount of time in microseconds that any database has held the local write lock.

serverStatus.locks...timeAcquiringMicros

The `timeAcquiringMicros` (page 348) document reports the amount of time in microseconds that operations have spent waiting to acquire a lock in all databases in this `mongod` (page 503) instance.

serverStatus.locks...timeAcquiringMicros.R

The `R` field reports the amount of time in microseconds that any database has spent waiting for the global read lock.

serverStatus.locks...timeAcquiringMicros.W

The `W` field reports the amount of time in microseconds that any database has spent waiting for the global write lock.

`serverStatus.locks.admin`

The `admin` (page 348) document contains two sub-documents that report data regarding lock use in the *admin database*.

`serverStatus.locks.admin.timeLockedMicros`

The `timeLockedMicros` (page 349) document reports the amount of time in microseconds that locks have existed in the context of the *admin database*.

`serverStatus.locks.admin.timeLockedMicros.r`

The `r` field reports the amount of time in microseconds that the *admin database* has held the read lock.

`serverStatus.locks.admin.timeLockedMicros.w`

The `w` field reports the amount of time in microseconds that the *admin database* has held the write lock.

`serverStatus.locks.admin.timeAcquiringMicros`

The `timeAcquiringMicros` (page 349) document reports on the amount of field time in microseconds that operations have spent waiting to acquire a lock for the *admin database*.

`serverStatus.locks.admin.timeAcquiringMicros.r`

The `r` field reports the amount of time in microseconds that operations have spent waiting to acquire a read lock on the *admin database*.

`serverStatus.locks.admin.timeAcquiringMicros.w`

The `w` field reports the amount of time in microseconds that operations have spent waiting to acquire a write lock on the *admin database*.

`serverStatus.locks.local`

The `local` (page 349) document contains two sub-documents that report data regarding lock use in the `local` database. The `local` database contains a number of instance specific data, including the *oplog* for replication.

`serverStatus.locks.local.timeLockedMicros`

The `timeLockedMicros` (page 349) document reports on the amount of time in microseconds that locks have existed in the context of the `local` database.

`serverStatus.locks.local.timeLockedMicros.r`

The `r` field reports the amount of time in microseconds that the `local` database has held the read lock.

`serverStatus.locks.local.timeLockedMicros.w`

The `w` field reports the amount of time in microseconds that the `local` database has held the write lock.

`serverStatus.locks.local.timeAcquiringMicros`

The `timeAcquiringMicros` (page 349) document reports on the amount of time in microseconds that operations have spent waiting to acquire a lock for the `local` database.

`serverStatus.locks.local.timeAcquiringMicros.r`

The `r` field reports the amount of time in microseconds that operations have spent waiting to acquire a read lock on the `local` database.

`serverStatus.locks.local.timeAcquiringMicros.w`

The `w` field reports the amount of time in microseconds that operations have spent waiting to acquire a write lock on the `local` database.

`serverStatus.locks.<database>`

For each additional database `locks` (page 348) includes a document that reports on the lock use for this database. The names of these documents reflect the database name itself.

`serverStatus.locks.<database>.timeLockedMicros`

The `timeLockedMicros` (page 349) document reports on the amount of time in microseconds that locks have existed in the context of the `<database>` database.

`serverStatus.locks.<database>.timeLockedMicros.r`

The `r` field reports the amount of time in microseconds that the `<database>` database has held the read lock.

`serverStatus.locks.<database>.timeLockedMicros.w`

The `w` field reports the amount of time in microseconds that the `<database>` database has held the write lock.

`serverStatus.locks.<database>.timeAcquiringMicros`

The `timeAcquiringMicros` (page 350) document reports on the amount of time in microseconds that operations have spent waiting to acquire a lock for the `<database>` database.

`serverStatus.locks.<database>.timeAcquiringMicros.r`

The `r` field reports the amount of time in microseconds that operations have spent waiting to acquire a read lock on the `<database>` database.

`serverStatus.locks.<database>.timeAcquiringMicros.w`

The `w` field reports the amount of time in microseconds that operations have spent waiting to acquire a write lock on the `<database>` database.

globalLock For an example of the `globalLock` output, see the *globalLock* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.globalLock`

The `globalLock` (page 350) data structure contains information regarding the database's current lock state, historical lock status, current operation queue, and the number of active clients.

`serverStatus.globalLock.totalTime`

The value of `totalTime` (page 350) represents the time, in microseconds, since the database last started and creation of the `globalLock` (page 350). This is roughly equivalent to total server uptime.

`serverStatus.globalLock.lockTime`

The value of `lockTime` (page 350) represents the time, in microseconds, since the database last started, that the `globalLock` (page 350) has been *held*.

Consider this value in combination with the value of `totalTime` (page 350). MongoDB aggregates these values in the `ratio` (page 350) value. If the `ratio` (page 350) value is small but `totalTime` (page 350) is high the `globalLock` (page 350) has typically been held frequently for shorter periods of time, which may be indicative of a more normal use pattern. If the `lockTime` (page 350) is higher and the `totalTime` (page 350) is smaller (relatively) then fewer operations are responsible for a greater portion of server's use (relatively).

`serverStatus.globalLock.ratio`

Changed in version 2.2: `ratio` (page 350) was removed. See `locks` (page 348).

The value of `ratio` (page 350) displays the relationship between `lockTime` (page 350) and `totalTime` (page 350).

Low values indicate that operations have held the `globalLock` (page 350) frequently for shorter periods of time. High values indicate that operations have held `globalLock` (page 350) infrequently for longer periods of time.

`serverStatus.globalLock.currentQueue`

The `currentQueue` (page 350) data structure value provides more granular information concerning the number of operations queued because of a lock.

`serverStatus.globalLock.currentQueue.total`

The value of `total` (page 350) provides a combined total of operations queued waiting for the lock.

A consistently small queue, particularly of shorter operations should cause no concern. Also, consider this value in light of the size of queue waiting for the read lock (e.g. `readers` (page 350)) and write lock (e.g. `writers` (page 350)) individually.

`serverStatus.globalLock.currentQueue.readers`

The value of `readers` (page 350) is the number of operations that are currently queued and waiting for the read lock. A consistently small read-queue, particularly of shorter operations should cause no concern.

`serverStatus.globalLock.currentQueue.writers`

The value of `writers` (page 350) is the number of operations that are currently queued and waiting for the write lock. A consistently small write-queue, particularly of shorter operations is no cause for concern.

globalLock.activeClients

`serverStatus.globalLock.activeClients`

The `activeClients` (page 351) data structure provides more granular information about the number of connected clients and the operation types (e.g. read or write) performed by these clients.

Use this data to provide context for the `currentQueue` (page 350) data.

`serverStatus.globalLock.activeClients.total`

The value of `total` (page 351) is the total number of active client connections to the database. This combines clients that are performing read operations (e.g. `readers` (page 351)) and clients that are performing write operations (e.g. `writers` (page 351)).

`serverStatus.globalLock.activeClients.readers`

The value of `readers` (page 351) contains a count of the active client connections performing read operations.

`serverStatus.globalLock.activeClients.writers`

The value of `writers` (page 351) contains a count of active client connections performing write operations.

mem For an example of the `mem` output, see the *mem* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.mem`

The `mem` (page 351) data structure holds information regarding the target system architecture of `mongod` (page 503) and current memory use.

`serverStatus.mem.bits`

The value of `bits` (page 351) is either 64 or 32, depending on which target architecture specified during the `mongod` (page 503) compilation process. In most instances this is 64, and this value does not change over time.

`serverStatus.mem.resident`

The value of `resident` (page 351) is roughly equivalent to the amount of RAM, in megabytes (MB), currently used by the database process. In normal use this value tends to grow. In dedicated database servers this number tends to approach the total amount of system memory.

`serverStatus.mem.virtual`

`virtual` (page 351) displays the quantity, in megabytes (MB), of virtual memory used by the `mongod` (page 503) process. With *journaling* enabled, the value of `virtual` (page 351) is at least twice the value of `mapped` (page 351).

If `virtual` (page 351) value is significantly larger than `mapped` (page 351) (e.g. 3 or more times), this may indicate a memory leak.

`serverStatus.mem.supported`

`supported` (page 351) is true when the underlying system supports extended memory information. If this value is false and the system does not support extended memory information, then other `mem` (page 351) values may not be accessible to the database server.

`serverStatus.mem.mapped`

The value of `mapped` (page 351) provides the amount of mapped memory, in megabytes (MB), by the database. Because MongoDB uses memory-mapped files, this value is likely to be to be roughly equivalent to the total size of your database or databases.

`serverStatus.mem.mappedWithJournal`

`mappedWithJournal` (page 351) provides the amount of mapped memory, in megabytes (MB), including

the memory used for journaling. This value will always be twice the value of `mapped` (page 351). This field is only included if journaling is enabled.

connections For an example of the `connections` output, see the *connections* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.connections

The `connections` (page 352) sub document data regarding the current connection status and availability of the database server. Use these values to assess the current load and capacity requirements of the server.

serverStatus.connections.current

The value of `current` (page 352) corresponds to the number of connections to the database server from clients. This number includes the current shell session. Consider the value of `available` (page 352) to add more context to this datum.

This figure will include the current shell connection as well as any inter-node connections to support a *replica set* or *sharded cluster*.

serverStatus.connections.available

`available` (page 352) provides a count of the number of unused available connections that the database can provide. Consider this value in combination with the value of `current` (page 352) to understand the connection load on the database, and the <http://docs.mongodb.org/manualreference/ulimit> document for more information about system thresholds on available connections.

serverStatus.connections.totalCreated

`totalCreated` (page 352) provides a count of **all** connections created to the server. This number includes connections that have since closed.

extra_info For an example of the `extra_info` output, see the *extra_info* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.extra_info

The `extra_info` (page 352) data structure holds data collected by the `mongod` (page 503) instance about the underlying system. Your system may only report a subset of these fields.

serverStatus.extra_info.note

The field `note` (page 352) reports that the data in this structure depend on the underlying platform, and has the text: “fields vary by platform.”

serverStatus.extra_info.heap_usage_bytes

The `heap_usage_bytes` (page 352) field is only available on Unix/Linux systems, and reports the total size in bytes of heap space used by the database process.

serverStatus.extra_info.page_faults

The `page_faults` (page 352) Reports the total number of page faults that require disk operations. Page faults refer to operations that require the database server to access data which isn’t available in active memory. The `page_faults` (page 352) counter may increase dramatically during moments of poor performance and may correlate with limited memory environments and larger data sets. Limited and sporadic page faults do not necessarily indicate an issue.

indexCounters For an example of the `indexCounters` output, see the *indexCounters* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.indexCounters

Changed in version 2.2: Previously, data in the `indexCounters` (page 352) document reported sampled data, and were only useful in relative comparison to each other, because they could not reflect absolute index use. In 2.2 and later, these data reflect actual index use.

Changed in version 2.4: Fields previously in the `btree` sub-document of `indexCounters` (page 352) are now fields in the `indexCounters` (page 352) document.

The `indexCounters` (page 352) data structure reports information regarding the state and use of indexes in MongoDB.

`serverStatus.indexCounters.accesses`

`accesses` (page 353) reports the number of times that operations have accessed indexes. This value is the combination of the `hits` (page 353) and `misses` (page 353). Higher values indicate that your database has indexes and that queries are taking advantage of these indexes. If this number does not grow over time, this might indicate that your indexes do not effectively support your use.

`serverStatus.indexCounters.hits`

The `hits` (page 353) value reflects the number of times that an index has been accessed and `mongod` (page 503) is able to return the index from memory.

A higher value indicates effective index use. `hits` (page 353) values that represent a greater proportion of the `accesses` (page 353) value, tend to indicate more effective index configuration.

`serverStatus.indexCounters.misses`

The `misses` (page 353) value represents the number of times that an operation attempted to access an index that was not in memory. These “misses,” do not indicate a failed query or operation, but rather an inefficient use of the index. Lower values in this field indicate better index use and likely overall performance as well.

`serverStatus.indexCounters.resets`

The `resets` (page 353) value reflects the number of times that the index counters have been reset since the database last restarted. Typically this value is 0, but use this value to provide context for the data specified by other `indexCounters` (page 352) values.

`serverStatus.indexCounters.missRatio`

The `missRatio` (page 353) value is the ratio of `hits` (page 353) to `misses` (page 353). This value is typically 0 or approaching 0.

backgroundFlushing For an example of the `backgroundFlushing` output, see the *backgroundFlushing* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.backgroundFlushing`

`mongod` (page 503) periodically flushes writes to disk. In the default configuration, this happens every 60 seconds. The `backgroundFlushing` (page 353) data structure contains data regarding these operations. Consider these values if you have concerns about write performance and *journaling* (page 357).

`serverStatus.backgroundFlushing.flushes`

`flushes` (page 353) is a counter that collects the number of times the database has flushed all writes to disk. This value will grow as database runs for longer periods of time.

`serverStatus.backgroundFlushing.total_ms`

The `total_ms` (page 353) value provides the total number of milliseconds (ms) that the `mongod` (page 503) processes have spent writing (i.e. flushing) data to disk. Because this is an absolute value, consider the value of `flushes` (page 353) and `average_ms` (page 353) to provide better context for this datum.

`serverStatus.backgroundFlushing.average_ms`

The `average_ms` (page 353) value describes the relationship between the number of flushes and the total amount of time that the database has spent writing data to disk. The larger `flushes` (page 353) is, the more likely this value is likely to represent a “normal,” time; however, abnormal data can skew this value.

Use the `last_ms` (page 353) to ensure that a high average is not skewed by transient historical issue or a random write distribution.

`serverStatus.backgroundFlushing.last_ms`

The value of the `last_ms` (page 353) field is the amount of time, in milliseconds, that the last flush operation

took to complete. Use this value to verify that the current performance of the server and is in line with the historical data provided by `average_ms` (page 353) and `total_ms` (page 353).

`serverStatus.backgroundFlushing.last_finished`

The `last_finished` (page 354) field provides a timestamp of the last completed flush operation in the *ISODate* format. If this value is more than a few minutes old relative to your server's current time and accounting for differences in time zone, restarting the database may result in some data loss.

Also consider ongoing operations that might skew this value by routinely block write operations.

cursors Deprecatd since version 2.6: See the `serverStatus.metrics.cursor` (page 360) field instead.

For an example of the `cursors` output, see the *cursors* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.cursors`

The `cursors` (page 354) data structure contains data regarding cursor state and use.

`serverStatus.cursors.note`

A note specifying to use the `serverStatus.metrics.cursor` (page 363) field instead of `serverStatus.cursors` (page 354).

`serverStatus.cursors.totalOpen`

`totalOpen` (page 354) provides the number of cursors that MongoDB is maintaining for clients. Because MongoDB exhausts unused cursors, typically this value small or zero. However, if there is a queue, stale tailable cursors, or a large number of operations this value may rise.

`serverStatus.cursors.clientCursors_size`

Deprecatd since version 1.x: See `totalOpen` (page 354) for this datum.

`serverStatus.cursors.timedOut`

`timedOut` (page 354) provides a counter of the total number of cursors that have timed out since the server process started. If this number is large or growing at a regular rate, this may indicate an application error.

`serverStatus.cursors.totalNoTimeout`

`totalNoTimeout` (page 354) provides the number of open cursors with the option `DBQuery.Option.noTimeout` (page 78) set to prevent timeout after a period of inactivity.

`serverStatus.cursors.pinned`

`serverStatus.cursors.pinned` (page 354) provides the number of "pinned" open cursors.

network For an example of the `network` output, see the *network* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.network`

The `network` (page 354) data structure contains data regarding MongoDB's network use.

`serverStatus.network.bytesIn`

The value of the `bytesIn` (page 354) field reflects the amount of network traffic, in bytes, received *by* this database. Use this value to ensure that network traffic sent to the `mongod` (page 503) process is consistent with expectations and overall inter-application traffic.

`serverStatus.network.bytesOut`

The value of the `bytesOut` (page 354) field reflects the amount of network traffic, in bytes, sent *from* this database. Use this value to ensure that network traffic sent by the `mongod` (page 503) process is consistent with expectations and overall inter-application traffic.

`serverStatus.network.numRequests`

The `numRequests` (page 354) field is a counter of the total number of distinct requests that the server has

received. Use this value to provide context for the `bytesIn` (page 354) and `bytesOut` (page 354) values to ensure that MongoDB's network utilization is consistent with expectations and application use.

repl For an example of the `repl` output, see the *repl* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.repl`

The `repl` (page 355) data structure contains status information for MongoDB's replication (i.e. "replica set") configuration. These values only appear when the current host has replication enabled.

See <http://docs.mongodb.org/manualreplication> for more information on replication.

`serverStatus.repl.setName`

The `setName` (page 355) field contains a string with the name of the current replica set. This value reflects the `--replSet` command line argument, or `replSetName` value in the configuration file.

See <http://docs.mongodb.org/manualreplication> for more information on replication.

`serverStatus.repl.ismaster`

The value of the `ismaster` (page 355) field is either `true` or `false` and reflects whether the current node is the master or primary node in the replica set.

See <http://docs.mongodb.org/manualreplication> for more information on replication.

`serverStatus.repl.secondary`

The value of the `secondary` (page 355) field is either `true` or `false` and reflects whether the current node is a secondary node in the replica set.

See <http://docs.mongodb.org/manualreplication> for more information on replication.

`serverStatus.repl.hosts`

`hosts` (page 355) is an array that lists the other nodes in the current replica set. Each member of the replica set appears in the form of `hostname:port`.

See <http://docs.mongodb.org/manualreplication> for more information on replication.

opcountersRepl For an example of the `opcountersRepl` output, see the *opcountersRepl* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.opcountersRepl`

The `opcountersRepl` (page 355) data structure, similar to the `opcounters` (page 356) data structure, provides an overview of database replication operations by type and makes it possible to analyze the load on the replica in more granular manner. These values only appear when the current host has replication enabled.

These values will differ from the `opcounters` (page 356) values because of how MongoDB serializes operations during replication. See <http://docs.mongodb.org/manualreplication> for more information on replication.

These numbers will grow over time in response to database use. Analyze these values over time to track database utilization.

`serverStatus.opcountersRepl.insert`

`insert` (page 355) provides a counter of the total number of replicated insert operations since the `mongod` (page 503) instance last started.

`serverStatus.opcountersRepl.query`

`query` (page 355) provides a counter of the total number of replicated queries since the `mongod` (page 503) instance last started.

`serverStatus.opcountersRepl.update`

`update` (page 355) provides a counter of the total number of replicated update operations since the `mongod` (page 503) instance last started.

`serverStatus.opcountersRepl.delete`

`delete` (page 356) provides a counter of the total number of replicated delete operations since the `mongod` (page 503) instance last started.

`serverStatus.opcountersRepl.getmore`

`getmore` (page 356) provides a counter of the total number of “getmore” operations since the `mongod` (page 503) instance last started. This counter can be high even if the query count is low. Secondary nodes send `getMore` operations as part of the replication process.

`serverStatus.opcountersRepl.command`

`command` (page 356) provides a counter of the total number of replicated commands issued to the database since the `mongod` (page 503) instance last started.

opcounters For an example of the `opcounters` output, see the *opcounters* section of the <http://docs.mongodb.org/manualreference/server-status> page.

`serverStatus.opcounters`

The `opcounters` (page 356) data structure provides an overview of database operations by type and makes it possible to analyze the load on the database in more granular manner.

These numbers will grow over time and in response to database use. Analyze these values over time to track database utilization.

Note: The data in `opcounters` (page 356) treats operations that affect multiple documents, such as bulk insert or multi-update operations, as a single operation. See `document` (page 360) for more granular document-level operation tracking.

`serverStatus.opcounters.insert`

`insert` (page 356) provides a counter of the total number of insert operations since the `mongod` (page 503) instance last started.

`serverStatus.opcounters.query`

`query` (page 356) provides a counter of the total number of queries since the `mongod` (page 503) instance last started.

`serverStatus.opcounters.update`

`update` (page 356) provides a counter of the total number of update operations since the `mongod` (page 503) instance last started.

`serverStatus.opcounters.delete`

`delete` (page 356) provides a counter of the total number of delete operations since the `mongod` (page 503) instance last started.

`serverStatus.opcounters.getmore`

`getmore` (page 356) provides a counter of the total number of “getmore” operations since the `mongod` (page 503) instance last started. This counter can be high even if the query count is low. Secondary nodes send `getMore` operations as part of the replication process.

`serverStatus.opcounters.command`

`command` (page 356) provides a counter of the total number of commands issued to the database since the `mongod` (page 503) instance last started.

asserts For an example of the `asserts` output, see the *asserts* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.asserts

The `asserts` (page 356) document reports the number of asserts on the database. While assert errors are typically uncommon, if there are non-zero values for the `asserts` (page 356), you should check the log file for the `mongod` (page 503) process for more information. In many cases these errors are trivial, but are worth investigating.

serverStatus.asserts.regular

The `regular` (page 357) counter tracks the number of regular assertions raised since the server process started. Check the log file for more information about these messages.

serverStatus.asserts.warning

The `warning` (page 357) counter tracks the number of warnings raised since the server process started. Check the log file for more information about these warnings.

serverStatus.asserts.msg

The `msg` (page 357) counter tracks the number of message assertions raised since the server process started. Check the log file for more information about these messages.

serverStatus.asserts.user

The `user` (page 357) counter reports the number of “user asserts” that have occurred since the last time the server process started. These are errors that user may generate, such as out of disk space or duplicate key. You can prevent these assertions by fixing a problem with your application or deployment. Check the MongoDB log for more information.

serverStatus.asserts.rollovers

The `rollovers` (page 357) counter displays the number of times that the rollover counters have rolled over since the last time the server process started. The counters will rollover to zero after 2^{30} assertions. Use this value to provide context to the other values in the `asserts` (page 356) data structure.

writeBacksQueued For an example of the `writeBacksQueued` output, see the *writeBacksQueued* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.writeBacksQueued

The value of `writeBacksQueued` (page 357) is `true` when there are operations from a `mongos` (page 518) instance queued for retrying. Typically this option is `false`.

See also:

writeBacks

Journaling (dur) New in version 1.8.

For an example of the Journaling (dur) output, see the *journaling* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.dur

The `dur` (page 357) (for “durability”) document contains data regarding the `mongod` (page 503)’s journaling-related operations and performance. `mongod` (page 503) must be running with journaling for these data to appear in the output of “`serverStatus` (page 347)”.

MongoDB reports the data in `dur` (page 357) based on 3 second intervals of data, collected between 3 and 6 seconds in the past.

See also:

<http://docs.mongodb.org/manualcore/journaling> for more information about journaling operations.

serverStatus.dur.commits

The `commits` (page 357) provides the number of transactions written to the *journal* during the last *journal group commit interval*.

serverStatus.dur.journaleMB

The `journaleMB` (page 358) provides the amount of data in megabytes (MB) written to *journal* during the last *journal group commit interval*.

serverStatus.dur.writeToDataFilesMB

The `writeToDataFilesMB` (page 358) provides the amount of data in megabytes (MB) written from *journal* to the data files during the last *journal group commit interval*.

serverStatus.dur.compression

New in version 2.0.

The `compression` (page 358) represents the compression ratio of the data written to the *journal*:

(`journaleMB` / `uncompressed_size_of_data`)

serverStatus.dur.commitsInWriteLock

The `commitsInWriteLock` (page 358) provides a count of the commits that occurred while a write lock was held. Commits in a write lock indicate a MongoDB node under a heavy write load and call for further diagnosis.

serverStatus.dur.earlyCommits

The `earlyCommits` (page 358) value reflects the number of times MongoDB requested a commit before the scheduled *journal group commit interval*. Use this value to ensure that your *journal group commit interval* is not too long for your deployment.

serverStatus.dur.timeMS

The `timeMS` (page 358) document provides information about the performance of the `mongod` (page 503) instance during the various phases of journaling in the last *journal group commit interval*.

serverStatus.dur.timeMS.dt

The `dt` (page 358) value provides, in milliseconds, the amount of time over which MongoDB collected the `timeMS` (page 358) data. Use this field to provide context to the other `timeMS` (page 358) field values.

serverStatus.dur.timeMS.prepLogBuffer

The `prepLogBuffer` (page 358) value provides, in milliseconds, the amount of time spent preparing to write to the journal. Smaller values indicate better journal performance.

serverStatus.dur.timeMS.writeToJournal

The `writeToJournal` (page 358) value provides, in milliseconds, the amount of time spent actually writing to the journal. File system speeds and device interfaces can affect performance.

serverStatus.dur.timeMS.writeToDataFiles

The `writeToDataFiles` (page 358) value provides, in milliseconds, the amount of time spent writing to data files after journaling. File system speeds and device interfaces can affect performance.

serverStatus.dur.timeMS.remapPrivateView

The `remapPrivateView` (page 358) value provides, in milliseconds, the amount of time spent remapping copy-on-write memory mapped views. Smaller values indicate better journal performance.

recordStats For an example of the `recordStats` output, see the *recordStats* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.recordStats

The `recordStats` (page 358) document provides fine grained reporting on page faults on a per database level.

MongoDB uses a read lock on each database to return `recordStats` (page 358). To minimize this overhead, you can disable this section, as in the following operation:

```
db.serverStatus( { recordStats: 0 } )
```

serverStatus.recordStats.accessesNotInMemory

`accessesNotInMemory` (page 359) reflects the number of times `mongod` (page 503) needed to access a memory page that was *not* resident in memory for *all* databases managed by this `mongod` (page 503) instance.

serverStatus.recordStats.pageFaultExceptionsThrown

`pageFaultExceptionsThrown` (page 359) reflects the number of page fault exceptions thrown by `mongod` (page 503) when accessing data for *all* databases managed by this `mongod` (page 503) instance.

serverStatus.recordStats.local.accessesNotInMemory

`accessesNotInMemory` (page 359) reflects the number of times `mongod` (page 503) needed to access a memory page that was *not* resident in memory for the `local` database.

serverStatus.recordStats.local.pageFaultExceptionsThrown

`pageFaultExceptionsThrown` (page 359) reflects the number of page fault exceptions thrown by `mongod` (page 503) when accessing data for the `local` database.

serverStatus.recordStats.admin.accessesNotInMemory

`accessesNotInMemory` (page 359) reflects the number of times `mongod` (page 503) needed to access a memory page that was *not* resident in memory for the *admin database*.

serverStatus.recordStats.admin.pageFaultExceptionsThrown

`pageFaultExceptionsThrown` (page 359) reflects the number of page fault exceptions thrown by `mongod` (page 503) when accessing data for the *admin database*.

serverStatus.recordStats.<database>.accessesNotInMemory

`accessesNotInMemory` (page 359) reflects the number of times `mongod` (page 503) needed to access a memory page that was *not* resident in memory for the `<database>` database.

serverStatus.recordStats.<database>.pageFaultExceptionsThrown

`pageFaultExceptionsThrown` (page 359) reflects the number of page fault exceptions thrown by `mongod` (page 503) when accessing data for the `<database>` database.

workingSet New in version 2.4.

Note: The `workingSet` (page 359) data is only included in the output of `serverStatus` (page 347) if explicitly enabled. To return the `workingSet` (page 359), use one of the following commands:

```
db.serverStatus( { workingSet: 1 } )
db.runCommand( { serverStatus: 1, workingSet: 1 } )
```

For an example of the `workingSet` output, see the *workingSet* section of the <http://docs.mongodb.org/manualreference/server-status> page.

serverStatus.workingSet

`workingSet` (page 359) is a document that contains values useful for estimating the size of the working set, which is the amount of data that MongoDB uses actively. `workingSet` (page 359) uses an internal data structure that tracks pages accessed by `mongod` (page 503).

serverStatus.workingSet.note

`note` (page 359) is a field that holds a string warning that the `workingSet` (page 359) document is an estimate.

serverStatus.workingSet.pagesInMemory

`pagesInMemory` (page 359) contains a count of the total number of pages accessed by `mongod` (page 503) over the period displayed in `overSeconds` (page 360). The default page size is 4 kilobytes; to convert this value to the amount of data in memory multiply this value by 4 kilobytes.

If your total working set is less than the size of physical memory, over time the value of `pagesInMemory` (page 359) will reflect your data size.

Use `pagesInMemory` (page 359) in conjunction with `overSeconds` (page 360) to help estimate the actual size of the working set.

`serverStatus.workingSet.computationTimeMicros`

`computationTimeMicros` (page 360) reports the amount of time the `mongod` (page 503) instance used to compute the other fields in the `workingSet` (page 359) section.

Reporting on `workingSet` (page 359) may impact the performance of other operations on the `mongod` (page 503) instance because MongoDB must collect some data within the context of a lock. Ensure that automated monitoring tools consider this metric when determining the frequency of collection for `workingSet` (page 359).

`serverStatus.workingSet.overSeconds`

`overSeconds` (page 360) returns the amount of time elapsed between the newest and oldest pages tracked in the `pagesInMemory` (page 359) data point.

If `overSeconds` (page 360) is decreasing, or if `pagesInMemory` (page 359) equals physical RAM and `overSeconds` (page 360) is very small, the working set may be much *larger* than physical RAM.

When `overSeconds` (page 360) is large, MongoDB's data set is equal to or *smaller* than physical RAM.

metrics For an example of the metrics output, see the *metrics* section of the <http://docs.mongodb.org/manualreference/server-status> page.

New in version 2.4.

`serverStatus.metrics`

The `metrics` (page 360) document holds a number of statistics that reflect the current use and state of a running `mongod` (page 503) instance.

`serverStatus.metrics.document`

The `document` (page 360) holds a document of that reflect document access and modification patterns and data use. Compare these values to the data in the `opcounters` (page 356) document, which track total number of operations.

`serverStatus.metrics.document.deleted`

`deleted` (page 360) reports the total number of documents deleted.

`serverStatus.metrics.document.inserted`

`inserted` (page 360) reports the total number of documents inserted.

`serverStatus.metrics.document.returned`

`returned` (page 360) reports the total number of documents returned by queries.

`serverStatus.metrics.document.updated`

`updated` (page 360) reports the total number of documents updated.

`serverStatus.metrics.getLastError`

`getLastError` (page 360) is a document that reports on `getLastError` (page 235) use.

`serverStatus.metrics.getLastError.wtime`

`wtime` (page 360) is a sub-document that reports `getLastError` (page 235) operation counts with a `w` argument greater than 1.

`serverStatus.metrics.getLastError.wtime.num`

`num` (page 360) reports the total number of `getLastError` (page 235) operations with a specified write concern (i.e. `w`) that wait for one or more members of a replica set to acknowledge the write operation (i.e. a `w` value greater than 1.)

`serverStatus.metrics.getLastError.wtime.totalMillis`

`totalMillis` (page 360) reports the total amount of time in milliseconds that the `mongod` (page 503) has spent performing `getLastError` (page 235) operations with write concern (i.e. `w`) that wait for one or more members of a replica set to acknowledge the write operation (i.e. a `w` value greater than 1.)

`serverStatus.metrics.getLastError.wtimeouts`

`wtimeouts` (page 361) reports the number of times that *write concern* operations have timed out as a result of the `wtimeout` threshold to `getLastError` (page 235).

`serverStatus.metrics.operation`

`operation` (page 361) is a sub-document that holds counters for several types of update and query operations that MongoDB handles using special operation types.

`serverStatus.metrics.operation.fastmod`

`fastmod` (page 361) reports the number of update operations that neither cause documents to grow nor require updates to the index. For example, this counter would record an update operation that use the `$inc` (page 412) operator to increment the value of a field that is not indexed.

`serverStatus.metrics.operation.idhack`

`idhack` (page 361) reports the number of queries that contain the `_id` field. For these queries, MongoDB will use default index on the `_id` field and skip all query plan analysis.

`serverStatus.metrics.operation.scanAndOrder`

`scanAndOrder` (page 361) reports the total number of queries that return sorted numbers that cannot perform the sort operation using an index.

`serverStatus.metrics.queryExecutor`

`queryExecutor` (page 361) is a document that reports data from the query execution system.

`serverStatus.metrics.queryExecutor.scanned`

`scanned` (page 361) reports the total number of index items scanned during queries and query-plan evaluation. This counter is the same as `nscanned` (page 83) in the output of `explain()` (page 80).

`serverStatus.metrics.record`

`record` (page 361) is a document that reports data related to record allocation in the on-disk memory files.

`serverStatus.metrics.record.moves`

`moves` (page 361) reports the total number of times documents move within the on-disk representation of the MongoDB data set. Documents move as a result of operations that increase the size of the document beyond their allocated record size.

`serverStatus.metrics.repl`

`repl` (page 361) holds a sub-document that reports metrics related to the replication process. `repl` (page 361) document appears on all `mongod` (page 503) instances, even those that aren't members of *replica sets*.

`serverStatus.metrics.repl.apply`

`apply` (page 361) holds a sub-document that reports on the application of operations from the replication *oplog*.

`serverStatus.metrics.repl.apply.batches`

`batches` (page 361) reports on the *oplog* application process on *secondaries* members of replica sets. See *replica-set-internals-multi-threaded-replication* for more information on the *oplog* application processes

`serverStatus.metrics.repl.apply.batches.num`

`num` (page 361) reports the total number of batches applied across all databases.

`serverStatus.metrics.repl.apply.batches.totalMillis`

`totalMillis` (page 361) reports the total amount of time the `mongod` (page 503) has spent applying operations from the *oplog*.

`serverStatus.metrics.repl.apply.ops`

`ops` (page 361) reports the total number of *oplog* operations applied.

`serverStatus.metrics.repl.buffer`

MongoDB buffers oplog operations from the replication sync source buffer before applying oplog entries in a batch. `buffer` (page 361) provides a way to track the oplog buffer. See *replica-set-internals-multi-threaded-replication* for more information on the oplog application process.

`serverStatus.metrics.repl.buffer.count`

`count` (page 362) reports the current number of operations in the oplog buffer.

`serverStatus.metrics.repl.buffer.maxSizeBytes`

`maxSizeBytes` (page 362) reports the maximum size of the buffer. This value is a constant setting in the `mongod` (page 503), and is not configurable.

`serverStatus.metrics.repl.buffer.sizeBytes`

`sizeBytes` (page 362) reports the current size of the contents of the oplog buffer.

`serverStatus.metrics.repl.network`

`network` (page 362) reports network use by the replication process.

`serverStatus.metrics.repl.network.bytes`

`bytes` (page 362) reports the total amount of data read from the replication sync source.

`serverStatus.metrics.repl.network.getmores`

`getmores` (page 362) reports on the `getmore` operations, which are requests for additional results from the oplog *cursor* as part of the oplog replication process.

`serverStatus.metrics.repl.network.getmores.num`

`num` (page 362) reports the total number of `getmore` operations, which are operations that request an additional set of operations from the replication sync source.

`serverStatus.metrics.repl.network.getmores.totalMillis`

`totalMillis` (page 362) reports the total amount of time required to collect data from `getmore` operations.

Note: This number can be quite large, as MongoDB will wait for more data even if the `getmore` operation does not initial return data.

`serverStatus.metrics.repl.network.ops`

`ops` (page 362) reports the total number of operations read from the replication source.

`serverStatus.metrics.repl.network.readersCreated`

`readersCreated` (page 362) reports the total number of oplog query processes created. MongoDB will create a new oplog query any time an error occurs in the connection, including a timeout, or a network operation. Furthermore, `readersCreated` (page 362) will increment every time MongoDB selects a new source fore replication.

`serverStatus.metrics.repl.oplog`

`oplog` (page 362) is a document that reports on the size and use of the *oplog* by this `mongod` (page 503) instance.

`serverStatus.metrics.repl.oplog.insert`

`insert` (page 362) is a document that reports insert operations into the *oplog*.

`serverStatus.metrics.repl.oplog.insert.num`

`num` (page 362) reports the total number of items inserted into the *oplog*.

`serverStatus.metrics.repl.oplog.insert.totalMillis`

`totalMillis` (page 362) reports the total amount of time spent for the `mongod` (page 503) to insert data into the *oplog*.

`serverStatus.metrics.repl.oplog.insertBytes`

`insertBytes` (page 362) the total size of documents inserted into the oplog.

`serverStatus.metrics.repl.preload`

`preload` (page 362) reports on the “pre-fetch” stage, where MongoDB loads documents and indexes into RAM to improve replication throughput.

See *replica-set-internals-multi-threaded-replication* for more information about the *pre-fetch* stage of the replication process.

`serverStatus.metrics.repl.preload.docs`

`docs` (page 363) is a sub-document that reports on the documents loaded into memory during the *pre-fetch* stage.

`serverStatus.metrics.repl.preload.docs.num`

`num` (page 363) reports the total number of documents loaded during the *pre-fetch* stage of replication.

`serverStatus.metrics.repl.preload.docs.totalMillis`

`totalMillis` (page 363) reports the total amount of time spent loading documents as part of the *pre-fetch* stage of replication.

`serverStatus.metrics.repl.preload.indexes`

`indexes` (page 363) is a sub-document that reports on the index items loaded into memory during the *pre-fetch* stage of replication.

See *replica-set-internals-multi-threaded-replication* for more information about the *pre-fetch* stage of replication.

`serverStatus.metrics.repl.preload.indexes.num`

`num` (page 363) reports the total number of index entries loaded by members before updating documents as part of the *pre-fetch* stage of replication.

`serverStatus.metrics.repl.preload.indexes.totalMillis`

`totalMillis` (page 363) reports the total amount of time spent loading index entries as part of the *pre-fetch* stage of replication.

`serverStatus.metrics.ttl`

`ttl` (page 363) is a sub-document that reports on the operation of the resource use of the `ttl` index process.

`serverStatus.metrics.ttl.deletedDocuments`

`deletedDocuments` (page 363) reports the total number of documents deleted from collections with a `ttl` index.

`serverStatus.metrics.ttl.passes`

`passes` (page 363) reports the number of times the background process removes documents from collections with a `ttl` index.

`serverStatus.metrics.cursor`

New in version 2.6.

The `cursor` (page 363) is a document that contains data regarding cursor state and use.

`serverStatus.metrics.cursor.timedOut`

New in version 2.6.

`timedOut` (page 363) provides the total number of cursors that have timed out since the server process started. If this number is large or growing at a regular rate, this may indicate an application error.

`serverStatus.metrics.cursor.open`

New in version 2.6.

The `open` (page 363) is an embedded document that contains data regarding open cursors.

`serverStatus.metrics.cursor.open.noTimeout`

New in version 2.6.

`noTimeout` (page 363) provides the number of open cursors with the option `DBQuery.Option.noTimeout` (page 78) set to prevent timeout after a period of inactivity.

`serverStatus.metrics.cursor.open.pinned`

New in version 2.6.

`serverStatus.metrics.cursor.open.pinned` (page 364) provides the number of “pinned” open cursors.

`serverStatus.metrics.cursor.open.total`

New in version 2.6.

`total` (page 364) provides the number of cursors that MongoDB is maintaining for clients. Because MongoDB exhausts unused cursors, typically this value small or zero. However, if there is a queue, stale tailable cursors, or a large number of operations this value may rise.

features

features

`features` (page 364) is an internal command that returns the build-level feature settings.

isSelf

_isSelf

`_isSelf` (page 364) is an internal command.

2.2.3 Internal Commands

Internal Commands

Name	Description
<code>handshake</code> (page 365)	Internal command.
<code>_recvChunkAbort</code> (page 365)	Internal command that supports chunk migrations in sharded clusters. Do not call directly.
<code>_recvChunkCommit</code> (page 365)	Internal command that supports chunk migrations in sharded clusters. Do not call directly.
<code>_recvChunkStart</code> (page 365)	Internal command that facilitates chunk migrations in sharded clusters.. Do not call directly.
<code>_recvChunkStatus</code> (page 365)	Internal command that returns data to support chunk migrations in sharded clusters. Do not call directly.
<code>_replSetFresh</code>	Internal command that supports replica set election operations.
<code>mapreduce.shardedfinish</code> (page 365)	Internal command that supports <i>map-reduce</i> in <i>sharded cluster</i> environments.
<code>_transferMods</code> (page 365)	Internal command that supports chunk migrations. Do not call directly.
<code>replSetHeartbeat</code> (page 366)	Internal command that supports replica set operations.
<code>replSetGetRBID</code> (page 366)	Internal command that supports replica set operations.
<code>_migrateClone</code> (page 366)	Internal command that supports chunk migration. Do not call directly.
<code>replSetElect</code> (page 366)	Internal command that supports replica set functionality.
<code>writeBacksQueued</code> (page 366)	Internal command that supports chunk migrations in sharded clusters.
<code>writebacklisten</code> (page 367)	Internal command that supports chunk migrations in sharded clusters.

handshake**handshake**

`handshake` (page 365) is an internal command.

recvChunkAbort**_recvChunkAbort**

`_recvChunkAbort` (page 365) is an internal command. Do not call directly.

recvChunkCommit**_recvChunkCommit**

`_recvChunkCommit` (page 365) is an internal command. Do not call directly.

recvChunkStart**_recvChunkStart**

`_recvChunkStart` (page 365) is an internal command. Do not call directly.

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed.

recvChunkStatus**_recvChunkStatus**

`_recvChunkStatus` (page 365) is an internal command. Do not call directly.

replSetFresh**replSetFresh**

`replSetFresh` (page 365) is an internal command that supports replica set functionality.

mapreduce.shardedfinish**mapreduce.shardedfinish**

Provides internal functionality to support *map-reduce* in *sharded* environments.

See also:

“`mapReduce` (page 208)”

transferMods**_transferMods**

`_transferMods` (page 365) is an internal command. Do not call directly.

replSetHeartbeat

replSetHeartbeat

`replSetHeartbeat` (page 366) is an internal command that supports replica set functionality.

replSetGetRBID

replSetGetRBID

`replSetGetRBID` (page 366) is an internal command that supports replica set functionality.

migrateClone

_migrateClone

`_migrateClone` (page 366) is an internal command. Do not call directly.

replSetElect

replSetElect

`replSetElect` (page 366) is an internal command that support replica set functionality.

writeBacksQueued

writeBacksQueued

`writeBacksQueued` (page 366) is an internal command that returns a document reporting there are operations in the write back queue for the given `mongos` (page 518) and information about the queues.

`writeBacksQueued.hasOpsQueued`
Boolean.

`hasOpsQueued` (page 366) is true if there are write Back operations queued.

`writeBacksQueued.totalOpsQueued`
Integer.

`totalOpsQueued` (page 366) reflects the number of operations queued.

`writeBacksQueued.queues`
Document.

`queues` (page 366) holds a sub-document where the fields are all write back queues. These field hold a document with two fields that reports on the state of the queue. The fields in these documents are:

`writeBacksQueued.queues.n`
`n` (page 366) reflects the size, by number of items, in the queues.

`writeBacksQueued.queues.minutesSinceLastCall`
The number of minutes since the last time the `mongos` (page 518) touched this queue.

The command document has the following prototype form:

```
{writeBacksQueued: 1}
```

To call `writeBacksQueued` (page 366) from the `mongo` (page 527) shell, use the following `db.runCommand()` (page 118) form:

```
db.runCommand({writeBacksQueued: 1})
```

Consider the following example output:

```
{
  "hasOpsQueued" : true,
  "totalOpsQueued" : 7,
  "queues" : {
    "50b4f09f6671b11ff1944089" : { "n" : 0, "minutesSinceLastCall" : 1 },
    "50b4f09fc332bflc5aeaaf59" : { "n" : 0, "minutesSinceLastCall" : 0 },
    "50b4f09f6671b1d51df98cb6" : { "n" : 0, "minutesSinceLastCall" : 0 },
    "50b4f0c67ccf1e5c6effb72e" : { "n" : 0, "minutesSinceLastCall" : 0 },
    "50b4faf12319f193cfdec0d1" : { "n" : 0, "minutesSinceLastCall" : 4 },
    "50b4f013d2c1f8d62453017e" : { "n" : 0, "minutesSinceLastCall" : 0 },
    "50b4f0f12319f193cfdec0d1" : { "n" : 0, "minutesSinceLastCall" : 1 }
  },
  "ok" : 1
}
```

writebacklisten

writebacklisten

`writebacklisten` (page 367) is an internal command.

2.2.4 Testing Commands

Testing Commands

Name	Description
<code>testDistLockWithSkew</code>	Internal command. Do not call this directly.
<code>testDistLockWithSyncCluster</code>	Internal command. Do not call this directly.
<code>captrunc</code> (page 368)	Internal command. Truncates capped collections.
<code>emptycapped</code> (page 368)	Internal command. Removes all documents from a capped collection.
<code>godinsert</code> (page 369)	Internal command for testing.
<code>_hashBSONElement</code> (page 369)	Internal command. Computes the MD5 hash of a BSON element.
<code>_journalLatencyTest</code>	Tests the time required to write and perform a file system sync for a file in the journal directory.
<code>sleep</code> (page 370)	Internal command for testing. Forces MongoDB to block all operations.
<code>replSetTest</code> (page 371)	Internal command for testing replica set functionality.
<code>forceerror</code> (page 371)	Internal command for testing. Forces a user assertion exception.
<code>skewClockCommand</code>	Internal command. Do not call this command directly.
<code>configureFailPoint</code> (page 372)	Internal command for testing. Configures failure points.

testDistLockWithSkew

`_testDistLockWithSkew`

`_testDistLockWithSkew` (page 367) is an internal command. Do not call directly.

Note: `_testDistLockWithSkew` (page 367) is an internal command that is not enabled by default. `_testDistLockWithSkew` (page 367) must be enabled by using `--setParameter`

`enableTestCommands=1` on the `mongod` (page 503) command line. `_testDistLockWithSkew` (page 367) cannot be enabled during run-time.

testDistLockWithSyncCluster

_testDistLockWithSyncCluster

`_testDistLockWithSyncCluster` (page 368) is an internal command. Do not call directly.

Note: `_testDistLockWithSyncCluster` (page 368) is an internal command that is not enabled by default. `_testDistLockWithSyncCluster` (page 368) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `_testDistLockWithSyncCluster` (page 368) cannot be enabled during run-time.

captrunc

Definition

captrunc

`captrunc` (page 368) is a command that truncates capped collections for diagnostic and testing purposes and is not part of the stable client facing API. The command takes the following form:

```
{ captrunc: "<collection>", n: <integer>, inc: <true|false> }.
```

`captrunc` (page 368) has the following fields:

field string captrunc The name of the collection to truncate.

field integer n The number of documents to remove from the collection.

field boolean inc Specifies whether to truncate the nth document.

Note: `captrunc` (page 368) is an internal command that is not enabled by default. `captrunc` (page 368) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `captrunc` (page 368) cannot be enabled during run-time.

Examples The following command truncates 10 older documents from the collection `records`:

```
db.runCommand({captrunc: "records" , n: 10})
```

The following command truncates 100 documents and the 101st document:

```
db.runCommand({captrunc: "records", n: 100, inc: true})
```

emptycapped

emptycapped

The `emptycapped` command removes all documents from a capped collection. Use the following syntax:

```
{ emptycapped: "events" }
```

This command removes all records from the capped collection named `events`.

Warning: This command obtains a write lock on the affected database and will block other operations until it has completed.

Note: `emptycapped` (page 368) is an internal command that is not enabled by default. `emptycapped` (page 368) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `emptycapped` (page 368) cannot be enabled during run-time.

godinsert

godinsert

`godinsert` (page 369) is an internal command for testing purposes only.

Note: This command obtains a write lock on the affected database and will block other operations until it has completed.

Note: `godinsert` (page 369) is an internal command that is not enabled by default. `godinsert` (page 369) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `godinsert` (page 369) cannot be enabled during run-time.

_hashBSONElement

Description

_hashBSONElement

New in version 2.4.

An internal command that computes the MD5 hash of a BSON element. The `_hashBSONElement` (page 369) command returns 8 bytes from the 16 byte MD5 hash.

The `_hashBSONElement` (page 369) command has the following form:

```
db.runCommand({ _hashBSONElement: <key> , seed: <seed> })
```

The `_hashBSONElement` (page 369) command has the following fields:

field BSONElement key The BSON element to hash.

field integer seed A seed used when computing the hash.

Note: `_hashBSONElement` (page 369) is an internal command that is not enabled by default. `_hashBSONElement` (page 369) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `_hashBSONElement` (page 369) cannot be enabled during run-time.

Output The `_hashBSONElement` (page 369) command returns a document that holds the following fields:

`_hashBSONElement.key`

The original BSON element.

`_hashBSONElement.seed`

The seed used for the hash, defaults to 0.

`_hashBSONElement.out`

The decimal result of the hash.

`_hashBSONElement.ok`

Holds the 1 if the function returns successfully, and 0 if the operation encountered an error.

Example Invoke a `mongod` (page 503) instance with test commands enabled:

```
mongod --setParameter enableTestCommands=1
```

Run the following to compute the hash of an `ISODate` string:

```
db.runCommand({_hashBSONElement: ISODate("2013-02-12T22:12:57.211Z")})
```

The command returns the following document:

```
{
  "key" : ISODate("2013-02-12T22:12:57.211Z"),
  "seed" : 0,
  "out" : NumberLong("-4185544074338741873"),
  "ok" : 1
}
```

Run the following to hash the same `ISODate` string but this time to specify a seed value:

```
db.runCommand({_hashBSONElement: ISODate("2013-02-12T22:12:57.211Z"), seed:2013})
```

The command returns the following document:

```
{
  "key" : ISODate("2013-02-12T22:12:57.211Z"),
  "seed" : 2013,
  "out" : NumberLong("7845924651247493302"),
  "ok" : 1
}
```

journalLatencyTest

journalLatencyTest

`journalLatencyTest` (page 370) is an administrative command that tests the length of time required to write and perform a file system sync (e.g. *fsync*) for a file in the journal directory. You must issue the `journalLatencyTest` (page 370) command against the *admin database* in the form:

```
{ journalLatencyTest: 1 }
```

The value (i.e. 1 above), does not affect the operation of the command.

Note: `journalLatencyTest` (page 370) is an internal command that is not enabled by default. `journalLatencyTest` (page 370) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `journalLatencyTest` (page 370) cannot be enabled during run-time.

sleep

Definition

sleep

Forces the database to block all operations. This is an internal command for testing purposes.

The `sleep` (page 370) command takes the following prototype form:

```
{ sleep: 1, w: <true|false>, secs: <seconds> }
```

The `sleep` (page 370) command has the following fields:

field boolean w If true, obtains a global write lock. Otherwise obtains a read lock.

field integer secs The number of seconds to sleep.

Behavior The command places the `mongod` (page 503) instance in a *write lock* state for 100 seconds. Without arguments, `sleep` (page 370) causes a “read lock” for 100 seconds.

Warning: `sleep` (page 370) claims the lock specified in the `w` argument and blocks *all* operations on the `mongod` (page 503) instance for the specified amount of time.

Note: `sleep` (page 370) is an internal command that is not enabled by default. `sleep` (page 370) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `sleep` (page 370) cannot be enabled during run-time.

replSetTest**replSetTest**

`replSetTest` (page 371) is internal diagnostic command used for regression tests that supports replica set functionality.

Note: `replSetTest` (page 371) is an internal command that is not enabled by default. `replSetTest` (page 371) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `replSetTest` (page 371) cannot be enabled during run-time.

forceerror**forceerror**

The `forceerror` (page 371) command is for testing purposes only. Use `forceerror` (page 371) to force a user assertion exception. This command always returns an `ok` value of 0.

skewClockCommand**_skewClockCommand**

`_skewClockCommand` (page 371) is an internal command. Do not call directly.

Note: `_skewClockCommand` (page 371) is an internal command that is not enabled by default. `_skewClockCommand` (page 371) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `_skewClockCommand` (page 371) cannot be enabled during run-time.

configureFailPoint

Definition

configureFailPoint

Configures a failure point that you can turn on and off while MongoDB runs. `configureFailPoint` (page 372) is an internal command for testing purposes that takes the following form:

```
{configureFailPoint: "<failure_point>", mode: <behavior> }
```

You must issue `configureFailPoint` (page 372) against the *admin database*. `configureFailPoint` (page 372) has the following fields:

field string configureFailPoint The name of the failure point.

field document,string mode Controls the behavior of a failure point. The possible values are `alwaysOn`, `off`, or a document in the form of `{times: n}` that specifies the number of times the failure point remains on before it deactivates. The maximum value for the number is a 32-bit signed integer.

field document data Passes in additional data for use in configuring the fail point. For example, to imitate a slow connection pass in a document that contains a delay time.

Note: `configureFailPoint` (page 372) is an internal command that is not enabled by default. `configureFailPoint` (page 372) must be enabled by using `--setParameter enableTestCommands=1` on the `mongod` (page 503) command line. `configureFailPoint` (page 372) cannot be enabled during run-time.

Example

```
db.adminCommand( { configureFailPoint: "blocking_thread", mode: {times: 21} } )
```

2.2.5 Auditing Commands

System Events Auditing Commands

Name	Description
<code>logApplicationMessage</code> (page 372)	Posts a custom message to the audit log.

logApplicationMessage

logApplicationMessage

The `logApplicationMessage` (page 372) command allows users to post a custom message to the audit log. If running with authorization, users must have `clusterAdmin` role, or roles that inherit from `clusterAdmin`, to run the command.

Note: The audit system is available only in MongoDB Enterprise¹⁸.

The `logApplicationMessage` (page 372) has the following syntax:

```
{ logApplicationMessage: <string> }
```

MongoDB associates these custom messages with the *audit operation* `applicationMessage`, and the messages are subject to any *filtering*.

¹⁸<http://www.mongodb.com/products/mongodb-enterprise>

2.3 Operators

Query and Projection Operators (page 373) Query operators provide ways to locate data within the database and projection operators modify how data is presented.

Update Operators (page 412) Update operators are operators that enable you to modify the data in your database or add additional data.

Aggregation Framework Operators (page 437) Aggregation pipeline operations have a collection of operators available to define and manipulate documents in pipeline stages.

Query Modifiers (page 477) Query modifiers determine the way that queries will be executed.

2.3.1 Query and Projection Operators

Query Selectors

Comparison

Comparison Query Operators

Name	Description
<code>\$gt</code> (page 373)	Matches values that are greater than the value specified in the query.
<code>\$gte</code> (page 374)	Matches values that are greater than or equal to the value specified in the query.
<code>\$in</code> (page 374)	Matches any of the values that exist in an array specified in the query.
<code>\$lt</code> (page 375)	Matches values that are less than the value specified in the query.
<code>\$lte</code> (page 375)	Matches values that are less than or equal to the value specified in the query.
<code>\$ne</code> (page 376)	Matches all values that are not equal to the value specified in the query.
<code>\$nin</code> (page 376)	Matches values that do not exist in an array specified to the query.

`$gt` `$gt`

Syntax: `{field: { $gt: value } }`

`$gt` (page 373) selects those documents where the value of the `field` is greater than (i.e. `>`) the specified value.

Consider the following example:

```
db.inventory.find( { qty: { $gt: 20 } } )
```

This query will select all documents in the `inventory` collection where the `qty` field value is greater than 20.

Consider the following example that uses the `$gt` (page 373) operator with a field from an embedded document:

```
db.inventory.update( { "carrier.fee": { $gt: 2 } }, { $set: { price: 9.99 } } )
```

This `update()` (page 69) operation will set the value of the `price` field in the first document found containing the embedded document `carrier` whose `fee` field value is greater than 2.

To set the value of the `price` field in *all* documents containing the embedded document `carrier` whose `fee` field value is greater than 2, specify the `multi:true` option in the `update()` (page 69) method:

```
db.inventory.update( { "carrier.fee": { $gt: 2 } },
  { $set: { price: 9.99 },
    { multi: true }
  )
```

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

\$gte

Syntax: {field: { \$gte: value } }

`$gte` (page 374) selects the documents where the value of the `field` is greater than or equal to (i.e. `>=`) a specified value (e.g. `value`.)

Consider the following example:

```
db.inventory.find( { qty: { $gte: 20 } } )
```

This query would select all documents in `inventory` where the `qty` field value is greater than or equal to 20.

Consider the following example which uses the `$gte` (page 374) operator with a field from an embedded document:

```
db.inventory.update( { "carrier.fee": { $gte: 2 } }, { $set: { price: 9.99 } } )
```

This `update()` (page 69) operation will set the value of the `price` field that contain the embedded document `carrier` whose `fee` field value is greater than or equal to 2.

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

\$in

The `$in` (page 374) operator selects the documents where the value of a field equals any value in the specified array. To specify an `$in` (page 374) expression, use the following prototype:

```
{ field: { $in: [<value1>, <value2>, ... <valueN> ] } }
```

If the `field` holds an array, then the `$in` (page 374) operator selects the documents whose `field` holds an array that contains at least one element that matches a value in the specified array (e.g. `<value1>`, `<value2>`, etc.)

Changed in version 2.6: MongoDB 2.6 removes the combinatorial limit for the `$in` (page 374) operator that exists for [earlier versions](http://docs.mongodb.org/v2.4/reference/operator/query/in)¹⁹ of the operator.

Examples

Use the \$in Operator to Match Values Consider the following example:

```
db.inventory.find( { qty: { $in: [ 5, 15 ] } } )
```

This query selects all documents in the `inventory` collection where the `qty` field value is either 5 or 15. Although you can express this query using the `$or` (page 377) operator, choose the `$in` (page 374) operator rather than the `$or` (page 377) operator when performing equality checks on the same field.

¹⁹<http://docs.mongodb.org/v2.4/reference/operator/query/in>

Use the \$in Operator to Match Values in an Array The collection `inventory` contains documents that include the field `tags`, as in the following:

```
{ _id: 1, item: "abc", qty: 10, tags: [ "school", "clothing" ], sale: false }
```

Then, the following `update()` (page 69) operation will set the `sale` field value to `true` where the `tags` field holds an array with at least one element matching either `"appliances"` or `"school"`.

```
db.inventory.update(
  { tags: { $in: [ "appliances", "school" ] } },
  { $set: { sale: true } }
)
```

Use the \$in Operator with a Regular Expression The `$in` (page 374) operator can specify matching values using regular expressions or the `$regex` (page 386) operator expressions.

Consider the following example:

```
db.inventory.find( { tags: { $in: [ /^be/, /^st/ ] } } )
```

This query selects all documents in the `inventory` collection where the `tags` field holds an array that contains at least one element that starts with either `be` or `st`.

See also:

`find()` (page 34), `update()` (page 69), `$or` (page 377), `$set` (page 416).

\$lt

\$lt

Syntax: { field: { \$lt: value } }

`$lt` (page 375) selects the documents where the value of the `field` is less than (i.e. `<`) the specified value.

Consider the following example:

```
db.inventory.find( { qty: { $lt: 20 } } )
```

This query will select all documents in the `inventory` collection where the `qty` field value is less than 20.

Consider the following example which uses the `$lt` (page 375) operator with a field from an embedded document:

```
db.inventory.update( { "carrier.fee": { $lt: 20 } }, { $set: { price: 9.99 } } )
```

This `update()` (page 69) operation will set the `price` field value in the documents that contain the embedded document `carrier` whose `fee` field value is less than 20.

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

\$lte

\$lte

Syntax: { field: { \$lte: value } }

`$lte` (page 375) selects the documents where the value of the `field` is less than or equal to (i.e. `<=`) the specified value.

Consider the following example:

```
db.inventory.find( { qty: { $lte: 20 } } )
```

This query will select all documents in the `inventory` collection where the `qty` field value is less than or equal to 20.

Consider the following example which uses the `$lte` (page 375) operator with a field from an embedded document:

```
db.inventory.update( { "carrier.fee": { $lte: 5 } }, { $set: { price: 9.99 } } )
```

This `update()` (page 69) operation will set the `price` field value in the documents that contain the embedded document `carrier` whose `fee` field value is less than or equal to 5.

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

\$ne

\$ne

Syntax: { field: { \$ne: value } }

`$ne` (page 376) selects the documents where the value of the `field` is not equal (i.e. `!=`) to the specified value. This includes documents that do not contain the `field`.

Consider the following example:

```
db.inventory.find( { qty: { $ne: 20 } } )
```

This query will select all documents in the `inventory` collection where the `qty` field value does not equal 20, including those documents that do not contain the `qty` field.

Consider the following example which uses the `$ne` (page 376) operator with a field in an embedded document:

```
db.inventory.update( { "carrier.state": { $ne: "NY" } }, { $set: { qty: 20 } } )
```

This `update()` (page 69) operation will set the `qty` field value in the documents that contain the embedded document `carrier` whose `state` field value does not equal “NY”, or where the `state` field or the `carrier` embedded document do not exist.

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

\$nin

\$nin

Syntax: { field: { \$nin: [<value1>, <value2> ... <valueN>] } }

`$nin` (page 376) selects the documents where:

- the `field` value is not in the specified array **or**
- the `field` does not exist.

Consider the following query:

```
db.inventory.find( { qty: { $nin: [ 5, 15 ] } } )
```

This query will select all documents in the `inventory` collection where the `qty` field value does **not** equal 5 nor 15. The selected documents will include those documents that do *not* contain the `qty` field.

If the `field` holds an array, then the `$nin` (page 376) operator selects the documents whose `field` holds an array with **no** element equal to a value in the specified array (e.g. `<value1>`, `<value2>`, etc.).

Consider the following query:

```
db.inventory.update( { tags: { $nin: [ "appliances", "school" ] } }, { $set: { sale: false } } )
```

This `update()` (page 69) operation will set the `sale` field value in the `inventory` collection where the `tags` field holds an array with **no** elements matching an element in the array `["appliances", "school"]` or where a document does not contain the `tags` field.

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416).

Logical

Logical Query Operators

Name	Description
<code>\$or</code> (page 377)	Joins query clauses with a logical OR returns all documents that match the conditions clause.
<code>\$and</code> (page 378)	Joins query clauses with a logical AND returns all documents that match the condition clauses.
<code>\$not</code> (page 379)	Inverts the effect of a query expression and returns documents that do <i>not</i> match the expression.
<code>\$nor</code> (page 380)	Joins query clauses with a logical NOR returns all documents that fail to match both clauses.

`$or`

`$or`

The `$or` (page 377) operator performs a logical OR operation on an array of *two or more* `<expressions>` and selects the documents that satisfy *at least* one of the `<expressions>`. The `$or` (page 377) has the following syntax:

```
{ $or: [ { <expression1> }, { <expression2> }, ... , { <expressionN> } ] }
```

Consider the following example:

```
db.inventory.find( { $or: [ { quantity: { $lt: 20 } }, { price: 10 } ] } )
```

This query will select all documents in the `inventory` collection where either the `quantity` field value is less than 20 **or** the `price` field value equals 10.

Behaviors

`$or` Clauses and Indexes When evaluating the clauses in the `$or` (page 377) expression, MongoDB either performs a collection scan or, if all the clauses are supported by indexes, MongoDB performs index scans. That is, for MongoDB to use indexes to evaluate an `$or` (page 377) expression, all the clauses in the `$or` (page 377) expression must be supported by indexes. Otherwise, MongoDB will perform a collection scan.

When using indexes with `$or` (page 377) queries, each clause of an `$or` (page 377) will execute in parallel. These clauses can each use their own index. Consider the following query:

```
db.inventory.find( { $or: [ { quantity: { $lt: 20 } }, { price: 10 } ] } )
```

To support this query, rather than a compound index, you would create one index on `quantity` and another index on `price`:

```
db.inventory.ensureIndex( { quantity: 1 } )
db.inventory.ensureIndex( { price: 1 } )
```

MongoDB can use all but the `geoHaystack` index to support `$or` (page 377) clauses.

\$or and text Queries Changed in version 2.6.

If `$or` (page 377) includes a `$text` (page 387) query, all clauses in the `$or` (page 377) array must be supported by an index. This is because a `$text` (page 387) query *must* use an index, and `$or` (page 377) can only use indexes if all its clauses are supported by indexes. If the `$text` (page 387) query cannot use an index, the query will return an error.

As each clause of an `$or` (page 377) will execute in parallel, each clause can have a separate index.

\$or and GeoSpatial Queries Changed in version 2.6.

`$or` (page 377) supports *geospatial clauses* (page 392) with the following exception for the near clause (near clause includes `$nearSphere` (page 396) and `$near` (page 394)). `$or` (page 377) cannot contain a near clause with any other clause.

\$or and Sort Operations Changed in version 2.6.

When executing `$or` (page 377) queries with a `sort()` (page 93), MongoDB can now use indexes that support the `$or` (page 377) clauses. Previous versions did not use the indexes.

\$or versus \$in When using `$or` (page 377) with `<expressions>` that are equality checks for the value of the same field, use the `$in` (page 374) operator instead of the `$or` (page 377) operator.

For example, to select all documents in the `inventory` collection where the `quantity` field value equals either 20 or 50, use the `$in` (page 374) operator:

```
db.inventory.find ( { quantity: { $in: [20, 50] } } )
```

Nested \$or Clauses You may nest `$or` (page 377) operations.

See also:

`$and` (page 378), `find()` (page 34), `sort()` (page 93), `$in` (page 374)

\$and

\$and

New in version 2.0.

Syntax: { \$and: [{ <expression1> }, { <expression2> } , ... , { <expressionN> }] }

`$and` (page 378) performs a logical AND operation on an array of *two or more* expressions (e.g. `<expression1>`, `<expression2>`, etc.) and selects the documents that satisfy *all* the expressions in the array. The `$and` (page 378) operator uses *short-circuit evaluation*. If the first expression (e.g. `<expression1>`) evaluates to `false`, MongoDB will not evaluate the remaining expressions.

Note: MongoDB provides an implicit AND operation when specifying a comma separated list of expressions.

Using an explicit AND with the `$and` (page 378) operator is necessary when the same field or operator has to be specified in multiple expressions.

Examples

AND Queries With Multiple Expressions Specifying the Same Field Consider the following example:

```
db.inventory.find( { $and: [ { price: { $ne: 1.99 } }, { price: { $exists: true } } ] } )
```

This query will select all documents in the `inventory` collection where:

- the `price` field value is not equal to 1.99 **and**
- the `price` field exists.

This query can also be constructed with an implicit AND operation by combining the operator expressions for the `price` field. For example, this query can be written as:

```
db.inventory.find( { price: { $ne: 1.99, $exists: true } } )
```

AND Queries With Multiple Expressions Specifying the Same Operator Consider the following example:

```
db.inventory.find( {
  $and : [
    { $or : [ { price : 0.99 }, { price : 1.99 } ] },
    { $or : [ { sale : true }, { qty : { $lt : 20 } } ] }
  ]
} )
```

This query will return all select all documents where:

- the `price` field value equals 0.99 or 1.99, **and**
- the `sale` field value is equal to `true` **or** the `qty` field value is less than 20.

This query cannot be constructed using an implicit AND operation, because it uses the `$or` (page 377) operator more than once.

See also:

`find()` (page 34), `update()` (page 69), `$ne` (page 376), `$exists` (page 381), `$set` (page 416).

\$not

\$not

Syntax: { field: { \$not: { <operator-expression> } } }

`$not` (page 379) performs a logical NOT operation on the specified <operator-expression> and selects the documents that do *not* match the <operator-expression>. This includes documents that do not contain the field.

Consider the following query:

```
db.inventory.find( { price: { $not: { $gt: 1.99 } } } )
```

This query will select all documents in the `inventory` collection where:

- the `price` field value is less than or equal to 1.99 **or**
- the `price` field does not exist

`{ $not: { $gt: 1.99 } }` is different from the `$lte` (page 375) operator. `{ $lte: 1.99 }` returns *only* the documents where `price` field exists and its value is less than or equal to 1.99.

Remember that the `$not` (page 379) operator only affects *other operators* and cannot check fields and documents independently. So, use the `$not` (page 379) operator for logical disjunctions and the `$ne` (page 376) operator to test the contents of fields directly.

Consider the following behaviors when using the `$not` (page 379) operator:

- The operation of the `$not` (page 379) operator is consistent with the behavior of other operators but may yield unexpected results with some data types like arrays.
- The `$not` (page 379) operator does **not** support operations with the `$regex` (page 386) operator. Instead use <http://docs.mongodb.org/manual/> or in your driver interfaces, use your language's regular expression capability to create regular expression objects.

Consider the following example which uses the pattern match expression <http://docs.mongodb.org/manual/>:

```
db.inventory.find( { item: { $not: /^p.*/ } } )
```

The query will select all documents in the `inventory` collection where the `item` field value does *not* start with the letter `p`.

If you are using Python, you can write the above query with the PyMongo driver and Python's `python:re.compile()` method to compile a regular expression, as follows:

```
import re
for noMatch in db.inventory.find( { "item": { "$not": re.compile("^p.*") } } ):
    print noMatch
```

See also:

`find()` (page 34), `update()` (page 69), `$set` (page 416), `$gt` (page 373), `$regex` (page 386), [PyMongo²⁰](#), *driver*.

`$nor`

`$nor`

`$nor` (page 380) performs a logical NOR operation on an array of one or more query expression and selects the documents that **fail** all the query expressions in the array. The `$nor` (page 380) has the following syntax:

```
{ $nor: [ { <expression1> }, { <expression2> }, ... { <expressionN> } ] }
```

See also:

`find()` (page 34), `update()` (page 69), `$or` (page 377), `$set` (page 416), and `$exists` (page 381).

Examples

`$nor` Query with Two Expressions Consider the following query which uses only the `$nor` (page 380) operator:

```
db.inventory.find( { $nor: [ { price: 1.99 }, { sale: true } ] } )
```

This query will return all documents that:

- contain the `price` field whose value is *not* equal to 1.99 and contain the `sale` field whose value is *not* equal to `true` **or**

²⁰<http://api.mongodb.org/pythoncurrent>

- contain the `price` field whose value is *not* equal to `1.99` *but* do *not* contain the `sale` field **or**
- do *not* contain the `price` field *but* contain the `sale` field whose value is *not* equal to `true` **or**
- do *not* contain the `price` field *and* do *not* contain the `sale` field

\$nor and Additional Comparisons Consider the following query:

```
db.inventory.find( { $nor: [ { price: 1.99 }, { qty: { $lt: 20 } }, { sale: true } ] } )
```

This query will select all documents in the `inventory` collection where:

- the `price` field value does *not* equal `1.99` **and**
- the `qty` field value is *not* less than `20` **and**
- the `sale` field value is *not* equal to `true`

including those documents that do not contain these field(s).

The exception in returning documents that do not contain the field in the `$nor` (page 380) expression is when the `$nor` (page 380) operator is used with the `$exists` (page 381) operator.

\$nor and \$exists Compare that with the following query which uses the `$nor` (page 380) operator with the `$exists` (page 381) operator:

```
db.inventory.find( { $nor: [ { price: 1.99 }, { price: { $exists: false } },
                             { sale: true }, { sale: { $exists: false } } ] } )
```

This query will return all documents that:

- contain the `price` field whose value is *not* equal to `1.99` and contain the `sale` field whose value is *not* equal to `true`

Element

Element Query Operators	Name	Description
	<code>\$exists</code> (page 381)	Matches documents that have the specified field.
	<code>\$type</code> (page 383)	Selects documents if a field is of the specified type.

\$exists

Definition

\$exists

Syntax: { field: { \$exists: <boolean> } }

When <boolean> is `true`, `$exists` (page 381) matches the documents that contain the field, including documents where the field value is `null`. If <boolean> is `false`, the query returns only the documents that do not contain the field.

MongoDB `$exists` does **not** correspond to SQL operator `exists`. For SQL `exists`, refer to the `$in` (page 374) operator.

See also:

`$nin` (page 376), `$in` (page 374), and *faq-developers-query-for-nulls*.

Examples

Exists and Not Equal To Consider the following example:

```
db.inventory.find( { qty: { $exists: true, $nin: [ 5, 15 ] } } )
```

This query will select all documents in the `inventory` collection where the `qty` field exists *and* its value does not equal 5 or 15.

Null Values Given a collection named `records` with the following documents:

```
{ a: 5, b: 5, c: null }
{ a: 3, b: null, c: 8 }
{ a: null, b: 3, c: 9 }
{ a: 1, b: 2, c: 3 }
{ a: 2, c: 5 }
{ a: 3, b: 2 }
{ a: 4 }
{ b: 2, c: 4 }
{ b: 2 }
{ c: 6 }
```

Consider the output of the following queries:

Query:

```
db.records.find( { a: { $exists: true } } )
```

Result:

```
{ a: 5, b: 5, c: null }
{ a: 3, b: null, c: 8 }
{ a: null, b: 3, c: 9 }
{ a: 1, b: 2, c: 3 }
{ a: 2, c: 5 }
{ a: 3, b: 2 }
{ a: 4 }
```

Query:

```
db.records.find( { b: { $exists: false } } )
```

Result:

```
{ a: 2, c: 5 }
{ a: 4 }
{ c: 6 }
```

Query:

```
db.records.find( { c: { $exists: false } } )
```

Result:

```
{ a: 3, b: 2 }
{ a: 4 }
{ b: 2 }
```

\$type
\$type

Syntax: { field: { \$type: <BSON type> } }

\$type (page 383) selects the documents where the *value* of the `field` is the specified *BSON* type.

Consider the following example:

```
db.inventory.find( { price: { $type : 1 } } )
```

This query will select all documents in the `inventory` collection where the `price` field value is a Double.

If the `field` holds an array, the **\$type** (page 383) operator performs the type check against the array elements and **not** the `field`.

Consider the following example where the `tags` field holds an array:

```
db.inventory.find( { tags: { $type : 4 } } )
```

This query will select all documents in the `inventory` collection where the `tags` array contains an element that is itself an array.

If instead you want to determine whether the `tags` field is an array type, use the **\$where** (page 391) operator:

```
db.inventory.find( { $where : "Array.isArray(this.tags)" } )
```

See the [SERVER-1475](https://jira.mongodb.org/browse/SERVER-1475)²¹ for more information about the array type.

Refer to the following table for the available *BSON* types and their corresponding numbers.

Type	Number
Double	1
String	2
Object	3
Array	4
Binary data	5
Undefined (deprecated)	6
Object id	7
Boolean	8
Date	9
Null	10
Regular Expression	11
JavaScript	13
Symbol	14
JavaScript (with scope)	15
32-bit integer	16
Timestamp	17
64-bit integer	18
Min key	255
Max key	127

`MinKey` and `MaxKey` compare less than and greater than all other possible *BSON* element values, respectively, and exist primarily for internal use.

Note: To query if a field value is a `MinKey`, you must use the **\$type** (page 383) with `-1` as in the following example:

²¹<https://jira.mongodb.org/browse/SERVER-1475>

```
db.collection.find( { field: { $type: -1 } } )
```

Example

Consider the following example operation sequence that demonstrates both type comparison *and* the special MinKey and MaxKey values:

```
db.test.insert( {x : 3});
db.test.insert( {x : 2.9} );
db.test.insert( {x : new Date() } );
db.test.insert( {x : true } );
db.test.insert( {x : MaxKey } )
db.test.insert( {x : MinKey } )

db.test.find().sort({x:1})
{ "_id" : ObjectId("4b04094b7c65b846e2090112"), "x" : { $minKey : 1 } }
{ "_id" : ObjectId("4b03155dce8de6586fb002c7"), "x" : 2.9 }
{ "_id" : ObjectId("4b03154cce8de6586fb002c6"), "x" : 3 }
{ "_id" : ObjectId("4b031566ce8de6586fb002c9"), "x" : true }
{ "_id" : ObjectId("4b031563ce8de6586fb002c8"), "x" : "Tue Jul 25 2012 18:42:03 GMT-0500 (EST)" }
{ "_id" : ObjectId("4b0409487c65b846e2090111"), "x" : { $maxKey : 1 } }
```

To query for the minimum value of a *shard key* of a *sharded cluster*, use the following operation when connected to the `mongos` (page 518):

```
use config
db.chunks.find( { "min.shardKey": { $type: -1 } } )
```

Warning: Storing values of the different types in the same field in a collection is *strongly* discouraged.

See also:

`find()` (page 34), `insert()` (page 52), `$where` (page 391), *BSON*, *shard key*, *sharded cluster*.

Evaluation

Evaluation Query Operators

Name	Description
<code>\$mod</code> (page 384)	Performs a modulo operation on the value of a field and selects documents with result.
<code>\$regex</code> (page 386)	Selects documents where values match a specified regular expression.
<code>\$text</code> (page 387)	Performs text search.
<code>\$where</code> (page 391)	Matches documents that satisfy a JavaScript expression.

`$mod`

`$mod`

Select documents where the value of a field divided by a divisor has the specified remainder (i.e. perform a modulo operation to select documents). To specify a `$mod` (page 384) expression, use the following syntax:

```
{ field: { $mod: [ divisor, remainder ] } }
```

Changed in version 2.6: The `$mod` (page 384) operator errors when passed an array with fewer or more elements. In previous versions, if passed an array with one element, the `$mod` (page 384) operator uses 0 as the remainder value, and if passed an array with more than two elements, the `$mod` (page 384) ignores all but the first two elements. Previous versions do return an error when passed an empty array. See *Not Enough Elements Error* (page 385) and *Too Many Elements Error* (page 386) for details.

Examples

Use `$mod` to Select Documents Consider a collection `inventory` with the following documents:

```
{ "_id" : 1, "item" : "abc123", "qty" : 0 }
{ "_id" : 2, "item" : "xyz123", "qty" : 5 }
{ "_id" : 3, "item" : "ijk123", "qty" : 12 }
```

Then, the following query selects those documents in the `inventory` collection where value of the `qty` field modulo 4 equals 0:

```
db.inventory.find( { qty: { $mod: [ 4, 0 ] } } )
```

The query returns the following documents:

```
{ "_id" : 1, "item" : "abc123", "qty" : 0 }
{ "_id" : 3, "item" : "ijk123", "qty" : 12 }
```

Not Enough Elements Error The `$mod` (page 384) operator errors when passed an array with fewer than two elements.

Array with Single Element The following operation incorrectly passes the `$mod` (page 384) operator an array that contains a single element:

```
db.inventory.find( { qty: { $mod: [ 4 ] } } )
```

The statement results in the following error:

```
error: {
  "$err" : "bad query: BadValue malformed mod, not enough elements",
  "code" : 16810
}
```

Changed in version 2.6: In previous versions, if passed an array with one element, the `$mod` (page 384) operator uses the specified element as the divisor and 0 as the remainder value.

Empty Array The following operation incorrectly passes the `$mod` (page 384) operator an empty array:

```
db.inventory.find( { qty: { $mod: [ ] } } )
```

The statement results in the following error:

```
error: {
  "$err" : "bad query: BadValue malformed mod, not enough elements",
  "code" : 16810
}
```

Changed in version 2.6: Previous versions returned the following error:

```
error: { "$err" : "mod can't be 0", "code" : 10073 }
```

Too Many Elements Error The `$mod` (page 384) operator errors when passed an array with more than two elements.

For example, the following operation attempts to use the `$mod` (page 384) operator with an array that contains four elements:

```
error: {
  "$err" : "bad query: BadValue malformed mod, too many elements",
  "code" : 16810
}
```

Changed in version 2.6: In previous versions, if passed an array with more than two elements, the `$mod` (page 384) ignores all but the first two elements.

`$regex`

`$regex`

The `$regex` (page 386) operator provides regular expression capabilities for pattern matching *strings* in queries. MongoDB uses Perl compatible regular expressions (i.e. “PCRE.”)

You can specify regular expressions using regular expression objects or using the `$regex` (page 386) operator. The following examples are equivalent:

```
db.collection.find( { field: /acme.*corp/i } );
db.collection.find( { field: { $regex: 'acme.*corp', $options: 'i' } } );
```

These expressions match all documents in `collection` where the value of `field` matches the case-insensitive regular expression `acme.*corp`.

`$regex` (page 386) uses “Perl Compatible Regular Expressions” (PCRE) as the matching engine.

`$options`

`$regex` (page 386) provides four option flags:

- `i` toggles case insensitivity, and allows all letters in the pattern to match upper and lower cases.
- `m` toggles multiline regular expression. Without this option, all regular expression match within one line.

If there are no newline characters (e.g. `\n`) or no start/end of line construct, the `m` option has no effect.

- `x` toggles an “extended” capability. When set, `$regex` (page 386) ignores all white space characters unless escaped or included in a character class.

Additionally, it ignores characters between an un-escaped `#` character and the next new line, so that you may include comments in complicated patterns. This only applies to data characters; white space characters may never appear within special character sequences in a pattern.

The `x` option does not affect the handling of the VT character (i.e. code 11.)

New in version 1.9.0.

- `s` allows the dot (e.g. `.`) character to match all characters *including* newline characters.

`$regex` (page 386) only provides the `i` and `m` options for the native JavaScript regular expression objects (e.g. `http://docs.mongodb.org/manual/acme.*corp/i`). To use `x` and `s` you must use the “`$regex` (page 386)” operator with the “`$options` (page 386)” syntax.

To combine a regular expression match with other operators, you need to use the “`$regex` (page 386)” operator. For example:

```
db.collection.find( { field: { $regex: /acme.*corp/i, $nin: [ 'acmeblahcorp' ] } } );
```

This expression returns all instances of `field` in `collection` that match the case insensitive regular expression `acme.*corp` that *don't* match `acmeblahcorp`.

If an index exists for the field, then MongoDB matches the regular expression against the values in the index, which can be faster than a collection scan. Further optimization can occur if the regular expression is a “prefix expression”, which means that all potential matches start with the same string. This allows MongoDB to construct a “range” from that prefix and only match against those values from the index that fall within that range.

A regular expression is a “prefix expression” if it starts with a caret (^) or a left anchor (\A), followed by a string of simple symbols. For example, the regex `http://docs.mongodb.org/manual^abc.*` will be optimized by matching only against the values from the index that start with `abc`.

Additionally, while `http://docs.mongodb.org/manual^a/`, `http://docs.mongodb.org/manual^a.*`, and `http://docs.mongodb.org/manual^a.*$` match equivalent strings, they have different performance characteristics. All of these expressions use an index if an appropriate index exists; however, `http://docs.mongodb.org/manual^a.*`, and `http://docs.mongodb.org/manual^a.*$` are slower. `http://docs.mongodb.org/manual^a/` can stop scanning after matching the prefix.

\$text

\$text

New in version 2.6.

`$text` (page 387) performs a text search on the content of the fields indexed with a `text` index. A `$text` (page 387) expression has the following syntax:

```
{ $text: { $search: <string>, $language: <string> } }
```

The `$text` (page 387) operator accepts a text query document with the following fields:

field string `$search` A string of terms that MongoDB parses and uses to query the text index. MongoDB performs a logical OR search of the terms unless specified as a phrase. See *Behavior* (page 387) for more information on the field.

field string `$language` The language that determines the list of stop words for the search and the rules for the stemmer and tokenizer. If not specified, the search uses the default language of the index. For supported languages, see *text-search-languages*.

The `$text` (page 387) operator, by default, does *not* return results sorted in terms of the results' score. For more information, see the *Text Score* (page 388) documentation.

Behavior

Restrictions

- A query can specify, at most, one `$text` (page 387) expression.
- The `$text` (page 387) query can not appear in `$nor` (page 380) expressions.
- To use a `$text` (page 387) query in an `$or` (page 377) expression, all clauses in the `$or` (page 377) array must be indexed.
- You cannot use `hint()` (page 86) if the query includes a `$text` (page 387) query expression.

- You cannot specify `$natural` (page 483) sort order if the query includes a `$text` (page 387) expression.
- You cannot combine the `$text` (page 387) expression, which requires a special *text index*, with a query operator that requires a different type of special index. For example you cannot combine `$text` (page 387) expression with the `$near` (page 394) operator.

\$search Field In the `$search` field, specify a string of words that the `text` operator parses and uses to query the `text` index. The `text` operator treats most punctuation in the string as delimiters, except a hyphen – that negates term or an escaped double quotes `\` that specifies a phrase.

Phrases To match on a phrase, as opposed to individual terms, enclose the phrase in escaped double quotes (`\`), as in:

```
"\"ssl certificate\""
```

If the `$search` string includes a phrase and individual terms, text search will only match the documents that include the phrase. More specifically, the search performs a logical AND of the phrase with the individual terms in the search string.

For example, passed a `$search` string:

```
"\"ssl certificate\" authority key"
```

The `$text` (page 387) operator searches for the phrase `"ssl certificate"` **and** `("authority" or "key" or "ssl" or "certificate")`.

Negations Prefixing a word with a hyphen sign (`-`) negates a word:

- The negated word excludes documents that contain the negated word from the result set.
- When passed a search string that only contains negated words, text search will not match any documents.
- A hyphenated word, such as `pre-market`, is not a negation. The `$text` (page 387) operator treats the hyphen as a delimiter.

The `$text` (page 387) operator adds all negations to the query with the logical AND operator.

Match Operation The `$text` (page 387) operator ignores language-specific stop words, such as `the` and `and` in English.

The `$text` (page 387) operator matches on the complete *stemmed* word. So if a document field contains the word `blueberry`, a search on the term `blue` will not match. However, `blueberry` or `blueberries` will match.

For non-diacritics, text search is case insensitive; i.e. case insensitive for `[A-z]`.

Text Score The `$text` (page 387) operator assigns a score to each document that contains the search term in the indexed fields. The score represents the relevance of a document to a given text search query. The score can be part of a `sort()` (page 93) method specification as well as part of the projection expression. The `{ $meta: "textScore" }` expression provides information on the processing of the `$text` (page 387) operation. See `$meta` (page 410) projection operator for details on accessing the score for projection or sort.

Examples The following examples assume a collection `articles` that has a text index on the field `subject`:

```
db.articles.ensureIndex( { subject: "text" } )
```


Search for a Single Word The following query searches for the term `coffee`:

```
db.articles.find( { $text: { $search: "coffee" } } )
```

This query returns documents that contain the term `coffee` in the indexed `subject` field.

Match Any of the Search Terms If the search string is a space-delimited string, `$text` (page 387) operator performs a logical OR search on each term and returns documents that contains any of the terms.

The following query searches specifies a `$search` string of three terms delimited by space, "bake coffee cake":

```
db.articles.find( { $text: { $search: "bake coffee cake" } } )
```

This query returns documents that contain either `bake` **or** `coffee` **or** `cake` in the indexed `subject` field.

Search for a Phrase To match the exact phrase as a single term, escape the quotes.

The following query searches for the phrase `coffee cake`:

```
db.articles.find( { $text: { $search: "\"coffee cake\"" } } )
```

This query returns documents that contain the phrase `coffee cake`.

See also:

[Phrases](#) (page 388)

Exclude Documents That Contain a Term A *negated* term is a term that is prefixed by a minus sign `-`. If you negate a term, the `$text` (page 387) operator will exclude the documents that contain those terms from the results.

The following example searches for documents that contain the words `bake` or `coffee` but do **not** contain the term `cake`:

```
db.articles.find( { $text: { $search: "bake coffee -cake" } } )
```

See also:

[Negations](#) (page 388)

Return the Text Search Score The following query searches for the term `cake` and returns the score assigned to each matching document:

```
db.articles.find(
  { $text: { $search: "cake" } },
  { score: { $meta: "textScore" } }
)
```

In the result set, the returned documents includes an *additional* field `score` that contains the document's score associated with the text search. ²²

See also:

[Text Score](#) (page 388)

²² The behavior and requirements of the `$meta` (page 410) operator differs from that of the `$meta` (page 468) aggregation operator. See the `$meta` (page 468) aggregation operator for details.

Sort by Text Search Score To sort by the text score, include the **same** `$meta` (page 410) expression in **both** the projection document and the sort expression. ¹ The following query searches for the term `cake` and sorts the results by the descending score:

```
db.articles.find(
  { $text: { $search: "cake" } },
  { score: { $meta: "textScore" } }
).sort( { score: { $meta: "textScore" } } )
```

In the result set, the returned documents includes an additional field `score` that contains the document's score associated with the text search.

See also:

[Text Score](#) (page 388)

Return Top 3 Matching Documents Use the `limit()` (page 86) method in conjunction with a `sort()` (page 93) to return the top three matching documents. The following query searches for the term `cake` and sorts the results by the descending score:

```
db.articles.find(
  { $text: { $search: "cake" } },
  { score: { $meta: "textScore" } }
).sort( { score: { $meta: "textScore" } } ).limit(3)
```

See also:

[Text Score](#) (page 388)

Text Search with Additional Query and Sort Expressions The following query searches for documents with status equal to "A" that contain the terms `coffee` or `cake` in the indexed field `subject` and specifies a sort order of ascending date, descending text score:

```
db.articles.find(
  { status: "A", $text: { $search: "coffee cake" } },
  { score: { $meta: "textScore" } }
).sort( { date: 1, score: { $meta: "textScore" } } )
```

Search a Different Language Use the optional `$language` field in the `$text` (page 387) expression to specify a language that determines the list of stop words and the rules for the stemmer and tokenizer for the search string.

The following query specifies `es` for Spanish as the language that determines the tokenization, stemming, and stop words:

```
db.articles.find(
  { $text: { $search: "leche", $language: "es" } }
)
```

The `$text` (page 387) expression can also accept the language by name, `spanish`. See *text-search-languages* for the supported languages.

See also:

<http://docs.mongodb.org/manual/tutorial/text-search-in-aggregation>

\$where

\$where

Use the `$where` (page 391) operator to pass either a string containing a JavaScript expression or a full JavaScript function to the query system. The `$where` (page 391) provides greater flexibility, but requires that the database processes the JavaScript expression or function for *each* document in the collection. Reference the document in the JavaScript expression or function using either `this` or `obj`.

Warning:

- Do not write to the database within the `$where` (page 391) JavaScript function.
- `$where` (page 391) evaluates JavaScript and cannot take advantage of indexes. Therefore, query performance improves when you express your query using the standard MongoDB operators (e.g., `$gt` (page 373), `$in` (page 374)).
- In general, you should use `$where` (page 391) only when you can't express your query using another operator. If you must use `$where` (page 391), try to include at least one other standard query operator to filter the result set. Using `$where` (page 391) alone requires a table scan.

Consider the following examples:

```
db.myCollection.find( { $where: "this.credits == this.debits" } );
db.myCollection.find( { $where: "obj.credits == obj.debits" } );

db.myCollection.find( { $where: function() { return (this.credits == this.debits) } } );
db.myCollection.find( { $where: function() { return obj.credits == obj.debits; } } );
```

Additionally, if the query consists only of the `$where` (page 391) operator, you can pass in just the JavaScript expression or JavaScript functions, as in the following examples:

```
db.myCollection.find( "this.credits == this.debits || this.credits > this.debits" );

db.myCollection.find( function() { return (this.credits == this.debits || this.credits > this.debits) } );
```

You can include both the standard MongoDB operators and the `$where` (page 391) operator in your query, as in the following examples:

```
db.myCollection.find( { active: true, $where: "this.credits - this.debits < 0" } );
db.myCollection.find( { active: true, $where: function() { return obj.credits - obj.debits < 0; } } );
```

Using normal non-`$where` (page 391) query statements provides the following performance advantages:

- MongoDB will evaluate non-`$where` (page 391) components of query before `$where` (page 391) statements. If the non-`$where` (page 391) statements match no documents, MongoDB will not perform any query evaluation using `$where` (page 391).
- The non-`$where` (page 391) query statements may use an *index*.

Note: Changed in version 2.4.

In MongoDB 2.4, `map-reduce operations` (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions **cannot** access certain global functions or properties, such as `db`, that are available in the `mongo` (page 527) shell.

When upgrading to MongoDB 2.4, you will need to refactor your code if your `map-reduce operations` (page 208), `group` (page 204) commands, or `$where` (page 391) operator expressions include any global shell functions or properties that are no longer available, such as `db`.

The following JavaScript functions and properties **are available** to `map-reduce operations` (page 208), the `group` (page 204) command, and `$where` (page 391) operator expressions in MongoDB 2.4:

Available Properties	Available Functions	
args MaxKey MinKey	assert() BinData() DBPointer() DBRef() doassert() emit() gc() HexData() hex_md5() isNumber() isObject() ISODate() isString()	Map() MD5() NumberInt() NumberLong() ObjectId() print() printjson() printjsononeline() sleep() Timestamp() tojson() tojsononeline() tojsonObject() UUID() version()

Geospatial

Geospatial Query Operators

Operators

	Name	Description
Query Selectors	<code>\$geoWithin</code> (page 392)	Selects geometries within a bounding <i>GeoJSON</i> geometry.
	<code>\$geoIntersects</code> (page 393)	Selects geometries that intersect with a <i>GeoJSON</i> geometry.
	<code>\$near</code> (page 394)	Returns geospatial objects in proximity to a point.
	<code>\$nearSphere</code> (page 396)	Returns geospatial objects in proximity to a point on a sphere.

`$geoWithin`

`$geoWithin`

New in version 2.4: `$geoWithin` (page 392) replaces `$within` (page 393) which is deprecated.

The `$geoWithin` (page 392) operator is a geospatial query operator that queries for a defined point, line or shape that exists entirely within another defined shape. When determining inclusion, MongoDB considers the border of a shape to be part of the shape, subject to the precision of floating point numbers.

The `$geoWithin` (page 392) operator queries for inclusion in a *GeoJSON* polygon or a shape defined by legacy coordinate pairs.

The `$geoWithin` (page 392) operator does not return sorted results. As a result MongoDB can return `$geoWithin` (page 392) queries more quickly than geospatial `$near` (page 394) or `$nearSphere` (page 396) queries, which sort results.

The 2dsphere and 2d indexes both support the `$geoWithin` (page 392) operator.

Changed in version 2.2.3: `$geoWithin` (page 392) does not require a geospatial index. However, a geospatial index will improve query performance.

If querying for geometries that exist within a GeoJSON *polygon* on a sphere, pass the polygon to `$geoWithin` (page 392) using the `$geometry` (page 397) operator.

For a polygon with only an exterior ring use following syntax:

```
db.<collection>.find( { <location field> :
  { $geoWithin :
    { $geometry :
      { type : "Polygon" ,
        coordinates : [ [ [ <lng1>, <lat1> ] , [ <lng2>, <lat2> ] ... ]
      } } } } )
```

Important: Specify coordinates in longitude, latitude order.

For a polygon with an exterior and interior ring use following syntax:

```
db.<collection>.find( { <location field> :
  { $geoWithin :
    { $geometry :
      { type : "Polygon" ,
        coordinates : [ [ [ <lng1>, <lat1> ] , [ <lng2>, <lat2> ] ... ]
                      [ [ <lngA>, <latA> ] , [ <lngB>, <latB> ] ... ]
      } } } } )
```

The following example selects all indexed points and shapes that exist entirely within a GeoJSON polygon:

```
db.places.find( { loc :
  { $geoWithin :
    { $geometry :
      { type : "Polygon" ,
        coordinates: [ [ [ 0 , 0 ] , [ 3 , 6 ] , [ 6 , 1 ] , [ 0 , 0 ] ] ]
      } } } } )
```

If querying for inclusion in a shape defined by legacy coordinate pairs on a plane, use the following syntax:

```
db.<collection>.find( { <location field> :
  { $geoWithin :
    { <shape operator> : <coordinates>
    } } } )
```

For the syntax of shape operators, see: `$box` (page 399), `$polygon` (page 399), `$center` (page 397) (defines a circle), and `$centerSphere` (page 398) (defines a circle on a sphere).

Note: Any geometry specified with *GeoJSON* to `$geoWithin` (page 392) queries, **must** fit within a single hemisphere. MongoDB interprets geometries larger than half of the sphere as queries for the smaller of the complementary geometries.

\$within

Deprecated since version 2.4: `$geoWithin` (page 392) replaces `$within` (page 393) in MongoDB 2.4.

\$geoIntersects

\$geoIntersects

New in version 2.4.

The `$geoIntersects` (page 393) operator is a geospatial query operator that selects all locations that intersect with a *GeoJSON* object. A location intersects a GeoJSON object if the intersection is non-empty. This includes documents that have a shared edge. The `$geoIntersects` (page 393) operator uses spherical geometry.

The 2dsphere geospatial index supports `$geoIntersects` (page 393).

To query for intersection, pass the GeoJSON object to `$geoIntersects` (page 393) through the `$geometry` (page 397) operator. Use the following syntax:

```
db.<collection>.find( { <location field> :
                      { $geoIntersects :
                        { $geometry :
                          { type : "<GeoJSON object type>" ,
                            coordinates : [ <coordinates> ]
                          }
                        }
                      } } )
```

Important: Specify coordinates in this order: “**longitude, latitude.**”

The following example uses `$geoIntersects` (page 393) to select all indexed points and shapes that intersect with the polygon defined by the `coordinates` array.

```
db.places.find( { loc :
                  { $geoIntersects :
                    { $geometry :
                      { type : "Polygon" ,
                        coordinates: [ [ [ 0 , 0 ] , [ 3 , 6 ] , [ 6 , 1 ] , [ 0 , 0 ] ] ]
                      }
                    }
                  } } )
```

Note: Any geometry specified with *GeoJSON* to `$geoIntersects` (page 393) queries, **must** fit within a single hemisphere. MongoDB interprets geometries larger than half of the sphere as queries for the smaller of the complementary geometries.

`$near`

Definition

`$near`

Changed in version 2.4.

Specifies a point for which a *geospatial* query returns the closest documents first. The query sorts the documents from nearest to farthest.

The `$near` (page 394) operator can query for a *GeoJSON* point or for a point defined by legacy coordinate pairs.

The optional `$maxDistance` (page 397) operator limits a `$near` (page 394) query to return only those documents that fall within a maximum distance of a point. If you query for a GeoJSON point, specify `$maxDistance` (page 397) in meters. If you query for legacy coordinate pairs, specify `$maxDistance` (page 397) in radians.

Considerations The `$near` (page 394) operator requires a geospatial index: a 2dsphere index for GeoJSON points; a 2d index for legacy coordinate pairs. By default, queries that use a 2d index return a limit of 100 documents; however you may use `limit()` (page 86) to change the number of results.

You cannot combine the `$near` (page 394) operator, which requires a special *geospatial index*, with a query operator or command that uses a different type of special index. For example you cannot combine `$near` (page 394) with the `text` (page 235) command.

Behavior `geoNear` (page 216) always returns the documents sorted by distance. Any other sort order requires to sort the documents in memory, which can be inefficient. To return results in a different sort order, use the `$geoWithin` (page 392) operator in combination with `sort()`

Examples For queries on GeoJSON data, use the following syntax:

```
db.<collection>.find(
  { <location field> :
    { $near :
      {
        $geometry : {
          type : "Point" ,
          coordinates : [ <longitude> , <latitude> ] },
          $maxDistance : <distance in meters>
        }
      }
    }
  }
)
```

Important: Specify coordinates in this order: “**longitude, latitude.**”

The following example selects the documents with coordinates nearest to [40 , 5] and limits the maximum distance to 500 meters from the specified GeoJSON point:

```
db.places.find(
  {
    loc:
      { $near :
        {
          $geometry : { type : "Point" , coordinates: [ 40 , 5 ] },
          $maxDistance : 500
        }
      }
  }
)
```

For queries on legacy coordinate pairs, use the following syntax:

```
db.<collection>.find( { <location field> :
  { $near : [ <x> , <y> ] ,
    $maxDistance: <distance>
  } } )
```

Important: If you use longitude and latitude, specify **longitude first**.

The following example query returns documents with location values that are 10 or fewer units from the point [40 , 5].

For GeoJSON point object, specify the `$maxDistance` in meters, not radians.

```
db.places.find( { loc :
  { $near : [ 40 , 5 ] ,
    $maxDistance : 10
  } } )
```

Note: You can further limit the number of results using `cursor.limit()` (page 86).

Specifying a batch size (i.e. `batchSize()` (page 78)) in conjunction with queries that use the `$near` (page 394) is not defined. See [SERVER-5236](https://jira.mongodb.org/browse/SERVER-5236)²³ for more information.

`$nearSphere`

`$nearSphere`

New in version 1.8.

Specifies a point for which a *geospatial* query returns the closest documents first. The query sorts the documents from nearest to farthest. MongoDB calculates distances for `$nearSphere` (page 396) using spherical geometry.

The `$nearSphere` (page 396) operator queries for points defined by either *GeoJSON* objects or legacy coordinate pairs.

The optional `$maxDistance` (page 397) operator limits a `$nearSphere` (page 396) query to return only those documents that fall within a maximum distance of a point. If you use `$maxDistance` (page 397) on GeoJSON points, the distance is measured in meters. If you use `$maxDistance` (page 397) on legacy coordinate pairs, the distance is measured in radians.

The `$nearSphere` (page 396) operator requires a geospatial index. The `2dsphere` and `2d` indexes both support `$nearSphere` (page 396) with both legacy coordinate pairs and GeoJSON points. Queries that use a `2d` index return a at most 100 documents.

Important: If you use longitude and latitude, specify **longitude first**.

For queries on GeoJSON data, use the following syntax:

```
db.<collection>.find( { <location field> :
  { $nearSphere :
    { $geometry :
      { type : "Point" ,
        coordinates : [ <longitude> , <latitude> ] } ,
      $maxDistance : <distance in meters>
    }
  }
} )
```

For queries on legacy coordinate pairs, use the following syntax:

```
db.<collection>.find( { <location field> :
  { $nearSphere: [ <x> , <y> ] ,
    $maxDistance: <distance in radians>
  }
} )
```

The following example selects the 100 documents with legacy coordinates pairs nearest to [40 , 5], as calculated by spherical geometry:

```
db.places.find( { loc :
  { $nearSphere : [ 40 , 5 ] ,
    $maxDistance : 10
  }
} )
```

²³<https://jira.mongodb.org/browse/SERVER-5236>

Geometry Specifiers	Name	Description
	<code>\$geometry</code> (page 397)	Specifies a geometry in <i>GeoJSON</i> format to geospatial query operators.
	<code>\$maxDistance</code> (page 397)	Specifies a distance to limit the results of <code>\$near</code> (page 394) and <code>\$nearSphere</code> (page 396) queries.
	<code>\$center</code> (page 397)	Specifies a circle using legacy coordinate pairs to <code>\$geoWithin</code> (page 392) queries when using planar geometry.
	<code>\$centerSphere</code> (page 398)	Specifies a circle using either legacy coordinate pairs or <i>GeoJSON</i> format for <code>\$geoWithin</code> (page 392) queries when using spherical geometry.
	<code>\$box</code> (page 399)	Specifies a rectangular box using legacy coordinate pairs for <code>\$geoWithin</code> (page 392) queries.
	<code>\$polygon</code> (page 399)	Specifies a polygon to using legacy coordinate pairs for <code>\$geoWithin</code> (page 392) queries.
	<code>\$uniqueDocs</code> (page 400)	Deprecated. Modifies a <code>\$geoWithin</code> (page 392) and <code>\$near</code> (page 394) queries to ensure that even if a document matches the query multiple times, the query returns the document only once.

\$geometry**\$geometry**

New in version 2.4.

The `$geometry` (page 397) operator specifies a *GeoJSON* for a geospatial query operators. For details on using `$geometry` (page 397) with an operator, see the operator:

- `$geoWithin` (page 392)
- `$geoIntersects` (page 393)
- `$near` (page 394)

\$maxDistance**\$maxDistance**

The `$maxDistance` (page 397) operator constrains the results of a geospatial `$near` (page 394) or `$nearSphere` (page 396) query to the specified distance. The measuring units for the maximum distance are determined by the coordinate system in use. For *GeoJSON* point object, specify the distance in meters, not radians.

Changed in version 2.6: Specify a non-negative number for `$maxDistance` (page 397).

The 2d and 2dsphere geospatial indexes both support `$maxDistance` (page 397).

The following example query returns documents with location values that are 10 or fewer units from the point [100 , 100].

```
db.places.find( { loc : { $near : [ 100 , 100 ] ,
                             $maxDistance: 10 }
                } )
```

MongoDB orders the results by their distance from [100 , 100]. The operation returns the first 100 results, unless you modify the query with the `cursor.limit()` (page 86) method.

\$center**\$center**

New in version 1.4.

The `$center` (page 397) operator specifies a circle for a *geospatial* `$geoWithin` (page 392) query. The query returns legacy coordinate pairs that are within the bounds of the circle. The operator does *not* return GeoJSON objects.

The query calculates distances using flat (planar) geometry.

The 2d geospatial index supports the `$center` (page 397) operator.

To use the `$center` (page 397) operator, specify an array that contains:

- The grid coordinates of the circle's center point
- The circle's radius, as measured in the units used by the coordinate system

Important: If you use longitude and latitude, specify **longitude first**.

Use the following syntax:

```
{ <location field> : { $geoWithin : { $center : [ [ <x>, <y> ] , <radius> ] } } }
```

The following example query returns all documents that have coordinates that exist within the circle centered on [-74 , 40.74] and with a radius of 10:

```
db.places.find( { loc: { $geoWithin :  
                        { $center : [ [-74, 40.74], 10 ] }  
                      } } )
```

Changed in version 2.2.3: Applications can use `$center` (page 397) *without* having a geospatial index. However, geospatial indexes support much faster queries than the unindexed equivalents. Before 2.2.3, a geospatial index *must* exist on a field holding coordinates before using any of the geospatial query operators.

`$centerSphere`

`$centerSphere`

New in version 1.8.

The `$centerSphere` (page 398) operator defines a circle for a *geospatial* query that uses spherical geometry. The query returns documents that are within the bounds of the circle.

You can use the `$centerSphere` (page 398) operator on both *GeoJSON* objects and legacy coordinate pairs.

The 2d and 2dsphere geospatial indexes both support `$centerSphere` (page 398).

To use `$centerSphere` (page 398), specify an array that contains:

- The grid coordinates of the circle's center point
- The circle's radius measured in radians. To calculate radians, see <http://docs.mongodb.org/manual/tutorial/calculate-distances-using-spherical-geometry/>

Use the following syntax:

```
db.<collection>.find( { <location field> :  
                        { $geoWithin :  
                          { $centerSphere : [ [ <x>, <y> ] , <radius> ] }  
                        } } )
```

Important: If you use longitude and latitude, specify **longitude first**.

The following example queries grid coordinates and returns all documents within a 10 mile radius of longitude 88 W and latitude 30 N. The query converts the distance to radians by dividing by the approximate radius of the earth, 3959 miles:

```
db.places.find( { loc : { $geoWithin :
                      { $centerSphere :
                        [ [ 88 , 30 ] , 10 / 3959 ]
                      }
                    } } )
```

Changed in version 2.2.3: Applications can use `$centerSphere` (page 398) *without* having a geospatial index. However, geospatial indexes support much faster queries than the unindexed equivalents. Before 2.2.3, a geospatial index *must* exist on a field holding coordinates before using any of the geospatial query operators.

\$box

\$box

New in version 1.4.

The `$box` (page 399) operator specifies a rectangle for a *geospatial* `$geoWithin` (page 392) query. The query returns documents that are within the bounds of the rectangle, according to their point-based location data. The `$box` (page 399) operator returns documents based on *grid coordinates* and does *not* query for GeoJSON shapes.

The query calculates distances using flat (planar) geometry. The 2d geospatial index supports the `$box` (page 399) operator.

To use the `$box` (page 399) operator, you must specify the bottom left and top right corners of the rectangle in an array object. Use the following syntax:

```
{ <location field> : { $geoWithin : { $box :
                                   [ [ <bottom left coordinates> ] ,
                                   [ <upper right coordinates> ] ] } } }
```

Important: If you use longitude and latitude, specify **longitude first**.

The following example query returns all documents that are within the box having points at: [0 , 0], [0 , 100], [100 , 0], and [100 , 100].

```
db.places.find( { loc : { $geoWithin : { $box :
                                   [ [ 0 , 0 ] ,
                                   [ 100 , 100 ] ] } } } )
```

Changed in version 2.2.3: Applications can use `$box` (page 399) *without* having a geospatial index. However, geospatial indexes support much faster queries than the unindexed equivalents. Before 2.2.3, a geospatial index *must* exist on a field holding coordinates before using any of the geospatial query operators.

\$polygon

\$polygon

New in version 1.9.

The `$polygon` (page 399) operator specifies a polygon for a *geospatial* `$geoWithin` (page 392) query on legacy coordinate pairs. The query returns pairs that are within the bounds of the polygon. The operator does *not* query for GeoJSON objects.

The `$polygon` (page 399) operator calculates distances using flat (planar) geometry.

The 2d geospatial index supports the `$polygon` (page 399) operator.

To define the polygon, specify an array of coordinate points. Use the following syntax:

```
{ <location field> : { $geoWithin : { $polygon : [ [ <x1> , <y1> ] ,
                                                    [ <x2> , <y2> ] ,
                                                    [ <x3> , <y3> ] ] } } }
```

Important: If you use longitude and latitude, specify **longitude first**.

The last point specified is always implicitly connected to the first. You can specify as many points, and therefore sides, as you like.

The following query returns all documents that have coordinates that exist within the polygon defined by [0 , 0], [3 , 6], and [6 , 0]:

```
db.places.find( { loc : { $geoWithin : { $polygon : [ [ 0 , 0 ] ,  
                                                    [ 3 , 6 ] ,  
                                                    [ 6 , 0 ] ] } } } )
```

Changed in version 2.2.3: Applications can use `$polygon` (page 399) *without* having a geospatial index. However, geospatial indexes support much faster queries than the unindexed equivalents. Before 2.2.3, a geospatial index *must* exist on a field holding coordinates before using any of the geospatial query operators.

\$uniqueDocs

Definition

\$uniqueDocs

Deprecated since version 2.6: Geospatial queries no longer return duplicate results. The `$uniqueDocs` (page 400) operator has no impact on results.

Returns a document only once for a geospatial query even if the document matches the query multiple times.

Geospatial Query Compatibility While numerous combinations of query operators are possible, the following table shows the recommended operators for different types of queries. The table uses the `$geoWithin` (page 392), `$geoIntersects` (page 393) and `$near` (page 394) operators.

Query Document	Geometry of the Query Condition	Surface Type for Query Calculation	Units for Query Calculation	Supported by this Index
Returns points, lines and polygons				
<pre>{ \$geoWithin : { \$geometry : <GeoJSON Polygon> } }</pre>	polygon	sphere	meters	2dsphere
<pre>{ \$geoIntersects : { \$geometry : <GeoJSON> } }</pre>	point, line or polygon	sphere	meters	2dsphere
<pre>{ \$near : { \$geometry : <GeoJSON Point>, \$maxDistance : d } }</pre>	point	sphere	meters	2dsphere The index is required.
Returns points only				
<pre>{ \$geoWithin : { \$box : [[x1, y1], [x2, y2]] } }</pre>	rectangle	flat	flat units	2d
<pre>{ \$geoWithin : { \$polygon : [[x1, y1], [x1, y2], [x2, y2], [x2, y1]] } }</pre>	polygon	flat	flat units	2d
<pre>{ \$geoWithin : { \$center : [[x1, y1], r], } }</pre>	circular region	flat	flat units	2d
<pre>{ \$geoWithin : { \$centerSphere : [[x, y], radius] } }</pre>	circular region	sphere	radians	2d 2dsphere
<pre>{ \$near : [x1, y1], \$maxDistance : d }</pre>	point	flat / flat units	flat units	2d The index is required.

Array

Query Operator Array	Name	Description
	<code>\$all</code> (page 402)	Matches arrays that contain all elements specified in the query.
	<code>\$elemMatch</code> (page 405)	Selects documents if element in the array field matches all the specified <code>\$elemMatch</code> (page 405) condition.
	<code>\$size</code> (page 405)	Selects documents if the array field is a specified size.

`$all`

`$all`

The `$all` (page 402) operator selects the documents where the value of a field is an array that contains all the specified elements. To specify an `$all` (page 402) expression, use the following prototype:

```
{ <field>: { $all: [ <value1> , <value2> ... ] } }
```

Behavior

Equivalent to `$and` Operation

 Changed in version 2.6.

The `$all` (page 402) is equivalent to an `$and` (page 378) operation of the specified values; i.e. the following statement:

```
{ tags: { $all: [ "ssl" , "security" ] } }
```

is equivalent to:

```
{ $and: [ { tags: "ssl" }, { tags: "security" } ] }
```

Nested Array

 Changed in version 2.6.

When passed an array of a nested array (e.g. `[["A"]]`), `$all` (page 402) can now match documents where the field contains the nested array as an element (e.g. `field: [["A"], ...]`), or the field equals the nested array (e.g. `field: ["A"]`).

For example, consider the following query²⁴:

```
db.articles.find( { tags: { $all: [ [ "ssl", "security" ] ] } } )
```

The query is equivalent to:

```
db.articles.find( { $and: [ { tags: [ "ssl", "security" ] } ] } )
```

which is equivalent to:

```
db.articles.find( { tags: [ "ssl", "security" ] } )
```

As such, the `$all` (page 402) expression can match documents where the `tags` field is an array that contains the nested array `["ssl", "security"]` or is an array that equals the nested array:

```
tags: [ [ "ssl", "security" ], ... ]
tags: [ "ssl", "security" ]
```

This behavior for `$all` (page 402) allows for more matches than previous versions of MongoDB. Earlier version could only match documents where the field contains the nested array.

²⁴ The `$all` (page 402) expression with a *single* element is for illustrative purposes since the `$all` (page 402) expression is unnecessary if matching only a single element. Instead, when matching a single element, a “contains” expression (i.e. `arrayField: element`) is more suitable.

Performance Queries that use the `$all` (page 402) operator must scan all the documents that match the first element in the `$all` (page 402) expression. As a result, even with an index to support the query, the operation may be long running, particularly when the first element in the `$all` (page 402) expression is not very selective.

Examples The following examples use the `inventory` collection that contains the documents:

```
{
  _id: ObjectId("5234cc89687ea597eabee675"),
  code: "xyz",
  tags: [ "school", "book", "bag", "headphone", "appliance" ],
  qty: [
    { size: "S", num: 10, color: "blue" },
    { size: "M", num: 45, color: "blue" },
    { size: "L", num: 100, color: "green" }
  ]
}

{
  _id: ObjectId("5234cc8a687ea597eabee676"),
  code: "abc",
  tags: [ "appliance", "school", "book" ],
  qty: [
    { size: "6", num: 100, color: "green" },
    { size: "6", num: 50, color: "blue" },
    { size: "8", num: 100, color: "brown" }
  ]
}

{
  _id: ObjectId("5234ccb7687ea597eabee677"),
  code: "efg",
  tags: [ "school", "book" ],
  qty: [
    { size: "S", num: 10, color: "blue" },
    { size: "M", num: 100, color: "blue" },
    { size: "L", num: 100, color: "green" }
  ]
}

{
  _id: ObjectId("52350353b2eff1353b349de9"),
  code: "ijk",
  tags: [ "electronics", "school" ],
  qty: [
    { size: "M", num: 100, color: "green" }
  ]
}
```

Use `$all` to Match Values The following operation uses the `$all` (page 402) operator to query the `inventory` collection for documents where the value of the `tags` field is an array whose elements include `appliance`, `school`, and `book`:

```
db.inventory.find( { tags: { $all: [ "appliance", "school", "book" ] } } )
```

The above query returns the following documents:

```
{
  _id: ObjectId("5234cc89687ea597eabee675"),
  code: "xyz",
  tags: [ "school", "book", "bag", "headphone", "appliance" ],
  qty: [
    { size: "S", num: 10, color: "blue" },
    { size: "M", num: 45, color: "blue" },
    { size: "L", num: 100, color: "green" }
  ]
}

{
  _id: ObjectId("5234cc8a687ea597eabee676"),
  code: "abc",
  tags: [ "appliance", "school", "book" ],
  qty: [
    { size: "6", num: 100, color: "green" },
    { size: "6", num: 50, color: "blue" },
    { size: "8", num: 100, color: "brown" }
  ]
}
```

Use \$all with \$elemMatch If the field contains an array of documents, you can use the `$all` (page 402) with the `$elemMatch` (page 405) operator.

The following operation queries the `inventory` collection for documents where the value of the `qty` field is an array whose elements match the `$elemMatch` (page 405) criteria:

```
db.inventory.find( {
  qty: { $all: [
    { "$elemMatch" : { size: "M", num: { $gt: 50 } } },
    { "$elemMatch" : { num : 100, color: "green" } }
  ] }
} )
```

The query returns the following documents:

```
{
  "_id" : ObjectId("5234ccb7687ea597eabee677"),
  "code" : "efg",
  "tags" : [ "school", "book"],
  "qty" : [
    { "size" : "S", "num" : 10, "color" : "blue" },
    { "size" : "M", "num" : 100, "color" : "blue" },
    { "size" : "L", "num" : 100, "color" : "green" }
  ]
}

{
  "_id" : ObjectId("52350353b2eff1353b349de9"),
  "code" : "ijk",
  "tags" : [ "electronics", "school" ],
  "qty" : [
    { "size" : "M", "num" : 100, "color" : "green" }
  ]
}
```

The `$all` (page 402) operator exists to support queries on arrays. But you may use the `$all` (page 402) operator to

select against a non-array field, as in the following example:

```
db.inventory.find( { qty: { $all: [ 50 ] } } )
```

However, use the following form to express the same query:

```
db.inventory.find( { qty: 50 } )
```

Both queries will select all documents in the `inventory` collection where the value of the `qty` field equals 50.

Note: In most cases, MongoDB does not treat arrays as sets. This operator provides a notable exception to this approach.

See also:

[find\(\)](#) (page 34), [update\(\)](#) (page 69), and [\\$set](#) (page 416).

\$elemMatch (query) **See also:**

[\\$elemMatch \(projection\)](#) (page 408)

\$elemMatch

New in version 1.4.

The [\\$elemMatch](#) (page 405) operator matches more than one component within an array element. For example,

```
db.collection.find( { array: { $elemMatch: { value1: 1, value2: { $gt: 1 } } } } );
```

returns all documents in `collection` where the array `array` satisfies all of the conditions in the [\\$elemMatch](#) (page 405) expression.

That is, where the value of `value1` is 1 and the value of `value2` is greater than 1. Matching arrays must have at least one element that matches all specified criteria. Therefore, the following document would not match the above query:

```
{ array: [ { value1:1, value2:0 }, { value1:2, value2:2 } ] }
```

while the following document would match this query:

```
{ array: [ { value1:1, value2:0 }, { value1:1, value2:2 } ] }
```

\$size

\$size

The [\\$size](#) (page 405) operator matches any array with the number of elements specified by the argument. For example:

```
db.collection.find( { field: { $size: 2 } } );
```

returns all documents in `collection` where `field` is an array with 2 elements. For instance, the above expression will return `{ field: [red, green] }` and `{ field: [apple, lime] }` but *not* `{ field: fruit }` or `{ field: [orange, lemon, grapefruit] }`. To match fields with only one element within an array use [\\$size](#) (page 405) with a value of 1, as follows:

```
db.collection.find( { field: { $size: 1 } } );
```

[\\$size](#) (page 405) does not accept ranges of values. To select documents based on fields with different numbers of elements, create a counter field that you increment when you add elements to a field.

Queries cannot use indexes for the `$size` (page 405) portion of a query, although the other portions of a query can use indexes if applicable.

Projection Operators

Projection Operators

Name	Description
<code>\$</code> (page 406)	Projects the first element in an array that matches the query condition.
<code>\$elemMatch</code> (page 408)	Projects only the first element from an array that matches the specified <code>\$elemMatch</code> (page 408) condition.
<code>\$meta</code> (page 410)	Projects the document's score assigned during <code>\$text</code> (page 387) operation.
<code>\$slice</code> (page 411)	Limits the number of elements projected from an array. Supports skip and limit slices.

`$` (projection)

Definition

`$`

The positional `$` (page 406) operator limits the contents of the `<array>` field that is included in the query results to contain the **first** matching element. To specify an array element to update, see the *positional \$ operator for updates* (page 420).

Used in the *projection* document of the `find()` (page 34) method or the `findOne()` (page 43) method:

- The `$` (page 406) projection operator limits the content of the `<array>` field to the **first** element that matches the *query document*.
- The `<array>` field **must** appear in the *query document*

```
db.collection.find( { <array>: <value> ... },
  { "<array>.$": 1 } )
db.collection.find( { <array.field>: <value> ... },
  { "<array>.$": 1 } )
```

The `<value>` can be documents that contains *query operator expressions* (page 373).

- Only **one** positional `$` (page 406) operator can appear in the projection document.
- Only **one** array field can appear in the *query document*; i.e. the following query is **incorrect**:

```
db.collection.find( { <array>: <value>, <someOtherArray>: <value2> },
  { "<array>.$": 1 } )
```

Behavior

Array Field Limitation Since only **one** array field can appear in the query document, if the array contains documents, to specify criteria on multiple fields of these documents, use the `$elemMatch` (page 405) operator. For example:

```
db.students.find( { grades: { $elemMatch: {
  mean: { $gt: 70 },
  grade: { $gt: 90 }
} } },
  { "grades.$": 1 } )
```

Sorts and the Positional Operator When the `find()` (page 34) method includes a `sort()` (page 93), the `find()` (page 34) method applies the `sort()` (page 93) to order the matching documents **before** it applies the positional `$` (page 406) projection operator.

If an array field contains multiple documents with the same field name and the `find()` (page 34) method includes a `sort()` (page 93) on that repeating field, the returned documents may not reflect the sort order because the sort was applied to the elements of the array before the `$` (page 406) projection operator.

Examples

Project Array Values A collection `students` contains the following documents:

```
{ "_id" : 1, "semester" : 1, "grades" : [ 70, 87, 90 ] }
{ "_id" : 2, "semester" : 1, "grades" : [ 90, 88, 92 ] }
{ "_id" : 3, "semester" : 1, "grades" : [ 85, 100, 90 ] }
{ "_id" : 4, "semester" : 2, "grades" : [ 79, 85, 80 ] }
{ "_id" : 5, "semester" : 2, "grades" : [ 88, 88, 92 ] }
{ "_id" : 6, "semester" : 2, "grades" : [ 95, 90, 96 ] }
```

In the following query, the projection `{ "grades.$": 1 }` returns only the first element greater than or equal to 85 for the `grades` field.

```
db.students.find( { semester: 1, grades: { $gte: 85 } },
                  { "grades.$": 1 } )
```

The operation returns the following documents:

```
{ "_id" : 1, "grades" : [ 87 ] }
{ "_id" : 2, "grades" : [ 90 ] }
{ "_id" : 3, "grades" : [ 85 ] }
```

Although the array field `grades` may contain multiple elements that are greater than or equal to 85, the `$` (page 406) projection operator returns only the first matching element from the array.

Project Array Documents A `students` collection contains the following documents where the `grades` field is an array of documents; each document contain the three field names `grade`, `mean`, and `std`:

```
{ "_id" : 7, semester: 3, "grades" : [ { grade: 80, mean: 75, std: 8 },
                                       { grade: 85, mean: 90, std: 5 },
                                       { grade: 90, mean: 85, std: 3 } ] }

{ "_id" : 8, semester: 3, "grades" : [ { grade: 92, mean: 88, std: 8 },
                                       { grade: 78, mean: 90, std: 5 },
                                       { grade: 88, mean: 85, std: 3 } ] }
```

In the following query, the projection `{ "grades.$": 1 }` returns only the first element with the `mean` greater than 70 for the `grades` field:

```
db.students.find(
  { "grades.mean": { $gt: 70 } },
  { "grades.$": 1 }
)
```

The operation returns the following documents:

```
{ "_id" : 7, "grades" : [ { "grade" : 80, "mean" : 75, "std" : 8 } ] }
{ "_id" : 8, "grades" : [ { "grade" : 92, "mean" : 88, "std" : 8 } ] }
```

Further Reading `$elemMatch` (projection) (page 408)

`$elemMatch` (projection) See also:

`$elemMatch` (query) (page 405)

`$elemMatch`

New in version 2.2.

The `$elemMatch` (page 408) projection operator limits the contents of an array field that is included in the query results to contain only the array element that matches the `$elemMatch` (page 408) condition.

Note:

- The elements of the array are documents.
 - If multiple elements match the `$elemMatch` (page 408) condition, the operator returns the **first** matching element in the array.
 - The `$elemMatch` (page 408) projection operator is similar to the positional `$` (page 406) projection operator.
-

The examples on the `$elemMatch` (page 408) projection operator assumes a collection `school` with the following documents:

```
{
  _id: 1,
  zipcode: "63109",
  students: [
    { name: "john", school: 102, age: 10 },
    { name: "jess", school: 102, age: 11 },
    { name: "jeff", school: 108, age: 15 }
  ]
}

{
  _id: 2,
  zipcode: "63110",
  students: [
    { name: "ajax", school: 100, age: 7 },
    { name: "achilles", school: 100, age: 8 },
  ]
}

{
  _id: 3,
  zipcode: "63109",
  students: [
    { name: "ajax", school: 100, age: 7 },
    { name: "achilles", school: 100, age: 8 },
  ]
}

{
  _id: 4,
  zipcode: "63109",
  students: [
    { name: "barney", school: 102, age: 7 },
  ]
}
```

Example

The following `find()` (page 34) operation queries for all documents where the value of the `zipcode` field is 63109. The `$elemMatch` (page 408) projection returns only the **first** matching element of the `students` array where the `school` field has a value of 102:

```
db.schools.find( { zipcode: "63109" },
                 { students: { $elemMatch: { school: 102 } } } )
```

The operation returns the following documents:

```
{ "_id" : 1, "students" : [ { "name" : "john", "school" : 102, "age" : 10 } ] }
{ "_id" : 3 }
{ "_id" : 4, "students" : [ { "name" : "barney", "school" : 102, "age" : 7 } ] }
```

- For the document with `_id` equal to 1, the `students` array contains multiple elements with the `school` field equal to 102. However, the `$elemMatch` (page 408) projection returns only the first matching element from the array.
- The document with `_id` equal to 3 does not contain the `students` field in the result since no element in its `students` array matched the `$elemMatch` (page 408) condition.

The `$elemMatch` (page 408) projection can specify criteria on multiple fields:

Example

The following `find()` (page 34) operation queries for all documents where the value of the `zipcode` field is 63109. The projection includes the **first** matching element of the `students` array where the `school` field has a value of 102 **and** the `age` field is greater than 10:

```
db.schools.find( { zipcode: "63109" },
                 { students: { $elemMatch: { school: 102, age: { $gt: 10 } } } } )
```

The operation returns the three documents that have `zipcode` equal to 63109:

```
{ "_id" : 1, "students" : [ { "name" : "jess", "school" : 102, "age" : 11 } ] }
{ "_id" : 3 }
{ "_id" : 4 }
```

Documents with `_id` equal to 3 and `_id` equal to 4 do not contain the `students` field since no element matched the `$elemMatch` (page 408) criteria.

When the `find()` (page 34) method includes a `sort()` (page 93), the `find()` (page 34) method applies the `sort()` (page 93) to order the matching documents **before** it applies the projection.

If an array field contains multiple documents with the same field name and the `find()` (page 34) method includes a `sort()` (page 93) on that repeating field, the returned documents may not reflect the sort order because the `sort()` (page 93) was applied to the elements of the array before the `$elemMatch` (page 408) projection.

Example

The following query includes a `sort()` (page 93) to order by descending `students.age` field:

```
db.schools.find(
    { zipcode: 63109 },
    { students: { $elemMatch: { school: 102 } } }
).sort( { "students.age": -1 } )
```

The operation applies the `sort()` (page 93) to order the documents that have the field `zipcode` equal to 63109 and then applies the projection. The operation returns the three documents in the following order:

```
{ "_id" : 1, "students" : [ { "name" : "john", "school" : 102, "age" : 10 } ] }
{ "_id" : 3 }
{ "_id" : 4, "students" : [ { "name" : "barney", "school" : 102, "age" : 7 } ] }
```

See also:

`$ (projection)` (page 406) operator

`$meta`

`$meta`

New in version 2.6.

The `$meta` (page 410) projection operator returns for each matching document the metadata (e.g. `"textScore"`) associated with the query. The `$meta` (page 410) expression can be a part of the *projection* document as well as a `sort()` (page 93) expression.

A `$meta` (page 410) expression has the following syntax:

```
{ <projectedFieldName>: { $meta: <metaDataKeyword> } }
```

The `$meta` (page 410) expression can specify the following keyword as the `<metaDataKeyword>`:

Key-word	Description	Sort Order
"textScore"	Returns the score associated with the corresponding query: <i>\$text</i> query for each matching document. The text score signifies how well the document matched the stemmed term or terms. If not used in conjunction with a query: <i>\$text</i> query, returns a score of 0.0	Descending

Behaviors

Projected Field Name The `<projectedFieldName>` cannot include a dot (.) in the name.

If the specified `<projectedFieldName>` already exists in the matching documents, in the result set, the existing fields will return with the `$meta` (page 410) values instead of with the stored values.

Projection The `$meta` (page 410) expression can be used in the *projection* document, as in:

```
db.collection.find(
  <query>,
  { score: { $meta: "textScore" } }
)
```

The `$meta` (page 410) expression specifies the inclusion of the field to the result set and does not specify an exclusion of the other fields.

The `$meta` (page 410) expression can be a part of a projection document that specifies exclusions of other fields or that specifies inclusions of other fields.

The metadata returns information on the processing of the `<query>` operation. As such, the returned metadata, assigned to the `<projectedFieldName>`, has no meaning inside a `<query>` expression; i.e. specifying a condition on the `<projectedFieldName>` as part of the `<query>` is similar to specifying a condition on a non-existing field if no field exists in the documents with the `<projectedFieldName>`.

Sort The `$meta` (page 410) expression can be part of a `sort()` (page 93) expression, as in:

```
db.collection.find(
  <query>,
  { score: { $meta: "textScore" } }
).sort( { score: { $meta: "textScore" } } )
```

To include a `$meta` (page 410) expression in a `sort()` (page 93) expression, the *same* `$meta` (page 410) expression, including the `<projectedFieldName>`, must appear in the projection document. The specified metadata determines the sort order. For example, the "textScore" metadata sorts in descending order.

For additional examples, see *Text Search with Additional Query and Sort Expressions* (page 390).

\$meta Aggregation Operator The behavior and requirements of the `$meta` (page 410) operator differs from that of the `$meta` (page 468) aggregation operator. See the `$meta` (page 468) aggregation operator for details.

Examples For examples of "textScore" projections and sorts, see `$text` (page 387).

\$slice (projection)

\$slice

The `$slice` (page 411) operator controls the number of items of an array that a query returns. For information on limiting the size of an array during an update with `$push` (page 425), see the `$slice` (page 428) modifier instead.

Consider the following prototype query:

```
db.collection.find( { field: value }, { array: { $slice: count } } );
```

This operation selects the document `collection` identified by a field named `field` that holds `value` and returns the number of elements specified by the value of `count` from the array stored in the `array` field. If `count` has a value greater than the number of elements in `array` the query returns all elements of the array.

`$slice` (page 411) accepts arguments in a number of formats, including negative values and arrays. Consider the following examples:

```
db.posts.find( {}, { comments: { $slice: 5 } } )
```

Here, `$slice` (page 411) selects the first five items in an array in the `comments` field.

```
db.posts.find( {}, { comments: { $slice: -5 } } )
```

This operation returns the last five items in array.

The following examples specify an array as an argument to `$slice` (page 411). Arrays take the form of `[skip , limit]`, where the first value indicates the number of items in the array to skip and the second value indicates the number of items to return.

```
db.posts.find( {}, { comments: { $slice: [ 20, 10 ] } } )
```

Here, the query will only return 10 items, after skipping the first 20 items of that array.

```
db.posts.find( {}, { comments: { $slice: [ -20, 10 ] } } )
```

This operation returns 10 items as well, beginning with the item that is 20th from the last item of the array.

2.3.2 Update Operators

Update Operators

Fields

Field Update Operators

Name	Description
<code>\$inc</code> (page 412)	Increments the value of the field by the specified amount.
<code>\$mul</code> (page 413)	Multiplies the value of the field by the specified amount.
<code>\$rename</code> (page 414)	Renames a field.
<code>\$setOnInsert</code> (page 415)	Sets the value of a field upon document creation during an upsert. Has no effect on operations that modify existing documents.
<code>\$set</code> (page 416)	Sets the value of a field in a document.
<code>\$unset</code> (page 417)	Removes the specified field from a document.
<code>\$min</code> (page 417)	Only updates the field if the specified value is less than the existing field value.
<code>\$max</code> (page 418)	Only updates the field if the specified value is greater than the existing field value.
<code>\$currentDate</code> (page 419)	Sets the value of a field to current date, either as a Date or a Timestamp.

`$inc`

`$inc`

The `$inc` (page 412) operator increments a value of a field by a specified amount. If the field does not exist, `$inc` (page 412) adds the field and sets the field to the specified amount. `$inc` (page 412) accepts positive and negative incremental amounts. Consider the following syntax:

```
{ $inc: { <field1>: <amount1>, ... } }
```

The following example increments the value of `quantity` by 5 for the *first* matching document in the `products` collection where `sku` equals `abc123`:

```
db.products.update( { sku: "abc123" },
  { $inc: { quantity: 5 } } )
```

To update all matching documents in the collection, specify `multi:true` option in the `update()` (page 69) method. For example:

```
db.records.update( { age: 20 }, { $inc: { age: 1 } }, { multi: true } );
```

The `update()` (page 69) operation increments the value of the `age` field by 1 for all documents in the `records` collection that have an `age` field equal to 20.

The `$inc` (page 412) operator can operate on multiple fields in a document. The following `update()` (page 69) operation uses the `$inc` (page 412) operator to modify both the `quantity` field and the `sales` field for the *first* matching document in the `products` collection where `sku` equals `abc123`:

```
db.products.update( { sku: "abc123" },
  { $inc: { quantity: -2, sales: 2 } } )
```

In the above example, the `$inc` (page 412) operator expression specifies `-2` for the `quantity` field to *decrease* the value of the `quantity` field (i.e. increment by `-2`) and specifies `2` for the `sales` field to increase the value of the `sales` field by 2.

\$mul**\$mul**

New in version 2.6.

Multiply the value of a field by a number. To specify a `$mul` (page 413) expression, use the following prototype:

```
{ $mul: { field: <number> } }
```

The field to update must contain a numeric value. If the field does not exist in a document, `$mul` (page 413) creates the field and sets the value to zero of the same numeric type as the multiplier.

Multiplication with values of mixed numeric types (32-bit integer, 64-bit integer, float) may result in conversion of numeric type. See *Multiplication Type Conversion Rules* for details.

Examples

Multiply the Value of a Field Consider a collection `products` with the following document:

```
{ _id: 1, item: "ABC", price: 10.99 }
```

The following `db.collection.update()` (page 69) operation updates the document, using the `$mul` (page 413) operator to multiply the value in the `price` field by 1.25:

```
db.products.update(
  { _id: 1 },
  { $mul: { price: 1.25 } }
)
```

The operation results in the following document, where the new value of the `price` field 13.7375 reflects the original value 10.99 multiplied by 1.25:

```
{ _id: 1, item: "ABC", price: 13.7375 }
```

Apply \$mul Operator to a Non-existing Field Consider a collection `products` with the following document:

```
{ _id: 2, item: "Unknown" }
```

The following `db.collection.update()` (page 69) operation updates the document, applying the `$mul` (page 413) operator to the field `price` that does not exist in the document:

```
db.products.update(
  { _id: 2 },
  { $mul: { price: NumberLong(100) } }
)
```

The operation results in the following document with a `price` field set to value 0 of numeric type *shell-type-long*, the same type as the multiplier:

```
{ "_id" : 2, "item" : "Unknown", "price" : NumberLong(0) }
```

Multiply Mixed Numeric Types Consider a collection `products` with the following document:

```
{ _id: 3, item: "XYZ", price: NumberLong(10) }
```

The following `db.collection.update()` (page 69) operation uses the `$mul` (page 413) operator to multiply the value in the `price` field `NumberLong(10)` by `NumberInt(5)`:

```
db.products.update(
  { _id: 3 },
  { $mul: { price: NumberInt(5) } }
)
```

The operation results in the following document:

```
{ "_id" : 3, "item" : "XYZ", "price" : NumberLong(50) }
```

The value in the `price` field is of type *shell-type-long*. See *Multiplication Type Conversion Rules* for details.

\$rename

\$rename

New in version 1.7.2.

Syntax: {`$rename`: { <old name1>: <new name1>, <old name2>: <new name2>, ... } }

The `$rename` (page 414) operator updates the name of a field. The new field name must differ from the existing field name.

Consider the following example:

```
db.students.update( { _id: 1 }, { $rename: { 'nickname': 'alias', 'cell': 'mobile' } } )
```

This operation renames the field `nickname` to `alias`, and the field `cell` to `mobile`.

Behavior The `$rename` (page 414) operator logically performs an `$unset` (page 417) of both the old name and the new name, and then performs a `$set` (page 416) operation with the new name. As such, the operation may not preserve the order of the fields in the document; i.e. the renamed field may move within the document.

If the document already has a field with the *new* field name, the `$rename` (page 414) operator removes that field and renames the field with the *old* field name to the *new* field name.

For fields in embedded documents, the `$rename` (page 414) operator can rename these fields as well as move the fields in and out of embedded documents. `$rename` (page 414) does not work if these fields are in array elements.

Examples A collection `students` the following document where a field `nmae` appears misspelled, i.e. should be `name`:

```
{ "_id": 1,
  "alias": [ "The American Cincinnatus", "The American Fabius" ],
  "mobile": "555-555-5555",
  "nmae": { "first" : "george", "last" : "washington" }
}
```

The examples in this section successively updates this document.

Rename a Field To rename a field, call the `$rename` (page 414) operator with the current name of the field and the new name:

```
db.students.update( { _id: 1 }, { $rename: { "nmae": "name" } } )
```

This operation renames the field `nmae` to `name`:

```
{
  "_id": 1,
  "alias": [ "The American Cincinnatus", "The American Fabius" ],
  "mobile": "555-555-5555",
  "name": { "first" : "george", "last" : "washington" }
}
```

Rename a Field in an Embedded Document To rename a field in an embedded document, call the `$rename` (page 414) operator using the *dot notation* to refer to the field. If the field is to remain in the same embedded document, also use the dot notation in the new name, as in the following:

```
db.students.update( { _id: 1 }, { $rename: { "name.first": "name.fname" } } )
```

This operation renames the embedded field `first` to `fname`:

```
{
  "_id" : 1,
  "alias" : [ "The American Cincinnatus", "The American Fabius" ],
  "mobile" : "555-555-5555",
  "name" : { "fname" : "george", "last" : "washington" }
}
```

Rename a Field That Does Not Exist When renaming a field and the existing field name refers to a field that does not exist, the `$rename` (page 414) operator does nothing, as in the following:

```
db.students.update( { _id: 1 }, { $rename: { 'wife': 'spouse' } } )
```

This operation does nothing because there is no field named `wife`.

\$setOnInsert

\$setOnInsert

New in version 2.4.

If an *upsert* results in an insert of a document, then `$setOnInsert` (page 415) assigns the specified values to the fields in the document. You can specify an upsert by specifying the *upsert* option for either the `db.collection.update()` (page 69) or `db.collection.findAndModify()` (page 39) methods. If the upsert results in an update, `$setOnInsert` (page 415) has no effect.

```
db.collection.update(
  <query>,
  { $setOnInsert: { <field1>: <value1>, ... } },
  { upsert: true }
)
```

Examples

Upsert Results in an Insert A collection named `products` contains no documents.

Then, the following *upsert* (page 69) operation performs an insert and applies the `$setOnInsert` (page 415) to set the field `defaultQty` to 100:

```
db.products.update(
  { _id: 1 },
  { $setOnInsert: { defaultQty: 100 } },
```

```
{ upsert: true }
)
```

The `products` collection contains the newly-inserted document:

```
{ "_id" : 1, "defaultQty" : 100 }
```

Upsert Results in an Update If the `db.collection.update()` (page 69) or the `db.collection.findAndModify()` (page 39) method has the `upsert` flag and performs an update and not an insert, `$setOnInsert` (page 415) has no effect.

A collection named `products` has the following document:

```
{ "_id" : 1, "defaultQty" : 100 }
```

The following `update()` (page 69) with the `upsert` flag operation performs an update:

```
db.products.update(
  { _id: 1 },
  {
    $setOnInsert: { defaultQty: 500, inStock: true },
    $set: { item: "apple" }
  },
  { upsert: true }
)
```

Because the `update()` (page 69) with `upsert` only performs an update, MongoDB ignores the `$setOnInsert` (page 415) operation and only applies the `$set` (page 416) operation.

The `products` collection now contains the following modified document:

```
{ "_id" : 1, "defaultQty" : 100, "item" : "apple" }
```

\$set

\$set

Syntax: { \$set: { <field1>: <value1>, ... } }

Use the `$set` (page 416) operator to replace the value of a field to the specified value. If the field does not exist, the `$set` (page 416) operator will add the field with the specified value.

The following example uses the `$set` (page 416) operator to update the value of the `quantity` field to 500 and the `instock` field to `true` for the *first* document where the field `sku` has the value `abc123`:

```
db.products.update( { sku: "abc123" },
  { $set: {
    quantity: 500,
    instock: true
  } }
)
```

To update all matching documents in the collection, specify `multi: true` option in the `update()` (page 69) method, as in the following example which sets the value of the field `instock` to `true` for all documents in the `products` collection where the `quantity` field is greater than (i.e. `$gt` (page 373)) 0 :

```
db.products.update( { quantity: { $gt: 0 } },
  { $set: { instock: true } },
  { multi: true }
)
```

```
    { multi: true }
  )
```

\$unset

\$unset

The `$unset` (page 417) operator deletes a particular field. The specified value in the `$unset` (page 417) expression (i.e. "" below) does not impact the operation. If the field does not exist, then `$unset` (page 417) has no effect. Consider the following syntax:

```
{ $unset: { <field1>: "", ... } }
```

For example, the following `update()` (page 69) operation uses the `$unset` (page 417) operator to remove the fields `quantity` and `instock` from the *first* document found in the `products` collection where the field `sku` has a value of `unknown`.

```
db.products.update( { sku: "unknown" },
  { $unset: {
    quantity: "",
    instock: ""
  } }
)
```

To remove the fields from *all* documents in the collection where the field `sku` has a value of `unknown`, specify the `multi: true` option in the `update()` (page 69) method, as in the following example:

```
db.products.update( { sku: "unknown" },
  { $unset: {
    quantity: "",
    instock: ""
  } },
  { multi: true }
)
```

\$min

\$min

The `$min` (page 417) updates the value of the field to a specified value *if* the specified value is **less than** the current value of the field. If the field does not exist, the `$min` (page 417) operator sets the field to the specified value. The `$min` (page 417) operator can compare values of different types, using the *BSON comparison order*.

Examples

Use \$min to Compare Numbers Consider the following document in the collection `scores`:

```
{ _id: 1, highScore: 800, lowScore: 200 }
```

The `lowScore` for the document currently has the value 200. The following operation uses `$min` (page 417) to compare 200 to the specified value 150 and updates the value of `lowScore` to 150 since 150 is less than 200:

```
db.scores.update( { _id: 1 }, { $min: { lowScore: 150 } } )
```

The `scores` collection now contains the following modified document:

```
{ _id: 1, highScore: 800, lowScore: 150 }
```

The next operation has no effect since the current value of the field `lowScore`, i.e 150, is less than 200:

```
db.scores.update( { _id: 1 }, { $min: { lowScore: 250 } } )
```

The document remains unchanged in the `scores` collection:

```
{ _id: 1, highScore: 800, lowScore: 150 }
```

Use `$min` to Compare Dates Consider the following document in the collection `tags`:

```
{
  _id: 1,
  desc: "crafts",
  dateEntered: ISODate("2013-10-01T05:00:00Z"),
  dateExpired: ISODate("2013-10-01T16:38:16Z")
}
```

The following operation compares the current value of the `dateEntered` field, i.e. `ISODate("2013-10-01T05:00:00Z")`, with the specified date `new Date("2013-09-25")` to determine whether to update the field:

```
db.tags.update( { _id: 1 },
  {
    $min: {
      dateEntered: new Date("2013-09-25")
    }
  }
)
```

The operation updates the `dateEntered` field:

```
{
  _id: 1,
  desc: "crafts",
  dateEntered: ISODate("2013-09-25T00:00:00Z"),
  dateExpired: ISODate("2013-10-01T16:38:16Z")
}
```

`$max`

`$max`

The `$max` (page 418) operator updates the value of the field to a specified value *if* the specified value is **greater than** the current value of the field. If the field does not exist, the `$max` (page 418) operator sets the field to the specified value. The `$max` (page 418) operator can compare values of different types, using the *BSON comparison order*.

Examples

Use `$max` to Compare Numbers Consider the following document in the collection `scores`:

```
{ _id: 1, highScore: 800, lowScore: 200 }
```

The `highScore` for the document currently has the value 800. The following operation uses `$max` (page 480) to compare the 800 and the specified value 950 and updates the value of `highScore` to 950 since 950 is greater than 800:

```
db.scores.update( { _id: 1 }, { $max: { highScore: 950 } } )
```

The scores collection now contains the following modified document:

```
{ _id: 1, highScore: 950, lowScore: 200 }
```

The next operation has no effect since the current value of the field `highScore`, i.e. 950, is greater than 870:

```
db.scores.update( { _id: 1 }, { $max: { highScore: 870 } } )
```

The document remains unchanged in the `scores` collection:

```
{ _id: 1, highScore: 950, lowScore: 200 }
```

Use `$max` to Compare Dates Consider the following document in the collection `tags`:

```
{
  _id: 1,
  desc: "crafts",
  dateEntered: ISODate("2013-10-01T05:00:00Z"),
  dateExpired: ISODate("2013-10-01T16:38:16.163Z")
}
```

The following operation compares the current value of the `dateExpired` field, i.e. `ISODate("2013-10-01T16:38:16.163Z")`, with the specified date `new Date("2013-09-30")` to determine whether to update the field:

```
db.tags.update( { _id: 1 },
  {
    $max: {
      dateExpired: new Date("2013-09-30"),
    }
  }
)
```

The operation does *not* update the `dateExpired` field:

```
{
  _id: 1,
  desc: "decorative arts",
  dateEntered: ISODate("2013-10-01T05:00:00Z"),
  dateExpired: ISODate("2013-10-01T16:38:16.163Z")
}
```

`$currentDate`

`$currentDate`

The `$currentDate` (page 419) operator sets the value of a field to the current date, either as a *Date* or a *timestamp*. The default type is *date*.

The `$currentDate` (page 419) operator can take as its operand either

- a boolean `true` which creates a *Date*, or
- a document which explicitly specifies the type, i.e. `{ $type: "timestamp" }` or `{ $type: "date" }`. The operator is *case-sensitive* and accepts only the lowercase `"timestamp"` or the lowercase `"date"`.

Example Consider the following document in the `users` collection:

```
{ _id: 1, status: "a", lastModified: ISODate("2013-10-02T01:11:18.965Z") }
```

The following updates the `lastModified` field to the current date and the `lastModifiedTS` field to the current timestamp as well as setting the `status` field to `"D"`.

```
db.users.update( { _id: 1 },
  {
    $currentDate: {
      lastModified: true,
      lastModifiedTS: { $type: "timestamp" }
    },
    $set: { status: "D" }
  }
)
```

Following this operation, the updated document would resemble:

```
{
  _id: 1,
  status: "D",
  lastModified: ISODate("2013-10-02T01:11:53.976Z"),
  lastModifiedTS: Timestamp(1380676313, 1)
}
```

Array

Array Update Operators

Update Operators	Name	Description
	\$ (page 420)	Acts as a placeholder to update the first element that matches the query condition in an update.
	\$addToSet (page 422)	Adds elements to an array only if they do not already exist in the set.
	\$pop (page 423)	Removes the first or last item of an array.
	\$pullAll (page 424)	Removes all matching values from an array.
	\$pull (page 424)	Removes all array elements that match a specified query.
	\$pushAll (page 425)	<i>Deprecated.</i> Adds several items to an array.
	\$push (page 425)	Adds an item to an array.

\$ (update)

Definition

\$

```
Syntax: { "<array>.$" : value }
```

The positional `$` operator identifies an element in an `array` field to update without explicitly specifying the position of the element in the array. To project, or return, an array element from a read operation, see the `$` (page 406) projection operator.

When used with the `update()` (page 69) method,

- the positional \$ operator acts as a placeholder for the **first** element that matches the *query document*, and
- the array field **must** appear as part of the *query document*.

```
db.collection.update( { <array>: value ... }, { <update operator>: { "<array>.$" : value } } )
```

Behavior

Upserts Do not use the positional operator \$ with *upsert* operations because inserts will use the \$ as a field name in the inserted document.

Nested Arrays The positional \$ operator cannot be used for queries which traverse more than one array, such as queries that traverse arrays nested within other arrays, because the replacement for the \$ placeholder is a single value

Unsets When used with the \$unset (page 417) operator, the positional \$ operator does not remove the matching element from the array but rather sets it to null.

Negations If the query matches the array using a negation operator, such as \$ne (page 376), \$not (page 379), or \$nin (page 376), then you cannot use the positional operator to update values from this array.

However, if the negated portion of the query is inside of an \$elemMatch (page 405) expression, then you *can* use the positional operator to update this field.

Examples

Update Values in an Array Consider a collection `students` with the following documents:

```
{ "_id" : 1, "grades" : [ 80, 85, 90 ] }
{ "_id" : 2, "grades" : [ 88, 90, 92 ] }
{ "_id" : 3, "grades" : [ 85, 100, 90 ] }
```

To update 80 to 82 in the `grades` array in the first document, use the positional \$ operator if you do not know the position of the element in the array:

```
db.students.update( { _id: 1, grades: 80 }, { $set: { "grades.$" : 82 } } )
```

Remember that the positional \$ operator acts as a placeholder for the **first match** of the update *query document*.

Update Documents in an Array The positional \$ operator facilitates updates to arrays that contain embedded documents. Use the positional \$ operator to access the fields in the embedded documents with the *dot notation* on the \$ operator.

```
db.collection.update( { <query selector> }, { <update operator>: { "array.$.field" : value } } )
```

Consider the following document in the `students` collection whose `grades` field value is an array of embedded documents:

```
{ "_id" : 4, "grades" : [ { grade: 80, mean: 75, std: 8 },
                        { grade: 85, mean: 90, std: 5 },
                        { grade: 90, mean: 85, std: 3 } ] }
```

Use the positional \$ operator to update the value of the `std` field in the embedded document with the grade of 85:

```
db.students.update( { _id: 4, "grades.grade": 85 }, { $set: { "grades.$.std" : 6 } } )
```

Further Reading [update\(\)](#) (page 69), [\\$set](#) (page 416) and [\\$unset](#) (page 417)

\$addToSet

Definition

\$addToSet

The [\\$addToSet](#) (page 422) operator adds a value to an array only *if* the value is *not* already in the array. If the value *is* in the array, [\\$addToSet](#) (page 422) does not modify the array.

```
db.collection.update( <query>, { $addToSet: { <field>: <value> } } );
```

For example, if a collection `inventory` has the following document:

```
{ _id: 1, item: "filter", tags: [ "electronics", "camera" ] }
```

The following operation adds the element `"accessories"` to the `tags` array since `"accessories"` does not exist in the array:

```
db.inventory.update(
  { _id: 1 },
  { $addToSet: { tags: "accessories" } }
)
```

However, the following operation has no effect as `"camera"` is already an element of the `tags` array:

```
db.inventory.update(
  { _id: 1 },
  { $addToSet: { tags: "camera" } }
)
```

Behavior

- [\\$addToSet](#) (page 422) only ensures that there are no duplicate items *added* to the set and does not affect existing duplicate elements. [\\$addToSet](#) (page 422) does not guarantee a particular ordering of elements in the modified set.
- If the field is absent in the document to update, [\\$addToSet](#) (page 422) adds the array field with the value as its element.
- If the field is **not** an array, the operation will fail.
- If the value is an array, [\\$addToSet](#) (page 422) appends the whole array as a *single* element. To add each element of the value separately, use [\\$addToSet](#) (page 422) with the [\\$each](#) (page 427) modifier. See [Modifiers](#) (page 422) for details.

Modifiers You can use the [\\$addToSet](#) (page 422) operator with the [\\$each](#) (page 427) modifier. The [\\$each](#) (page 427) modifier allows to [\\$addToSet](#) (page 422) operator to add multiple values to the array field.

A collection `inventory` has the following document:

```
{ _id: 2, item: "cable", tags: [ "electronics", "supplies" ] }
```

Then the following operation uses the `$addToSet` (page 422) operator with the `$each` (page 427) modifier to add multiple elements to the `tags` array:

```
db.inventory.update(
  { _id: 2 },
  { $addToSet: { tags: { $each: [ "camera",
                                "electronics",
                                "accessories" ] } } }
)
```

The operation adds only "camera" and "accessories" to the `tags` array since "electronics" already exists in the array:

```
{ _id: 2,
  item: "cable",
  tags: [ "electronics", "supplies", "camera", "accessories" ] }
```

See also:

`$push` (page 425)

`$pop`

`$pop`

New in version 1.1.

The `$pop` (page 423) operator removes the first or last element of an array. Pass `$pop` (page 423) a value of `-1` to remove the first element of an array and `1` to remove the last element in an array.

```
db.collection.update( <query>,
  { $pop: { <field>: <-1 | 1> } }
)
```

Behavior The `$pop` (page 423) operation fails if the `<field>` is not an array.

If the `$pop` (page 423) operator removes the last item in the `<field>`, the `<field>` will then hold an empty array.

Examples

Remove the First Item of an Array Given the following document in a collection `students`:

```
{ _id: 1, scores: [ 8, 9, 10 ] }
```

The following example removes the *first* element (8) in the `scores` array:

```
db.students.update( { _id: 1 }, { $pop: { scores: -1 } } )
```

After the operation, the updated document has the first item 8 removed from its `scores` array:

```
{ _id: 1, scores: [ 9, 10 ] }
```

Remove the Last Item of an Array Given the following document in a collection `students`:

```
{ _id: 1, scores: [ 9, 10 ] }
```

The following example removes the *last* element (10) in the `scores` array by specifying 1 in the `$pop` (page 423) expression:

```
db.students.update( { _id: 1 }, { $pop: { scores: 1 } } )
```

After the operation, the updated document has the last item 10 removed from its `scores` array:

```
{ _id: 1, scores: [ 9 ] }
```

\$pullAll

\$pullAll

The `$pullAll` (page 424) operator removes all instances of the specified values from an existing array.

```
db.collection.update( <query>,
                      { $pullAll: { <arrayField>: [ <value1>, <value2> ... ] } }
                      )
```

Unlike the `$pull` (page 424) operator that removes elements by specifying a query, `$pullAll` (page 424) removes elements that match the listed values.

For example, given the following document in the `survey` collection:

```
{ _id: 1, scores: [ 0, 2, 5, 5, 1, 0 ] }
```

The following operation removes all instances of the value 0 and 5 from the `scores` array:

```
db.survey.update( { _id: 1 }, { $pullAll: { scores: [ 0, 5 ] } } )
```

After the operation, the updated document has all instances of 0 and 5 removed from the `scores` field:

```
{ "_id" : 1, "scores" : [ 2, 1 ] }
```

\$pull

\$pull

The `$pull` (page 424) operator removes from an existing array all instances of a value or values that match a specified query.

```
db.collection.update( <query>,
                      { $pull: { <arrayField>: <query2> } }
                      )
```

To specify the query to remove values from the array, use *query operators* (page 373).

Examples

Remove All Items That Equals a Specified Value Given the following documents in the `cpuinfo` collection:

```
{ _id: 1, flags: [ "vme", "de", "msr", "tsc", "pse", "msr" ] }
{ _id: 2, flags: [ "msr", "pse", "tsc" ] }
```

The following operation removes the value "msr" from the `flags` array:

```
db.cpuinfo.update(
  { flags: "msr" },
  { $pull: { flags: "msr" } },
  { multi: true }
)
```

After the operation, the documents no longer have any "msr" values:

```
{ _id: 1, flags: [ "vme", "de", "tsc", "pse" ] }
{ _id: 2, flags: [ "pse", "tsc" ] }
```

Remove All Items Greater Than a Specified Value Given the following document in the `profiles` collection:

```
{ _id: 1, votes: [ 3, 5, 6, 7, 7, 8 ] }
```

The following operation will remove all items from the `votes` array that are greater than or equal (`$gte` (page 374)) 6:

```
db.profiles.update( { _id: 1 }, { $pull: { votes: { $gte: 6 } } } )
```

After the update operation, the document only has values less than 6:

```
{ _id: 1, votes: [ 3, 5 ] }
```

\$pushAll

\$pushAll

Deprecated since version 2.4: Use the `$push` (page 425) operator with `$each` (page 427) instead.

The `$pushAll` (page 425) operator is similar to the `$push` (page 425) but adds the ability to append several values to an array at once.

```
db.collection.update( { field: value }, { $pushAll: { field1: [ value1, value2, value3 ] } } );
```

Here, `$pushAll` (page 425) appends the values in `[value1, value2, value3]` to the array in `field1` in the document matched by the statement `{ field: value }` in `collection`.

If you specify a single value, `$pushAll` (page 425) will behave as `$push` (page 425).

\$push

\$push

The `$push` (page 425) operator appends a specified value to an array.

```
db.collection.update( <query>,
  { $push: { <field>: <value> } }
)
```

The following example appends 89 to the `scores` array for the first document where the `_id` field equals 1:

```
db.students.update(
  { _id: 1 },
  { $push: { scores: 89 } }
)
```

Note:

- If the field is absent in the document to update, `$push` (page 425) adds the array field with the value as its element.
- If the field is **not** an array, the operation will fail.
- If the value is an array, `$push` (page 425) appends the whole array as a *single* element. To add each element of the value separately, use `$push` (page 425) with the `$each` (page 427) modifier.

Changed in version 2.4: MongoDB adds support for the `$each` (page 427) modifier to the `$push` (page 425) operator. Before 2.4, use `$pushAll` (page 425) for similar functionality.

The following example appends each element of [90, 92, 85] to the `scores` array for the document where the `name` field equals `joe`:

```
db.students.update(
  { name: "joe" },
  { $push: { scores: { $each: [ 90, 92, 85 ] } } }
)
```

Modifiers New in version 2.4.

You can use the `$push` (page 425) operator with the following modifiers:

- `$each` (page 427) appends multiple values to the array field,
Changed in version 2.6: When used in conjunction with the other modifiers, the `$each` (page 427) modifier no longer needs to be first.
- `$slice` (page 428), which is only available when used with `$each` (page 427), limits the number of array elements,
- `$sort` (page 431), which is only available when used with `$each` (page 427), orders elements of the array, and
Changed in version 2.6: In previous versions, `$sort` (page 431) is only available when used with both `$each` (page 427) and `$slice` (page 428).
- `$position` (page 433), which is only available when used with `$each` (page 427), specifies the location in the array at which to insert the new elements. Without the `$position` (page 433) modifier, the `$push` (page 425) appends the elements to the end of the array.

New in version 2.6.

The processing of the `push` operation with modifiers occur in the following order, regardless of the order in which the modifiers appear:

1. Update array to add elements in the correct position.
2. Apply sort, if specified.
3. Slice the array, if specified.
4. Store the array.

Examples

Use `$push` Operator with Multiple Modifiers A collection `students` has the following document:

```
{
  "_id" : 5,
  "quizzes" : [
    { wk: 1, "score" : 10 },
    { wk: 2, "score" : 8 },
    { wk: 3, "score" : 5 },
    { wk: 4, "score" : 6 }
  ]
}
```

The following `$push` (page 425) operation uses:

- the `$each` (page 427) modifier to add multiple documents to the `quizzes` array,

- the `$sort` (page 431) modifier to sort all the elements of the modified `quizzes` array by the `score` field in descending order, and
- the `$slice` (page 428) modifier to keep only the **first** three sorted elements of the `quizzes` array.

```
db.students.update( { _id: 5 },
  { $push: { quizzes: { $each: [ { wk: 5, score: 8 },
                                { wk: 6, score: 7 },
                                { wk: 7, score: 6 } ],
                                $sort: { score: -1 },
                                $slice: 3
                              }
        }
  }
)
```

The result of the operation is keep only the three highest scoring quizzes:

```
{ "_id" : 5,
  "quizzes" : [
    { "wk" : 1, "score" : 10 },
    { "wk" : 2, "score" : 8 },
    { "wk" : 5, "score" : 8 }
  ]
}
```

Update Operator Modifiers

Name	Description
<code>\$each</code> (page 427)	Modifies the <code>\$push</code> (page 425) and <code>\$addToSet</code> (page 422) operators to append items for array updates.
<code>\$slice</code> (page 428)	Modifies the <code>\$push</code> (page 425) operator to limit the size of updated arrays.
<code>\$sort</code> (page 431)	Modifies the <code>\$push</code> (page 425) operator to reorder documents stored in an array.
<code>\$position</code> (page 433)	Modifies the <code>\$push</code> (page 425) operator to specify the position in the array to add elements.

`$each`

`$each`

The `$each` (page 427) modifier is available for use with the `$addToSet` (page 422) operator and the `$push` (page 425) operator.

Use the `$each` (page 427) modifier with the `$addToSet` (page 422) operator to add multiple values to an array `<field>` if the values do not exist in the `<field>`.

```
db.collection.update( <query>,
  {
    $addToSet: { <field>: { $each: [ <value1>, <value2> ... ] } }
  }
)
```

Use the `$each` (page 427) modifier with the `$push` (page 425) operator to append multiple values to an array `<field>`.

```
db.collection.update( <query>,
  {
    $push: { <field>: { $each: [ <value1>, <value2> ... ] } }
  }
)
```

Changed in version 2.4: MongoDB adds support for the `$each` (page 427) modifier to the `$push` (page 425) operator. The `$push` (page 425) operator can use `$each` (page 427) modifier with other modifiers. See `$push` (page 425) for details.

Examples

Use `$each` with `$push` Operator The following example appends each element of [90, 92, 85] to the `scores` array for the document where the `name` field equals `joe`:

```
db.students.update(
  { name: "joe" },
  { $push: { scores: { $each: [ 90, 92, 85 ] } } }
)
```

Use `$each` with `$addToSet` Operator A collection `inventory` has the following document:

```
{ _id: 2, item: "cable", tags: [ "electronics", "supplies" ] }
```

Then the following operation uses the `$addToSet` (page 422) operator with the `$each` (page 427) modifier to add multiple elements to the `tags` array:

```
db.inventory.update(
  { _id: 2 },
  { $addToSet: { tags: { $each: [ "camera",
                                "electronics",
                                "accessories" ] } } }
)
```

The operation adds only "camera" and "accessories" to the `tags` array since "electronics" already exists in the array:

```
{ _id: 2,
  item: "cable",
  tags: [ "electronics", "supplies", "camera", "accessories" ] }
```

`$slice`

`$slice`

New in version 2.4.

The `$slice` (page 428) modifier limits the number of array elements during a `$push` (page 425) operation. To project, or return, a specified number of array elements from a read operation, see the `$slice` (page 411) projection operator instead.

To use the `$slice` (page 428) modifier, it must appear with the `$each` (page 427) modifier.

Tip

You can pass an empty array [] to the `$each` (page 427) modifier such that only the `$slice` (page 428) modifier has an effect.

Changed in version 2.6: The `$slice` (page 428) can slice from the beginning of the array. Trying to use the `$slice` (page 428) modifier without the `$each` (page 427) modifier results in an error. The order in which the modifiers appear is immaterial. Previous versions required the `$each` (page 427) modifier to appear as the first modifier if used in conjunction with `$slice` (page 428).


```

•db.collection.update( <query>,
                      { $push: {
                          <field>: {
                              $each: [ <value1>, <value2>, ... ],
                              $slice: <num>
                          }
                      }
                    }
                  )

```

The <num> can be:

- zero** to update the array <field> to an empty array,
- negative** to update the array <field> to contain only the last <num> elements, or
- positive** to update the array <field> contain only the first <num> elements.

New in version 2.6.

Examples

Slice from the End of the Array A collection `students` contains the following document:

```
{ "_id" : 1, "scores" : [ 40, 50, 60 ] }
```

The following operation adds new elements to the `scores` array, and then uses the `$slice` (page 428) modifier to trim the array to the last five elements:

```

db.students.update( { _id: 1 },
                  { $push: { scores: {
                              $each: [ 80, 78, 86 ],
                              $slice: -5
                            }
                      }
                }
              )

```

The result of the operation is slice the elements of the updated `scores` array to the last five elements:

```
{ "_id" : 1, "scores" : [ 50, 60, 80, 78, 86 ] }
```

Slice from the Front of the Array A collection `students` contains the following document:

```
{ "_id" : 2, "scores" : [ 89, 90 ] }
```

The following operation adds new elements to the `scores` array, and then uses the `$slice` (page 428) modifier to trim the array to the first three elements.

```

db.students.update( { _id: 2 },
                  { $push: { scores: { $each: [ 100, 20 ], $slice: 3 } } }
                )

```

The result of the operation is to slice the elements of the updated `scores` array to the first three elements:

```
{ "_id" : 2, "scores" : [ 89, 90, 100 ] }
```

Update Array Using Slice Only A collection `students` contains the following document:

```
{ "_id" : 3, "scores" : [ 89, 70, 100, 20 ] }
```

To update the `scores` field with just the effects of the `$slice` (page 428) modifier, specify the number of elements to slice (e.g. `-3`) for the `$slice` (page 428) modifier and an empty array `[]` for the `$each` (page 427) modifier, as in the following:

```
db.students.update(
  { _id: 3 },
  { $push: { scores: { $each: [], $slice: -3 } } }
)
```

The result of the operation is to slice the elements of the `scores` array to the last three elements:

```
{ "_id" : 3, "scores" : [ 70, 100, 20 ] }
```

Use \$slice with Other \$push Modifiers A collection `students` has the following document:

```
{
  "_id" : 5,
  "quizzes" : [
    { wk: 1, "score" : 10 },
    { wk: 2, "score" : 8 },
    { wk: 3, "score" : 5 },
    { wk: 4, "score" : 6 }
  ]
}
```

The following `$push` (page 425) operation uses:

- the `$each` (page 427) modifier to add multiple documents to the `quizzes` array,
- the `$sort` (page 431) modifier to sort all the elements of the modified `quizzes` array by the `score` field in descending order, and
- the `$slice` (page 428) modifier to keep only the **first** three sorted elements of the `quizzes` array.

```
db.students.update( { _id: 5 },
  { $push: { quizzes: { $each: [ { wk: 5, score: 8 },
                                { wk: 6, score: 7 },
                                { wk: 7, score: 6 } ],
                                $sort: { score: -1 },
                                $slice: 3
                              } } }
)
```

The result of the operation is keep only the three highest scoring quizzes:

```
{ "_id" : 5,
  "quizzes" : [
    { "wk" : 1, "score" : 10 },
    { "wk" : 2, "score" : 8 },
    { "wk" : 5, "score" : 8 }
  ]
}
```

The order of the modifiers is immaterial to the order in which the modifiers are processed. See *Modifiers* (page 426) for details.

\$sort**\$sort**

New in version 2.4.

The `$sort` (page 431) modifier orders the elements of an array during a `$push` (page 425) operation.

To use the `$sort` (page 431) modifier, it must appear with the `$each` (page 427) modifier.

Tip

You can pass an empty array `[]` to the `$each` (page 427) modifier such that only the `$sort` (page 431) modifier has an effect.

Changed in version 2.6: The `$sort` (page 431) modifier can sort array elements that are not documents. In previous versions, the `$sort` (page 431) modifier required the array elements be documents. If the array elements are documents, the modifier can sort by either the whole document or by a specific field in the documents. In previous versions, the `$sort` (page 431) modifier can only sort by a specific field in the documents. The `$sort` (page 431) no longer requires the `$slice` (page 428) modifier. Trying to use the `$sort` (page 431) modifier without the `$each` (page 427) modifier results in an error.

```
•db.collection.update( <query>,
                      { $push: {
                          <arrayField>: {
                              $each: [ <value1>,
                                      <value2>,
                                      ...
                                      ],
                              $sort: <sort specification>,
                          }
                      }
                      )
```

For `<sort specification>`:

- To sort array elements that are not documents, or if the array elements are documents, to sort by the whole documents, specify `1` for ascending or `-1` for descending.
- If the array elements are documents, to sort by a field in the documents, specify a sort document with the field and the direction, i.e. `{ field: 1 }` or `{ field: -1 }`. **Do not** reference the containing array field in the sort specification (e.g. `{ "arrayField.field": 1 }` is incorrect).

Examples

Sort Array of Documents by a Field in the Documents A collection `students` contains the following document:

```
{ "_id": 1,
  "quizzes": [
    { "id" : 1, "score" : 6 },
    { "id" : 2, "score" : 9 }
  ]
}
```

The following update appends additional documents to the `quizzes` array and then sorts all the elements of the array by the ascending `score` field:

```
db.students.update( { _id: 1 },
  { $push: { quizzes: { $each: [ { id: 3, score: 8 },
                                { id: 4, score: 7 },
                                { id: 5, score: 6 } ] },
    $sort: { score: 1 },
  }
})
```

Important: The sort document refers directly to the field in the documents and does not reference the containing array field `quizzes`; i.e. `{ score: 1 }` and **not** `{ "quizzes.score": 1 }`

After the update, the array elements are in order of ascending `score` field.:

```
{
  "_id" : 1,
  "quizzes" : [
    { "id" : 1, "score" : 6 },
    { "id" : 5, "score" : 6 },
    { "id" : 4, "score" : 7 },
    { "id" : 3, "score" : 8 },
    { "id" : 2, "score" : 9 }
  ]
}
```

Sort Array Elements That Are Not Documents A collection `students` contains the following document:

```
{ "_id" : 2, "tests" : [ 89, 70, 89, 50 ] }
```

The following operation adds two more elements to the `scores` array and sorts the elements:

```
db.students.update(
  { _id: 2 },
  { $push: { tests: { $each: [ 40, 60 ], $sort: 1 } } }
)
```

The updated document has the elements of the `scores` array in ascending order:

```
{ "_id" : 2, "tests" : [ 40, 50, 60, 70, 89, 89 ] }
```

Update Array Using Sort Only A collection `students` contains the following document:

```
{ "_id" : 3, "tests" : [ 89, 70, 100, 20 ] }
```

To update the `tests` field to sort its elements in descending order, specify the `{ $sort: -1 }` and specify an empty array `[]` for the `$each` (page 427) modifier, as in the following:

```
db.students.update(
  { _id: 3 },
  { $push: { tests: { $each: [ ], $sort: -1 } } }
)
```

The result of the operation is to update the `scores` field to sort its elements in descending order:

```
{ "_id" : 3, "tests" : [ 100, 89, 70, 20 ] }
```

Use \$sort with Other \$push Modifiers A collection `students` has the following document:

```
{
  "_id" : 5,
  "quizzes" : [
    { wk: 1, "score" : 10 },
    { wk: 2, "score" : 8 },
    { wk: 3, "score" : 5 },
    { wk: 4, "score" : 6 }
  ]
}
```

The following `$push` (page 425) operation uses:

- the `$each` (page 427) modifier to add multiple documents to the `quizzes` array,
- the `$sort` (page 431) modifier to sort all the elements of the modified `quizzes` array by the `score` field in descending order, and
- the `$slice` (page 428) modifier to keep only the **first** three sorted elements of the `quizzes` array.

```
db.students.update( { _id: 5 },
  { $push: { quizzes: { $each: [ { wk: 5, score: 8 },
                                { wk: 6, score: 7 },
                                { wk: 7, score: 6 } ],
                                $sort: { score: -1 },
                                $slice: 3
                              }
        }
  }
)
```

The result of the operation is keep only the three highest scoring quizzes:

```
{ "_id" : 5,
  "quizzes" : [
    { "wk" : 1, "score" : 10 },
    { "wk" : 2, "score" : 8 },
    { "wk" : 5, "score" : 8 }
  ]
}
```

The order of the modifiers is immaterial to the order in which the modifiers are processed. See *Modifiers* (page 426) for details.

\$position

\$position

New in version 2.6.

The `$position` (page 433) modifier specifies the location in the array at which the `$push` (page 425) operator insert elements. Without the `$position` (page 433) modifier, the `$push` (page 425) operator inserts elements to the end of the array. See *\$push modifiers* (page 426) for more information.

To use the `$position` (page 433) modifier, it must appear with the `$each` (page 427) modifier.

```
db.collection.update( <query>,
  { $push: {
    <field>: {
      $each: [ <value1>, <value2>, ... ],
      $position: <num>
    }
  }
}
```

```
    }  
  }  
)
```

The `<num>` is a non-negative number that corresponds to the position in the array, based on a zero-based index. If the number is greater or equal to the length of the array, the `$position` (page 433) modifier has no effect and the operator adds elements to the end of the array.

Examples

Add Elements at the Start of the Array Consider a collection `students` that contains the following document:

```
{ "_id" : 1, "scores" : [ 100 ] }
```

The following operation updates the `scores` field to add the elements 50, 60 and 70 to the beginning of the array:

```
db.students.update( { _id: 1 },  
  { $push: { scores: {  
    $each: [ 50, 60, 70 ],  
    $position: 0  
  } }  
})
```

The operation results in the following updated document:

```
{ "_id" : 1, "scores" : [ 50, 60, 70, 100 ] }
```

Add Elements to the Middle of the Array Consider a collection `students` that contains the following document:

```
{ "_id" : 1, "scores" : [ 50, 60, 70, 100 ] }
```

The following operation updates the `scores` field to add the elements 20 and 30 at the array index of 2:

```
db.students.update( { _id: 1 },  
  { $push: { scores: {  
    $each: [ 20, 30 ],  
    $position: 2  
  } }  
})
```

The operation results in the following updated document:

```
{ "_id" : 1, "scores" : [ 50, 60, 20, 30, 70, 100 ] }
```

Bitwise

Bitwise Update Operator

Name	Description
<code>\$bit</code> (page 435)	Performs bitwise AND, OR, and XOR updates of integer values.

\$bit
\$bit

Changed in version 2.6: Added support for bitwise `xor` operation.

The `$bit` (page 435) operator performs a bitwise update of a field. The `$bit` (page 435) operator supports bitwise `and`, bitwise `or`, and bitwise `xor` (i.e. exclusive or) operations. To specify a `$bit` (page 435) operator expression, use the following prototype:

```
{ $bit: { field: { <and|or|xor>: <int> } } }
```

Only use this operator with integer fields (either 32-bit integer or 64-bit integer).

Note: All numbers in the `mongo` (page 527) shell are doubles, not integers. Use the `NumberInt()` or the `NumberLong()` constructor to specify integers. See *shell-type-int* or *shell-type-long* for more information.

Examples

Bitwise AND Consider the following document inserted into the collection `switches`:

```
{ _id: 1, expdata: NumberInt(13) }
```

The following `update()` (page 69) operation updates the `expdata` field to the result of a bitwise `and` operation between the current value `NumberInt(13)` (i.e. 1101) and `NumberInt(10)` (i.e. 1010):

```
db.switches.update(
  { _id: 1 },
  { $bit: { expdata: { and: NumberInt(10) } } }
)
```

The bitwise `and` operation results in the integer 8 (i.e. 1000):

```
1101
1010
----
1000
```

And the updated document has the following value for `expdata`:

```
{ "_id" : 1, "expdata" : 8 }
```

The `mongo` (page 527) shell displays `NumberInt(8)` as 8.

Bitwise OR Consider the following document inserted into the collection `switches`:

```
{ _id: 2, expdata: NumberLong(3) }
```

The following `update()` (page 69) operation updates the `expdata` field to the result of a bitwise `or` operation between the current value `NumberLong(3)` (i.e. 0011) and `NumberInt(5)` (i.e. 0101):

```
db.switches.update(
  { _id: 2 },
  { $bit: { expdata: { or: NumberInt(5) } } }
)
```

The bitwise `or` operation results in the integer 7 (i.e. 0111):

```
0011
0101
----
0111
```

And the updated document has the following value for `expdata`:

```
{ "_id" : 2, "expdata" : NumberLong(7) }
```

Bitwise XOR Consider the following document in the collection `switches`:

```
{ _id: 3, expdata: NumberLong(1) }
```

The following `update()` (page 69) operation updates the `expdata` field to the result of a bitwise `xor` operation between the current value `NumberLong(1)` (i.e. 0001) and `NumberInt(5)` (i.e. 0101):

```
db.switches.update(
  { _id: 3 },
  { $bit: { expdata: { xor: NumberInt(5) } } }
)
```

The bitwise `xor` operation results in the integer 4:

```
0001
0101
----
0100
```

And the updated document has the following value for `expdata`:

```
{ "_id" : 3, "expdata" : NumberLong(4) }
```

Isolation

Isolation Update Operator

Name	Description
<code>\$isolated</code> (page 436)	Modifies behavior of multi-updates to increase the isolation of the operation

`$isolated`

`$isolated`

The `$isolated` (page 436) isolation operator **isolates** a write operation that affects multiple documents from other write operations.

Note: The `$isolated` (page 436) isolation operator does **not** provide “all-or-nothing” atomicity for write operations.

Consider the following example:

```
db.foo.update( { field1 : 1 , $isolated : 1 }, { $inc : { field2 : 1 } } , { multi: true } )
```

Without the `$isolated` (page 436) operator, multi-updates will allow other operations to interleave with these updates. If these interleaved operations contain writes, the update operation may produce unexpected results. By specifying `$isolated` (page 436) you can guarantee isolation for the entire multi-update.

Warning: `$isolated` (page 436) does not work with *sharded clusters*.

See also:

See `db.collection.update()` (page 69) for more information about the `db.collection.update()` (page 69) method.

\$atomic

Deprecated since version 2.2: The `$isolated` (page 436) operator replaces `$atomic`.

2.3.3 Aggregation Framework Operators

Pipeline Operators

Note: The *aggregation pipeline* cannot operate on values of the following types: `Symbol`, `MinKey`, `MaxKey`, `DBRef`, `Code`, and `CodeWScope`.

Pipeline operators appear in an array. Documents pass through the operators in a sequence.

Name	Description
<code>\$project</code> (page 438)	Reshapes a document stream. <code>\$project</code> (page 438) can rename, add, or remove fields as well as create computed values and sub-documents.
<code>\$match</code> (page 440)	Filters the document stream, and only allows matching documents to pass into the next pipeline stage. <code>\$match</code> (page 440) uses standard MongoDB queries.
<code>\$redact</code> (page 441)	Restricts the content of a returned document on a per-field level.
<code>\$limit</code> (page 445)	Restricts the number of documents in an aggregation pipeline.
<code>\$skip</code> (page 445)	Skips over a specified number of documents from the pipeline and returns the rest.
<code>\$unwind</code> (page 446)	Takes an array of documents and returns them as a stream of documents.
<code>\$group</code> (page 447)	Groups documents together for the purpose of calculating aggregate values based on a collection of documents.
<code>\$sort</code> (page 449)	Takes all input documents and returns them in a stream of sorted documents.
<code>\$geoNear</code> (page 451)	Returns an ordered stream of documents based on proximity to a geospatial point.
<code>\$out</code> (page 453)	Writes documents from the pipeline to a collection. The <code>\$out</code> (page 453) operator must be the last stage in the pipeline.

Pipeline Aggregation Operators

Name	Description
<code>\$project</code> (page 438)	Reshapes a document stream. <code>\$project</code> (page 438) can rename, add, or remove fields as well as create computed values and sub-documents.
<code>\$match</code> (page 440)	Filters the document stream, and only allows matching documents to pass into the next pipeline stage. <code>\$match</code> (page 440) uses standard MongoDB queries.
<code>\$redact</code> (page 441)	Restricts the content of a returned document on a per-field level.
<code>\$limit</code> (page 445)	Restricts the number of documents in an aggregation pipeline.
<code>\$skip</code> (page 445)	Skips over a specified number of documents from the pipeline and returns the rest.
<code>\$unwind</code> (page 446)	Takes an array of documents and returns them as a stream of documents.
<code>\$group</code> (page 447)	Groups documents together for the purpose of calculating aggregate values based on a collection of documents.
<code>\$sort</code> (page 449)	Takes all input documents and returns them in a stream of sorted documents.
<code>\$geoNear</code> (page 451)	Returns an ordered stream of documents based on proximity to a geospatial point.
<code>\$out</code> (page 453)	Writes documents from the pipeline to a collection. The <code>\$out</code> (page 453) operator must be the last stage in the pipeline.

`$project` (aggregation)**`$project`**

Reshapes a document stream by renaming, adding, or removing fields. Also use `$project` (page 438) to create computed values or sub-documents. Use `$project` (page 438) to:

- Include fields from the original document.
- Insert computed fields.
- Changed in version 2.6: You can use variables in the calculation of computed fields. See `$let` (page 471) and `$map` (page 471). The system variables `$$CURRENT` (page 493) and `$$ROOT` (page 493) are also available directly.
- Rename fields.
- Create and populate fields that hold sub-documents.

Use `$project` (page 438) to quickly select the fields that you want to include or exclude from the response. Consider the following aggregation framework operation.

```
db.article.aggregate(
  { $project : {
    title : 1 ,
    author : 1 ,
  }}
);
```

This operation includes the `title` field and the `author` field in the document that returns from the aggregation *pipeline*.

Note: The `_id` field is always included by default. You may explicitly exclude `_id` as follows:

```
db.article.aggregate(
  { $project : {
    _id : 0 ,
    title : 1 ,
    author : 1
  }}
);
```

Here, the projection excludes the `_id` field but includes the `title` and `author` fields.

Projections can also add computed fields to the document stream passing through the pipeline. A computed field can use any of the *expression operators* (page 454) or for text search, use the `$meta` (page 468) operator. Consider the following example:

```
db.article.aggregate(
  { $project : {
    title : 1,
    doctoredPageViews : { $add:["$pageViews", 10] }
  }}
);
```

Here, the field `doctoredPageViews` represents the value of the `pageViews` field after adding 10 to the original field using the `$add` (page 464).

Note: You must enclose the expression that defines the computed field in braces, so that the expression is a valid object.

You may also use `$project` (page 438) to rename fields. Consider the following example:

```
db.article.aggregate(
  { $project : {
    title : 1 ,
    page_views : "$pageViews" ,
    bar : "$other.foo"
  }}
);
```

This operation renames the `pageViews` field to `page_views`, and renames the `foo` field in the `other` sub-document as the top-level field `bar`. The field references used for renaming fields are direct expressions and do not use an operator or surrounding braces. All aggregation field references can use dotted paths to refer to fields in nested documents.

Finally, you can use the `$project` (page 438) to create and populate new sub-documents. Consider the following example that creates a new object-valued field named `stats` that holds a number of values:

```
db.article.aggregate(
  { $project : {
    title : 1 ,
    stats : {
      pv : "$pageViews",
      foo : "$other.foo",
      dpv : { $add:["$pageViews", 10] }
    }
  }}
);
```

This projection includes the `title` field and places `$project` (page 438) into “inclusive” mode. Then, it creates the `stats` documents with the following fields:

- `pv` which includes and renames the `pageViews` from the top level of the original documents.
- `foo` which includes the value of `other.foo` from the original documents.
- `dpv` which is a computed field that adds 10 to the value of the `pageViews` field in the original document using the `$add` (page 464) aggregation expression.

\$match (aggregation)

\$match

`$match` (page 440) pipes the documents that match its conditions to the next operator in the pipeline.

The `$match` (page 440) query syntax is identical to the *read operation query* syntax.

Examples

Equality Match The following operation uses `$match` (page 440) to perform a simple equality match:

```
db.articles.aggregate(  
  [ { $match : { author : "dave" } } ]  
);
```

The `$match` (page 440) selects the documents where the `author` field equals `dave`, and the aggregation returns the following:

```
{  
  "result" : [  
    {  
      "_id" : ObjectId("512bc95fe835e68f199c8686"),  
      "author": "dave",  
      "score" : 80  
    },  
    {  
      "_id" : ObjectId("512bc962e835e68f199c8687"),  
      "author" : "dave",  
      "score" : 85  
    }  
  ],  
  "ok" : 1  
}
```

Perform a Count The following example selects documents to process using the `$match` (page 440) pipeline operator and then pipes the results to the `$group` (page 447) pipeline operator to compute a count of the documents:

```
db.articles.aggregate( [  
  { $match : { score : { $gt : 70, $lte : 90 } } },  
  { $group: { _id: null, count: { $sum: 1 } } }  
] );
```

In the aggregation pipeline, `$match` (page 440) selects the documents where the `score` is greater than 70 and less than or equal to 90. These documents are then piped to the `$group` (page 447) to perform a count. The aggregation returns the following:

```
{  
  "result" : [  
    {  
      "_id" : null,  
      "count" : 3  
    }  
  ]  
}
```

```

    }
  ],
  "ok" : 1
}

```

Behavior

Pipeline Optimization

- Place the `$match` (page 440) as early in the aggregation *pipeline* as possible. Because `$match` (page 440) limits the total number of documents in the aggregation pipeline, earlier `$match` (page 440) operations minimize the amount of processing down the pipe.
- If you place a `$match` (page 440) at the very beginning of a pipeline, the query can take advantage of *indexes* like any other `db.collection.find()` (page 34) or `db.collection.findOne()` (page 43).

Restrictions

- You cannot use `$where` (page 391) in `$match` (page 440) queries as part of the aggregation pipeline.
- To use `$text` (page 387) in the `$match` (page 440) stage, the `$match` (page 440) stage has to be the first stage of the pipeline.

\$redact (aggregation)

Definition

\$redact

New in version 2.6.

Restricts the contents of the documents based on information stored in the documents themselves.

The `$redact` (page 441) pipeline operator takes an expression that evaluates to `$$DESCEND` (page 442), `$$PRUNE` (page 442), or `$$KEEP` (page 442).

For example, the following `$redact` (page 441) pipeline uses the `$cond` (page 475) expression ²⁵:

```

db.<collection>.aggregate(
  [
    { $redact:
      {
        $cond:
          {
            if: <boolean-expression>,
            then: <"$$DESCEND" | "$$PRUNE" | "$$KEEP">,
            else: <"$$DESCEND" | "$$PRUNE" | "$$KEEP">
          }
        }
      }
    ]
  )

```

²⁵ The `$cond` (page 475) expression supports an alternate syntax that accepts an array instead of a document form. See `$cond` (page 475) for details.

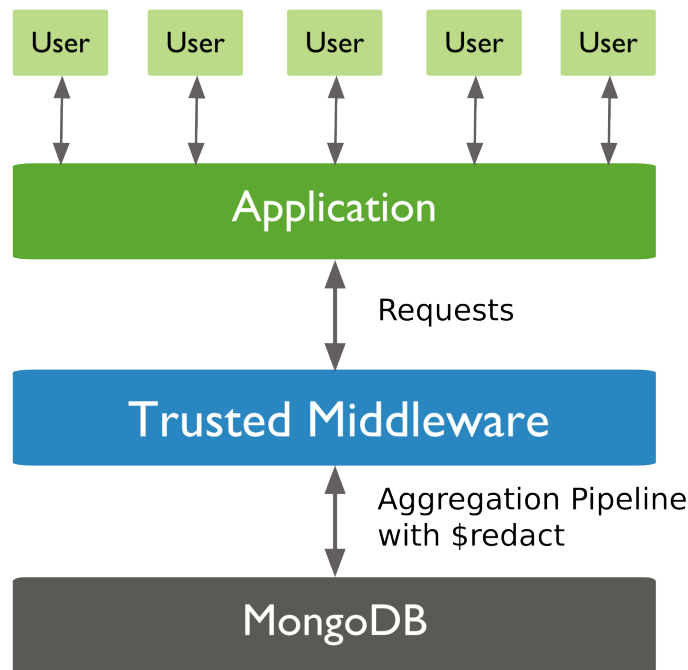


Figure 2.2: Diagram of security architecture with middleware and redaction.

In the example `$cond` (page 475) expression, the `<boolean-expression>` uses a field or fields in the document to specify the conditions for either returning or omitting content.

Tip

To handle documents that are missing field(s) used in `<boolean-expression>`, include `$ifNull` (page 476) in the expression.

System Variable	Description
<code>\$\$DESCEND</code>	<code>\$redact</code> (page 441) returns the <i>non-subdocument</i> fields at the current document/subdocument level. For subdocuments or subdocuments in arrays, apply the <code>\$cond</code> (page 475) expression to the subdocuments to determine access for these subdocuments.
<code>\$\$PRUNE</code>	<code>\$redact</code> (page 441) excludes all fields at this current document/subdocument level, without further inspection of any of the excluded fields. This applies even if the excluded field contains subdocuments that may have different access levels.
<code>\$\$KEEP</code>	<code>\$redact</code> (page 441) returns or keeps all fields at this current document/subdocument level, without further inspection of the fields at this level. This applies even if the included field contains subdocuments that may have different access levels.

See also:

`$cond` (page 475).

Examples The examples in this section use the `db.collection.aggregate()` (page 22) helper provided in the 2.6 version of the `mongo` (page 527) shell.

Evaluate Access at Every Document/Sub-document Level A `forecasts` collection contains documents of the following form where the `tags` field lists the different access values for that document/subdocument level; i.e. a value of ["G", "STLW"] specifies either "G" or "STLW" can access the data:

```
{
  _id: 1,
  title: "123 Department Report",
  tags: [ "G", "STLW" ],
  year: 2014,
  subsections: [
    {
      subtitle: "Section 1: Overview",
      tags: [ "SI", "G" ],
      content: "Section 1: This is the content of section 1."
    },
    {
      subtitle: "Section 2: Analysis",
      tags: [ "STLW" ],
      content: "Section 2: This is the content of section 2."
    },
    {
      subtitle: "Section 3: Budgeting",
      tags: [ "TK" ],
      content: {
        text: "Section 3: This is the content of section3.",
        tags: [ "HCS" ]
      }
    }
  ]
}
```

A user has access to view information with either the tag "STLW" or "G". To run a query on all documents with year 2014 for this user, include a `$redact` (page 441) stage as in the following:

```
var userAccess = [ "STLW", "G" ];
db.forecasts.aggregate(
  [
    { $match: { year: 2014 } },
    { $redact:
      {
        $cond:
          {
            if: { $gt: [ { $size: { $setIntersection: [ "$tags", userAccess ] } }, 0 ] },
            then: "$$DESCEND",
            else: "$$PRUNE"
          }
        }
      }
    ]
  )
```

The aggregation operation returns the following “redacted” document:

```
{
  "_id" : 1,
  "title" : "123 Department Report",
  "tags" : [ "G", "STLW" ],
  "year" : 2014,
  "subsections" : [
```

```
{
  "subtitle" : "Section 1: Overview",
  "tags" : [ "SI", "G" ],
  "content" : "Section 1: This is the content of section 1."
},
{
  "subtitle" : "Section 2: Analysis",
  "tags" : [ "STLW" ],
  "content" : "Section 2: This is the content of section 2."
}
]
```

See also:

[\\$size](#) (page 470), [\\$setIntersection](#) (page 461)

Exclude All Fields at a Given Level A collection `accounts` contains the following document:

```
{
  _id: 1,
  level: 1,
  acct_id: "xyz123",
  cc: {
    level: 5,
    type: "yy",
    num: 0000000000000,
    exp_date: ISODate("2015-11-01T00:00:00.000Z"),
    billing_addr: {
      level: 5,
      addr1: "123 ABC Street",
      city: "Some City"
    },
    shipping_addr: [
      {
        level: 3,
        addr1: "987 XYZ Ave",
        city: "Some City"
      },
      {
        level: 3,
        addr1: "PO Box 0123",
        city: "Some City"
      }
    ]
  },
  status: "A"
}
```

In this example document, the `level` field determines the access level required to view the data.

To run a query on all documents with status `A` and exclude *all* fields contained in a document/subdocument at level 5, include a [\\$redact](#) (page 441) stage that specifies the system variable `"$$PRUNE"` in the `then` field:

```
db.accounts.aggregate(
[
  { $match: { status: "A" } },
  { $redact:
    {
```



```

    $cond: {
      if: { $eq: [ "$level", 5 ] },
      then: "$$PRUNE",
      else: "$$DESCEND"
    }
  }
]
)

```

The `$redact` (page 441) stage evaluates the `level` field to determine access. If the `level` field equals 5, then exclude all fields at that level, even if the excluded field contains subdocuments that may have different `level` values, such as the `shipping_addr` field.

The aggregation operation returns the following “redacted” document:

```

{
  "_id" : 1,
  "level" : 1,
  "acct_id" : "xyz123",
  "status" : "A"
}

```

The result set shows that the `$redact` (page 441) stage excluded the field `cc` as a whole, including the `shipping_addr` field which contained subdocuments that had `level` field values equal to 3 and not 5.

See also:

<http://docs.mongodb.org/manual/tutorial/implement-field-level-redaction> for steps to set up multiple combinations of access for the same data.

\$limit (aggregation)

\$limit

Restricts the number of *documents* that pass through the `$limit` (page 445) in the *pipeline*.

`$limit` (page 445) takes a single numeric (positive whole number) value as a parameter. Once the specified number of documents pass through the pipeline operator, no more will. Consider the following example:

```

db.article.aggregate(
  { $limit : 5 }
);

```

This operation returns only the first 5 documents passed to it from by the pipeline. `$limit` (page 445) has no effect on the content of the documents it passes.

Note: When a `$sort` (page 449) immediately precedes a `$limit` (page 445) in the pipeline, the `$sort` (page 449) operation only maintains the top *n* results as it progresses, where *n* is the specified limit, and MongoDB only needs to store *n* items in memory. This optimization still applies when `allowDiskUse` is `true` and the *n* items exceed the *aggregation memory limit*.

Changed in version 2.4: Before MongoDB 2.4, `$sort` (page 449) would sort all the results in memory, and then limit the results to *n* results.

\$skip (aggregation)

\$skip

Skips over the specified number of *documents* that pass through the `$skip` (page 445) in the *pipeline* before passing all of the remaining input.

`$skip` (page 445) takes a single numeric (positive whole number) value as a parameter. Once the operation has skipped the specified number of documents, it passes all the remaining documents along the *pipeline* without alteration. Consider the following example:

```
db.article.aggregate(  
  { $skip : 5 }  
);
```

This operation skips the first 5 documents passed to it by the pipeline. `$skip` (page 445) has no effect on the content of the documents it passes along the pipeline.

\$unwind (aggregation)

\$unwind

Peels off the elements of an array individually, and returns a stream of documents. `$unwind` (page 446) returns one document for every member of the unwound array within every source document. Take the following aggregation command:

```
db.article.aggregate(  
  { $project : {  
    author : 1 ,  
    title : 1 ,  
    tags : 1  
  } },  
  { $unwind : "$tags" }  
);
```

Note: The dollar sign (i.e. `$`) must precede the field specification handed to the `$unwind` (page 446) operator.

In the above aggregation `$project` (page 438) selects (inclusively) the `author`, `title`, and `tags` fields, as well as the `_id` field implicitly. Then the pipeline passes the results of the projection to the `$unwind` (page 446) operator, which will unwind the `tags` field. This operation may return a sequence of documents that resemble the following for a collection that contains one document holding a `tags` field with an array of 3 items.

```
{  
  "result" : [  
    {  
      "_id" : ObjectId("4e6e4ef557b77501a49233f6"),  
      "title" : "this is my title",  
      "author" : "bob",  
      "tags" : "fun"  
    },  
    {  
      "_id" : ObjectId("4e6e4ef557b77501a49233f6"),  
      "title" : "this is my title",  
      "author" : "bob",  
      "tags" : "good"  
    },  
    {  
      "_id" : ObjectId("4e6e4ef557b77501a49233f6"),  
      "title" : "this is my title",  
      "author" : "bob",  
      "tags" : "fun"  
    }  
  ],  
  "OK" : 1  
}
```

A single document becomes 3 documents: each document is identical except for the value of the `tags` field. Each value of `tags` is one of the values in the original “tags” array.

Note: `$unwind` (page 446) has the following behaviors:

- `$unwind` (page 446) is most useful in combination with `$group` (page 447).
 - You may undo the effects of `unwind` operation with the `$group` (page 447) pipeline operator.
 - If you specify a target field for `$unwind` (page 446) that does not exist in an input document, the pipeline ignores the input document, and will generate no result documents.
 - If you specify a target field for `$unwind` (page 446) that is not an array, `db.collection.aggregate()` (page 22) generates an error.
 - If you specify a target field for `$unwind` (page 446) that holds an empty array (`[]`) in an input document, the pipeline ignores the input document, and will not generate any result documents.
-

`$group` (aggregation)

`$group`

Groups documents together for the purpose of calculating aggregate values based on a collection of documents. In practice, `$group` (page 447) often supports tasks such as average page views for each page in a website on a daily basis.

Important: The output of `$group` (page 447) is not ordered.

The output of `$group` (page 447) depends on how you define groups. Begin by specifying an identifier (i.e. an `_id` field) for the group you’re creating with this pipeline. For this `_id` field, you can specify various expressions, including a single field from the documents in the pipeline, a computed value from a previous stage, a document that consists of multiple fields, and other valid expressions, such as constant or subdocument fields. You can use `$project` (page 438) operators in expressions for the `_id` field.

The following example of an `_id` field specifies a document that consists of multiple fields:

```
{ $group: { _id : { author: '$author', pageViews: '$pageViews', posted: '$posted' } } }
```

Every `$group` (page 447) expression **must** specify an `_id` field. In addition to the `_id` field, `$group` (page 447) expression can include computed fields. These other fields must use one of the following *accumulators*:

- `$addToSet` (page 455)
- `$first` (page 456)
- `$last` (page 456)
- `$max` (page 456)
- `$min` (page 456)
- `$avg` (page 457)
- `$push` (page 458)
- `$sum` (page 460)

With the exception of the `_id` field, `$group` (page 447) cannot output nested documents.

Tip

Use `$project` (page 438) as needed to rename the grouped field after a `$group` (page 447) operation.

Variables Changed in version 2.6.

You can use variables in expressions for the `$group` (page 447) phase. See `$let` (page 471) and `$map` (page 471).

The system variables `$$CURRENT` (page 493) and `$$ROOT` (page 493) are also available directly. See *Group Documents by author* (page 449) for an example.

`$group` Operator and Memory The `$group` (page 447) stage has a limit of 100 megabytes of RAM. By default, if the stage exceeds this limit, `$group` (page 447) will produce an error. However, to allow for the handling of large datasets, set the `allowDiskUse` option to `true` to enable `$group` (page 447) operations to write to temporary files. See the `allowDiskUse` option in `db.collection.aggregate()` (page 22) method and the `aggregate` (page 198) command for details.

Changed in version 2.6: MongoDB introduces a limit of 100 megabytes of RAM for the `$group` (page 447) stage.

Examples

Calculate Count and Sum Consider the following example:

```
db.article.aggregate(  
  { $group : {  
    _id : "$author",  
    docsPerAuthor : { $sum : 1 },  
    viewsPerAuthor : { $sum : "$pageViews" }  
  }}  
);
```

This aggregation pipeline groups by the `author` field and computes two fields, `docsPerAuthor` and `viewsPerAuthor`, per each group. The `docsPerAuthor` field is a counter field that uses the `$sum` (page 460) operator to add 1 for each document with a given author. The `viewsPerAuthor` field is the sum of the values in the `pageViews` field for each group.

Pivot Data A collection `books` contains the following documents:

```
{ "_id" : 8751, "title" : "The Banquet", "author" : "Dante", "copies" : 2 }  
{ "_id" : 8752, "title" : "Divine Comedy", "author" : "Dante", "copies" : 1 }  
{ "_id" : 8645, "title" : "Eclogues", "author" : "Dante", "copies" : 2 }  
{ "_id" : 7000, "title" : "The Odyssey", "author" : "Homer", "copies" : 10 }  
{ "_id" : 7020, "title" : "Iliad", "author" : "Homer", "copies" : 10 }
```

Group title by author The following aggregation operation pivots the data in the `books` collection to have titles grouped by authors.

```
db.books.aggregate(  
  [  
    { $group : { _id : "$author", books: { $push: "$title" } } }  
  ]  
)
```

The operation returns the following documents:

```
{ "_id" : "Homer", "books" : [ "The Odyssey", "Iliad" ] }  
{ "_id" : "Dante", "books" : [ "The Banquet", "Divine Comedy", "Eclogues" ] }
```

Group Documents by author The following aggregation operation uses the `$$ROOT` (page 493) system variable to group the documents by authors. The resulting documents must not exceed the `BSON Document Size` (page 604) limit.

```
db.books.aggregate(
  [
    { $group : { _id : "$author", books: { $push: "$$ROOT" } } }
  ]
)
```

The operation returns the following documents:

```
{
  "_id" : "Homer",
  "books" :
  [
    { "_id" : 7000, "title" : "The Odyssey", "author" : "Homer", "copies" : 10 },
    { "_id" : 7020, "title" : "Iliad", "author" : "Homer", "copies" : 10 }
  ]
}

{
  "_id" : "Dante",
  "books" :
  [
    { "_id" : 8751, "title" : "The Banquet", "author" : "Dante", "copies" : 2 },
    { "_id" : 8752, "title" : "Divine Comedy", "author" : "Dante", "copies" : 1 },
    { "_id" : 8645, "title" : "Eclogues", "author" : "Dante", "copies" : 2 }
  ]
}
```

See also:

Push Current Document Into the Returned Array Field (page 459)

\$sort (aggregation)

\$sort

The `$sort` (page 449) *pipeline* operator sorts all input documents and returns them to the pipeline in sorted order.

Consider the following prototype form:

```
db.<collection-name>.aggregate(
  [
    { $sort : { <sort-key> } }
  ]
)
```

This sorts the documents in the collection named `<collection-name>`, according to the key and specification in the `{ <sort-key> }` document. The sort key has the following syntax:

```
{ field: value }
```

The `{ <sort-key> }` document can specify *ascending or descending sort on existing fields* (page 450) or *sort on computed metadata* (page 450).

Behaviors

Ascending/Descending Sort Specify in the { <sort-key> } document the field or fields to sort by and a value of 1 or -1 to specify an ascending or descending sort respectively, as in the following example:

```
db.users.aggregate(  
  [  
    { $sort : { age : -1, posts: 1 } }  
  ]  
)
```

This operation sorts the documents in the `users` collection, in descending order according by the `age` field and then in ascending order according to the value in the `posts` field.

When comparing values of different *BSON* types, MongoDB uses the following comparison order, from lowest to highest:

1. MinKey (internal type)
2. Null
3. Numbers (ints, longs, doubles)
4. Symbol, String
5. Object
6. Array
7. BinData
8. ObjectId
9. Boolean
10. Date, Timestamp
11. Regular Expression
12. MaxKey (internal type)

MongoDB treats some types as equivalent for comparison purposes. For instance, numeric types undergo conversion before comparison.

The comparison treats a non-existent field as it would an empty BSON Object. As such, a sort on the `a` field in documents { } and { `a`: null } would treat the documents as equivalent in sort order.

With arrays, a less-than comparison or an ascending sort compares the smallest element of arrays, and a greater-than comparison or a descending sort compares the largest element of the arrays. As such, when comparing a field whose value is a single-element array (e.g. [1]) with non-array fields (e.g. 2), the comparison is between 1 and 2. A comparison of an empty array (e.g. []) treats the empty array as less than `null` or a missing field.

Metadata Sort Specify in the { <sort-key> } document, a new field name for the computed metadata and specify the `$meta` (page 468) expression as its value, as in the following example:

```
db.users.aggregate(  
  [  
    { $match: { $text: { $search: "operating" } } },  
    { $sort: { score: { $meta: "textScore" }, posts: -1 } }  
  ]  
)
```

This operation uses the `$text` (page 387) operator to match the documents, and then sorts first by the `"textScore"` metadata and then by descending order of the `posts` field. The specified metadata determines the sort order. For

example, the "textScore" metadata sorts in descending order. See `$meta` (page 468) for more information on metadata.

`$sort` Operator and Memory

`$sort` + `$limit` Memory Optimization When a `$sort` (page 449) immediately precedes a `$limit` (page 445) in the pipeline, the `$sort` (page 449) operation only maintains the top `n` results as it progresses, where `n` is the specified limit, and MongoDB only needs to store `n` items in memory. This optimization still applies when `allowDiskUse` is `true` and the `n` items exceed the *aggregation memory limit*.

Changed in version 2.4: Before MongoDB 2.4, `$sort` (page 449) would sort all the results in memory, and then limit the results to `n` results.

Optimizations are subject to change between releases.

`$sort` and Memory Restrictions The `$sort` (page 449) stage has a limit of 100 megabytes of RAM. By default, if the stage exceeds this limit, `$sort` (page 449) will produce an error. To allow for the handling of large datasets, set the `allowDiskUse` option to `true` to enable `$sort` (page 449) operations to write to temporary files. See the `allowDiskUse` option in `db.collection.aggregate()` (page 22) method and the `aggregate` (page 198) command for details.

Changed in version 2.6: The memory limit for `$sort` (page 449) changed from 10 percent of RAM to 100 megabytes of RAM.

`$sort` Operator and Performance `$sort` (page 449) operator can take advantage of an index when placed at the **beginning** of the pipeline or placed **before** the following aggregation operators: `$project` (page 438), `$unwind` (page 446), and `$group` (page 447).

`$geoNear` (aggregation)

Definition

`$geoNear`

New in version 2.4.

`$geoNear` (page 451) returns documents in order of nearest to farthest from a specified point and pass the documents through the aggregation *pipeline*.

The `$geoNear` (page 451) operator accepts a *document* that contains the following fields. Specify all distances in the same units as the document coordinate system:

`:field` GeoJSON point, `:term`: *legacy coordinate pairs* <*legacy coordinate pairs*> `near`:

The point for which to find the closest documents.

field string `distanceField` The output field that contains the calculated distance. To specify a field within a subdocument, use *dot notation*.

field number `limit` The maximum number of documents to return. The default value is 100. See also the `num` option.

field number `num` The `num` option provides the same function as the `limit` option. Both define the maximum number of documents to return. If both options are included, the `num` value overrides the `limit` value.

field number maxDistance A distance from the center point. Specify the distance in radians. MongoDB limits the results to those documents that fall within the specified distance from the center point.

field document query Limits the results to the documents that match the query. The query syntax is the usual MongoDB *read operation query* syntax.

field Boolean spherical If `true`, MongoDB references points using a spherical surface. The default value is `false`.

field number distanceMultiplier The factor to multiply all distances returned by the query. For example, use the `distanceMultiplier` to convert radians, as returned by a spherical query, to kilometers by multiplying by the radius of the Earth.

field string includeLocs This specifies the output field that identifies the location used to calculate the distance. This option is useful when a location field contains multiple locations. To specify a field within a subdocument, use *dot notation*.

field Boolean uniqueDocs If this value is `true`, the query returns a matching document once, even if more than one of the document's location fields match the query.

Deprecated since version 2.6: Geospatial queries no longer return duplicate results. The `$uniqueDocs` (page 400) operator has no impact on results.

Behavior When using `$geoNear` (page 451), consider that:

- You can only use `$geoNear` (page 451) as the first stage of a pipeline.
- You must include the `distanceField` option. The `distanceField` option specifies the field that will contain the calculated distance.
- The collection must have a `geospatial` index.
- The `$geoNear` (page 451) requires that a collection have *at most* only one `2d` index and/or only one `2dsphere` index.

Generally, the options for `$geoNear` (page 451) are similar to the `geoNear` (page 216) command with the following exceptions:

- `distanceField` is a mandatory field for the `$geoNear` (page 451) pipeline operator; the option does not exist in the `geoNear` (page 216) command.
- `includeLocs` accepts a string in the `$geoNear` (page 451) pipeline operator and a `boolean` in the `geoNear` (page 216) command.

Example The following aggregation finds at most 5 unique documents with a location at most .008 from the center [40.72, -73.99] and have type equal to public:

```
db.places.aggregate([
  {
    $geoNear: {
      near: [40.724, -73.997],
      distanceField: "dist.calculated",
      maxDistance: 0.008,
      query: { type: "public" },
      includeLocs: "dist.location",
      num: 5
    }
  }
])
```


The aggregation returns the following:

```
{
  "result" : [
    {
      "_id" : 7,
      "name" : "Washington Square",
      "type" : "public",
      "location" : [
        [ 40.731, -73.999 ],
        [ 40.732, -73.998 ],
        [ 40.730, -73.995 ],
        [ 40.729, -73.996 ]
      ],
      "dist" : {
        "calculated" : 0.0050990195135962296,
        "location" : [ 40.729, -73.996 ]
      }
    },
    {
      "_id" : 8,
      "name" : "Sara D. Roosevelt Park",
      "type" : "public",
      "location" : [
        [ 40.723, -73.991 ],
        [ 40.723, -73.990 ],
        [ 40.715, -73.994 ],
        [ 40.715, -73.994 ]
      ],
      "dist" : {
        "calculated" : 0.006082762530298062,
        "location" : [ 40.723, -73.991 ]
      }
    }
  ],
  "ok" : 1
}
```

The matching documents in the `result` field contain two new fields:

- `dist.calculated` field that contains the calculated distance, and
- `dist.location` field that contains the location used in the calculation.

\$out (aggregation) New in version 2.6.

Definition

\$out

Takes the documents returned by the aggregation pipeline and writes them to a specified collection. The `$out` (page 453) operator lets the aggregation framework return result sets of any size. The `$out` (page 453) operator must be *the last stage* in the pipeline.

The command has the following syntax, where `<output-collection>` is collection that will hold the output of the aggregation operation. `$out` (page 453) is only permissible at the end of the pipeline:

```
db.<collection>.aggregate( [
  { <operation> },
  { <operation> },
  ...,
```

```
    { $out : "<output-collection>" }  
  ] )
```

Important:

- You cannot specify a sharded collection as the output collection. The input collection for a pipeline can be sharded.
 - The `$out` (page 453) operator cannot write results to a capped collection.
-

Behaviors

Create New Collection The `$out` (page 453) operation creates a new collection in the current database if one does not already exist. The collection is not visible until the aggregation completes. If the aggregation fails, MongoDB does not create the collection.

Replace Existing Collection If the collection specified by the `$out` (page 453) operation already exists, then upon completion of the aggregation, the `$out` (page 453) stage atomically replaces the existing collection with the new results collection. The `$out` (page 453) operation does not change any indexes that existed on the previous collection. If the aggregation fails, the `$out` (page 453) operation makes no changes to the pre-existing collection.

Index Constraints The pipeline will fail to complete if the documents produced by the pipeline would violate any unique indexes, including the index on the `_id` field of the original output collection.

Example A collection `books` contains the following documents:

```
{ "_id" : 8751, "title" : "The Banquet", "author" : "Dante", "copies" : 2 }  
{ "_id" : 8752, "title" : "Divine Comedy", "author" : "Dante", "copies" : 1 }  
{ "_id" : 8645, "title" : "Eclogues", "author" : "Dante", "copies" : 2 }  
{ "_id" : 7000, "title" : "The Odyssey", "author" : "Homer", "copies" : 10 }  
{ "_id" : 7020, "title" : "Iliad", "author" : "Homer", "copies" : 10 }
```

The following aggregation operation pivots the data in the `books` collection to have titles grouped by authors and then writes the results to the `authors` collection.

```
db.books.aggregate( [  
    { $group : { _id : "$author", books : { $push : "$title" } } },  
    { $out : "authors" }  
  ] )
```

After the operation, the `authors` collection contains the following documents:

```
{ "_id" : "Homer", "books" : [ "The Odyssey", "Iliad" ] }  
{ "_id" : "Dante", "books" : [ "The Banquet", "Divine Comedy", "Eclogues" ] }
```

Expression Operators

Expression operators calculate values within the *Pipeline Operators* (page 437).

\$group Operators**Group Aggregation Operators**

Name	Description
<code>\$addToSet</code> (page 455)	Returns an array of all the <i>unique</i> values for the selected field among for each document in that group.
<code>\$first</code> (page 456)	Returns the first value in a group.
<code>\$last</code> (page 456)	Returns the last value in a group.
<code>\$max</code> (page 456)	Returns the highest value in a group.
<code>\$min</code> (page 456)	Returns the lowest value in a group.
<code>\$avg</code> (page 457)	Returns an average of all the values in a group.
<code>\$push</code> (page 458)	Returns an array of <i>all</i> values for the selected field among for each document in that group.
<code>\$sum</code> (page 460)	Returns the sum of all the values in a group.

\$addToSet (aggregation)**\$addToSet**

Returns an array of all the values found in the selected field among the documents in that group. *Every unique value only appears once* in the result set. There is no ordering guarantee for the output documents.

Example In the `mongo` (page 527) shell, insert documents into a collection named `products` using the following operation:

```
db.products.insert( [
  { "type" : "phone", "price" : 389.99, "stocked" : 270000 },
  { "type" : "phone", "price" : 376.99 , "stocked" : 97000},
  { "type" : "phone", "price" : 389.99 , "stocked" : 97000},
  { "type" : "chair", "price" : 59.99, "stocked" : 108 }
] )
```

Use the following `db.collection.aggregate()` (page 22) operation in the `mongo` (page 527) shell with the `$addToSet` (page 455) operator:

```
db.products.aggregate( {
  $group : {
    _id : "$type",
    price: { $addToSet: "$price" },
    stocked: { $addToSet: "$stocked" },
  } }
)
```

This aggregation pipeline returns documents grouped on the value of the `type` field, with *sets* constructed from the unique values of the `price` and `stocked` fields in the input documents. Consider the following aggregation results:

```
{
  "_id" : "chair",
  "price" : [
    59.99
  ],
  "stocked" : [
    108
  ]
},
{
  "_id" : "phone",
  "price" : [
    376.99,
```

```
        389.99
    ],
    "stocked" : [
        97000,
        270000,
    ]
}
```

\$first (aggregation)

\$first

Returns the first value it encounters for its group.

Note: Only use `$first` (page 456) when the `$group` (page 447) follows a `$sort` (page 449) operation. Otherwise, the result of this operation is unpredictable.

\$last (aggregation)

\$last

Returns the last value it encounters for its group.

Note: Only use `$last` (page 456) when the `$group` (page 447) follows an `$sort` (page 449) operation. Otherwise, the result of this operation is unpredictable.

\$max (aggregation)

\$max

Returns the highest value among all values of the field in all documents selected by this group.

\$min (aggregation)

\$min

The `$min` (page 456) operator returns the lowest non-null value of a field in the documents for a `$group` (page 447) operation.

Changed in version 2.4: If some, **but not all**, documents for the `$min` (page 456) operation have either a `null` value for the field or are missing the field, the `$min` (page 456) operator only considers the non-null and the non-missing values for the field. If **all** documents for the `$min` (page 456) operation have `null` value for the field or are missing the field, the `$min` (page 456) operator returns `null` for the minimum value.

Before 2.4, if any of the documents for the `$min` (page 456) operation were missing the field, the `$min` (page 456) operator would not return any value. If any of the documents for the `$min` (page 456) had the value `null`, the `$min` (page 456) operator would return a `null`.

Example

The `users` collection contains the following documents:

```
{ "_id" : "abc001", "age" : 25 }
{ "_id" : "abe001", "age" : 35 }
{ "_id" : "efg001", "age" : 20 }
{ "_id" : "xyz001", "age" : 15 }
```

- To find the minimum value of the `age` field from all the documents, use the `$min` (page 456) operator:

```
db.users.aggregate( [ { $group: { _id:0, minAge: { $min: "$age" } } } ] )
```

The operation returns the value of the age field in the minAge field:

```
{ "result" : [ { "_id" : 0, "minAge" : 15 } ], "ok" : 1 }
```

- To find the minimum value of the age field for only those documents with `_id` starting with the letter a, use the `$min` (page 456) operator after a `$match` (page 440) operation:

```
db.users.aggregate( [ { $match: { _id: /^a/ } },
                      { $group: { _id: 0, minAge: { $min: "$age" } } }
                    ] )
```

The operation returns the minimum value of the age field for the two documents with `_id` starting with the letter a:

```
{ "result" : [ { "_id" : 0, "minAge" : 25 } ], "ok" : 1 }
```

Example

The `users` collection contains the following documents where some of the documents are either missing the age field or the age field contains null:

```
{ "_id" : "abc001", "age" : 25 }
{ "_id" : "abe001", "age" : 35 }
{ "_id" : "efg001", "age" : 20 }
{ "_id" : "xyz001", "age" : 15 }
{ "_id" : "xxx001" }
{ "_id" : "zzz001", "age" : null }
```

- The following operation finds the minimum value of the age field in all the documents:

```
db.users.aggregate( [ { $group: { _id:0, minAge: { $min: "$age" } } } ] )
```

Because only some documents for the `$min` (page 456) operation are missing the age field or have age field equal to null, `$min` (page 456) only considers the non-null and the non-missing values and the operation returns the following document:

```
{ "result" : [ { "_id" : 0, "minAge" : 15 } ], "ok" : 1 }
```

- The following operation finds the minimum value of the age field for only those documents where the `_id` equals "xxx001" or "zzz001":

```
db.users.aggregate( [ { $match: { _id: { $in: [ "xxx001", "zzz001" ] } } },
                      { $group: { _id: 0, minAge: { $min: "$age" } } }
                    ] )
```

The `$min` (page 456) operation returns null for the minimum age since **all** documents for the `$min` (page 456) operation have null value for the field age or are missing the field:

```
{ "result" : [ { "_id" : 0, "minAge" : null } ], "ok" : 1 }
```

`$avg` (aggregation)

`$avg`

Returns the average of all the values of the field in all documents selected by this group.

\$push (aggregation)**\$push**

Returns an array of all the values found in the selected field among the documents in that group. A value may appear *more than once* in the result set if more than one field in the grouped documents has that value.

Example The following examples use the following collection named `users` as the input for the aggregation pipeline:

```
{ "_id": 1, "user" : "Jan", "age" : 25, "score": 80 }
{ "_id": 2, "user" : "Mel", "age" : 35, "score": 70 }
{ "_id": 3, "user" : "Ty", "age" : 20, "score": 102 }
{ "_id": 4, "user" : "Lee", "age" : 25, "score": 45 }
```

Push Values of a Single Field Into the Returned Array Field To group by `age` and return all the `user` values for each age, use the `$push` (page 458) operator.

```
db.users.aggregate(
  {
    $group: {
      _id: "$age",
      users: { $push: "$user" }
    }
  }
)
```

For each age, the operation returns the field `users` that contains an array of all the `user` values associated with that age:

```
{
  "result" : [
    {
      "_id" : 20,
      "users" : [
        "Ty"
      ]
    },
    {
      "_id" : 35,
      "users" : [
        "Mel"
      ]
    },
    {
      "_id" : 25,
      "users" : [
        "Jan",
        "Lee"
      ]
    }
  ],
  "ok" : 1
}
```

Push Documents Into the Returned Array Field The `$push` (page 458) operator can return an array of documents.

To group by age and return all the user and associated score values for each age, use the `$push` (page 458) operator.

```
db.users.aggregate(
  {
    $group: {
      _id: "$age",
      users: { $push: { userid: "$user", score: "$score" } }
    }
  }
)
```

For each age, the operation returns the field `users` that contains an array of documents. These documents contain the fields `userid` and `score` that hold respectively the user value and the score value associated with that age:

```
{
  "result" : [
    {
      "_id" : 20,
      "users" : [
        {
          "userid" : "Ty",
          "score" : 102
        }
      ]
    },
    {
      "_id" : 35,
      "users" : [
        {
          "userid" : "Mel",
          "score" : 70
        }
      ]
    },
    {
      "_id" : 25,
      "users" : [
        {
          "userid" : "Jan",
          "score" : 80
        },
        {
          "userid" : "Lee",
          "score" : 45
        }
      ]
    }
  ],
  "ok" : 1
}
```

Push Current Document Into the Returned Array Field The `$push` (page 458) operator can use the system variable `$$ROOT` (page 493) to push the current document being processed into the array. The resulting documents must not exceed the `BSON Document Size` (page 604) limit.

To group by age and return the documents containing that age, use the `$push` (page 458) operator with `$$ROOT` (page 493),

```
db.users.aggregate(  
  {  
    $group:  
    {  
      _id: "$age",  
      users: { $push: "$$ROOT" }  
    }  
  }  
)
```

The operation returns the following documents:

```
{  
  "_id" : 20,  
  "users" : [ { "_id" : 3, "user" : "Ty", "age" : 20, "score" : 102 } ]  
}  
{  
  "_id" : 35,  
  "users" : [ { "_id" : 2, "user" : "Mel", "age" : 35, "score" : 70 } ]  
}  
{  
  "_id" : 25,  
  "users" :  
    [  
      { "_id" : 1, "user" : "Jan", "age" : 25, "score" : 80 },  
      { "_id" : 4, "user" : "Lee", "age" : 25, "score" : 45 }  
    ]  
}
```

\$sum (aggregation)

\$sum

Returns the sum of all the values for a specified field in the grouped documents.

Alternately, if you specify a value as an argument, `$sum` (page 460) will increment this field by the specified value for every document in the grouping. Typically, specify a value of 1 in order to count members of the group.

Boolean Operators

These operators accept Booleans as arguments and return Booleans as results.

The operators convert non-Booleans to Boolean values according to the BSON standards. Here, `null`, `undefined`, and `0` values become `false`, while non-zero numeric values, and all other types, such as strings, dates, objects become `true`.

Boolean Aggregation Operators

Name	Description
<code>\$and</code> (page 460)	Returns true only when <i>all</i> values in its input array are true.
<code>\$or</code> (page 461)	Returns true when <i>any</i> value in its input array are true.
<code>\$not</code> (page 461)	Returns the boolean value that is the opposite of the input value.

\$and (aggregation)

\$and

Takes an array one or more values and returns `true` if *all* of the values in the array are `true`. Otherwise `$and` (page 460) returns `false`.

Note: `$and` (page 460) uses short-circuit logic: the operation stops evaluation after encountering the first `false` expression.

\$or (aggregation)

\$or

Takes an array of one or more values and returns `true` if *any* of the values in the array are `true`. Otherwise `$or` (page 461) returns `false`.

Note: `$or` (page 461) uses short-circuit logic: the operation stops evaluation after encountering the first `true` expression.

\$not (aggregation)

\$not

Returns the boolean opposite value passed to it. When passed a `true` value, `$not` (page 461) returns `false`; when passed a `false` value, `$not` (page 461) returns `true`.

Set Operators

These operators provide operations on sets.

Set Operators (Aggregation)

Name	Description
<code>\$setEquals</code> (page 461)	Returns <code>true</code> if two sets have the same elements.
<code>\$setIntersection</code> (page 461)	Returns the common elements of the input sets.
<code>\$setDifference</code> (page 462)	Returns elements of a set that do not appear in a second set.
<code>\$setUnion</code> (page 462)	Returns a set that holds all elements of the input sets.
<code>\$setIsSubset</code> (page 462)	Returns <code>true</code> if all elements of a set appear in a second set.
<code>\$anyElementTrue</code> (page 462)	Returns <code>true</code> if <i>any</i> elements of a set evaluate to <code>true</code> , and <code>false</code> otherwise.
<code>\$allElementsTrue</code> (page 462)	Returns <code>true</code> if <i>all</i> elements of a set evaluate to <code>true</code> , and <code>false</code> otherwise.

\$setEquals (aggregation)

\$setEquals

New in version 2.6.

Takes two arrays and returns `true` when they contain the same elements, and `false` otherwise.

Important: `$setEquals` (page 461) takes arrays as arguments and treats these arrays as sets. `$setEquals` (page 461) ignores duplicate entries in input arrays and produce arrays that contain unique entries.

\$setIntersection (aggregation)

\$setIntersection

New in version 2.6.

Takes any number of arrays and returns an array that contains the elements that appear in every input array.

Important: `$setIntersection` (page 461) takes arrays as arguments and treats these arrays as sets. `$setIntersection` (page 461) ignores duplicate entries in input arrays and produce arrays that contain unique entries.

\$setDifference (aggregation)

\$setDifference

New in version 2.6.

Takes two arrays and returns an array containing the elements that only exist in the first array.

Important: `$setDifference` (page 462) takes arrays as arguments and treats these arrays as sets. `$setDifference` (page 462) ignores duplicate entries in input arrays and produce arrays that contain unique entries.

\$setUnion (aggregation)

\$setUnion

New in version 2.6.

Takes any number of arrays and returns an array containing the elements that appear in any input array.

Important: `$setUnion` (page 462) takes arrays as arguments and treats these arrays as sets. `$setUnion` (page 462) ignores duplicate entries in input arrays and produce arrays that contain unique entries.

\$setIsSubset (aggregation)

\$setIsSubset

New in version 2.6.

Takes two arrays and returns `true` when the first array is a subset of the second and `false` otherwise.

Important: `$setIsSubset` (page 462) takes arrays as arguments and treats these arrays as sets. `$setIsSubset` (page 462) ignores duplicate entries in input arrays and produce arrays that contain unique entries.

\$allElementsTrue (aggregation)

\$allElementsTrue

New in version 2.6.

Takes a single expression that returns an array and returns `true` if all its values are `true` and `false` otherwise. An empty array returns `true`.

\$anyElementTrue (aggregation)

\$anyElementTrue

New in version 2.6.

Takes a single expression that returns an array and returns `true` if any of its values are `true` and `false` otherwise. An empty array returns `false`.

Comparison Operators

These operators perform comparisons between two values and return a Boolean, in most cases reflecting the result of the comparison.

All comparison operators take an array with a pair of values. You may compare numbers, strings, and dates. Except for `$cmp` (page 463), all comparison operators return a Boolean value. `$cmp` (page 463) returns an integer.

Comparison Aggregation Operators

Name	Description
\$cmp (page 463)	Compares two values and returns the result of the comparison as an integer.
\$eq (page 463)	Takes two values and returns true if the values are equivalent.
\$gt (page 463)	Takes two values and returns true if the first is larger than the second.
\$gte (page 463)	Takes two values and returns true if the first is larger than or equal to the second.
\$lt (page 463)	Takes two values and returns true if the second value is larger than the first.
\$lte (page 463)	Takes two values and returns true if the second value is larger than or equal to the first.
\$ne (page 464)	Takes two values and returns true if the values are <i>not</i> equivalent.

\$cmp (aggregation)**\$cmp**

Takes two values in an array and returns an integer. The returned value is:

- A negative number if the first value is less than the second.
- A positive number if the first value is greater than the second.
- 0 if the two values are equal.

\$eq (aggregation)**\$eq**

Takes two values in an array and returns a boolean. The returned value is:

- `true` when the values are equivalent.
- `false` when the values are **not** equivalent.

\$gt (aggregation)**\$gt**

Takes two values in an array and returns a boolean. The returned value is:

- `true` when the first value is *greater than* the second value.
- `false` when the first value is *less than or equal to* the second value.

\$gte (aggregation)**\$gte**

Takes two values in an array and returns a boolean. The returned value is:

- `true` when the first value is *greater than or equal to* the second value.
- `false` when the first value is *less than* the second value.

\$lt (aggregation)**\$lt**

Takes two values in an array and returns a boolean. The returned value is:

- `true` when the first value is *less than* the second value.
- `false` when the first value is *greater than or equal to* the second value.

\$lte (aggregation)**\$lte**

Takes two values in an array and returns a boolean. The returned value is:

- `true` when the first value is *less than or equal to* the second value.

- `false` when the first value is *greater than* the second value.

\$ne (aggregation)

\$ne

Takes two values in an array returns a boolean. The returned value is:

- `true` when the values are **not equivalent**.
- `false` when the values are **equivalent**.

Arithmetic Operators

Arithmetic operators support only numbers.

Arithmetic Aggregation Operators

Name	Description
\$add (page 464)	Computes the sum of an array of numbers.
\$divide (page 464)	Takes two numbers and divides the first number by the second.
\$mod (page 464)	Takes two numbers and calculates the modulo of the first number divided by the second.
\$multiply (page 464)	Computes the product of an array of numbers.
\$subtract (page 464)	Takes an array that contains two numbers or two dates and subtracts the second value from the first.

\$add (aggregation)

\$add

Takes an array of one or more numbers and adds them together, returning the sum.

\$divide (aggregation)

\$divide

Takes an array that contains a pair of numbers and returns the value of the first number divided by the second number.

\$mod (aggregation)

\$mod

Takes an array that contains a pair of numbers and returns the *remainder* of the first number divided by the second number.

See also:

[\\$mod](#) (page 384)

\$multiply (aggregation)

\$multiply

Takes an array of one or more numbers and multiplies them, returning the resulting product.

\$subtract (aggregation)

\$subtract

Takes an array that contains two numbers or two dates and subtracts the second value from the first. With dates, the operator returns the difference in milliseconds.

String Operators

String operators that manipulate strings.

String Aggregation Operators

Name	Description
<code>\$concat</code> (page 465)	Concatenates two strings.
<code>\$strcasecmp</code> (page 467)	Compares two strings and returns an integer that reflects the comparison.
<code>\$substr</code> (page 468)	Takes a string and returns portion of that string.
<code>\$toLowerCase</code> (page 468)	Converts a string to lowercase.
<code>\$toUpperCase</code> (page 468)	Converts a string to uppercase.

`$concat` (aggregation)

`$concat`

New in version 2.4.

Takes an array of strings, concatenates the strings, and returns the concatenated string. `$concat` (page 465) can only accept an array of strings.

Use `$concat` (page 465) with the following syntax:

```
{ $concat: [ <string>, <string>, ... ] }
```

If array element has a value of `null` or refers to a field that is missing, `$concat` (page 465) will return `null`.

Example

Project new concatenated values.

A collection `menu` contains the documents that stores information on menu items separately in the `section`, the `category` and the `type` fields, as in the following:

```
{ _id: 1, item: { sec: "dessert", category: "pie", type: "apple" } }
{ _id: 2, item: { sec: "dessert", category: "pie", type: "cherry" } }
{ _id: 3, item: { sec: "main", category: "pie", type: "shepherd's" } }
{ _id: 4, item: { sec: "main", category: "pie", type: "chicken pot" } }
```

The following operation uses `$concat` (page 465) to concatenate the `type` field from the sub-document `item`, a space, and the `category` field from the sub-document `item` to project a new `food` field:

```
db.menu.aggregate( { $project: { food:
                        { $concat: [ "$item.type",
                                      " ",
                                      "$item.category"
                                    ]
                        }
                    }
                }
            )
```

The operation returns the following result set where the `food` field contains the concatenated strings:

```
{
  "result" : [
    { "_id" : 1, "food" : "apple pie" },
    { "_id" : 2, "food" : "cherry pie" },
    { "_id" : 3, "food" : "shepherd's pie" },
    { "_id" : 4, "food" : "chicken pot pie" }
  ],
}
```

```
"ok" : 1
}
```

Example

Group by a concatenated string.

A collection `menu` contains the documents that stores information on menu items separately in the `section`, the `category` and the `type` fields, as in the following:

```
{ _id: 1, item: { sec: "dessert", category: "pie", type: "apple" } }
{ _id: 2, item: { sec: "dessert", category: "pie", type: "cherry" } }
{ _id: 3, item: { sec: "main", category: "pie", type: "shepherd's" } }
{ _id: 4, item: { sec: "main", category: "pie", type: "chicken pot" } }
```

The following aggregation uses `$concat` (page 465) to concatenate the `sec` field from the sub-document `item`, the string `" : "`, and the `category` field from the sub-document `item` to group by the new concatenated string and perform a count:

```
db.menu.aggregate( { $group: { _id:
                          { $concat: [ "$item.sec",
                                        ": ",
                                        "$item.category"
                                      ]
                          },
                          count: { $sum: 1 }
                        }
                    }
                )
```

The aggregation returns the following document:

```
{
  "result" : [
    { "_id" : "main: pie", "count" : 2 },
    { "_id" : "dessert: pie", "count" : 2 }
  ],
  "ok" : 1
}
```

Example

Concatenate null or missing values.

A collection `menu` contains the documents that stores information on menu items separately in the `section`, the `category` and the `type` fields. Not all documents have the all three fields. For example, the document with `_id` equal to 5 is missing the `category` field:

```
{ _id: 1, item: { sec: "dessert", category: "pie", type: "apple" } }
{ _id: 2, item: { sec: "dessert", category: "pie", type: "cherry" } }
{ _id: 3, item: { sec: "main", category: "pie", type: "shepherd's" } }
{ _id: 4, item: { sec: "main", category: "pie", type: "chicken pot" } }
{ _id: 5, item: { sec: "beverage", type: "coffee" } }
```

The following aggregation uses the `$concat` (page 465) to concatenate the `type` field from the sub-document `item`, a space, and the `category` field from the sub-document `item`:

```

db.menu.aggregate( { $project: { food:
                                { $concat: [ "$item.type",
                                              " ",
                                              "$item.category"
                                            ]
                                }
                              }
                    }
)

```

Because the document with `_id` equal to 5 is missing the `type` field in the `item` sub-document, `$concat` (page 465) returns the value `null` as the concatenated value for the document:

```

{
  "result" : [
    { "_id" : 1, "food" : "apple pie" },
    { "_id" : 2, "food" : "cherry pie" },
    { "_id" : 3, "food" : "shepherd's pie" },
    { "_id" : 4, "food" : "chicken pot pie" },
    { "_id" : 5, "food" : null }
  ],
  "ok" : 1
}

```

To handle possible missing fields, you can use `$ifNull` (page 476) with `$concat` (page 465), as in the following example which substitutes `<unknown type>` if the field `type` is null or missing, and `<unknown category>` if the field `category` is null or is missing:

```

db.menu.aggregate( { $project: { food:
                                { $concat: [ { $ifNull: ["$item.type", "<unknown type>"]
                                              " ",
                                              { $ifNull: ["$item.category", "<unknown cate
                                            ]
                                }
                              }
                    }
)

```

The aggregation returns the following result set:

```

{
  "result" : [
    { "_id" : 1, "food" : "apple pie" },
    { "_id" : 2, "food" : "cherry pie" },
    { "_id" : 3, "food" : "shepherd's pie" },
    { "_id" : 4, "food" : "chicken pot pie" },
    { "_id" : 5, "food" : "coffee <unknown category>" }
  ],
  "ok" : 1
}

```

\$strcasecmp (aggregation)

\$strcasecmp

Takes in two strings. Returns a number. `$strcasecmp` (page 467) is positive if the first string is “greater than” the second and negative if the first string is “less than” the second. `$strcasecmp` (page 467) returns 0 if the strings are identical.

Note: `$strcasecmp` (page 467) may not make sense when applied to glyphs outside the Roman alphabet.

`$strcasecmp` (page 467) internally capitalizes strings before comparing them to provide a case-*insensitive* comparison. Use `$cmp` (page 463) for a case sensitive comparison.

\$substr (aggregation)

\$substr

`$substr` (page 468) takes a string and two numbers. The first number represents the number of bytes in the string to skip, and the second number specifies the number of bytes to return from the string.

Note: `$substr` (page 468) is not encoding aware and if used improperly may produce a result string containing an invalid UTF-8 character sequence.

\$toLower (aggregation)

\$toLower

Takes a single string and converts that string to lowercase, returning the result. All uppercase letters become lowercase.

Note: `$toLower` (page 468) may not make sense when applied to glyphs outside the Roman alphabet.

\$toUpper (aggregation)

\$toUpper

Takes a single string and converts that string to uppercase, returning the result. All lowercase letters become uppercase.

Note: `$toUpper` (page 468) may not make sense when applied to glyphs outside the Roman alphabet.

Text Search Operators

Operators to support text search.

Text Search Aggregation Operators

Name	Description
<code>\$meta</code> (page 468)	Access metadata for <code>\$sort</code> (page 449) stage or <code>\$project</code> (page 438)

\$meta (aggregation)

\$meta

New in version 2.6.

The `$meta` (page 468) operator returns the metadata associated with a document in a pipeline operations, e.g. "textScore" when performing text search.

A `$meta` (page 468) expression has the following syntax:

```
{ <projectedFieldName>: { $meta: <metaDataKeyword> } }
```

The `$meta` (page 468) expression can specify the following keyword as the `<metaDataKeyword>`:

Key-word	Description	Sort Order
"textScore"	Returns the score associated with the corresponding query: <i>\$text</i> query for each matching document. The text score signifies how well the document matched the stemmed term or terms. If not used in conjunction with a query: <i>\$text</i> query, returns a score of 0.0	Descending

Behaviors The *\$meta* (page 468) expression can be a part of the *\$project* (page 438) stage and the *\$sort* (page 449) stage.

Projected Field Name If the specified <projectedFieldName> already exists in the matching documents, in the result set, the existing fields will return with the *\$meta* (page 468) values instead of with the stored values.

Projection The *\$meta* (page 468) expression can be used in the *\$project* (page 438) stage, as in:

```
db.articles.aggregate(
  [
    { $match: { $text: { $search: "cake" } } },
    { $project: { title: 1, score: { $meta: "textScore" } } }
  ]
)
```

The inclusion of the *\$meta* (page 468) aggregation expression in the *\$project* (page 438) pipeline specifies both the inclusion of the metadata *as well as* the exclusion of the fields, other than *_id*, that are *not* explicitly included in the projection document. This differs from the behavior of the *\$meta* (page 410) projection operator in a *db.collection.find()* (page 34) operation which only signifies the inclusion of the metadata and does *not* signify an exclusion of other fields.

Sort To use the metadata to sort, specify the *\$meta* (page 468) expression in *\$sort* (page 449) stage, as in:

```
db.articles.aggregate(
  [
    { $match: { $text: { $search: "cake tea" } } },
    { $sort: { score: { $meta: "textScore" } } },
    { $project: { title: 1, _id: 0 } }
  ]
)
```

The specified metadata determines the sort order. For example, the "textScore" metadata sorts in descending order.

Examples For examples of "textScore" projections and sorts, see <http://docs.mongodb.org/manual/tutorial/text>

Array Operators

Operators that manipulate arrays.

Array Aggregation Operators

Name	Description
<i>\$size</i> (page 470)	Returns the size of the array.

\$size (aggregation) New in version 2.6.

Definition

\$size

Counts and returns the total the number of items in an array. Consider the following syntax:

```
{ <field>: { $size: <array> } }
```

Example Given a `survey` collection that records town occupants by household, and that includes documents similar to the following:

```
{
  "_id" : ObjectId("524d82e535edde4707c684c5"),
  "household" : "Carter",
  "st_num" : 300,
  "street" : "North Bond Street",
  "occupants" : [
    "Amy",
    "Donnel",
    "Jack",
    "James"
  ]
}
```

The following aggregation pipeline operation counts the number of residents on each street. The pipeline groups documents according to the `street` field and uses the `$size` (page 405) operator to count the entries in each household's `occupants` array. The pipeline uses `$sum` (page 460) to add the number of residents on each street.

```
db.survey.aggregate(
  [
    {
      $group: {
        _id: "$street",
        numResidents: { $sum: { $size: "$occupants" } }
      }
    }
  ]
)
```

Projection Expressions

Operators that increase the flexibility within aggregation projection and projection-like expressions. These operators are available in the `$project` (page 438), `$group` (page 447), and `$redact` (page 441) pipeline stages.

Aggregation Projection Expressions

Name	Description
<code>\$map</code> (page 471)	Applies a sub-expression to each item in an array and returns the result of the sub-expression.
<code>\$let</code> (page 471)	Defines variables for use within the scope of an aggregation expression.
<code>\$literal</code> (page 472)	Forces the aggregation pipeline to return a literal value without evaluating the expression.

\$map (aggregation)

Definition**\$map**

`$map` (page 471) applies a sub-expression to each item in an array and returns an array with the result of the sub-expression.

`$map` (page 471) is available in the `$project` (page 438), `$group` (page 447), and `$redact` (page 441) pipeline stages.

Example Given an input document that resembles the following:

```
{ skews: [ 1, 1, 2, 3, 5, 8 ] }
```

And the following `$project` (page 438) statement:

```
{ $project: { adjustments: { $map: { input: "$skews",
                                   as: "adj",
                                   in: { $add: [ "$$adj", 12 ] } } } } }
```

The `$map` (page 471) would transform the input document into the following output document:

```
{ adjustments: [ 13, 13, 14, 15, 17, 20 ] }
```

See also:

`$let` (page 471)

\$let (aggregation)**Definition****\$let**

Binds *variables* (page 492) for use in subexpressions. To access the variable in the subexpressions, use a string with the variable name prefixed with double dollar signs (`$$`).

The `$let` (page 471) expression has the following syntax:

```
{
  $let:
    {
      vars: { <var1>: <value1>, ... },
      in: { <expression using "$$var1", ...> }
    }
}
```

Returns The value of the subexpression evaluated with the bound variables.

See *Variables in Aggregation* (page 492) for more information on using variables in the aggregation pipeline.

Behavior In the `vars: { <var1>: <value1>, ... }` assignment block, the order of the assignment does not matter, and using `$$var` to access a variable's value refers to the existing value, if any, of the variable. Even if the variable is being reassigned, `$$var` would refer to the current and not the reassigned value in the assignment block.

For example, the following `$let` (page 471) expression is invalid since in the `vars: { low: 1, high: "$$low" }` assignment block, `$$low` refers to the pre-assignment value of the variable `low`, which is undefined:

```
{
  $let:
  {
    vars: { low: 1, high: "$$low" },
    in: { $gt: [ "$$low", "$$high" ] }
  }
}
```

`$let` (page 471) can access variables defined outside its expression block, including *system variables* (page 492). If you modify the values of externally defined variables in the `vars` block, the new values take effect only in the `in` subexpression, and the variables retain to their previous values outside the `in` subexpression.

Examples

Project Values Calculated Using Variables A sales collection has the following documents:

```
{ _id: 1, price: 10, tax: 0.50, applyDiscount: true }
{ _id: 2, price: 10, tax: 0.25, applyDiscount: false }
```

The following aggregation uses `$let` (page 471) in the the `$project` (page 438) pipeline stage to calculate and return the `finalTotal` for each document:

```
db.sales.aggregate( [
  {
    $project: {
      finalTotal: {
        $let: {
          vars: {
            total: { $add: [ '$price', '$tax' ] },
            discounted: { $cond: { if: '$applyDiscount', then: 0.9, else: 1 } }
          },
          in: { $multiply: [ "$$total", "$$discounted" ] }
        }
      }
    }
  }
] )
```

The aggregation returns the following results:

```
{ "_id" : 1, "finalTotal" : 9.450000000000001 }
{ "_id" : 2, "finalTotal" : 10.25 }
```

See also:

`$map` (page 471)

`$literal` (aggregation)

Definition

`$literal`

Wraps an expression to prevent the aggregation pipeline from evaluating the expression.

Examples

Treat \$ as a Literal In various aggregation expressions ²⁶, the dollar sign \$ evaluates to a field path; i.e. provides access to the field. For example, the \$eq expression \$eq: ["\$price", "\$1"] performs an equality check between the value in the field named price and the value in the field named 1 in the document.

The following example uses a \$literal (page 472) expression to treat a string that contains a dollar sign "\$1" as a constant value.

A collection records has the following documents:

```
{ "_id" : 1, "item" : "abc123", price: "$2.50" }
{ "_id" : 2, "item" : "xyz123", price: "1" }
{ "_id" : 3, "item" : "ijk123", price: "$1" }
```

```
db.records.aggregate( [
  { $project: { costsOneDollar: { $eq: [ "$price", { $literal: "$1" } ] } } }
] )
```

This operation projects a field named costsOneDollar that holds a boolean value, indicating whether the value of the price field is equal to the string "\$1":

```
{ "_id" : 1, "costsOneDollar" : false }
{ "_id" : 2, "costsOneDollar" : false }
{ "_id" : 3, "costsOneDollar" : true }
```

Project a New Field with Value 1 The \$project (page 438) stage uses the expression <field>: 1 to include the <field> in the output. The following example uses the \$literal (page 472) to return a new field set to the value of 1.

A collection bids has the following documents:

```
{ "_id" : 1, "item" : "abc123", condition: "new" }
{ "_id" : 2, "item" : "xyz123", condition: "new" }
```

The following aggregation evaluates the expression item: 1 to mean return the existing field item in the output, but uses the { \$literal: 1 } (page 472) expression to return a new field startAt set to the value 1:

```
db.bids.aggregate( [
  { $project: { item: 1, startAt: { $literal: 1 } } }
] )
```

The operation results in the following documents:

```
{ "_id" : 1, "item" : "abc123", "startAt" : 1 }
{ "_id" : 2, "item" : "xyz123", "startAt" : 1 }
```

Date Operators

Date operators take a “Date” typed value as a single argument and return a number.

²⁶ The \$match (page 440) expressions do not evaluate \$ as the field path.

Date Aggregation Operators	Name	Description
	<code>\$dayOfYear</code> (page 474)	Converts a date to a number between 1 and 366.
	<code>\$dayOfMonth</code> (page 474)	Converts a date to a number between 1 and 31.
	<code>\$dayOfWeek</code> (page 474)	Converts a date to a number between 1 and 7.
	<code>\$year</code> (page 474)	Converts a date to the full year.
	<code>\$month</code> (page 474)	Converts a date into a number between 1 and 12.
	<code>\$week</code> (page 474)	Converts a date into a number between 0 and 53
	<code>\$hour</code> (page 474)	Converts a date into a number between 0 and 23.
	<code>\$minute</code> (page 474)	Converts a date into a number between 0 and 59.
	<code>\$second</code> (page 475)	Converts a date into a number between 0 and 59. May be 60 to account for leap seconds.
	<code>\$millisecond</code> (page 475)	Returns the millisecond portion of a date as an integer between 0 and 999.

`$dayOfYear` (aggregation)**`$dayOfYear`**

Takes a date and returns the day of the year as a number between 1 and 366.

`$dayOfMonth` (aggregation)**`$dayOfMonth`**

Takes a date and returns the day of the month as a number between 1 and 31.

`$dayOfWeek` (aggregation)**`$dayOfWeek`**

Takes a date and returns the day of the week as a number between 1 (Sunday) and 7 (Saturday.)

`$year` (aggregation)**`$year`**

Takes a date and returns the full year.

`$month` (aggregation)**`$month`**

Takes a date and returns the month as a number between 1 and 12.

`$week` (aggregation)**`$week`**

Takes a date and returns the week of the year as a number between 0 and 53.

Weeks begin on Sundays, and week 1 begins with the first Sunday of the year. Days preceding the first Sunday of the year are in week 0. This behavior is the same as the “%U” operator to the `strftime` standard library function.

`$hour` (aggregation)**`$hour`**

Takes a date and returns the hour between 0 and 23.

`$minute` (aggregation)**`$minute`**

Takes a date and returns the minute between 0 and 59.

\$second (aggregation)**\$second**

Takes a date and returns the second between 0 and 59, but can be 60 to account for leap seconds.

\$millisecond (aggregation)**\$millisecond**

Takes a date and returns the millisecond portion of the date as an integer between 0 and 999.

Conditional Expressions

	Name	Description
Conditional Aggregation Operators	\$cond (page 475)	A ternary operator that evaluates one expression, and depending on the result, returns the value of one of the following expressions.
	\$ifNull (page 476)	Evaluates an expression and returns a value.

\$cond (aggregation)**Definition****\$cond**

[\\$cond](#) (page 475) is a ternary operator that takes three expressions and evaluates the first expression to determine which of the subsequent expressions to return. [\\$cond](#) (page 475) accepts input either as an array with three items, or as an object.

New in version 2.6: [\\$cond](#) (page 475) now accepts expressions in the form of documents.

Syntax

Document New in version 2.6: [\\$cond](#) (page 475) adds support for the document format.

When [\\$cond](#) (page 475) takes a document, the document has three fields: `if`, `then`, and `else`. Consider the following example:

```
{ $cond: { if: <boolean-expression>,
           then: <true-case>,
           else: <false-case> } }
```

The `if` field takes an expression that evaluates to a Boolean value. If the expression evaluates to `true`, then [\\$cond](#) (page 475) evaluates and returns the value of the `then` field. Otherwise, [\\$cond](#) (page 475) evaluates and returns the value of the `else` field.

The expressions in the `if`, `then`, and `else` fields may be valid MongoDB *aggregation expressions* (page 437). You cannot use JavaScript in the expressions.

Array When you specify [\\$cond](#) (page 475) as an array of three expressions, the first expression evaluates to a Boolean value. If the first expression evaluates to `true`, then [\\$cond](#) (page 475) evaluates and returns the value of the second expression. If the first expression evaluates to `false`, then [\\$cond](#) (page 475) evaluates and returns the third expression.

Use the [\\$cond](#) (page 475) operator with the following syntax:

```
{ $cond: [ <boolean-expression>, <true-case>, <false-case> ] }
```

All three values in the array specified to `$cond` (page 475) must be valid MongoDB *aggregation expressions* (page 437) or document fields. Do not use JavaScript in any aggregation statements, including `$cond` (page 475).

Examples

Specify `$cond` Expression as a Document The following aggregation pipeline operation returns a `weightedCount` for each `item_id`. The `$sum` (page 460) operator uses the `$cond` (page 475) expression to add 2 if the value stored in the `level` field is `E` and 1 otherwise.

```
db.survey.aggregate([
  {
    $group: {
      _id: "$item_id",
      weightedCount: { $sum: { $cond: { if: { $eq: [ "$level", "E" ] } ,
                                     then: 2,
                                     else: 1
                                   } } }
    }
  }
])
```

Specify a `$cond` Expression using an Array The following aggregation on the `survey` collection groups by the `item_id` field and returns a `weightedCount` for each `item_id`. The `$sum` (page 460) operator uses the `$cond` (page 475) expression to add either 2 if the value stored in the `level` field is `E` and 1 otherwise.

```
db.survey.aggregate([
  {
    $group: {
      _id: "$item_id",
      weightedCount: { $sum: { $cond: [ { $eq: [ "$level", "E" ] } , 2, 1 ] } }
    }
  }
])
```

`$ifNull` (aggregation)

`$ifNull`

Takes an array with two expressions. `$ifNull` (page 476) returns the first expression if it evaluates to a non-null value. Otherwise, `$ifNull` (page 476) returns the second expression's value.

Use the `$ifNull` (page 476) operator with the following syntax:

```
{ $ifNull: [ <expression>, <replacement-if-null> ] }
```

Both values in the array specified to `$ifNull` (page 476) must be valid MongoDB *aggregation expressions* (page 437) or document fields. Do not use JavaScript in any aggregation statements, including `$ifNull` (page 476).

Example

The following aggregation on the `offSite` collection groups by the `location` field and returns a count for each location. If the `location` field contains `null` or does not exist, the `$ifNull` (page 476) returns `"Unspecified"` as the value. MongoDB assigns the returned value to `_id` in the aggregated document.

```
db.offSite.aggregate(  
  [  
    {  
      $group: {  
        _id: { $ifNull: [ "$location", "Unspecified" ] },  
        count: { $sum: 1 }  
      }  
    }  
  ]  
)
```

2.3.4 Query Modifiers

Introduction

In addition to the *MongoDB Query Operators* (page 373), there are a number of “meta” operators that let you modify the output or behavior of a query. On the server, MongoDB treats the query and the options as a single object. The `mongo` (page 527) shell and driver interfaces may provide *cursor methods* (page 77) that wrap these options. When possible, use these methods; otherwise, you can add these options using either of the following syntax:

```
db.collection.find( { <query> } )._addSpecial( <option> )  
db.collection.find( { $query: { <query> }, <option> } )
```

Operators

Modifiers

Many of these operators have corresponding *methods in the shell* (page 77). These methods provide a straightforward and user-friendly interface and are the preferred way to add these options.

Name	Description
<code>\$comment</code> (page 478)	Adds a comment to the query to identify queries in the <i>database profiler</i> output.
<code>\$explain</code> (page 478)	Forces MongoDB to report on query execution plans. See <code>explain()</code> (page 80).
<code>\$hint</code> (page 479)	Forces MongoDB to use a specific index. See <code>hint()</code> (page 86)
<code>\$maxScan</code> (page 479)	Limits the number of documents scanned.
<code>\$maxTimeMS</code> (page 480)	Specifies a cumulative time limit in milliseconds for processing operations on a cursor. See <code>maxTimeMS()</code> (page 87).
<code>\$max</code> (page 480)	Specifies an <i>exclusive</i> upper limit for the index to use in a query. See <code>max()</code> (page 88).
<code>\$min</code> (page 481)	Specifies an <i>inclusive</i> lower limit for the index to use in a query. See <code>min()</code> (page 89).
<code>\$orderby</code> (page 481)	Returns a cursor with documents sorted according to a sort specification. See <code>sort()</code> (page 93).
<code>\$returnKey</code> (page 482)	Forces the cursor to only return fields included in the index.
<code>\$showDiskLoc</code> (page 482)	Modifies the documents returned to include references to the on-disk location of each document.
<code>\$snapshot</code> (page 482)	Forces the query to use the index on the <code>_id</code> field. See <code>snapshot()</code> (page 92).
<code>\$query</code> (page 483)	Wraps a query document.

\$comment**\$comment**

The `$comment` (page 478) makes it possible to attach a comment to a query. Because these comments propagate to the `profile` (page 334) log, adding `$comment` (page 478) modifiers can make your profile data much easier to interpret and trace. Use one of the following forms:

```
db.collection.find( { <query> } )._addSpecial( "$comment", <comment> )
db.collection.find( { $query: { <query> }, $comment: <comment> } )
```

\$explain**\$explain**

The `$explain` (page 478) operator provides information on the query plan. It returns a document that describes the process and indexes used to return the query. This may provide useful insight when attempting to optimize a query. For details on the output, see *cursor.explain()* (page 80).

You can specify the `$explain` (page 478) operator in either of the following forms:

```
db.collection.find()._addSpecial( "$explain", 1 )
db.collection.find( { $query: {}, $explain: 1 } )
```

You also can specify `$explain` (page 478) through the `explain()` (page 80) method in the `mongo` (page 527) shell:

```
db.collection.find().explain()
```

Behavior `$explain` (page 478) runs the actual query to determine the result. Although there are some differences between running the query with `$explain` (page 478) and running without, generally, the performance will be similar between the two. So, if the query is slow, the `$explain` (page 478) operation is also slow.

Additionally, the `$explain` (page 478) operation reevaluates a set of candidate query plans, which may cause the `$explain` (page 478) operation to perform differently than a normal query. As a result, these operations generally provide an accurate account of *how* MongoDB would perform the query, but do not reflect the length of these queries.

See also:

- [explain\(\)](#) (page 80)
- <http://docs.mongodb.org/manualadministration/optimization> page for information regarding optimization strategies.
- <http://docs.mongodb.org/manualtutorial/manage-the-database-profiler> tutorial for information regarding the database profile.
- [Current Operation Reporting](#) (page 103)

\$hint**\$hint**

The [\\$hint](#) (page 479) operator forces the *query optimizer* to use a specific index to fulfill the query. Specify the index either by the index name or by document.

Use [\\$hint](#) (page 479) for testing query performance and indexing strategies. The [mongo](#) (page 527) shell provides a helper method [hint\(\)](#) (page 86) for the [\\$hint](#) (page 479) operator.

Consider the following operation:

```
db.users.find().hint( { age: 1 } )
```

This operation returns all documents in the collection named `users` using the index on the `age` field.

You can also specify a hint using either of the following forms:

```
db.users.find()._addSpecial( "$hint", { age : 1 } )
db.users.find( { $query: {} }, $hint: { age : 1 } )
```

Note: When the query specifies the [\\$hint](#) (page 479) in the following form:

```
db.users.find( { $query: {} }, $hint: { age : 1 } )
```

Then, in order to include the [\\$explain](#) (page 478) option, you must add the [\\$explain](#) (page 478) option to the document, as in the following:

```
db.users.find( { $query: {} }, $hint: { age : 1 }, $explain: 1 )
```

When an *index filter* exists for the query shape, MongoDB ignores the [\\$hint](#) (page 479). The [explain.filterSet](#) (page 84) field of the [explain\(\)](#) (page 80) output indicates whether MongoDB applied an index filter for the query.

\$maxScan**\$maxScan**

Constrains the query to only scan the specified number of documents when fulfilling the query. Use one of the following forms:

```
db.collection.find( { <query> } )._addSpecial( "$maxScan" , <number> )
db.collection.find( { $query: { <query> } }, $maxScan: <number> )
```

Use this modifier to prevent potentially long running queries from disrupting performance by scanning through too much data.

\$maxTimeMS

\$maxTimeMS

New in version 2.6: The `$maxTimeMS` (page 480) operator specifies a cumulative time limit in milliseconds for processing operations on the cursor. MongoDB interrupts the operation at the earliest following *interrupt point*.

The `mongo` (page 527) shell provides the `cursor.maxTimeMS()` (page 87) method

```
db.collection.find().maxTimeMS(100)
```

You can also specify the option in either of the following forms:

```
db.collection.find( $query: { } , $maxTimeMS: 100 )
db.collection.find().addSpecial("$maxTimeMS", 100)
```

Interrupted operations return an error message similar to the following:

```
error: { "$err" : "operation exceeded time limit", "code" : 50 }
```

\$max

\$max

Specify a `$max` (page 480) value to specify the *exclusive* upper bound for a specific index in order to constrain the results of `find()` (page 34). The `mongo` (page 527) shell provides the `max()` (page 88) wrapper method:

```
db.collection.find( { <query> } ).max( { field1: <max value>, ... fieldN: <max valueN> } )
```

You can also specify the option with either of the two forms:

```
db.collection.find( { <query> } ).addSpecial( "$max", { field1: <max value1>, ... fieldN: <max valueN> } )
db.collection.find( { $query: { <query> }, $max: { field1: <max value1>, ... fieldN: <max valueN> } }
```

The `$max` (page 480) specifies the upper bound for *all* keys of a specific index *in order*.

Consider the following operations on a collection named `collection` that has an index `{ age: 1 }`:

```
db.collection.find( { <query> } ).max( { age: 100 } )
```

This operation limits the query to those documents where the field `age` is less than 100 using the index `{ age: 1 }`.

You can explicitly specify the corresponding index with `hint()` (page 86). Otherwise, MongoDB selects the index using the fields in the `indexBounds`; however, if multiple indexes exist on same fields with different sort orders, the selection of the index may be ambiguous.

Consider a collection named `collection` that has the following two indexes:

```
{ age: 1, type: -1 }
{ age: 1, type: 1 }
```

Without explicitly using `hint()` (page 86), MongoDB may select either index for the following operation:

```
db.collection.find().max( { age: 50, type: 'B' } )
```

Use `$max` (page 480) alone or in conjunction with `$min` (page 481) to limit results to a specific range for the *same* index, as in the following example:

```
db.collection.find().min( { age: 20 } ).max( { age: 25 } )
```

Note: Because `max()` (page 88) requires an index on a field, and forces the query to use this index, you may prefer the `$lt` (page 375) operator for the query if possible. Consider the following example:

```
db.collection.find( { _id: 7 } ).max( { age: 25 } )
```

The query uses the index on the age field, even if the index on `_id` may be better.

\$min

\$min

Specify a `$min` (page 481) value to specify the *inclusive* lower bound for a specific index in order to constrain the results of `find()` (page 34). The `mongo` (page 527) shell provides the `min()` (page 89) wrapper method:

```
db.collection.find( { <query> } ).min( { field1: <min value>, ... fieldN: <min valueN> } )
```

You can also specify the option with either of the two forms:

```
db.collection.find( { <query> } )._addSpecial( "$min", { field1: <min value1>, ... fieldN: <min valueN> } )
db.collection.find( { $query: { <query> }, $min: { field1: <min value1>, ... fieldN: <min valueN> } } )
```

The `$min` (page 481) specifies the lower bound for *all* keys of a specific index *in order*.

Consider the following operations on a collection named `collection` that has an index `{ age: 1 }`:

```
db.collection.find().min( { age: 20 } )
```

These operations limit the query to those documents where the field `age` is at least 20 using the index `{ age: 1 }`.

You can explicitly specify the corresponding index with `hint()` (page 86). Otherwise, MongoDB selects the index using the fields in the `indexBounds`; however, if multiple indexes exist on same fields with different sort orders, the selection of the index may be ambiguous.

Consider a collection named `collection` that has the following two indexes:

```
{ age: 1, type: -1 }
{ age: 1, type: 1 }
```

Without explicitly using `hint()` (page 86), it is unclear which index the following operation will select:

```
db.collection.find().min( { age: 20, type: 'C' } )
```

You can use `$min` (page 481) in conjunction with `$max` (page 480) to limit results to a specific range for the *same* index, as in the following example:

```
db.collection.find().min( { age: 20 } ).max( { age: 25 } )
```

Note: Because `min()` (page 89) requires an index on a field, and forces the query to use this index, you may prefer the `$gte` (page 374) operator for the query if possible. Consider the following example:

```
db.collection.find( { _id: 7 } ).min( { age: 25 } )
```

The query will use the index on the age field, even if the index on `_id` may be better.

\$orderby

\$orderby

The `$orderby` (page 481) operator sorts the results of a query in ascending or descending order.

The `mongo` (page 527) shell provides the `cursor.sort()` (page 93) method:

```
db.collection.find().sort( { age: -1 } )
```

You can also specify the option in either of the following forms:

```
db.collection.find().__addSpecial( "$orderby", { age : -1 } )
db.collection.find( { $query: {} , $orderby: { age : -1 } } )
```

These examples return all documents in the collection named `collection` sorted by the `age` field in descending order. Specify a value to `$orderby` (page 481) of negative one (e.g. `-1`, as above) to sort in descending order or a positive value (e.g. `1`) to sort in ascending order.

Behavior The sort function requires that the entire sort be able to complete within 32 megabytes. When the sort option consumes more than 32 megabytes, MongoDB will return an error.

To avoid this error, create an index to support the sort operation or use `$orderby` (page 481) in conjunction with `$maxScan` (page 479) and/or `cursor.limit()` (page 86). The `cursor.limit()` (page 86) increases the speed and reduces the amount of memory required to return this query by way of an optimized algorithm. The specified limit must result in a number of documents that fall within the 32 megabyte limit.

\$returnKey

\$returnKey

Only return the index field or fields for the results of the query. If `$returnKey` (page 482) is set to `true` and the query does not use an index to perform the read operation, the returned documents will not contain any fields. Use one of the following forms:

```
db.collection.find( { <query> } ).__addSpecial( "$returnKey", true )
db.collection.find( { $query: { <query> } , $returnKey: true } )
```

\$showDiskLoc

\$showDiskLoc

`$showDiskLoc` (page 482) option adds a field `$diskLoc` to the returned documents. The `$diskLoc` field contains the disk location information.

The `mongo` (page 527) shell provides the `cursor.showDiskLoc()` (page 91) method:

```
db.collection.find().showDiskLoc()
```

You can also specify the option in either of the following forms:

```
db.collection.find( { <query> } ).__addSpecial("$showDiskLoc" , true)
db.collection.find( { $query: { <query> } , $showDiskLoc: true } )
```

\$snapshot

\$snapshot

The `$snapshot` (page 482) operator prevents the cursor from returning a document more than once because an intervening write operation results in a move of the document.

Even in snapshot mode, objects inserted or deleted during the lifetime of the cursor may or may not be returned.

The `mongo` (page 527) shell provides the `cursor.snapshot()` (page 92) method:

```
db.collection.find().snapshot()
```

You can also specify the option in either of the following forms:

```
db.collection.find()._addSpecial( "$snapshot", true )
db.collection.find( { $query: {}, $snapshot: true } )
```

The `$snapshot` (page 482) operator traverses the index on the `_id` field ²⁷.

Warning:

- You cannot use `$snapshot` (page 482) with *sharded collections*.
- Do not use `$snapshot` (page 482) with `$hint` (page 479) or `$orderby` (page 481) (or the corresponding `cursor.hint()` (page 86) and `cursor.sort()` (page 93) methods.)

\$query

\$query

The `$query` (page 483) operator provides an interface to describe queries. Consider the following operation:

```
db.collection.find( { $query: { age : 25 } } )
```

This is equivalent to the more familiar `db.collection.find()` (page 34) method:

```
db.collection.find( { age : 25 } )
```

These operations return only those documents in the collection named `collection` where the `age` field equals 25.

Note: Do not mix query forms. If you use the `$query` (page 483) format, do not append *cursor methods* (page 77) to the `find()` (page 34). To modify the query use the *meta-query operators* (page 477), such as `$explain` (page 478).

Therefore, the following two operations are equivalent:

```
db.collection.find( { $query: { age : 25 }, $explain: true } )
db.collection.find( { age : 25 } ).explain()
```

See also:

For more information about queries in MongoDB see <http://docs.mongodb.org/manualcore/read-operations>, `db.collection.find()` (page 34), and <http://docs.mongodb.org/manualtutorial/getting-started>.

Sort Order

Name	Description
<code>\$natural</code> (page 483)	A special sort order that orders documents using the order of documents on disk.

\$natural

\$natural

Use the `$natural` (page 483) operator to use *natural order* for the results of a sort operation. Natural order refers to the order of documents in the file on disk.

The `$natural` (page 483) operator uses the following syntax to return documents in the order they exist on disk:

```
db.collection.find().sort( { $natural: 1 } )
```

²⁷ You can achieve the `$snapshot` (page 482) isolation behavior using any *unique* index on invariable fields.

Use `-1` to return documents in the reverse order as they occur on disk:

```
db.collection.find().sort( { $natural: -1 } )
```

You cannot specify `$natural` (page 483) sort order if the query includes a `$text` (page 387) expression.

See also:

`cursor.sort()` (page 93)

2.4 Aggregation Reference

***Aggregation Operator Quick Reference** (page 484)* Quick reference card for aggregation pipeline.

***Aggregation Framework Operators** (page 437)* Aggregation pipeline operations have a collection of operators available to define and manipulate documents in pipeline stages.

***Aggregation Commands Comparison** (page 488)* A comparison of `group` (page 204), `mapReduce` (page 208) and `aggregate` (page 198) that explores the strengths and limitations of each aggregation modality.

***SQL to Aggregation Mapping Chart** (page 500)* An overview common aggregation operations in SQL and MongoDB using the aggregation pipeline and operators in MongoDB and common SQL statements.

***Aggregation Interfaces** (page 492)* The data aggregation interfaces document the invocation format and output for MongoDB's aggregation commands and methods.

***Variables in Aggregation** (page 492)* Use of variables in aggregation pipeline expressions.

2.4.1 Aggregation Operator Quick Reference

Pipeline Operators

Note: The *aggregation pipeline* cannot operate on values of the following types: `Symbol`, `MinKey`, `MaxKey`, `DBRef`, `Code`, and `CodeWScope`.

Pipeline operators appear in an array. Documents pass through the operators in a sequence.

Name	Description
<code>\$project</code> (page 438)	Reshapes a document stream. <code>\$project</code> (page 438) can rename, add, or remove fields as well as create computed values and sub-documents.
<code>\$match</code> (page 440)	Filters the document stream, and only allows matching documents to pass into the next pipeline stage. <code>\$match</code> (page 440) uses standard MongoDB queries.
<code>\$redact</code> (page 441)	Restricts the content of a returned document on a per-field level.
<code>\$limit</code> (page 445)	Restricts the number of documents in an aggregation pipeline.
<code>\$skip</code> (page 445)	Skips over a specified number of documents from the pipeline and returns the rest.
<code>\$unwind</code> (page 446)	Takes an array of documents and returns them as a stream of documents.
<code>\$group</code> (page 447)	Groups documents together for the purpose of calculating aggregate values based on a collection of documents.
<code>\$sort</code> (page 449)	Takes all input documents and returns them in a stream of sorted documents.
<code>\$geoNear</code> (page 451)	Returns an ordered stream of documents based on proximity to a geospatial point.
<code>\$out</code> (page 453)	Writes documents from the pipeline to a collection. The <code>\$out</code> (page 453) operator must be the last stage in the pipeline.

Expression Operators

Expression operators calculate values within the *Pipeline Operators* (page 437).

`$group` Operators

Name	Description
<code>\$addToSet</code> (page 455)	Returns an array of all the <i>unique</i> values for the selected field among for each document in that group.
<code>\$first</code> (page 456)	Returns the first value in a group.
<code>\$last</code> (page 456)	Returns the last value in a group.
<code>\$max</code> (page 456)	Returns the highest value in a group.
<code>\$min</code> (page 456)	Returns the lowest value in a group.
<code>\$avg</code> (page 457)	Returns an average of all the values in a group.
<code>\$push</code> (page 458)	Returns an array of <i>all</i> values for the selected field among for each document in that group.
<code>\$sum</code> (page 460)	Returns the sum of all the values in a group.

Boolean Operators

These operators accept Booleans as arguments and return Booleans as results.

The operators convert non-Booleans to Boolean values according to the BSON standards. Here, `null`, `undefined`, and `0` values become `false`, while non-zero numeric values, and all other types, such as strings, dates, objects become `true`.

Name	Description
<code>\$and</code> (page 460)	Returns true only when <i>all</i> values in its input array are true.
<code>\$or</code> (page 461)	Returns true when <i>any</i> value in its input array are true.
<code>\$not</code> (page 461)	Returns the boolean value that is the opposite of the input value.

Set Operators

These operators provide operations on sets.

Name	Description
<code>\$setEquals</code> (page 461)	Returns true if two sets have the same elements.
<code>\$setIntersection</code> (page 461)	Returns the common elements of the input sets.
<code>\$setDifference</code> (page 462)	Returns elements of a set that do not appear in a second set.
<code>\$setUnion</code> (page 462)	Returns a set that holds all elements of the input sets.
<code>\$setIsSubset</code> (page 462)	Returns true if all elements of a set appear in a second set.
<code>\$anyElementTrue</code> (page 462)	Returns true if <i>any</i> elements of a set evaluate to true, and false otherwise.
<code>\$allElementsTrue</code> (page 462)	Returns true if <i>all</i> elements of a set evaluate to true, and false otherwise.

Comparison Operators

These operators perform comparisons between two values and return a Boolean, in most cases reflecting the result of the comparison.

All comparison operators take an array with a pair of values. You may compare numbers, strings, and dates. Except for `$cmp` (page 463), all comparison operators return a Boolean value. `$cmp` (page 463) returns an integer.

Name	Description
<code>\$cmp</code> (page 463)	Compares two values and returns the result of the comparison as an integer.
<code>\$eq</code> (page 463)	Takes two values and returns true if the values are equivalent.
<code>\$gt</code> (page 463)	Takes two values and returns true if the first is larger than the second.
<code>\$gte</code> (page 463)	Takes two values and returns true if the first is larger than or equal to the second.
<code>\$lt</code> (page 463)	Takes two values and returns true if the second value is larger than the first.
<code>\$lte</code> (page 463)	Takes two values and returns true if the second value is larger than or equal to the first.
<code>\$ne</code> (page 464)	Takes two values and returns true if the values are <i>not</i> equivalent.

Arithmetic Operators

Arithmetic operators support only numbers.

Name	Description
<code>\$add</code> (page 464)	Computes the sum of an array of numbers.
<code>\$divide</code> (page 464)	Takes two numbers and divides the first number by the second.
<code>\$mod</code> (page 464)	Takes two numbers and calculates the modulo of the first number divided by the second.
<code>\$multiply</code> (page 464)	Computes the product of an array of numbers.
<code>\$subtract</code> (page 464)	Takes an array that contains two numbers or two dates and subtracts the second value from the first.

String Operators

String operators that manipulate strings.

Name	Description
<code>\$concat</code> (page 465)	Concatenates two strings.
<code>\$strcasecmp</code> (page 467)	Compares two strings and returns an integer that reflects the comparison.
<code>\$substr</code> (page 468)	Takes a string and returns portion of that string.
<code>\$toLowerCase</code> (page 468)	Converts a string to lowercase.
<code>\$toUpperCase</code> (page 468)	Converts a string to uppercase.

Text Search Operators

Operators to support text search.

Name	Description
<code>\$meta</code> (page 468)	Access metadata for <code>\$sort</code> (page 449) stage or <code>\$project</code> (page 438) stage.

Array Operators

Operators that manipulate arrays.

Name	Description
<code>\$size</code> (page 470)	Returns the size of the array.

Projection Expressions

Operators that increase the flexibility within aggregation projection and projection-like expressions. These operators are available in the `$project` (page 438), `$group` (page 447), and `$redact` (page 441) pipeline stages.

Name	Description
<code>\$map</code> (page 471)	Applies a sub-expression to each item in an array and returns the result of the sub-expression.
<code>\$let</code> (page 471)	Defines variables for use within the scope of an aggregation expression.
<code>\$literal</code> (page 472)	Forces the aggregation pipeline to return a literal value without evaluating the expression.

Date Operators

Date operators take a “Date” typed value as a single argument and return a number.

Name	Description
<code>\$dayOfYear</code> (page 474)	Converts a date to a number between 1 and 366.
<code>\$dayOfMonth</code> (page 474)	Converts a date to a number between 1 and 31.
<code>\$dayOfWeek</code> (page 474)	Converts a date to a number between 1 and 7.
<code>\$year</code> (page 474)	Converts a date to the full year.
<code>\$month</code> (page 474)	Converts a date into a number between 1 and 12.
<code>\$week</code> (page 474)	Converts a date into a number between 0 and 53.
<code>\$hour</code> (page 474)	Converts a date into a number between 0 and 23.
<code>\$minute</code> (page 474)	Converts a date into a number between 0 and 59.
<code>\$second</code> (page 475)	Converts a date into a number between 0 and 59. May be 60 to account for leap seconds.
<code>\$millisecond</code> (page 475)	Returns the millisecond portion of a date as an integer between 0 and 999.

Conditional Expressions

Name	Description
<code>\$cond</code> (page 475)	A ternary operator that evaluates one expression, and depending on the result returns the value of one following expressions.
<code>\$ifNull</code> (page 476)	

2.4.2 Aggregation Commands Comparison

The following table provides a brief overview of the features of the MongoDB aggregation commands.

	aggregate (page 198)	mapReduce (page 208)	group (page 204)
Description	New in version 2.2. Designed with specific goals of improving performance and usability for aggregation tasks. Uses a “pipeline” approach where objects are transformed as they pass through a series of pipeline operators such as \$group (page 447), \$match (page 440), and \$sort (page 449). See Aggregation Framework Operators (page 437) for more information on the pipeline operators.	Implements the Map-Reduce aggregation for processing large data sets.	Provides grouping functionality. Is slower than the aggregate (page 198) command and has less functionality than the mapReduce (page 208) command.
Key Features	Pipeline operators can be repeated as needed. Pipeline operators need not produce one output document for every input document. Can also generate new documents or filter out documents.	In addition to grouping operations, can perform complex aggregation tasks as well as perform incremental aggregation on continuously growing datasets. See http://docs.mongodb.org/manual/tutorial/map-reduce-examples/ and http://docs.mongodb.org/manual/tutorial/perform-incremental-map-reduce/	Can either group by existing fields or with a custom <code>keyf</code> JavaScript function, can group by calculated fields. See group (page 204) for information and example using the <code>keyf</code> function.
Flexibility	Limited to the operators and expressions supported by the aggregation pipeline. However, can add computed fields, create new virtual sub-objects, and extract sub-fields into the top-level of results by using the \$project (page 438) pipeline operator. See \$project (page 438) for more information as well as Aggregation Framework Operators (page 437) for more information on all the available pipeline operators.	Custom <code>map</code> , <code>reduce</code> and <code>finalize</code> JavaScript functions offer flexibility to aggregation logic. See mapReduce (page 208) for details and restrictions on the functions.	Custom <code>reduce</code> and <code>finalize</code> JavaScript functions offer flexibility to grouping logic. See group (page 204) for details and restrictions on these functions.
Output Results	Returns results in various options (inline as a document that contains the result set, a cursor to the result set) or stores the results in a collection. The result is subject to the BSON Document size (page 604) limit if returned inline as a document that contains the result set. Changed in version 2.6: Can return results as a cursor or store the results to a collection.	Returns results in various options (inline, new collection, merge, replace, reduce). See mapReduce (page 208) for details on the output options. Changed in version 2.2: Provides much better support for sharded map-reduce output than previous versions.	Returns results inline as an array of grouped items. The result set must fit within the maximum BSON document size limit (page 604). Changed in version 2.2: The returned array can contain at most 20,000 elements; i.e. at most 20,000 unique groupings. Previous versions had a limit of 10,000 elements.
Sharding	Supports non-sharded and sharded input collections.	Supports non-sharded and sharded input collections.	Does not support sharded collection.
Notes		Prior to 2.4, JavaScript code executed in a single thread.	Prior to 2.4, JavaScript code executed in a single thread.
2.4. Aggregation Reference			
More Information	See http://docs.mongodb.org/manual/core/aggregate and aggregate (page 198).	See http://docs.mongodb.org/manual/core/map-reduce and mapReduce (page 208).	See group (page 204).

2.4.3 SQL to Aggregation Mapping Chart

The aggregation pipeline allows MongoDB to provide native aggregation capabilities that corresponds to many common data aggregation operations in SQL. If you're new to MongoDB you might want to consider the <http://docs.mongodb.org/manual/faq> section for a selection of common questions.

The following table provides an overview of common SQL aggregation terms, functions, and concepts and the corresponding MongoDB *aggregation operators* (page 437):

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	<code>\$match</code> (page 440)
GROUP BY	<code>\$group</code> (page 447)
HAVING	<code>\$match</code> (page 440)
SELECT	<code>\$project</code> (page 438)
ORDER BY	<code>\$sort</code> (page 449)
LIMIT	<code>\$limit</code> (page 445)
SUM()	<code>\$sum</code> (page 460)
COUNT()	<code>\$sum</code> (page 460)
join	No direct corresponding operator; <i>however</i> , the <code>\$unwind</code> (page 446) operator allows for somewhat similar functionality, but with fields embedded within the document.

Examples

The following table presents a quick reference of SQL aggregation statements and the corresponding MongoDB statements. The examples in the table assume the following conditions:

- The SQL examples assume *two* tables, `orders` and `order_lineitem` that join by the `order_lineitem.order_id` and the `orders.id` columns.
- The MongoDB examples assume *one* collection `orders` that contain documents of the following prototype:

```
{
  cust_id: "abc123",
  ord_date: ISODate("2012-11-02T17:04:11.102Z"),
  status: 'A',
  price: 50,
  items: [ { sku: "xxx", qty: 25, price: 1 },
            { sku: "yyy", qty: 25, price: 1 } ]
}
```

- The MongoDB statements prefix the names of the fields from the *documents* in the collection `orders` with a `$` character when they appear as operands to the aggregation operations.

SQL Example	MongoDB Example	Description
<pre>SELECT COUNT(*) AS count FROM orders</pre>	<pre>db.orders.aggregate([{ \$group: { _id: null, count: { \$sum: 1 } } }])</pre>	Count all records from orders
<pre>SELECT SUM(price) AS total FROM orders</pre>	<pre>db.orders.aggregate([{ \$group: { _id: null, total: { \$sum: "\$price" } } }])</pre>	Sum the price field from orders
<pre>SELECT cust_id, SUM(price) AS total FROM orders GROUP BY cust_id</pre>	<pre>db.orders.aggregate([{ \$group: { _id: "\$cust_id", total: { \$sum: "\$price" } } }])</pre>	For each unique cust_id, sum the price field.
<pre>SELECT cust_id, SUM(price) AS total FROM orders GROUP BY cust_id ORDER BY total</pre>	<pre>db.orders.aggregate([{ \$group: { _id: "\$cust_id", total: { \$sum: "\$price" } } }, { \$sort: { total: 1 } }])</pre>	For each unique cust_id, sum the price field, results sorted by sum.
<pre>SELECT cust_id, ord_date, SUM(price) AS total FROM orders GROUP BY cust_id, ord_date</pre>	<pre>db.orders.aggregate([{ \$group: { _id: { cust_id: "\$cust_id", ord_date: "\$ord_date" }, total: { \$sum: "\$price" } } }])</pre>	For each unique cust_id, ord_date grouping, sum the price field.
<pre>SELECT cust_id, count(*) FROM orders GROUP BY cust_id HAVING count(*) > 1</pre>	<pre>db.orders.aggregate([{ \$group: { _id: "\$cust_id", count: { \$sum: 1 } } }, { \$filter: { _id: "\$cust_id", count: { \$gt: 1 } } }])</pre>	For cust_id with multiple records, return the cust_id and the corresponding record count.

2.4.4 Aggregation Interfaces

Aggregation Commands

Name	Description
<code>aggregate</code> (page 198)	Performs aggregation tasks such as group using the aggregation framework.
<code>count</code> (page 201)	Counts the number of documents in a collection.
<code>distinct</code> (page 203)	Displays the distinct values found for a specified key in a collection.
<code>group</code> (page 204)	Groups documents in a collection by the specified key and performs simple aggregation.
<code>mapReduce</code> (page 208)	Performs map-reduce aggregation for large data sets.

Aggregation Methods

Name	Description
<code>db.collection.aggregate()</code> (page 22)	Provides access to the aggregation pipeline.
<code>db.collection.group()</code> (page 47)	Groups documents in a collection by the specified key and performs simple aggregation.
<code>db.collection.mapReduce()</code> (page 55)	Performs map-reduce aggregation for large data sets.

2.4.5 Variables in Aggregation

Aggregation expressions can use both user-defined and system variables.

Variables can hold any BSON type data. To access the value of the variable, use a string with the variable name prefixed with double dollar signs (\$\$).

If the variable references an object, to access a specific field in the object, use the dot notation; i.e. "\$\$<variable>.<field>".

User Variables

User variable names can contain the ascii characters `[_a-zA-Z0-9]` and any non-ascii character.

User variable names must begin with a lowercase ascii letter `[a-z]` or a non-ascii character.

System Variables

MongoDB offers the following system variables:

Variable	Description
ROOT	References the root document, i.e. the top-level document, currently being processed in the aggregation pipeline stage.
CURRENT	References the start of the field path being processed in the aggregation pipeline stage. Unless documented otherwise, all stages start with CURRENT (page 493) the same as ROOT (page 493). CURRENT (page 493) is modifiable. However, since <code>\$<field></code> is equivalent to <code>\$\$CURRENT.<field></code> , rebinding CURRENT (page 493) changes the meaning of <code>\$</code> accesses.
DESCEND	One of the allowed results of a <code>\$redact</code> (page 441) expression.
PRUNE	One of the allowed results of a <code>\$redact</code> (page 441) expression.
KEEP	One of the allowed results of a <code>\$redact</code> (page 441) expression.

See also:

`$let` (page 471), `$redact` (page 441)

MongoDB and SQL Interface Comparisons

3.1 SQL to MongoDB Mapping Chart

In addition to the charts that follow, you might want to consider the <http://docs.mongodb.org/manualfaq> section for a selection of common questions about MongoDB.

3.1.1 Terminology and Concepts

The following table presents the various SQL terminology and concepts and the corresponding MongoDB terminology and concepts.

SQL Terms/Concepts	MongoDB Terms/Concepts
database	<i>database</i>
table	<i>collection</i>
row	<i>document</i> or <i>BSON</i> document
column	<i>field</i>
index	<i>index</i>
table joins	embedded documents and linking
primary key	<i>primary key</i>
Specify any unique column or column combination as primary key.	In MongoDB, the primary key is automatically set to the <i>_id</i> field.
aggregation (e.g. group by)	aggregation pipeline See the <i>SQL to Aggregation Mapping Chart</i> (page 500).

3.1.2 Executables

The following table presents the MySQL/Oracle executables and the corresponding MongoDB executables.

	MySQL/Oracle	MongoDB
Database Server	mysqld/oracle	<i>mongod</i> (page 503)
Database Client	mysql/sqlplus	<i>mongo</i> (page 527)

3.1.3 Examples

The following table presents the various SQL statements and the corresponding MongoDB statements. The examples in the table assume the following conditions:

- The SQL examples assume a table named `users`.
- The MongoDB examples assume a collection named `users` that contain documents of the following prototype:

```
{
  _id: ObjectId("509a8fb2f3f4948bd2f983a0"),
  user_id: "abc123",
  age: 55,
  status: 'A'
}
```

Create and Alter

The following table presents the various SQL statements related to table-level actions and the corresponding MongoDB statements.

SQL Schema Statements	MongoDB Schema Statements
<pre> CREATE TABLE users (id MEDIUMINT NOT NULL AUTO_INCREMENT, user_id Varchar(30), age Number, status char(1), PRIMARY KEY (id)) </pre>	<p>Implicitly created on first <code>insert()</code> (page 52) operation. The primary key <code>_id</code> is automatically added if <code>_id</code> field is not specified.</p> <pre> db.users.insert({ user_id: "abc123", age: 55, status: "A" }) </pre> <p>However, you can also explicitly create a collection:</p> <pre> db.createCollection("users") </pre>
<pre> ALTER TABLE users ADD join_date DATETIME </pre>	<p>Collections do not describe or enforce the structure of its documents; i.e. there is no structural alteration at the collection level.</p> <p>However, at the document level, <code>update()</code> (page 69) operations can add fields to existing documents using the <code>\$set</code> (page 416) operator.</p> <pre> db.users.update({ }, { \$set: { join_date: new Date() } }, { multi: true }) </pre>
<pre> ALTER TABLE users DROP COLUMN join_date </pre>	<p>Collections do not describe or enforce the structure of its documents; i.e. there is no structural alteration at the collection level.</p> <p>However, at the document level, <code>update()</code> (page 69) operations can remove fields from documents using the <code>\$unset</code> (page 417) operator.</p> <pre> db.users.update({ }, { \$unset: { join_date: "" } }, { multi: true }) </pre>
<pre> CREATE INDEX idx_user_id_asc ON users(user_id) </pre>	<pre> db.users.ensureIndex({ user_id: 1 }) </pre>
<pre> CREATE INDEX idx_user_id_asc_age_desc ON users(user_id, age DESC) </pre>	<pre> db.users.ensureIndex({ user_id: 1, age: -1 }) </pre>
<pre> DROP TABLE users </pre>	<pre> db.users.drop() </pre>

For more information, see `db.collection.insert()` (page 52), `db.createCollection()` (page 102), `db.collection.update()` (page 69), `$set` (page 416), `$unset` (page 417), `db.collection.ensureIndex()` (page 30), indexes, `db.collection.drop()` (page 29), and <http://docs.mongodb.org/manualcore/data-models>.

Insert

The following table presents the various SQL statements related to inserting records into tables and the corresponding MongoDB statements.

SQL INSERT Statements	MongoDB insert() Statements
<pre>INSERT INTO users (user_id, age, status) VALUES ("bcd001", 45, "A")</pre>	<pre>db.users.insert ({ user_id: "bcd001", age: 45, status: "A" })</pre>

For more information, see `db.collection.insert()` (page 52).

Select

The following table presents the various SQL statements related to reading records from tables and the corresponding MongoDB statements.

SQL SELECT Statements	MongoDB find() Statements
<pre> SELECT * FROM users SELECT id, user_id, status FROM users SELECT user_id, status FROM users SELECT * FROM users WHERE status = "A" SELECT user_id, status FROM users WHERE status = "A" SELECT * FROM users WHERE status != "A" SELECT * FROM users WHERE status = "A" AND age = 50 SELECT * FROM users WHERE status = "A" OR age = 50 SELECT * FROM users WHERE age > 25 SELECT * FROM users WHERE age < 25 SELECT * FROM users WHERE age > 25 AND age <= 50 </pre>	<pre> db.users.find() db.users.find({ }, { user_id: 1, status: 1 }) db.users.find({ }, { user_id: 1, status: 1, _id: 0 }) db.users.find({ status: "A" }) db.users.find({ status: "A" }, { user_id: 1, status: 1, _id: 0 }) db.users.find({ status: { \$ne: "A" } }) db.users.find({ status: "A", age: 50 }) db.users.find({ \$or: [{ status: "A" } , { age: 50 }] }) db.users.find({ age: { \$gt: 25 } }) db.users.find({ age: { \$lt: 25 } }) db.users.find({ age: { \$gt: 25, \$lte: 50 } }) </pre>
3.1 SQL to MongoDB Mapping Chart <pre> SELECT * FROM users WHERE user_id like "%bc%" </pre>	<pre> db.users.find({ user_id: /bc/ }) </pre>

For more information, see `db.collection.find()` (page 34), `db.collection.distinct()` (page 29), `db.collection.findOne()` (page 43), `$ne` (page 376), `$and` (page 378), `$or` (page 377), `$gt` (page 373), `$lt` (page 375), `$exists` (page 381), `$lte` (page 375), `$regex` (page 386), `limit()` (page 86), `skip()` (page 92), `explain()` (page 80), `sort()` (page 93), and `count()` (page 79).

Update Records

The following table presents the various SQL statements related to updating existing records in tables and the corresponding MongoDB statements.

SQL Update Statements	MongoDB update() Statements
<pre>UPDATE users SET status = "C" WHERE age > 25</pre>	<pre>db.users.update({ age: { \$gt: 25 } }, { \$set: { status: "C" } }, { multi: true })</pre>
<pre>UPDATE users SET age = age + 3 WHERE status = "A"</pre>	<pre>db.users.update({ status: "A" }, { \$inc: { age: 3 } }, { multi: true })</pre>

For more information, see `db.collection.update()` (page 69), `$set` (page 416), `$inc` (page 412), and `$gt` (page 373).

Delete Records

The following table presents the various SQL statements related to deleting records from tables and the corresponding MongoDB statements.

SQL Delete Statements	MongoDB remove() Statements
<pre>DELETE FROM users WHERE status = "D"</pre>	<pre>db.users.remove({ status: "D" })</pre>
<pre>DELETE FROM users</pre>	<pre>db.users.remove({})</pre>

For more information, see `db.collection.remove()` (page 62).

3.2 SQL to Aggregation Mapping Chart

The aggregation pipeline allows MongoDB to provide native aggregation capabilities that corresponds to many common data aggregation operations in SQL. If you're new to MongoDB you might want to consider the <http://docs.mongodb.org/manualfaq> section for a selection of common questions.

The following table provides an overview of common SQL aggregation terms, functions, and concepts and the corresponding MongoDB *aggregation operators* (page 437):

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	<code>\$match</code> (page 440)
GROUP BY	<code>\$group</code> (page 447)
HAVING	<code>\$match</code> (page 440)
SELECT	<code>\$project</code> (page 438)
ORDER BY	<code>\$sort</code> (page 449)
LIMIT	<code>\$limit</code> (page 445)
SUM()	<code>\$sum</code> (page 460)
COUNT()	<code>\$sum</code> (page 460)
join	No direct corresponding operator; <i>however</i> , the <code>\$unwind</code> (page 446) operator allows for somewhat similar functionality, but with fields embedded within the document.

3.2.1 Examples

The following table presents a quick reference of SQL aggregation statements and the corresponding MongoDB statements. The examples in the table assume the following conditions:

- The SQL examples assume *two* tables, `orders` and `order_lineitem` that join by the `order_lineitem.order_id` and the `orders.id` columns.
- The MongoDB examples assume *one* collection `orders` that contain documents of the following prototype:

```
{
  cust_id: "abc123",
  ord_date: ISODate("2012-11-02T17:04:11.102Z"),
  status: 'A',
  price: 50,
  items: [ { sku: "xxx", qty: 25, price: 1 },
            { sku: "yyy", qty: 25, price: 1 } ]
}
```

- The MongoDB statements prefix the names of the fields from the *documents* in the collection `orders` with a `$` character when they appear as operands to the aggregation operations.

SQL Example	MongoDB Example	Description
SELECT COUNT(*) AS count FROM orders	db.orders.aggregate([{ \$group: { _id: null, count: { \$sum: 1 } } }])	Count all records from orders
SELECT SUM(price) AS total FROM orders	db.orders.aggregate([{ \$group: { _id: null, total: { \$sum: "\$price" } } }])	Sum the price field from orders
SELECT cust_id, SUM(price) AS total FROM orders GROUP BY cust_id	db.orders.aggregate([{ \$group: { _id: "\$cust_id", total: { \$sum: "\$price" } } }])	For each unique cust_id, sum the price field.
SELECT cust_id, SUM(price) AS total FROM orders GROUP BY cust_id ORDER BY total	db.orders.aggregate([{ \$group: { _id: "\$cust_id", total: { \$sum: "\$price" } } }, { \$sort: { total: 1 } }])	For each unique cust_id, sum the price field, results sorted by sum.
SELECT cust_id, ord_date, SUM(price) AS total FROM orders GROUP BY cust_id, ord_date	db.orders.aggregate([{ \$group: { _id: { cust_id: "\$cust_id", ord_date: "\$ord_date" }, total: { \$sum: "\$price" } } }])	For each unique cust_id, ord_date grouping, sum the price field.
SELECT cust_id, count (*) FROM orders GROUP BY cust_id HAVING count (*) > 1	db.orders.aggregate([{ \$group: { _id: "\$cust_id", count: { \$sum: 1 } } }, { \$sort: { count: 1 } }])	For cust_id with multiple records, return the cust_id and the corresponding record count.

Program and Tool Reference Pages

4.1 MongoDB Package Components

4.1.1 Core Processes

The core components in the MongoDB package are: `mongod` (page 503), the core database process; `mongos` (page 518) the controller and query router for *sharded clusters*; and `mongo` (page 527) the interactive MongoDB Shell.

`mongod`

Synopsis

`mongod` (page 503) is the primary daemon process for the MongoDB system. It handles data requests, manages data format, and performs background management operations.

This document provides a complete overview of all command line options for `mongod` (page 503). These options are primarily useful for testing purposes. In common operation, use the `configuration file options` to control the behavior of your database, which is fully capable of all operations described below.

Options

`mongod`

Core Options

`mongod`

command line option!—help, -h

`--help, -h`

Returns information on the options and use of `mongod` (page 503).

command line option!—version

`--version`

Returns the `mongod` (page 503) release number.

command line option!—config <filename>, -f

--config <filename>, **-f**

Specifies a configuration file for runtime configuration options. The configuration file is the preferred method for runtime configuration of `mongod` (page 503). The options are equivalent to the command-line configuration options. See <http://docs.mongodb.org/manualreference/configuration-options> for more information.

Ensure the configuration file uses ASCII encoding. The `mongod` (page 503) instance does not support configuration files with non-ASCII encoding, including UTF-8.

command line option!–verbose, -v

--verbose, **-v**

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!–quiet

--quiet

Runs the `mongod` (page 503) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!–port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!–bind_ip <ip address>

--bind_ip <ip address>

Default: All interfaces. .. *versionchanged:: 2.6.0* The `deb` and `rpm` packages include a default configuration file that sets `{role}` to `127.0.0.1`.

Specifies the IP address that `mongod` (page 503) binds to in order to listen for connections from applications. You may attach `mongod` (page 503) to any interface. When attaching `mongod` (page 503) to a publicly accessible interface, ensure that you have implemented proper authentication and firewall restrictions to protect the integrity of your database.

command line option!–maxConns <number>

--maxConns <number>

The maximum number of simultaneous connections that `mongod` (page 503) will accept. This setting has no effect if it is higher than your operating system's configured maximum connection tracking threshold.

Changed in version 2.6: MongoDB removed the upward limit on the `maxIncomingConnections` setting.

command line option!–syslog

--syslog

Sends all logging output to the host's *syslog* system rather than to standard output or to a log file. , as with `--logpath` (page 520).

The `--syslog` (page 520) option is not supported on Windows.

command line option!–syslogFacility <string>

--syslogFacility <string>

Default: user

Specifies the facility level used when logging messages to syslog. The value you specify must be supported by your operating system's implementation of syslog. To use this option, you must enable the `--syslog` (page 520) option.

command line option!—logpath <path>

--logpath <path>

Sends all diagnostic logging information to a log file instead of to standard output or to the host's *syslog* system. MongoDB creates the log file at the path you specify.

By default, MongoDB overwrites the log file when the process restarts. To instead append to the log file, set the `--logappend` (page 520) option.

command line option!—logappend

--logappend

Appends new entries to the end of the log file rather than overwriting the content of the log when the `mongod` (page 503) instance restarts.

command line option!—timeStampFormat <string>

--timeStampFormat <string>

Default: iso8601-local

The time format for timestamps in log messages. Specify one of the following values:

Value	Description
ctime	Displays timestamps as Wed Dec 31 18:17:54.811.
iso8601-utc	Displays timestamps in Coordinated Universal Time (UTC) in the ISO-8601 format. For example, for New York at the start of the Epoch: 1970-01-01T00:00:00.000Z
iso8601-local	Displays timestamps in local time in the ISO-8601 format. For example, for New York at the start of the Epoch: 1969-12-31T19:00:00.000+0500

command line option!—diaglog <value>

--diaglog <value>

Default: 0

Deprecated since version 2.6.

`--diaglog` (page 505) is for internal use and not intended for most users.

Creates a very verbose *diagnostic log* for troubleshooting and recording various errors. MongoDB writes these log files in the `dbPath` directory in a series of files that begin with the string `diaglog` and end with the initiation time of the logging as a hex string.

The specified value configures the level of verbosity:

Value	Setting
0	Off. No logging.
1	Log write operations.
2	Log read operations.
3	Log both read and write operations.
7	Log write and some read operations.

You can use the `mongosniff` (page 581) tool to replay this output for investigation. Given a typical `diaglog` file located at `/data/db/diaglog.4f76a58c`, you might use a command in the following form to read these files:

```
mongosniff --source DIAGLOG /data/db/diaglog.4f76a58c
```

.. warning::

Setting the diagnostic level to ``0`` will cause :program:`mongod` to stop writing data to the :term:`diagnostic log` file. However, the :program:`mongod` instance will **continue** to keep the file open, even **if** it is no longer writing data to the file. If you want to rename, move, or delete the diagnostic log you must cleanly shut down the :program:`mongod` instance before doing so.

command line option!-traceExceptions

--traceExceptions

For internal diagnostic use only.

command line option!-pidfilepath <path>

--pidfilepath <path>

Specifies a file location to hold the process ID of the [mongod](#) (page 503) process. This is useful for tracking the [mongod](#) (page 503) process in combination with the [--fork](#) (page 521) option. Without a specified [--pidfilepath](#) (page 520) option, the process creates no PID file.

command line option!-keyFile <file>

--keyFile <file>

Specifies the path to a key file to that stores the shared secret that MongoDB processes use to authenticate to each other in a [sharded cluster](#) or [replica set](#). [--keyFile](#) (page 521) implies [--auth](#) (page 507). See [inter-process-auth](#) for more information.

command line option!-setParameter <options>

--setParameter <options>

Specifies one of the MongoDB parameters described in <http://docs.mongodb.org/manualreference/parameters>. You can specify multiple `setParameter` fields.

command line option!-httpinterface

--httpinterface

New in version 2.6.

Enables the HTTP interface. Enabling the interface can increase network exposure.

Leave the HTTP interface *disabled* for production deployments. If you *do* enable this interface, you should only allow trusted clients to access this port. See [security-firewalls](#).

Note: In MongoDB Enterprise, the HTTP Console does not support Kerberos Authentication.

command line option!-nohttpinterface

--nohttpinterface

Deprecated since version 2.6: MongoDB disables the HTTP interface by default.

Disables the HTTP interface.

Do not use in conjunction with [--rest](#) (page 507) or [--jsonp](#) (page 527).

Note: In MongoDB Enterprise, the HTTP Console does not support Kerberos Authentication.

command line option!-nounixsocket

--noinxsocket

Disables listening on the UNIX domain socket. The `mongod` (page 503) process always listens on the UNIX socket unless one of the following is true:

- `--noinxsocket` (page 521) is set
- `bindIp` is not set
- `bindIp` does not specify `127.0.0.1`

New in version 2.6: `mongod` (page 503) installed from official `.deb` and `.rpm` packages have the `bind_ip` configuration set to `127.0.0.1` by default.

command line option!—unixSocketPrefix <path>

--unixSocketPrefix <path>

Default: /tmp

The path for the UNIX socket. If this option has no value, the `mongod` (page 503) process creates a socket with `http://docs.mongodb.org/manual/tmp` as a prefix. MongoDB creates and listens on a UNIX socket unless one of the following is true:

- `--noinxsocket` (page 521) is set
- `bindIp` is not set
- `bindIp` does not specify `127.0.0.1`

command line option!—fork

--fork

Enables a *daemon* mode that runs the `mongod` (page 503) process in the background. By default `mongod` (page 503) does not run as a daemon: typically you will run `mongod` (page 503) as a daemon, either by using `--fork` (page 521) or by using a controlling process that handles the daemonization process (e.g. as with `upstart` and `systemd`).

command line option!—auth

--auth

Requires database authentication for users connecting from remote hosts.

Configure users via the *mongo shell* (page 527). If no users exist, the localhost interface will continue to have access to the database until you create the first user.

See *Security and Authentication* for more information.

command line option!—noauth

--noauth

Disables authentication. Currently the default. Exists for future compatibility and clarity.

command line option!—ipv6

--ipv6

Enables IPv6 support and allows the `mongod` (page 503) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!—jsonp

--jsonp

Permits *JSONP* access via an HTTP interface. Enabling the interface can increase network exposure. The `--jsonp` (page 527) option enables the HTTP interface, even if the `HTTP interface` option is disabled.

command line option!—rest

--rest

Enables the simple *REST* API. Enabling the *REST* API enables the HTTP interface, even if the HTTP interface option is disabled, and as a result can increase network exposure.

command line option!-slowms <value>

--slowms <value>

Default: 100

The threshold in milliseconds at which the database profiler considers a query slow. MongoDB records all slow queries to the log, even when the database profiler is off. When the profiler is on, it writes to the `system.profile` collection. See the `profile` (page 334) command for more information on the database profiler.

command line option!-profile <level>

--profile <level>

Default: 0

Changes the level of database profiling, which inserts information about operation performance into standard output or a log file. Specify one of the following levels:

Level	Setting
0	Off. No profiling.
1	On. Only includes slow operations.
2	On. Includes all operations.

Database profiling can impact database performance. Enable this option only after careful consideration.

command line option!-cpu

--cpu

Forces the `mongod` (page 503) process to report the percentage of CPU time in write lock. The process generates output every four seconds and writes the data to standard output or, if you are using the `systemLog.path` option, to the log file.

command line option!-sysinfo

--sysinfo

Returns diagnostic system information and then exits. The information provides the page size, the number of physical pages, and the number of available physical pages.

command line option!-dbpath <path>

--dbpath <path>

Default: /data/db on Linux and OS X, \data\db on Windows

The directory where the `mongod` (page 503) instance stores its data. If you installed MongoDB using a package management system, check the `/etc/mongodb.conf` file provided by your packages to see which directory is specified.

command line option!-directoryperdb

--directoryperdb

Stores each database's files in its own folder in the *data directory*. When applied to an existing system, the `directoryPerDB` option alters the storage pattern of the data directory.

Use this option in conjunction with your file system and device configuration so that MongoDB will store data on a number of distinct disk devices to increase write throughput or disk capacity.

Warning: To enable this option for an **existing** system, migrate the database-specific data files to the new directory structure before enabling `directoryPerDB`. Database-specific data files begin with the name of an existing database and end with either “ns” or a number. For example, the following data directory includes files for the `local` and `test` databases:

```
journal
mongod.lock
local.0
local.1
local.ns
test.0
test.1
test.ns
```

After migration, the data directory would have the following structure:

```
journal
mongod.lock
local/local.0
local/local.1
local/local.ns
test/test.0
test/test.1
test/test.ns
```

command line option!—noIndexBuildRetry

--noIndexBuildRetry

Stops the `mongod` (page 503) from rebuilding incomplete indexes on the next start up. This applies in cases where the `mongod` (page 503) restarts after it has shut down or stopped in the middle of an index build. In such cases, the `mongod` (page 503) always removes any incomplete indexes, and then also, by default, attempts to rebuild them. To stop the `mongod` (page 503) from rebuilding incomplete indexes on start up, include this option on the command-line.

command line option!—noprealloc

--noprealloc

Disables the preallocation of data files. This shortens the start up time in some cases and can cause significant performance penalties during normal operations.

command line option!—nssize <value>

--nssize <value>

Default: 16

Specifies the default size for namespace files, which are files that end in `.ns`. Each collection and index counts as a namespace.

Use this setting to control size for newly created namespace files. This option has no impact on existing files. The maximum size for a namespace file is 2047 megabytes. The default value of 16 megabytes provides for approximately 24,000 namespaces.

command line option!—quota

--quota

Enables a maximum limit for the number data files each database can have. When running with the `--quota` (page 509) option, MongoDB has a maximum of 8 data files per database. Adjust the quota with `--quotaFiles` (page 509).

command line option!—quotaFiles <number>

--quotaFiles <number>

Default: 8

Modifies the limit on the number of data files per database. `--quotaFiles` (page 509) option requires that you set `--quota` (page 509).

command line option!–smallfiles

--smallfiles

Sets MongoDB to use a smaller default file size. The `--smallfiles` (page 510) option reduces the initial size for data files and limits the maximum size to 512 megabytes. `--smallfiles` (page 510) also reduces the size of each *journal* file from 1 gigabyte to 128 megabytes. Use `--smallfiles` (page 510) if you have a large number of databases that each holds a small quantity of data.

The `--smallfiles` (page 510) option can lead the `mongod` (page 503) instance to create a large number of files, which can affect performance for larger databases.

command line option!–syncdelay <value>

--syncdelay <value>

Default: 60

Controls how much time can pass before MongoDB flushes data to the data files via an *fsync* operation. **Do not set this value on production systems.** In almost every situation, you should use the default setting.

Warning: If you set `--syncdelay` (page 510) to 0, MongoDB will not sync the memory mapped files to disk.

The `mongod` (page 503) process writes data very quickly to the journal and lazily to the data files. `syncPeriodSecs` has no effect on the journal files or journaling.

The `serverStatus` (page 347) command reports the background flush thread's status via the `backgroundFlushing` (page 353) field.

command line option!–upgrade

--upgrade

Upgrades the on-disk data format of the files specified by the `--dbpath` (page 546) to the latest version, if needed.

This option only affects the operation of the `mongod` (page 503) if the data files are in an old format.

In most cases you should not set this value, so you can exercise the most control over your upgrade process. See the MongoDB [release notes](http://www.mongodb.org/downloads)¹ (on the download page) for more information about the upgrade process.

command line option!–repair

--repair

Runs a repair routine on all databases. This is equivalent to shutting down and running the `repairDatabase` (page 319) database command on all databases.

Warning: During normal operations, only use the `repairDatabase` (page 319) command and wrappers including `db.repairDatabase()` (page 117) in the `mongo` (page 527) shell and `mongod --repair`, to compact database files and/or reclaim disk space. Be aware that these operations remove and do not save any corrupt data during the repair process.

If you are trying to repair a *replica set* member, and you have access to an intact copy of your data (e.g. a recent backup or an intact member of the *replica set*), you should restore from that intact copy, and **not** use `repairDatabase` (page 319).

¹<http://www.mongodb.org/downloads>

When using *journaling*, there is almost never any need to run `repairDatabase` (page 319). In the event of an unclean shutdown, the server will be able restore the data files to a pristine state automatically.

Changed in version 2.1.2.

If you run the repair option *and* have data in a journal file, the `mongod` (page 503) instance refuses to start. In these cases you should start the `mongod` (page 503) without the `--repair` (page 542) option, which allows the `mongod` (page 503) to recover data from the journal. This completes more quickly and is more likely to produce valid data files. To continue the repair operation despite the journal files, shut down the `mongod` (page 503) cleanly and restart with the `--repair` (page 542) option.

The `--repair` (page 542) option copies data from the source data files into new data files in the `repairPath` and then replaces the original data files with the repaired data files. If `repairPath` is on the same device as `dbPath`, you may interrupt a `mongod` (page 503) running the `--repair` (page 542) option without affecting the integrity of the data set.

command line option!-repairpath <path>

--repairpath <path>

Default: A `_tmp` directory within the path specified by the `dbPath` option.

Specifies the root directory containing MongoDB data files to use for the `--repair` (page 542) operation.

command line option!-objcheck

--objcheck

Forces the `mongod` (page 503) to validate all requests from clients upon receipt to ensure that clients never insert invalid documents into the database. For objects with a high degree of sub-document nesting, the `--objcheck` (page 581) option can have a small impact on performance. You can set `--noobjcheck` (page 547) to disable object checking at runtime.

Changed in version 2.4: MongoDB enables the `--objcheck` (page 581) option by default in order to prevent any client from inserting malformed or invalid BSON into a MongoDB database.

command line option!-noobjcheck

--noobjcheck

New in version 2.4.

Disables the default document validation that MongoDB performs on all incoming BSON documents.

command line option!-noscripting

--noscripting

Disables the scripting engine.

command line option!-notablesan

--notablesan

Forbids operations that require a table scan.

command line option!-journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongod` (page 503) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!-nojournal

--nojournal

Disables the durability journaling. The `mongod` (page 503) instance enables journaling by default in 64-bit versions after v2.0.

command line option!-journalOptions <arguments>

--journalOptions <arguments>

Provides functionality for testing. Not for general use, and will affect data file integrity in the case of abnormal system shutdown.

command line option!-journalCommitInterval <value>

--journalCommitInterval <value>

Default: 100 or 30

The maximum amount of time the `mongod` (page 503) process allows between journal operations. Values can range from 2 to 300 milliseconds. Lower values increase the durability of the journal, at the expense of disk performance.

The default journal commit interval is 100 milliseconds if a single block device (e.g. physical volume, RAID device, or LVM volume) contains both the journal and the data files.

If the journal is on a different block device than the data files the default journal commit interval is 30 milliseconds.

To force `mongod` (page 503) to commit to the journal more frequently, you can specify `j:true`. When a write operation with `j:true` is pending, `mongod` (page 503) will reduce `commitIntervalMs` to a third of the set value.

command line option!-shutdown

--shutdown

Used in *control scripts*, the `--shutdown` (page 512) option cleanly and safely terminates the `mongod` (page 503) process. When invoking `mongod` (page 503) with this option you must set the `--dbpath` (page 546) option either directly or by way of the configuration file and the `--config` (page 519) option.

The `--shutdown` (page 512) option is available only on Linux systems.

Replication Options command line option!-replSet <setname>

--replSet <setname>

Configures replication. Specify a replica set name as an argument to this set. All hosts in the replica set must have the same set name.

If your application connects to more than one replica set, each set should have a distinct name. Some drivers group replica set connections by replica set name.

command line option!-oplogSize <value>

--oplogSize <value>

Specifies a maximum size in megabytes for the replication operation log (i.e., the *oplog*). The `mongod` (page 503) process creates an *oplog* based on the maximum amount of space available. For 64-bit systems, the oplog is typically 5% of available disk space. Once the `mongod` (page 503) has created the oplog for the first time, changing the `--oplogSize` (page 512) option will not affect the size of the oplog.

command line option!-replIndexPrefetch

--replIndexPrefetch

Default: all

New in version 2.2.

Determines which indexes *secondary* members of a *replica set* load into memory before applying operations from the oplog. By default secondaries load all indexes related to an operation into memory before applying operations from the oplog. This option can have one of the following values:

Value	Description
none	Secondaries do not load indexes into memory.
all	Secondaries load all indexes related to an operation.
_id_only	Secondaries load no additional indexes into memory beyond the already existing <code>_id</code> index.

Master-Slave Replication These options provide access to conventional master-slave database replication. While this functionality remains accessible in MongoDB, replica sets are the preferred configuration for database replication.

--master

Configures the `mongod` (page 503) to run as a replication *master*.

command line option!—slave

--slave

Configures the `mongod` (page 503) to run as a replication *slave*.

command line option!—source <host><:port>

--source <host><:port>

For use with the `--slave` (page 513) option, the `--source` option designates the server that this instance will replicate.

command line option!—only <arg>

--only <arg>

For use with the `--slave` (page 513) option, the `--only` option specifies only a single *database* to replicate.

command line option!—slavedelay <value>

--slavedelay <value>

For use with the `--slave` (page 513) option, the `--slavedelay` (page 513) option configures a “delay” in seconds, for this slave to wait to apply operations from the *master* node.

command line option!—autoresync

--autoresync

For use with the `--slave` (page 513) option. When set, the `--autoresync` (page 513) option allows this slave to automatically resync if it is more than 10 seconds behind the master. This setting may be problematic if the `--oplogSize` (page 512) specifies a too small oplog.

If the *oplog* is not large enough to store the difference in changes between the master’s current state and the state of the slave, this instance will forcibly resync itself unnecessarily. If you don’t specify `--autoresync` (page 513), the slave will not attempt an automatic resync more than once in a ten minute period.

command line option!—fastsync

--fastsync

In the context of *replica set* replication, set this option if you have seeded this member with a snapshot of the *dbpath* of another member of the set. Otherwise the `mongod` (page 503) will attempt to perform an initial sync, as though the member were a new member.

Warning: If the data is not perfectly synchronized *and* the `mongod` (page 503) starts with *fastsync*, then the secondary or slave will be permanently out of sync with the primary, which may cause significant consistency problems.

Sharded Cluster Options command line option!—configsvr

--configsvr

Declares that this `mongod` (page 503) instance serves as the *config database* of a sharded cluster. When running with this option, clients will not be able to write data to any database other than `config` and `admin`. The default port for a `mongod` (page 503) with this option is 27019 and the default `--dbpath` (page 546) directory is `/data/configdb`, unless specified.

Changed in version 2.2: The `--configsvr` (page 513) option also sets `--smallfiles` (page 510).

Changed in version 2.4: The `--configsvr` (page 513) option creates a local *oplog*.

Do not use the `--configsvr` (page 513) option with `--replSet` (page 512) or `--shardsvr` (page 514). Config servers cannot be a shard server or part of a *replica set*.

command line option!—shardsvr

--shardsvr

Configures this `mongod` (page 503) instance as a shard in a partitioned cluster. The default port for these instances is 27018. The only effect of `--shardsvr` (page 514) is to change the port number.

command line option!—moveParanoia

--moveParanoia

New in version 2.4.

During chunk migrations, the `--moveParanoia` (page 514) option forces the `mongod` (page 503) instances to save to the `moveChunk` directory of the `storage.dbPath` all the documents migrated from this shard. MongoDB does not delete data stored in `moveChunk`.

Prior to 2.4, `--moveParanoia` (page 514) was the default behavior of MongoDB.

SSL Options

See

<http://docs.mongodb.org/manual/tutorial/configure-ssl> for full documentation of MongoDB's support.

command line option!—sslOnNormalPorts

--sslOnNormalPorts

Deprecated since version 2.6.

Enables SSL for `mongod` (page 503).

With `--sslOnNormalPorts` (page 523), a `mongod` (page 503) requires SSL encryption for all connections on the default MongoDB port, or the port specified by `--port` (page 577). By default, `--sslOnNormalPorts` (page 523) is disabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!—sslMode <mode>

--sslMode <mode>

New in version 2.6.

Enables SSL or mixed SSL on a port. The argument to the `--sslMode` (page 523) option can be one of the following:

Value	Description
disabled	The server does not use SSL.
allowSSL	Connections between servers do not use SSL. For incoming connections, the server accepts both SSL and non-SSL.
preferSSL	Connections between servers use SSL. For incoming connections, the server accepts both SSL and non-SSL.
requireSSL	The server uses and accepts only SSL encrypted connections.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.2.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

When SSL is enabled, you must specify `--sslPEMKeyFile` (page 577).

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.2.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile`). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongod` (page 503) will redact the password from all logging and reporting output.

Changed in version 2.6: If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongod` (page 503) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-clusterAuthMode <option>

--clusterAuthMode <option>

Default: keyFile

New in version 2.6.

The authentication mode used for cluster authentication. If you use *internal x.509 authentication*, specify so here. This option can have one of the following values:

Value	Description
keyFile	Use a keyfile for authentication. Accept only keyfiles.
sendKeyFile	For rolling upgrade purposes. Send a keyfile for authentication but can accept both keyfiles and x.509 certificates.
sendX509	For rolling upgrade purposes. Send the x.509 certificate for authentication but can accept both keyfiles and x.509 certificates.
x509	Recommended. Send the x.509 certificate for authentication and accept only x.509 certificates.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslClusterFile <filename>

--sslClusterFile <filename>

New in version 2.6.

Specifies the .pem file that contains the x.509 certificate-key file for *membership authentication* for the cluster or replica set.

If `--sslClusterFile` (page 524) does not specify the .pem file for internal cluster authentication, the cluster uses the .pem file specified in the `--sslPEMKeyFile` (page 577) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslClusterPassword` <value>

--sslClusterPassword <value>

New in version 2.6.

Specifies the password to de-crypt the x.509 certificate-key file specified with `--sslClusterFile`. Use the `--sslClusterPassword` (page 524) option only if the certificate-key file is encrypted. In all cases, the `mongod` (page 503) will redact the password from all logging and reporting output.

If the x.509 key file is encrypted and you do not specify the `--sslClusterPassword` (page 524) option, the `mongod` (page 503) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslCAFile` <filename>

--sslCAFile <filename>

New in version 2.4.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslCRLFile` <filename>

--sslCRLFile <filename>

New in version 2.4.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslAllowInvalidCertificates`

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for SSL certificates on other servers in the cluster and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslWeakCertificateValidation`

--sslWeakCertificateValidation

New in version 2.4.

Disables the requirement for SSL certificate validation that `--sslCAFile` enables. With the `--sslWeakCertificateValidation` (page 525) option, the `mongod` (page 503) will accept connections when the client does not present a certificate when establishing the connection.

If the client presents a certificate and the `mongod` (page 503) has `--sslWeakCertificateValidation` (page 525) enabled, the `mongod` (page 503) will validate the certificate using the root certificate chain specified by `--sslCAFile` and reject clients with invalid certificates.

Use the `--sslWeakCertificateValidation` (page 525) option if you have a mixed deployment that includes clients that do not or cannot present certificates to the `mongod` (page 503).

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.4.

Directs the `mongod` (page 503) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

Audit Options command line option!-auditDestination

--auditDestination

New in version 2.6.

Enables auditing. The `--auditDestination` (page 526) option can have one of the following values:

Value	Description
syslog	Output the audit events to syslog in JSON format. Not available on Windows. Audit messages have a syslog severity level of <code>info</code> and a facility level of <code>user</code> . The syslog message limit can result in the truncation of audit messages. The auditing system will neither detect the truncation nor error upon its occurrence.
console	Output the audit events to <code>stdout</code> in JSON format.
file	Output the audit events to the file specified in <code>--auditPath</code> (page 526) in the format specified in <code>--auditFormat</code> (page 526).

Note: The `audit` system is available only in MongoDB Enterprise².

command line option!-auditFormat

--auditFormat

New in version 2.6.

Specifies the format of the output file if `--auditDestination` (page 526) is `file`. The `--auditFormat` (page 526) option can have one of the following values:

Value	Description
JSON	Output the audit events in JSON format to the file specified in <code>--auditPath</code> (page 526).
BSON	Output the audit events in BSON binary format to the file specified in <code>--auditPath</code> (page 526).

Printing audit events to a file in JSON format degrades server performance more than printing to a file in BSON format.

Note: The `audit` system is available only in MongoDB Enterprise³.

²<http://www.mongodb.com/products/mongodb-enterprise>

command line option!–auditPath

--auditPath

New in version 2.6.

Specifies the output file for auditing if `--auditDestination` (page 526) has value of file. The `--auditPath` (page 526) option can take either a full path name or a relative path name.

Note: The `audit` system is available only in MongoDB Enterprise⁴.

command line option!–auditFilter

--auditFilter

New in version 2.6.

Specifies the filter to limit the *types of operations* the audit system records. The option takes a document of the form:

```
{ atype: <expression> }
```

For authentication operations, the option can also take a document of the form:

```
{ atype: <expression>, "param.db": <database> }
```

Note: The `audit` system is available only in MongoDB Enterprise⁵.

SNMP Options command line option!–snmp-subagent

--snmp-subagent

Runs SNMP as a subagent. For more information, see [http://docs.mongodb.org/manual/tutorial/monitor-with-](http://docs.mongodb.org/manual/tutorial/monitor-with-snmp/)

command line option!–snmp-master

--snmp-master

Runs SNMP as a master. For more information, see [http://docs.mongodb.org/manual/tutorial/monitor-with-](http://docs.mongodb.org/manual/tutorial/monitor-with-snmp/)

mongos

Synopsis

`mongos` (page 518) for “MongoDB Shard,” is a routing service for MongoDB shard configurations that processes queries from the application layer, and determines the location of this data in the *sharded cluster*, in order to complete these operations. From the perspective of the application, a `mongos` (page 518) instance behaves identically to any other MongoDB instance.

Options

mongos

³<http://www.mongodb.com/products/mongodb-enterprise>

⁴<http://www.mongodb.com/products/mongodb-enterprise>

⁵<http://www.mongodb.com/products/mongodb-enterprise>

Core Options**mongos**

command line option!–help, -h

--help, -h

Returns information on the options and use of `mongos` (page 518).

command line option!–version

--version

Returns the `mongos` (page 518) release number.

command line option!–config <filename>, -f

--config <filename>, -f

Specifies a configuration file for runtime configuration options. The configuration file is the preferred method for runtime configuration of `mongos` (page 518). The options are equivalent to the command-line configuration options. See <http://docs.mongodb.org/manualreference/configuration-options> for more information.

Ensure the configuration file uses ASCII encoding. The `mongos` (page 518) instance does not support configuration files with non-ASCII encoding, including UTF-8.

command line option!–verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!–quiet

--quiet

Runs the `mongos` (page 518) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!–port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!–bind_ip <ip address>

--bind_ip <ip address>

Default: All interfaces. .. versionchanged:: 2.6.0 The deb and rpm packages include a default configuration file that sets `{{role}}` to `127.0.0.1`.

Specifies the IP address that `mongos` (page 518) binds to in order to listen for connections from applications. You may attach `mongos` (page 518) to any interface. When attaching `mongos` (page 518) to a publicly accessible interface, ensure that you have implemented proper authentication and firewall restrictions to protect the integrity of your database.

command line option!–maxConns <number>

--maxConns <number>

Specifies the maximum number of simultaneous connections that `mongos` (page 518) will accept. This setting

will have no effect if the value of this setting is higher than your operating system's configured maximum connection tracking threshold.

This setting is particularly useful for `mongos` (page 518) if you have a client that creates a number of connections but allows them to timeout rather than close the connections. When you set `maxIncomingConnections`, ensure the value is slightly higher than the size of the connection pool or the total number of connections to prevent erroneous connection spikes from propagating to the members of a *sharded cluster*.

Changed in version 2.6: MongoDB removed the upward limit on the `maxIncomingConnections` setting.

command line option!-syslog

--syslog

Sends all logging output to the host's *syslog* system rather than to standard output or to a log file. , as with `--logpath` (page 520).

The `--syslog` (page 520) option is not supported on Windows.

command line option!-syslogFacility <string>

--syslogFacility <string>

Default: user

Specifies the facility level used when logging messages to syslog. The value you specify must be supported by your operating system's implementation of syslog. To use this option, you must enable the `--syslog` (page 520) option.

command line option!-logpath <path>

--logpath <path>

Sends all diagnostic logging information to a log file instead of to standard output or to the host's *syslog* system. MongoDB creates the log file at the path you specify.

By default, MongoDB overwrites the log file when the process restarts. To instead append to the log file, set the `--logappend` (page 520) option.

command line option!-logappend

--logappend

Appends new entries to the end of the log file rather than overwriting the content of the log when the `mongos` (page 518) instance restarts.

command line option!-timestampFormat <string>

--timestampFormat <string>

Default: iso8601-local

The time format for timestamps in log messages. Specify one of the following values:

Value	Description
ctime	Displays timestamps as Wed Dec 31 18:17:54.811.
iso8601-utc	Displays timestamps in Coordinated Universal Time (UTC) in the ISO-8601 format. For example, for New York at the start of the Epoch: 1970-01-01T00:00:00.000Z
iso8601-local	Displays timestamps in local time in the ISO-8601 format. For example, for New York at the start of the Epoch: 1969-12-31T19:00:00.000+0500

command line option!-pidfilepath <path>

--pidfilepath <path>

Specifies a file location to hold the process ID of the `mongos` (page 518) process. This is useful for tracking the `mongos` (page 518) process in combination with the `--fork` (page 521) option. Without a specified `--pidfilepath` (page 520) option, the process creates no PID file.

command line option!-keyFile <file>

--keyFile <file>

Specifies the path to a key file to that stores the shared secret that MongoDB processes use to authenticate to each other in a *sharded cluster* or *replica set*. `--keyFile` (page 521) implies `--auth` (page 507). See *inter-process-auth* for more information.

command line option!-setParameter <options>

--setParameter <options>

Specifies one of the MongoDB parameters described in <http://docs.mongodb.org/manualreference/parameters>. You can specify multiple `setParameter` fields.

command line option!-httpinterface

--httpinterface

New in version 2.6.

Enables the HTTP interface. Enabling the interface can increase network exposure.

Leave the HTTP interface *disabled* for production deployments. If you *do* enable this interface, you should only allow trusted clients to access this port. See *security-firewalls*.

Note: In MongoDB Enterprise, the HTTP Console does not support Kerberos Authentication.

command line option!-noinxsocket

--noinxsocket

Disables listening on the UNIX domain socket. The `mongos` (page 518) process always listens on the UNIX socket unless one of the following is true:

- `--noinxsocket` (page 521) is set
- `bindIp` is not set
- `bindIp` does not specify `127.0.0.1`

New in version 2.6: `mongos` (page 518) installed from official `.deb` and `.rpm` packages have the `bind_ip` configuration set to `127.0.0.1` by default.

command line option!-unixSocketPrefix <path>

--unixSocketPrefix <path>

Default: `/tmp`

The path for the UNIX socket. If this option has no value, the `mongos` (page 518) process creates a socket with <http://docs.mongodb.org/manualtmp> as a prefix. MongoDB creates and listens on a UNIX socket unless one of the following is true:

- `--noinxsocket` (page 521) is set
- `bindIp` is not set
- `bindIp` does not specify `127.0.0.1`

command line option!-fork

--fork

Enables a *daemon* mode that runs the `mongos` (page 518) process in the background. By default `mongos` (page 518) does not run as a daemon: typically you will run `mongos` (page 518) as a daemon, either by using `--fork` (page 521) or by using a controlling process that handles the daemonization process (e.g. as with `upstart` and `systemd`).

Sharded Cluster Options command line option!—configdb <config1>,<config2>,<config3>

--configdb <config1>,<config2>,<config3>

Specifies the *configuration database* for the *sharded cluster*. You must specify either 1 or 3 configuration servers, in a comma separated list.

All *mongos* (page 518) instances **must** specify the hosts in the *--configdb* (page 522) option in the in the same order.

If your configuration databases reside in more that one data center, order the hosts so that the config database that is closest to the majority of your *mongos* (page 518) instances is first servers in the list.

Warning: Never remove a config server from this setting, even if the config server is not available or offline.

command line option!—localThreshold

--localThreshold

Default: 15

Affects the logic that *mongos* (page 518) uses when selecting *replica set* members to pass read operations to from clients. Specify a value in milliseconds. The default value of 15 corresponds to the default value in all of the client drivers.

When *mongos* (page 518) receives a request that permits reads to *secondary* members, the *mongos* (page 518) will:

- Find the member of the set with the lowest ping time.
- Construct a list of replica set members that is within a ping time of 15 milliseconds of the nearest suitable member of the set.

If you specify a value for the *--localThreshold* (page 522) option, *mongos* (page 518) will construct the list of replica members that are within the latency allowed by this value.

- Select a member to read from at random from this list.

The ping time used for a member compared by the *--localThreshold* (page 522) setting is a moving average of recent ping times, calculated at most every 10 seconds. As a result, some queries may reach members above the threshold until the *mongos* (page 518) recalculates the average.

See the *replica-set-read-preference-behavior-member-selection* section of the *read preference* documentation for more information.

command line option!—upgrade

--upgrade

Updates the meta data format used by the *config database*.

command line option!—chunkSize <value>

--chunkSize <value>

Default: 64

Determines the size in megabytes of each *chunk* in the *sharded cluster*. A size of 64 megabytes is ideal in most deployments: larger chunk size can lead to uneven data distribution; smaller chunk size can lead to inefficient movement of chunks between nodes.

This option affects chunk size *only* when you initialize the cluster for the first time. If you later modify the option, the new value has no effect. See the <http://docs.mongodb.org/manual/tutorial/modify-chunk-size-in-sharded-cluster> procedure if you need to change the chunk size on an existing sharded cluster.

command line option!--noAutoSplit

--noAutoSplit

Prevents `mongos` (page 518) from automatically inserting metadata splits in a *sharded collection*. If set on all `mongos` (page 518) instances, this prevents MongoDB from creating new chunks as the data in a collection grows.

Because any `mongos` (page 518) in a cluster can create a split, to totally disable splitting in a cluster you must set `--noAutoSplit` (page 523) on all `mongos` (page 518).

Warning: With `--noAutoSplit` (page 523) enabled, the data in your sharded cluster may become imbalanced over time. Enable with caution.

SSL Options

See

<http://docs.mongodb.org/manual/tutorial/configure-ssl> for full documentation of MongoDB's support.

command line option!--sslOnNormalPorts

--sslOnNormalPorts

Deprecated since version 2.6.

Enables SSL for `mongos` (page 518).

With `--sslOnNormalPorts` (page 523), a `mongos` (page 518) requires SSL encryption for all connections on the default MongoDB port, or the port specified by `--port` (page 577). By default, `--sslOnNormalPorts` (page 523) is disabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!--sslMode <mode>

--sslMode <mode>

New in version 2.6.

Enables SSL or mixed SSL on a port. The argument to the `--sslMode` (page 523) option can be one of the following:

Value	Description
<code>disabled</code>	The server does not use SSL.
<code>allowSSL</code>	Connections between servers do not use SSL. For incoming connections, the server accepts both SSL and non-SSL.
<code>preferSSL</code>	Connections between servers use SSL. For incoming connections, the server accepts both SSL and non-SSL.
<code>requireSSL</code>	The server uses and accepts only SSL encrypted connections.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!--sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.2.

Specifies the `.pem` file that contains both the SSL certificate and key. Specify the file name of the `.pem` file using relative or absolute paths.

When SSL is enabled, you must specify `--sslPEMKeyFile` (page 577).

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.2.

Specifies the password to de-crypt the certificate-key file (i.e. `--sslPEMKeyFile`). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongos` (page 518) will redact the password from all logging and reporting output.

Changed in version 2.6: If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongos` (page 518) will prompt for a passphrase. See `ssl-certificate-password`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-clusterAuthMode <option>

--clusterAuthMode <option>

Default: keyFile

New in version 2.6.

The authentication mode used for cluster authentication. If you use *internal x.509 authentication*, specify so here. This option can have one of the following values:

Value	Description
keyFile	Use a keyfile for authentication. Accept only keyfiles.
sendKeyFile	For rolling upgrade purposes. Send a keyfile for authentication but can accept both keyfiles and x.509 certificates.
sendX509	For rolling upgrade purposes. Send the x.509 certificate for authentication but can accept both keyfiles and x.509 certificates.
x509	Recommended. Send the x.509 certificate for authentication and accept only x.509 certificates.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslClusterFile <filename>

--sslClusterFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the x.509 certificate-key file for *membership authentication* for the cluster or replica set.

If `--sslClusterFile` (page 524) does not specify the `.pem` file for internal cluster authentication, the cluster uses the `.pem` file specified in the `--sslPEMKeyFile` (page 577) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslClusterPassword <value>

--sslClusterPassword <value>

New in version 2.6.

Specifies the password to de-crypt the x.509 certificate-key file specified with `--sslClusterFile`. Use the `--sslClusterPassword` (page 524) option only if the certificate-key file is encrypted. In all cases, the `mongos` (page 518) will redact the password from all logging and reporting output.

If the x.509 key file is encrypted and you do not specify the `--sslClusterPassword` (page 524) option, the `mongos` (page 518) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.4.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.4.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslWeakCertificateValidation

--sslWeakCertificateValidation

New in version 2.4.

Disables the requirement for SSL certificate validation that `--sslCAFile` enables. With the `--sslWeakCertificateValidation` (page 525) option, the `mongos` (page 518) will accept connections when the client does not present a certificate when establishing the connection.

If the client presents a certificate and the `mongos` (page 518) has `--sslWeakCertificateValidation` (page 525) enabled, the `mongos` (page 518) will validate the certificate using the root certificate chain specified by `--sslCAFile` and reject clients with invalid certificates.

Use the `--sslWeakCertificateValidation` (page 525) option if you have a mixed deployment that includes clients that do not or cannot present certificates to the `mongos` (page 518).

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for SSL certificates on other servers in the cluster and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.4.

Directs the `mongos` (page 518) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMode` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

Audit Options command line option!–auditDestination

--auditDestination

New in version 2.6.

Enables auditing. The `--auditDestination` (page 526) option can have one of the following values:

Value	Description
syslog	Output the audit events to syslog in JSON format. Not available on Windows. Audit messages have a syslog severity level of <code>info</code> and a facility level of <code>user</code> . The syslog message limit can result in the truncation of audit messages. The auditing system will neither detect the truncation nor error upon its occurrence.
console	Output the audit events to <code>stdout</code> in JSON format.
file	Output the audit events to the file specified in <code>--auditPath</code> (page 526) in the format specified in <code>--auditFormat</code> (page 526).

Note: The `audit` system is available only in MongoDB Enterprise⁶.

command line option!–auditFormat

--auditFormat

New in version 2.6.

Specifies the format of the output file if `--auditDestination` (page 526) is `file`. The `--auditFormat` (page 526) option can have one of the following values:

Value	Description
JSON	Output the audit events in JSON format to the file specified in <code>--auditPath</code> (page 526).
BSON	Output the audit events in BSON binary format to the file specified in <code>--auditPath</code> (page 526).

Printing audit events to a file in JSON format degrades server performance more than printing to a file in BSON format.

Note: The `audit` system is available only in MongoDB Enterprise⁷.

command line option!–auditPath

--auditPath

New in version 2.6.

Specifies the output file for auditing if `--auditDestination` (page 526) has value of `file`. The `--auditPath` (page 526) option can take either a full path name or a relative path name.

Note: The `audit` system is available only in MongoDB Enterprise⁸.

command line option!–auditFilter

--auditFilter

New in version 2.6.

⁶<http://www.mongodb.com/products/mongodb-enterprise>

⁷<http://www.mongodb.com/products/mongodb-enterprise>

⁸<http://www.mongodb.com/products/mongodb-enterprise>

Specifies the filter to limit the *types of operations* the audit system records. The option takes a document of the form:

```
{ atype: <expression> }
```

For authentication operations, the option can also take a document of the form:

```
{ atype: <expression>, "param.db": <database> }
```

Note: The `audit` system is available only in [MongoDB Enterprise](#)⁹.

Additional Options `command line option!-ipv6`

--ipv6

Enables IPv6 support and allows the `mongos` (page 518) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!`-jsonp`

--jsonp

Permits *JSONP* access via an HTTP interface. Enabling the interface can increase network exposure. The `--jsonp` (page 527) option enables the HTTP interface, even if the `HTTP interface` option is disabled.

command line option!`-noscripting`

--noscripting

Disables the scripting engine.

mongo

Description

mongo

`mongo` (page 527) is an interactive JavaScript shell interface to MongoDB, which provides a powerful interface for systems administrators as well as a way for developers to test queries and operations directly with the database. `mongo` (page 527) also provides a fully functional JavaScript environment for use with a MongoDB. This document addresses the basic invocation of the `mongo` (page 527) shell and an overview of its usage.

Options

Core Options

mongo

command line option!`-shell`

--shell

Enables the shell interface. If you invoke the `mongo` (page 527) command and specify a JavaScript file as an argument, or use `--eval` (page 528) to specify JavaScript on the command line, the `--shell` (page 527) option provides the user with a shell prompt after the file finishes executing.

command line option!`-nodb`

--nodb

Prevents the shell from connecting to any database instances. Later, to connect to a database within the shell, see *mongo-shell-new-connections*.

⁹<http://www.mongodb.com/products/mongodb-enterprise>

command line option!-norc

--norc

Prevents the shell from sourcing and evaluating `~/ .mongorc.js` on start up.

command line option!-quiet

--quiet

Silences output from the shell during the connection process.

command line option!-port <port>

--port <port>

Specifies the port where the `mongod` (page 503) or `mongos` (page 518) instance is listening. If `--port` (page 577) is not specified, `mongo` (page 527) attempts to connect to port 27017.

command line option!-host <hostname>

--host <hostname>

Specifies the name of the host machine where the `mongod` (page 503) or `mongos` (page 518) is running. If this is not specified, `mongo` (page 527) attempts to connect to a MongoDB process running on the localhost.

command line option!-eval <javascript>

--eval <javascript>

Evaluates a JavaScript expression that is specified as an argument. `mongo` (page 527) does not load its own environment when evaluating code. As a result many options of the shell environment are not available.

command line option!-username <username>, -u

--username <username>, **-u**

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, **-p**

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongo` (page 527).

command line option!-version

--version

Returns the `mongo` (page 527) release number.

command line option!-verbose

--verbose

Increases the verbosity of the output of the shell during the connection process.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongo` (page 527) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

<db address>

Specifies the “database address” of the database to connect to. For example:

```
mongo admin
```

The above command will connect the `mongo` (page 527) shell to the `admin database` on the local machine. You may specify a remote database instance, with the resolvable hostname or IP address. Separate the database name from the hostname using a `http://docs.mongodb.org/manual` character. See the following examples:

```
mongo mongodb1.example.net
mongo mongodb1/admin
mongo 10.8.8.10/test
```

<file.js>

Specifies a JavaScript file to run and then exit. Generally this should be the last option specified.

Optional

To specify a JavaScript file to execute *and* allow `mongo` (page 527) to prompt you for a password using `--password` (page 578), pass the filename as the first parameter with `--username` (page 578) and `--password` (page 578) as the last options, as in the following:

```
mongo file.js --username username --password
```

Use the `--shell` (page 527) option to return to a shell after the file finishes running.

Authentication Options `command line option!-authenticationDatabase <dbname>`

`--authenticationDatabase <dbname>`

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongo` (page 527) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

`command line option!-authenticationMechanism <name>`

`--authenticationMechanism <name>`

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the PLAIN and MONGODB-X509 authentication mechanisms.

Specifies the authentication mechanism the `mongo` (page 527) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use PLAIN for authenticating in-database users. PLAIN transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ¹² .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ¹³ .

¹⁰<http://www.mongodb.com/products/mongodb-enterprise>

¹¹<http://www.mongodb.com/products/mongodb-enterprise>

¹²<http://www.mongodb.com/products/mongodb-enterprise>

¹³<http://www.mongodb.com/products/mongodb-enterprise>

SSL Options command line option!–ssl

--ssl

New in version 2.2.

Enables connection to a [mongod](#) (page 503) or [mongos](#) (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!–sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.4.

Specifies the `.pem` file that contains both the SSL certificate and key. Specify the file name of the `.pem` file using relative or absolute paths.

This option is required when using the `--ssl` option to connect to a [mongod](#) (page 503) or [mongos](#) (page 518) that has `CAFile` enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!–sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.4.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile`). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the [mongo](#) (page 527) will redact the password from all logging and reporting output.

Changed in version 2.6: If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the [mongo](#) (page 527) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!–sslCAFile <filename>

--sslCAFile <filename>

New in version 2.4.

Specifies the `.pem` file that contains the root certificate chain from the Certificate Authority. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!–sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.4.

Specifies the `.pem` file that contains the Certificate Revocation List. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!–sslFIPSMode

--sslFIPSMODE

New in version 2.4.

Directs the `mongo` (page 527) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

Files

`~/ .dbshell` `mongo` (page 527) maintains a history of commands in the `.dbshell` file.

Note: `mongo` (page 527) does not record interaction related to authentication in the history file, including `authenticate` (page 249) and `db.addUser()` (page 143).

Warning: Versions of Windows `mongo.exe` earlier than 2.2.0 will save the `.dbshell` file in the `mongo.exe` working directory.

`~/ .mongorc.js` `mongo` (page 527) will read the `.mongorc.js` file from the home directory of the user invoking `mongo` (page 527). In the file, users can define variables, customize the `mongo` (page 527) shell prompt, or update information that they would like updated every time they launch a shell. If you use the shell to evaluate a JavaScript file or expression either on the command line with `--eval` (page 528) or by specifying *a .js file to mongo* (page 529), `mongo` (page 527) will read the `.mongorc.js` file *after* the JavaScript has finished processing.

Specify the `--norc` (page 528) option to disable reading `.mongorc.js`.

`/etc/mongorc.js` Global `mongorc.js` file which the `mongo` (page 527) shell evaluates upon start-up. If a user also has a `.mongorc.js` file located in the `HOME` (page 532) directory, the `mongo` (page 527) shell evaluates the global `/etc/mongorc.js` file *before* evaluating the user's `.mongorc.js` file.

`/etc/mongorc.js` must have read permission for the user running the shell. The `--norc` (page 528) option for `mongo` (page 527) suppresses only the user's `.mongorc.js` file.

On Windows, the global `mongorc.js` `</etc/mongorc.js>` exists in the `%ProgramData%\MongoDB` directory.

`http://docs.mongodb.org/manual/tmp/mongo_edit<time_t>.js` Created by `mongo` (page 527) when editing a file. If the file exists, `mongo` (page 527) will append an integer from 1 to 10 to the time value to attempt to create a unique file.

`%TEMP%mongo_edit<time_t>.js` Created by `mongo.exe` on Windows when editing a file. If the file exists, `mongo` (page 527) will append an integer from 1 to 10 to the time value to attempt to create a unique file.

Environment

EDITOR

Specifies the path to an editor to use with the `edit` shell command. A JavaScript variable `EDITOR` will override the value of `EDITOR` (page 532).

HOME

Specifies the path to the home directory where `mongo` (page 527) will read the `.mongorc.js` file and write the `.dbshell` file.

HOMEDRIVE

On Windows systems, `HOMEDRIVE` (page 532) specifies the path the directory where `mongo` (page 527) will read the `.mongorc.js` file and write the `.dbshell` file.

HOMEPATH

Specifies the Windows path to the home directory where `mongo` (page 527) will read the `.mongorc.js` file and write the `.dbshell` file.

Keyboard Shortcuts

The `mongo` (page 527) shell supports the following keyboard shortcuts: ¹⁴

Keybinding	Function
Up arrow	Retrieve previous command from history
Down-arrow	Retrieve next command from history
Home	Go to beginning of the line
End	Go to end of the line
Tab	Autocomplete method/command
Left-arrow	Go backward one character
Right-arrow	Go forward one character
Ctrl-left-arrow	Go backward one word
Ctrl-right-arrow	Go forward one word
Meta-left-arrow	Go backward one word
Meta-right-arrow	Go forward one word
Ctrl-A	Go to the beginning of the line
Ctrl-B	Go backward one character
Ctrl-C	Exit the <code>mongo</code> (page 527) shell
Ctrl-D	Delete a char (or exit the <code>mongo</code> (page 527) shell)
Ctrl-E	Go to the end of the line
Ctrl-F	Go forward one character
Ctrl-G	Abort
Ctrl-J	Accept/evaluate the line
Ctrl-K	Kill/erase the line
Ctrl-L or type <code>cls</code>	Clear the screen
Ctrl-M	Accept/evaluate the line
Ctrl-N	Retrieve next command from history
Ctrl-P	Retrieve previous command from history
Ctrl-R	Reverse-search command history
Ctrl-S	Forward-search command history
Ctrl-T	Transpose characters
Ctrl-U	Perform Unix line-discard
Ctrl-W	Perform Unix word-rubout
Continued on next page	

¹⁴ MongoDB accommodates multiple keybinding. Since 2.0, `mongo` (page 527) includes support for basic emacs keybindings.

Table 4.1 – continued from previous page

Keybinding	Function
Ctrl-Y	Yank
Ctrl-Z	Suspend (job control works in linux)
Ctrl-H	Backward-delete a character
Ctrl-I	Complete, same as Tab
Meta-B	Go backward one word
Meta-C	Capitalize word
Meta-D	Kill word
Meta-F	Go forward one word
Meta-L	Change word to lowercase
Meta-U	Change word to uppercase
Meta-Y	Yank-pop
Meta-Backspace	Backward-kill word
Meta-<	Retrieve the first command in command history
Meta->	Retrieve the last command in command history

Use

Typically users invoke the shell with the `mongo` (page 527) command at the system prompt. Consider the following examples for other scenarios.

To connect to a database on a remote host using authentication and a non-standard port, use the following form:

```
mongo --username <user> --password <pass> --host <host> --port 28015
```

Alternatively, consider the following short form:

```
mongo -u <user> -p <pass> --host <host> --port 28015
```

Replace `<user>`, `<pass>`, and `<host>` with the appropriate values for your situation and substitute or omit the `--port` (page 577) as needed.

To execute a JavaScript file without evaluating the `~/.mongorc.js` file before starting a shell session, use the following form:

```
mongo --shell --norc alternate-environment.js
```

To execute a JavaScript file with authentication, with password prompted rather than provided on the command-line, use the following form:

```
mongo script-file.js -u <user> -p
```

To print return a query as *JSON*, from the system prompt using the `--eval` option, use the following form:

```
mongo --eval 'db.collection.find().forEach(printjson)'
```

Use single quotes (e.g. `'`) to enclose the JavaScript, as well as the additional JavaScript required to generate this output.

4.1.2 Windows Services

The `mongod.exe` (page 534) and `mongos.exe` (page 535) describe the options available for configuring MongoDB when running as a Windows Service. The `mongod.exe` (page 534) and `mongos.exe` (page 535) binaries provide a superset of the `mongod` (page 503) and `mongos` (page 518) options.

`mongod.exe`

Synopsis

`mongod.exe` (page 534) is the build of the MongoDB daemon (i.e. `mongod` (page 503)) for the Windows platform. `mongod.exe` (page 534) has all of the features of `mongod` (page 503) on Unix-like platforms and is completely compatible with the other builds of `mongod` (page 503). In addition, `mongod.exe` (page 534) provides several options for interacting with the Windows platform itself.

This document *only* references options that are unique to `mongod.exe` (page 534). All `mongod` (page 503) options are available. See the `mongod` (page 503) and the <http://docs.mongodb.org/manualreference/configuration-options> documents for more information regarding `mongod.exe` (page 534).

To install and use `mongod.exe` (page 534), read the <http://docs.mongodb.org/manualtutorial/install-mongodb-document>.

Options

`mongod.exe`

`mongod.exe`

command line option!-install

--install

Installs `mongod.exe` (page 534) as a Windows Service and exits.

command line option!-remove

--remove

Removes the `mongod.exe` (page 534) Windows Service. If `mongod.exe` (page 534) is running, this operation will stop and then remove the service.

`--remove` (page 535) requires the `--serviceName` (page 536) if you configured a non-default `--serviceName` (page 536) during the `--install` (page 535) operation.

command line option!-reinstall

--reinstall

Removes `mongod.exe` (page 534) and reinstalls `mongod.exe` (page 534) as a Windows Service.

command line option!-serviceName name

--serviceName name

Default: MongoDB

Set the service name of `mongod.exe` (page 534) when running as a Windows Service. Use this name with the `net start <name>` and `net stop <name>` operations.

You must use `--serviceName` (page 536) in conjunction with either the `--install` (page 535) or `--remove` (page 535) install option.

command line option!-serviceDisplayName <name>

--serviceDisplayName <name>

Default: MongoDB

Sets the name listed for MongoDB on the Services administrative application.

command line option!-serviceDescription <description>

--serviceDescription <description>

Default: MongoDB Server

Sets the `mongod.exe` (page 534) service description.

You must use `--serviceDescription` (page 536) in conjunction with the `--install` (page 535) option.

For descriptions that contain spaces, you must enclose the description in quotes.

command line option!—serviceUser <user>

--serviceUser <user>

Runs the `mongod.exe` (page 534) service in the context of a certain user. This user must have “Log on as a service” privileges.

You must use `--serviceUser` (page 536) in conjunction with the `--install` (page 535) option.

command line option!—servicePassword <password>

--servicePassword <password>

Sets the password for <user> for `mongod.exe` (page 534) when running with the `--serviceUser` (page 536) option.

You must use `--servicePassword` (page 536) in conjunction with the `--install` (page 535) option.

mongos.exe

Synopsis

`mongos.exe` (page 535) is the build of the MongoDB Shard (i.e. `mongos` (page 518)) for the Windows platform. `mongos.exe` (page 535) has all of the features of `mongos` (page 518) on Unix-like platforms and is completely compatible with the other builds of `mongos` (page 518). In addition, `mongos.exe` (page 535) provides several options for interacting with the Windows platform itself.

This document only references options that are unique to `mongos.exe` (page 535). All `mongos` (page 518) options are available. See the `mongos` (page 518) and the <http://docs.mongodb.org/manualreference/configuration-options> documents for more information regarding `mongos.exe` (page 535).

To install and use `mongos.exe` (page 535), read the <http://docs.mongodb.org/manualtutorial/install-mongodb-document>.

Options

`mongos.exe`

`mongos.exe`

command line option!—install

--install

Installs `mongos.exe` (page 535) as a Windows Service and exits.

command line option!—remove

--remove

Removes the `mongos.exe` (page 535) Windows Service. If `mongos.exe` (page 535) is running, this operation will stop and then remove the service.

`--remove` (page 535) requires the `--serviceName` (page 536) if you configured a non-default `--serviceName` (page 536) during the `--install` (page 535) operation.

command line option!`--reinstall`

`--reinstall`

Removes `mongos.exe` (page 535) and reinstalls `mongos.exe` (page 535) as a Windows Service.

command line option!`--serviceName` *name*

`--serviceName` *name*

Default: MongoS

Set the service name of `mongos.exe` (page 535) when running as a Windows Service. Use this name with the `net start <name>` and `net stop <name>` operations.

You must use `--serviceName` (page 536) in conjunction with either the `--install` (page 535) or `--remove` (page 535) install option.

command line option!`--serviceDisplayName` *<name>*

`--serviceDisplayName` *<name>*

Default: Mongo DB Router

Sets the name listed for MongoDB on the Services administrative application.

command line option!`--serviceDescription` *<description>*

`--serviceDescription` *<description>*

Default: Mongo DB Sharding Router

Sets the `mongos.exe` (page 535) service description.

You must use `--serviceDescription` (page 536) in conjunction with the `--install` (page 535) option.

For descriptions that contain spaces, you must enclose the description in quotes.

command line option!`--serviceUser` *<user>*

`--serviceUser` *<user>*

Runs the `mongos.exe` (page 535) service in the context of a certain user. This user must have “Log on as a service” privileges.

You must use `--serviceUser` (page 536) in conjunction with the `--install` (page 535) option.

command line option!`--servicePassword` *<password>*

`--servicePassword` *<password>*

Sets the password for *<user>* for `mongos.exe` (page 535) when running with the `--serviceUser` (page 536) option.

You must use `--servicePassword` (page 536) in conjunction with the `--install` (page 535) option.

4.1.3 Binary Import and Export Tools

`mongodump` (page 537) provides a method for creating *BSON* dump files from the `mongod` (page 503) instances, while `mongorestore` (page 543) makes it possible to restore these dumps. `bsondump` (page 549) converts *BSON* dump files into *JSON*. The `mongooplog` (page 551) utility provides the ability to stream *oplog* entries outside of normal replication.

`mongodump`

Synopsis

`mongodump` (page 537) is a utility for creating a binary export of the contents of a database. Consider using this

utility as part an effective backup strategy. Use `mongodump` (page 537) in conjunction with `mongorestore` (page 543) to restore databases.

`mongodump` (page 537) can read data from either `mongod` (page 503) or `mongos` (page 518) instances, in addition to reading directly from MongoDB data files without an active `mongod` (page 503).

See also:

`mongorestore` (page 543), <http://docs.mongodb.org/manual/tutorial/backup-sharded-cluster-with-dat>, and <http://docs.mongodb.org/manual/core/backups>.

Behavior

`mongodump` (page 537) does *not* dump the content of the `local` database.

The data format used by `mongodump` (page 537) from version 2.2 or later is *incompatible* with earlier versions of `mongod` (page 503). Do not use recent versions of `mongodump` (page 537) to back up older data stores.

When running `mongodump` (page 537) against a `mongos` (page 518) instance where the *sharded cluster* consists of *replica sets*, the *read preference* of the operation will prefer reads from *secondary* members of the set.

Changed in version 2.2: When used in combination with `fsync` (page 312) or `db.fsyncLock()` (page 109), `mongod` (page 503) may block some reads, including those from `mongodump` (page 537), when queued write operation waits behind the `fsync` (page 312) lock.

Required Access

Backup Collections To backup all the databases in a cluster via `mongodump` (page 537), you should have the backup role. The backup role provides all the needed privileges for backing up all database. The role confers no additional access, in keeping with the policy of *least privilege*.

To backup a given database, you must have `read` access on the database. Several roles provide this access, including the backup role.

To backup the `system.profile` collection in a database, you must have `read` access on certain system collections in the database. Several roles provide this access, including the `clusterAdmin` and `dbAdmin` roles.

Backup Users Changed in version 2.6.

To backup users and *user-defined roles* for a given database, you must have access to the `admin` database. MongoDB stores the user data and role definitions for all databases in the `admin` database.

Specifically, to backup a given database's users, you must have the `find` *action* on the `admin` database's `admin.system.users` (page 601) collection. The `backup` and `userAdminAnyDatabase` roles both provide this privilege.

To backup the user-defined roles on a database, you must have the `find` *action* on the `admin` database's `admin.system.roles` (page 601) collection. Both the `backup` and `userAdminAnyDatabase` roles provide this privilege.

Options

`mongodump`

`mongodump`

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongodump` (page 537).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

--quiet

Runs the `mongodump` (page 537) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `mongodump` (page 537) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, -h

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongodump` (page 537) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongodump` (page 537) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has CAFile enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongodump` (page 537) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongodump` (page 537) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMode

--sslFIPSMode

New in version 2.6.

Directs the `mongodump` (page 537) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMode` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, -u

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, -p

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongodump` (page 537) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the `PLAIN` and `MONGODB-X509` authentication mechanisms.

Specifies the authentication mechanism the `mongodump` (page 537) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use <code>PLAIN</code> for authenticating in-database users. <code>PLAIN</code> transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ¹⁷ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ¹⁸ .

command line option!-dbpath <path>

--dbpath <path>

Specifies the directory of the MongoDB data files. The `--dbpath` (page 546) option lets the `mongodump` (page 537) attach directly to the local data files without going through a running `mongod` (page 503). When run

¹⁵<http://www.mongodb.com/products/mongodb-enterprise>

¹⁶<http://www.mongodb.com/products/mongodb-enterprise>

¹⁷<http://www.mongodb.com/products/mongodb-enterprise>

¹⁸<http://www.mongodb.com/products/mongodb-enterprise>

with `--dbpath` (page 546), the `mongodump` (page 537) locks access to the data files. No `mongod` (page 503) can access the files while the `mongodump` (page 537) process runs.

command line option!—directoryperdb

--directoryperdb

When used in conjunction with the corresponding option in `mongod` (page 503), allows the `mongodump` (page 537) to access data from MongoDB instances that use an on-disk format where every database has a distinct directory. This option is only relevant when specifying the `--dbpath` (page 546) option.

command line option!—journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongodump` (page 537) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!—db <database>, -d

--db <database>, -d

Specifies a database to backup. If you do not specify a database, `mongodump` (page 537) copies all databases in this instance into the dump files.

command line option!—collection <collection>, -c

--collection <collection>, -c

Specifies a collection to backup. If you do not specify a collection, this option copies all collections in the specified database or instance to the dump files.

command line option!—out <path>, -o

--out <path>, -o

Specifies the directory where `mongodump` (page 537) saves the output of the database dump. By default, `mongodump` (page 537) saves output files in a directory named `dump` in the current working directory.

To send the database dump to standard output, specify “-” instead of a path. Write to standard output if you want process the output before saving it, such as to use `gzip` to compress the dump. When writing standard output, `mongodump` (page 537) does not write the metadata that writes in a `<dbname>.metadata.json` file when writing to files directly.

command line option!—query <json>, -q

--query <json>, -q

Provides a *JSON document* as a query that optionally limits the documents included in the output of `mongodump` (page 537).

command line option!—oplog

--oplog

Ensures that `mongodump` (page 537) creates a dump of the database that includes a partial *oplog* containing operations from the duration of the `mongodump` (page 537) operation. This oplog produces an effective point-in-time snapshot of the state of a `mongod` (page 503) instance. To restore to a specific point-in-time backup, use the output created with this option in conjunction with `mongorestore --oplogReplay`.

Without `--oplog` (page 541), if there are write operations during the dump operation, the dump will not reflect a single moment in time. Changes made to the database during the update process can affect the output of the backup.

`--oplog` (page 541) has no effect when running `mongodump` (page 537) against a `mongos` (page 518) instance to dump the entire contents of a sharded cluster. However, you can use `--oplog` (page 541) to dump individual shards.

`--oplog` (page 541) only works against nodes that maintain an *oplog*. This includes all members of a replica set, as well as *master* nodes in master/slave replication deployments.

`--oplog` (page 541) does not dump the oplog collection.

command line option!—repair

--repair

Runs a repair option in addition to dumping the database. The repair option attempts to repair a database that may be in an invalid state as a result of an improper shutdown or `mongod` (page 503) crash.

The `--repair` (page 542) option uses aggressive data-recovery algorithms that may produce a large amount of duplication.

command line option!—forceTableScan

--forceTableScan

Forces `mongodump` (page 537) to scan the data store directly: typically, `mongodump` (page 537) saves entries as they appear in the index of the `_id` field. Use `--forceTableScan` (page 567) to skip the index and scan the data directly. Typically there are two cases where this behavior is preferable to the default:

- 1.If you have key sizes over 800 bytes that would not be present in the `_id` index.
- 2.Your database uses a custom `_id` field.

When you run with `--forceTableScan` (page 567), `mongodump` (page 537) does not use `$snapshot` (page 482). As a result, the dump produced by `mongodump` (page 537) can reflect the state of the database at many different points in time.

Important: Use `--forceTableScan` (page 567) with extreme caution and consideration.

command line option!—dumpDbUsersAndRoles

--dumpDbUsersAndRoles

Includes user and role definitions when performing `mongodump` (page 537) on a specific database. This option applies only when you specify a database in the `--db` (page 547) option. MongoDB always includes user and role definitions when `mongodump` (page 537) applies to an entire instance and not just a specific database.

Use

See the <http://docs.mongodb.org/manual/tutorial/backup-with-mongodump> for a larger overview of `mongodump` (page 537) usage. Also see the *mongorestore* (page 543) document for an overview of the *mongorestore* (page 543), which provides the related inverse functionality.

The following command creates a dump file that contains only the collection named `collection` in the database named `test`. In this case the database is running on the local interface on port 27017:

```
mongodump --collection collection --db test
```

In the next example, `mongodump` (page 537) creates a backup of the database instance stored in the `/srv/mongodb` directory on the local machine. This requires that no `mongod` (page 503) instance is using the `/srv/mongodb` directory.

```
mongodump --dbpath /srv/mongodb
```

In the final example, `mongodump` (page 537) creates a database dump located at <http://docs.mongodb.org/manual/opt/backup/mongodump-2011-10-24>, from a database running on port 37017 on the host `mongodb1.example.net` and authenticating using the username `user` and the password `pass`, as follows:

```
mongodump --host mongodbl.example.net --port 37017 --username user --password pass --out /opt/backup/
```

mongorestore

Synopsis

The `mongorestore` (page 543) program writes data from a binary database dump created by `mongodump` (page 537) to a MongoDB instance. `mongorestore` (page 543) can create a new database or add data to an existing database.

`mongorestore` (page 543) can write data to either *mongod* or *mongos* (page 518) instances, in addition to writing directly to MongoDB data files without an active `mongod` (page 503).

Behavior

If you restore to an existing database, `mongorestore` (page 543) will only insert into the existing database, and does not perform updates of any kind. If existing documents have the same value `_id` field in the target database and collection, `mongorestore` (page 543) will *not* overwrite those documents.

Remember the following properties of `mongorestore` (page 543) behavior:

- `mongorestore` (page 543) recreates indexes recorded by `mongodump` (page 537).
- all operations are inserts, not updates.
- `mongorestore` (page 543) does not wait for a response from a `mongod` (page 503) to ensure that the MongoDB process has received or recorded the operation.

The `mongod` (page 503) will record any errors to its log that occur during a restore operation, but `mongorestore` (page 543) will not receive errors.

The data format used by `mongodump` (page 537) from version 2.2 or later is *incompatible* with earlier versions of `mongod` (page 503). Do not use recent versions of `mongodump` (page 537) to back up older data stores.

Required Access to Restore User Data

Changed in version 2.6.

To restore users and *user-defined roles* on a given database, you must have access to the `admin` database. MongoDB stores the user data and role definitions for all databases in the `admin` database.

Specifically, to restore users to a given database, you must have the `insert` *action* on the `admin` database's `admin.system.users` (page 601) collection. The `restore` role provides this privilege.

To restore user-defined roles to a database, you must have the `insert` action on the `admin` database's `admin.system.roles` (page 601) collection. The `restore` role provides this privilege.

Options

mongorestore

mongorestore

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongorestore` (page 543).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

--quiet

Runs the `mongorestore` (page 543) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `mongorestore` (page 543) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, -h

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongorestore` (page 543) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

`<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...`

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongorestore` (page 543) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has CAFile enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongorestore` (page 543) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongorestore` (page 543) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMode

--sslFIPSMode

New in version 2.6.

Directs the `mongorestore` (page 543) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMode` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, -u

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, -p

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongorestore` (page 543) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the `PLAIN` and `MONGODB-X509` authentication mechanisms.

Specifies the authentication mechanism the `mongorestore` (page 543) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use <code>PLAIN</code> for authenticating in-database users. <code>PLAIN</code> transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ²¹ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ²² .

command line option!-dbpath <path>

--dbpath <path>

Specifies the directory of the MongoDB data files. The `--dbpath` (page 546) option lets the `mongorestore` (page 543) attach directly to the local data files without going through a running `mongod` (page 503). When

¹⁹<http://www.mongodb.com/products/mongodb-enterprise>

²⁰<http://www.mongodb.com/products/mongodb-enterprise>

²¹<http://www.mongodb.com/products/mongodb-enterprise>

²²<http://www.mongodb.com/products/mongodb-enterprise>

run with `--dbpath` (page 546), the `mongorestore` (page 543) locks access to the data files. No `mongod` (page 503) can access the files while the `mongorestore` (page 543) process runs.

command line option!—directoryperdb

--directoryperdb

When used in conjunction with the corresponding option in `mongod` (page 503), allows the `mongorestore` (page 543) to access data from MongoDB instances that use an on-disk format where every database has a distinct directory. This option is only relevant when specifying the `--dbpath` (page 546) option.

command line option!—journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongorestore` (page 543) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!—db <database>, -d

--db <database>, -d

Specifies a database for `mongorestore` (page 543) to restore data *into*. If the database does not exist, `mongorestore` (page 543) creates the database. If you do not specify a <db>, `mongorestore` (page 543) creates new databases that correspond to the databases where data originated and data may be overwritten. Use this option to restore data into a MongoDB instance that already has data.

`--db` (page 547) does *not* control which *BSON* files `mongorestore` (page 543) restores. You must use the `mongorestore` (page 543) *path option* (page 548) to limit that restored data.

command line option!—collection <collection>, -c

--collection <collection>, -c

Specifies a single collection for `mongorestore` (page 543) to restore. If you do not specify `--collection` (page 547), `mongorestore` (page 543) takes the collection name from the input filename. If the input file has an extension, MongoDB omits the extension of the file from the collection name.

command line option!—objcheck

--objcheck

Forces `mongorestore` (page 543) to validate all requests from clients upon receipt to ensure that clients never insert invalid documents into the database. For objects with a high degree of sub-document nesting, `--objcheck` (page 581) can have a small impact on performance. You can set `--noobjcheck` (page 547) to disable object checking at run-time.

Changed in version 2.4: MongoDB enables `--objcheck` (page 581) by default, to prevent any client from inserting malformed or invalid BSON into a MongoDB database.

command line option!—noobjcheck

--noobjcheck

New in version 2.4.

Disables the default document validation that MongoDB performs on all incoming BSON documents.

command line option!—filter <JSON>

--filter <JSON>

Limits the documents that `mongorestore` (page 543) imports to only those documents that match the JSON document specified as '`<JSON>`'. Be sure to include the document in single quotes to avoid interaction with your system's shell environment. For an example of `--filter` (page 547), see *backup-restore-filter*.

command line option!—drop

--drop

Modifies the restoration procedure to drop every collection from the target database before restoring the collection from the dumped backup.

command line option!-oplogReplay

--oplogReplay

Replays the *oplog* after restoring the dump to ensure that the current state of the database reflects the point-in-time backup captured with the “*mongodump --oplog*” command. For an example of *--oplogReplay* (page 548), see *backup-restore-oplogreplay*.

command line option!-oplogLimit <timestamp>

--oplogLimit <timestamp>

New in version 2.2.

Prevents *mongorestore* (page 543) from applying *oplog* entries with timestamp newer than or equal to <timestamp>. Specify <timestamp> values in the form of <time_t>:<ordinal>, where <time_t> is the seconds since the UNIX epoch, and <ordinal> represents a counter of operations in the oplog that occurred in the specified second.

You must use *--oplogLimit* (page 548) in conjunction with the *--oplogReplay* (page 548) option.

command line option!-keepIndexVersion

--keepIndexVersion

Prevents *mongorestore* (page 543) from upgrading the index to the latest version during the restoration process.

command line option!-noIndexRestore

--noIndexRestore

New in version 2.2.

Prevents *mongorestore* (page 543) from restoring and building indexes as specified in the corresponding *mongodump* (page 537) output.

command line option!-noOptionsRestore

--noOptionsRestore

New in version 2.2.

Prevents *mongorestore* (page 543) from setting the collection options, such as those specified by the *collMod* (page 316) *database command*, on restored collections.

command line option!-w <number of replicas per write>

--w <number of replicas per write>

New in version 2.2.

Specifies the *write concern* for each write operation that *mongorestore* (page 543) writes to the target database. By default, *mongorestore* (page 543) does not wait for a response for *write acknowledgment*.

<path>

The final argument of the *mongorestore* (page 543) command is a directory path. This argument specifies the location of the database dump from which to restore.

Use

See <http://docs.mongodb.org/manual/tutorial/backup-with-mongodump> for a larger overview of *mongorestore* (page 543) usage. Also see the *mongodump* (page 536) document for an overview of the *mongodump* (page 537), which provides the related inverse functionality.

Consider the following example:

```
mongorestore --collection people --db accounts dump/accounts/people.bson
```

Here, `mongorestore` (page 543) reads the database dump in the `dump/` sub-directory of the current directory, and restores *only* the documents in the collection named `people` from the database named `accounts`. `mongorestore` (page 543) restores data to the instance running on the localhost interface on port 27017.

In the next example, `mongorestore` (page 543) restores a backup of the database instance located in `dump` to a database instance stored in the `/srv/mongodb` on the local machine. This requires that there are no active `mongod` (page 503) instances attached to `/srv/mongodb` data directory.

```
mongorestore --dbpath /srv/mongodb
```

In the final example, `mongorestore` (page 543) restores a database dump located at `http://docs.mongodb.org/manual/opt/backup/mongodump-2011-10-24`, to a database running on port 37017 on the host `mongodb1.example.net`. The `mongorestore` (page 543) command authenticates to the MongoDB instance using the username `user` and the password `pass`, as follows:

```
mongorestore --host mongodb1.example.net --port 37017 --username user --password pass /opt/backup/mor
```

bsondump

Synopsis

The `bsondump` (page 549) converts *BSON* files into human-readable formats, including *JSON*. For example, `bsondump` (page 549) is useful for reading the output files generated by `mongodump` (page 537).

Important: `bsondump` (page 549) is a diagnostic tool for inspecting BSON files, not a tool for data ingestion or other application use.

Options

bsondump

bsondump

command line option!-help, -h

--help, -h

Returns information on the options and use of `bsondump` (page 549).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

--quiet

Runs the `bsondump` (page 549) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity

- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `bsondump` (page 549) release number.

command line option!-objcheck

--objcheck

Validates each *BSON* object before outputting it in *JSON* format. By default, `bsondump` (page 549) enables `--objcheck` (page 581). For objects with a high degree of sub-document nesting, `--objcheck` (page 581) can have a small impact on performance. You can set `--noobjcheck` (page 547) to disable object checking.

Changed in version 2.4: MongoDB enables `--objcheck` (page 581) by default, to prevent any client from inserting malformed or invalid BSON into a MongoDB database.

command line option!-noobjcheck

--noobjcheck

New in version 2.4.

Disables the default document validation that MongoDB performs on all incoming BSON documents.

command line option!-filter <JSON>

--filter <JSON>

Limits the documents that `bsondump` (page 549) exports to only those documents that match the *JSON document* specified as '`<JSON>`'. Be sure to include the document in single quotes to avoid interaction with your system's shell environment.

command line option!-type <=json|=debug>

--type <=json|=debug>

Changes the operation of `bsondump` (page 549) from outputting "*JSON*" (the default) to a debugging format.

<bsonFilename>

The final argument to `bsondump` (page 549) is a document containing *BSON*. This data is typically generated by `bsondump` (page 549) or by MongoDB in a *rollback* operation.

Use

By default, `bsondump` (page 549) outputs data to standard output. To create corresponding *JSON* files, you will need to use the shell redirect. See the following command:

```
bsondump collection.bson > collection.json
```

Use the following command (at the system shell) to produce debugging output for a *BSON* file:

```
bsondump --type=debug collection.bson
```

mongooplog

New in version 2.2.

Synopsis

`mongooplog` (page 551) is a simple tool that polls operations from the *replication oplog* of a remote server, and applies them to the local server. This capability supports certain classes of real-time migrations that require that the source server remain online and in operation throughout the migration process.

Typically this command will take the following form:

```
mongooplog --from mongodb0.example.net --host mongodb1.example.net
```

This command copies oplog entries from the `mongod` (page 503) instance running on the host `mongodb0.example.net` and duplicates operations to the host `mongodb1.example.net`. If you do not need to keep the `--from` host running during the migration, consider using `mongodump` (page 537) and `mongorestore` (page 543) or another backup operation, which may be better suited to your operation.

Note: If the `mongod` (page 503) instance specified by the `--from` argument is running with authentication, then `mongooplog` (page 551) will not be able to copy oplog entries.

See also:

`mongodump` (page 537), `mongorestore` (page 543), <http://docs.mongodb.org/manualcore/backups>, <http://docs.mongodb.org/manualcore/replica-set-oplog>.

Options

`mongooplog`

`mongooplog`

command line option!-help, -h

`--help, -h`

Returns information on the options and use of `mongooplog` (page 551).

command line option!-verbose, -v

`--verbose, -v`

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

`--quiet`

Runs the `mongooplog` (page 551) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

`--version`

Returns the `mongooplog` (page 551) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, **-h**

Specifies a resolvable hostname for the `mongod` (page 503) instance to which `mongooplog` (page 551) will apply *oplog* operations retrieved from the server specified by the `--from` option.

By default `mongooplog` (page 551) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

```
<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...
```

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!`--port`

--port

Specifies the port number of the `mongod` (page 503) instance where `mongooplog` (page 551) will apply *oplog* entries. Specify this option only if the MongoDB instance to connect to is not running on the standard port of 27017. You may also specify a port number using the `--host` command.

command line option!`--ipv6`

--ipv6

Enables IPv6 support and allows the `mongooplog` (page 551) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!`--ssl`

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslCAFile` <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the root certificate chain from the Certificate Authority. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslPEMKeyFile` <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains both the SSL certificate and key. Specify the file name of the `.pem` file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has `CAFile` enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!`--sslPEMKeyPassword` <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongooplog` (page 551) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongooplog` (page 551) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.6.

Directs the `mongooplog` (page 551) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, -u

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, -p

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongooplog` (page 551) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the `PLAIN` and `MONGODB-X509` authentication mechanisms.

Specifies the authentication mechanism the `mongooplog` (page 551) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use <code>PLAIN</code> for authenticating in-database users. <code>PLAIN</code> transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ²⁵ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ²⁶ .

command line option!-dbpath <path>

--dbpath <path>

Specifies a directory, containing MongoDB data files, to which `mongooplog` (page 551) will apply operations from the `oplog` of the database specified with the `--from` option.

When used, the `--dbpath` (page 546) option enables `mongo` (page 527) to attach directly to local data files and write data without a running `mongod` (page 503) instance.

To run with `--dbpath` (page 546), `mongooplog` (page 551) needs to restrict access to the data directory: as a result, no `mongod` (page 503) can be access the same path while the process runs.

command line option!-directoryperdb

--directoryperdb

When used in conjunction with the corresponding option in `mongod` (page 503), allows the `mongooplog` (page 551) to access data from MongoDB instances that use an on-disk format where every database has a distinct directory. This option is only relevant when specifying the `--dbpath` (page 546) option.

command line option!-journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongooplog` (page 551) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!-db <database>, -d

--db <database>, **-d**

Specifies the name of the database on which to run the `mongooplog` (page 551).

²³<http://www.mongodb.com/products/mongodb-enterprise>

²⁴<http://www.mongodb.com/products/mongodb-enterprise>

²⁵<http://www.mongodb.com/products/mongodb-enterprise>

²⁶<http://www.mongodb.com/products/mongodb-enterprise>

command line option!-collection <collection>, -c

--collection <collection>, -c
Specifies the collection to export.

command line option!-seconds <number>, -s

--seconds <number>, -s
Specify a number of seconds of operations for [mongooplog](#) (page 551) to pull from the *remote host*. Unless specified the default value is 86400 seconds, or 24 hours.

command line option!-from <host[:port]>

--from <host[:port]>
Specify the host for [mongooplog](#) (page 551) to retrieve *oplog* operations from. [mongooplog](#) (page 551) *requires* this option.

Unless you specify the **--host** option, [mongooplog](#) (page 551) will apply the operations collected with this option to the oplog of the [mongod](#) (page 503) instance running on the localhost interface connected to port 27017.

command line option!-oplogns <namespace>

--oplogns <namespace>
Specify a namespace in the **--from** host where the oplog resides. The default value is `local.oplog.rs`, which is the where *replica set* members store their operation log. However, if you've copied *oplog* entries into another database or collection, use this option to copy oplog entries stored in another location. Namespaces take the form of `[database].[collection]`.

Use

Consider the following prototype [mongooplog](#) (page 551) command:

```
mongooplog --from mongodb0.example.net --host mongodb1.example.net
```

Here, entries from the *oplog* of the [mongod](#) (page 503) running on port 27017. This only pull entries from the last 24 hours.

Use the **--seconds** argument to capture a greater or smaller amount of time. Consider the following example:

```
mongooplog --from mongodb0.example.net --seconds 172800
```

In this operation, [mongooplog](#) (page 551) captures 2 full days of operations. To migrate 12 hours of *oplog* entries, use the following form:

```
mongooplog --from mongodb0.example.net --seconds 43200
```

For the previous two examples, [mongooplog](#) (page 551) migrates entries to the [mongod](#) (page 503) process running on the localhost interface connected to the 27017 port. [mongooplog](#) (page 551) can also operate directly on MongoDB's data files if no [mongod](#) (page 503) is running on the *target* host. Consider the following example:

```
mongooplog --from mongodb0.example.net --dbpath /srv/mongodb --journal
```

Here, [mongooplog](#) (page 551) imports *oplog* operations from the [mongod](#) (page 503) host connected to port 27017. This migrates operations to the MongoDB data files stored in the `/srv/mongodb` directory. Additionally [mongooplog](#) (page 551) will use the durability *journal* to ensure that the data files remain valid.

4.1.4 Data Import and Export Tools

`mongoimport` (page 556) provides a method for taking data in *JSON*, *CSV*, or *TSV* and importing it into a `mongod` (page 503) instance. `mongoexport` (page 562) provides a method to export data from a `mongod` (page 503) instance into JSON, CSV, or TSV.

Note: The conversion between BSON and other formats lacks full type fidelity. Therefore you cannot use `mongoimport` (page 556) and `mongoexport` (page 562) for round-trip import and export operations.

`mongoimport`

Synopsis

The `mongoimport` (page 556) tool provides a route to import content from a JSON, CSV, or TSV export created by `mongoexport` (page 562), or potentially, another third-party export tool. See the <http://docs.mongodb.org/manualcore/import-export> document for a more in depth usage overview, and the `mongoexport` (page 562) document for more information regarding `mongoexport` (page 562), which provides the inverse “exporting” capability.

Considerations

Do not use `mongoimport` (page 556) and `mongoexport` (page 562) for full instance, production backups because they will not reliably capture data type information. Use `mongodump` (page 537) and `mongorestore` (page 543) as described in <http://docs.mongodb.org/manualcore/backups> for this kind of functionality.

Options

`mongoimport`

`mongoimport`

command line option!-help, -h

`--help, -h`

Returns information on the options and use of `mongoimport` (page 556).

command line option!-verbose, -v

`--verbose, -v`

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

`--quiet`

Runs the `mongoimport` (page 556) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `mongoimport` (page 556) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, -h

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongoimport` (page 556) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

```
<replica_set_name>/<hostname1><:port>, <hostname2><:port>, ...
```

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongoimport` (page 556) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the root certificate chain from the Certificate Authority. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains both the SSL certificate and key. Specify the file name of the `.pem` file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has CAFile enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongoimport` (page 556) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongoimport` (page 556) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the Certificate Revocation List. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.6.

Directs the `mongoimport` (page 556) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, **-u**

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, **-p**

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongoimport` (page 556) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the PLAIN and MONGODB-X509 authentication mechanisms.

Specifies the authentication mechanism the `mongoimport` (page 556) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use PLAIN for authenticating in-database users. PLAIN transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ²⁹ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ³⁰ .

command line option!-dbpath <path>

--dbpath <path>

Specifies the directory of the MongoDB data files. The `--dbpath` (page 546) option lets the `mongoimport` (page 556) attach directly to the local data files without going through a running `mongod` (page 503). When run with `--dbpath` (page 546), the `mongoimport` (page 556) locks access to the data files. No `mongod` (page 503) can access the files while the `mongoimport` (page 556) process runs.

command line option!-directoryperdb

--directoryperdb

When used in conjunction with the corresponding option in `mongod` (page 503), allows the `mongoimport` (page 556) to access data from MongoDB instances that use an on-disk format where every database has a distinct directory. This option is only relevant when specifying the `--dbpath` (page 546) option.

command line option!-journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongoimport` (page 556) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!-db <database>, -d

--db <database>, -d

Specifies the name of the database on which to run the `mongoimport` (page 556).

command line option!-collection <collection>, -c

²⁷<http://www.mongodb.com/products/mongodb-enterprise>

²⁸<http://www.mongodb.com/products/mongodb-enterprise>

²⁹<http://www.mongodb.com/products/mongodb-enterprise>

³⁰<http://www.mongodb.com/products/mongodb-enterprise>

--collection <collection>, **-c**
Specifies the collection to import.

New in version 2.6: If you do not specify **--collection** (page 547), **mongoimport** (page 556) takes the collection name from the input filename. MongoDB omits the extension of the file from the collection name, if the input file has an extension.

command line option!-fields <field1[,field2]>, **-f**

--fields <field1[,field2]>, **-f**

Specify a comma separated list of field names when importing *csv* or *tsv* files that do not have field names in the first (i.e. header) line of the file.

command line option!-fieldFile <filename>

--fieldFile <filename>

As an alternative to **--fields** (page 560), the **--fieldFile** (page 560) option allows you to specify a file that holds a list of field names if your *csv* or *tsv* file does not include field names in the first line of the file (i.e. header). Place one field per line.

command line option!-ignoreBlanks

--ignoreBlanks

Ignores empty fields in *csv* and *tsv* exports. If not specified, **mongoimport** (page 556) creates fields without values in imported documents.

command line option!-type <json|csv|tsv>

--type <json|csv|tsv>

Specifies the file type to import. The default format is *JSON*, but it's possible to import *csv* and *tsv* files.

command line option!-file <filename>

--file <filename>

Specifies the location and name of a file containing the data to import. If you do not specify a file, **mongoimport** (page 556) reads data from standard input (e.g. "stdin").

command line option!-drop

--drop

Modifies the import process so that the target instance drops every collection before importing the collection from the input.

command line option!-headerline

--headerline

If using **--type csv** or **--type tsv**, uses the first line as field names. Otherwise, **mongoimport** (page 556) will import the first line as a distinct document.

command line option!-upsert

--upsert

Modifies the import process to update existing objects in the database if they match an imported object, while inserting all other objects.

If you do not specify a field or fields using the **--upsertFields** (page 560) **mongoimport** (page 556) will upsert on the basis of the `_id` field.

command line option!-upsertFields <field1[,field2]>

--upsertFields <field1[,field2]>

Specifies a list of fields for the query portion of the *upsert*. Use this option if the `_id` fields in the existing documents don't match the field in the document, but another field or field combination can uniquely identify documents as a basis for performing upsert operations.

To ensure adequate performance, indexes should exist for this field or fields.

command line option!-stopOnError

--stopOnError

New in version 2.2.

Forces `mongoimport` (page 556) to halt the import operation at the first error rather than continuing the operation despite errors.

command line option!-jsonArray

--jsonArray

Accepts the import of data expressed with multiple MongoDB documents within a single *JSON* array.

Used in conjunction with `mongoexport --jsonArray` to import data written as a single *JSON* array. Limited to imports of 16 MB or smaller.

Use

In this example, `mongoimport` (page 556) imports the *csv* formatted data in the `http://docs.mongodb.org/manualopt/backups/contacts.csv` into the collection `contacts` in the `users` database on the MongoDB instance running on the localhost port numbered 27017. `mongoimport` (page 556) determines the name of files using the first line in the CSV file, because of the `--headerline`:

```
mongoimport --db users --collection contacts --type csv --headerline --file /opt/backups/contacts.csv
```

Since `mongoimport` (page 556) uses the input file name, without the extension, as the collection name if `-c` or `--collection` is unspecified. The following following example is equivalent:

```
mongoimport --db users --type csv --headerline --file /opt/backups/contacts.csv
```

In the following example, `mongoimport` (page 556) imports the data in the *JSON* formatted file `contacts.json` into the collection `contacts` on the MongoDB instance running on the localhost port number 27017.

```
mongoimport --collection contacts --file contacts.json
```

In the next example, `mongoimport` (page 556) takes data passed to it on standard input (i.e. with a `|` pipe.) and imports it into the MongoDB datafiles located at `/srv/mongodb/`. if the import process encounters an error, the `mongoimport` (page 556) will halt because of the `--stopOnError` option.

```
mongoimport --db sales --collection contacts --stopOnError --dbpath /srv/mongodb/
```

In the final example, `mongoimport` (page 556) imports data from the file `http://docs.mongodb.org/manualopt/backups/mdb1-examplenet.json` into the collection `contacts` within the database `marketing` on a remote MongoDB database. This `mongoimport` (page 556) accesses the `mongod` (page 503) instance running on the host `mongodb1.example.net` over port 37017, which requires the username `user` and the password `pass`.

```
mongoimport --host mongodb1.example.net --port 37017 --username user --password pass --collection con
```

Type Fidelity

Warning: `mongoimport` (page 556) and `mongoexport` (page 562) do not reliably preserve all rich *BSON* data types because *JSON* can only represent a subset of the types supported by *BSON*. As a result, data exported or imported with these tools may lose some measure of fidelity. See <http://docs.mongodb.org/manualreference/mongodb-extended-json> for more information.

JSON can only represent a subset of the types supported by *BSON*. To preserve type information, `mongoimport` (page 556) accepts strict mode representation for certain ptypes.

For example, to preserve type information for *BSON* types `data_date` and `data_numberlong` during `mongoimport` (page 556), the data should be in strict mode representation, as in the following:

```
{ "_id" : 1, "volume" : { "$numberLong" : "2980000" }, "date" : { "$date" : "2014-03-13T13:47:42.483Z" }}
```

For the `data_numberlong` type, `mongoimport` (page 556) converts into a float during the import.

See <http://docs.mongodb.org/manualreference/mongodb-extended-json> for a complete list of these types and the representations used.

`mongoexport`

Synopsis

`mongoexport` (page 562) is a utility that produces a *JSON* or *CSV* export of data stored in a MongoDB instance. See the <http://docs.mongodb.org/manualcore/import-export> document for a more in depth usage overview, and the `mongoimport` (page 556) document for more information regarding the `mongoimport` (page 556) utility, which provides the inverse “importing” capability.

Considerations

Do not use `mongoimport` (page 556) and `mongoexport` (page 562) for full-scale production backups because they may not reliably capture data type information. Use `mongodump` (page 537) and `mongorestore` (page 543) as described in <http://docs.mongodb.org/manualcore/backups> for this kind of functionality.

Options

`mongoexport`

`mongoexport`

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongoexport` (page 562).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the -v form by including the option multiple times, (e.g. -vvvvvv.)

command line option!-quiet

--quiet

Runs the `mongoexport` (page 562) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!—version

--version

Returns the `mongoexport` (page 562) release number.

command line option!—host <hostname><:port>, -h

--host <hostname><:port>, -h

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongoexport` (page 562) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

<replica_set_name>/<hostname1><:port>, <hostname2><:port>, ...

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!—port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!—ipv6

--ipv6

Enables IPv6 support and allows the `mongoexport` (page 562) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!—ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!—sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the root certificate chain from the Certificate Authority. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has `CAFile` enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-crypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongoexport` (page 562) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongoexport` (page 562) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.6.

Directs the `mongoexport` (page 562) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, **-u**

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the **--password** and **--authenticationDatabase** options.

command line option! **--password** <password>, **-p**

--password <password>, **-p**

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the **--username** and **--authenticationDatabase** options.

command line option! **--authenticationDatabase** <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the [mongoexport](#) (page 562) assumes that the database specified as the argument to the **--db** (page 547) option holds the user's credentials.

command line option! **--authenticationMechanism** <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the PLAIN and MONGODB-X509 authentication mechanisms.

Specifies the authentication mechanism the [mongoexport](#) (page 562) instance uses to authenticate to the [mongod](#) (page 503) or [mongos](#) (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use PLAIN for authenticating in-database users. PLAIN transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ³³ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ³⁴ .

command line option! **--dbpath** <path>

--dbpath <path>

Specifies the directory of the MongoDB data files. The **--dbpath** (page 546) option lets the [mongoexport](#) (page 562) attach directly to the local data files without going through a running [mongod](#) (page 503). When run with **--dbpath** (page 546), the [mongoexport](#) (page 562) locks access to the data files. No [mongod](#) (page 503) can access the files while the [mongoexport](#) (page 562) process runs.

command line option! **--directoryperdb**

--directoryperdb

When used in conjunction with the corresponding option in [mongod](#) (page 503), allows [mongoexport](#) (page 562) to export data from MongoDB instances that have every database's files saved in discrete directories on the disk. This option is only relevant when specifying the **--dbpath** (page 546) option.

command line option! **--journal**

³¹<http://www.mongodb.com/products/mongodb-enterprise>

³²<http://www.mongodb.com/products/mongodb-enterprise>

³³<http://www.mongodb.com/products/mongodb-enterprise>

³⁴<http://www.mongodb.com/products/mongodb-enterprise>

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongoexport` (page 562) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!-db <database>, -d

--db <database>, -d

Specifies the name of the database on which to run the `mongoexport` (page 562).

command line option!-collection <collection>, -c

--collection <collection>, -c

Specifies the collection to export.

command line option!-fields <field1[,field2]>, -f

--fields <field1[,field2]>, -f

Specifies a field or fields to *include* in the export. Use a comma separated list of fields to specify multiple fields.

For `--csv` output formats, `mongoexport` (page 562) includes only the specified field(s), and the specified field(s) can be a field within a sub-document.

For *JSON* output formats, `mongoexport` (page 562) includes only the specified field(s) **and** the `_id` field, and if the specified field(s) is a field within a sub-document, the `mongoexport` (page 562) includes the sub-document with all its fields, not just the specified field within the document.

command line option!-fieldFile <filename>

--fieldFile <filename>

An alternative to `--fields`. The `--fieldFile` (page 560) option allows you to specify in a file the field or fields to *include* in the export and is **only valid** with the `--csv` option. The file must have only one field per line, and the line(s) must end with the LF character (0x0A).

`mongoexport` (page 562) includes only the specified field(s). The specified field(s) can be a field within a sub-document.

command line option!-query <JSON>, -q

--query <JSON>, -q

Provides a *JSON document* as a query that optionally limits the documents returned in the export. Specify JSON in strict format.

For example, given a collection named `records` in the database `test` with the following documents:

```
{ "_id" : ObjectId("51f0188846a64a1ed98fde7c"), "a" : 1 }
{ "_id" : ObjectId("520e61b0c6646578e3661b59"), "a" : 1, "b" : 2 }
{ "_id" : ObjectId("520e642bb7fa4ea22d6b1871"), "a" : 2, "b" : 3, "c" : 5 }
{ "_id" : ObjectId("520e6431b7fa4ea22d6b1872"), "a" : 3, "b" : 3, "c" : 6 }
{ "_id" : ObjectId("520e6445b7fa4ea22d6b1873"), "a" : 5, "b" : 6, "c" : 8 }
```

The following `mongoexport` (page 562) uses the `-q` (page ??) option to export only the documents with the field `a` greater than or equal to (`$gte` (page 374)) to 3:

```
mongoexport -d test -c records -q "{ a: { $gte: 3 } }" --out exportdir/myRecords.json
```

The resulting file contains the following documents:

```
{ "_id" : { "$oid" : "520e6431b7fa4ea22d6b1872" }, "a" : 3, "b" : 3, "c" : 6 }
{ "_id" : { "$oid" : "520e6445b7fa4ea22d6b1873" }, "a" : 5, "b" : 6, "c" : 8 }
```

You can sort the results with the `--sort` (page 567) option to `mongoexport` (page 562).

command line option!-csv

--csv

Changes the export format to a comma-separated-values (CSV) format. By default `mongoexport` (page 562) writes data using one *JSON* document for every MongoDB document.

If you specify `--csv` (page 566), then you must also use either the `--fields` (page 560) or the `--fieldFile` (page 560) option to declare the fields to export from the collection.

command line option!-out <file>, -o

--out <file>, -o

Specifies a file to write the export to. If you do not specify a file name, the `mongoexport` (page 562) writes data to standard output (e.g. `stdout`).

command line option!-jsonArray

--jsonArray

Modifies the output of `mongoexport` (page 562) to write the entire contents of the export as a single *JSON* array. By default `mongoexport` (page 562) writes data using one JSON document for every MongoDB document.

command line option!-slaveOk, -k

--slaveOk, -k

Allows `mongoexport` (page 562) to read data from secondary or slave nodes when using `mongoexport` (page 562) with a replica set. This option is only available if connected to a `mongod` (page 503) or `mongos` (page 518) and is not available when used with the “`mongoexport --dbpath`” option.

This is the default behavior.

command line option!-forceTableScan

--forceTableScan

New in version 2.2.

Forces `mongoexport` (page 562) to scan the data store directly: typically, `mongoexport` (page 562) saves entries as they appear in the index of the `_id` field. Use `--forceTableScan` (page 567) to skip the index and scan the data directly. Typically there are two cases where this behavior is preferable to the default:

- 1.If you have key sizes over 800 bytes that would not be present in the `_id` index.
- 2.Your database uses a custom `_id` field.

When you run with `--forceTableScan` (page 567), `mongoexport` (page 562) does not use `$snapshot` (page 482). As a result, the export produced by `mongoexport` (page 562) can reflect the state of the database at many different points in time.

Warning: Use `--forceTableScan` (page 567) with extreme caution and consideration.

command line option!-skip <number>

--skip <number>

Use `--skip` (page 567) to control where `mongoexport` (page 562) begins exporting documents. See `skip()` (page 92) for information about the underlying operation.

command line option!-limit <number>

--limit <number>

Specifies a maximum number of documents to include in the export. See `limit()` (page 86) for information about the underlying operation.

command line option!-sort <JSON>

--sort <JSON>

Specifies an ordering for exported results. If an index does **not** exist that can support the sort operation, the results must be *less than 32 megabytes*.

Use `--sort` (page 567) conjunction with `--skip` (page 567) and `--limit` (page 567) to limit number of exported documents.

```
mongoexport -d test -c records --sort '{a: 1}' --limit 100 --out export.0.json
mongoexport -d test -c records --sort '{a: 1}' --limit 100 --skip 100 --out export.1.json
mongoexport -d test -c records --sort '{a: 1}' --limit 100 --skip 200 --out export.2.json
```

See `sort()` (page 93) for information about the underlying operation.

Use

Export in CSV Format In the following example, `mongoexport` (page 562) exports the collection `contacts` from the `users` database from the `mongod` (page 503) instance running on the local-host port number 27017. This command writes the export data in *CSV* format into a file located at <http://docs.mongodb.org/manual/opt/backups/contacts.csv>. The `fields.txt` file contains a line-separated list of fields to export.

```
mongoexport --db users --collection contacts --csv --fieldFile fields.txt --out /opt/backups/contacts.csv
```

Export in JSON Format The next example creates an export of the collection `contacts` from the MongoDB instance running on the localhost port number 27017, with journaling explicitly enabled. This writes the export to the `contacts.json` file in *JSON* format.

```
mongoexport --db sales --collection contacts --out contacts.json --journal
```

Export Collection Directly From Data Files The following example exports the collection `contacts` from the `sales` database located in the MongoDB data files located at `/srv/mongodb/`. This operation writes the export to standard output in *JSON* format.

```
mongoexport --db sales --collection contacts --dbpath /srv/mongodb/
```

Warning: The above example will only succeed if there is no `mongod` (page 503) connected to the data files located in the `/srv/mongodb/` directory.

Export from Remote Host Running with Authentication The following example exports the collection `contacts` from the database `marketing`. This data resides on the MongoDB instance located on the host `mongodb1.example.net` running on port 37017, which requires the username `user` and the password `pass`.

```
mongoexport --host mongodb1.example.net --port 37017 --username user --password pass --collection contacts
```

Type Fidelity

Warning: `mongoimport` (page 556) and `mongoexport` (page 562) do not reliably preserve all rich *BSON* data types because *JSON* can only represent a subset of the types supported by *BSON*. As a result, data exported or imported with these tools may lose some measure of fidelity. See <http://docs.mongodb.org/manualreference/mongodb-extended-json> for more information.

JSON can only represent a subset of the types supported by BSON. To preserve type information, `mongoexport` (page 562) uses the `strict mode` representation for certain types.

For example, the following insert operation in the `mongo` (page 527) shell uses the `mongoShell` mode representation for the BSON types `data_date` and `data_numberlong`:

```
use test
db.traffic.insert( { _id: 1, volume: NumberLong(2980000), date: new Date() } )
```

Use `mongoexport` (page 562) to export the data:

```
mongoexport --db test --collection traffic --out traffic.json
```

The exported data is in `strict mode` representation to preserve type information:

```
{ "_id" : 1, "volume" : { "$numberLong" : "2980000" }, "date" : { "$date" : "2014-03-13T13:47:42.483"
```

See <http://docs.mongodb.org/manualreference/mongodb-extended-json> for a complete list of these types and the representations used.

4.1.5 Diagnostic Tools

`mongostat` (page 570), `mongotop` (page 576), and `mongosniff` (page 581) provide diagnostic information related to the current operation of a `mongod` (page 503) instance.

Note: Because `mongosniff` (page 581) depends on *libpcap*, most distributions of MongoDB do *not* include `mongosniff` (page 581).

`mongostat`

Synopsis

The `mongostat` (page 570) utility provides a quick overview of the status of a currently running `mongod` (page 503) or `mongos` (page 518) instance. `mongostat` (page 570) is functionally similar to the UNIX/Linux file system utility `vmstat`, but provides data regarding `mongod` (page 503) and `mongos` (page 518) instances.

See also:

For more information about monitoring MongoDB, see <http://docs.mongodb.org/manualadministration/monitoring>.

For more background on various other MongoDB status outputs see:

- *serverStatus* (page 346)
- *replSetGetStatus* (page 273)
- *dbStats* (page 331)
- *collStats* (page 325)

For an additional utility that provides MongoDB metrics see `mongotop` (page 576).

`mongostat` (page 570) connects to the `mongod` (page 503) instance running on the local host interface on TCP port 27017; however, `mongostat` (page 570) can connect to any accessible remote `mongod` (page 503) instance.

Options

mongostat

mongostat

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongostat` (page 570).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-version

--version

Returns the `mongostat` (page 570) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, -h

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongostat` (page 570) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

`<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...`

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongostat` (page 570) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has `CAFile` enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongostat` (page 570) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongostat` (page 570) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.6.

Directs the `mongostat` (page 570) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, **-u**

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the **--password** and **--authenticationDatabase** options.

command line option!-password <password>, -p

--password <password>, **-p**

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the **--username** and **--authenticationDatabase** options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the [mongostat](#) (page 570) assumes that the database specified as the argument to the **--db** (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the `PLAIN` and `MONGODB-X509` authentication mechanisms.

Specifies the authentication mechanism the [mongostat](#) (page 570) instance uses to authenticate to the [mongod](#) (page 503) or [mongos](#) (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use <code>PLAIN</code> for authenticating in-database users. <code>PLAIN</code> transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ³⁷ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ³⁸ .

command line option!-noheaders

--noheaders

Disables the output of column or field names.

command line option!-rowcount <number>, -n

--rowcount <number>, **-n**

Controls the number of rows to output. Use in conjunction with the `sleeptime` argument to control the duration of a [mongostat](#) (page 570) operation.

³⁵<http://www.mongodb.com/products/mongodb-enterprise>

³⁶<http://www.mongodb.com/products/mongodb-enterprise>

³⁷<http://www.mongodb.com/products/mongodb-enterprise>

³⁸<http://www.mongodb.com/products/mongodb-enterprise>

Unless `--rowcount` (page 572) is specified, `mongostat` (page 570) will return an infinite number of rows (e.g. value of 0.)

command line option!-http

--http

Configures `mongostat` (page 570) to collect data using the HTTP interface rather than a raw database connection.

command line option!-discover

--discover

Discovers and reports on statistics from all members of a *replica set* or *sharded cluster*. When connected to any member of a replica set, `--discover` (page 573) all non-*hidden members* of the replica set. When connected to a `mongos` (page 518), `mongostat` (page 570) will return data from all *shards* in the cluster. If a replica set provides a shard in the sharded cluster, `mongostat` (page 570) will report on non-hidden members of that replica set.

The `mongostat --host` option is not required but potentially useful in this case.

Changed in version 2.6: When running with `--discover` (page 573), `mongostat` (page 570) now respects `:option:-rowcount`.

command line option!-all

--all

Configures `mongostat` (page 570) to return all optional *fields* (page 573).

<sleeptime>

The final argument is the length of time, in seconds, that `mongostat` (page 570) waits in between calls. By default `mongostat` (page 570) returns one call every second.

`mongostat` (page 570) returns values that reflect the operations over a 1 second period. For values of `<sleeptime>` greater than 1, `mongostat` (page 570) averages data to reflect average operations per second.

Fields

`mongostat` (page 570) returns values that reflect the operations over a 1 second period. When `mongostat <sleeptime>` has a value greater than 1, `mongostat` (page 570) averages the statistics to reflect average operations per second.

`mongostat` (page 570) outputs the following fields:

inserts

The number of objects inserted into the database per second. If followed by an asterisk (e.g. *), the datum refers to a replicated operation.

query

The number of query operations per second.

update

The number of update operations per second.

delete

The number of delete operations per second.

getmore

The number of get more (i.e. cursor batch) operations per second.

command

The number of commands per second. On *slave* and *secondary* systems, `mongostat` (page 570) presents two values separated by a pipe character (e.g. |), in the form of `local|replicated` commands.

flushes

The number of *fsync* operations per second.

mapped

The total amount of data mapped in megabytes. This is the total data size at the time of the last `mongostat` (page 570) call.

size

The amount of virtual memory in megabytes used by the process at the time of the last `mongostat` (page 570) call.

non-mapped

The total amount of virtual memory excluding all mapped memory at the time of the last `mongostat` (page 570) call.

res

The amount of resident memory in megabytes used by the process at the time of the last `mongostat` (page 570) call.

faults

Changed in version 2.1.

The number of page faults per second.

Before version 2.1 this value was only provided for MongoDB instances running on Linux hosts.

locked

The percent of time in a global write lock.

Changed in version 2.2: The `locked_db` field replaces the `locked %` field to more appropriate data regarding the database specific locks in version 2.2.

locked_db

New in version 2.2.

The percent of time in the per-database context-specific lock. `mongostat` (page 570) will report the database that has spent the most time since the last `mongostat` (page 570) call with a write lock.

This value represents the amount of time that the listed database spent in a locked state *combined* with the time that the `mongod` (page 503) spent in the global lock. Because of this, and the sampling method, you may see some values greater than 100%.

idx miss

The percent of index access attempts that required a page fault to load a btree node. This is a sampled value.

qr

The length of the queue of clients waiting to read data from the MongoDB instance.

qw

The length of the queue of clients waiting to write data from the MongoDB instance.

ar

The number of active clients performing read operations.

aw

The number of active clients performing write operations.

netIn

The amount of network traffic, in *bytes*, received by the MongoDB instance.

This includes traffic from `mongostat` (page 570) itself.

netOut

The amount of network traffic, in *bytes*, sent by the MongoDB instance.

This includes traffic from `mongostat` (page 570) itself.

conn

The total number of open connections.

set

The name, if applicable, of the replica set.

repl

The replication status of the member.

Value	Replication Type
M	<i>master</i>
SEC	<i>secondary</i>
REC	recovering
UNK	unknown
SLV	<i>slave</i>
RTR	mongos process (“router”)

Usage

In the first example, `mongostat` (page 570) will return data every second for 20 seconds. `mongostat` (page 570) collects data from the `mongod` (page 503) instance running on the localhost interface on port 27017. All of the following invocations produce identical behavior:

```
mongostat --rowcount 20 1
mongostat --rowcount 20
mongostat -n 20 1
mongostat -n 20
```

In the next example, `mongostat` (page 570) returns data every 5 minutes (or 300 seconds) for as long as the program runs. `mongostat` (page 570) collects data from the `mongod` (page 503) instance running on the localhost interface on port 27017. Both of the following invocations produce identical behavior.

```
mongostat --rowcount 0 300
mongostat -n 0 300
mongostat 300
```

In the following example, `mongostat` (page 570) returns data every 5 minutes for an hour (12 times.) `mongostat` (page 570) collects data from the `mongod` (page 503) instance running on the localhost interface on port 27017. Both of the following invocations produce identical behavior.

```
mongostat --rowcount 12 300
mongostat -n 12 300
```

In many cases, using the `--discover` will help provide a more complete snapshot of the state of an entire group of machines. If a `mongos` (page 518) process connected to a *sharded cluster* is running on port 27017 of the local machine, you can use the following form to return statistics from all members of the cluster:

```
mongostat --discover
```

mongotop

Synopsis

`mongotop` (page 576) provides a method to track the amount of time a MongoDB instance spends reading and writing data. `mongotop` (page 576) provides statistics on a per-collection level. By default, `mongotop` (page 576) returns values every second.

See also:

For more information about monitoring MongoDB, see <http://docs.mongodb.org/manualadministration/monitoring>.

For additional background on various other MongoDB status outputs see:

- [serverStatus](#) (page 346)
- [replSetGetStatus](#) (page 273)
- [dbStats](#) (page 331)
- [collStats](#) (page 325)

For an additional utility that provides MongoDB metrics see [mongostat](#) (page 569).

Options

mongotop

mongotop

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongotop` (page 576).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

--quiet

Runs the `mongotop` (page 576) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `mongotop` (page 576) release number.

command line option!-host <hostname><:port>, -h

--host <hostname><:port>, **-h**

Default: localhost:27017

Specifies a resolvable hostname for the `mongod` (page 503) to which to connect. By default, the `mongotop` (page 576) attempts to connect to a MongoDB instance running on the localhost on port number 27017.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

```
<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...
```

You can always connect directly to a single MongoDB instance by specifying the host and port number directly.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongotop` (page 576) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the root certificate chain from the Certificate Authority. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains both the SSL certificate and key. Specify the file name of the `.pem` file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has `CAFile` enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongotop` (page 576) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongotop` (page 576) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the `.pem` file that contains the Certificate Revocation List. Specify the file name of the `.pem` file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMODE

--sslFIPSMODE

New in version 2.6.

Directs the `mongotop` (page 576) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMODE` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, **-u**

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, **-p**

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongotop` (page 576) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the `PLAIN` and `MONGODB-X509` authentication mechanisms.

Specifies the authentication mechanism the `mongotop` (page 576) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use <code>PLAIN</code> for authenticating in-database users. <code>PLAIN</code> transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ⁴¹ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ⁴² .

command line option!—locks

--locks

Toggles the mode of `mongotop` (page 576) to report on use of per-database *locks* (page 348). These data are useful for measuring concurrent operations and lock percentage.

<sleeptime>

The final argument is the length of time, in seconds, that `mongotop` (page 576) waits in between calls. By default `mongotop` (page 576) returns data every second.

Fields

`mongotop` (page 576) returns time values specified in milliseconds (ms.)

`mongotop` (page 576) only reports active namespaces or databases, depending on the `--locks` (page 579) option. If you don't see a database or collection, it has received no recent activity. You can issue a simple operation in the `mongo` (page 527) shell to generate activity to affect the output of `mongotop` (page 576).

`mongotop.ns`

Contains the database namespace, which combines the database name and collection.

Changed in version 2.2: If you use the `--locks` (page 579), the `ns` (page 579) field does not appear in the `mongotop` (page 576) output.

`mongotop.db`

New in version 2.2.

Contains the name of the database. The database named `.` refers to the global lock, rather than a specific database.

This field does not appear unless you have invoked `mongotop` (page 576) with the `--locks` (page 579) option.

`mongotop.total`

Provides the total amount of time that this `mongod` (page 503) spent operating on this namespace.

³⁹<http://www.mongodb.com/products/mongodb-enterprise>

⁴⁰<http://www.mongodb.com/products/mongodb-enterprise>

⁴¹<http://www.mongodb.com/products/mongodb-enterprise>

⁴²<http://www.mongodb.com/products/mongodb-enterprise>

mongotop.read

Provides the amount of time that this [mongod](#) (page 503) spent performing read operations on this namespace.

mongotop.write

Provides the amount of time that this [mongod](#) (page 503) spent performing write operations on this namespace.

mongotop.<timestamp>

Provides a time stamp for the returned data.

Use

By default [mongotop](#) (page 576) connects to the MongoDB instance running on the localhost port 27017. However, [mongotop](#) (page 576) can optionally connect to remote [mongod](#) (page 503) instances. See the [mongotop options](#) (page 576) for more information.

To force [mongotop](#) (page 576) to return less frequently specify a number, in seconds at the end of the command. In this example, [mongotop](#) (page 576) will return every 15 seconds.

```
mongotop 15
```

This command produces the following output:

```
connected to: 127.0.0.1
```

	ns	total	read	write	2012-08-13T15:45:40
test.system.namespaces		0ms	0ms	0ms	
local.system.replset		0ms	0ms	0ms	
local.system.indexes		0ms	0ms	0ms	
admin.system.indexes		0ms	0ms	0ms	
admin.		0ms	0ms	0ms	

	ns	total	read	write	2012-08-13T15:45:55
test.system.namespaces		0ms	0ms	0ms	
local.system.replset		0ms	0ms	0ms	
local.system.indexes		0ms	0ms	0ms	
admin.system.indexes		0ms	0ms	0ms	
admin.		0ms	0ms	0ms	

To return a [mongotop](#) (page 576) report every 5 minutes, use the following command:

```
mongotop 300
```

To report the use of per-database locks, use `mongotop --locks`, which produces the following output:

```
$ mongotop --locks
```

```
connected to: 127.0.0.1
```

	db	total	read	write	2012-08-13T16:33:34
	local	0ms	0ms	0ms	
	admin	0ms	0ms	0ms	
	.	0ms	0ms	0ms	

mongosniff**Synopsis**

[mongosniff](#) (page 581) provides a low-level operation tracing/sniffing view into database activity in real time.

Think of `mongosniff` (page 581) as a MongoDB-specific analogue of `tcpdump` for TCP/IP network traffic. Typically, `mongosniff` (page 581) is most frequently used in driver development.

Note: `mongosniff` (page 581) requires `libpcap` and is only available for Unix-like systems. Furthermore, the version distributed with the MongoDB binaries is dynamically linked against version 0.9 of `libpcap`. If your system has a different version of `libpcap`, you will need to compile `mongosniff` (page 581) yourself or create a symbolic link pointing to `libpcap.so.0.9` to your local version of `libpcap`. Use an operation that resembles the following:

```
ln -s /usr/lib/libpcap.so.1.1.1 /usr/lib/libpcap.so.0.9
```

Change the path's and name of the shared library as needed.

As an alternative to `mongosniff` (page 581), Wireshark, a popular network sniffing tool is capable of inspecting and parsing the MongoDB wire protocol.

Options

`mongosniff`

`mongosniff`

command line option!-help, -h

`--help, -h`

Returns information on the options and use of `mongosniff` (page 581).

command line option!-forward <host><:port>

`--forward <host><:port>`

Declares a host to forward all parsed requests that the `mongosniff` (page 581) intercepts to another `mongod` (page 503) instance and issue those operations on that database instance.

Specify the target host name and port in the <host><:port> format.

To connect to a replica set, specify the replica set seed name and the seed list of set members. Use the following format:

```
<replica_set_name>/<hostname1><:port>,<hostname2><:port>,...
```

command line option!-source <NET [interface]>, <FILE [filename]>, <DIAGLOG [filename]>

`--source <NET [interface]>`

Specifies source material to inspect. Use `--source NET [interface]` to inspect traffic from a network interface (e.g. `eth0` or `lo`.) Use `--source FILE [filename]` to read captured packets in *pcap* format.

You may use the `--source DIAGLOG [filename]` option to read the output files produced by the `--diaglog` option.

command line option!-objcheck

`--objcheck`

Displays invalid BSON objects only and nothing else. Use this option for troubleshooting driver development. This option has some performance impact on the performance of `mongosniff` (page 581).

<port>

Specifies alternate ports to sniff for traffic. By default, `mongosniff` (page 581) watches for MongoDB traffic on port 27017. Append multiple port numbers to the end of `mongosniff` (page 581) to monitor traffic on multiple ports.

Use

Use the following command to connect to a `mongod` (page 503) or `mongos` (page 518) running on port 27017 and 27018 on the localhost interface:

```
mongosniff --source NET lo 27017 27018
```

Use the following command to only log invalid *BSON* objects for the `mongod` (page 503) or `mongos` (page 518) running on the localhost interface and port 27018, for driver development and troubleshooting:

```
mongosniff --objcheck --source NET lo 27018
```

Build `mongosniff`

To build `mongosniff` yourself, Linux users can use the following procedure:

1. Obtain prerequisites using your operating systems package management software. Dependencies include:

- `libpcap` - to capture network packets.
- `git` - to download the MongoDB source code.
- `scons` and a C++ compiler - to build `mongosniff` (page 581).

2. Download a copy of the MongoDB source code using `git`:

```
git clone git://github.com/mongodb/mongo.git
```

3. Issue the following sequence of commands to change to the `mongo/` directory and build `mongosniff` (page 581):

```
cd mongo
scons mongosniff
```

Note: If you run `scons mongosniff` before installing `libpcap` you must run `scons clean` before you can build `mongosniff` (page 581).

`mongoperf`

Synopsis

`mongoperf` (page 583) is a utility to check disk I/O performance independently of MongoDB.

It times tests of random disk I/O and presents the results. You can use `mongoperf` (page 583) for any case apart from MongoDB. The `mmf` (page 584) `true` mode is completely generic. In that mode it is somewhat analogous to tools such as `bonnie++`⁴³ (albeit `mongoperf` is simpler).

Specify options to `mongoperf` (page 583) using a JavaScript document.

See also:

- `bonnie`⁴⁴
- `bonnie++`⁴⁵

⁴³<http://sourceforge.net/projects/bonnie/>

⁴⁴<http://www.textuality.com/bonnie/>

⁴⁵<http://sourceforge.net/projects/bonnie/>

- Output from an example run⁴⁶
- Checking Disk Performance with the mongoperf Utility⁴⁷

Options

mongoperf

mongoperf

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongoperf` (page 583).

<jsonconfig>

`mongoperf` (page 583) accepts configuration options in the form of a file that holds a *JSON* document. You must stream the content of this file into `mongoperf` (page 583), as in the following operation:

```
mongoperf < config
```

In this example `config` is the name of a file that holds a JSON document that resembles the following example:

```
{
  nThreads:<n>,
  fileSizeMB:<n>,
  sleepMicros:<n>,
  mmf:<bool>,
  r:<bool>,
  w:<bool>,
  recSizeKB:<n>,
  syncDelay:<n>
}
```

See the *Configuration Fields* (page 583) section for documentation of each of these fields.

Configuration Fields

mongoperf.nThreads

Type: Integer.

Default: 1

Defines the number of threads `mongoperf` (page 583) will use in the test. To saturate your system's storage system you will need multiple threads. Consider setting `nThreads` (page 583) to 16.

mongoperf.fileSizeMB

Type: Integer.

Default: 1 megabyte (i.e. 1024² bytes)

Test file size.

mongoperf.sleepMicros

Type: Integer.

Default: 0

⁴⁶<https://gist.github.com/1694664>

⁴⁷<http://blog.mongodb.org/post/40769806981/checking-disk-performance-with-the-mongoperf-utility>

`mongoperf` (page 583) will pause for the number of specified `sleepMicros` (page 583) divided by the `nThreads` (page 583) between each operation.

`mongoperf.mmfs`

Type: Boolean.

Default: `false`

Set `mmfs` (page 584) to `true` to use memory mapped files for the tests.

Generally:

- when `mmfs` (page 584) is `false`, `mongoperf` (page 583) tests direct, physical, I/O, without caching. Use a large file size to test heavy random I/O load and to avoid I/O coalescing.
- when `mmfs` (page 584) is `true`, `mongoperf` (page 583) runs tests of the caching system, and can use normal file system cache. Use `mmfs` (page 584) in this mode to test file system cache behavior with memory mapped files.

`mongoperf.r`

Type: Boolean.

Default: `false`

Set `r` (page 584) to `true` to perform reads as part of the tests.

Either `r` (page 584) or `w` (page 584) must be `true`.

`mongoperf.w`

Type: Boolean.

Default: `false`

Set `w` (page 584) to `true` to perform writes as part of the tests.

Either `r` (page 584) or `w` (page 584) must be `true`.

`mongoperf.recSizeKB`

New in version 2.4.

Type: Integer.

Default: 4 kb

The size of each write operation.

`mongoperf.syncDelay`

Type: Integer.

Default: 0

Seconds between disk flushes. `mongoperf.syncDelay` (page 584) is similar to `--syncdelay` for `mongod` (page 503).

The `syncDelay` (page 584) controls how frequently `mongoperf` (page 583) performs an asynchronous disk flush of the memory mapped file used for testing. By default, `mongod` (page 503) performs this operation every 60 seconds. Use `syncDelay` (page 584) to test basic system performance of this type of operation.

Only use `syncDelay` (page 584) in conjunction with `mmfs` (page 584) set to `true`.

The default value of 0 disables this.

Use

```
mongoperf < jsonconfigfile
```

Replace `jsonconfigfile` with the path to the `mongoperf` (page 583) configuration. You may also invoke `mongoperf` (page 583) in the following form:

```
echo "{nThreads:16,fileSizeMB:1000,r:true}" | ./mongoperf
```

In this operation:

- `mongoperf` (page 583) tests direct physical random read io's, using 16 concurrent reader threads.
- `mongoperf` (page 583) uses a 1 gigabyte test file.

Consider using `iostat`, as invoked in the following example to monitor I/O performance during the test.

```
iostat -xm 2
```

4.1.6 GridFS

`mongofiles` (page 585) provides a command-line interact to a MongoDB *GridFS* storage system.

mongofiles

mongofiles

Synopsis

The `mongofiles` (page 585) utility makes it possible to manipulate files stored in your MongoDB instance in *GridFS* objects from the command line. It is particularly useful as it provides an interface between objects stored in your file system and GridFS.

All `mongofiles` (page 585) commands have the following form:

```
mongofiles <options> <commands> <filename>
```

The components of the `mongofiles` (page 585) command are:

1. *Options* (page 586). You may use one or more of these options to control the behavior of `mongofiles` (page 585).
2. *Commands* (page 589). Use one of these commands to determine the action of `mongofiles` (page 585).
3. A filename which is either: the name of a file on your local's file system, or a GridFS object.

`mongofiles` (page 585), like `mongodump` (page 537), `mongoexport` (page 562), `mongoimport` (page 556), and `mongorestore` (page 543), can access data stored in a MongoDB data directory without requiring a running `mongod` (page 503) instance, if no other `mongod` (page 503) is running.

Important: For *replica sets*, `mongofiles` (page 585) can only read from the set's *primary*.

Options

mongofiles

command line option!-help, -h

--help, -h

Returns information on the options and use of `mongofiles` (page 585).

command line option!-verbose, -v

--verbose, -v

Increases the amount of internal reporting returned on standard output or in log files. Increase the verbosity with the `-v` form by including the option multiple times, (e.g. `-vvvvvv`.)

command line option!-quiet

--quiet

Runs the `mongofiles` (page 585) in a quiet mode that attempts to limit the amount of output. This option suppresses:

- output from *database commands*
- replication activity
- connection accepted events
- connection closed events

command line option!-version

--version

Returns the `mongofiles` (page 585) release number.

command line option!-host <hostname><:port>

--host <hostname><:port>

Specifies a resolvable hostname for the `mongod` (page 503) that holds your GridFS system. By default `mongofiles` (page 585) attempts to connect to a MongoDB process running on the localhost port number 27017.

Optionally, specify a port number to connect a MongoDB instance running on a port other than 27017.

command line option!-port <port>

--port <port>

Default: 27017

Specifies the TCP port on which the MongoDB instance listens for client connections.

command line option!-ipv6

--ipv6

Enables IPv6 support and allows the `mongofiles` (page 585) to connect to the MongoDB instance using an IPv6 network. All MongoDB programs and processes disable IPv6 support by default.

command line option!-ssl

--ssl

New in version 2.6.

Enables connection to a `mongod` (page 503) or `mongos` (page 518) that has SSL support enabled.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCAFile <filename>

--sslCAFile <filename>

New in version 2.6.

Specifies the .pem file that contains the root certificate chain from the Certificate Authority. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyFile <filename>

--sslPEMKeyFile <filename>

New in version 2.6.

Specifies the .pem file that contains both the SSL certificate and key. Specify the file name of the .pem file using relative or absolute paths.

This option is required when using the `--ssl` (page 577) option to connect to a `mongod` (page 503) or `mongos` (page 518) that has CAFile enabled *without* `weakCertificateValidation`.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslPEMKeyPassword <value>

--sslPEMKeyPassword <value>

New in version 2.6.

Specifies the password to de-encrypt the certificate-key file (i.e. `--sslPEMKeyFile` (page 577)). Use the `--sslPEMKeyPassword` (page 577) option only if the certificate-key file is encrypted. In all cases, the `mongofiles` (page 585) will redact the password from all logging and reporting output.

If the private key in the PEM file is encrypted and you do not specify the `--sslPEMKeyPassword` (page 577) option, the `mongofiles` (page 585) will prompt for a passphrase. See *ssl-certificate-password*.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslCRLFile <filename>

--sslCRLFile <filename>

New in version 2.6.

Specifies the .pem file that contains the Certificate Revocation List. Specify the file name of the .pem file using relative or absolute paths.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslAllowInvalidCertificates

--sslAllowInvalidCertificates

New in version 2.6.

Bypasses the validation checks for server certificates and allows the use of invalid certificates. When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificate.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-sslFIPSMode

--sslFIPSMode

New in version 2.6.

Directs the `mongofiles` (page 585) to use the FIPS mode of the installed OpenSSL library. Your system must have a FIPS compliant OpenSSL library to use the `--sslFIPSMode` (page 578) option.

The default distribution of MongoDB does not contain support for SSL. For more information on MongoDB and SSL, see <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

command line option!-username <username>, -u

--username <username>, -u

Specifies a username with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--password` and `--authenticationDatabase` options.

command line option!-password <password>, -p

--password <password>, -p

Specifies a password with which to authenticate to a MongoDB database that uses authentication. Use in conjunction with the `--username` and `--authenticationDatabase` options.

command line option!-authenticationDatabase <dbname>

--authenticationDatabase <dbname>

New in version 2.4.

Specifies the database that holds the user's credentials. If you do not specify an authentication database, the `mongofiles` (page 585) assumes that the database specified as the argument to the `--db` (page 547) option holds the user's credentials.

command line option!-authenticationMechanism <name>

--authenticationMechanism <name>

Default: MONGODB-CR

New in version 2.4.

Changed in version 2.6: Added support for the PLAIN and MONGODB-X509 authentication mechanisms.

Specifies the authentication mechanism the `mongofiles` (page 585) instance uses to authenticate to the `mongod` (page 503) or `mongos` (page 518).

Value	Description
MONGODB-CR	MongoDB challenge/response authentication.
MONGODB-X509	MongoDB SSL certificate authentication.
PLAIN	External authentication using LDAP. You can also use PLAIN for authenticating in-database users. PLAIN transmits passwords in plain text. This mechanism is available only in MongoDB Enterprise ⁵⁰ .
GSSAPI	External authentication using Kerberos. This mechanism is available only in MongoDB Enterprise ⁵¹ .

command line option!-dbpath <path>

--dbpath <path>

Specifies the directory of the MongoDB data files. The `--dbpath` (page 546) option lets the `mongofiles` (page 585) attach directly to the local data files without going through a running `mongod` (page 503). When

⁴⁸<http://www.mongodb.com/products/mongodb-enterprise>

⁴⁹<http://www.mongodb.com/products/mongodb-enterprise>

⁵⁰<http://www.mongodb.com/products/mongodb-enterprise>

⁵¹<http://www.mongodb.com/products/mongodb-enterprise>

run with `--dbpath` (page 546), the `mongofiles` (page 585) locks access to the data files. No `mongod` (page 503) can access the files while the `mongofiles` (page 585) process runs.

command line option!-directoryperdb

--directoryperdb

When used in conjunction with the corresponding option in `mongod` (page 503), allows the `mongofiles` (page 585) to access data from MongoDB instances that use an on-disk format where every database has a distinct directory. This option is only relevant when specifying the `--dbpath` (page 546) option.

command line option!-journal

--journal

Enables the durability *journal* to ensure data files remain valid and recoverable. This option applies only when you specify the `--dbpath` (page 546) option. The `mongofiles` (page 585) enables journaling by default on 64-bit builds of versions after 2.0.

command line option!-db <database>, -d

--db <database>, -d

Specifies the name of the database on which to run the `mongofiles` (page 585).

command line option!-collection <collection>, -c

--collection <collection>, -c

This option has no use in this context and a future release may remove it. See [SERVER-4931⁵²](#) for more information.

command line option!-local <filename>, -l

--local <filename>, -l

Specifies the local filesystem name of a file for get and put operations.

In the `mongofiles put` and `mongofiles get` commands, the required <filename> modifier refers to the name the object will have in GridFS. `mongofiles` (page 585) assumes that this reflects the file's name on the local file system. This setting overrides this default.

command line option!-type <MIME>

--type <MIME>

Provides the ability to specify a *MIME* type to describe the file inserted into GridFS storage. `mongofiles` (page 585) omits this option in the default operation.

Use only with `mongofiles put` operations.

command line option!-replace, -r

--replace, -r

Alters the behavior of `mongofiles put` to replace existing GridFS objects with the specified local file, rather than adding an additional object with the same name.

In the default operation, files will not be overwritten by a `mongofiles put` option.

Commands

list <prefix>

Lists the files in the GridFS store. The characters specified after `list` (e.g. <prefix>) optionally limit the list of returned items to files that begin with that string of characters.

search <string>

Lists the files in the GridFS store with names that match any portion of <string>.

⁵²<https://jira.mongodb.org/browse/SERVER-4931>

put <filename>

Copy the specified file from the local file system into GridFS storage.

Here, <filename> refers to the name the object will have in GridFS, and [mongofiles](#) (page 585) assumes that this reflects the name the file has on the local file system. If the local filename is different use the *`mongofiles --local`* option.

get <filename>

Copy the specified file from GridFS storage to the local file system.

Here, <filename> refers to the name the object will have in GridFS, and [mongofiles](#) (page 585) assumes that this reflects the name the file has on the local file system. If the local filename is different use the *`mongofiles --local`* option.

delete <filename>

Delete the specified file from GridFS storage.

Examples

To return a list of all files in a *GridFS* collection in the `records` database, use the following invocation at the system shell:

```
mongofiles -d records list
```

This [mongofiles](#) (page 585) instance will connect to the [mongod](#) (page 503) instance running on the 27017 localhost interface to specify the same operation on a different port or hostname, and issue a command that resembles one of the following:

```
mongofiles --port 37017 -d records list
mongofiles --hostname db1.example.net -d records list
mongofiles --hostname db1.example.net --port 37017 -d records list
```

Modify any of the following commands as needed if you're connecting the [mongod](#) (page 503) instances on different ports or hosts.

To upload a file named `32-corinth.lp` to the GridFS collection in the `records` database, you can use the following command:

```
mongofiles -d records put 32-corinth.lp
```

To delete the `32-corinth.lp` file from this GridFS collection in the `records` database, you can use the following command:

```
mongofiles -d records delete 32-corinth.lp
```

To search for files in the GridFS collection in the `records` database that have the string `corinth` in their names, you can use following command:

```
mongofiles -d records search corinth
```

To list all files in the GridFS collection in the `records` database that begin with the string `32`, you can use the following command:

```
mongofiles -d records list 32
```

To fetch the file from the GridFS collection in the `records` database named `32-corinth.lp`, you can use the following command:

```
mongofiles -d records get 32-corinth.lp
```

Internal Metadata

5.1 Config Database

The `config` database supports *sharded cluster* operation. See the <http://docs.mongodb.org/manual/sharding> section of this manual for full documentation of sharded clusters.

Important: Consider the schema of the `config` database *internal* and may change between releases of MongoDB. The `config` database is not a dependable API, and users should not write data to the `config` database in the course of normal operation or maintenance.

Warning: Modification of the `config` database on a functioning system may lead to instability or inconsistent data sets. If you must modify the `config` database, use `mongodump` (page 537) to create a full backup of the `config` database.

To access the `config` database, connect to a `mongos` (page 518) instance in a sharded cluster, and use the following helper:

```
use config
```

You can return a list of the collections, with the following helper:

```
show collections
```

5.1.1 Collections

config

`config.changelog`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `changelog` (page 593) collection stores a document for each change to the metadata of a sharded collection.

Example

The following example displays a single record of a chunk split from a `changelog` (page 593) collection:

```
{
  "_id" : "<hostname>-<timestamp>-<increment>",
  "server" : "<hostname>:<port>",
  "clientAddr" : "127.0.0.1:63381",
  "time" : ISODate("2012-12-11T14:09:21.039Z"),
  "what" : "split",
  "ns" : "<database>.<collection>",
  "details" : {
    "before" : {
      "min" : {
        "<database>" : { $minKey : 1 }
      },
      "max" : {
        "<database>" : { $maxKey : 1 }
      },
      "lastmod" : Timestamp(1000, 0),
      "lastmodEpoch" : ObjectId("000000000000000000000000")
    },
    "left" : {
      "min" : {
        "<database>" : { $minKey : 1 }
      },
      "max" : {
        "<database>" : "<value>"
      },
      "lastmod" : Timestamp(1000, 1),
      "lastmodEpoch" : ObjectId(<...>)
    },
    "right" : {
      "min" : {
        "<database>" : "<value>"
      },
      "max" : {
        "<database>" : { $maxKey : 1 }
      },
      "lastmod" : Timestamp(1000, 2),
      "lastmodEpoch" : ObjectId("<...>")
    }
  }
}
```

Each document in the `changelog` (page 593) collection contains the following fields:

`config.changelog._id`

The value of `changelog._id` is: `<hostname>-<timestamp>-<increment>`.

`config.changelog.server`

The hostname of the server that holds this data.

`config.changelog.clientAddr`

A string that holds the address of the client, a `mongos` (page 518) instance that initiates this change.

`config.changelog.time`

A *ISODate* timestamp that reflects when the change occurred.

`config.changelog.what`

Reflects the type of change recorded. Possible values are:

- dropCollection
- dropCollection.start
- dropDatabase
- dropDatabase.start
- moveChunk.start
- moveChunk.commit
- split
- multi-split

`config.changelog.ns`

Namespace where the change occurred.

`config.changelog.details`

A *document* that contains additional details regarding the change. The structure of the `details` (page 595) document depends on the type of change.

`config.chunks`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `chunks` (page 595) collection stores a document for each chunk in the cluster. Consider the following example of a document for a chunk named `records.pets-animal_"cat"`:

```
{
  "_id" : "mydb.foo-a_"cat",
  "lastmod" : Timestamp(1000, 3),
  "lastmodEpoch" : ObjectId("5078407bd58b175c5c225fdc"),
  "ns" : "mydb.foo",
  "min" : {
    "animal" : "cat"
  },
  "max" : {
    "animal" : "dog"
  },
  "shard" : "shard0004"
}
```

These documents store the range of values for the shard key that describe the chunk in the `min` and `max` fields. Additionally the `shard` field identifies the shard in the cluster that “owns” the chunk.

`config.collections`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `collections` (page 595) collection stores a document for each sharded collection in the cluster. Given a collection named `pets` in the `records` database, a document in the `collections` (page 595) collection would resemble the following:

```
{
  "_id" : "records.pets",
  "lastmod" : ISODate("1970-01-16T15:00:58.107Z"),
  "dropped" : false,
  "key" : {
    "a" : 1
  },
  "unique" : false,
  "lastmodEpoch" : ObjectId("5078407bd58b175c5c225fdc")
}
```

`config.databases`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `databases` (page 596) collection stores a document for each database in the cluster, and tracks if the database has sharding enabled. `databases` (page 596) represents each database in a distinct document. When a databases have sharding enabled, the `primary` field holds the name of the *primary shard*.

```
{ "_id" : "admin", "partitioned" : false, "primary" : "config" }
{ "_id" : "mydb", "partitioned" : true, "primary" : "shard0000" }
```

`config.lockpings`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `lockpings` (page 596) collection keeps track of the active components in the sharded cluster. Given a cluster with a `mongos` (page 518) running on `example.com:30000`, the document in the `lockpings` (page 596) collection would resemble:

```
{ "_id" : "example.com:30000:1350047994:16807", "ping" : ISODate("2012-10-12T18:32:54.892Z") }
```

`config.locks`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `locks` (page 596) collection stores a distributed lock. This ensures that only one `mongos` (page 518) instance can perform administrative tasks on the cluster at once. The `mongos` (page 518) acting as *balancer* takes a lock by inserting a document resembling the following into the `locks` collection.

```
{
  "_id" : "balancer",
  "process" : "example.net:40000:1350402818:16807",
  "state" : 2,
  "ts" : ObjectId("507daeef40e1879df62e5f3"),
  "when" : ISODate("2012-10-16T19:01:01.593Z"),
}
```



```

    "who" : "example.net:40000:1350402818:16807:Balancer:282475249",
    "why" : "doing balance round"
  }

```

If a `mongos` (page 518) holds the balancer lock, the `state` field has a value of 2, which means that balancer is active. The `when` field indicates when the balancer began the current operation.

Changed in version 2.0: The value of the `state` field was 1 before MongoDB 2.0.

`config.mongos`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `mongos` (page 597) collection stores a document for each `mongos` (page 518) instance affiliated with the cluster. `mongos` (page 518) instances send pings to all members of the cluster every 30 seconds so the cluster can verify that the `mongos` (page 518) is active. The `ping` field shows the time of the last ping, while the `up` field reports the uptime of the `mongos` (page 518) as of the last ping. The cluster maintains this collection for reporting purposes.

The following document shows the status of the `mongos` (page 518) running on `example.com:30000`.

```

{ "_id" : "example.com:30000", "ping" : ISODate("2012-10-12T17:08:13.538Z"), "up" : 13699, "wait"

```

`config.settings`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `settings` (page 597) collection holds the following sharding configuration settings:

- Chunk size. To change chunk size, see <http://docs.mongodb.org/manual/tutorial/modify-chunk-size/>
- Balancer status. To change status, see *sharding-balancing-disable-temporarily*.

The following is an example `settings` collection:

```

{ "_id" : "chunksize", "value" : 64 }
{ "_id" : "balancer", "stopped" : false }

```

`config.shards`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `shards` (page 597) collection represents each shard in the cluster in a separate document. If the shard is a replica set, the `host` field displays the name of the replica set, then a slash, then the hostname, as in the following example:

```

{ "_id" : "shard0000", "host" : "shard1/localhost:30000" }

```

If the shard has *tags* assigned, this document has a `tags` field, that holds an array of the tags, as in the following example:

```
{ "_id" : "shard0001", "host" : "localhost:30001", "tags": [ "NYC" ] }
```

`config.tags`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `tags` (page 598) collection holds documents for each tagged shard key range in the cluster. The documents in the `tags` (page 598) collection resemble the following:

```
{
  "_id" : { "ns" : "records.users", "min" : { "zipcode" : "10001" } },
  "ns" : "records.users",
  "min" : { "zipcode" : "10001" },
  "max" : { "zipcode" : "10281" },
  "tag" : "NYC"
}
```

`config.version`

Internal MongoDB Metadata

The `config` (page 593) database is internal: applications and administrators should not modify or depend upon its content in the course of normal operation.

The `version` (page 598) collection holds the current metadata version number. This collection contains only one document:

To access the `version` (page 598) collection you must use the `db.getCollection()` (page 110) method. For example, to display the collection's document:

```
mongos> db.getCollection("version").find()
{ "_id" : 1, "version" : 3 }
```

Note: Like all databases in MongoDB, the `config` database contains a `system.indexes` (page 601) collection contains metadata for all indexes in the database for information on indexes, see <http://docs.mongodb.org/manual/indexes>.

5.2 The `local` Database

5.2.1 Overview

Every `mongod` (page 503) instance has its own `local` database, which stores data used in the replication process, and other instance-specific data. The `local` database is invisible to replication: collections in the `local` database are not replicated.

In replication, the `local` database store stores internal replication data for each member of a *replica set*. The `local` stores the following collections:

Changed in version 2.4: When running with authentication (i.e. `authorization`), authenticating to the `local` database is **not** equivalent to authenticating to the `admin` database. In previous versions, authenticating to the `local` database provided access to all databases.

5.2.2 Collection on all `mongod` Instances

`local.startup_log`

On startup, each `mongod` (page 503) instance inserts a document into `startup_log` (page 599) with diagnostic information about the `mongod` (page 503) instance itself and host information. `startup_log` (page 599) is a capped collection. This information is primarily useful for diagnostic purposes.

Example

Consider the following prototype of a document from the `startup_log` (page 599) collection:

```
{
  "_id" : "<string>",
  "hostname" : "<string>",
  "startTime" : ISODate("<date>"),
  "startTimeLocal" : "<string>",
  "cmdLine" : {
    "dbpath" : "<path>",
    "<option>" : <value>
  },
  "pid" : <number>,
  "buildinfo" : {
    "version" : "<string>",
    "gitVersion" : "<string>",
    "sysInfo" : "<string>",
    "loaderFlags" : "<string>",
    "compilerFlags" : "<string>",
    "allocator" : "<string>",
    "versionArray" : [ <num>, <num>, <...> ],
    "javascriptEngine" : "<string>",
    "bits" : <number>,
    "debug" : <boolean>,
    "maxBsonObjectSize" : <number>
  }
}
```

Documents in the `startup_log` (page 599) collection contain the following fields:

`local.startup_log._id`

Includes the system hostname and a millisecond epoch value.

`local.startup_log.hostname`

The system's hostname.

`local.startup_log.startTime`

A UTC *ISODate* value that reflects when the server started.

`local.startup_log.startTimeLocal`

A string that reports the `startTime` (page 599) in the system's local time zone.

`local.startup_log.cmdLine`

A sub-document that reports the `mongod` (page 503) runtime options and their values.

`local.startup_log.pid`

The process identifier for this process.

local.startup_log.buildinfo

A sub-document that reports information about the build environment and settings used to compile this `mongod` (page 503). This is the same output as `buildInfo` (page 324). See `buildInfo` (page 325).

5.2.3 Collections on Replica Set Members

local.system.replset

`local.system.replset` (page 600) holds the replica set's configuration object as its single document. To view the object's configuration information, issue `rs.conf()` (page 165) from the `mongo` (page 527) shell. You can also query this collection directly.

local.oplog.rs

`local.oplog.rs` (page 600) is the capped collection that holds the *oplog*. You set its size at creation using the `oplogSizeMB` setting. To resize the oplog after replica set initiation, use the <http://docs.mongodb.org/manual/tutorial/change-oplog-size> procedure. For additional information, see the *replica-set-oplog-sizing* section.

local.replset.minvalid

This contains an object used internally by replica sets to track replication status.

local.slaves

This contains information about each member of the set and the latest point in time that this member has synced to. If this collection becomes out of date, you can refresh it by dropping the collection and allowing MongoDB to automatically refresh it during normal replication:

```
db.getSiblingDB("local").slaves.drop()
```

5.2.4 Collections used in Master/Slave Replication

In *master/slave* replication, the `local` database contains the following collections:

- On the master:

local.oplog.\$main

This is the oplog for the master-slave configuration.

local.slaves

This contains information about each slave.

- On each slave:

local.sources

This contains information about the slave's master server.

5.3 System Collections

5.3.1 Synopsis

MongoDB stores system information in collections that use the `<database>.system.* namespace`, which MongoDB reserves for internal use. Do not create collections that begin with `system`.

MongoDB also stores some additional instance-local metadata in the *local database* (page 598), specifically for replication purposes.

5.3.2 Collections

System collections include these collections stored in the `admin` database:

`admin.system.roles`

New in version 2.6.

The `admin.system.roles` (page 601) collection stores custom roles that administrators create and assign to users to provide access to specific resources.

`admin.system.users`

Changed in version 2.6.

The `admin.system.users` (page 601) collection stores the user's authentication credentials as well as any roles assigned to the user. Users may define authorization roles in the `admin.system.roles` (page 601) collection.

`admin.system.version`

New in version 2.6.

Stores the schema version of the user credential documents.

System collections also include these collections stored directly in each database:

`<database>.system.namespaces`

The `<database>.system.namespaces` (page 601) collection contains information about all of the database's collections. Additional namespace metadata exists in the `database.ns` files and is opaque to database users.

`<database>.system.indexes`

The `<database>.system.indexes` (page 601) collection lists all the indexes in the database. Add and remove data from this collection via the `ensureIndex()` (page 30) and `dropIndex()` (page 29)

`<database>.system.profile`

The `<database>.system.profile` (page 601) collection stores database profiling information. For information on profiling, see *database-profiling*.

`<database>.system.js`

The `<database>.system.js` (page 601) collection holds special JavaScript code for use in server side JavaScript. See <http://docs.mongodb.org/manual/tutorial/store-javascript-function-on-database/> for more information.

General System Reference

6.1 Exit Codes and Statuses

MongoDB will return one of the following codes and statuses when exiting. Use this guide to interpret logs and when troubleshooting issues with [mongod](#) (page 503) and [mongos](#) (page 518) instances.

- 0
Returned by MongoDB applications upon successful exit.
- 2
The specified options are in error or are incompatible with other options.
- 3
Returned by [mongod](#) (page 503) if there is a mismatch between hostnames specified on the command line and in the `local.sources` (page 600) collection. [mongod](#) (page 503) may also return this status if `oplog` collection in the `local` database is not readable.
- 4
The version of the database is different from the version supported by the [mongod](#) (page 503) (or `mongod.exe` (page 534)) instance. The instance exits cleanly. Restart [mongod](#) (page 503) with the `--upgrade` option to upgrade the database to the version supported by this [mongod](#) (page 503) instance.
- 5
Returned by [mongod](#) (page 503) if a `moveChunk` (page 296) operation fails to confirm a commit.
- 12
Returned by the `mongod.exe` (page 534) process on Windows when it receives a Control-C, Close, Break or Shutdown event.
- 14
Returned by MongoDB applications which encounter an unrecoverable error, an uncaught exception or uncaught signal. The system exits without performing a clean shut down.
- 20
Message: ERROR: wsastartup failed <reason>
Returned by MongoDB applications on Windows following an error in the WSASStartup function.
Message: NT Service Error
Returned by MongoDB applications for Windows due to failures installing, starting or removing the NT Service for the application.
- 45
Returned when a MongoDB application cannot open a file or cannot obtain a lock on a file.

- 47 MongoDB applications exit cleanly following a large clock skew (32768 milliseconds) event.
- 48 `mongod` (page 503) exits cleanly if the server socket closes. The server socket is on port 27017 by default, or as specified to the `--port` run-time option.
- 49 Returned by `mongod.exe` (page 534) or `mongos.exe` (page 535) on Windows when either receives a shut-down message from the *Windows Service Control Manager*.
- 100 Returned by `mongod` (page 503) when the process throws an uncaught exception.

6.2 MongoDB Limits and Thresholds

This document provides a collection of hard and soft limitations of the MongoDB system.

6.2.1 BSON Documents

BSON Document Size

The maximum BSON document size is 16 megabytes.

The maximum document size helps ensure that a single document cannot use excessive amount of RAM or, during transmission, excessive amount of bandwidth. To store documents larger than the maximum size, MongoDB provides the GridFS API. See `mongofiles` (page 585) and the documentation for your driver for more information about GridFS.

Nested Depth for BSON Documents

Changed in version 2.2.

MongoDB supports no more than 100 levels of nesting for *BSON documents*.

6.2.2 Namespaces

Namespace Length

Each namespace, including database and collection name, must be shorter than 123 bytes.

Number of Namespaces

The limitation on the number of namespaces is the size of the namespace file divided by 628.

A 16 megabyte namespace file can support approximately 24,000 namespaces. Each collection and index is a namespace.

Size of Namespace File

Namespace files can be no larger than 2047 megabytes.

By default namespace files are 16 megabytes. You can configure the size using the `nsSize` option.

6.2.3 Indexes

Index Key Limit

The *total size* of an index entry, which can include structural overhead depending on the BSON type, must be *less than* 1024 bytes.

Changed in version 2.6: MongoDB 2.6 implements a stronger enforcement of the limit on `index key` (page 604):

- MongoDB will **not** `create an index` (page 30) on a collection if the index entry for an existing document exceeds the `index key limit` (page 604). Previous versions of MongoDB would create the index but not index such documents.
- Reindexing operations will error if the index entry for an indexed field exceeds the `index key limit` (page 604). Reindexing operations occur as part of `compact` (page 313) and `repairDatabase` (page 319) commands as well as the `db.collection.reIndex()` (page 62) method.

Because these operations drop *all* the indexes from a collection and then recreate them sequentially, the error from the `index key limit` (page 604) prevents these operations from rebuilding any remaining indexes for the collection and, in the case of the `repairDatabase` (page 319) command, from continuing with the remainder of the process.

- MongoDB will not insert into an indexed collection any document with an indexed field whose corresponding index entry would exceed the `index key limit` (page 604), and instead, will return an error. Previous versions of MongoDB would insert but not index such documents.
- Updates to the indexed field will error if the updated value causes the index entry to exceed the `index key limit` (page 604).

If an existing document contains an indexed field whose index entry exceeds the limit, *any* update that results in the relocation of that document on disk will error.

- `mongorestore` (page 543) and `mongoimport` (page 556) will not insert documents that contain an indexed field whose corresponding index entry would exceed the `index key limit` (page 604).
- In MongoDB 2.6, secondary members of replica sets will continue to replicate documents with an indexed field whose corresponding index entry exceeds the `index key limit` (page 604) on initial sync but will print warnings in the logs.

Secondary members also allow index build and rebuild operations on a collection that contains an indexed field whose corresponding index entry exceeds the `index key limit` (page 604) but with warnings in the logs.

With *mixed version* replica sets where the secondaries are version 2.6 and the primary is version 2.4, secondaries will replicate documents inserted or updated on the 2.4 primary, but will print error messages in the log if the documents contain an indexed field whose corresponding index entry exceeds the `index key limit` (page 604).

- For existing sharded collections, `chunk migration` will fail if the chunk has a document that contains an indexed field whose index entry exceeds the `index key limit` (page 604).

Number of Indexes per Collection

A single collection can have *no more* than 64 indexes.

Index Name Length

Fully qualified index names, which includes the namespace and the dot separators (i.e. `<database name>.<collection name>.$<index name>`), cannot be longer than 128 characters.

By default, `<index name>` is the concatenation of the field names and index type. You can explicitly specify the `<index name>` to the `ensureIndex()` (page 30) method to ensure that the fully qualified index name does not exceed the limit.

Number of Indexed Fields in a Compound Index

There can be no more than 31 fields in a compound index.

Queries cannot use both text and Geospatial Indexes

You cannot combine the `text` (page 235) command, which requires a special *text index*, with a query operator

that requires a different type of special index. For example you cannot combine `text` (page 235) command with the `$near` (page 394) operator.

See also:

The unique indexes limit in *Sharding Operational Restrictions* (page 606).

6.2.4 Data

Maximum Number of Documents in a Capped Collection

Changed in version 2.4.

If you specify a maximum number of documents for a capped collection using the `max` parameter to `create` (page 304), the limit must be less than 2^{32} documents. If you do not specify a maximum number of documents when creating a capped collection, there is no limit on the number of documents.

Data Size

A single `mongod` (page 503) instance cannot manage a data set that exceeds maximum virtual memory address space provided by the underlying operating system.

Table 6.1: Virtual Memory Limitations

Operating System	Journalled	Not Journalled
Linux	64 terabytes	128 terabytes
Windows	4 terabytes	8 terabytes

Number of Collections in a Database

The maximum number of collections in a database is a function of the size of the namespace file and the number of indexes of collections in the database.

See *Number of Namespaces* (page 604) for more information.

6.2.5 Replica Sets

Number of Members of a Replica Set

Replica sets can have no more than 12 members.

Number of Voting Members of a Replica Set

Only 7 members of a replica set can have votes at any given time. See *can vote replica-set-non-voting-members* for more information

Maximum Size of Auto-Created Oplog

Changed in version 2.6.

If you do not explicitly specify an oplog size (i.e. with `oplogSizeMB` or `--oplogSize`) MongoDB will create an oplog that is no larger than 50 gigabytes.

6.2.6 Sharded Clusters

Sharded clusters have the restrictions and thresholds described here.

Sharding Operational Restrictions

Operations Unavailable in Sharded Environments

The `group` (page 204) does not work with sharding. Use `mapReduce` (page 208) or `aggregate` (page 198) instead.

`db.eval()` (page 108) is incompatible with sharded collections. You may use `db.eval()` (page 108) with un-sharded collections in a shard cluster.

`$where` (page 391) does not permit references to the `db` object from the `$where` (page 391) function. This is uncommon in un-sharded collections.

The `$isolated` (page 436) update modifier does not work in sharded environments.

`$snapshot` (page 482) queries do not work in sharded environments.

The `geoSearch` (page 219) command is not supported in sharded environments.

Covered Queries in Sharded Clusters

MongoDB does not support *covered queries* from sharded collections.

Sharding Existing Collection Data Size

For existing collections that hold documents, MongoDB supports enabling sharding on *any* collections that contains less than 256 gigabytes of data. MongoDB *may* be able to shard collections with as many as 400 gigabytes depending on the distribution of document sizes. The precise size of the limitation is a function of the chunk size and the data size.

Important: Sharded collections may have *any* size, after successfully enabling sharding.

Single Document Modification Operations in Sharded Collections

All `update()` (page 69) and `remove()` (page 62) operations for a sharded collection that specify the `justOne` or `multi: false` option must include the *shard key* or the `_id` field in the query specification. `update()` (page 69) and `remove()` (page 62) operations specifying `justOne` or `multi: false` in a sharded collection without the *shard key* or the `_id` field return an error.

Unique Indexes in Sharded Collections

MongoDB does not support unique indexes across shards, except when the unique index contains the full shard key as a prefix of the index. In these situations MongoDB will enforce uniqueness across the full key, not a single field.

See

<http://docs.mongodb.org/manual/tutorial/enforce-unique-keys-for-sharded-collections> for an alternate approach.

Shard Key Limitations

Shard Key Size

A shard key cannot exceed 512 bytes.

Shard Key Index Type

A *shard key* index can be an ascending index on the shard key, a compound index that start with the shard key and specify ascending order for the shard key, or a `hashed` index.

A *shard key* index cannot be an index that specifies a `multikey` index, a `text` index or a *geospatial index* on the *shard key* fields.

Shard Key is Immutable

You cannot change a shard key after sharding the collection. If you must change a shard key:

- Dump all data from MongoDB into an external format.
- Drop the original sharded collection.
- Configure sharding using the new shard key.

- Pre-split the shard key range to ensure initial even distribution.
- Restore the dumped data into MongoDB.

Shard Key Value in a Document is Immutable

After you insert a document into a sharded collection, you cannot change the document's value for the field or fields that comprise the shard key. The `update()` (page 69) operation will not modify the value of a shard key in an existing document.

Monotonically Increasing Shard Keys Can Limit Insert Throughput

For clusters with high insert volumes, a shard keys with monotonically increasing and decreasing keys can affect insert throughput. If your shard key is the `_id` field, be aware that the default values of the `_id` fields are *ObjectIds* which have generally increasing values.

When inserting documents with monotonically increasing shard keys, all inserts belong to the same *chunk* on a single *shard*. The system will eventually divide the chunk range that receives all write operations and migrate its contents to distribute data more evenly. However, at any moment the cluster can direct insert operations only to a single shard, which creates an insert throughput bottleneck.

If the operations on the cluster are predominately read operations and updates, this limitation may not affect the cluster.

To avoid this constraint, use a *hashed shard key* or select a field that does not increase or decrease monotonically.

Changed in version 2.4: *Hashed shard keys* and *hashed indexes* store hashes of keys with ascending values.

6.2.7 Operations

Sorted Documents

MongoDB will only return sorted results on fields without an index *if* the sort operation uses less than 32 megabytes of memory.

Aggregation Pipeline Operation

Changed in version 2.6.

Pipeline stages have a limit of 100 megabytes of RAM. If a stage exceeds this limit, MongoDB will produce an error. To allow for the handling of large datasets, use the `allowDiskUse` option to enable aggregation pipeline stages to write data to temporary files.

See also:

\$sort and Memory Restrictions (page 451) and *\$group Operator and Memory* (page 448).

2d Geospatial queries cannot use the \$or operator

See

\$or (page 377) and <http://docs.mongodb.org/manualcore/geospatial-indexes>.

Spherical Polygons must fit within a hemisphere.

Any geometry specified with *GeoJSON* to *\$geoIntersects* (page 393) or *\$geoWithin* (page 392) queries, **must** fit within a single hemisphere. MongoDB interprets geometries larger than half of the sphere as queries for the smaller of the complementary geometries.

Bulk Operation Size

A bulk operation can have at most 1000 operations.

6.2.8 Naming Restrictions

Database Name Case Sensitivity

MongoDB does not permit database names that differ only by the case of the characters.

Restrictions on Database Names for Windows

Changed in version 2.2: [Restrictions on Database Names for Windows](#) (page 672).

For MongoDB deployments running on Windows, MongoDB will not permit database names that include any of the following characters:

```
/\ . " * < > : | ?
```

Also, database names cannot contain the null character.

Restrictions on Database Names for Unix and Linux Systems

For MongoDB deployments running on Unix and Linux systems, MongoDB will not permit database names that include any of the following characters:

```
/\ . "
```

Also, database names cannot contain the null character.

Length of Database Names

Database names cannot be empty and must have fewer than 64 characters.

Restriction on Collection Names

New in version 2.2.

Collection names should begin with an underscore or a letter character, and *cannot*:

- contain the \$.
- be an empty string (e.g. "").
- contain the null character.
- begin with the `system.` prefix. (Reserved for internal use.)

In the `mongo` (page 527) shell, use `db.getCollection()` (page 110) to specify collection names that might interact with the shell or are not valid JavaScript.

Restrictions on Field Names

Field names cannot contain dots (i.e. `.`), dollar signs (i.e. `$`), or null characters. See [faq-dollar-sign-escaping](#) for an alternate approach.

6.3 Glossary

\$cmd A special virtual *collection* that exposes MongoDB's *database commands*. To use database commands, see *issue-commands*.

_id A field required in every MongoDB *document*. The `_id` field must have a unique value. You can think of the `_id` field as the document's *primary key*. If you create a new document without an `_id` field, MongoDB automatically creates the field and assigns a unique BSON *ObjectId*.

accumulator An *expression* in the *aggregation framework* that maintains state between documents in the aggregation *pipeline*. For a list of accumulator operations, see `$group` (page 447).

action An operation the user can perform on a resource. Actions and *resources* combine to create *privileges*. See *action*.

admin database A privileged database. Users must have access to the `admin` database to run certain administrative commands. For a list of administrative commands, see *Instance Administration Commands* (page 299).

aggregation Any of a variety of operations that reduces and summarizes large sets of data. MongoDB's `aggregate()` (page 22) and `mapReduce()` (page 55) methods are two examples of aggregation operations. For more information, see <http://docs.mongodb.org/manualcore/aggregation>.

aggregation framework The set of MongoDB operators that let you calculate aggregate values without having to use *map-reduce*. For a list of operators, see *Aggregation Reference* (page 484).

arbiter A member of a *replica set* that exists solely to vote in *elections*. Arbiters do not replicate data. See *replica-set-arbiter-configuration*.

authentication Verification of the user identity. See <http://docs.mongodb.org/manualcore/authentication>.

authorization Provisioning of access to databases and operations. See <http://docs.mongodb.org/manualcore/authorization>.

B-tree A data structure commonly used by database management systems to store indexes. MongoDB uses B-trees for its indexes.

balancer An internal MongoDB process that runs in the context of a *sharded cluster* and manages the migration of *chunks*. Administrators must disable the balancer for all maintenance operations on a sharded cluster. See *sharding-balancing*.

BSON A serialization format used to store *documents* and make remote procedure calls in MongoDB. “BSON” is a portmanteau of the words “binary” and “JSON”. Think of BSON as a binary representation of JSON (JavaScript Object Notation) documents. See <http://docs.mongodb.org/manualreference/bson-types> and <http://docs.mongodb.org/manualreference/mongodb-extended-json>.

BSON types The set of types supported by the *BSON* serialization format. For a list of BSON types, see <http://docs.mongodb.org/manualreference/bson-types>.

CAP Theorem Given three properties of computing systems, consistency, availability, and partition tolerance, a distributed computing system can provide any two of these features, but never all three.

capped collection A fixed-sized *collection* that automatically overwrites its oldest entries when it reaches its maximum size. The MongoDB *oplog* that is used in *replication* is a capped collection. See <http://docs.mongodb.org/manualcore/capped-collections>.

checksum A calculated value used to ensure data integrity. The *md5* algorithm is sometimes used as a checksum.

chunk A contiguous range of *shard key* values within a particular *shard*. Chunk ranges are inclusive of the lower boundary and exclusive of the upper boundary. MongoDB splits chunks when they grow beyond the configured chunk size, which by default is 64 megabytes. MongoDB migrates chunks when a shard contains too many chunks of a collection relative to other shards. See *sharding-data-partitioning* and <http://docs.mongodb.org/manualcore/sharded-cluster-mechanics>.

client The application layer that uses a database for data persistence and storage. *Drivers* provide the interface level between the application layer and the database server.

cluster See *sharded cluster*.

collection A grouping of MongoDB *documents*. A collection is the equivalent of an *RDBMS* table. A collection exists within a single *database*. Collections do not enforce a schema. Documents within a collection can have different fields. Typically, all documents in a collection have a similar or related purpose. See *faq-dev-namespace*.

compound index An *index* consisting of two or more keys. See *index-type-compound*.

config database An internal database that holds the metadata associated with a *sharded cluster*. Applications and administrators should not modify the `config` database in the course of normal operation. See *Config Database* (page 593).

- config server** A `mongod` (page 503) instance that stores all the metadata associated with a *sharded cluster*. A production sharded cluster requires three config servers, each on a separate machine. See *sharding-config-server*.
- control script** A simple shell script, typically located in the `/etc/rc.d` or `/etc/init.d` directory, and used by the system's initialization process to start, restart or stop a *daemon* process.
- CRUD** An acronym for the fundamental operations of a database: Create, Read, Update, and Delete. See <http://docs.mongodb.org/manual/crud>.
- CSV** A text-based data format consisting of comma-separated values. This format is commonly used to exchange data between relational databases since the format is well-suited to tabular data. You can import CSV files using `mongoimport` (page 556).
- cursor** A pointer to the result set of a *query*. Clients can iterate through a cursor to retrieve results. By default, cursors timeout after 10 minutes of inactivity. See *read-operations-cursors*.
- daemon** The conventional name for a background, non-interactive process.
- data directory** The file-system location where the `mongod` (page 503) stores data files. The `dbPath` option specifies the data directory.
- data-center awareness** A property that allows clients to address members in a system based on their locations. *Replica sets* implement data-center awareness using *tagging*. See */data-center-awareness*.
- database** A physical container for *collections*. Each database gets its own set of files on the file system. A single MongoDB server typically has multiple databases.
- database command** A MongoDB operation, other than an insert, update, remove, or query. For a list of database commands, see *Database Commands* (page 198). To use database commands, see *issue-commands*.
- database profiler** A tool that, when enabled, keeps a record on all long-running operations in a database's `system.profile` collection. The profiler is most often used to diagnose slow queries. See *database-profiling*.
- datum** A set of values used to define measurements on the earth. MongoDB uses the *WGS84* datum in certain *geospatial* calculations. See <http://docs.mongodb.org/manual/applications/geospatial-indexes>.
- dbpath** The location of MongoDB's data file storage. See `dbPath`.
- delayed member** A *replica set* member that cannot become primary and applies operations at a specified delay. The delay is useful for protecting data from human error (i.e. unintentionally deleted databases) or updates that have unforeseen effects on the production database. See *replica-set-delayed-members*.
- diagnostic log** A verbose log of operations stored in the *dbpath*. See the `--diaglog` option.
- document** A record in a MongoDB *collection* and the basic unit of data in MongoDB. Documents are analogous to *JSON* objects but exist in the database in a more type-rich format known as *BSON*. See <http://docs.mongodb.org/manual/core/document>.
- dot notation** MongoDB uses the dot notation to access the elements of an array and to access the fields of a subdocument. See *document-dot-notation*.
- draining** The process of removing or “shedding” *chunks* from one *shard* to another. Administrators must drain shards before removing them from the cluster. See <http://docs.mongodb.org/manual/tutorial/remove-shards-from-cluster>.
- driver** A client library for interacting with MongoDB in a particular language. See <http://docs.mongodb.org/manual/applications/drivers>.
- election** The process by which members of a *replica set* select a *primary* on startup and in the event of a failure. See *replica-set-elections*.

eventual consistency A property of a distributed system that allows changes to the system to propagate gradually. In a database system, this means that readable members are not required to reflect the latest writes at all times. In MongoDB, reads to a primary have *strict consistency*; reads to secondaries have *eventual consistency*.

expression In the context of *aggregation framework*, expressions are the stateless transformations that operate on the data that passes through a *pipeline*. See <http://docs.mongodb.org/manualcore/aggregation>.

failover The process that allows a *secondary* member of a *replica set* to become *primary* in the event of a failure. See *replica-set-failover*.

field A name-value pair in a *document*. A document has zero or more fields. Fields are analogous to columns in relational databases. See *document-structure*.

firewall A system level networking filter that restricts access based on, among other things, IP address. Firewalls form a part of an effective network security strategy. See *security-firewalls*.

fsync A system call that flushes all dirty, in-memory pages to disk. MongoDB calls `fsync()` on its database files at least every 60 seconds. See *fsync* (page 312).

geohash A geohash value is a binary representation of the location on a coordinate grid. See *geospatial-indexes-geohash*.

GeoJSON A *geospatial* data interchange format based on JavaScript Object Notation (*JSON*). GeoJSON is used in *geospatial* queries. For supported GeoJSON objects, see *geo-overview-location-data*. For the GeoJSON format specification, see <http://geojson.org/geojson-spec.html>.

geospatial Data that relates to geographical location. In MongoDB, you may store, index, and query data according to geographical parameters. See <http://docs.mongodb.org/manualapplications/geospatial-indexes>.

GridFS A convention for storing large files in a MongoDB database. All of the official MongoDB drivers support this convention, as does the *mongofiles* (page 585) program. See <http://docs.mongodb.org/manualcore/gridfs> and <http://docs.mongodb.org/manualreference/gridfs>.

hashed shard key A special type of *shard key* that uses a hash of the value in the shard key field to distribute documents among members of the *sharded cluster*. See *index-type-hashed*.

haystack index A *geospatial* index that enhances searches by creating “buckets” of objects grouped by a second criterion. See <http://docs.mongodb.org/manualcore/geohaystack>.

hidden member A *replica set* member that cannot become *primary* and are invisible to client applications. See *replica-set-hidden-members*.

idempotent The quality of an operation to produce the same result given the same input, whether run once or run multiple times.

index A data structure that optimizes queries. See <http://docs.mongodb.org/manualcore/indexes>.

initial sync The *replica set* operation that replicates data from an existing replica set member to a new or restored replica set member. See *replica-set-initial-sync*.

interrupt point A point in an operation’s lifecycle when it can safely abort. MongoDB only terminates an operation at designated interrupt points. See <http://docs.mongodb.org/manualtutorial/terminate-running-operations>.

IPv6 A revision to the IP (Internet Protocol) standard that provides a significantly larger address space to more effectively support the number of hosts on the contemporary Internet.

ISODate The international date format used by *mongo* (page 527) to display dates. The format is: YYYY-MM-DD HH:MM.SS.millis.

- JavaScript** A popular scripting language originally designed for web browsers. The MongoDB shell and certain server-side functions use a JavaScript interpreter. See <http://docs.mongodb.org/manualcore/server-side-javascript> for more information.
- journal** A sequential, binary transaction log used to bring the database into a valid state in the event of a hard shutdown. Journaling writes data first to the journal and then to the core data files. MongoDB enables journaling by default for 64-bit builds of MongoDB version 2.0 and newer. Journal files are pre-allocated and exist as files in the *data directory*. See <http://docs.mongodb.org/manualcore/journaling/>.
- JSON** JavaScript Object Notation. A human-readable, plain text format for expressing structured data with support in many programming languages. For more information, see <http://www.json.org>. Certain MongoDB tools render an approximation of MongoDB *BSON* documents in JSON format. See <http://docs.mongodb.org/manualreference/mongodb-extended-json>.
- JSON document** A *JSON* document is a collection of fields and values in a structured format. For sample JSON documents, see <http://json.org/example.html>.
- JSONP** *JSON* with Padding. Refers to a method of injecting JSON into applications. **Presents potential security concerns.**
- least privilege** An authorization policy that gives a user only the amount of access that is essential to that user's work and no more.
- legacy coordinate pairs** The format used for *geospatial* data prior to MongoDB version 2.4. This format stores geospatial data as points on a planar coordinate system (e.g. [*x*, *y*]). See <http://docs.mongodb.org/manualapplications/geospatial-indexes>.
- LineString** A LineString is defined by an array of two or more positions. A closed LineString with four or more positions is called a LinearRing, as described in the GeoJSON LineString specification: <http://geojson.org/geojson-spec.html#linestring>. To use a LineString in MongoDB, see *geospatial-indexes-store-geojson*.
- lock** MongoDB uses locks to ensure concurrency. MongoDB uses both *read locks* and *write locks*. For more information, see *faq-concurrency-locking*.
- LVM** Logical volume manager. LVM is a program that abstracts disk images from physical devices and provides a number of raw disk manipulation and snapshot capabilities useful for system management. For information on LVM and MongoDB, see *lvm-backup-and-restore*.
- map-reduce** A data processing and aggregation paradigm consisting of a “map” phase that selects data and a “reduce” phase that transforms the data. In MongoDB, you can run arbitrary aggregations over data using map-reduce. For map-reduce implementation, see <http://docs.mongodb.org/manualcore/map-reduce>. For all approaches to aggregation, see <http://docs.mongodb.org/manualcore/aggregation>.
- mapping type** A Structure in programming languages that associate keys with values, where keys may nest other pairs of keys and values (e.g. dictionaries, hashes, maps, and associative arrays). The properties of these structures depend on the language specification and implementation. Generally the order of keys in mapping types is arbitrary and not guaranteed.
- master** The database that receives all writes in a conventional master-*slave* replication. In MongoDB, *replica sets* replace master-slave replication for most use cases. For more information on master-slave replication, see <http://docs.mongodb.org/manualcore/master-slave>.
- md5** A hashing algorithm used to efficiently provide reproducible unique strings to identify and *checksum* data. MongoDB uses md5 to identify chunks of data for *GridFS*. See *filemd5* (page 308).
- MIB** Management Information Base. MongoDB uses MIB files to define the type of data tracked by SNMP in the MongoDB Enterprise edition.
- MIME** Multipurpose Internet Mail Extensions. A standard set of type and encoding definitions used to declare the encoding and type of data in multiple data storage, transmission, and email contexts. The *mongofiles* (page 585) tool provides an option to specify a MIME type to describe a file inserted into *GridFS* storage.

mongo The MongoDB shell. The `mongo` (page 527) process starts the MongoDB shell as a daemon connected to either a `mongod` (page 503) or `mongos` (page 518) instance. The shell has a JavaScript interface. See *mongo* (page 527) and *mongo Shell Methods* (page 21).

mongod The MongoDB database server. The `mongod` (page 503) process starts the MongoDB server as a daemon. The MongoDB server manages data requests and formats and manages background operations. See *mongod* (page 503).

MongoDB An open-source document-based database system. “MongoDB” derives from the word “humongous” because of the database’s ability to scale up with ease and hold very large amounts of data. MongoDB stores *documents* in *collections* within databases.

MongoDB Enterprise A commercial edition of MongoDB that includes additional features. For more information, see *MongoDB Subscriptions*¹.

mongos The routing and load balancing process that acts an interface between an application and a MongoDB *sharded cluster*. See *mongos* (page 518).

namespace The canonical name for a collection or index in MongoDB. The namespace is a combination of the database name and the name of the collection or index, like so: `[database-name].[collection-or-index-name]`. All documents belong to a namespace. See *faq-dev-namespace*.

natural order The order that a database stores documents on disk. Typically, the order of documents on disks reflects insertion order, except when a document moves internally because an update operation increases its size. In *capped collections*, insertion order and natural order are identical because documents do not move internally. MongoDB returns documents in forward natural order for a `find()` (page 34) query with no parameters. MongoDB returns documents in reverse natural order for a `find()` (page 34) query *sorted* (page 93) with a parameter of `$natural:-1`. See *\$natural* (page 483).

ObjectId A special 12-byte *BSON* type that guarantees uniqueness within the *collection*. The ObjectId is generated based on timestamp, machine ID, process ID, and a process-local incremental counter. MongoDB uses ObjectId values as the default values for *_id* fields.

operator A keyword beginning with a \$ used to express an update, complex query, or data transformation. For example, `$gt` is the query language’s “greater than” operator. For available operators, see *Operators* (page 373).

oplog A *capped collection* that stores an ordered history of logical writes to a MongoDB database. The oplog is the basic mechanism enabling *replication* in MongoDB. See <http://docs.mongodb.org/manualcore/replica-set-oplog>.

ordered query plan A query plan that returns results in the order consistent with the `sort()` (page 93) order. See *read-operations-query-optimization*.

orphaned document In a sharded cluster, orphaned documents are those documents on a shard that also exist in chunks on other shards as a result of failed migrations or incomplete migration cleanup due to abnormal shut-down. Delete orphaned documents using `cleanupOrphaned` (page 285) to reclaim disk space and reduce confusion.

padding The extra space allocated to document on the disk to prevent moving a document when it grows as the result of `update()` (page 69) operations. See *write-operations-padding-factor*.

padding factor An automatically-calibrated constant used to determine how much extra space MongoDB should allocate per document container on disk. A padding factor of 1 means that MongoDB will allocate only the amount of space needed for the document. A padding factor of 2 means that MongoDB will allocate twice the amount of space required by the document. See *write-operations-padding-factor*.

page fault The event that occurs when a process requests stored data (i.e. a page) from memory that the operating system has moved to disk. See *faq-storage-page-faults*.

¹<https://www.mongodb.com/products/mongodb-subscriptions>

- partition** A distributed system architecture that splits data into ranges. *Sharding* uses partitioning. See *sharding-data-partitioning*.
- passive member** A member of a *replica set* that cannot become primary because its `priority` is 0. See <http://docs.mongodb.org/manualcore/replica-set-priority-0-member>.
- pcap** A packet-capture format used by *mongosniff* (page 581) to record packets captured from network interfaces and display them as human-readable MongoDB operations. See *Options* (page 581).
- PID** A process identifier. UNIX-like systems assign a unique-integer PID to each running process. You can use a PID to inspect a running process and send signals to it. See *proc-file-system*.
- pipe** A communication channel in UNIX-like systems allowing independent processes to send and receive data. In the UNIX shell, piped operations allow users to direct the output of one command into the input of another.
- pipeline** A series of operations in an *aggregation* process. See <http://docs.mongodb.org/manualcore/aggregation>.
- Point** A single coordinate pair as described in the GeoJSON Point specification: <http://geojson.org/geojson-spec.html#point>. To use a Point in MongoDB, see *geospatial-indexes-store-geojson*.
- Polygon** An array of *LinearRing* coordinate arrays, as described in the GeoJSON Polygon specification: <http://geojson.org/geojson-spec.html#polygon>. For Polygons with multiple rings, the first must be the exterior ring and any others must be interior rings or holes.
- MongoDB does not permit the exterior ring to self-intersect. Interior rings must be fully contained within the outer loop and cannot intersect or overlap with each other. See *geospatial-indexes-store-geojson*.
- powerOf2Sizes** A per-collection setting that changes and normalizes the way MongoDB allocates space for each *document*, in an effort to maximize storage reuse and to reduce fragmentation. This is the default for TTL Collections. See *collMod* (page 316) and *usePowerOf2Sizes* (page 316).
- pre-splitting** An operation performed before inserting data that divides the range of possible shard key values into chunks to facilitate easy insertion and high write throughput. In some cases pre-splitting expedites the initial distribution of documents in *sharded cluster* by manually dividing the collection rather than waiting for the MongoDB *balancer* to do so. See <http://docs.mongodb.org/manualtutorial/create-chunks-in-sharded-cluster>.
- primary** In a *replica set*, the primary member is the current *master* instance, which receives all write operations. See *replica-set-primary-member*.
- primary key** A record's unique immutable identifier. In an *RDBMS*, the primary key is typically an integer stored in each row's `id` field. In MongoDB, the `_id` field holds a document's primary key which is usually a BSON *ObjectId*.
- primary shard** The *shard* that holds all the un-sharded collections. See *primary-shard*.
- priority** A configurable value that helps determine which members in a *replica set* are most likely to become *primary*. See *priority*.
- privilege** A combination of specified *resource* and *actions* permitted on the resource. See *privilege*.
- projection** A document given to a *query* that specifies which fields MongoDB returns in the result set. See *projection*. For a list of projection operators, see *Projection Operators* (page 406).
- query** A read request. MongoDB uses a *JSON*-like query language that includes a variety of *query operators* with names that begin with a `$` character. In the *mongo* (page 527) shell, you can issue queries using the `find()` (page 34) and `findOne()` (page 43) methods. See *read-operations-queries*.
- query optimizer** A process that generates query plans. For each query, the optimizer generates a plan that matches the query to the index that will return results as efficiently as possible. The optimizer reuses the query plan each time the query runs. If a collection changes significantly, the optimizer creates a new query plan. See *read-operations-query-optimization*.

query shape A combination of query predicate, sort, and projection specifications.

For the query predicate, only the structure of the predicate, including the field names, are significant; the values in the query predicate are insignificant. As such, a query predicate { type: 'food' } is equivalent to the query predicate { type: 'utensil' } for a query shape.

RDBMS Relational Database Management System. A database management system based on the relational model, typically using *SQL* as the query language.

read lock In the context of a reader-writer lock, a lock that while held allows concurrent readers but no writers. See *faq-concurrency-locking*.

read preference A setting that determines how clients direct read operations. Read preference affects all replica sets, including shards. By default, MongoDB directs reads to *primaries* for *strict consistency*. However, you may also direct reads to secondaries for *eventually consistent* reads. See *Read Preference*.

record size The space allocated for a document including the padding. For more information on padding, see *write-operations-padding-factor* and *compact* (page 313).

recovering A *replica set* member status indicating that a member is not ready to begin normal activities of a secondary or primary. Recovering members are unavailable for reads.

replica pairs The precursor to the MongoDB *replica sets*.

Deprecated since version 1.6.

replica set A cluster of MongoDB servers that implements master-slave replication and automated failover. MongoDB's recommended replication strategy. See <http://docs.mongodb.org/manualreplication>.

replication A feature allowing multiple database servers to share the same data, thereby ensuring redundancy and facilitating load balancing. See <http://docs.mongodb.org/manualreplication>.

replication lag The length of time between the last operation in the *primary's oplog* and the last operation applied to a particular *secondary*. In general, you want to keep replication lag as small as possible. See *Replication Lag*.

resident memory The subset of an application's memory currently stored in physical RAM. Resident memory is a subset of *virtual memory*, which includes memory mapped to physical RAM and to disk.

resource A database, collection, set of collections, or cluster. A *privilege* permits *actions* on a specified resource. See *resource*.

REST An API design pattern centered around the idea of resources and the *CRUD* operations that apply to them. Typically REST is implemented over HTTP. MongoDB provides a simple HTTP REST interface that allows HTTP clients to run commands against the server. See *rest-interface* and *rest-api*.

role A set of privileges that permit *actions* on specified *resources*. Roles assigned to a user determine the user's access to resources and operations. See <http://docs.mongodb.org/manualcore/security-introduction>.

rollback A process that reverts writes operations to ensure the consistency of all replica set members. See *replica-set-rollback*.

secondary A *replica set* member that replicates the contents of the master database. Secondary members may handle read requests, but only the *primary* members can handle write operations. See *replica-set-secondary-members*.

secondary index A database *index* that improves query performance by minimizing the amount of work that the query engine must perform to fulfill a query. See <http://docs.mongodb.org/manualindexes>.

set name The arbitrary name given to a replica set. All members of a replica set must have the same name specified with the `replSetName` setting or the `--replSet` option.

shard A single *mongod* (page 503) instance or *replica set* that stores some portion of a *sharded cluster's* total data set. In production, all shards should be replica sets. See <http://docs.mongodb.org/manualcore/sharded-cluster-shards>.

shard key The field MongoDB uses to distribute documents among members of a *sharded cluster*. See *shard-key*.

sharded cluster The set of nodes comprising a *sharded* MongoDB deployment. A sharded cluster consists of three config processes, one or more replica sets, and one or more *mongos* (page 518) routing processes. See <http://docs.mongodb.org/manualcore/sharded-cluster-components>.

sharding A database architecture that partitions data by key ranges and distributes the data among two or more database instances. Sharding enables horizontal scaling. See <http://docs.mongodb.org/manualsharding>.

shell helper A method in the *mongo* shell that provides a more concise syntax for a *database command* (page 198). Shell helpers improve the general interactive experience. See *Mongo Shell Methods* (page 21).

single-master replication A *replication* topology where only a single database instance accepts writes. Single-master replication ensures consistency and is the replication topology employed by MongoDB. See <http://docs.mongodb.org/manualcore/replica-set-primary>.

slave A read-only database that replicates operations from a *master* database in conventional master/slave replication. In MongoDB, *replica sets* replace master/slave replication for most use cases. However, for information on master/slave replication, see <http://docs.mongodb.org/manualcore/master-slave>.

split The division between *chunks* in a *sharded cluster*. See <http://docs.mongodb.org/manualcore/sharding-chunk->

SQL Structured Query Language (SQL) is a common special-purpose programming language used for interaction with a relational database, including access control, insertions, updates, queries, and deletions. There are some similar elements in the basic SQL syntax supported by different database vendors, but most implementations have their own dialects, data types, and interpretations of proposed SQL standards. Complex SQL is generally not directly portable between major *RDBMS* products. SQL is often used as metonym for relational databases.

SSD Solid State Disk. A high-performance disk drive that uses solid state electronics for persistence, as opposed to the rotating platters and movable read/write heads used by traditional mechanical hard drives.

stale Refers to the amount of time a *secondary* member of a *replica set* trails behind the current state of the *primary's oplog*. If a secondary becomes too stale, it can no longer use replication to catch up to the current state of the primary. See <http://docs.mongodb.org/manualcore/replica-set-oplog> and <http://docs.mongodb.org/manualcore/replica-set-sync> for more information.

standalone An instance of *mongod* (page 503) that is running as a single server and not as part of a *replica set*. To convert a standalone into a replica set, see <http://docs.mongodb.org/manualtutorial/convert-standalone-to-replica-set>.

strict consistency A property of a distributed system requiring that all members always reflect the latest changes to the system. In a database system, this means that any system that can provide data must reflect the latest writes at all times. In MongoDB, reads from a primary have *strict consistency*; reads from secondary members have *eventual consistency*.

sync The *replica set* operation where members replicate data from the *primary*. Sync first occurs when MongoDB creates or restores a member, which is called *initial sync*. Sync then occurs continually to keep the member updated with changes to the replica set's data. See <http://docs.mongodb.org/manualcore/replica-set-sync>.

syslog On UNIX-like systems, a logging process that provides a uniform standard for servers and processes to submit logging information. MongoDB provides an option to send output to the host's syslog system. See *syslogFacility*.

tag A label applied to a replica set member or shard and used by clients to issue data-center-aware operations. For more information on using tags with replica sets and with shards, see the following sections of this manual: *replica-set-read-preference-tag-sets* and *shards-tag-sets*.

TSV A text-based data format consisting of tab-separated values. This format is commonly used to exchange data between relational databases, since the format is well-suited to tabular data. You can import TSV files using

`mongoimport` (page 556).

TTL Stands for “time to live” and represents an expiration time or period for a given piece of information to remain in a cache or other temporary storage before the system deletes it or ages it out. MongoDB has a TTL collection feature. See <http://docs.mongodb.org/manual/tutorial/expire-data>.

unique index An index that enforces uniqueness for a particular field across a single collection. See *index-type-unique*.

unordered query plan A query plan that returns results in an order inconsistent with the `sort()` (page 93) order. See *read-operations-query-optimization*.

upsert An operation that will either update the first document matched by a query or insert a new document if none matches. The new document will have the fields implied by the operation. You perform upserts with the `update()` (page 69) operation. See *Upsert Parameter* (page 70).

virtual memory An application’s working memory, typically residing on both disk and in physical RAM.

WGS84 The default *datum* MongoDB uses to calculate geometry over an Earth-like sphere. MongoDB uses the WGS84 datum for *geospatial* queries on *GeoJSON* objects. See the “EPSG:4326: WGS 84” specification: <http://spatialreference.org/ref/epsg/4326/>.

working set The data that MongoDB uses most often. This data is preferably held in RAM, solid-state drive (SSD), or other fast media. See *faq-working-set*.

write concern Specifies whether a write operation has succeeded. Write concern allows your application to detect insertion errors or unavailable `mongod` (page 503) instances. For *replica sets*, you can configure write concern to confirm replication to a specified number of members. See <http://docs.mongodb.org/manual/core/write-concern>.

write lock A lock on the database for a given writer. When a process writes to the database, it takes an exclusive write lock to prevent other processes from writing or reading. For more information on locks, see <http://docs.mongodb.org/manual/faq/concurrency>.

writeBacks The process within the sharding system that ensures that writes issued to a *shard* that *is not* responsible for the relevant chunk get applied to the proper shard. For related information, see *faq-writebacklisten* and *writeBacksQueued* (page 357).

Release Notes

Always install the latest, stable version of MongoDB. See *release-version-numbers* for more information.

See the following release notes for an account of the changes in major versions. Release notes also include instructions for upgrade.

7.1 Current Stable Release

(2.6-series)

7.1.1 Release Notes for MongoDB 2.6

April 8, 2014

MongoDB 2.6 is now available. Key features include aggregation enhancements, text-search integration, query-engine improvements, a new write-operation protocol, and security enhancements.

MMS 1.4, which includes On-Prem Backup in addition to Monitoring, is now also available. See [MMS 1.4 documentation](#)¹ and the [MMS 1.4 release notes](#)² for more information.

Minor Releases

2.6 Changelog

2.6.1 – Changes

Stability [SERVER-13739](#)³ Repair database failure can delete database files

Build and Packaging

- [SERVER-13287](#)⁴ Addition of debug symbols has doubled compile time
- [SERVER-13563](#)⁵ Upgrading from 2.4.x to 2.6.0 via `yum` clobbers configuration file

¹<https://mms.mongodb.com/help-hosted/v1.4/>

²<https://mms.mongodb.com/help-hosted/v1.4/management/changelog/>

³<https://jira.mongodb.org/browse/SERVER-13739>

⁴<https://jira.mongodb.org/browse/SERVER-13287>

⁵<https://jira.mongodb.org/browse/SERVER-13563>

- [SERVER-13691](#)⁶ yum and apt “stable” repositories contain release candidate 2.6.1-rc0 packages
- [SERVER-13515](#)⁷ Cannot install MongoDB as a service on Windows

Querying

- [SERVER-13066](#)⁸ Negations over multikey fields do not use index
- [SERVER-13495](#)⁹ Concurrent GETMORE and KILLCURSORS operations can cause race condition and server crash
- [SERVER-13503](#)¹⁰ The `$where` (page 391) operator should not be allowed under `$elemMatch` (page 405)
- [SERVER-13537](#)¹¹ Large skip and limit values can cause crash in blocking sort stage
- [SERVER-13557](#)¹² Incorrect negation of `$elemMatch` value in 2.6
- [SERVER-13562](#)¹³ Queries that use tailable cursors do not stream results if `skip()` is applied
- [SERVER-13566](#)¹⁴ Using the `OplogReplay` flag with extra predicates can yield incorrect results
- [SERVER-13611](#)¹⁵ Missing sort order for compound index leads to unnecessary in-memory sort
- [SERVER-13618](#)¹⁶ Optimization for sorted `$in` queries not applied to reverse sort order
- [SERVER-13661](#)¹⁷ Increase the maximum allowed depth of query objects
- [SERVER-13664](#)¹⁸ Query with `$elemMatch` (page 405) using a compound multikey index can generate incorrect results
- [SERVER-13677](#)¹⁹ Query planner should traverse through `$all` while handling `$elemMatch` object predicates
- [SERVER-13766](#)²⁰ Dropping index or collection while `$or` query is yielding triggers fatal assertion

Geospatial

- [SERVER-13666](#)²¹ `$near` (page 394) queries with out-of-bounds points in legacy format can lead to crashes
- [SERVER-13540](#)²² The `geoNear` (page 216) command no longer returns distance in radians for legacy point
- [SERVER-13486](#)²³: The `geoNear` (page 216) command can create too large BSON objects for aggregation.

⁶<https://jira.mongodb.org/browse/SERVER-13691>

⁷<https://jira.mongodb.org/browse/SERVER-13515>

⁸<https://jira.mongodb.org/browse/SERVER-13066>

⁹<https://jira.mongodb.org/browse/SERVER-13495>

¹⁰<https://jira.mongodb.org/browse/SERVER-13503>

¹¹<https://jira.mongodb.org/browse/SERVER-13537>

¹²<https://jira.mongodb.org/browse/SERVER-13557>

¹³<https://jira.mongodb.org/browse/SERVER-13562>

¹⁴<https://jira.mongodb.org/browse/SERVER-13566>

¹⁵<https://jira.mongodb.org/browse/SERVER-13611>

¹⁶<https://jira.mongodb.org/browse/SERVER-13618>

¹⁷<https://jira.mongodb.org/browse/SERVER-13661>

¹⁸<https://jira.mongodb.org/browse/SERVER-13664>

¹⁹<https://jira.mongodb.org/browse/SERVER-13677>

²⁰<https://jira.mongodb.org/browse/SERVER-13766>

²¹<https://jira.mongodb.org/browse/SERVER-13666>

²²<https://jira.mongodb.org/browse/SERVER-13540>

²³<https://jira.mongodb.org/browse/SERVER-13486>

Replication

- [SERVER-13500](#)²⁴ Changing replica set configuration can crash running members
- [SERVER-13589](#)²⁵ Background index builds from a 2.6.0 primary fail to complete on 2.4.x secondaries
- [SERVER-13620](#)²⁶ Replicated data definition commands will fail on secondaries during background index build
- [SERVER-13496](#)²⁷ Creating index with same name but different spec in mixed version replicaset can abort replication

Sharding

- [SERVER-12638](#)²⁸ Initial sharding with hashed shard key can result in duplicate split points
- [SERVER-13518](#)²⁹ The `_id` field is no longer automatically generated by `mongos` (page 518) when missing
- [SERVER-13777](#)³⁰ Migrated ranges waiting for deletion do not report cursors still open

Security

- [SERVER-9358](#)³¹ Log rotation can overwrite previous log files
- [SERVER-13644](#)³² Sensitive credentials in startup options are not redacted and may be exposed
- [SERVER-13441](#)³³ Inconsistent error handling in user management shell helpers

Write Operations

- [SERVER-13466](#)³⁴ Error message in collection creation failure contains incorrect namespace
- [SERVER-13499](#)³⁵ Yield policy for batch-inserts should be the same as for batch-updates/deletes
- [SERVER-13516](#)³⁶ Array updates on documents with more than 128 BSON elements may crash `mongod` (page 503)

2.6.1 – May 5, 2014

- Fix to install MongoDB service on Windows with the `--install` option [SERVER-13515](#)³⁷.
- Allow direct upgrade from 2.4.x to 2.6.0 via `yum` [SERVER-13563](#)³⁸.
- Fix issues with background index builds on secondaries: [SERVER-13589](#)³⁹ and [SERVER-13620](#)⁴⁰.

²⁴<https://jira.mongodb.org/browse/SERVER-13500>

²⁵<https://jira.mongodb.org/browse/SERVER-13589>

²⁶<https://jira.mongodb.org/browse/SERVER-13620>

²⁷<https://jira.mongodb.org/browse/SERVER-13496>

²⁸<https://jira.mongodb.org/browse/SERVER-12638>

²⁹<https://jira.mongodb.org/browse/SERVER-13518>

³⁰<https://jira.mongodb.org/browse/SERVER-13777>

³¹<https://jira.mongodb.org/browse/SERVER-9358>

³²<https://jira.mongodb.org/browse/SERVER-13644>

³³<https://jira.mongodb.org/browse/SERVER-13441>

³⁴<https://jira.mongodb.org/browse/SERVER-13466>

³⁵<https://jira.mongodb.org/browse/SERVER-13499>

³⁶<https://jira.mongodb.org/browse/SERVER-13516>

³⁷<https://jira.mongodb.org/browse/SERVER-13515>

³⁸<https://jira.mongodb.org/browse/SERVER-13563>

³⁹<https://jira.mongodb.org/browse/SERVER-13589>

⁴⁰<https://jira.mongodb.org/browse/SERVER-13620>

- Redact credential information passed as startup options [SERVER-13644](#)⁴¹.
- [2.6.1 Changelog](#) (page 619).
- All issues closed in 2.6.1⁴².

Major Changes

The following changes in MongoDB affect both the standard and Enterprise editions:

Aggregation Enhancements

The aggregation pipeline adds the ability to return result sets of any size, either by returning a cursor or writing the output to a collection. Additionally, the aggregation pipeline supports variables and adds new operations to handle sets and redact data.

- The `db.collection.aggregate()` (page 22) now returns a cursor, which enables the aggregation pipeline to return result sets of any size.
- Aggregation pipelines now support an `explain` operation to aid analysis of aggregation operations.
- Aggregation can now use a more efficient external-disk-based sorting process.
- New pipeline stages:
 - `$out` (page 453) stage to output to a collection.
 - `$redact` (page 441) stage to allow additional control to accessing the data.
- New or modified operators:
 - *set expression operators* (page 461).
 - `$let` (page 471) and `$map` (page 471) operators to allow for the use of variables.
 - `$literal` (page 472) operator and `$size` (page 470) operator.
 - `$cond` (page 475) expression now accepts either an object or an array.

Text Search Integration

Text search is now enabled by default, and the query system, including the aggregation pipeline `$match` (page 440) stage, includes the `$text` (page 387) operator, which resolves text-search queries.

MongoDB 2.6 includes an updated `text` index format and deprecates the `text` (page 235) command.

Insert and Update Improvements

Improvements to the update and insert systems include additional operations and improvements that increase consistency of modified data.

- MongoDB preserves the order of the document fields following write operations *except* for the following cases:
 - The `_id` field is always the first field in the document.
 - Updates that include `renaming` (page 414) of field names may result in the reordering of fields in the document.

⁴¹<https://jira.mongodb.org/browse/SERVER-13644>

⁴²[https://jira.mongodb.org/issues/?jql=fix Version%20%3D%2022.4.10%22%20AND%20project%20%3D%20SERVER](https://jira.mongodb.org/issues/?jql=fix%20Version%20%3D%2022.4.10%22%20AND%20project%20%3D%20SERVER)

- New or enhanced update operators:
 - `$bit` (page 435) operator supports bitwise `xor` operation.
 - `$min` (page 417) and `$max` (page 418) operators that perform conditional update depending on the relative size of the specified value and the current value of a field.
 - `$push` (page 425) operator has enhanced support for the `$sort` (page 431), `$slice` (page 428), and `$each` (page 427) modifiers and supports a new `$position` (page 433) modifier.
 - `$currentDate` (page 419) operator to set the value of a field to the current date.
- The `$mul` (page 413) operator for multiplicative increments for insert and update operations.

See also:

Update Operator Syntax Validation (page 632)

New Write Operation Protocol

A new write protocol integrates write operations with write concerns. The protocol also provides improved support for bulk operations.

MongoDB 2.6 adds the write commands `insert` (page 220), `update` (page 222), and `delete` (page 226), which provide the basis for the improved bulk insert. All officially supported MongoDB drivers support the new write commands.

The `mongo` (page 527) shell now includes methods to perform bulk-write operations. See `Bulk()` (page 128) for more information.

See also:

Write Method Acknowledgements (page 629)

MSI Package for MongoDB Available for Windows

MongoDB now distributes MSI packages for Microsoft Windows. This is the recommended method for MongoDB installation under Windows.

Security Improvements

MongoDB 2.6 enhances support for secure deployments through improved SSL support, x.509-based authentication, an improved authorization system with more granular controls, as well as centralized credential storage, and improved user management tools.

Specifically these changes include:

- A new *authorization model* that provides the ability to create custom *user-defined-roles* and the ability to specify user privileges at a collection-level granularity.
- Global user management, which stores all user and user-defined role data in the `admin` database and provides a new set of commands for managing users and roles.
- `x.509 certificate authentication` for client authentication as well as for internal authentication of sharded and/or replica set cluster members. `x.509` authentication is only available for deployments using SSL.
- Enhanced SSL Support:
 - Rolling upgrades of clusters to use SSL.

- *mongodb-tools-support-ssl* support connections to `mongod` (page 503) and `mongos` (page 518) instances using SSL connections.
- *Prompt for passphrase* by `mongod` (page 503) or `mongos` (page 518) at startup.
- Require the use of strong SSL ciphers, with a minimum 128-bit key length for all connections. The strong-cipher requirement prevents an old or malicious client from forcing use of a weak cipher.
- MongoDB disables the http interface by default, limiting `network exposure`. To enable the interface, see `enabled`.

See also:

New Authorization Model (page 630), *SSL Certificate Hostname Validation* (page 631), and <http://docs.mongodb.org/manualadministration/security-checklist>.

Query Engine Improvements

- MongoDB can now use `index intersection` to fulfill queries supported by more than one index.
- *index-filters* to limit which indexes can become the winning plan for a query.
- *Query Plan Cache Methods* (page 123) methods to view and clear the `query plans` cached by the query optimizer.
- MongoDB can now use `count()` (page 79) with `hint()` (page 86). See `count()` (page 79) for details.

Improvements

Geospatial Enhancements

- *2dsphere indexes version 2*.
- Support for *geojson-multipoint*, *geojson-multilinestring*, *geojson-multipolygon*, and *geojson-geometrycollection*.
- Support for geospatial query clauses in `$or` (page 377) expressions.

See also:

2dsphere Index Version 2 (page 631), *\$maxDistance Changes* (page 634), *Deprecated \$uniqueDocs* (page 634), *Stronger Validation of Geospatial Queries* (page 634)

Index Build Enhancements

- *Background index build* allowed on secondaries. If you initiate a background index build on a *primary*, the secondaries will replicate the index build in the background.
- Automatic rebuild of interrupted index builds after a restart.
 - If a standalone or a primary instance terminates during an index build *without a clean shutdown*, `mongod` (page 503) now restarts the index build when the instance restarts. If the instance shuts down cleanly or if a user kills the index build, the interrupted index builds do not automatically restart upon the restart of the server.
 - If a secondary instance terminates during an index build, the `mongod` (page 503) instance will now restart the interrupted index build when the instance restarts.

To disable this behavior, use the `--noIndexBuildRetry` (page 509) command-line option.

- `ensureIndex()` (page 30) now wraps a new `createIndex` command.
- The `dropDups` option to `ensureIndex()` (page 30) and `createIndex` is deprecated.

See also:

Enforce Index Key Length Limit (page 627)

Enhanced Sharding and Replication Administration

- New `cleanupOrphaned` (page 285) command to remove *orphaned documents* from a shard.
- New `mergeChunks` (page 289) command to combine contiguous chunks located on a single shard. See `mergeChunks` (page 289) and <http://docs.mongodb.org/manual/tutorial/merge-chunks-in-sharded-cluster>.
- New `rs.printReplicationInfo()` (page 166) and `rs.printSlaveReplicationInfo()` (page 166) methods to provide a formatted report of the status of a replica set from the perspective of the primary and the secondary, respectively.

Configuration Options YAML File Format

MongoDB 2.6 supports a YAML-based configuration file format in addition to the previous configuration file format. See <http://docs.mongodb.org/manual/reference/configuration-options> for details.

Operational Changes**Storage**

`usePowerOf2Sizes` (page 316) is now the default allocation strategy for all new collections. The new allocation strategy uses more storage relative to total document size but results in lower levels of storage fragmentation and more predictable storage capacity planning over time.

To use the previous *exact-fit allocation strategy*:

- For a specific collection, use `collMod` (page 316) with `usePowerOf2Sizes` (page 316) set to `false`.
- For all new collections on an entire `mongod` (page 503) instance, set `newCollectionsUsePowerOf2Sizes` to `false`.

See <http://docs.mongodb.org/manual/core/storage> for more information about MongoDB's storage system.

Networking

- Removed upward limit for the `maxIncomingConnections` for `mongod` (page 503) and `mongos` (page 518). Previous versions capped the maximum possible `maxIncomingConnections` setting at 20,000 connections.
- Connection pools for a `mongos` (page 518) instance may be used by multiple MongoDB servers. This can reduce the number of connections needed for high-volume workloads and reduce resource consumption in sharded clusters.
- The C++ driver now monitors *replica set* health with the `isMaster` (page 280) command instead of `replSetGetStatus` (page 273). This allows the C++ driver to support systems that require authentication.

- New `cursor.maxTimeMS()` (page 87) and corresponding `maxTimeMS` option for commands to specify a time limit.

Tool Improvements

- `mongo` (page 527) shell supports a global `/etc/mongorc.js` (page 531).
- All MongoDB *executable files* (page 503) now support the `--quiet` option to suppress all logging output except for error messages.
- `mongoimport` (page 556) uses the input filename, without the file extension if any, as the collection name if run without the `-c` or `--collection` specification.
- `mongoexport` (page 562) can now constrain export data using `--skip` (page 567) and `--limit` (page 567), as well as order the documents in an export using the `--sort` (page 567) option.
- `mongostat` (page 570) can support the use of `--rowcount` (page 572) (`-n` (page ??)) with the `--discover` (page 573) option to produce the specified number of output lines.
- Add strict mode representation for `data_numberlong` for use by `mongoexport` (page 562) and `mongoimport` (page 556).

MongoDB Enterprise Features

The following changes are specific to MongoDB Enterprise Editions:

MongoDB Enterprise for Windows

MongoDB Enterprise for Windows is now available. It includes support for Kerberos, SSL, and SNMP.

MongoDB Enterprise for Windows does **not** include LDAP support for authentication. However, MongoDB Enterprise for Linux supports using LDAP authentication with an ActiveDirectory server.

MongoDB Enterprise for Windows includes OpenSSL version 1.0.1g.

Auditing

MongoDB Enterprise adds auditing capability for `mongod` (page 503) and `mongos` (page 518) instances. See *auditing* for details.

LDAP Support for Authentication

MongoDB Enterprise provides support for proxy authentication of users. This allows administrators to configure a MongoDB cluster to authenticate users by proxying authentication requests to a specified Lightweight Directory Access Protocol (LDAP) service. See <http://docs.mongodb.org/manual/tutorial/configure-ldap-sasl-openldap> and <http://docs.mongodb.org/manual/tutorial/configure-ldap-sasl-activedirectory> for details.

MongoDB Enterprise for Windows does **not** include LDAP support for authentication. However, MongoDB Enterprise for Linux supports using LDAP authentication with an ActiveDirectory server.

MongoDB does **not** support LDAP authentication in mixed sharded cluster deployments that contain both version 2.4 and version 2.6 shards. See *Upgrade MongoDB to 2.6* (page 637) for upgrade instructions.

Expanded SNMP Support

MongoDB Enterprise has greatly expanded its SNMP support to provide SNMP access to nearly the full range of metrics provided by `db.serverStatus()` (page 118).

See also:

SNMP Changes (page 632)

Additional Information

Changes Affecting Compatibility

Compatibility Changes in MongoDB 2.6 The following 2.6 changes can affect the compatibility with older versions of MongoDB. See *Release Notes for MongoDB 2.6* (page 619) for the full list of the 2.6 changes.

Index Changes

Enforce Index Key Length Limit

Description MongoDB 2.6 implements a stronger enforcement of the limit on `index key`.

Creating indexes will error if an index key in an existing document exceeds the limit:

- `db.collection.ensureIndex()` (page 30), `db.collection.reIndex()` (page 62), `compact` (page 313), and `repairDatabase` (page 319) will error and not create the index. Previous versions of MongoDB would create the index but not index such documents.
- Because `db.collection.reIndex()` (page 62), `compact` (page 313), and `repairDatabase` (page 319) drop *all* the indexes from a collection and then recreate them sequentially, the error from the index key limit prevents these operations from rebuilding any remaining indexes for the collection and, in the case of the `repairDatabase` (page 319) command, from continuing with the remainder of the process.

Inserts will error:

- `db.collection.insert()` (page 52) and other operations that perform inserts (e.g. `db.collection.save()` (page 66) and `db.collection.update()` (page 69) with `upsert` that result in inserts) will fail to insert if the new document has an indexed field whose corresponding index entry exceeds the limit. Previous versions of MongoDB would insert but not index such documents.
- `mongorestore` (page 543) and `mongoimport` (page 556) will fail to insert if the new document has an indexed field whose corresponding index entry exceeds the limit.

Updates will error:

- `db.collection.update()` (page 69) and `db.collection.save()` (page 66) operations on an indexed field will error if the updated value causes the index entry to exceed the limit.
- If an existing document contains an indexed field whose index entry exceeds the limit, updates on other fields that result in the relocation of a document on disk will error.

Chunk Migration will fail:

- Migrations will fail for a chunk that has a document with an indexed field whose index entry exceeds the limit.

- If left unfixed, the chunk will repeatedly fail migration, effectively ceasing chunk balancing for that collection. Or, if chunk splits occur in response to the migration failures, this response would lead to unnecessarily large number of chunks and an overly large config databases.

Secondary members of replica sets will warn:

- Secondaries will continue to replicate documents with an indexed field whose corresponding index entry exceeds the limit on initial sync but will print warnings in the logs.
- Secondaries allow index build and rebuild operations on a collection that contains an indexed field whose corresponding index entry exceeds the limit but with warnings in the logs.
- With *mixed version* replica sets where the secondaries are version 2.6 and the primary is version 2.4, secondaries will replicate documents inserted or updated on the 2.4 primary, but will print error messages in the log if the documents contain an indexed field whose corresponding index entry exceeds the limit.

Solution Run `db.upgradeCheckAllDBs()` (page 121) to find current keys that violate this limit and correct as appropriate. Preferably, run the test before upgrading; i.e. connect the 2.6 `mongo` (page 527) shell to your MongoDB 2.4 database and run the method.

If you have an existing data set and want to disable the default index key length validation so that you can upgrade before resolving these indexing issues, use the `failIndexKeyTooLong` parameter.

Index Specifications Validate Field Names

Description In MongoDB 2.6, create and re-index operations fail when the index key refers to an empty field, e.g. `"a..b"` : 1 or the field name starts with a dollar sign (\$).

- `db.collection.ensureIndex()` (page 30) will not create a new index with an invalid or empty key name.
- `db.collection.reIndex()` (page 62), `compact` (page 313), and `repairDatabase` (page 319) will error if an index exists with an invalid or empty key name.
- Chunk migration will fail if an index exists with an invalid or empty key name.

Previous versions of MongoDB allow the index.

Solution Run `db.upgradeCheckAllDBs()` (page 121) to find current keys that violate this limit and correct as appropriate. Preferably, run the test before upgrading; i.e. connect the 2.6 `mongo` (page 527) shell to your MongoDB 2.4 database and run the method.

ensureIndex and Existing Indexes

Description `db.collection.ensureIndex()` (page 30) now errors:

- if you try to create an existing index but with different options; e.g. in the following example, the second `db.collection.ensureIndex()` (page 30) will error.

```
db.mycollection.ensureIndex( { x: 1 } )
db.mycollection.ensureIndex( { x: 1 }, { unique: 1 } )
```

- if you specify an index name that already exists but the key specifications differ; e.g. in the following example, the second `db.collection.ensureIndex()` (page 30) will error.

```
db.mycollection.ensureIndex( { a: 1 }, { name: "myIdx" } )
db.mycollection.ensureIndex( { z: 1 }, { name: "myIdx" } )
```

Previous versions did not create the index but did not error.

Write Method Acknowledgements

Description The `mongo` (page 527) shell write methods `db.collection.insert()` (page 52), `db.collection.update()` (page 69), `db.collection.save()` (page 66) and `db.collection.remove()` (page 62) now integrate the write concern directly into the method rather than with a separate `getLastError` (page 235) command to provide *safe writes* whether run interactively in the `mongo` (page 527) shell or non-interactively in a script. In previous versions, these methods exhibited a “fire-and-forget” behavior.⁴³

- Existing scripts for the `mongo` (page 527) shell that used these methods will now observe safe writes which take **longer** than the previous “fire-and-forget” behavior.
- The write methods now return a `WriteResult` (page 188) object that contains the results of the operation, including any write errors and write concern errors, and obviates the need to call `getLastError` (page 235) command to get the status of the results. See `db.collection.insert()` (page 52), `db.collection.update()` (page 69), `db.collection.save()` (page 66) and `db.collection.remove()` (page 62) for details.

Solution Scripts that used these `mongo` (page 527) shell methods for bulk write operations with “fire-and-forget” behavior should use the `Bulk()` (page 128) methods.

For example, instead of:

```
for (var i = 1; i <= 1000000; i++) {
  db.test.insert( { x : i } );
}
```

In MongoDB 2.6, replace with `Bulk()` (page 128) operation:

```
var bulk = db.test.initializeUnorderedBulkOp();

for (var i = 1; i <= 1000000; i++) {
  bulk.insert( { x : i } );
}

bulk.execute( { w: 1 } );
```

Bulk method returns a `BulkWriteResult` (page 189) object that contains the result of the operation.

See also:

New Write Operation Protocol (page 623), `Bulk()` (page 128), `Bulk.execute()` (page 137), `db.collection.initializeUnorderedBulkOp()` (page 51), `db.collection.initializeOrderedBulkOp()` (page 51)

`db.collection.aggregate()` Change

Description The `db.collection.aggregate()` (page 22) method in the `mongo` (page 527) shell defaults to returning a cursor to the results set. This change enables the aggregation pipeline to return result sets of any size and requires cursor iteration to access the result set. For example:

```
var myCursor = db.orders.aggregate( [
  {
    $group: {
      _id: "$cust_id",
      total: { $sum: "$price" }
    }
  }
]
```

⁴³ In previous versions, when using the `mongo` (page 527) shell interactively, the `mongo` (page 527) shell automatically called the `getLastError` (page 235) command after a write method to provide “safe writes”. Scripts, however, would observe “fire-and-forget” behavior in previous versions unless the scripts included an **explicit** call to the `getLastError` (page 235) command after a write method.

```
    }  
  ] );  
  
myCursor.forEach( function(x) { printjson (x); } );
```

Previous versions returned a single document with a field `results` that contained an array of the result set, subject to the *BSON Document size* (page 604) limit. Accessing the result set in the previous versions of MongoDB required accessing the `results` field and iterating the array. For example:

```
var returnedDoc = db.orders.aggregate( [  
  {  
    $group: {  
      _id: "$cust_id",  
      total: { $sum: "$price" }  
    }  
  }  
] );  
  
var myArray = returnedDoc.result; // access the result field  
  
myArray.forEach( function(x) { printjson (x); } );
```

Solution Update scripts that currently expect `db.collection.aggregate()` (page 22) to return a document with a `results` array to handle cursors instead.

See also:

Aggregation Enhancements (page 622), `db.collection.aggregate()` (page 22),

Write Concern Validation

Description Specifying a write concern that includes `j: true` to a `mongod` (page 503) or `mongos` (page 518) instance running with `--nojournal` (page 511) option now errors. Previous versions would ignore the `j: true`.

Solution Either remove the `j: true` specification from the write concern when issued against a `mongod` (page 503) or `mongos` (page 518) instance with `--nojournal` (page 511) or run `mongod` (page 503) or `mongos` (page 518) with journaling.

Security Changes

New Authorization Model

Description MongoDB 2.6 *authorization model* changes how MongoDB stores and manages user privilege information:

- Before the upgrade, MongoDB 2.6 requires at least one user in the admin database.
- MongoDB versions using older models cannot create/modify users or create user-defined roles.

Solution Ensure that at least one user exists in the admin database. If no user exists in the admin database, add a user. Then upgrade to MongoDB 2.6. Finally, upgrade the user privilege model. See *Upgrade MongoDB to 2.6* (page 637).

Important: Before upgrading the authorization model, you should first upgrade MongoDB binaries to 2.6. For sharded clusters, ensure that **all** cluster components are 2.6. If there are users in any database, be sure you have at least one user in the `admin` database **before** upgrading the MongoDB binaries.

See also:

[Security Improvements](#) (page 623)

SSL Certificate Hostname Validation

Description The SSL certificate validation now checks the Common Name (CN) and the Subject Alternative Name (SAN) fields to ensure that either the CN or one of the SAN entries matches the hostname of the server. As a result, if you currently use SSL and *neither* the CN nor any of the SAN entries of your current SSL certificates match the hostnames, upgrading to version 2.6 will cause the SSL connections to fail.

Solution To allow for the continued use of these certificates, MongoDB provides the `allowInvalidCertificates` setting. The setting is available for:

- `mongod` (page 503) and `mongos` (page 518) to bypass the validation of SSL certificates on other servers in the cluster.
- `mongo` (page 527) shell, *MongoDB tools that support SSL*, and the C++ driver to bypass the validation of server certificates.

When using the `allowInvalidCertificates` setting, MongoDB logs as a warning the use of the invalid certificates.

Warning: The `allowInvalidCertificates` setting bypasses the other certificate validation, such as checks for expiration and valid signatures.

2dsphere Index Version 2

Description MongoDB 2.6 introduces a version 2 of the `2dsphere` index. If a document lacks a `2dsphere` index field (or the field is `null` or an empty array), MongoDB does not add an entry for the document to the `2dsphere` index. For inserts, MongoDB inserts the document but does not add to the `2dsphere` index.

Previous version would not insert documents where the `2dsphere` index field is a `null` or an empty array. For documents that lack the `2dsphere` index field, previous versions would insert and index the document.

Solution To revert to old behavior, create the `2dsphere` index with `{ "2dsphereIndexVersion" : 1 }` to create a version 1 index. However, version 1 index cannot use the new GeoJSON geometries.

See also:

[2dsphere-v2](#)

Log Messages**Timestamp Format Change**

Description Each message now starts with the timestamp formatted in `iso8601-local`, i.e. `YYYY-MM-DDTHH:mm:ss.mmm<+/-Offset>`. For example, `2014-03-04T20:13:38.944-0500`. Previous versions used `ctime` format.

Solution MongoDB adds a new option `--timestampFormat` (page 520) which supports timestamp format in `ctime` (page 520), `iso8601-utc` (page 520), and `iso8601-local` (page 520) (new default).

Package Configuration Changes

Default `bindIp` for RPM/DEB Packages

Description In the official MongoDB packages in RPM (Red Hat, CentOS, Fedora Linux, and derivatives) and DEB (Debian, Ubuntu, and derivatives), the default `bindIp` value attaches MongoDB components to the localhost interface *only*. These packages set this default in the default configuration file (i.e. `/etc/mongodb.conf`.)

Solution If you use one of these packages and have *not* modified the default `/etc/mongodb.conf` file, you will need to set `bindIp` before or during the upgrade.

There is no default `bindIp` setting in any other official MongoDB packages.

SNMP Changes

Description

- The IANA enterprise identifier for MongoDB changed from 37601 to 34601.
- MongoDB changed the MIB field name `globalopcounts` to `globalOpcounts`.

Solution

- Users of SNMP monitoring must modify their SNMP configuration (i.e. MIB) from 37601 to 34601.
- Update references to `globalopcounts` to `globalOpcounts`.

Remove Method Signature Change

Description `db.collection.remove()` (page 62) requires a query document as a parameter. In previous versions, the method invocation without a query document deleted all documents in a collection.

Solution For existing `db.collection.remove()` (page 62) invocations without a query document, modify the invocations to include an empty document `db.collection.remove({})`.

Update Operator Syntax Validation

Description

- *Update operators (e.g. \$set)* (page 412) must specify a non-empty operand expression. For example, the following expression is now invalid:

```
{ $set: { } }
```

- *Update operators (e.g. \$set)* (page 412) cannot repeat in the update statement. For example, the following expression is invalid:

```
{ $set: { a: 5 }, $set: { b: 5 } }
```

Updates Enforce Field Name Restrictions

Description

- Updates cannot use *update operators (e.g. \$set)* (page 412) to target fields with empty field names (i.e. "").
- Updates no longer support saving field names that contain a dot (.) or a field name that starts with a dollar sign (\$).

Solution

- For existing documents that currently have fields with empty names "", replace the whole document. See `db.collection.update()` (page 69) and `db.collection.save()` (page 66) for details on replacing an existing document.

- `Unset` (page 417) or `rename` (page 414) existing fields with names that contain a dot (.) or that start with a dollar sign (\$). Run `db.upgradeCheckAllDBs()` (page 121) to find fields whose names contain a dot or starts with a dollar sign.

See *New Write Operation Protocol* (page 623) for the changes to the write operation protocol, and *Insert and Update Improvements* (page 622) for the changes to the insert and update operations. Also consider the documentation of the *Restrictions on Field Names* (page 609).

Query and Sort Changes

Enforce Field Name Restrictions

Description Queries cannot specify conditions on fields with names that start with a dollar sign (\$).

Solution `Unset` (page 417) or `rename` (page 414) existing fields whose names start with a dollar sign (\$). Run `db.upgradeCheckAllDBs()` (page 121) to find fields whose names start with a dollar sign.

Sparse Index and Incomplete Results

Description If a `sparse` index results in an incomplete result set for queries and sort operations, MongoDB will not use that index unless a `hint()` (page 86) explicitly specifies the index.

For example, the query `{ x: { $exists: false } }` will no longer use a sparse index on the `x` field, unless explicitly hinted.

Solution To override the behavior to use the sparse index and return incomplete results, explicitly specify the index with a `hint()` (page 86).

See *sparse-index-incomplete-results* for an example that details the new behavior.

sort () Specification Values

Description The `sort()` (page 93) method **only** accepts the following values for the sort keys:

- 1 to specify ascending order for a field,
- -1 to specify descending order for a field, or
- `$meta` (page 410) expression to specify sort by the text search score.

Any other value will result in an error.

Previous versions also accepted either `true` or `false` for ascending.

Solution Update sort key values that use `true` or `false` to 1.

skip () and _id Queries

Description Equality match on the `_id` field obeys `skip()` (page 92).

Previous versions ignored `skip()` (page 92) when performing an equality match on the `_id` field.

explain () Retains Query Plan Cache

Description `explain()` (page 80) no longer clears the query plans cached for that *query shape*.

In previous versions, `explain()` (page 80) would have the side effect of clearing the query plan cache for that query shape.

See also:

[Query Plan Cache Methods](#) (page 123)

Geospatial Changes

`$maxDistance` Changes

Description

- For `$near` (page 394) queries on GeoJSON data, if the queries specify a `$maxDistance` (page 397), `$maxDistance` (page 397) must be inside of the `$near` (page 394) document.

In previous version, `$maxDistance` (page 397) could be either inside or outside the `$near` (page 394) document.

- `$maxDistance` (page 397) must be a positive value.

Solution

- Update any existing `$near` (page 394) queries on GeoJSON data that currently have the `$maxDistance` (page 397) outside the `$near` (page 394) document
- Update any existing queries where `$maxDistance` (page 397) is a negative value.

Deprecated `$uniqueDocs`

Description MongoDB 2.6 deprecates `$uniqueDocs` (page 400), and geospatial queries no longer return duplicated results when a document matches the query multiple times.

Stronger Validation of Geospatial Queries

Description MongoDB 2.6 enforces a stronger validation of geospatial queries, such as validating the options or GeoJSON specifications, and errors if the geospatial query is invalid. Previous versions allowed/ignored invalid options.

Query Operator Changes

`$not` Query Behavior Changes

Description

- Queries with `$not` (page 379) expressions on an indexed field now match:
 - Documents that are missing the indexed field. Previous versions would not return these documents using the index.
 - Documents whose indexed field value is a different type than that of the specified value. Previous versions would not return these documents using the index.

For example, if a collection `orders` contains the following documents:

```
{ _id: 1, status: "A", cust_id: "123", price: 40 }
{ _id: 2, status: "A", cust_id: "xyz", price: "N/A" }
{ _id: 3, status: "D", cust_id: "xyz" }
```

If the collection has an index on the `price` field:

```
db.orders.ensureIndex( { price: 1 } )
```

The following query uses the index to search for documents where `price` is not greater than or equal to 50:

```
db.orders.find( { price: { $not: { $gte: 50 } } } )
```

In 2.6, the query returns the following documents:

```
{ "_id" : 3, "status" : "D", "cust_id" : "xyz" }
{ "_id" : 1, "status" : "A", "cust_id" : "123", "price" : 40 }
{ "_id" : 2, "status" : "A", "cust_id" : "xyz", "price" : "N/A" }
```

In previous versions, indexed plans would only return matching documents where the type of the field matches the type of the query predicate:

```
{ "_id" : 1, "status" : "A", "cust_id" : "123", "price" : 40 }
```

If using a collection scan, previous versions would return the same results as those in 2.6.

- MongoDB 2.6 allows chaining of `$not` (page 379) expressions.

null Comparison Queries

Description

- `$lt` (page 375) and `$gt` (page 373) comparisons to `null` no longer match documents that are missing the field.
- `null` equality conditions on array elements (e.g. `"a.b": null`) no longer match document missing the nested field `a.b` (e.g. `a: [2, 3]`).
- `null` equality queries (i.e. `field: null`) now match fields with values undefined.

\$all Operator Behavior Change

Description

- The `$all` (page 402) operator is now equivalent to an `$and` (page 378) operation of the specified values. This change in behavior can allow for more matches than previous versions when passed an array of a single nested array (e.g. `[["A"]]`). When passed an array of a nested array, `$all` (page 402) can now match documents where the field contains the nested array as an element (e.g. `field: [["A"], ...]`), or the field equals the nested array (e.g. `field: ["A", "B"]`). Earlier version could only match documents where the field contains the nested array.
- The `$all` (page 402) operator returns no match if the array field contains nested arrays (e.g. `field: ["a", ["b"]]`) and `$all` (page 402) on the nested field is the element of the nested array (e.g. `"field.1": { $all: ["b"] }`). Previous versions would return a match.

\$mod Operator Enforces Strict Syntax

Description The `$mod` (page 384) operator now only accepts an array with exactly two elements, and errors when passed an array with fewer or more elements. See *Not Enough Elements Error* (page 385) and *Too Many Elements Error* (page 386) for details.

In previous versions, if passed an array with one element, the `$mod` (page 384) operator uses 0 as the second element, and if passed an array with more than two elements, the `$mod` (page 384) ignores all but the first two elements. Previous versions do return an error when passed an empty array.

Solution Ensure that the array passed to `$mod` (page 384) contains exactly two elements:

- If the array contains the a single element, add 0 as the second element.
- If the array contains more than two elements, remove the extra elements.

`$where` Must Be Top-Level

Description `$where` (page 391) expressions can now only be at top level and cannot be nested within another expression, such as `$elemMatch` (page 405).

Solution Update existing queries that nest `$where` (page 391).

`$exists` and `notablescan` If the MongoDB server has disabled collection scans, i.e. `notablescan`, then `$exists` (page 381) queries that have no *indexed solution* will error.

MinKey and MaxKey Queries

Description Equality match for either `MinKey` or `MaxKey` no longer match documents missing the field.

Text Search Compatibility MongoDB does not support the use of the `$text` (page 387) query operator in mixed sharded cluster deployments that contain both version 2.4 and version 2.6 shards. See *Upgrade MongoDB to 2.6* (page 637) for upgrade instructions.

Replica Set/Sharded Cluster Validation

Shard Name Checks on Metadata Refresh

Description For sharded clusters, MongoDB 2.6 disallows a shard from refreshing the metadata if the shard name has not been explicitly set.

For mixed sharded cluster deployments that contain both version 2.4 and version 2.6 shards, this change can cause errors when migrating chunks **from** version 2.4 shards **to** version 2.6 shards if the shard name is unknown to the version 2.6 shards. MongoDB does not support migrations in mixed sharded cluster deployments.

Solution Upgrade all components of the cluster to 2.6. See *Upgrade MongoDB to 2.6* (page 637).

Replica Set Vote Configuration Validation

Description MongoDB now deprecates giving any *replica set* member more than a single vote. During configuration, `local.system.replset.members[n].votes` should only have a value of 1 for voting members and 0 for non-voting members. MongoDB treats values other than 1 or 0 as a value of 1 and produces a warning message.

Solution Update `local.system.replset.members[n].votes` with values other than 1 or 0 to 1 or 0 as appropriate.

Other Resources

- [All backwards incompatible changes \(JIRA\)](#)⁴⁴.
- *Release Notes for MongoDB 2.6* (page 619).

⁴⁴<https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fixVersion+in+%28%222.5.0%22%2C+%222.5.1%22%2C+%222.6.0-rc2%22%29+AND+%22Backward+Breaking%22+in+%28+Rarely+%2C+sometimes%2C+yes+%29+ORDER+BY+votes+DESC%2C+issuetype>

Upgrade Process

Upgrade MongoDB to 2.6 In the general case, the upgrade from MongoDB 2.4 to 2.6 is a binary-compatible “drop-in” upgrade: shut down the `mongod` (page 503) instances and replace them with `mongod` (page 503) instances running 2.6. **However**, before you attempt any upgrade, familiarize yourself with the content of this document, particularly the *Upgrade Recommendations and Checklists* (page 637), the procedure for *upgrading sharded clusters* (page 638), and the considerations for *reverting to 2.4 after running 2.6* (page 642).

Upgrade Recommendations and Checklists

When upgrading, consider the following:

Upgrade Requirements To upgrade an existing MongoDB deployment to 2.6, you must be running 2.4. If you're running a version of MongoDB before 2.4, you *must* upgrade to 2.4 before upgrading to 2.6. See [Upgrade MongoDB to 2.4](#) (page 659) for the procedure to upgrade from 2.2 to 2.4.

If you use MMS Backup, ensure that you're running *at least* version `v20131216.1` of the Backup agent before upgrading.

Preparedness Before upgrading MongoDB always test your application in a staging environment before deploying the upgrade to your production environment.

To begin the upgrade procedure, connect a 2.6 `mongo` (page 527) shell to your MongoDB 2.4 `mongos` (page 518) or `mongod` (page 503) and run the `db.upgradeCheckAllDBs()` (page 121) to check your data set for compatibility. This is a preliminary automated check. Assess and resolve all issues identified by `db.upgradeCheckAllDBs()` (page 121).

Some changes in MongoDB 2.6 require manual checks and intervention. See *Compatibility Changes in MongoDB 2.6* (page 627) for an explanation of these changes. Resolve all incompatibilities in your deployment before continuing.

Authentication MongoDB 2.6 includes significant changes to the authorization model, which requires changes to the way that MongoDB stores users' credentials. As a result, in addition to upgrading MongoDB processes, if your deployment uses authentication and authorization, after upgrading all MongoDB process to 2.6 you must also upgrade the authorization model.

Before beginning the upgrade process for a deployment that uses authentication and authorization:

- Ensure that at least one user exists in the `admin` database.
- If your application performs CRUD operations on the `<database>.system.users` collection or uses a `db.addUser()` (page 143)-like method, then you **must** upgrade those drivers (i.e. client libraries) **before** `mongod` (page 503) or `mongos` (page 518) instances.

⁴⁵ <https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fixVersion+in+%282022.5.0%2C+%22.5.1%22%2C+%22.6.0%22%2C+%22.6.0-rc2%22%2C+%22.6.0-rc3%22%29+AND+%22Backward+Breaking%22+in+%28Rarely+%2C+sometimes%2C+yes+%29+ORDER+BY+vc>

- You must fully complete the upgrade procedure for *all* MongoDB processes before upgrading the authorization model.

See *Upgrade User Authorization Data to 2.6 Format* (page 640) for a complete discussion of the upgrade procedure for the authorization model including additional requirements and procedures.

Downgrade Limitations Once upgraded to MongoDB 2.6, you **cannot** downgrade to **any** version earlier than MongoDB 2.4. If you created `text` or `2dsphere` indexes while running 2.6, you can only downgrade to MongoDB 2.4.10 or later.

Upgrade MongoDB Processes

Upgrade Standalone mongod Instance to MongoDB 2.6 The following steps outline the procedure to upgrade a standalone `mongod` (page 503) from version 2.4 to 2.6. To upgrade from version 2.2 to 2.6, *upgrade to version 2.4* (page 659) *first*, and then follow the procedure to upgrade from 2.4 to 2.6.

1. Download binaries of the latest release in the 2.6 series from the [MongoDB Download Page](http://docs.mongodb.org/manual/installation)⁴⁶. See <http://docs.mongodb.org/manual/installation> for more information.
2. Shut down your `mongod` (page 503) instance. Replace the existing binary with the 2.6 `mongod` (page 503) binary and restart `mongod` (page 503).

Upgrade a Replica Set to 2.6 The following steps outline the procedure to upgrade a replica set from MongoDB 2.4 to MongoDB 2.6. To upgrade from MongoDB 2.2 to 2.6, *upgrade all members of the replica set to version 2.4* (page 659) *first*, and then follow the procedure to upgrade from MongoDB 2.4 to 2.6.

You can upgrade from MongoDB 2.4 to 2.6 using a “rolling” upgrade to minimize downtime by upgrading the members individually while the other members are available:

Step 1: Upgrade secondary members of the replica set. Upgrade the *secondary* members of the set one at a time by shutting down the `mongod` (page 503) and replacing the 2.4 binary with the 2.6 binary. After upgrading a `mongod` (page 503) instance, wait for the member to recover to `SECONDARY` state before upgrading the next instance. To check the member’s state, issue `rs.status()` (page 168) in the `mongo` (page 527) shell.

Step 2: Step down the replica set primary. Use `rs.stepDown()` (page 168) in the `mongo` (page 527) shell to step down the *primary* and force the set to *failover*. `rs.stepDown()` (page 168) expedites the failover procedure and is preferable to shutting down the primary directly.

Step 3: Upgrade the primary. When `rs.status()` (page 168) shows that the primary has stepped down and another member has assumed `PRIMARY` state, shut down the previous primary and replace the `mongod` (page 503) binary with the 2.6 binary and start the new instance.

Replica set failover is not instant but will render the set unavailable accept writes until the failover process completes. Typically this takes 30 seconds or more: schedule the upgrade procedure during a scheduled maintenance window.

Upgrade a Sharded Cluster to 2.6 Only upgrade sharded clusters to 2.6 if **all** members of the cluster are currently running instances of 2.4. The only supported upgrade path for sharded clusters running 2.2 is via 2.4. The upgrade process checks all components of the cluster and will produce warnings if any component is running version 2.2.

⁴⁶<http://www.mongodb.org/downloads>

Considerations The upgrade process does not require any downtime. However, while you upgrade the sharded cluster, ensure that clients do not make changes to the collection meta-data. For example, during the upgrade, do **not** do any of the following:

- `sh.enableSharding()` (page 175)
- `sh.shardCollection()` (page 178)
- `sh.addShard()` (page 173)
- `db.createCollection()` (page 102)
- `db.collection.drop()` (page 29)
- `db.dropDatabase()` (page 108)
- any operation that creates a database
- any other operation that modifies the cluster metadata in any way. See <http://docs.mongodb.org/manualreference/sharding> for a complete list of sharding commands. Note, however, that not all commands on the <http://docs.mongodb.org/manualreference/sharding> page modifies the cluster meta-data.

Upgrade Sharded Clusters *Optional but Recommended.* As a precaution, take a backup of the `config` database *before* upgrading the sharded cluster.

Step 1: Disable the Balancer. Turn off the *balancer* in the sharded cluster, as described in *sharding-balancing-disable-temporarily*.

Step 2: Upgrade the cluster's meta data. Start a single 2.6 `mongos` (page 518) instance with the `configDB` pointing to the cluster's config servers and with the `--upgrade` (page 522) option.

To run a `mongos` (page 518) with the `--upgrade` (page 522) option, you can upgrade an existing `mongos` (page 518) instance to 2.6, or if you need to avoid reconfiguring a production `mongos` (page 518) instance, you can use a new 2.6 `mongos` (page 518) that can reach all the config servers.

To upgrade the meta data, run:

```
mongos --configdb <config servers> --upgrade
```

You can include the `--logpath` (page 520) option to output the log messages to a file instead of the standard output.

The `mongos` (page 518) will exit upon completion of the `--upgrade` process.

The upgrade will prevent any chunk moves or splits from occurring during the upgrade process. If the data files have many sharded collections or if failed processes hold stale locks, acquiring the locks for all collections can take seconds or minutes. Watch the log for progress updates.

Step 3: Ensure `mongos --upgrade` process completes successfully. The `mongos` (page 518) will exit upon completion of the meta data upgrade process. If successful, the process will log the following messages:

```
upgrade of config server to v5 successful
Config database is at version v5
```

After a successful upgrade, restart the `mongos` (page 518) instance. If `mongos` (page 518) fails to start, check the log for more information.

If the `mongos` (page 518) instance loses its connection to the config servers during the upgrade or if the upgrade is otherwise unsuccessful, you may always safely retry the upgrade.

Step 4: Upgrade the remaining mongos instances to v2.6. Upgrade and restart **without** the `--upgrade` (page 522) option the other `mongos` (page 518) instances in the sharded cluster. After upgrading all the `mongos` (page 518), see *Complete Sharded Cluster Upgrade* (page 640) for information on upgrading the other cluster components.

Complete Sharded Cluster Upgrade After you have successfully upgraded *all* `mongos` (page 518) instances, you can upgrade the other instances in your MongoDB deployment.

Warning: Do not upgrade `mongod` (page 503) instances until after you have upgraded *all* `mongos` (page 518) instances.

While the balancer is still disabled, upgrade the components of your sharded cluster in the following order:

- Upgrade all 3 `mongod` (page 503) config server instances, leaving the *first* system in the `mongos` `--configdb` argument to upgrade *last*.
- Upgrade each shard, one at a time, upgrading the `mongod` (page 503) secondaries before running `replSetStepDown` (page 277) and upgrading the primary of each shard.

When this process is complete, *re-enable the balancer*.

Upgrade Procedure Once upgraded to MongoDB 2.6, you **cannot** downgrade to **any** version earlier than MongoDB 2.4. If you have `text` indexes, you can only downgrade to MongoDB 2.4.10 or later.

Except as described on this page, moving between 2.4 and 2.6 is a drop-in replacement:

Step 1: Stop the existing mongod instance. For example, on Linux, run 2.4 `mongod` (page 503) with the `--shutdown` (page 512) option as follows:

```
mongod --dbpath /var/mongod/data --shutdown
```

Replace `/var/mongod/data` with your MongoDB `dbPath`. See also the *terminate-mongod-processes* for alternate methods of stopping a `mongod` (page 503) instance.

Step 2: Start the new mongod instance. Ensure you start the 2.6 `mongod` (page 503) with the same `dbPath`:

```
mongod --dbpath /var/mongod/data
```

Replace `/var/mongod/data` with your MongoDB `dbPath`.

Upgrade User Authorization Data to 2.6 Format MongoDB 2.6 includes significant changes to the authorization model, which requires changes to the way that MongoDB stores users' credentials. As a result, in addition to upgrading MongoDB processes, if your deployment uses authentication and authorization, after upgrading all MongoDB process to 2.6 you must also upgrade the authorization model.

Considerations

Complete all other Upgrade Requirements Before upgrading the authorization model, you should first upgrade MongoDB binaries to 2.6. For sharded clusters, ensure that **all** cluster components are 2.6. If there are users in any database, be sure you have at least one user in the `admin` database **before** upgrading the MongoDB binaries.

Timing Because downgrades are more difficult after you upgrade the user authorization model, once you upgrade the MongoDB binaries to version 2.6, allow your MongoDB deployment to run a day or two **without** upgrading the user authorization model.

This allows 2.6 some time to “burn in” and decreases the likelihood of downgrades occurring after the user privilege model upgrade. The user authentication and access control will continue to work as it did in 2.4, **but** it will be impossible to create or modify users or to use user-defined roles until you run the authorization upgrade.

If you decide to upgrade the user authorization model immediately instead of waiting the recommended “burn in” period, then for sharded clusters, you must wait at least 10 seconds after upgrading the sharded clusters to run the authorization upgrade script.

Replica Sets For a replica set, it is only necessary to run the upgrade process on the *primary* as the changes will automatically replicate to the secondaries.

Sharded Clusters For a sharded cluster, connect to a *mongos* (page 518) and run the upgrade procedure to upgrade the cluster’s authorization data. By default, the procedure will upgrade the authorization data of the shards as well.

To override this behavior, run the upgrade command with the additional parameter `upgradeShards: false`. If you choose to override, you must run the upgrade procedure on the *mongos* (page 518) first, and then run the procedure on the *primary* members of each shard.

For a sharded cluster, do **not** run the upgrade process directly against the `config` servers. Instead, perform the upgrade process using one *mongos* (page 518) instance to interact with the config database.

Requirements To upgrade the authorization model, you must have a user in the `admin` database with the role `userAdminAnyDatabase`.

Procedure

Step 1: Connect to MongoDB instance. Connect and authenticate to the *mongod* (page 503) instance for a single deployment or a *mongos* (page 518) for a sharded cluster as an `admin` database user with the role `userAdminAnyDatabase`.

Step 2: Upgrade authorization schema. Use the `authSchemaUpgrade` (page 250) command in the `admin` database to update the user data using the *mongo* (page 527) shell.

Run `authSchemaUpgrade` command.

```
res = db.getSiblingDB("admin").runCommand({authSchemaUpgrade: 1});
print(tojson(res));
```

In case of error, you may safely rerun the `authSchemaUpgrade` (page 250) command.

Sharded cluster `authSchemaUpgrade` consideration. For a sharded cluster, `authSchemaUpgrade` (page 250) will upgrade the authorization data of the shards as well and the upgrade is complete. You can, however, override this behavior by including `upgradeShards: false` in the command, as in the following example:

```
res = db.getSiblingDB("admin").runCommand({authSchemaUpgrade: 1,
upgradeShards: false});
```

If you override the behavior, after running `authSchemaUpgrade` (page 250) on a `mongos` (page 518) instance, you will need to connect to the primary for each shard and repeat the upgrade process after upgrading on the `mongos` (page 518).

Result All users in a 2.6 system are stored in the `admin.system.users` (page 601) collection. To manipulate these users, use the *user management methods* (page 141).

The upgrade procedure copies the version 2.4 `admin.system.users` collection to `admin.system.backup_users`.

The upgrade procedure leaves the version 2.4 `<database>.system.users` collection(s) intact.

Downgrade MongoDB from 2.6 Before you attempt any downgrade, familiarize yourself with the content of this document, particularly the *Downgrade Recommendations and Checklist* (page 642) and the procedure for *downgrading sharded clusters* (page 645).

Downgrade Recommendations and Checklist When downgrading, consider the following:

Downgrade Path Once upgraded to MongoDB 2.6, you **cannot** downgrade to **any** version earlier than MongoDB 2.4. If you created `text` or `2dsphere` indexes while running 2.6, you can only downgrade to MongoDB 2.4.10 or later.

Preparedness

- *Remove or downgrade version 2 text indexes* (page 644) before downgrading MongoDB 2.6 to 2.4.
- *Remove or downgrade version 2 2dsphere indexes* (page 644) before downgrading MongoDB 2.6 to 2.4.
- *Downgrade 2.6 User Authorization Model* (page 642). If you have upgraded to the 2.6 user authorization model, you must downgrade the user model to 2.4 before downgrading MongoDB 2.6 to 2.4.

Procedures Follow the downgrade procedures:

- To downgrade sharded clusters, see *Downgrade a 2.6 Sharded Cluster* (page 645).
- To downgrade replica sets, see *Downgrade a 2.6 Replica Set* (page 645).
- To downgrade a standalone MongoDB instance, see *Downgrade 2.6 Standalone mongod Instance* (page 644).

Downgrade 2.6 User Authorization Model If you have upgraded to the 2.6 user authorization model, you **must first** downgrade the user authorization model to 2.4 **before** before downgrading MongoDB 2.6 to 2.4.

Considerations

- For a replica set, it is only necessary to run the downgrade process on the *primary* as the changes will automatically replicate to the secondaries.
- For sharded clusters, although the procedure lists the downgrade of the cluster's authorization data first, you may downgrade the authorization data of the cluster or shards first.
- You *must* have the `admin.system.backup_users` and `admin.system.new_users` collections created during the upgrade process.
- **Important.** The downgrade process returns the user data to its state prior to upgrading to 2.6 authorization model. Any changes made to the user/role data using the 2.6 users model will be lost.

Procedure The following downgrade procedure requires `<database>.system.users` collections used in version 2.4. to be intact for non-admin databases:

Step 1: Connect and authenticate to MongoDB instance. Connect and authenticate to the `mongod` (page 503) instance for a single deployment or a `mongos` (page 518) for a sharded cluster as a user with the following privileges:

```
{ resource: { db: "admin", collection: "system.new_users" }, actions: [ "find", "insert", "update" ]
{ resource: { db: "admin", collection: "system.backup_users" }, actions: [ "find" ] }
{ resource: { db: "admin", collection: "system.users" }, actions: [ "find", "remove", "insert" ] }
{ resource: { db: "admin", collection: "system.version" }, actions: [ "find", "update" ] }
```

Step 2: Create backup of 2.6 admin.system.users collection. Copy all documents in the `admin.system.users` collection to the `admin.system.new_users` collection:

```
db.getSiblingDB("admin").system.users.find().forEach( function(userDoc) {
    db.getSiblingDB("admin").system.new_users.save( userDoc );
}
);
```

Step 3: Update the version document for the authSchema.

```
db.getSiblingDB("admin").system.version.update(
    { _id: "authSchema" },
    { $set: { currentVersion: 2 } }
);
```

Step 4: Remove existing documents from the admin.system.users collection.

```
db.getSiblingDB("admin").system.users.remove( {} )
```

Step 5: Copy documents from the admin.system.backup_users collection. Copy all documents from the `admin.system.backup_users`, created during the 2.6 upgrade, to `admin.system.users`.

```
db.getSiblingDB("admin").system.backup_users.find().forEach(
    function (userDoc) {
        db.getSiblingDB("admin").system.users.insert( userDoc );
    }
);
```

Step 6: Update the version document for the authSchema.

```
db.getSiblingDB("admin").system.version.update(
    { _id: "authSchema" },
    { $set: { currentVersion: 1 } }
);
```

For a sharded cluster, repeat the downgrade process by connecting to the *primary* replica set member for each shard.

Note: The cluster's `mongos` (page 518) instances will fail to detect the authorization model downgrade until the user cache is refreshed. You can run `invalidateUserCache` (page 272) on each `mongos` (page 518) instance to refresh immediately, or you can wait until the cache is refreshed automatically at the end of the user cache invalidation interval.

Result The downgrade process returns the user data to its state prior to upgrading to 2.6 authorization model. Any changes made to the user/role data using the 2.6 users model will be lost.

Downgrade Updated Indexes

Text Index Version Check If you have *version 2* text indexes (i.e. the default version for text indexes in MongoDB 2.6), drop the *version 2* text indexes before downgrading MongoDB. After the downgrade, enable text search and recreate the dropped text indexes.

To determine the version of your text indexes, run `db.collection.getIndexes()` (page 45) to view index specifications. For text indexes, the method returns the version information in the field `textIndexVersion`. For example, the following shows that the text index on the `quotes` collection is version 2.

```
{
  "v" : 1,
  "key" : {
    "_fts" : "text",
    "_ftsx" : 1
  },
  "name" : "quote_text_translation.quote_text",
  "ns" : "test.quotes",
  "weights" : {
    "quote" : 1,
    "translation.quote" : 1
  },
  "default_language" : "english",
  "language_override" : "language",
  "textIndexVersion" : 2
}
```

2dsphere Index Version Check If you have *version 2* 2dsphere indexes (i.e. the default version for 2dsphere indexes in MongoDB 2.6), drop the *version 2* 2dsphere indexes before downgrading MongoDB. After the downgrade, recreate the 2dsphere indexes.

To determine the version of your 2dsphere indexes, run `db.collection.getIndexes()` (page 45) to view index specifications. For 2dsphere indexes, the method returns the version information in the field `2dsphereIndexVersion`. For example, the following shows that the 2dsphere index on the `locations` collection is version 2.

```
{
  "v" : 1,
  "key" : {
    "geo" : "2dsphere"
  },
  "name" : "geo_2dsphere",
  "ns" : "test.locations",
  "sparse" : true,
  "2dsphereIndexVersion" : 2
}
```

Downgrade MongoDB Processes

Downgrade 2.6 Standalone mongod Instance The following steps outline the procedure to downgrade a standalone `mongod` (page 503) from version 2.6 to 2.4.

1. Download binaries of the latest release in the 2.4 series from the [MongoDB Download Page](http://docs.mongodb.org/manual/installation)⁴⁷. See <http://docs.mongodb.org/manual/installation> for more information.
2. Shut down your `mongod` (page 503) instance. Replace the existing binary with the 2.4 `mongod` (page 503) binary and restart `mongod` (page 503).

Downgrade a 2.6 Replica Set The following steps outline a “rolling” downgrade process for the replica set. The “rolling” downgrade process minimizes downtime by downgrading the members individually while the other members are available:

Step 1: Downgrade each secondary member, one at a time. For each *secondary* in a replica set:

Replace and restart secondary `mongod` instances. First, shut down the `mongod` (page 503), then replace these binaries with the 2.4 binary and restart `mongod` (page 503). See *terminate-mongod-processes* for instructions on safely terminating `mongod` (page 503) processes.

Allow secondary to recover. Wait for the member to recover to `SECONDARY` state before upgrading the next secondary.

To check the member’s state, use the `rs.status()` (page 168) method in the `mongo` (page 527) shell.

Step 2: Step down the primary. Use `rs.stepDown()` (page 168) in the `mongo` (page 527) shell to step down the *primary* and force the normal *failover* procedure.

```
rs.stepDown()
```

`rs.stepDown()` (page 168) expedites the failover procedure and is preferable to shutting down the primary directly.

Step 3: Replace and restart former primary `mongod`. When `rs.status()` (page 168) shows that the primary has stepped down and another member has assumed `PRIMARY` state, shut down the previous primary and replace the `mongod` (page 503) binary with the 2.4 binary and start the new instance.

Replica set failover is not instant but will render the set unavailable to writes and interrupt reads until the failover process completes. Typically this takes 10 seconds or more. You may wish to plan the downgrade during a predetermined maintenance window.

Downgrade a 2.6 Sharded Cluster

Requirements While the downgrade is in progress, you cannot make changes to the collection meta-data. For example, during the downgrade, do **not** do any of the following:

- `sh.enableSharding()` (page 175)
- `sh.shardCollection()` (page 178)
- `sh.addShard()` (page 173)
- `db.createCollection()` (page 102)
- `db.collection.drop()` (page 29)
- `db.dropDatabase()` (page 108)

⁴⁷<http://www.mongodb.org/downloads>

- any operation that creates a database
- any other operation that modifies the cluster meta-data in any way. See <http://docs.mongodb.org/manualreference/sharding> for a complete list of sharding commands. Note, however, that not all commands on the <http://docs.mongodb.org/manualreference/sharding> page modifies the cluster meta-data.

Procedure The downgrade procedure for a sharded cluster reverses the order of the upgrade procedure.

1. Turn off the *balancer* in the sharded cluster, as described in *sharding-balancing-disable-temporarily*.
2. Downgrade each shard, one at a time. For each shard,
 - (a) Downgrade the `mongod` (page 503) secondaries *before* downgrading the primary.
 - (b) To downgrade the primary, run `replSetStepDown` (page 277) and downgrade.
3. Downgrade all 3 `mongod` (page 503) config server instances, leaving the *first* system in the `mongos --configdb` argument to downgrade *last*.
4. Downgrade and restart each `mongos` (page 518), one at a time. The downgrade process is a binary drop-in replacement.
5. Turn on the balancer, as described in *sharding-balancing-enable*.

Downgrade Procedure Once upgraded to MongoDB 2.6, you **cannot** downgrade to **any** version earlier than MongoDB 2.4. If you have `text` indexes, you can only downgrade to MongoDB 2.4.10 or later.

Except as described on this page, moving between 2.4 and 2.6 is a drop-in replacement:

Step 1: Stop the existing `mongod` instance. For example, on Linux, run 2.6 `mongod` (page 503) with the `--shutdown` (page 512) option as follows:

```
mongod --dbpath /var/mongod/data --shutdown
```

Replace `/var/mongod/data` with your MongoDB `dbPath`. See also the *terminate-mongod-processes* for alternate methods of stopping a `mongod` (page 503) instance.

Step 2: Start the new `mongod` instance. Ensure you start the 2.4 `mongod` (page 503) with the same `dbPath`:

```
mongod --dbpath /var/mongod/data
```

Replace `/var/mongod/data` with your MongoDB `dbPath`.

See *Upgrade MongoDB to 2.6* (page 637) for full upgrade instructions.

Download

To download MongoDB 2.6, go to the [downloads](http://www.mongodb.org/downloads) page⁴⁸.

⁴⁸<http://www.mongodb.org/downloads>

Other Resources

- All JIRA issues resolved in 2.6⁴⁹.
- All Third Party License Notices⁵⁰.

7.2 Previous Stable Releases

7.2.1 Release Notes for MongoDB 2.4

March 19, 2013

MongoDB 2.4 includes enhanced geospatial support, switch to V8 JavaScript engine, security enhancements, and text search (beta) and hashed index.

Minor Releases

2.4 Changelog

2.4.10 - Changes

- Indexes: Fixed issue that can cause index corruption when building indexes concurrently ([SERVER-12990](#)⁵¹)
- Indexes: Fixed issue that can cause index corruption when shutting down secondary node during index build ([SERVER-12956](#)⁵²)
- Indexes: Mongod now recognizes incompatible “future” text and geo index versions and exits gracefully ([SERVER-12914](#)⁵³)
- Indexes: Fixed issue that can cause secondaries to fail replication when building the same index multiple times concurrently ([SERVER-12662](#)⁵⁴)
- Indexes: Fixed issue that can cause index corruption on the tenth index in a collection if the index build fails ([SERVER-12481](#)⁵⁵)
- Indexes: Introduced versioning for text and geo indexes to ensure backwards compatibility ([SERVER-12175](#)⁵⁶)
- Indexes: Disallowed building indexes on the system.indexes collection, which can lead to initial sync failure on secondaries ([SERVER-10231](#)⁵⁷)
- Sharding: Avoid frequent immediate balancer retries when config servers are out of sync ([SERVER-12908](#)⁵⁸)
- Sharding: Add indexes to locks collection on config servers to avoid long queries in case of large numbers of collections ([SERVER-12548](#)⁵⁹)

⁴⁹<https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fixVersion+in+%28%222.5.0%22%2C+%222.5.1%22%2C+%222.6.0-rc2%22%2C+%222.6.0-rc3%22%29>

⁵⁰<https://github.com/mongodb/mongo/blob/v2.6/distsrc/THIRD-PARTY-NOTICES>

⁵¹<https://jira.mongodb.org/browse/SERVER-12990>

⁵²<https://jira.mongodb.org/browse/SERVER-12956>

⁵³<https://jira.mongodb.org/browse/SERVER-12914>

⁵⁴<https://jira.mongodb.org/browse/SERVER-12662>

⁵⁵<https://jira.mongodb.org/browse/SERVER-12481>

⁵⁶<https://jira.mongodb.org/browse/SERVER-12175>

⁵⁷<https://jira.mongodb.org/browse/SERVER-10231>

⁵⁸<https://jira.mongodb.org/browse/SERVER-12908>

⁵⁹<https://jira.mongodb.org/browse/SERVER-12548>

- Sharding: Fixed issue that can corrupt the config metadata cache when sharding collections concurrently ([SERVER-12515](#)⁶⁰)
- Sharding: Don't move chunks created on collections with a hashed shard key if the collection already contains data ([SERVER-9259](#)⁶¹)
- Replication: Fixed issue where node appears to be down in a replica set during a compact operation ([SERVER-12264](#)⁶²)
- Replication: Fixed issue that could cause delays in elections when a node is not vetoing an election ([SERVER-12170](#)⁶³)
- Replication: Step down all primaries if multiple primaries are detected in replica set to ensure correct election result ([SERVER-10793](#)⁶⁴)
- Replication: Upon clock skew detection, secondaries will switch to sync directly from the primary to avoid sync cycles ([SERVER-8375](#)⁶⁵)
- Runtime: The SIGXCPU signal is now caught and mongod writes a log message and exits gracefully ([SERVER-12034](#)⁶⁶)
- Runtime: Fixed issue where mongod fails to start on Linux when /sys/dev/block directory is not readable ([SERVER-9248](#)⁶⁷)
- Windows: No longer zero-fill newly allocated files on systems other than Windows 7 or Windows Server 2008 R2 ([SERVER-8480](#)⁶⁸)
- GridFS: Chunk size is decreased to 255 KB (from 256 KB) to avoid overhead with usePowerOf2Sizes option ([SERVER-13331](#)⁶⁹)
- SNMP: Fixed MIB file validation under smilint ([SERVER-12487](#)⁷⁰)
- Shell: Fixed issue in V8 memory allocation that could cause long-running shell commands to crash ([SERVER-11871](#)⁷¹)
- Shell: Fixed memory leak in the md5sumFile shell utility method ([SERVER-11560](#)⁷²)

Previous Releases

- All 2.4.9 improvements⁷³.
- All 2.4.8 improvements⁷⁴.
- All 2.4.7 improvements⁷⁵.
- All 2.4.6 improvements⁷⁶.

⁶⁰<https://jira.mongodb.org/browse/SERVER-12515>

⁶¹<https://jira.mongodb.org/browse/SERVER-9259>

⁶²<https://jira.mongodb.org/browse/SERVER-12264>

⁶³<https://jira.mongodb.org/browse/SERVER-12170>

⁶⁴<https://jira.mongodb.org/browse/SERVER-10793>

⁶⁵<https://jira.mongodb.org/browse/SERVER-8375>

⁶⁶<https://jira.mongodb.org/browse/SERVER-12034>

⁶⁷<https://jira.mongodb.org/browse/SERVER-9248>

⁶⁸<https://jira.mongodb.org/browse/SERVER-8480>

⁶⁹<https://jira.mongodb.org/browse/SERVER-13331>

⁷⁰<https://jira.mongodb.org/browse/SERVER-12487>

⁷¹<https://jira.mongodb.org/browse/SERVER-11871>

⁷²<https://jira.mongodb.org/browse/SERVER-11560>

⁷³<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.9%22%20AND%20project%20%3D%20SERVER>

⁷⁴<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.8%22%20AND%20project%20%3D%20SERVER>

⁷⁵<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.7%22%20AND%20project%20%3D%20SERVER>

⁷⁶<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.6%22%20AND%20project%20%3D%20SERVER>

- All 2.4.5 improvements⁷⁷.
- All 2.4.4 improvements⁷⁸.
- All 2.4.3 improvements⁷⁹.
- All 2.4.2 improvements⁸⁰.
- All 2.4.1 improvements⁸¹.

2.4.10 – April 4, 2014

- Performs fast file allocation on Windows when available [SERVER-8480](#)⁸².
- Start elections if more than one primary is detected [SERVER-10793](#)⁸³.
- Changes to allow safe downgrading from v2.6 to v2.4 [SERVER-12914](#)⁸⁴, [SERVER-12175](#)⁸⁵.
- Fixes for edge cases in index creation [SERVER-12481](#)⁸⁶, [SERVER-12956](#)⁸⁷.
- *2.4.10 Changelog* (page 647).
- All 2.4.10 improvements⁸⁸.

2.4.9 – January 10, 2014

- Fix for instances where `mongos` (page 518) incorrectly reports a successful write [SERVER-12146](#)⁸⁹.
- Make non-primary read preferences consistent with `slaveOK` versioning logic [SERVER-11971](#)⁹⁰.
- Allow new sharded cluster connections to read from secondaries when primary is down [SERVER-7246](#)⁹¹.
- All 2.4.9 improvements⁹².

2.4.8 – November 1, 2013

- Increase future compatibility for 2.6 authorization features [SERVER-11478](#)⁹³.
- Fix dbhash cache issue for config servers [SERVER-11421](#)⁹⁴.
- All 2.4.8 improvements⁹⁵.

⁷⁷<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.5%22%20AND%20project%20%3D%20SERVER>

⁷⁸<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.4%22%20AND%20project%20%3D%20SERVER>

⁷⁹<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.3%22%20AND%20project%20%3D%20SERVER>

⁸⁰<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.2%22%20AND%20project%20%3D%20SERVER>

⁸¹<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.1%22%20AND%20project%20%3D%20SERVER>

⁸²<https://jira.mongodb.org/browse/SERVER-8480>

⁸³<https://jira.mongodb.org/browse/SERVER-10793>

⁸⁴<https://jira.mongodb.org/browse/SERVER-12914>

⁸⁵<https://jira.mongodb.org/browse/SERVER-12175>

⁸⁶<https://jira.mongodb.org/browse/SERVER-12481>

⁸⁷<https://jira.mongodb.org/browse/SERVER-12956>

⁸⁸<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.10%22%20AND%20project%20%3D%20SERVER>

⁸⁹<https://jira.mongodb.org/browse/SERVER-12146>

⁹⁰<https://jira.mongodb.org/browse/SERVER-11971>

⁹¹<https://jira.mongodb.org/browse/SERVER-7246>

⁹²<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.9%22%20AND%20project%20%3D%20SERVER>

⁹³<https://jira.mongodb.org/browse/SERVER-11478>

⁹⁴<https://jira.mongodb.org/browse/SERVER-11421>

⁹⁵<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.8%22%20AND%20project%20%3D%20SERVER>

2.4.7 – October 21, 2013

- Fixed over-aggressive caching of V8 Isolates [SERVER-10596](#)⁹⁶.
- Removed extraneous initial count during mapReduce [SERVER-9907](#)⁹⁷.
- Cache results of dbhash command [SERVER-11021](#)⁹⁸.
- Fixed memory leak in aggregation [SERVER-10554](#)⁹⁹.
- All 2.4.7 improvements¹⁰⁰.

2.4.6 – August 20, 2013

- Fix for possible loss of documents during the chunk migration process if a document in the chunk is very large [SERVER-10478](#)¹⁰¹.
- Fix for C++ client shutdown issues [SERVER-8891](#)¹⁰².
- Improved replication robustness in presence of high network latency [SERVER-10085](#)¹⁰³.
- Improved Solaris support [SERVER-9832](#)¹⁰⁴, [SERVER-9786](#)¹⁰⁵, and [SERVER-7080](#)¹⁰⁶.
- All 2.4.6 improvements¹⁰⁷.

2.4.5 – July 3, 2013

- Fix for CVE-2013-4650 Improperly grant user system privileges on databases other than local [SERVER-9983](#)¹⁰⁸.
- Fix for CVE-2013-3969 Remotely triggered segmentation fault in Javascript engine [SERVER-9878](#)¹⁰⁹.
- Fix to prevent identical background indexes from being built [SERVER-9856](#)¹¹⁰.
- Config server performance improvements [SERVER-9864](#)¹¹¹ and [SERVER-5442](#)¹¹².
- Improved initial sync resilience to network failure [SERVER-9853](#)¹¹³.
- All 2.4.5 improvements¹¹⁴.

⁹⁶<https://jira.mongodb.org/browse/SERVER-10596>

⁹⁷<https://jira.mongodb.org/browse/SERVER-9907>

⁹⁸<https://jira.mongodb.org/browse/SERVER-11021>

⁹⁹<https://jira.mongodb.org/browse/SERVER-10554>

¹⁰⁰[https://jira.mongodb.org/issues/?jql=fix Version%20%3D%2022.4.7%22%20AND%20project%20%3D%20SERVER](https://jira.mongodb.org/issues/?jql=fix%20Version%20%3D%2022.4.7%22%20AND%20project%20%3D%20SERVER)

¹⁰¹<https://jira.mongodb.org/browse/SERVER-10478>

¹⁰²<https://jira.mongodb.org/browse/SERVER-8891>

¹⁰³<https://jira.mongodb.org/browse/SERVER-10085>

¹⁰⁴<https://jira.mongodb.org/browse/SERVER-9832>

¹⁰⁵<https://jira.mongodb.org/browse/SERVER-9786>

¹⁰⁶<https://jira.mongodb.org/browse/SERVER-7080>

¹⁰⁷[https://jira.mongodb.org/issues/?jql=fix Version%20%3D%2022.4.6%22%20AND%20project%20%3D%20SERVER](https://jira.mongodb.org/issues/?jql=fix%20Version%20%3D%2022.4.6%22%20AND%20project%20%3D%20SERVER)

¹⁰⁸<https://jira.mongodb.org/browse/SERVER-9983>

¹⁰⁹<https://jira.mongodb.org/browse/SERVER-9878>

¹¹⁰<https://jira.mongodb.org/browse/SERVER-9856>

¹¹¹<https://jira.mongodb.org/browse/SERVER-9864>

¹¹²<https://jira.mongodb.org/browse/SERVER-5442>

¹¹³<https://jira.mongodb.org/browse/SERVER-9853>

¹¹⁴[https://jira.mongodb.org/issues/?jql=fix Version%20%3D%2022.4.5%22%20AND%20project%20%3D%20SERVER](https://jira.mongodb.org/issues/?jql=fix%20Version%20%3D%2022.4.5%22%20AND%20project%20%3D%20SERVER)

2.4.4 – June 4, 2013

- Performance fix for Windows version [SERVER-9721](#)¹¹⁵
- Fix for config upgrade failure [SERVER-9661](#)¹¹⁶.
- Migration to Cyrus SASL library for MongoDB Enterprise [SERVER-8813](#)¹¹⁷.
- All 2.4.4 improvements¹¹⁸.

2.4.3 – April 23, 2013

- Fix for mongo shell ignoring modified object's `_id` field [SERVER-9385](#)¹¹⁹.
- Fix for race condition in log rotation [SERVER-4739](#)¹²⁰.
- Fix for `copydb` command with authorization in a sharded cluster [SERVER-9093](#)¹²¹.
- All 2.4.3 improvements¹²².

2.4.2 – April 17, 2013

- Several V8 memory leak and performance fixes [SERVER-9267](#)¹²³ and [SERVER-9230](#)¹²⁴.
- Fix for upgrading sharded clusters [SERVER-9125](#)¹²⁵.
- Fix for high volume connection crash [SERVER-9014](#)¹²⁶.
- All 2.4.2 improvements¹²⁷

2.4.1 – April 17, 2013

- Fix for losing index changes during initial sync [SERVER-9087](#)¹²⁸
- All 2.4.1 improvements¹²⁹.

Major New Features

The following changes in MongoDB affect both standard and Enterprise editions:

¹¹⁵<https://jira.mongodb.org/browse/SERVER-9721>

¹¹⁶<https://jira.mongodb.org/browse/SERVER-9661>

¹¹⁷<https://jira.mongodb.org/browse/SERVER-8813>

¹¹⁸<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.4%22%20AND%20project%20%3D%20SERVER>

¹¹⁹<https://jira.mongodb.org/browse/SERVER-9385>

¹²⁰<https://jira.mongodb.org/browse/SERVER-4739>

¹²¹<https://jira.mongodb.org/browse/SERVER-9093>

¹²²<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.3%22%20AND%20project%20%3D%20SERVER>

¹²³<https://jira.mongodb.org/browse/SERVER-9267>

¹²⁴<https://jira.mongodb.org/browse/SERVER-9230>

¹²⁵<https://jira.mongodb.org/browse/SERVER-9125>

¹²⁶<https://jira.mongodb.org/browse/SERVER-9014>

¹²⁷<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.2%22%20AND%20project%20%3D%20SERVER>

¹²⁸<https://jira.mongodb.org/browse/SERVER-9087>

¹²⁹<https://jira.mongodb.org/issues/?jql=fixVersion%20%3D%2022.4.1%22%20AND%20project%20%3D%20SERVER>

Text Search

Add support for text search of content in MongoDB databases as a *beta* feature. See <http://docs.mongodb.org/manualcore/index-text> for more information.

Geospatial Support Enhancements

- Add new `2dsphere` index. The new index supports [GeoJSON](#)¹³⁰ objects `Point`, `LineString`, and `Polygon`. See <http://docs.mongodb.org/manualcore/2dsphere> and <http://docs.mongodb.org/manualapplications/geospatial-indexes>.
- Introduce operators `$geometry` (page 397), `$geoWithin` (page 392) and `$geoIntersects` (page 393) to work with the GeoJSON data.

Hashed Index

Add new *hashed index* to index documents using hashes of field values. When used to index a shard key, the hashed index ensures an evenly distributed shard key. See also *sharding-hashed-sharding*.

Improvements to the Aggregation Framework

- Improve support for geospatial queries. See the `$geoWithin` (page 392) operator and the `$geoNear` (page 451) pipeline stage.
- Improve sort efficiency when the `$sort` (page 449) stage immediately precedes a `$limit` (page 445) in the pipeline.
- Add new operators `$millisecond` (page 475) and `$concat` (page 465) and modify how `$min` (page 456) operator processes `null` values.

Changes to Update Operators

- Add new `$setOnInsert` (page 415) operator for use with `upsert` (page 69) .
- Enhance functionality of the `$push` (page 425) operator, supporting its use with the `$each` (page 427), the `$sort` (page 431), and the `$slice` (page 428) modifiers.

Additional Limitations for Map-Reduce and \$where Operations

The `mapReduce` (page 208) command, `group` (page 204) command, and the `$where` (page 391) operator expressions cannot access certain global functions or properties, such as `db`, that are available in the `mongo` (page 527) shell. See the individual command or operator for details.

Improvements to `serverStatus` Command

Provide additional metrics and customization for the `serverStatus` (page 347) command. See `db.serverStatus()` (page 118) and `serverStatus` (page 347) for more information.

¹³⁰<http://geojson.org/geojson-spec.html>

Security Enhancements

- Introduce a role-based access control system [/reference/user-privileges](http://docs.mongodb.org/manual/reference/user-privileges)¹³¹ using new <http://docs.mongodb.org/manual/reference/privilege-documents>.
- Enforce uniqueness of the user in user privilege documents per database. Previous versions of MongoDB did not enforce this requirement, and existing databases may have duplicates.
- Support encrypted connections using SSL certificates signed by a Certificate Authority. See <http://docs.mongodb.org/manual/tutorial/configure-ssl>.

For more information on security and risk management strategies, see MongoDB Security Practices and Procedures.

Performance Improvements

V8 JavaScript Engine

JavaScript Changes in MongoDB 2.4 Consider the following impacts of *V8 JavaScript Engine* (page 653) in MongoDB 2.4:

Tip

Use the new `interpreterVersion()` method in the `mongo` (page 527) shell and the `javascriptEngine` (page 325) field in the output of `db.serverBuildInfo()` (page 118) to determine which JavaScript engine a MongoDB binary uses.

Improved Concurrency Previously, MongoDB operations that required the JavaScript interpreter had to acquire a lock, and a single `mongod` (page 503) could only run a single JavaScript operation at a time. The switch to V8 improves concurrency by permitting multiple JavaScript operations to run at the same time.

Modernized JavaScript Implementation (ES5) The 5th edition of *ECMAScript*¹³², abbreviated as ES5, adds many new language features, including:

- `standardized JSON`¹³³,
- `strict mode`¹³⁴,
- `function.bind()`¹³⁵,
- `array extensions`¹³⁶, and
- `getters and setters`.

With V8, MongoDB supports the ES5 implementation of Javascript with the following exceptions.

Note: The following features do not work as expected on documents **returned from MongoDB queries**:

- `Object.seal()` throws an exception on documents returned from MongoDB queries.
- `Object.freeze()` throws an exception on documents returned from MongoDB queries.

¹³¹<http://docs.mongodb.org/v2.4/reference/user-privileges>

¹³²<http://www.ecma-international.org/publications/standards/Ecma-262.htm>

¹³³<http://www.ecma-international.org/ecma-262/5.1/#sec-15.12.1>

¹³⁴<http://www.ecma-international.org/ecma-262/5.1/#sec-4.2.2>

¹³⁵<http://www.ecma-international.org/ecma-262/5.1/#sec-15.3.4.5>

¹³⁶<http://www.ecma-international.org/ecma-262/5.1/#sec-15.4.4.16>

- `Object.preventExtensions()` incorrectly allows the addition of new properties on documents returned from MongoDB queries.
- `enumerable` properties, when added to documents returned from MongoDB queries, are not saved during write operations.

See [SERVER-8216¹³⁷](#), [SERVER-8223¹³⁸](#), [SERVER-8215¹³⁹](#), and [SERVER-8214¹⁴⁰](#) for more information.

For objects that have not been returned from MongoDB queries, the features work as expected.

Removed Non-Standard SpiderMonkey Features V8 does **not** support the following *non-standard SpiderMonkey*¹⁴¹ JavaScript extensions, previously supported by MongoDB's use of SpiderMonkey as its JavaScript engine.

E4X Extensions V8 does not support the *non-standard E4X*¹⁴² extensions. E4X provides a native `XML`¹⁴³ object to the JavaScript language and adds the syntax for embedding literal XML documents in JavaScript code.

You need to use alternative XML processing if you used any of the following constructors/methods:

- `XML()`
- `Namespace()`
- `QName()`
- `XMLList()`
- `isXMLName()`

Destructuring Assignment V8 does not support the non-standard destructuring assignments. Destructuring assignment “extract[s] data from arrays or objects using a syntax that mirrors the construction of array and object literals.” - [Mozilla docs](#)¹⁴⁴

Example

The following destructuring assignment is **invalid** with V8 and throws a `SyntaxError`:

```
original = [4, 8, 15];
var [b, ,c] = a; // <== destructuring assignment
print(b) // 4
print(c) // 15
```

Iterator(), StopIteration(), and Generators V8 does not support `Iterator()`, `StopIteration()`, and `generators`¹⁴⁵.

InternalError() V8 does not support `InternalError()`. Use `Error()` instead.

¹³⁷<https://jira.mongodb.org/browse/SERVER-8216>

¹³⁸<https://jira.mongodb.org/browse/SERVER-8223>

¹³⁹<https://jira.mongodb.org/browse/SERVER-8215>

¹⁴⁰<https://jira.mongodb.org/browse/SERVER-8214>

¹⁴¹<https://developer.mozilla.org/en-US/docs/SpiderMonkey>

¹⁴²<https://developer.mozilla.org/en-US/docs/E4X>

¹⁴³https://developer.mozilla.org/en-US/docs/E4X/Processing_XML_with_E4X

¹⁴⁴[https://developer.mozilla.org/en-US/docs/JavaScript/New_in_JavaScript/1.7#Destructuring_assignment_\(Merge_into_own_page.2Fsection\)](https://developer.mozilla.org/en-US/docs/JavaScript/New_in_JavaScript/1.7#Destructuring_assignment_(Merge_into_own_page.2Fsection))

¹⁴⁵https://developer.mozilla.org/en-US/docs/JavaScript/Guide/Iterators_and_Generators

for each...in Construct V8 does not support the use of `for each...in`¹⁴⁶ construct. Use `for (var x in y)` construct instead.

Example

The following `for each (var x in y)` construct is **invalid** with V8:

```
var o = { name: 'MongoDB', version: 2.4 };

for each (var value in o) {
  print(value);
}
```

Instead, in version 2.4, you can use the `for (var x in y)` construct:

```
var o = { name: 'MongoDB', version: 2.4 };

for (var prop in o) {
  var value = o[prop];
  print(value);
}
```

You can also use the array *instance* method `forEach()` with the ES5 method `Object.keys()`:

```
Object.keys(o).forEach(function (key) {
  var value = o[key];
  print(value);
});
```

Array Comprehension V8 does not support *Array comprehensions*¹⁴⁷.

Use other methods such as the Array *instance* methods `map()`, `filter()`, or `forEach()`.

Example

With V8, the following array comprehension is **invalid**:

```
var a = { w: 1, x: 2, y: 3, z: 4 }

var arr = [i * i for each (i in a) if (i > 2)]
printjson(arr)
```

Instead, you can implement using the Array *instance* method `forEach()` and the ES5 method `Object.keys()`:

```
var a = { w: 1, x: 2, y: 3, z: 4 }

var arr = [];
Object.keys(a).forEach(function (key) {
  var val = a[key];
  if (val > 2) arr.push(val * val);
})
printjson(arr)
```

Note: The new logic uses the Array *instance* method `forEach()` and not the *generic* method

¹⁴⁶https://developer.mozilla.org/en-US/docs/JavaScript/Reference/Statements/for_each...in

¹⁴⁷https://developer.mozilla.org/en-US/docs/JavaScript/Guide/Predefined_Core_Objects#Array_comprehensions

`Array.forEach()`; V8 does **not** support *Array generic* methods. See *Array Generic Methods* (page 657) for more information.

Multiple Catch Blocks V8 does not support multiple `catch` blocks and will throw a `SyntaxError`.

Example

The following multiple catch blocks is **invalid** with V8 and will throw `"SyntaxError: Unexpected token if"`:

```
try {
  something()
} catch (err if err instanceof SomeError) {
  print('some error')
} catch (err) {
  print('standard error')
}
```

Conditional Function Definition V8 will produce different outcomes than SpiderMonkey with *conditional function definitions*¹⁴⁸.

Example

The following conditional function definition produces different outcomes in SpiderMonkey versus V8:

```
function test () {
  if (false) {
    function go () {};
  }
  print(typeof go)
}
```

With SpiderMonkey, the conditional function outputs `undefined`, whereas with V8, the conditional function outputs `function`.

If your code defines functions this way, it is highly recommended that you refactor the code. The following example refactors the conditional function definition to work in both SpiderMonkey and V8.

```
function test () {
  var go;
  if (false) {
    go = function () {}
  }
  print(typeof go)
}
```

The refactored code outputs `undefined` in both SpiderMonkey and V8.

Note: ECMAScript prohibits conditional function definitions. To force V8 to throw an `Error`, *enable strict mode*¹⁴⁹.

```
function test () {
  'use strict';
```

¹⁴⁸<https://developer.mozilla.org/en-US/docs/JavaScript/Guide/Functions>

¹⁴⁹<http://www.nczonline.net/blog/2012/03/13/its-time-to-start-using-javascript-strict-mode/>

```

    if (false) {
      function go () {}
    }
  }
}

```

The JavaScript code throws the following syntax error:

SyntaxError: In strict mode code, functions can only be declared at top level or immediately within a

String Generic Methods V8 does not support [String generics](#)¹⁵⁰. String generics are a set of methods on the String class that mirror instance methods.

Example

The following use of the generic method `String.toLowerCase()` is **invalid** with V8:

```

var name = 'MongoDB';

var lower = String.toLowerCase(name);

```

With V8, use the String instance method `toLowerCase()` available through an *instance* of the String class instead:

```

var name = 'MongoDB';

var lower = name.toLowerCase();
print(name + ' becomes ' + lower);

```

With V8, use the String *instance* methods instead of following *generic* methods:

<code>String.charAt()</code>	<code>String.quote()</code>	<code>String.toLocaleLowerCase()</code>
<code>String.charCodeAt()</code>	<code>String.replace()</code>	<code>String.toLocaleUpperCase()</code>
<code>String.concat()</code>	<code>String.search()</code>	<code>String.toLowerCase()</code>
<code>String.endsWith()</code>	<code>String.slice()</code>	<code>String.toUpperCase()</code>
<code>String.indexOf()</code>	<code>String.split()</code>	<code>String.trim()</code>
<code>String.lastIndexOf()</code>	<code>String.startsWith()</code>	<code>String.trimLeft()</code>
<code>String.localeCompare()</code>	<code>String.substr()</code>	<code>String.trimRight()</code>
<code>String.match()</code>	<code>String.substring()</code>	

Array Generic Methods V8 does not support [Array generic methods](#)¹⁵¹. Array generics are a set of methods on the Array class that mirror instance methods.

Example

The following use of the generic method `Array.every()` is **invalid** with V8:

```

var arr = [4, 8, 15, 16, 23, 42];

function isEven (val) {
  return 0 === val % 2;
}

```

¹⁵⁰https://developer.mozilla.org/en-US/docs/JavaScript/Reference/Global_Objects/String#String_generic_methods

¹⁵¹https://developer.mozilla.org/en-US/docs/JavaScript/Reference/Global_Objects/Array#Array_generic_methods

```
var allEven = Array.every(arr, isEven);
print(allEven);
```

With V8, use the `Array` instance method `every()` available through an *instance* of the `Array` class instead:

```
var allEven = arr.every(isEven);
print(allEven);
```

With V8, use the `Array` *instance* methods instead of the following *generic* methods:

<code>Array.concat()</code>	<code>Array.lastIndexOf()</code>	<code>Array.slice()</code>
<code>Array.every()</code>	<code>Array.map()</code>	<code>Array.some()</code>
<code>Array.filter()</code>	<code>Array.pop()</code>	<code>Array.sort()</code>
<code>Array.forEach()</code>	<code>Array.push()</code>	<code>Array.splice()</code>
<code>Array.indexOf()</code>	<code>Array.reverse()</code>	<code>Array.unshift()</code>
<code>Array.join()</code>	<code>Array.shift()</code>	

Array Instance Method `toSource()` V8 does not support the `Array` instance method `toSource()`¹⁵². Use the `Array` instance method `toString()` instead.

`uneval()` V8 does not support the non-standard method `uneval()`. Use the standardized `JSON.stringify()`¹⁵³ method instead.

Change default JavaScript engine from SpiderMonkey to V8. The change provides improved concurrency for JavaScript operations, modernized JavaScript implementation, and the removal of non-standard SpiderMonkey features, and affects all JavaScript behavior including the commands `mapReduce` (page 208), `group` (page 204), and `eval` (page 238) and the query operator `$where` (page 391).

See *JavaScript Changes in MongoDB 2.4* (page 653) for more information about all changes .

BSON Document Validation Enabled by Default for `mongod` and `mongorestore`

Enable basic *BSON* object validation for `mongod` (page 503) and `mongorestore` (page 543) when writing to MongoDB data files. See `wireObjectCheck` for details.

Index Build Enhancements

- Add support for multiple concurrent index builds in the background by a single `mongod` (page 503) instance. See *building indexes in the background* for more information on background index builds.
- Allow the `db.killOp()` (page 114) method to terminate a foreground index build.
- Improve index validation during index creation. See *Compatibility and Index Type Changes in MongoDB 2.4* (page 666) for more information.

Set Parameters as Command Line Options

Provide `--setParameter` as a command line option for `mongos` (page 518) and `mongod` (page 503). See `mongod` (page 503) and `mongos` (page 518) for list of available options for `setParameter`.

¹⁵²https://developer.mozilla.org/en-US/docs/JavaScript/Reference/Global_Objects/Array/toSource

¹⁵³https://developer.mozilla.org/en-US/docs/JavaScript/Reference/Global_Objects/JSON/stringify

Changed Replication Behavior for Chunk Migration

By default, each document move during *chunk migration* in a *sharded cluster* propagates to at least one secondary before the balancer proceeds with its next operation. See *chunk-migration-replication*.

Improved Chunk Migration Queue Behavior

Increase performance for moving multiple chunks off an overloaded shard. The balancer no longer waits for the current migration's delete phase to complete before starting the next chunk migration. See *chunk-migration-queuing* for details.

Enterprise

The following changes are specific to MongoDB Enterprise Editions:

SASL Library Change

In 2.4.4, MongoDB Enterprise uses Cyrus SASL. Earlier 2.4 Enterprise versions use GNU SASL (`libsasl`). To upgrade to 2.4.4 MongoDB Enterprise or greater, you **must** install all package dependencies related to this change, including the appropriate Cyrus SASL GSSAPI library. See <http://docs.mongodb.org/manual/tutorial/install-mongodb-enterprise> for details of the dependencies.

New Modular Authentication System with Support for Kerberos

In 2.4, the MongoDB Enterprise now supports authentication via a Kerberos mechanism. See <http://docs.mongodb.org/manual/tutorial/control-access-to-mongodb-with-kerberos-authentication/> for more information. For drivers that provide support for Kerberos authentication to MongoDB, refer to *kerberos-and-drivers*.

For more information on security and risk management strategies, see MongoDB Security Practices and Procedures.

Additional Information

Platform Notes

For OS X, MongoDB 2.4 only supports OS X versions 10.6 (Snow Leopard) and later. There are no other platform support changes in MongoDB 2.4. See the [downloads page](#)¹⁵⁴ for more information on platform support.

Upgrade Process

Upgrade MongoDB to 2.4 In the general case, the upgrade from MongoDB 2.2 to 2.4 is a binary-compatible “drop-in” upgrade: shut down the `mongod` (page 503) instances and replace them with `mongod` (page 503) instances running 2.4. **However**, before you attempt any upgrade please familiarize yourself with the content of this document, particularly the procedure for *upgrading sharded clusters* (page 660) and the considerations for *reverting to 2.2 after running 2.4* (page 665).

¹⁵⁴<http://www.mongodb.org/downloads/>

Upgrade Recommendations and Checklist When upgrading, consider the following:

- For all deployments using authentication, upgrade the drivers (i.e. client libraries), before upgrading the `mongod` (page 503) instance or instances.
- To upgrade to 2.4 sharded clusters *must* upgrade following the *meta-data upgrade procedure* (page 660).
- If you're using 2.2.0 and running with `authorization` enabled, you will need to upgrade first to 2.2.1 and then upgrade to 2.4. See *Rolling Upgrade Limitation for 2.2.0 Deployments Running with auth Enabled* (page 664).
- If you have `system.users` documents (i.e. for authorization) that you created before 2.4 you *must* ensure that there are no duplicate values for the `user` field in the `system.users` collection in *any* database. If you *do* have documents with duplicate user fields, you must remove them before upgrading.

See *Security Enhancements* (page 653) for more information.

Upgrade Standalone mongod Instance to MongoDB 2.4

1. Download binaries of the latest release in the 2.4 series from the [MongoDB Download Page](http://docs.mongodb.org/manual/installation)¹⁵⁵. See <http://docs.mongodb.org/manual/installation> for more information.
2. Shutdown your `mongod` (page 503) instance. Replace the existing binary with the 2.4 `mongod` (page 503) binary and restart `mongod` (page 503).

Upgrade a Replica Set from MongoDB 2.2 to MongoDB 2.4 You can upgrade to 2.4 by performing a “rolling” upgrade of the set by upgrading the members individually while the other members are available to minimize downtime. Use the following procedure:

1. Upgrade the *secondary* members of the set one at a time by shutting down the `mongod` (page 503) and replacing the 2.2 binary with the 2.4 binary. After upgrading a `mongod` (page 503) instance, wait for the member to recover to `SECONDARY` state before upgrading the next instance. To check the member's state, issue `rs.status()` (page 168) in the `mongo` (page 527) shell.
2. Use the `mongo` (page 527) shell method `rs.stepDown()` (page 168) to step down the *primary* to allow the normal *failover* procedure. `rs.stepDown()` (page 168) expedites the failover procedure and is preferable to shutting down the primary directly.

Once the primary has stepped down and another member has assumed `PRIMARY` state, as observed in the output of `rs.status()` (page 168), shut down the previous primary and replace `mongod` (page 503) binary with the 2.4 binary and start the new process.

Note: Replica set failover is not instant but will render the set unavailable to read or accept writes until the failover process completes. Typically this takes 10 seconds or more. You may wish to plan the upgrade during a predefined maintenance window.

Upgrade a Sharded Cluster from MongoDB 2.2 to MongoDB 2.4

Important: Only upgrade sharded clusters to 2.4 if **all** members of the cluster are currently running instances of 2.2. The only supported upgrade path for sharded clusters running 2.0 is via 2.2.

Overview Upgrading a *sharded cluster* from MongoDB version 2.2 to 2.4 (or 2.3) requires that you run a 2.4 `mongos` (page 518) with the `--upgrade` option, described in this procedure. The upgrade process does not require downtime.

¹⁵⁵<http://www.mongodb.org/downloads>

The upgrade to MongoDB 2.4 adds epochs to the meta-data for all collections and chunks in the existing cluster. MongoDB 2.2 processes are capable of handling epochs, even though 2.2 did not require them. This procedure applies only to upgrades from version 2.2. Earlier versions of MongoDB do not correctly handle epochs. See *Cluster Meta-data Upgrade* (page 661) for more information.

After completing the meta-data upgrade you can fully upgrade the components of the cluster. With the balancer disabled:

- Upgrade all `mongos` (page 518) instances in the cluster.
- Upgrade all 3 `mongod` (page 503) config server instances.
- Upgrade the `mongod` (page 503) instances for each shard, one at a time.

See *Upgrade Sharded Cluster Components* (page 664) for more information.

Cluster Meta-data Upgrade

Considerations Beware of the following properties of the cluster upgrade process:

- Before you start the upgrade, ensure that the amount of free space on the filesystem for the *config database* (page 593) is at least 4 to 5 times the amount of space currently used by the *config database* (page 593) data files.

Additionally, ensure that all indexes in the *config database* (page 593) are `{v:1}` indexes. If a critical index is a `{v:0}` index, chunk splits can fail due to known issues with the `{v:0}` format. `{v:0}` indexes are present on clusters created with MongoDB 2.0 or earlier.

The duration of the metadata upgrade depends on the network latency between the node performing the upgrade and the three config servers. Ensure low latency between the upgrade process and the config servers.

- While the upgrade is in progress, you cannot make changes to the collection meta-data. For example, during the upgrade, do **not** perform:
 - `sh.enableSharding()` (page 175),
 - `sh.shardCollection()` (page 178),
 - `sh.addShard()` (page 173),
 - `db.createCollection()` (page 102),
 - `db.collection.drop()` (page 29),
 - `db.dropDatabase()` (page 108),
 - any operation that creates a database, or
 - any other operation that modifies the cluster meta-data in any way. See <http://docs.mongodb.org/manualreference/sharding> for a complete list of sharding commands. Note, however, that not all commands on the <http://docs.mongodb.org/manualreference/sharding> page modifies the cluster meta-data.
- Once you upgrade to 2.4 and complete the upgrade procedure **do not** use 2.0 `mongod` (page 503) and `mongos` (page 518) processes in your cluster. 2.0 process may re-introduce old meta-data formats into cluster meta-data.

The upgraded config database will require more storage space than before, to make backup and working copies of the `config.chunks` (page 595) and `config.collections` (page 595) collections. As always, if storage requirements increase, the `mongod` (page 503) might need to pre-allocate additional data files. See *faq-tools-for-measuring-storage-use* for more information.

Meta-data Upgrade Procedure Changes to the meta-data format for sharded clusters, stored in the *config database* (page 593), require a special meta-data upgrade procedure when moving to 2.4.

Do not perform operations that modify meta-data while performing this procedure. See *Upgrade a Sharded Cluster from MongoDB 2.2 to MongoDB 2.4* (page 660) for examples of prohibited operations.

1. Before you start the upgrade, ensure that the amount of free space on the filesystem for the *config database* (page 593) is at least 4 to 5 times the amount of space currently used by the *config database* (page 593) data files.

Additionally, ensure that all indexes in the *config database* (page 593) are `{v:1}` indexes. If a critical index is a `{v:0}` index, chunk splits can fail due to known issues with the `{v:0}` format. `{v:0}` indexes are present on clusters created with MongoDB 2.0 or earlier.

The duration of the metadata upgrade depends on the network latency between the node performing the upgrade and the three config servers. Ensure low latency between the upgrade process and the config servers.

To check the version of your indexes, use `db.collection.getIndexes()` (page 45).

If any index **on the config database** is `{v:0}`, you should rebuild those indexes by connecting to the *mongos* (page 518) and either: rebuild all indexes using the `db.collection.reIndex()` (page 62) method, or drop and rebuild specific indexes using `db.collection.dropIndex()` (page 29) and then `db.collection.ensureIndex()` (page 30). If you need to upgrade the `_id` index to `{v:1}` use `db.collection.reIndex()` (page 62).

You may have `{v:0}` indexes on other databases in the cluster.

2. Turn off the *balancer* in the *sharded cluster*, as described in *sharding-balancing-disable-temporarily*.

Optional

For additional security during the upgrade, you can make a backup of the config database using *mongodump* (page 537) or other backup tools.

3. Ensure there are no version 2.0 *mongod* (page 503) or *mongos* (page 518) processes still active in the sharded cluster. The automated upgrade process checks for 2.0 processes, but network availability can prevent a definitive check. Wait 5 minutes after stopping or upgrading version 2.0 *mongos* (page 518) processes to confirm that none are still active.
4. Start a single 2.4 *mongos* (page 518) process with `configDB` pointing to the sharded cluster's *config servers* and with the `--upgrade` option. The upgrade process happens before the process becomes a daemon (i.e. before `--fork`.)

You can upgrade an existing *mongos* (page 518) instance to 2.4 or you can start a new *mongos* instance that can reach all config servers if you need to avoid reconfiguring a production *mongos* (page 518).

Start the *mongos* (page 518) with a command that resembles the following:

```
mongos --configdb <config servers> --upgrade
```

Without the `--upgrade` option 2.4 *mongos* (page 518) processes will fail to start until the upgrade process is complete.

The upgrade will prevent any chunk moves or splits from occurring during the upgrade process. If there are very many sharded collections or there are stale locks held by other failed processes, acquiring the locks for all collections can take seconds or minutes. See the log for progress updates.

5. When the *mongos* (page 518) process starts successfully, the upgrade is complete. If the *mongos* (page 518) process fails to start, check the log for more information.

If the *mongos* (page 518) terminates or loses its connection to the config servers during the upgrade, you may always safely retry the upgrade.

However, if the upgrade failed during the short critical section, the `mongos` (page 518) will exit and report that the upgrade will require manual intervention. To continue the upgrade process, you must follow the *Resync after an Interruption of the Critical Section* (page 663) procedure.

Optional

If the `mongos` (page 518) logs show the upgrade waiting for the upgrade lock, a previous upgrade process may still be active or may have ended abnormally. After 15 minutes of no remote activity `mongos` (page 518) will force the upgrade lock. If you can verify that there are no running upgrade processes, you may connect to a 2.2 `mongos` (page 518) process and force the lock manually:

```
mongo <mongos.example.net>
```

```
db.getMongo().getCollection("config.locks").findOne({ _id : "configUpgrade" })
```

If the process specified in the `process` field of this document is *verifiably* offline, run the following operation to force the lock.

```
db.getMongo().getCollection("config.locks").update({ _id : "configUpgrade" }, { $set : { state :
```

It is always more safe to wait for the `mongos` (page 518) to verify that the lock is inactive, if you have any doubts about the activity of another upgrade operation. In addition to the `configUpgrade`, the `mongos` (page 518) may need to wait for specific collection locks. Do not force the specific collection locks.

-
6. Upgrade and restart other `mongos` (page 518) processes in the sharded cluster, *without* the `--upgrade` option.

See *Upgrade Sharded Cluster Components* (page 664) for more information.

7. *Re-enable the balancer.* You can now perform operations that modify cluster meta-data.

Once you have upgraded, *do not* introduce version 2.0 MongoDB processes into the sharded cluster. This can reintroduce old meta-data formats into the config servers. The meta-data change made by this upgrade process will help prevent errors caused by cross-version incompatibilities in future versions of MongoDB.

Resync after an Interruption of the Critical Section During the short critical section of the upgrade that applies changes to the meta-data, it is unlikely but possible that a network interruption can prevent all three config servers from verifying or modifying data. If this occurs, the *config servers* must be re-synced, and there may be problems starting new `mongos` (page 518) processes. The *sharded cluster* will remain accessible, but avoid all cluster meta-data changes until you resync the config servers. Operations that change meta-data include: adding shards, dropping databases, and dropping collections.

Note: Only perform the following procedure *if* something (e.g. network, power, etc.) interrupts the upgrade process during the short critical section of the upgrade. Remember, you may always safely attempt the *meta data upgrade procedure* (page 662).

To resync the config servers:

1. Turn off the *balancer* in the sharded cluster and stop all meta-data operations. If you are in the middle of an upgrade process (*Upgrade a Sharded Cluster from MongoDB 2.2 to MongoDB 2.4* (page 660)), you have already disabled the balancer.
2. Shut down two of the three config servers, preferably the last two listed in the `configDB` string. For example, if your `configDB` string is `configA:27019,configB:27019,configC:27019`, shut down `configB` and `configC`. Shutting down the last two config servers ensures that most `mongos` (page 518) instances will have uninterrupted access to cluster meta-data.
3. `mongodump` (page 537) the data files of the active config server (`configA`).

4. Move the data files of the deactivated config servers (`configB` and `configC`) to a backup location.
5. Create new, empty *data directories*.
6. Restart the disabled config servers with `--dbpath` pointing to the now-empty data directory and `--port` pointing to an alternate port (e.g. 27020).
7. Use `mongorestore` (page 543) to repopulate the data files on the disabled documents from the active config server (`configA`) to the restarted config servers on the new port (`configB:27020, configC:27020`). These config servers are now re-synced.
8. Restart the restored config servers on the old port, resetting the port back to the old settings (`configB:27019` and `configC:27019`).
9. In some cases connection pooling may cause spurious failures, as the `mongos` (page 518) disables old connections only after attempted use. 2.4 fixes this problem, but to avoid this issue in version 2.2, you can restart all `mongos` (page 518) instances (one-by-one, to avoid downtime) and use the `rs.stepDown()` (page 168) method before restarting each of the shard *replica set primaries*.
10. The sharded cluster is now fully resynced; however before you attempt the upgrade process again, you must manually reset the upgrade state using a version 2.2 `mongos` (page 518). Begin by connecting to the 2.2 `mongos` (page 518) with the `mongo` (page 527) shell:

```
mongo <mongos.example.net>
```

Then, use the following operation to reset the upgrade process:

```
db.getMongo().getCollection("config.version").update({ _id : 1 }, { $unset : { upgradeState : 1
```

11. Finally retry the upgrade process, as in *Upgrade a Sharded Cluster from MongoDB 2.2 to MongoDB 2.4* (page 660).

Upgrade Sharded Cluster Components After you have successfully completed the meta-data upgrade process described in *Meta-data Upgrade Procedure* (page 662), and the 2.4 `mongos` (page 518) instance starts, you can upgrade the other processes in your MongoDB deployment.

While the balancer is still disabled, upgrade the components of your sharded cluster in the following order:

- Upgrade all `mongos` (page 518) instances in the cluster, in any order.
- Upgrade all 3 `mongod` (page 503) config server instances, upgrading the *first* system in the `mongos --configdb` argument *last*.
- Upgrade each shard, one at a time, upgrading the `mongod` (page 503) secondaries before running `replSetStepDown` (page 277) and upgrading the primary of each shard.

When this process is complete, you can now *re-enable the balancer*.

Rolling Upgrade Limitation for 2.2.0 Deployments Running with auth Enabled MongoDB *cannot* support deployments that mix 2.2.0 and 2.4.0, or greater, components. MongoDB version 2.2.1 and later processes *can* exist in mixed deployments with 2.4-series processes. Therefore you cannot perform a rolling upgrade from MongoDB 2.2.0 to MongoDB 2.4.0. To upgrade a cluster with 2.2.0 components, use one of the following procedures.

1. Perform a rolling upgrade of all 2.2.0 processes to the latest 2.2-series release (e.g. 2.2.3) so that there are no processes in the deployment that predate 2.2.1. When there are no 2.2.0 processes in the deployment, perform a rolling upgrade to 2.4.0.
2. Stop all processes in the cluster. Upgrade all processes to a 2.4-series release of MongoDB, and start all processes at the same time.

Upgrade from 2.3 to 2.4 If you used a `mongod` (page 503) from the 2.3 or 2.4-rc (release candidate) series, you can safely transition these databases to 2.4.0 or later; *however*, if you created `2dsphere` or `text` indexes using a `mongod` (page 503) before v2.4-rc2, you will need to rebuild these indexes. For example:

```
db.records.dropIndex( { loc: "2dsphere" } )
db.records.dropIndex( "records_text" )

db.records.ensureIndex( { loc: "2dsphere" } )
db.records.ensureIndex( { records: "text" } )
```

Downgrade MongoDB from 2.4 to Previous Versions For some cases the on-disk format of data files used by 2.4 and 2.2 `mongod` (page 503) is compatible, and you can upgrade and downgrade if needed. However, several new features in 2.4 are incompatible with previous versions:

- `2dsphere` indexes are incompatible with 2.2 and earlier `mongod` (page 503) instances.
- `text` indexes are incompatible with 2.2 and earlier `mongod` (page 503) instances.
- using a hashed index as a shard key are incompatible with 2.2 and earlier `mongos` (page 518) instances.
- hashed indexes are incompatible with 2.0 and earlier `mongod` (page 503) instances.

Important: Collections sharded using hashed shard keys, should **not** use 2.2 `mongod` (page 503) instances, which cannot correctly support cluster operations for these collections.

If you completed the *meta-data upgrade for a sharded cluster* (page 660), you can safely downgrade to 2.2 MongoDB processes. **Do not** use 2.0 processes after completing the upgrade procedure.

Note: In sharded clusters, once you have completed the *meta-data upgrade procedure* (page 660), you cannot use 2.0 `mongod` (page 503) or `mongos` (page 518) instances in the same cluster.

If you complete the meta-data upgrade, you can safely downgrade components in any order. When upgrade again, always upgrade `mongos` (page 518) instances before `mongod` (page 503) instances.

Do not create `2dsphere` or `text` indexes in a cluster that has 2.2 components.

Considerations and Compatibility If you upgrade to MongoDB 2.4, and then need to run MongoDB 2.2 with the same data files, consider the following limitations.

- If you use a hashed index as the shard key index, which is only possible under 2.4 you will not be able to query data in this sharded collection. Furthermore, a 2.2 `mongos` (page 518) cannot properly route an insert operation for a collections sharded using a hashed index for the shard key index: any data that you insert using a 2.2 `mongos` (page 518), will not arrive on the correct shard and will not be reachable by future queries.
- If you *never* create an `2dsphere` or `text` index, you can move between a 2.4 and 2.2 `mongod` (page 503) for a given data set; however, after you create the first `2dsphere` or `text` index with a 2.4 `mongod` (page 503) you will need to run a 2.2 `mongod` (page 503) with the `--upgrade` option and drop any `2dsphere` or `text` index.

Upgrade and Downgrade Procedures

Basic Downgrade and Upgrade Except as described below, moving between 2.2 and 2.4 is a drop-in replacement:

- stop the existing `mongod` (page 503), using the `--shutdown` option as follows:

```
mongod --dbpath /var/mongod/data --shutdown
```

Replace `/var/mongod/data` with your MongoDB `dbPath`.

- start the new `mongod` (page 503) processes with the same `dbPath` setting, for example:

```
mongod --dbpath /var/mongod/data
```

Replace `/var/mongod/data` with your MongoDB `dbPath`.

Downgrade to 2.2 After Creating a 2dsphere or text Index If you have created 2dsphere or text indexes while running a 2.4 `mongod` (page 503) instance, you can downgrade at any time, by starting the 2.2 `mongod` (page 503) with the `--upgrade` option as follows:

```
mongod --dbpath /var/mongod/data/ --upgrade
```

Then, you will need to drop any existing 2dsphere or text indexes using `db.collection.dropIndex()` (page 29), for example:

```
db.records.dropIndex( { loc: "2dsphere" } )
db.records.dropIndex( "records_text" )
```

Warning: `--upgrade` will run `repairDatabase` (page 319) on any database where you have created a 2dsphere or text index, which will rebuild *all* indexes.

Troubleshooting Upgrade/Downgrade Operations If you do not use `--upgrade`, when you attempt to start a 2.2 `mongod` (page 503) and you have created a 2dsphere or text index, `mongod` (page 503) will return the following message:

```
'need to upgrade database index_plugin_upgrade with pdfile version 4.6, new version: 4.5 Not upgrading'
```

While running 2.4, to check the data file version of a MongoDB database, use the following operation in the shell:

```
db.getSiblingDB('<databaseName>').stats().dataFileVersion
```

The major data file ¹⁵⁶ version for both 2.2 and 2.4 is 4, the minor data file version for 2.2 is 5 and the minor data file version for 2.4 is 6 **after** you create a 2dsphere or text index.

Compatibility and Index Type Changes in MongoDB 2.4 In 2.4 MongoDB includes two new features related to indexes that users upgrading to version 2.4 must consider, particularly with regard to possible downgrade paths. For more information on downgrades, see [Downgrade MongoDB from 2.4 to Previous Versions](#) (page 665).

New Index Types In 2.4 MongoDB adds two new index types: 2dsphere and text. These index types do not exist in 2.2, and for each database, creating a 2dsphere or text index, will upgrade the data-file version and make that database incompatible with 2.2.

If you intend to downgrade, you should always drop all 2dsphere and text indexes before moving to 2.2.

You can use the [downgrade procedure](#) (page 665) to downgrade these databases and run 2.2 if needed, however this will run a full database repair (as with `repairDatabase` (page 319)) for all affected databases.

¹⁵⁶ The data file version (i.e. pdfile version) is independent and unrelated to the release version of MongoDB.

Index Type Validation In MongoDB 2.2 and earlier you could specify invalid index types that did not exist. In these situations, MongoDB would create an ascending (e.g. 1) index. Invalid indexes include index types specified by strings that do not refer to an existing index type, and all numbers other than 1 and -1.¹⁵⁷

In 2.4, creating any invalid index will result in an error. Furthermore, you cannot create a 2dsphere or text index on a collection if its containing database has any invalid index types.¹

Example

If you attempt to add an invalid index in MongoDB 2.4, as in the following:

```
db.coll.ensureIndex( { field: "1" } )
```

MongoDB will return the following error document:

```
{
  "err" : "Unknown index plugin '1' in index { field: \"1\" }"
  "code" : 16734,
  "n" : <number>,
  "connectionId" : <number>,
  "ok" : 1
}
```

See *Upgrade MongoDB to 2.4* (page 659) for full upgrade instructions.

Other Resources

- MongoDB Downloads¹⁵⁸.
- All JIRA issues resolved in 2.4¹⁵⁹.
- All Backwards incompatible changes¹⁶⁰.
- All Third Party License Notices¹⁶¹.

7.2.2 Release Notes for MongoDB 2.2

Upgrading

MongoDB 2.2 is a production release series and succeeds the 2.0 production release series.

MongoDB 2.0 data files are compatible with 2.2-series binaries without any special migration process. However, always perform the upgrade process for replica sets and sharded clusters using the procedures that follow.

Synopsis

- `mongod` (page 503), 2.2 is a drop-in replacement for 2.0 and 1.8.

¹⁵⁷ In 2.4, indexes that specify a type of "1" or "-1" (the strings "1" and "-1") will continue to exist, despite a warning on start-up. **However**, a *secondary* in a replica set cannot complete an initial sync from a primary that has a "1" or "-1" index. Avoid all indexes with invalid types.

¹⁵⁸ <http://mongodb.org/downloads>

¹⁵⁹ <https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fix+Version+in+%28%222.3.2%22,+%222.3.1%22,+%222.3.0%22,+%222.4.0-rc1%22,+%222.4.0-rc2%22,+%222.4.0-rc3%22%29>

¹⁶⁰ <https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fix+Version+in+%28%222.3.2%22%2C+%222.3.1%22%2C+%222.3.0%22%2C+%222.4.0-rc1%22%2C+%222.4.0-rc2%22%2C+%222.4.0-rc3%22%29+AND+%22Backward+Breaking%22+in+%28+Rarely+%2C+sometimes%2C+yes%29%29>

¹⁶¹ <https://github.com/mongodb/mongo/blob/v2.4/distsrc/THIRD-PARTY-NOTICES>

- Check your `driver` documentation for information regarding required compatibility upgrades, and always run the recent release of your driver.

Typically, only users running with authentication, will need to upgrade drivers before continuing with the upgrade to 2.2.

- For all deployments using authentication, upgrade the drivers (i.e. client libraries), before upgrading the `mongod` (page 503) instance or instances.
- For all upgrades of sharded clusters:
 - turn off the balancer during the upgrade process. See the *sharding-balancing-disable-temporarily* section for more information.
 - upgrade all `mongos` (page 518) instances before upgrading any `mongod` (page 503) instances.

Other than the above restrictions, 2.2 processes can interoperate with 2.0 and 1.8 tools and processes. You can safely upgrade the `mongod` (page 503) and `mongos` (page 518) components of a deployment one by one while the deployment is otherwise operational. Be sure to read the detailed upgrade procedures below before upgrading production systems.

Upgrading a Standalone `mongod`

1. Download binaries of the latest release in the 2.2 series from the [MongoDB Download Page](#)¹⁶².
2. Shutdown your `mongod` (page 503) instance. Replace the existing binary with the 2.2 `mongod` (page 503) binary and restart MongoDB.

Upgrading a Replica Set

You can upgrade to 2.2 by performing a “rolling” upgrade of the set by upgrading the members individually while the other members are available to minimize downtime. Use the following procedure:

1. Upgrade the *secondary* members of the set one at a time by shutting down the `mongod` (page 503) and replacing the 2.0 binary with the 2.2 binary. After upgrading a `mongod` (page 503) instance, wait for the member to recover to `SECONDARY` state before upgrading the next instance. To check the member’s state, issue `rs.status()` (page 168) in the `mongo` (page 527) shell.
2. Use the `mongo` (page 527) shell method `rs.stepDown()` (page 168) to step down the *primary* to allow the normal *failover* procedure. `rs.stepDown()` (page 168) expedites the failover procedure and is preferable to shutting down the primary directly.

Once the primary has stepped down and another member has assumed `PRIMARY` state, as observed in the output of `rs.status()` (page 168), shut down the previous primary and replace `mongod` (page 503) binary with the 2.2 binary and start the new process.

Note: Replica set failover is not instant but will render the set unavailable to read or accept writes until the failover process completes. Typically this takes 10 seconds or more. You may wish to plan the upgrade during a predefined maintenance window.

Upgrading a Sharded Cluster

Use the following procedure to upgrade a sharded cluster:

- *Disable the balancer.*

¹⁶²<http://downloads.mongodb.org/>

- Upgrade all `mongos` (page 518) instances *first*, in any order.
- Upgrade all of the `mongod` (page 503) config server instances using the *stand alone* (page 668) procedure. To keep the cluster online, be sure that at all times at least one config server is up.
- Upgrade each shard's replica set, using the *upgrade procedure for replica sets* (page 668) detailed above.
- re-enable the balancer.

Note: Balancing is not currently supported in *mixed* 2.0.x and 2.2.0 deployments. Thus you will want to reach a consistent version for all shards within a reasonable period of time, e.g. same-day. See [SERVER-6902](#)¹⁶³ for more information.

Changes

Major Features

Aggregation Framework The aggregation framework makes it possible to do aggregation operations without needing to use *map-reduce*. The `aggregate` (page 198) command exposes the aggregation framework, and the `aggregate()` (page 22) helper in the `mongo` (page 527) shell provides an interface to these operations. Consider the following resources for background on the aggregation framework and its use:

- Documentation: <http://docs.mongodb.org/manualcore/aggregation>
- Reference: *Aggregation Reference* (page 484)
- Examples: <http://docs.mongodb.org/manualapplications/aggregation>

TTL Collections TTL collections remove expired data from a collection, using a special index and a background thread that deletes expired documents every minute. These collections are useful as an alternative to *capped collections* in some cases, such as for data warehousing and caching cases, including: machine generated event data, logs, and session information that needs to persist in a database for only a limited period of time.

For more information, see the <http://docs.mongodb.org/manualtutorial/expire-data> tutorial.

Concurrency Improvements MongoDB 2.2 increases the server's capacity for concurrent operations with the following improvements:

1. DB Level Locking¹⁶⁴
2. Improved Yielding on Page Faults¹⁶⁵
3. Improved Page Fault Detection on Windows¹⁶⁶

To reflect these changes, MongoDB now provides changed and improved reporting for concurrency and use, see *locks* (page 348) and *recordStats* (page 358) in *server status* (page 346) and see `db.currentOp()` (page 103), *mongotop* (page 576), and *mongostat* (page 569).

¹⁶³<https://jira.mongodb.org/browse/SERVER-6902>

¹⁶⁴<https://jira.mongodb.org/browse/SERVER-4328>

¹⁶⁵<https://jira.mongodb.org/browse/SERVER-3357>

¹⁶⁶<https://jira.mongodb.org/browse/SERVER-4538>

Improved Data Center Awareness with Tag Aware Sharding MongoDB 2.2 adds additional support for geographic distribution or other custom partitioning for sharded collections in *clusters*. By using this “tag aware” sharding, you can automatically ensure that data in a sharded database system is always on specific shards. For example, with tag aware sharding, you can ensure that data is closest to the application servers that use that data most frequently.

Shard tagging controls data location, and is complementary but separate from replica set tagging, which controls *read preference* and *write concern*. For example, shard tagging can pin all “USA” data to one or more logical shards, while replica set tagging can control which *mongod* (page 503) instances (e.g. “production” or “reporting”) the application uses to service requests.

See the documentation for the following helpers in the *mongo* (page 527) shell that support tagged sharding configuration:

- `sh.addShardTag()` (page 174)
- `sh.addTagRange()` (page 174)
- `sh.removeShardTag()` (page 178)

Also, see <http://docs.mongodb.org/manualcore/tag-aware-sharding> and <http://docs.mongodb.org/manualtutorial/administer-shard-tags>.

Fully Supported Read Preference Semantics All MongoDB clients and drivers now support full *read preferences*, including consistent support for a full range of *read preference modes* and *tag sets*. This support extends to the *mongos* (page 518) and applies identically to single replica sets and to the replica sets for each shard in a *sharded cluster*.

Additional read preference support now exists in the *mongo* (page 527) shell using the `readPref()` (page 91) cursor method.

Compatibility Changes

Authentication Changes MongoDB 2.2 provides more reliable and robust support for authentication clients, including drivers and *mongos* (page 518) instances.

If your cluster runs with authentication:

- For all drivers, use the latest release of your driver and check its release notes.
- In sharded environments, to ensure that your cluster remains available during the upgrade process you **must** use the *upgrade procedure for sharded clusters* (page 668).

findAndModify Returns Null Value for Upserts that Perform Inserts In version 2.2, for *upsert* that perform inserts with the `new` option set to `false`, `findAndModify` (page 229) commands will now return the following output:

```
{ 'ok': 1.0, 'value': null }
```

In the *mongo* (page 527) shell, *upsert findAndModify* (page 229) operations that perform inserts (with `new` set to `false`.) only output a `null` value.

In version 2.0 these operations would return an empty document, e.g. `{ }`.

See: [SERVER-6226](#)¹⁶⁷ for more information.

¹⁶⁷<https://jira.mongodb.org/browse/SERVER-6226>

mongodump 2.2 Output Incompatible with Pre-2.2 mongorestore If you use the `mongodump` (page 537) tool from the 2.2 distribution to create a dump of a database, you must use a 2.2 (or later) version of `mongorestore` (page 543) to restore that dump.

See: [SERVER-6961](https://jira.mongodb.org/browse/SERVER-6961)¹⁶⁸ for more information.

ObjectId().toString() Returns String Literal ObjectId("...") In version 2.2, the `toString()` (page 187) method returns the string representation of the `ObjectId()` object and has the format `ObjectId("...")`.

Consider the following example that calls the `toString()` (page 187) method on the `ObjectId("507c7f79bcf86cd7994f6c0e")` object:

```
ObjectId("507c7f79bcf86cd7994f6c0e").toString()
```

The method now returns the *string* `ObjectId("507c7f79bcf86cd7994f6c0e")`.

Previously, in version 2.0, the method would return the *hexadecimal string* `507c7f79bcf86cd7994f6c0e`.

If compatibility between versions 2.0 and 2.2 is required, use `ObjectId().str`, which holds the hexadecimal string value in both versions.

ObjectId().valueOf() Returns hexadecimal string In version 2.2, the `valueOf()` (page 187) method returns the value of the `ObjectId()` object as a lowercase hexadecimal string.

Consider the following example that calls the `valueOf()` (page 187) method on the `ObjectId("507c7f79bcf86cd7994f6c0e")` object:

```
ObjectId("507c7f79bcf86cd7994f6c0e").valueOf()
```

The method now returns the *hexadecimal string* `507c7f79bcf86cd7994f6c0e`.

Previously, in version 2.0, the method would return the *object* `ObjectId("507c7f79bcf86cd7994f6c0e")`.

If compatibility between versions 2.0 and 2.2 is required, use `ObjectId().str` attribute, which holds the hexadecimal string value in both versions.

Behavioral Changes

Restrictions on Collection Names In version 2.2, collection names cannot:

- contain the \$.
- be an empty string (i.e. "").

This change does not affect collections created with now illegal names in earlier versions of MongoDB.

These new restrictions are in addition to the existing restrictions on collection names which are:

- A collection name should begin with a letter or an underscore.
- A collection name cannot contain the null character.
- Begin with the `system.` prefix. MongoDB reserves `system.` for system collections, such as the `system.indexes` collection.
- The maximum size of a collection name is 128 characters, including the name of the database. However, for maximum flexibility, collections should have names less than 80 characters.

¹⁶⁸<https://jira.mongodb.org/browse/SERVER-6961>

Collections names may have any other valid UTF-8 string.

See the [SERVER-4442](#)¹⁶⁹ and the *faq-restrictions-on-collection-names* FAQ item.

Restrictions on Database Names for Windows Database names running on Windows can no longer contain the following characters:

/ \ . " * < > : | ?

The names of the data files include the database name. If you attempt to upgrade a database instance with one or more of these characters, `mongod` (page 503) will refuse to start.

Change the name of these databases before upgrading. See [SERVER-4584](#)¹⁷⁰ and [SERVER-6729](#)¹⁷¹ for more information.

_id Fields and Indexes on Capped Collections All *capped collections* now have an `_id` field by default, *if* they exist outside of the `local` database, and now have indexes on the `_id` field. This change only affects capped collections created with 2.2 instances and does not affect existing capped collections.

See: [SERVER-5516](#)¹⁷² for more information.

New \$elemMatch Projection Operator The `$elemMatch` (page 408) operator allows applications to narrow the data returned from queries so that the query operation will only return the first matching element in an array. See the *\$elemMatch (projection)* (page 408) documentation and the [SERVER-2238](#)¹⁷³ and [SERVER-828](#)¹⁷⁴ issues for more information.

Windows Specific Changes

Windows XP is Not Supported As of 2.2, MongoDB does not support Windows XP. Please upgrade to a more recent version of Windows to use the latest releases of MongoDB. See [SERVER-5648](#)¹⁷⁵ for more information.

Service Support for mongos.exe You may now run `mongos.exe` (page 535) instances as a Windows Service. See the *mongos.exe* (page 535) reference and *tutorial-mongod-as-windows-service* and [SERVER-1589](#)¹⁷⁶ for more information.

Log Rotate Command Support MongoDB for Windows now supports log rotation by way of the `logRotate` (page 322) database command. See [SERVER-2612](#)¹⁷⁷ for more information.

New Build Using SlimReadWrite Locks for Windows Concurrency Labeled “2008+” on the [Downloads Page](#)¹⁷⁸, this build for 64-bit versions of Windows Server 2008 R2 and for Windows 7 or newer, offers increased performance over the standard 64-bit Windows build of MongoDB. See [SERVER-3844](#)¹⁷⁹ for more information.

¹⁶⁹<https://jira.mongodb.org/browse/SERVER-4442>

¹⁷⁰<https://jira.mongodb.org/browse/SERVER-4584>

¹⁷¹<https://jira.mongodb.org/browse/SERVER-6729>

¹⁷²<https://jira.mongodb.org/browse/SERVER-5516>

¹⁷³<https://jira.mongodb.org/browse/SERVER-2238>

¹⁷⁴<https://jira.mongodb.org/browse/SERVER-828>

¹⁷⁵<https://jira.mongodb.org/browse/SERVER-5648>

¹⁷⁶<https://jira.mongodb.org/browse/SERVER-1589>

¹⁷⁷<https://jira.mongodb.org/browse/SERVER-2612>

¹⁷⁸<http://www.mongodb.org/downloads>

¹⁷⁹<https://jira.mongodb.org/browse/SERVER-3844>

Tool Improvements

Index Definitions Handled by `mongodump` and `mongorestore` When you specify the `--collection` option to `mongodump` (page 537), `mongodump` (page 537) will now backup the definitions for all indexes that exist on the source database. When you attempt to restore this backup with `mongorestore` (page 543), the target `mongod` (page 503) will rebuild all indexes. See [SERVER-808](#)¹⁸⁰ for more information.

`mongorestore` (page 543) now includes the `--noIndexRestore` option to provide the preceding behavior. Use `--noIndexRestore` to prevent `mongorestore` (page 543) from building previous indexes.

`mongooplog` for Replaying Oplogs The `mongooplog` (page 551) tool makes it possible to pull *oplog* entries from `mongod` (page 503) instance and apply them to another `mongod` (page 503) instance. You can use `mongooplog` (page 551) to achieve point-in-time backup of a MongoDB data set. See the [SERVER-3873](#)¹⁸¹ case and the *mongooplog* (page 550) documentation.

Authentication Support for `mongotop` and `mongostat` `mongotop` (page 576) and `mongostat` (page 570) now contain support for username/password authentication. See [SERVER-3875](#)¹⁸² and [SERVER-3871](#)¹⁸³ for more information regarding this change. Also consider the documentation of the following options for additional information:

- `mongotop --username`
- `mongotop --password`
- `mongostat --username`
- `mongostat --password`

Write Concern Support for `mongoimport` and `mongorestore` `mongoimport` (page 556) now provides an option to halt the import if the operation encounters an error, such as a network interruption, a duplicate key exception, or a write error. The `--stopOnError` option will produce an error rather than silently continue importing data. See [SERVER-3937](#)¹⁸⁴ for more information.

In `mongorestore` (page 543), the `--w` option provides support for configurable write concern.

`mongodump` Support for Reading from Secondaries You can now run `mongodump` (page 537) when connected to a *secondary* member of a *replica set*. See [SERVER-3854](#)¹⁸⁵ for more information.

`mongoimport` Support for full 16MB Documents Previously, `mongoimport` (page 556) would only import documents that were less than 4 megabytes in size. This issue is now corrected, and you may use `mongoimport` (page 556) to import documents that are at least 16 megabytes in size. See [SERVER-4593](#)¹⁸⁶ for more information.

`Timestamp()` Extended JSON format MongoDB extended JSON now includes a new `Timestamp()` type to represent the Timestamp type that MongoDB uses for timestamps in the *oplog* among other contexts.

This permits tools like `mongooplog` (page 551) and `mongodump` (page 537) to query for specific timestamps. Consider the following `mongodump` (page 537) operation:

¹⁸⁰<https://jira.mongodb.org/browse/SERVER-808>

¹⁸¹<https://jira.mongodb.org/browse/SERVER-3873>

¹⁸²<https://jira.mongodb.org/browse/SERVER-3875>

¹⁸³<https://jira.mongodb.org/browse/SERVER-3871>

¹⁸⁴<https://jira.mongodb.org/browse/SERVER-3937>

¹⁸⁵<https://jira.mongodb.org/browse/SERVER-3854>

¹⁸⁶<https://jira.mongodb.org/browse/SERVER-4593>

```
mongodump --db local --collection oplog.rs --query '{"ts":{"$gt":{"$timestamp" : {"t": 1344969612000,
```

See [SERVER-3483](#)¹⁸⁷ for more information.

Shell Improvements

Improved Shell User Interface 2.2 includes a number of changes that improve the overall quality and consistency of the user interface for the `mongo` (page 527) shell:

- Full Unicode support.
- Bash-like line editing features. See [SERVER-4312](#)¹⁸⁸ for more information.
- Multi-line command support in shell history. See [SERVER-3470](#)¹⁸⁹ for more information.
- Windows support for the `edit` command. See [SERVER-3998](#)¹⁹⁰ for more information.

Helper to load Server-Side Functions The `db.loadServerScripts()` (page 115) loads the contents of the current database's `system.js` collection into the current `mongo` (page 527) shell session. See [SERVER-1651](#)¹⁹¹ for more information.

Support for Bulk Inserts If you pass an array of *documents* to the `insert()` (page 52) method, the `mongo` (page 527) shell will now perform a bulk insert operation. See [SERVER-3819](#)¹⁹² and [SERVER-2395](#)¹⁹³ for more information.

Note: For bulk inserts on sharded clusters, the `getLastError` (page 235) command alone is insufficient to verify success. Applications should must verify the success of bulk inserts in application logic.

Operations

Support for Logging to Syslog See the [SERVER-2957](#)¹⁹⁴ case and the documentation of the `syslogFacility` run-time option or the `mongod --syslog` and `mongos --syslog` command line-options.

touch Command Added the `touch` (page 321) command to read the data and/or indexes from a collection into memory. See: [SERVER-2023](#)¹⁹⁵ and `touch` (page 321) for more information.

indexCounters No Longer Report Sampled Data `indexCounters` now report actual counters that reflect index use and state. In previous versions, these data were sampled. See [SERVER-5784](#)¹⁹⁶ and `indexCounters` for more information.

¹⁸⁷<https://jira.mongodb.org/browse/SERVER-3483>

¹⁸⁸<https://jira.mongodb.org/browse/SERVER-4312>

¹⁸⁹<https://jira.mongodb.org/browse/SERVER-3470>

¹⁹⁰<https://jira.mongodb.org/browse/SERVER-3998>

¹⁹¹<https://jira.mongodb.org/browse/SERVER-1651>

¹⁹²<https://jira.mongodb.org/browse/SERVER-3819>

¹⁹³<https://jira.mongodb.org/browse/SERVER-2395>

¹⁹⁴<https://jira.mongodb.org/browse/SERVER-2957>

¹⁹⁵<https://jira.mongodb.org/browse/SERVER-2023>

¹⁹⁶<https://jira.mongodb.org/browse/SERVER-5784>

Padding Specifiable on compact Command See the documentation of the `compact` (page 313) and the [SERVER-4018](#)¹⁹⁷ issue for more information.

Added Build Flag to Use System Libraries The Boost library, version 1.49, is now embedded in the MongoDB code base.

If you want to build MongoDB binaries using system Boost libraries, you can pass `scons` using the `--use-system-boost` flag, as follows:

```
scons --use-system-boost
```

When building MongoDB, you can also pass `scons` a flag to compile MongoDB using only system libraries rather than the included versions of the libraries. For example:

```
scons --use-system-all
```

See the [SERVER-3829](#)¹⁹⁸ and [SERVER-5172](#)¹⁹⁹ issues for more information.

Memory Allocator Changed to TCMalloc To improve performance, MongoDB 2.2 uses the TCMalloc memory allocator from Google Perftools. For more information about this change see the [SERVER-188](#)²⁰⁰ and [SERVER-4683](#)²⁰¹. For more information about TCMalloc, see the documentation of [TCMalloc](#)²⁰² itself.

Replication

Improved Logging for Replica Set Lag When *secondary* members of a replica set fall behind in replication, `mongod` (page 503) now provides better reporting in the log. This makes it possible to track replication in general and identify what process may produce errors or halt replication. See [SERVER-3575](#)²⁰³ for more information.

Replica Set Members can Sync from Specific Members The new `replSetSyncFrom` (page 278) command and new `rs.syncFrom()` (page 169) helper in the `mongo` (page 527) shell make it possible for you to manually configure from which member of the set a replica will poll *oplog* entries. Use these commands to override the default selection logic if needed. Always exercise caution with `replSetSyncFrom` (page 278) when overriding the default behavior.

Replica Set Members will not Sync from Members Without Indexes Unless `buildIndexes: false` To prevent inconsistency between members of replica sets, if the member of a replica set has `buildIndexes` set to `true`, other members of the replica set will *not* sync from this member, unless they also have `buildIndexes` set to `true`. See [SERVER-4160](#)²⁰⁴ for more information.

New Option To Configure Index Pre-Fetching during Replication By default, when replicating options, *secondaries* will pre-fetch *indexes* associated with a query to improve replication throughput in most cases. The `replication.secondaryIndexPrefetch` setting and `--replIndexPrefetch` option allow administrators to disable this feature or allow the `mongod` (page 503) to pre-fetch only the index on the `_id` field. See [SERVER-6718](#)²⁰⁵ for more information.

¹⁹⁷<https://jira.mongodb.org/browse/SERVER-4018>

¹⁹⁸<https://jira.mongodb.org/browse/SERVER-3829>

¹⁹⁹<https://jira.mongodb.org/browse/SERVER-5172>

²⁰⁰<https://jira.mongodb.org/browse/SERVER-188>

²⁰¹<https://jira.mongodb.org/browse/SERVER-4683>

²⁰²<http://goog-perftools.sourceforge.net/doc/tcmalloc.html>

²⁰³<https://jira.mongodb.org/browse/SERVER-3575>

²⁰⁴<https://jira.mongodb.org/browse/SERVER-4160>

²⁰⁵<https://jira.mongodb.org/browse/SERVER-6718>

Map Reduce Improvements

In 2.2 Map Reduce received the following improvements:

- Improved support for sharded MapReduce²⁰⁶, and
- MapReduce will retry jobs following a config error²⁰⁷.

Sharding Improvements

Index on Shard Keys Can Now Be a Compound Index If your shard key uses the prefix of an existing index, then you do not need to maintain a separate index for your shard key in addition to your existing index. This index, however, cannot be a multi-key index. See the *sharding-shard-key-indexes* documentation and [SERVER-1506](#)²⁰⁸ for more information.

Migration Thresholds Modified The *migration thresholds* have changed in 2.2 to permit more even distribution of *chunks* in collections that have smaller quantities of data. See the *sharding-migration-thresholds* documentation for more information.

Licensing Changes

Added License notice for Google Perftools (TCMalloc Utility). See the [License Notice](#)²⁰⁹ and the [SERVER-4683](#)²¹⁰ for more information.

Resources

- [MongoDB Downloads](#)²¹¹.
- [All JIRA issues resolved in 2.2](#)²¹².
- [All backwards incompatible changes](#)²¹³.
- [All third party license notices](#)²¹⁴.
- [What's New in MongoDB 2.2 Online Conference](#)²¹⁵.

7.2.3 Release Notes for MongoDB 2.0

Upgrading

Although the major version number has changed, MongoDB 2.0 is a standard, incremental production release and works as a drop-in replacement for MongoDB 1.8.

²⁰⁶<https://jira.mongodb.org/browse/SERVER-4521>

²⁰⁷<https://jira.mongodb.org/browse/SERVER-4158>

²⁰⁸<https://jira.mongodb.org/browse/SERVER-1506>

²⁰⁹<https://github.com/mongodb/mongo/blob/v2.2/distsrc/THIRD-PARTY-NOTICES#L231>

²¹⁰<https://jira.mongodb.org/browse/SERVER-4683>

²¹¹<http://mongodb.org/downloads>

²¹²<https://jira.mongodb.org/secure/IssueNavigator.jspa?reset=true&jqlQuery=project+%3D+SERVER+AND+fixVersion+in+%28%222.1.0%22%2C+%222.1.1%22%2C+%222.2.0-rc1%22%2C+%222.2.0-rc2%22%29+ORDER+BY+component+ASC%2C+key+DESC>

²¹³<https://jira.mongodb.org/secure/IssueNavigator.jspa?requestId=11225>

²¹⁴<https://github.com/mongodb/mongo/blob/v2.2/distsrc/THIRD-PARTY-NOTICES>

²¹⁵<http://www.mongodb.com/events/webinar/mongodb-online-conference-sept>

Preparation

Read through all release notes before upgrading, and ensure that no changes will affect your deployment.

If you create new indexes in 2.0, then downgrading to 1.8 is possible but you must reindex the new collections.

`mongoimport` (page 556) and `mongoexport` (page 562) now correctly adhere to the CSV spec for handling CSV input/output. This may break existing import/export workflows that relied on the previous behavior. For more information see [SERVER-1097](https://jira.mongodb.org/browse/SERVER-1097)²¹⁶.

Journaling²¹⁷ is **enabled by default** in 2.0 for 64-bit builds. If you still prefer to run without journaling, start `mongod` (page 503) with the `--nojournal` run-time option. Otherwise, MongoDB creates journal files during startup. The first time you start `mongod` (page 503) with journaling, you will see a delay as `mongod` (page 503) creates new files. In addition, you may see reduced write throughput.

2.0 `mongod` (page 503) instances are interoperable with 1.8 `mongod` (page 503) instances; however, for best results, upgrade your deployments using the following procedures:

Upgrading a Standalone `mongod`

1. Download the v2.0.x binaries from the [MongoDB Download Page](https://www.mongodb.org/downloads)²¹⁸.
2. Shutdown your `mongod` (page 503) instance. Replace the existing binary with the 2.0.x `mongod` (page 503) binary and restart MongoDB.

Upgrading a Replica Set

1. Upgrade the *secondary* members of the set one at a time by shutting down the `mongod` (page 503) and replacing the 1.8 binary with the 2.0.x binary from the [MongoDB Download Page](https://www.mongodb.org/downloads)²¹⁹.
2. To avoid losing the last few updates on failover you can temporarily halt your application (failover should take less than 10 seconds), or you can set *write concern* in your application code to confirm that each update reaches multiple servers.
3. Use the `rs.stepDown()` (page 168) to step down the primary to allow the normal *failover* procedure.

`rs.stepDown()` (page 168) and `replSetStepDown` (page 277) provide for shorter and more consistent failover procedures than simply shutting down the primary directly.

When the primary has stepped down, shut down its instance and upgrade by replacing the `mongod` (page 503) binary with the 2.0.x binary.

Upgrading a Sharded Cluster

1. Upgrade all *config server* instances *first*, in any order. Since config servers use two-phase commit, *shard* configuration metadata updates will halt until all are up and running.
2. Upgrade `mongos` (page 518) routers in any order.

²¹⁶<https://jira.mongodb.org/browse/SERVER-1097>

²¹⁷<http://www.mongodb.org/display/DOCS/Journaling>

²¹⁸<http://downloads.mongodb.org/>

²¹⁹<http://downloads.mongodb.org/>

Changes

Compact Command

A `compact` (page 313) command is now available for compacting a single collection and its indexes. Previously, the only way to compact was to repair the entire database.

Concurrency Improvements

When going to disk, the server will yield the write lock when writing data that is not likely to be in memory. The initial implementation of this feature now exists:

See [SERVER-2563](https://jira.mongodb.org/browse/SERVER-2563)²²⁰ for more information.

The specific operations yield in 2.0 are:

- Updates by `_id`
- Removes
- Long cursor iterations

Default Stack Size

MongoDB 2.0 reduces the default stack size. This change can reduce total memory usage when there are many (e.g., 1000+) client connections, as there is a thread per connection. While portions of a thread's stack can be swapped out if unused, some operating systems do this slowly enough that it might be an issue. The default stack size is lesser of the system setting or 1MB.

Index Performance Enhancements

v2.0 includes significant improvements to the `index`. Indexes are often 25% smaller and 25% faster (depends on the use case). When upgrading from previous versions, the benefits of the new index type are realized only if you create a new index or re-index an old one.

Dates are now signed, and the max index key size has increased slightly from 819 to 1024 bytes.

All operations that create a new index will result in a 2.0 index by default. For example:

- Reindexing results on an older-version index results in a 2.0 index. However, reindexing on a secondary does *not* work in versions prior to 2.0. Do not reindex on a secondary. For a workaround, see [SERVER-3866](https://jira.mongodb.org/browse/SERVER-3866)²²¹.
- The `repairDatabase` (page 319) command converts indexes to a 2.0 indexes.

To convert all indexes for a given collection to the *2.0 type* (page 678), invoke the `compact` (page 313) command.

Once you create new indexes, downgrading to 1.8.x will require a re-index of any indexes created using 2.0. See <http://docs.mongodb.org/manual/tutorial/roll-back-to-v1.8-index>.

Sharding Authentication

Applications can now use authentication with *sharded clusters*.

²²⁰<https://jira.mongodb.org/browse/SERVER-2563>

²²¹<https://jira.mongodb.org/browse/SERVER-3866>

Replica Sets

Hidden Nodes in Sharded Clusters In 2.0, `mongos` (page 518) instances can now determine when a member of a replica set becomes “hidden” without requiring a restart. In 1.8, `mongos` (page 518) if you reconfigured a member as hidden, you *had* to restart `mongos` (page 518) to prevent queries from reaching the hidden member.

Priorities Each *replica set* member can now have a priority value consisting of a floating-point from 0 to 1000, inclusive. Priorities let you control which member of the set you prefer to have as *primary* the member with the highest priority that can see a majority of the set will be elected primary.

For example, suppose you have a replica set with three members, A, B, and C, and suppose that their priorities are set as follows:

- A’s priority is 2.
- B’s priority is 3.
- C’s priority is 1.

During normal operation, the set will always chose B as primary. If B becomes unavailable, the set will elect A as primary.

For more information, see the [priority](#) documentation.

Data-Center Awareness You can now “tag” *replica set* members to indicate their location. You can use these tags to design custom *write rules* across data centers, racks, specific servers, or any other architecture choice.

For example, an administrator can define rules such as “very important write” or `customerData` or “audit-trail” to replicate to certain servers, racks, data centers, etc. Then in the application code, the developer would say:

```
db.foo.insert(doc, {w : "very important write"})
```

which would succeed if it fulfilled the conditions the DBA defined for “very important write”.

For more information, see [Tagging](#)²²².

Drivers may also support tag-aware reads. Instead of specifying `slaveOk`, you specify `slaveOk` with tags indicating which data-centers to read from. For details, see the <http://docs.mongodb.org/manualapplications/drivers> documentation.

w: majority You can also set `w` to `majority` to ensure that the write propagates to a majority of nodes, effectively committing it. The value for “majority” will automatically adjust as you add or remove nodes from the set.

For more information, see <http://docs.mongodb.org/manualcore/write-concern>.

Reconfiguration with a Minority Up If the majority of servers in a set has been permanently lost, you can now force a reconfiguration of the set to bring it back online.

For more information see <http://docs.mongodb.org/manualtutorial/reconfigure-replica-set-with-unava>.

Primary Checks for a Caught up Secondary before Stepping Down To minimize time without a *primary*, the `rs.stepDown()` (page 168) method will now fail if the primary does not see a *secondary* within 10 seconds of its latest optime. You can force the primary to step down anyway, but by default it will return an error message.

See also <http://docs.mongodb.org/manualtutorial/force-member-to-be-primary>.

²²²<http://www.mongodb.org/display/DOCS/Data+Center+Awareness#DataCenterAwareness-Tagging%28version2.0%29>

Extended Shutdown on the Primary to Minimize Interruption When you call the `shutdown` (page 321) command, the *primary* will refuse to shut down unless there is a *secondary* whose optime is within 10 seconds of the primary. If such a secondary isn't available, the primary will step down and wait up to a minute for the secondary to be fully caught up before shutting down.

Note that to get this behavior, you must issue the `shutdown` (page 321) command explicitly; sending a signal to the process will not trigger this behavior.

You can also force the primary to shut down, even without an up-to-date secondary available.

Maintenance Mode When `repair` or `compact` (page 313) runs on a *secondary*, the secondary will automatically drop into “recovering” mode until the operation finishes. This prevents clients from trying to read from it while it's busy.

Geospatial Features

Multi-Location Documents Indexing is now supported on documents which have multiple location objects, embedded either inline or in nested sub-documents. Additional command options are also supported, allowing results to return with not only distance but the location used to generate the distance.

For more information, see [Multi-location Documents](#)²²³.

Polygon searches Polygonal `$within` (page 393) queries are also now supported for simple polygon shapes. For details, see the `$within` (page 393) operator documentation.

Journaling Enhancements

- Journaling is now enabled by default for 64-bit platforms. Use the `--nojournal` command line option to disable it.
- The journal is now compressed for faster commits to disk.
- A new `--journalCommitInterval` run-time option exists for specifying your own group commit interval. The default settings do not change.
- A new `{ getLastError: { j: true } }` (page 235) option is available to wait for the group commit. The group commit will happen sooner when a client is waiting on `{ j: true }`. If journaling is disabled, `{ j: true }` is a no-op.

New ContinueOnError Option for Bulk Insert

Set the `continueOnError` option for bulk inserts, in the driver, so that bulk insert will continue to insert any remaining documents even if an insert fails, as is the case with duplicate key exceptions or network interruptions. The `getLastError` (page 235) command will report whether any inserts have failed, not just the last one. If multiple errors occur, the client will only receive the most recent `getLastError` (page 235) results.

See [OP_INSERT](#)²²⁴.

Note: For bulk inserts on sharded clusters, the `getLastError` (page 235) command alone is insufficient to verify success. Applications should must verify the success of bulk inserts in application logic.

²²³<http://www.mongodb.org/display/DOCS/Geospatial+Indexing#GeospatialIndexing-MultilocationDocuments>

²²⁴<http://www.mongodb.org/display/DOCS/Mongo+Wire+Protocol#MongoWireProtocol-OPINSERT>

Map Reduce

Output to a Sharded Collection Using the new `sharded` flag, it is possible to send the result of a map/reduce to a sharded collection. Combined with the `reduce` or `merge` flags, it is possible to keep adding data to very large collections from map/reduce jobs.

For more information, see [MapReduce Output Options](#)²²⁵ and *mapReduce* (page 208).

Performance Improvements Map/reduce performance will benefit from the following:

- Larger in-memory buffer sizes, reducing the amount of disk I/O needed during a job
- Larger javascript heap size, allowing for larger objects and less GC
- Supports pure JavaScript execution with the `jsMode` flag. See *mapReduce* (page 208).

New Querying Features

Additional regex options: `s` Allows the dot (`.`) to match all characters including new lines. This is in addition to the currently supported `i`, `m` and `x`. See [Regular Expressions](#)²²⁶ and `$regex` (page 386).

`$and` A special boolean `$and` (page 378) query operator is now available.

Command Output Changes

The output of the `validate` (page 335) command and the documents in the `system.profile` collection have both been enhanced to return information as BSON objects with keys for each value rather than as free-form strings.

Shell Features

Custom Prompt You can define a custom prompt for the `mongo` (page 527) shell. You can change the prompt at any time by setting the prompt variable to a string or a custom JavaScript function returning a string. For examples, see [Custom Prompt](#)²²⁷.

Default Shell Init Script On startup, the shell will check for a `.mongorc.js` file in the user's home directory. The shell will execute this file after connecting to the database and before displaying the prompt.

If you would like the shell not to run the `.mongorc.js` file automatically, start the shell with `--norc`.

For more information, see *mongo* (page 527).

Most Commands Require Authentication

In 2.0, when running with authentication (e.g. `authorization`) *all* database commands require authentication, *except* the following commands.

- `isMaster` (page 280)

²²⁵<http://www.mongodb.org/display/DOCS/MapReduce#MapReduce-Outputoptions>

²²⁶<http://www.mongodb.org/display/DOCS/Advanced+Queries#AdvancedQueries-RegularExpressions>

²²⁷<http://www.mongodb.org/display/DOCS/Overview+-+The+MongoDB+Interactive+Shell#Overview-TheMongoDBInteractiveShell-CustomPrompt>

- [authenticate](#) (page 249)
- [getnonce](#) (page 250)
- [buildInfo](#) (page 324)
- [ping](#) (page 334)
- [isdbgrid](#) (page 298)

Resources

- [MongoDB Downloads](#)²²⁸
- [All JIRA Issues resolved in 2.0](#)²²⁹
- [All Backward Incompatible Changes](#)²³⁰

7.2.4 Release Notes for MongoDB 1.8

Upgrading

MongoDB 1.8 is a standard, incremental production release and works as a drop-in replacement for MongoDB 1.6, except:

- *Replica set* members should be upgraded in a particular order, as described in *Upgrading a Replica Set* (page 683).
- The `mapReduce` (page 208) command has changed in 1.8, causing incompatibility with previous releases. `mapReduce` (page 208) no longer generates temporary collections (thus, `keepTemp` has been removed). Now, you must always supply a value for `out`. See the `out` field options in the `mapReduce` (page 208) document. If you use MapReduce, this also likely means you need a recent version of your client driver.

Preparation

Read through all release notes before upgrading and ensure that no changes will affect your deployment.

Upgrading a Standalone `mongod`

1. Download the v1.8.x binaries from the [MongoDB Download Page](#)²³¹.
2. Shutdown your `mongod` (page 503) instance.
3. Replace the existing binary with the 1.8.x `mongod` (page 503) binary.
4. Restart MongoDB.

²²⁸<http://mongodb.org/downloads>

²²⁹<https://jira.mongodb.org/secure/IssueNavigator.jspa?mode=hide&requestId=11002>

²³⁰<https://jira.mongodb.org/secure/IssueNavigator.jspa?requestId=11023>

²³¹<http://downloads.mongodb.org/>

Upgrading a Replica Set

1.8.x *secondaries* **can** replicate from 1.6.x *primaries*.

1.6.x secondaries **cannot** replicate from 1.8.x primaries.

Thus, to upgrade a *replica set* you must replace all of your secondaries first, then the primary.

For example, suppose you have a replica set with a primary, an *arbiter* and several secondaries. To upgrade the set, do the following:

1. For the arbiter:
 - (a) Shut down the arbiter.
 - (b) Restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³².
2. Change your config (optional) to prevent election of a new primary.

It is possible that, when you start shutting down members of the set, a new primary will be elected. To prevent this, you can give all of the secondaries a priority of 0 before upgrading, and then change them back afterwards. To do so:

- (a) Record your current config. Run `rs.config()` (page 165) and paste the results into a text file.
- (b) Update your config so that all secondaries have priority 0. For example:

```
config = rs.conf()
{
  "_id" : "foo",
  "version" : 3,
  "members" : [
    {
      "_id" : 0,
      "host" : "ubuntu:27017"
    },
    {
      "_id" : 1,
      "host" : "ubuntu:27018"
    },
    {
      "_id" : 2,
      "host" : "ubuntu:27019",
      "arbiterOnly" : true
    },
    {
      "_id" : 3,
      "host" : "ubuntu:27020"
    },
    {
      "_id" : 4,
      "host" : "ubuntu:27021"
    }
  ]
}
config.version++
3
rs.isMaster()
{
  "setName" : "foo",
```

²³²<http://downloads.mongodb.org/>

```
"ismaster" : false,
"secondary" : true,
"hosts" : [
  "ubuntu:27017",
  "ubuntu:27018"
],
"arbiters" : [
  "ubuntu:27019"
],
"primary" : "ubuntu:27018",
"ok" : 1
}
// for each secondary
config.members[0].priority = 0
config.members[3].priority = 0
config.members[4].priority = 0
rs.reconfig(config)
```

3. For each secondary:

- (a) Shut down the secondary.
- (b) Restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³³.

4. If you changed the config, change it back to its original state:

```
config = rs.conf()
config.version++
config.members[0].priority = 1
config.members[3].priority = 1
config.members[4].priority = 1
rs.reconfig(config)
```

5. Shut down the primary (the final 1.6 server), and then restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³⁴.

Upgrading a Sharded Cluster

1. Turn off the balancer:

```
mongo <a_mongos_hostname>
use config
db.settings.update({_id:"balancer"},{$set : {stopped:true}}, true)
```

2. For each *shard*:

- If the shard is a *replica set*, follow the directions above for *Upgrading a Replica Set* (page 683).
- If the shard is a single *mongod* (page 503) process, shut it down and then restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³⁵.

3. For each *mongos* (page 518):

- (a) Shut down the *mongos* (page 518) process.
- (b) Restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³⁶.

²³³<http://downloads.mongodb.org/>

²³⁴<http://downloads.mongodb.org/>

²³⁵<http://downloads.mongodb.org/>

²³⁶<http://downloads.mongodb.org/>

4. For each config server:
 - (a) Shut down the config server process.
 - (b) Restart it with the 1.8.x binary from the [MongoDB Download Page](#)²³⁷.

5. Turn on the balancer:

```
use config
db.settings.update({_id:"balancer"},{$set : {stopped:false}})
```

Returning to 1.6

If for any reason you must move back to 1.6, follow the steps above in reverse. Please be careful that you have not inserted any documents larger than 4MB while running on 1.8 (where the max size has increased to 16MB). If you have you will get errors when the server tries to read those documents.

Journaling Returning to 1.6 after using 1.8 Journaling works fine, as journaling does not change anything about the data file format. Suppose you are running 1.8.x with journaling enabled and you decide to switch back to 1.6. There are two scenarios:

- If you shut down cleanly with 1.8.x, just restart with the 1.6 mongod binary.
- If 1.8.x shut down uncleanly, start 1.8.x up again and let the journal files run to fix any damage (incomplete writes) that may have existed at the crash. Then shut down 1.8.x cleanly and restart with the 1.6 mongod binary.

Changes

Journaling

MongoDB now supports write-ahead <http://docs.mongodb.org/manualcore/journaling> to facilitate fast crash recovery and durability in the storage engine. With journaling enabled, a [mongod](#) (page 503) can be quickly restarted following a crash without needing to repair the [collections](#). The aggregation framework makes it possible to do aggregation

Sparse and Covered Indexes

Sparse Indexes are indexes that only include documents that contain the fields specified in the index. Documents missing the field will not appear in the index at all. This can significantly reduce index size for indexes of fields that contain only a subset of documents within a [collection](#).

Covered Indexes enable MongoDB to answer queries entirely from the index when the query only selects fields that the index contains.

Incremental MapReduce Support

The [mapReduce](#) (page 208) command supports new options that enable incrementally updating existing [collections](#). Previously, a MapReduce job could output either to a temporary collection or to a named permanent collection, which it would overwrite with new data.

You now have several options for the output of your MapReduce jobs:

²³⁷<http://downloads.mongodb.org/>

- You can merge MapReduce output into an existing collection. Output from the Reduce phase will replace existing keys in the output collection if it already exists. Other keys will remain in the collection.
- You can now re-reduce your output with the contents of an existing collection. Each key output by the reduce phase will be reduced with the existing document in the output collection.
- You can replace the existing output collection with the new results of the MapReduce job (equivalent to setting a permanent output collection in previous releases)
- You can compute MapReduce inline and return results to the caller without persisting the results of the job. This is similar to the temporary collections generated in previous releases, except results are limited to 8MB.

For more information, see the `out` field options in the [mapReduce](#) (page 208) document.

Additional Changes and Enhancements

1.8.1

- Sharding migrate fix when moving larger chunks.
- Durability fix with background indexing.
- Fixed mongos concurrency issue with many incoming connections.

1.8.0

- All changes from 1.7.x series.

1.7.6

- Bug fixes.

1.7.5

- Journaling.
- Extent allocation improvements.
- Improved *replica set* connectivity for [mongos](#) (page 518).
- `getLastError` (page 235) improvements for *sharding*.

1.7.4

- [mongos](#) (page 518) routes `slaveOk` queries to *secondaries* in *replica sets*.
- New [mapReduce](#) (page 208) output options.
- *index-type-sparse*.

1.7.3

- Initial *covered index* support.
- Distinct can use data from indexes when possible.
- [mapReduce](#) (page 208) can merge or reduce results into an existing collection.
- [mongod](#) (page 503) tracks and [mongostat](#) (page 570) displays network usage. See [mongostat](#) (page 569).

- Sharding stability improvements.

1.7.2

- `$rename` (page 414) operator allows renaming of fields in a document.
- `db.eval()` (page 108) not to block.
- Geo queries with sharding.
- `mongostat --discover` option
- Chunk splitting enhancements.
- Replica sets network enhancements for servers behind a nat.

1.7.1

- Many sharding performance enhancements.
- Better support for `$elemMatch` (page 408) on primitives in embedded arrays.
- Query optimizer enhancements on range queries.
- Window service enhancements.
- Replica set setup improvements.
- `$pull` (page 424) works on primitives in arrays.

1.7.0

- Sharding performance improvements for heavy insert loads.
- Slave delay support for replica sets.
- `getLastErrorDefaults` for replica sets.
- Auto completion in the shell.
- Spherical distance for geo search.
- All fixes from 1.6.1 and 1.6.2.

Release Announcement Forum Pages

- [1.8.1²³⁸](#), [1.8.0²³⁹](#)
- [1.7.6²⁴⁰](#), [1.7.5²⁴¹](#), [1.7.4²⁴²](#), [1.7.3²⁴³](#), [1.7.2²⁴⁴](#), [1.7.1²⁴⁵](#), [1.7.0²⁴⁶](#)

²³⁸<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/v09MbEm62Y>

²³⁹<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/JeHQOnam6Qk>

²⁴⁰<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/3t6GNZ1qGYc>

²⁴¹<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/S5R0Tx9wkEg>

²⁴²<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/9Om3Vuw-y9c>

²⁴³<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/DfNUrdbmflI>

²⁴⁴<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/df7mwK6Xixo>

²⁴⁵<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/HUR9zYtTpA8>

²⁴⁶<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/TUnJCg9161A>

Resources

- [MongoDB Downloads](#)²⁴⁷
- [All JIRA Issues resolved in 1.8](#)²⁴⁸

7.2.5 Release Notes for MongoDB 1.6

Upgrading

MongoDB 1.6 is a drop-in replacement for 1.4. To upgrade, simply shutdown `mongod` (page 503) then restart with the new binaries.

Please note that you should upgrade to the latest version of whichever driver you're using. Certain drivers, including the Ruby driver, will require the upgrade, and all the drivers will provide extra features for connecting to replica sets.

Sharding

<http://docs.mongodb.org/manualsharding> is now production-ready, making MongoDB horizontally scalable, with no single point of failure. A single instance of `mongod` (page 503) can now be upgraded to a distributed cluster with zero downtime when the need arises.

- <http://docs.mongodb.org/manualsharding>
- <http://docs.mongodb.org/manualtutorial/deploy-shard-cluster>
- <http://docs.mongodb.org/manualtutorial/convert-replica-set-to-replicated-shard-cluster>

Replica Sets

Replica sets, which provide automated failover among a cluster of *n* nodes, are also now available.

Please note that replica pairs are now deprecated; we strongly recommend that replica pair users upgrade to replica sets.

- <http://docs.mongodb.org/manualreplication>
- <http://docs.mongodb.org/manualtutorial/deploy-replica-set>
- <http://docs.mongodb.org/manualtutorial/convert-standalone-to-replica-set>

Other Improvements

- The `w` option (and `wtimeout`) forces writes to be propagated to *n* servers before returning success (this works especially well with replica sets)
- *\$or queries* (page 377)
- Improved concurrency
- *\$slice* (page 411) operator for returning subsets of arrays
- 64 indexes per collection (formerly 40 indexes per collection)
- 64-bit integers can now be represented in the shell using `NumberLong`

²⁴⁷<http://mongodb.org/downloads>

²⁴⁸<https://jira.mongodb.org/secure/IssueNavigator.jspa?mode=hide&requestId=10172>

- The `findAndModify` (page 229) command now supports upserts. It also allows you to specify fields to return
- `$showDiskLoc` option to see disk location of a document
- Support for IPv6 and UNIX domain sockets

Installation

- Windows service improvements
- The C++ client is a separate tarball from the binaries

1.6.x Release Notes

- 1.6.5²⁴⁹

1.5.x Release Notes

- 1.5.8²⁵⁰
- 1.5.7²⁵¹
- 1.5.6²⁵²
- 1.5.5²⁵³
- 1.5.4²⁵⁴
- 1.5.3²⁵⁵
- 1.5.2²⁵⁶
- 1.5.1²⁵⁷
- 1.5.0²⁵⁸

You can see a full list of all changes on [JIRA](#)²⁵⁹.

Thank you everyone for your support and suggestions!

7.2.6 Release Notes for MongoDB 1.4

Upgrading

We're pleased to announce the 1.4 release of MongoDB. 1.4 is a drop-in replacement for 1.2. To upgrade you just need to shutdown `mongod` (page 503), then restart with the new binaries. (Users upgrading from release 1.0 should review the [1.2 release notes](#) (page 691), in particular the instructions for upgrading the DB format.)

²⁴⁹https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/06_QCC05Fpk

²⁵⁰<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/uJfF1QN6Thk>

²⁵¹<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/OYvz40RWs90>

²⁵²https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/4l0N2U_H0cQ

²⁵³<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/oO749nvTARY>

²⁵⁴https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/380V_Ec_q1c

²⁵⁵<https://groups.google.com/forum/?hl=en&fromgroups=#!topic/mongodb-user/hsUQL9CxTQw>

²⁵⁶<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/94EE3HVidAA>

²⁵⁷<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/7SBPQ2RSfdM>

²⁵⁸<https://groups.google.com/forum/?fromgroups=#!topic/mongodb-user/VAhJcjDGTy0>

²⁵⁹<https://jira.mongodb.org/secure/IssueNavigator.jspa?mode=hide&requestId=10107>

Release 1.4 includes the following improvements over release 1.2:

Core Server Enhancements

- concurrency improvements
- indexing memory improvements
- *background index creation*
- better detection of regular expressions so the index can be used in more cases

Replication and Sharding

- better handling for restarting slaves offline for a while
- fast new slaves from snapshots (`--fastsync`)
- configurable slave delay (`--slavedelay`)
- replication handles clock skew on master
- *\$inc* (page 412) replication fixes
- sharding alpha 3 - notably 2-phase commit on config servers

Deployment and Production

- *configure “slow threshold” for profiling*
- ability to do *fsync + lock* (page 312) for backing up raw files
- option for separate directory per db (`--directoryperdb`)
- `http://localhost:28017/_status` to get `serverStatus` via http
- REST interface is off by default for security (`--rest` to enable)
- can rotate logs with a db command, *logRotate* (page 322)
- enhancements to *serverStatus* (page 346) command (`db.serverStatus()`) - counters and *replication lag* stats
- new *mongostat* (page 569) tool

Query Language Improvements

- *\$all* (page 402) with regex
- *\$not* (page 379)
- partial matching of array elements *\$elemMatch* (page 408)
- *\$* operator for updating arrays
- *\$addToSet* (page 422)
- *\$unset* (page 417)
- *\$pull* (page 424) supports object matching
- *\$set* (page 416) with array indexes

Geo

- 2d geospatial search
- geo *\$center* (page 397) and *\$box* (page 399) searches

7.2.7 Release Notes for MongoDB 1.2.x

New Features

- More indexes per collection
- Faster index creation
- Map/Reduce
- Stored JavaScript functions
- Configurable fsync time
- Several small features and fixes

DB Upgrade Required

There are some changes that will require doing an upgrade if your previous version is $\leq 1.0.x$. If you're already using a version $\geq 1.1.x$ then these changes aren't required. There are 2 ways to do it:

- `--upgrade`
 - stop your *mongod* (page 503) process
 - run `./mongod --upgrade`
 - start *mongod* (page 503) again
- use a slave
 - start a slave on a different port and data directory
 - when its synced, shut down the master, and start the new slave on the regular port.

Ask in the forums or IRC for more help.

Replication Changes

- There have been minor changes in replication. If you are upgrading a master/slave setup from $\leq 1.1.2$ you have to update the slave first.

mongoimport

- `mongoimport json` has been removed and is replaced with *mongoimport* (page 556) that can do json/csv/tsv

field filter changing

- We've changed the semantics of the field filter a little bit. Previously only objects with those fields would be returned. Now the field filter only changes the output, not which objects are returned. If you need that behavior, you can use *\$exists* (page 381)

7.3 Other MongoDB Release Notes

7.3.1 Default Write Concern Change

These release notes outline a change to all driver interfaces released in November 2012. See release notes for specific drivers for additional information.

Changes

As of the releases listed below, there are two major changes to all drivers:

1. All drivers will add a new top-level connection class that will increase consistency for all MongoDB client interfaces.

This change is non-backward breaking: existing connection classes will remain in all drivers for a time, and will continue to operate as expected. However, those previous connection classes are now deprecated as of these releases, and will eventually be removed from the driver interfaces.

The new top-level connection class is named `MongoClient`, or similar depending on how host languages handle namespacing.

2. The default write concern on the new `MongoClient` class will be to acknowledge all write operations²⁶⁰. This will allow your application to receive acknowledgment of all write operations.

See the documentation of *Write Concern* for more information about write concern in MongoDB.

Please migrate to the new `MongoClient` class expeditiously.

Releases

The following driver releases will include the changes outlined in *Changes* (page 692). See each driver's release notes for a full account of each release as well as other related driver-specific changes.

- C#, version 1.7
- Java, version 2.10.0
- Node.js, version 1.2
- Perl, version 0.501.1
- PHP, version 1.4
- Python, version 2.4
- Ruby, version 1.8

²⁶⁰ The drivers will call `getLastError` (page 235) without arguments, which is logically equivalent to the `w: 1` option; however, this operation allows *replica set* users to override the default write concern with the `getLastErrorDefaults` setting in the <http://docs.mongodb.org/manualreference/replica-configuration>.

Symbols

- {-}all
 - command line option 573
- {-}auditDestination
 - command line option 517, 526
- {-}auditFilter
 - command line option 518, 526
- {-}auditFormat
 - command line option 517, 526
- {-}auditPath
 - command line option 518, 526
- {-}auth
 - command line option 507
- {-}authenticationDatabase <dbname>
 - command line option 529, 540, 546, 553, 558, 565, 572, 578, 588
- {-}authenticationMechanism <name>
 - command line option 529, 540, 546, 554, 559, 565, 572, 578, 588
- {-}autoresync
 - command line option 513
- {-}bind_ip <ip address>
 - command line option 504, 519
- {-}chunkSize <value>
 - command line option 522
- {-}clusterAuthMode <option>
 - command line option 515, 524
- {-}collection <collection>, -c
 - command line option 541, 547, 555, 559, 566, 589
- {-}config <filename>, -f
 - command line option 503, 519
- {-}configdb <config1>,<config2>,<config3>
 - command line option 522
- {-}configsrv
 - command line option 513
- {-}cpu
 - command line option 508
- {-}csv
 - command line option 566
- {-}db <database>, -d
 - command line option 541, 547, 554, 559, 566, 589
- {-}dbpath <path>
 - command line option 508, 540, 546, 554, 559, 565, 588
- {-}diaglog <value>
 - command line option 505
- {-}directoryperdb
 - command line option 508, 541, 547, 554, 559, 565, 589
- {-}discover
 - command line option 573
- {-}drop
 - command line option 547, 560
- {-}dumpDbUsersAndRoles
 - command line option 542
- {-}eval <javascript>
 - command line option 528
- {-}fastsync
 - command line option 513
- {-}fieldFile <filename>
 - command line option 560, 566
- {-}fields <field1[,field2]>,-f
 - command line option 560, 566
- {-}file <filename>
 - command line option 560
- {-}filter <JSON>
 - command line option 547, 550
- {-}forceTableScan
 - command line option 542, 567
- {-}fork
 - command line option 507, 521
- {-}forward <host><:port>
 - command line option 581
- {-}from <host[:port]>
 - command line option 555
- {-}headerline
 - command line option 560
- {-}help, -h
 - command line option 503, 519, 528, 537, 543, 549, 551, 556, 562, 570, 576, 581, 583, 586
- {-}host <hostname>

- command line option 528
- {-}host <hostname><:port>
 - command line option 586
- {-}host <hostname><:port>, -h
 - command line option 538, 544, 551, 557, 563, 570, 576
- {-}http
 - command line option 573
- {-}httpinterface
 - command line option 506, 521
- {-}ignoreBlanks
 - command line option 560
- {-}install
 - command line option 534, 535
- {-}ipv6
 - command line option 507, 527, 528, 538, 544, 552, 557, 563, 570, 577, 586
- {-}journal
 - command line option 511, 541, 547, 554, 559, 565, 589
- {-}journalCommitInterval <value>
 - command line option 512
- {-}journalOptions <arguments>
 - command line option 512
- {-}jsonArray
 - command line option 561, 567
- {-}jsonp
 - command line option 507, 527
- {-}keepIndexVersion
 - command line option 548
- {-}keyFile <file>
 - command line option 506, 521
- {-}limit <number>
 - command line option 567
- {-}local <filename>, -l
 - command line option 589
- {-}localThreshold
 - command line option 522
- {-}locks
 - command line option 579
- {-}logappend
 - command line option 505, 520
- {-}logpath <path>
 - command line option 505, 520
- {-}master
 - command line option 513
- {-}maxConns <number>
 - command line option 504, 519
- {-}moveParanoia
 - command line option 514
- {-}noAutoSplit
 - command line option 523
- {-}noIndexBuildRetry
 - command line option 509
- {-}noIndexRestore
 - command line option 548
- {-}noOptionsRestore
 - command line option 548
- {-}noauth
 - command line option 507
- {-}nodb
 - command line option 527
- {-}noheaders
 - command line option 572
- {-}nohttpinterface
 - command line option 506
- {-}nojournal
 - command line option 511
- {-}noobjcheck
 - command line option 511, 547, 550
- {-}noprealloc
 - command line option 509
- {-}norc
 - command line option 528
- {-}noscripting
 - command line option 511, 527
- {-}notablesan
 - command line option 511
- {-}nounixsocket
 - command line option 506, 521
- {-}nssize <value>
 - command line option 509
- {-}objcheck
 - command line option 511, 547, 550, 581
- {-}only <arg>
 - command line option 513
- {-}oplog
 - command line option 541
- {-}oplogLimit <timestamp>
 - command line option 548
- {-}oplogReplay
 - command line option 548
- {-}oplogSize <value>
 - command line option 512
- {-}oplogns <namespace>
 - command line option 555
- {-}out <file>, -o
 - command line option 567
- {-}out <path>, -o
 - command line option 541
- {-}password <password>, -p
 - command line option 528, 540, 546, 553, 558, 565, 572, 578, 588
- {-}pidfilepath <path>
 - command line option 506, 520
- {-}port
 - command line option 552
- {-}port <port>

- command line option 504, 519, 528, 538, 544, 557, 563, 570, 577, 586
- {-}profile <level>
 - command line option 508
- {-}query <JSON>, -q
 - command line option 566
- {-}query <json>, -q
 - command line option 541
- {-}quiet
 - command line option 504, 519, 528, 538, 544, 549, 551, 556, 562, 576, 586
- {-}quota
 - command line option 509
- {-}quotaFiles <number>
 - command line option 509
- {-}reinstall
 - command line option 534, 536
- {-}remove
 - command line option 534, 535
- {-}repair
 - command line option 510, 542
- {-}repairpath <path>
 - command line option 511
- {-}replIndexPrefetch
 - command line option 512
- {-}replSet <setname>
 - command line option 512
- {-}replace, -r
 - command line option 589
- {-}rest
 - command line option 507
- {-}rowcount <number>, -n
 - command line option 572
- {-}seconds <number>, -s
 - command line option 555
- {-}serviceDescription <description>
 - command line option 534, 536
- {-}serviceDisplayName <name>
 - command line option 534, 536
- {-}serviceName name
 - command line option 534, 536
- {-}servicePassword <password>
 - command line option 535, 536
- {-}serviceUser <user>
 - command line option 535, 536
- {-}setParameter <options>
 - command line option 506, 521
- {-}shardsvr
 - command line option 514
- {-}shell
 - command line option 527
- {-}shutdown
 - command line option 512
- {-}skip <number>
 - command line option 567
- {-}slave
 - command line option 513
- {-}slaveOk, -k
 - command line option 567
- {-}slavedelay <value>
 - command line option 513
- {-}slowms <value>
 - command line option 508
- {-}smallfiles
 - command line option 510
- {-}snmp-master
 - command line option 518
- {-}snmp-subagent
 - command line option 518
- {-}sort <JSON>
 - command line option 567
- {-}source <NET [interface]>, <FILE [filename]>, <DIA-
GLOG [filename]>
 - command line option 581
- {-}source <host><:port>
 - command line option 513
- {-}ssl
 - command line option 530, 538, 544, 552, 557, 563, 570, 577, 586
- {-}sslAllowInvalidCertificates
 - command line option 516, 525, 531, 539, 545, 553, 558, 564, 571, 578, 587
- {-}sslCAFile <filename>
 - command line option 516, 525, 530, 539, 545, 552, 557, 563, 570, 577, 587
- {-}sslCRLFile <filename>
 - command line option 516, 525, 530, 539, 545, 553, 558, 564, 571, 578, 587
- {-}sslClusterFile <filename>
 - command line option 515, 524
- {-}sslClusterPassword <value>
 - command line option 516, 524
- {-}sslFIPSMODE
 - command line option 517, 525, 530, 539, 545, 553, 558, 564, 571, 578, 587
- {-}sslMode <mode>
 - command line option 514, 523
- {-}sslOnNormalPorts
 - command line option 514, 523
- {-}sslPEMKeyFile <filename>
 - command line option 515, 523, 530, 539, 545, 552, 557, 564, 571, 577, 587
- {-}sslPEMKeyPassword <value>
 - command line option 515, 524, 530, 539, 545, 552, 558, 564, 571, 577, 587
- {-}sslWeakCertificateValidation
 - command line option 516, 525
- {-}stopOnError

- command line option 561
- {-}syncdelay <value>
 - command line option 510
- {-}sysinfo
 - command line option 508
- {-}syslog
 - command line option 504, 520
- {-}syslogFacility <string>
 - command line option 504, 520
- {-}timeStampFormat <string>
 - command line option 505, 520
- {-}traceExceptions
 - command line option 506
- {-}type <=json|=debug>
 - command line option 550
- {-}type <MIME>
 - command line option 589
- {-}type <json|csv|tsv>
 - command line option 560
- {-}unixSocketPrefix <path>
 - command line option 507, 521
- {-}upgrade
 - command line option 510, 522
- {-}upsert
 - command line option 560
- {-}upsertFields <field1[,field2]>
 - command line option 560
- {-}username <username>, -u
 - command line option 528, 540, 546, 553, 558, 564, 572, 578, 588
- {-}verbose
 - command line option 528
- {-}verbose, -v
 - command line option 504, 519, 538, 544, 549, 551, 556, 562, 570, 576, 586
- {-}version
 - command line option 503, 519, 528, 538, 544, 550, 551, 557, 563, 570, 576, 586
- {-}w <number of replicas per write>
 - command line option 548
- 503–531, 534–567, 570–573, 576–579, 581, 583, 586–589
- \$ (projection operator), 406
- \$ (update operator), 420
- \$add (aggregation framework transformation expression), 464
- \$addToSet (aggregation framework group expression), 455
- \$addToSet (update operator), 422
- \$all (query), 402
- \$allElementsTrue (aggregation framework transformation expression), 462
- \$and (aggregation framework transformation expression), 460
- \$and (query), 378
- \$anyElementTrue (aggregation framework transformation expression), 462
- \$atomic (update operator), 437
- \$avg (aggregation framework group expression), 457
- \$bit (update operator), 435
- \$box (query), 399
- \$center (query), 397
- \$centerSphere (query), 398
- \$cmd, 609
- \$cmp (aggregation framework transformation expression), 463
- \$comment (operator), 478
- \$concat (aggregation framework transformation expression), 465
- \$cond (aggregation framework transformation expression), 475
- \$currentDate (update operator), 419
- \$dayOfMonth (aggregation framework transformation expression), 474
- \$dayOfWeek (aggregation framework transformation expression), 474
- \$dayOfYear (aggregation framework transformation expression), 474
- \$divide (aggregation framework transformation expression), 464
- \$each (update operator), 427
- \$elemMatch (projection operator), 408
- \$elemMatch (query), 405
- \$eq (aggregation framework transformation expression), 463
- \$exists (query), 381
- \$explain (operator), 478
- \$first (aggregation framework group expression), 456
- \$geoIntersects (query), 393
- \$geoNear (aggregation framework pipeline operator), 451
- \$geoWithin (query), 392
- \$geometry (query), 397
- \$group (aggregation framework pipeline operator), 447
- \$gt (aggregation framework transformation expression), 463
- \$gt (query), 373
- \$gte (aggregation framework transformation expression), 463
- \$gte (query), 374
- \$hint (operator), 479
- \$hour (aggregation framework transformation expression), 474
- \$ifNull (aggregation framework transformation expression), 476
- \$in (query), 374
- \$inc (update operator), 412
- \$isolated (update operator), 436
- \$last (aggregation framework group expression), 456

- \$let (aggregation framework transformation expression), 471
- \$limit (aggregation framework pipeline operator), 445
- \$literal (aggregation framework transformation expression), 472
- \$lt (aggregation framework transformation expression), 463
- \$lt (query), 375
- \$lte (aggregation framework transformation expression), 463
- \$lte (query), 375
- \$map (aggregation framework transformation expression), 471
- \$match (aggregation framework pipeline operator), 440
- \$max (aggregation framework group expression), 456
- \$max (operator), 480
- \$max (update operator), 418
- \$maxDistance (query), 397
- \$maxScan (operator), 479
- \$maxTimeMS (operator), 480
- \$meta (aggregation framework transformation expression), 468
- \$meta (projection operator), 410
- \$millisecond (aggregation framework transformation expression), 475
- \$min (aggregation framework group expression), 456
- \$min (operator), 481
- \$min (update operator), 417
- \$minute (aggregation framework transformation expression), 474
- \$mod (aggregation framework transformation expression), 464
- \$mod (query), 384
- \$month (aggregation framework transformation expression), 474
- \$mul (update operator), 413
- \$multiply (aggregation framework transformation expression), 464
- \$natural (operator), 483
- \$ne (aggregation framework transformation expression), 464
- \$ne (query), 376
- \$near (query), 394
- \$nearSphere (query), 396
- \$nin (query), 376
- \$nor (query), 380
- \$not (aggregation framework transformation expression), 461
- \$not (query), 379
- \$options (operator), 386
- \$or (aggregation framework transformation expression), 461
- \$or (query), 377
- \$orderby (query), 481
- \$out (aggregation framework pipeline operator), 453
- \$polygon (query), 399
- \$pop (update operator), 423
- \$position (update operator), 433
- \$project (aggregation framework pipeline operator), 438
- \$pull (update operator), 424
- \$pullAll (update operator), 424
- \$push (aggregation framework group expression), 458
- \$push (update operator), 425
- \$pushAll (update operator), 425
- \$query (operator), 483
- \$redact (aggregation framework pipeline operator), 441
- \$regex (query), 386
- \$rename (update operator), 414
- \$returnKey (operator), 482
- \$second (aggregation framework transformation expression), 475
- \$set (update operator), 416
- \$setDifference (aggregation framework transformation expression), 462
- \$setEquals (aggregation framework transformation expression), 461
- \$setIntersection (aggregation framework transformation expression), 461
- \$setIsSubset (aggregation framework transformation expression), 462
- \$setOnInsert (update operator), 415
- \$setUnion (aggregation framework transformation expression), 462
- \$showDiskLoc (operator), 482
- \$size (aggregation framework transformation expression), 470
- \$size (query), 405
- \$skip (aggregation framework pipeline operator), 445
- \$slice (projection operator), 411
- \$slice (update operator), 428
- \$snapshot (operator), 482
- \$sort (aggregation framework pipeline operator), 449
- \$sort (update operator), 431
- \$strcasecmp (aggregation framework transformation expression), 467
- \$substr (aggregation framework transformation expression), 468
- \$subtract (aggregation framework transformation expression), 464
- \$sum (aggregation framework group expression), 460
- \$text (query), 387
- \$toLower (aggregation framework transformation expression), 468
- \$toUpper (aggregation framework transformation expression), 468
- \$type (query), 383
- \$uniqueDocs (query), 400
- \$unset (update operator), 417

\$unwind (aggregation framework pipeline operator), 446
\$week (aggregation framework transformation expression), 474
\$where (query), 391
\$within (query), 393
\$year (aggregation framework transformation expression), 474
_hashBSONElement (database command), 369
_hashBSONElement.key (MongoDB reporting output), 369
_hashBSONElement.ok (MongoDB reporting output), 370
_hashBSONElement.out (MongoDB reporting output), 369
_hashBSONElement.seed (MongoDB reporting output), 369
_id, 609
_isSelf (database command), 364
_isWindows (shell method), 196
_migrateClone (database command), 366
_rand (shell method), 197
_recvChunkAbort (database command), 365
_recvChunkCommit (database command), 365
_recvChunkStart (database command), 365
_recvChunkStatus (database command), 365
_skewClockCommand (database command), 371
_srand (shell method), 198
_startMongoProgram (shell method), 185
_testDistLockWithSkew (database command), 367
_testDistLockWithSyncCluster (database command), 368
_transferMods (database command), 365
<database>.system.indexes (MongoDB reporting output), 601
<database>.system.js (MongoDB reporting output), 601
<database>.system.namespaces (MongoDB reporting output), 601
<database>.system.profile (MongoDB reporting output), 601
0 (error code), 603
100 (error code), 604
12 (error code), 603
14 (error code), 603
2 (error code), 603
20 (error code), 603
2d Geospatial queries cannot use the \$or operator (MongoDB system limit), 608
3 (error code), 603
4 (error code), 603
45 (error code), 603
47 (error code), 603
48 (error code), 604
49 (error code), 604
5 (error code), 603

A

accumulator, 609
action, 609
addShard (database command), 284
admin database, 610
admin.system.roles (MongoDB reporting output), 601
admin.system.users (MongoDB reporting output), 601
admin.system.version (MongoDB reporting output), 601
aggregate (database command), 198
aggregation, 610
aggregation framework, 610
Aggregation Pipeline Operation (MongoDB system limit), 608
applyOps (database command), 279
arbiter, 610
authenticate (database command), 249
authentication, 610
authorization, 610
authSchemaUpgrade (database command), 250
availableQueryOptions (database command), 324

B

B-tree, 610
balancer, 610
batchType (MongoDB reporting output), 139
BSON, 610
BSON Document Size (MongoDB system limit), 604
BSON types, 610
bsondump (program), 549
buildInfo (database command), 324
buildInfo (MongoDB reporting output), 325
buildInfo allocator (MongoDB reporting output), 325
buildInfo.bits (MongoDB reporting output), 325
buildInfo.compilerFlags (MongoDB reporting output), 325
buildInfo.debug (MongoDB reporting output), 325
buildInfo.gitVersion (MongoDB reporting output), 325
buildInfo.javascriptEngine (MongoDB reporting output), 325
buildInfo.loaderFlags (MongoDB reporting output), 325
buildInfo.maxBsonObjectSize (MongoDB reporting output), 325
buildInfo.sysInfo (MongoDB reporting output), 325
buildInfo.versionArray (MongoDB reporting output), 325
Bulk (shell method), 128
Bulk Operation Size (MongoDB system limit), 608
Bulk.execute (shell method), 137
Bulk.find (shell method), 130
Bulk.find.remove (shell method), 131
Bulk.find.removeOne (shell method), 130
Bulk.find.replaceOne (shell method), 131
Bulk.find.update (shell method), 134
Bulk.find.updateOne (shell method), 132
Bulk.find.upsert (shell method), 134

- Bulk.getOperations (shell method), 138
 - Bulk.insert (shell method), 129
 - Bulk.toJson (shell method), 140
 - Bulk.toString (shell method), 140
 - BulkWriteResult (shell method), 189
 - BulkWriteResult.nInserted (MongoDB reporting output), 189
 - BulkWriteResult.nMatched (MongoDB reporting output), 189
 - BulkWriteResult.nModified (MongoDB reporting output), 189
 - BulkWriteResult.nRemoved (MongoDB reporting output), 189
 - BulkWriteResult.nUpserted (MongoDB reporting output), 189
 - BulkWriteResult.upserted (MongoDB reporting output), 190
 - BulkWriteResult.upserted._id (MongoDB reporting output), 190
 - BulkWriteResult.upserted.index (MongoDB reporting output), 190
 - BulkWriteResult.writeConcernError (MongoDB reporting output), 190
 - BulkWriteResult.writeConcernError.code (MongoDB reporting output), 190
 - BulkWriteResult.writeConcernError.errInfo (MongoDB reporting output), 190
 - BulkWriteResult.writeConcernError.errmsg (MongoDB reporting output), 190
 - BulkWriteResult.writeErrors (MongoDB reporting output), 190
 - BulkWriteResult.writeErrors.code (MongoDB reporting output), 190
 - BulkWriteResult.writeErrors.errmsg (MongoDB reporting output), 190
 - BulkWriteResult.writeErrors.index (MongoDB reporting output), 190
 - BulkWriteResult.writeErrors.op (MongoDB reporting output), 190
- ## C
- CAP Theorem, 610
 - capped collection, 610
 - captrunc (database command), 368
 - cat (shell method), 194
 - cd (shell method), 195
 - checkShardingIndex (database command), 288
 - checksum, 610
 - chunk, 610
 - clean (database command), 313
 - cleanupOrphaned (database command), 285
 - cleanupOrphaned.ok (MongoDB reporting output), 286
 - cleanupOrphaned.stoppedAtKey (MongoDB reporting output), 286
 - clearRawMongoProgramOutput (shell method), 184
 - client, 610
 - clone (database command), 305
 - cloneCollection (database command), 306
 - cloneCollectionAsCapped (database command), 306
 - closeAllDatabases (database command), 307
 - cluster, 610
 - collection, 610
 - system, 600
 - collMod (database command), 316
 - collStats (database command), 325
 - collStats.avgObjSize (MongoDB reporting output), 326
 - collStats.count (MongoDB reporting output), 326
 - collStats.flags (MongoDB reporting output), 327
 - collStats.indexSizes (MongoDB reporting output), 327
 - collStats.lastExtentSize (MongoDB reporting output), 326
 - collStats.nindexes (MongoDB reporting output), 326
 - collStats.ns (MongoDB reporting output), 326
 - collStats.numExtents (MongoDB reporting output), 326
 - collStats.paddingFactor (MongoDB reporting output), 326
 - collStats.size (MongoDB reporting output), 326
 - collStats.storageSize (MongoDB reporting output), 326
 - collStats.systemFlags (MongoDB reporting output), 327
 - collStats.totalIndexSize (MongoDB reporting output), 327
 - collStats.userFlags (MongoDB reporting output), 327
 - compact (database command), 313
 - compound index, 610
 - config (MongoDB reporting output), 593
 - config database, 610
 - config server, 611
 - config.changelog (MongoDB reporting output), 593
 - config.changelog._id (MongoDB reporting output), 594
 - config.changelog.clientAddr (MongoDB reporting output), 594
 - config.changelog.details (MongoDB reporting output), 595
 - config.changelog.ns (MongoDB reporting output), 595
 - config.changelog.server (MongoDB reporting output), 594
 - config.changelog.time (MongoDB reporting output), 594
 - config.changelog.what (MongoDB reporting output), 594
 - config.chunks (MongoDB reporting output), 595
 - config.collections (MongoDB reporting output), 595
 - config.databases (MongoDB reporting output), 596
 - config.lockpings (MongoDB reporting output), 596
 - config.locks (MongoDB reporting output), 596
 - config.mongos (MongoDB reporting output), 597
 - config.settings (MongoDB reporting output), 597
 - config.shards (MongoDB reporting output), 597
 - config.tags (MongoDB reporting output), 598
 - config.version (MongoDB reporting output), 598

- configureFailPoint (database command), 372
- connect (shell method), 193
- connPoolStats (database command), 328
- connPoolStats.createdByType (MongoDB reporting output), 329
- connPoolStats.createdByType.master (MongoDB reporting output), 329
- connPoolStats.createdByType.set (MongoDB reporting output), 329
- connPoolStats.createdByType.sync (MongoDB reporting output), 329
- connPoolStats.hosts (MongoDB reporting output), 328
- connPoolStats.hosts.[host].available (MongoDB reporting output), 328
- connPoolStats.hosts.[host].created (MongoDB reporting output), 328
- connPoolStats.numAScopedConnection (MongoDB reporting output), 329
- connPoolStats.numDBClientConnection (MongoDB reporting output), 329
- connPoolStats.replicaSets (MongoDB reporting output), 328
- connPoolStats.replicaSets.shard (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host[n].addr (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host[n].hidden (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host[n].ismaster (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host[n].ok (MongoDB reporting output), 328
- connPoolStats.replicaSets.[shard].host[n].pingTimeMillis (MongoDB reporting output), 329
- connPoolStats.replicaSets.[shard].host[n].secondary (MongoDB reporting output), 329
- connPoolStats.replicaSets.[shard].host[n].tags (MongoDB reporting output), 329
- connPoolStats.replicaSets.[shard].master (MongoDB reporting output), 329
- connPoolStats.replicaSets.[shard].nextSlave (MongoDB reporting output), 329
- connPoolStats.totalAvailable (MongoDB reporting output), 329
- connPoolStats.totalCreated (MongoDB reporting output), 329
- connPoolSync (database command), 313
- control script, **611**
- convertToCapped (database command), 307
- copydb (database command), 300
- copydbgetnonce (database command), 250
- copyDbpath (shell method), 195
- count (database command), 201
- Covered Queries in Sharded Clusters (MongoDB system limit), 607
- create (database command), 304
- createIndexes (database command), 308
- createIndexes.code (MongoDB reporting output), 311
- createIndexes.createdCollectionAutomatically (MongoDB reporting output), 311
- createIndexes.errmsg (MongoDB reporting output), 311
- createIndexes.note (MongoDB reporting output), 311
- createIndexes.numIndexesAfter (MongoDB reporting output), 311
- createIndexes.numIndexesBefore (MongoDB reporting output), 311
- createIndexes.ok (MongoDB reporting output), 311
- createRole (database command), 259
- createUser (database command), 250
- CRUD, **611**
- CSV, **611**
- CURRENT (system variable available in aggregation), 493
- currentOp.active (MongoDB reporting output), 105
- currentOp.client (MongoDB reporting output), 106
- currentOp.connectionId (MongoDB reporting output), 106
- currentOp.desc (MongoDB reporting output), 106
- currentOp.killed (MongoDB reporting output), 106
- currentOp.locks (MongoDB reporting output), 106
- currentOp.locks.^ (MongoDB reporting output), 106
- currentOp.locks.^<database> (MongoDB reporting output), 106
- currentOp.locks.^local (MongoDB reporting output), 106
- currentOp.lockStats (MongoDB reporting output), 107
- currentOp.msg (MongoDB reporting output), 106
- currentOp.ns (MongoDB reporting output), 105
- currentOp.numYields (MongoDB reporting output), 107
- currentOp.op (MongoDB reporting output), 105
- currentOp.opid (MongoDB reporting output), 105
- currentOp.progress (MongoDB reporting output), 106
- currentOp.progress.done (MongoDB reporting output), 106
- currentOp.progress.total (MongoDB reporting output), 106
- currentOp.query (MongoDB reporting output), 106
- currentOp.secs_running (MongoDB reporting output), 105
- currentOp.threadId (MongoDB reporting output), 106
- currentOp.timeAcquiringMicros (MongoDB reporting output), 107
- currentOp.timeAcquiringMicros.R (MongoDB reporting output), 107
- currentOp.timeAcquiringMicros.r (MongoDB reporting output), 107

- currentOp.timeAcquiringMicros.W (MongoDB reporting output), 107
- currentOp.timeAcquiringMicros.w (MongoDB reporting output), 107
- currentOp.timeLockedMicros (MongoDB reporting output), 107
- currentOp.timeLockedMicros.R (MongoDB reporting output), 107
- currentOp.timeLockedMicros.r (MongoDB reporting output), 107
- currentOp.timeLockedMicros.W (MongoDB reporting output), 107
- currentOp.timeLockedMicros.w (MongoDB reporting output), 107
- currentOp.waitForLock (MongoDB reporting output), 106
- cursor, **611**
- cursor.addOption (shell method), 77
- cursor.batchSize (shell method), 78
- cursor.count (shell method), 79
- cursor.explain (shell method), 80
- cursor.forEach (shell method), 85
- cursor.hasNext (shell method), 85
- cursor.hint (shell method), 86
- cursor.limit (shell method), 86
- cursor.map (shell method), 87
- cursor.max (shell method), 88
- cursor.maxTimeMS (shell method), 87
- cursor.min (shell method), 89
- cursor.next (shell method), 91
- cursor.objsLeftInBatch (shell method), 91
- cursor.readPref (shell method), 91
- cursor.showDiskLoc (shell method), 91
- cursor.size (shell method), 92
- cursor.skip (shell method), 92
- cursor.snapshot (shell method), 92
- cursor.sort (shell method), 93
- cursor.toArray (shell method), 96
- cursorInfo (database command), 333
- D**
- daemon, **611**
- data directory, **611**
- Data Size (MongoDB system limit), 606
- data-center awareness, **611**
- database, **611**
- local, 598
- database command, **611**
- Database Name Case Sensitivity (MongoDB system limit), 609
- database profiler, **611**
- dataSize (database command), 333
- Date (shell method), 186
- datum, **611**
- db.addUser (shell method), 97, 143
- db.auth (shell method), 99
- db.changeUserPassword (shell method), 99, 147
- db.cloneCollection (shell method), 99
- db.cloneDatabase (shell method), 100
- db.collection.aggregate (shell method), 22
- db.collection.copyTo (shell method), 26
- db.collection.count (shell method), 25
- db.collection.createIndex (shell method), 27
- db.collection.dataSize (shell method), 28
- db.collection.distinct (shell method), 28
- db.collection.drop (shell method), 29
- db.collection.dropIndex (shell method), 29
- db.collection.dropIndexes (shell method), 30
- db.collection.ensureIndex (shell method), 30
- db.collection.find (shell method), 34
- db.collection.findAndModify (shell method), 39
- db.collection.findOne (shell method), 43
- db.collection.getIndexes (shell method), 45
- db.collection.getIndexStats (shell method), 27
- db.collection.getPlanCache (shell method), 123
- db.collection.getShardDistribution (shell method), 45
- db.collection.getShardVersion (shell method), 47
- db.collection.group (shell method), 47
- db.collection.indexStats (shell method), 28
- db.collection.initializeOrderedBulkOp (shell method), 50
- db.collection.initializeUnorderedBulkOp (shell method), 51
- db.collection.insert (shell method), 52
- db.collection.isCapped (shell method), 55
- db.collection.mapReduce (shell method), 55
- db.collection.reIndex (shell method), 62
- db.collection.remove (shell method), 62
- db.collection.renameCollection (shell method), 65
- db.collection.save (shell method), 66
- db.collection.stats (shell method), 68
- db.collection.storageSize (shell method), 68
- db.collection.totalIndexSize (shell method), 69
- db.collection.totalSize (shell method), 68
- db.collection.update (shell method), 69
- db.collection.validate (shell method), 76
- db.commandHelp (shell method), 100
- db.copyDatabase (shell method), 100
- db.createCollection (shell method), 102
- db.createRole (shell method), 152
- db.createUser (shell method), 141
- db.currentOp (shell method), 103
- db.dropAllRoles (shell method), 156
- db.dropAllUsers (shell method), 147
- db.dropDatabase (shell method), 108
- db.dropRole (shell method), 155
- db.dropUser (shell method), 148
- db.eval (shell method), 108
- db.fsyncLock (shell method), 109

- db.fsyncUnlock (shell method), 110
- db.getCollection (shell method), 110
- db.getCollectionNames (shell method), 110
- db.getLastError (shell method), 110
- db.getLastErrorObj (shell method), 111
- db.getMongo (shell method), 111
- db.getName (shell method), 111
- db.getPrevError (shell method), 111
- db.getProfilingLevel (shell method), 111
- db.getProfilingStatus (shell method), 112
- db.getReplicationInfo (shell method), 112
- db.getReplicationInfo.errmsg (MongoDB reporting output), 112
- db.getReplicationInfo.logSizeMB (MongoDB reporting output), 112
- db.getReplicationInfo.now (MongoDB reporting output), 112
- db.getReplicationInfo.oplogMainRowCount (MongoDB reporting output), 112
- db.getReplicationInfo.tFirst (MongoDB reporting output), 112
- db.getReplicationInfo.timeDiff (MongoDB reporting output), 112
- db.getReplicationInfo.timeDiffHours (MongoDB reporting output), 112
- db.getReplicationInfo.tLast (MongoDB reporting output), 112
- db.getReplicationInfo.usedMB (MongoDB reporting output), 112
- db.getRole (shell method), 162
- db.getRoles (shell method), 163
- db.getSiblingDB (shell method), 113
- db.getUser (shell method), 151
- db.getUsers (shell method), 151
- db.grantPrivilegesToRole (shell method), 156
- db.grantRolesToRole (shell method), 160
- db.grantRolesToUser (shell method), 148
- db.help (shell method), 113
- db.hostInfo (shell method), 113
- db.isMaster (shell method), 114
- db.killOp (shell method), 114
- db.listCommands (shell method), 115
- db.loadServerScripts (shell method), 115
- db.logout (shell method), 115
- db.printCollectionStats (shell method), 115
- db.printReplicationInfo (shell method), 116
- db.printShardingStatus (shell method), 116
- db.printSlaveReplicationInfo (shell method), 116
- db.removeUser (shell method), 117, 147
- db.repairDatabase (shell method), 117
- db.resetError (shell method), 117
- db.revokePrivilegesFromRole (shell method), 158
- db.revokeRolesFromRole (shell method), 161
- db.revokeRolesFromUser (shell method), 150
- db.runCommand (shell method), 118
- db.serverBuildInfo (shell method), 118
- db.serverStatus (shell method), 118
- db.setProfilingLevel (shell method), 119
- db.shutdownServer (shell method), 119
- db.stats (shell method), 119
- db.updateRole (shell method), 153
- db.updateUser (shell method), 145
- db.upgradeCheck (shell method), 120
- db.upgradeCheckAllDBs (shell method), 121
- db.version (shell method), 120
- dbHash (database command), 324
- dbpath, **611**
- DBQuery.Option.awaitData (MongoDB reporting output), 78
- DBQuery.Option.exhaust (MongoDB reporting output), 78
- DBQuery.Option.noTimeout (MongoDB reporting output), 78
- DBQuery.Option.oplogReplay (MongoDB reporting output), 78
- DBQuery.Option.partial (MongoDB reporting output), 78
- DBQuery.Option.slaveOk (MongoDB reporting output), 78
- DBQuery.Option.tailable (MongoDB reporting output), 78
- dbStats (database command), 331
- dbStats.avgObjSize (MongoDB reporting output), 332
- dbStats.collections (MongoDB reporting output), 332
- dbStats.dataFileVersion (MongoDB reporting output), 332
- dbStats.dataFileVersion.major (MongoDB reporting output), 332
- dbStats.dataFileVersion.minor (MongoDB reporting output), 332
- dbStats.dataSize (MongoDB reporting output), 332
- dbStats.db (MongoDB reporting output), 332
- dbStats.fileSize (MongoDB reporting output), 332
- dbStats.indexes (MongoDB reporting output), 332
- dbStats.indexSize (MongoDB reporting output), 332
- dbStats.nsSizeMB (MongoDB reporting output), 332
- dbStats.numExtents (MongoDB reporting output), 332
- dbStats.objects (MongoDB reporting output), 332
- dbStats.storageSize (MongoDB reporting output), 332
- delayed member, **611**
- delete (database command), 226
- delete.n (MongoDB reporting output), 228
- delete.ok (MongoDB reporting output), 228
- delete.writeConcernError (MongoDB reporting output), 229
- delete.writeConcernError.code (MongoDB reporting output), 229
- delete.writeConcernError.errmsg (MongoDB reporting output), 229

delete.writeErrors (MongoDB reporting output), 228
 delete.writeErrors.code (MongoDB reporting output), 229
 delete.writeErrors.errmsg (MongoDB reporting output), 229
 delete.writeErrors.index (MongoDB reporting output), 229
 DESCEND (system variable available in aggregation), 493
 diagLogging (database command), 333
 diagnostic log, 611
 distinct (database command), 203
 document, 611
 space allocation, 316
 dot notation, 611
 draining, 611
 driver, 611
 driverOIDTest (database command), 324
 drop (database command), 304
 dropAllRolesFromDatabase (database command), 263
 dropAllUsersFromDatabase (database command), 254
 dropDatabase (database command), 304
 dropIndexes (database command), 311
 dropRole (database command), 262
 dropUser (database command), 253

E

EDITOR, 532
 election, 611
 emptycapped (database command), 368
 enableSharding (database command), 288
 environment variable
 EDITOR, 532
 HOME, 531, 532
 HOMEDRIVE, 532
 HOMEPATH, 532
 eval (database command), 238
 eventual consistency, 612
 expireAfterSeconds, 317
 explain.allPlans (MongoDB reporting output), 84
 explain.clauses (MongoDB reporting output), 84
 explain.clusteredType (MongoDB reporting output), 85
 explain.cursor (MongoDB reporting output), 83
 explain.filterSet (MongoDB reporting output), 84
 explain.indexBounds (MongoDB reporting output), 84
 explain.indexOnly (MongoDB reporting output), 84
 explain.isMultiKey (MongoDB reporting output), 83
 explain.millis (MongoDB reporting output), 84
 explain.millisShardAvg (MongoDB reporting output), 85
 explain.millisShardTotal (MongoDB reporting output), 85
 explain.n (MongoDB reporting output), 83
 explain.nChunkSkips (MongoDB reporting output), 84
 explain.nscanned (MongoDB reporting output), 83

explain.nscannedAllPlans (MongoDB reporting output), 83
 explain.nscannedObjects (MongoDB reporting output), 83
 explain.nscannedObjectsAllPlans (MongoDB reporting output), 83
 explain.numQueries (MongoDB reporting output), 85
 explain.numShards (MongoDB reporting output), 85
 explain.nYields (MongoDB reporting output), 84
 explain.oldPlan (MongoDB reporting output), 84
 explain.scanAndOrder (MongoDB reporting output), 83
 explain.server (MongoDB reporting output), 84
 explain.shards (MongoDB reporting output), 85
 expression, 612

F

failover, 612
 features (database command), 364
 field, 612
 filemd5 (database command), 308
 findAndModify (database command), 229
 firewall, 612
 flushRouterConfig (database command), 283
 forceerror (database command), 371
 fsync, 612
 fsync (database command), 312
 fuzzFile (shell method), 195

G

geohash, 612
 GeoJSON, 612
 geoNear (database command), 216
 geoNear.ok (MongoDB reporting output), 219
 geoNear.results (MongoDB reporting output), 218
 geoNear.results[n].dis (MongoDB reporting output), 218
 geoNear.results[n].obj (MongoDB reporting output), 218
 geoNear.stats (MongoDB reporting output), 218
 geoNear.stats.avgDistance (MongoDB reporting output), 219
 geoNear.stats.maxDistance (MongoDB reporting output), 219
 geoNear.stats.nscanned (MongoDB reporting output), 218
 geoNear.stats.objectsLoaded (MongoDB reporting output), 219
 geoNear.stats.time (MongoDB reporting output), 219
 geoSearch (database command), 219
 geospatial, 612
 geoWalk (database command), 219
 getCmdLineOpts (database command), 333
 getHostName (shell method), 195
 getLastError (database command), 235
 getLastError.code (MongoDB reporting output), 236

getLastError.connectionId (MongoDB reporting output), 236

getLastError.err (MongoDB reporting output), 236

getLastError.lastOp (MongoDB reporting output), 236

getLastError.n (MongoDB reporting output), 236

getLastError.ok (MongoDB reporting output), 236

getLastError.shards (MongoDB reporting output), 236

getLastError.singleShard (MongoDB reporting output), 236

getLastError.updatedExisting (MongoDB reporting output), 236

getLastError.upserted (MongoDB reporting output), 236

getLastError.waited (MongoDB reporting output), 237

getLastError.wnote (MongoDB reporting output), 236

getLastError.wtime (MongoDB reporting output), 237

getLastError.wtimeout (MongoDB reporting output), 237

getLog (database command), 344

getMemInfo (shell method), 195

getnonce (database command), 250

getoptime (database command), 282

getParameter (database command), 318

getPrevError (database command), 237

getShardMap (database command), 289

getShardVersion (database command), 289

godinsert (database command), 369

grantPrivilegesToRole (database command), 264

grantRolesToRole (database command), 267

grantRolesToUser (database command), 255

GridFS, 612

group (database command), 204

H

handshake (database command), 365

hashed shard key, 612

haystack index, 612

hidden member, 612

HOME, 531

HOMEDRIVE, 532

hostInfo (database command), 344

hostInfo (MongoDB reporting output), 345

hostInfo.extra (MongoDB reporting output), 346

hostInfo.extra.alwaysFullSync (MongoDB reporting output), 346

hostInfo.extra.cpuFeatures (MongoDB reporting output), 346

hostInfo.extra.cpuFrequencyMHz (MongoDB reporting output), 346

hostInfo.extra.kernelVersion (MongoDB reporting output), 346

hostInfo.extra.libcVersion (MongoDB reporting output), 346

hostInfo.extra.maxOpenFiles (MongoDB reporting output), 346

hostInfo.extra.nfsAsync (MongoDB reporting output), 346

hostInfo.extra.numPages (MongoDB reporting output), 346

hostInfo.extra.pageSize (MongoDB reporting output), 346

hostInfo.extra.scheduler (MongoDB reporting output), 346

hostInfo.extra.versionString (MongoDB reporting output), 346

hostInfo.os (MongoDB reporting output), 345

hostInfo.os.name (MongoDB reporting output), 346

hostInfo.os.type (MongoDB reporting output), 346

hostInfo.os.version (MongoDB reporting output), 346

hostInfo.system (MongoDB reporting output), 345

hostInfo.system.cpuAddrSize (MongoDB reporting output), 345

hostInfo.system.cpuArch (MongoDB reporting output), 345

hostInfo.system.currentTime (MongoDB reporting output), 345

hostInfo.system.hostname (MongoDB reporting output), 345

hostInfo.system.memSizeMB (MongoDB reporting output), 345

hostInfo.system.numaEnabled (MongoDB reporting output), 345

hostInfo.system.numCores (MongoDB reporting output), 345

hostname (shell method), 196

I

idempotent, 612

index, 612

index (collection flag), 317

Index Key Limit (MongoDB system limit), 604

Index Name Length (MongoDB system limit), 605

indexStats (database command), 339

indexStats.bucketBodyBytes (MongoDB reporting output), 339

indexStats.depth (MongoDB reporting output), 339

indexStats.index (MongoDB reporting output), 339

indexStats.isIdIndex (MongoDB reporting output), 339

indexStats.keyPattern (MongoDB reporting output), 339

indexStats.overall (MongoDB reporting output), 339

indexStats.overall.bsonRatio (MongoDB reporting output), 340

indexStats.overall.fillRatio (MongoDB reporting output), 340

indexStats.overall.keyCount (MongoDB reporting output), 340

indexStats.overall.keyNodeRatio (MongoDB reporting output), 340

- indexStats.overall.numBuckets (MongoDB reporting output), 340
 - indexStats.overall.usedKeyCount (MongoDB reporting output), 340
 - indexStats.perLevel (MongoDB reporting output), 340
 - indexStats.perLevel.bsonRatio (MongoDB reporting output), 340
 - indexStats.perLevel.fillRatio (MongoDB reporting output), 340
 - indexStats.perLevel.keyCount (MongoDB reporting output), 340
 - indexStats.perLevel.keyNodeRatio (MongoDB reporting output), 340
 - indexStats.perLevel.numBuckets (MongoDB reporting output), 340
 - indexStats.perLevel.usedKeyCount (MongoDB reporting output), 340
 - indexStats.storageNs (MongoDB reporting output), 339
 - indexStats.version (MongoDB reporting output), 339
 - initial sync, **612**
 - insert (database command), 220
 - insert.n (MongoDB reporting output), 221
 - insert.ok (MongoDB reporting output), 221
 - insert.writeConcernError (MongoDB reporting output), 221
 - insert.writeConcernError.code (MongoDB reporting output), 222
 - insert.writeConcernError.errmsg (MongoDB reporting output), 222
 - insert.writeErrors (MongoDB reporting output), 221
 - insert.writeErrors.code (MongoDB reporting output), 221
 - insert.writeErrors.errmsg (MongoDB reporting output), 221
 - insert.writeErrors.index (MongoDB reporting output), 221
 - internals
 - config database, 593
 - interrupt point, **612**
 - invalidateUserCache (database command), 272
 - IPv6, **612**
 - isdbgrid (database command), 298
 - isMaster (database command), 280
 - isMaster.arbiterOnly (MongoDB reporting output), 281
 - isMaster.arbiters (MongoDB reporting output), 281
 - isMaster.hidden (MongoDB reporting output), 281
 - isMaster.hosts (MongoDB reporting output), 281
 - isMaster.ismaster (MongoDB reporting output), 280
 - isMaster.localTime (MongoDB reporting output), 280
 - isMaster.maxBsonObjectSize (MongoDB reporting output), 280
 - isMaster.maxMessageSizeBytes (MongoDB reporting output), 280
 - isMaster.maxWireVersion (MongoDB reporting output), 281
 - isMaster.me (MongoDB reporting output), 282
 - isMaster.minWireVersion (MongoDB reporting output), 280
 - isMaster.msg (MongoDB reporting output), 281
 - isMaster.passive (MongoDB reporting output), 281
 - isMaster.passives (MongoDB reporting output), 281
 - isMaster.primary (MongoDB reporting output), 281
 - isMaster.secondary (MongoDB reporting output), 281
 - isMaster.setName (MongoDB reporting output), 281
 - isMaster.tags (MongoDB reporting output), 281
 - ISODate, **612**
- ## J
- JavaScript, **613**
 - journal, **613**
 - journalLatencyTest (database command), 370
 - JSON, **613**
 - JSON document, **613**
 - JSONP, **613**
- ## K
- KEEP (system variable available in aggregation), 493
- ## L
- least privilege, **613**
 - legacy coordinate pairs, **613**
 - Length of Database Names (MongoDB system limit), 609
 - LineString, **613**
 - listCommands (database command), 324
 - listDatabases (database command), 323
 - listFiles (shell method), 196
 - listShards (database command), 288
 - load (shell method), 196
 - local database, 598
 - local.oplog.\$main (MongoDB reporting output), 600
 - local.oplog.rs (MongoDB reporting output), 600
 - local.replset.minvalid (MongoDB reporting output), 600
 - local.slaves (MongoDB reporting output), 600
 - local.sources (MongoDB reporting output), 600
 - local.startup_log (MongoDB reporting output), 599
 - local.startup_log._id (MongoDB reporting output), 599
 - local.startup_log.buildinfo (MongoDB reporting output), 599
 - local.startup_log.cmdLine (MongoDB reporting output), 599
 - local.startup_log.hostname (MongoDB reporting output), 599
 - local.startup_log.pid (MongoDB reporting output), 599
 - local.startup_log.startTime (MongoDB reporting output), 599
 - local.startup_log.startTimeLocal (MongoDB reporting output), 599
 - local.system.replset (MongoDB reporting output), 600
 - lock, **613**

logApplicationMessage (database command), 372
logout (database command), 249
logRotate (database command), 322
ls (shell method), 197
LVM, 613

M

map-reduce, 613
mapping type, 613
mapReduce (database command), 208
mapreduce.shardedfinish (database command), 365
master, 613
Maximum Number of Documents in a Capped Collection (MongoDB system limit), 606
Maximum Size of Auto-Created Oplog (MongoDB system limit), 606
md5, 613
md5sumFile (shell method), 197
medianKey (database command), 296
mergeChunks (database command), 289
MIB, 613
MIME, 613
mkdir (shell method), 197
mongo, 614
mongo (program), 527
Mongo (shell method), 193
Mongo.getDB (shell method), 191
Mongo.getReadPrefMode (shell method), 191
Mongo.getReadPrefTagSet (shell method), 192
Mongo.setReadPref (shell method), 192
Mongo.setSlaveOk (shell method), 193
mongod, 614
mongod (program), 503
mongod.exe (program), 534
MongoDB, 614
MongoDB Enterprise, 614
mongodump (program), 537
mongoexport (program), 562
mongofiles (program), 585, 586
mongoimport (program), 556
mongooplog (program), 551
mongoperf (program), 583
mongoperf.fileSizeMB (setting), 583
mongoperf.mmf (setting), 584
mongoperf.nThreads (setting), 583
mongoperf.r (setting), 584
mongoperf.recSizeKB (setting), 584
mongoperf.sleepMicros (setting), 583
mongoperf.syncDelay (setting), 584
mongoperf.w (setting), 584
mongorestore (program), 543
mongos, 614
mongos (program), 518, 519
mongos.exe (program), 535

mongosniff (program), 581
mongostat (program), 570
mongotop (program), 576
mongotop.<timestamp> (MongoDB reporting output), 580
mongotop.db (MongoDB reporting output), 579
mongotop.ns (MongoDB reporting output), 579
mongotop.read (MongoDB reporting output), 579
mongotop.total (MongoDB reporting output), 579
mongotop.write (MongoDB reporting output), 580
Monotonically Increasing Shard Keys Can Limit Insert Throughput (MongoDB system limit), 608
moveChunk (database command), 296
movePrimary (database command), 297

N

namespace, 614
 local, 598
 system, 600
Namespace Length (MongoDB system limit), 604
natural order, 614
Nested Depth for BSON Documents (MongoDB system limit), 604
netstat (database command), 334
Number of Collections in a Database (MongoDB system limit), 606
Number of Indexed Fields in a Compound Index (MongoDB system limit), 605
Number of Indexes per Collection (MongoDB system limit), 605
Number of Members of a Replica Set (MongoDB system limit), 606
Number of Namespaces (MongoDB system limit), 604
Number of Voting Members of a Replica Set (MongoDB system limit), 606

O

ObjectId, 614
ObjectId.getTimestamp (shell method), 187
ObjectId.toString (shell method), 187
ObjectId.valueOf (shell method), 187
operations (MongoDB reporting output), 139
Operations Unavailable in Sharded Environments (MongoDB system limit), 606
operator, 614
oplog, 614
ordered query plan, 614
originalZeroIndex (MongoDB reporting output), 139
orphaned document, 614

P

padding, 614
padding factor, 614
page fault, 614

parallelCollectionScan (database command), 240
 partition, **615**
 passive member, **615**
 pcap, **615**
 PID, **615**
 ping (database command), 334
 pipe, **615**
 pipeline, **615**
 PlanCache.clear (shell method), 127
 PlanCache.clearPlansByQuery (shell method), 126
 PlanCache.getPlansByQuery (shell method), 125
 PlanCache.help (shell method), 124
 PlanCache.listQueryShapes (shell method), 124
 planCacheClear (database command), 247
 planCacheClearFilters (database command), 243
 planCacheListFilters (database command), 241
 planCacheListFilters.filters (MongoDB reporting output), 241
 planCacheListFilters.filters.indexes (MongoDB reporting output), 242
 planCacheListFilters.filters.projection (MongoDB reporting output), 242
 planCacheListFilters.filters.query (MongoDB reporting output), 241
 planCacheListFilters.filters.sort (MongoDB reporting output), 241
 planCacheListFilters.ok (MongoDB reporting output), 242
 planCacheListPlans (database command), 246
 planCacheListQueryShapes (database command), 245
 planCacheSetFilter (database command), 242
 Point, **615**
 Polygon, **615**
 powerOf2Sizes, **615**
 pre-splitting, **615**
 primary, **615**
 primary key, **615**
 primary shard, **615**
 priority, **615**
 privilege, **615**
 profile (database command), 334
 projection, **615**
 PRUNE (system variable available in aggregation), 493
 pwd (shell method), 197

Q

Queries cannot use both text and Geospatial Indexes (MongoDB system limit), 605
 query, **615**
 query optimizer, **615**
 query shape, **616**
 quit (shell method), 197

R

rawMongoProgramOutput (shell method), 185
 RDBMS, **616**
 read lock, **616**
 read preference, **616**
 record size, **616**
 recovering, **616**
 reIndex (database command), 317
 removeFile (shell method), 198
 removeShard (database command), 288
 renameCollection (database command), 299
 repairDatabase (database command), 319
 replica pairs, **616**
 replica set, **616**
 local database, 598
 replication, **616**
 replication lag, **616**
 replSetElect (database command), 366
 replSetFreeze (database command), 273
 replSetFresh (database command), 365
 replSetGetRBID (database command), 366
 replSetGetStatus (database command), 273
 replSetGetStatus.date (MongoDB reporting output), 274
 replSetGetStatus.members (MongoDB reporting output), 274
 replSetGetStatus.members.health (MongoDB reporting output), 274
 replSetGetStatus.members.lastHeartbeat (MongoDB reporting output), 275
 replSetGetStatus.members.lastHeartbeatMessage (MongoDB reporting output), 275
 replSetGetStatus.members.lastHeartbeatRecv (MongoDB reporting output), 275
 replSetGetStatus.members.name (MongoDB reporting output), 274
 replSetGetStatus.members.optime (MongoDB reporting output), 274
 replSetGetStatus.members.optime.i (MongoDB reporting output), 275
 replSetGetStatus.members.optime.t (MongoDB reporting output), 274
 replSetGetStatus.members.optimeDate (MongoDB reporting output), 275
 replSetGetStatus.members.pingMs (MongoDB reporting output), 275
 replSetGetStatus.members.self (MongoDB reporting output), 274
 replSetGetStatus.members.state (MongoDB reporting output), 274
 replSetGetStatus.members.stateStr (MongoDB reporting output), 274
 replSetGetStatus.members.uptime (MongoDB reporting output), 274

replSetGetStatus.myState (MongoDB reporting output), 274

replSetGetStatus.set (MongoDB reporting output), 274

replSetGetStatus.syncingTo (MongoDB reporting output), 275

replSetHeartbeat (database command), 366

replSetInitiate (database command), 275

replSetMaintenance (database command), 276

replSetReconfig (database command), 276

replSetStepDown (database command), 277

replSetSyncFrom (database command), 278

replSetTest (database command), 371

resetDbpath (shell method), 195

resetError (database command), 237

resident memory, 616

resource, 616

REST, 616

Restriction on Collection Names (MongoDB system limit), 609

Restrictions on Database Names for Unix and Linux Systems (MongoDB system limit), 609

Restrictions on Database Names for Windows (MongoDB system limit), 609

Restrictions on Field Names (MongoDB system limit), 609

resync (database command), 279

revokePrivilegesFromRole (database command), 265

revokeRolesFromRole (database command), 268

revokeRolesFromUser (database command), 256

role, 616

rolesInfo (database command), 270

rolesInfo.db (MongoDB reporting output), 270

rolesInfo.indirectRoles (MongoDB reporting output), 271

rolesInfo.isBuiltin (MongoDB reporting output), 271

rolesInfo.privileges (MongoDB reporting output), 271

rolesInfo.role (MongoDB reporting output), 270

rolesInfo.roles (MongoDB reporting output), 270

rollback, 616

ROOT (system variable available in aggregation), 493

rs.add (shell method), 164

rs.addArb (shell method), 165

rs.conf (shell method), 165

rs.config (shell method), 165

rs.freeze (shell method), 165

rs.help (shell method), 166

rs.initiate (shell method), 166

rs.printReplicationInfo (shell method), 166

rs.printSlaveReplicationInfo (shell method), 166

rs.reconfig (shell method), 167

rs.remove (shell method), 168

rs.slaveOk (shell method), 168

rs.status (shell method), 168

rs.stepDown (shell method), 168

rs.syncFrom (shell method), 169

run (shell method), 185

runMongoProgram (shell method), 185

runProgram (shell method), 185

S

secondary, 616

secondary index, 616

serverStatus (database command), 347

serverStatus.asserts (MongoDB reporting output), 356

serverStatus.asserts.msg (MongoDB reporting output), 357

serverStatus.asserts.regular (MongoDB reporting output), 357

serverStatus.asserts.rollovers (MongoDB reporting output), 357

serverStatus.asserts.user (MongoDB reporting output), 357

serverStatus.asserts.warning (MongoDB reporting output), 357

serverStatus.backgroundFlushing (MongoDB reporting output), 353

serverStatus.backgroundFlushing.average_ms (MongoDB reporting output), 353

serverStatus.backgroundFlushing.flushes (MongoDB reporting output), 353

serverStatus.backgroundFlushing.last_finished (MongoDB reporting output), 354

serverStatus.backgroundFlushing.last_ms (MongoDB reporting output), 353

serverStatus.backgroundFlushing.total_ms (MongoDB reporting output), 353

serverStatus.connections (MongoDB reporting output), 352

serverStatus.connections.available (MongoDB reporting output), 352

serverStatus.connections.current (MongoDB reporting output), 352

serverStatus.connections.totalCreated (MongoDB reporting output), 352

serverStatus.cursors (MongoDB reporting output), 354

serverStatus.cursors.clientCursors_size (MongoDB reporting output), 354

serverStatus.cursors.note (MongoDB reporting output), 354

serverStatus.cursors.pinned (MongoDB reporting output), 354

serverStatus.cursors.timedOut (MongoDB reporting output), 354

serverStatus.cursors.totalNoTimeout (MongoDB reporting output), 354

serverStatus.cursors.totalOpen (MongoDB reporting output), 354

serverStatus.dur (MongoDB reporting output), 357

- serverStatus.dur.commits (MongoDB reporting output), 357
- serverStatus.dur.commitsInWriteLock (MongoDB reporting output), 358
- serverStatus.dur.compression (MongoDB reporting output), 358
- serverStatus.dur.earlyCommits (MongoDB reporting output), 358
- serverStatus.dur.journaldMB (MongoDB reporting output), 358
- serverStatus.dur.timeMS (MongoDB reporting output), 358
- serverStatus.dur.timeMS.dt (MongoDB reporting output), 358
- serverStatus.dur.timeMS.prepLogBuffer (MongoDB reporting output), 358
- serverStatus.dur.timeMS.remapPrivateView (MongoDB reporting output), 358
- serverStatus.dur.timeMS.writeToDataFiles (MongoDB reporting output), 358
- serverStatus.dur.timeMS.writeToJournal (MongoDB reporting output), 358
- serverStatus.dur.writeToDataFilesMB (MongoDB reporting output), 358
- serverStatus.extra_info (MongoDB reporting output), 352
- serverStatus.extra_info.heap_usage_bytes (MongoDB reporting output), 352
- serverStatus.extra_info.note (MongoDB reporting output), 352
- serverStatus.extra_info.page_faults (MongoDB reporting output), 352
- serverStatus.globalLock (MongoDB reporting output), 350
- serverStatus.globalLock.activeClients (MongoDB reporting output), 351
- serverStatus.globalLock.activeClients.readers (MongoDB reporting output), 351
- serverStatus.globalLock.activeClients.total (MongoDB reporting output), 351
- serverStatus.globalLock.activeClients.writers (MongoDB reporting output), 351
- serverStatus.globalLock.currentQueue (MongoDB reporting output), 350
- serverStatus.globalLock.currentQueue.readers (MongoDB reporting output), 350
- serverStatus.globalLock.currentQueue.total (MongoDB reporting output), 350
- serverStatus.globalLock.currentQueue.writers (MongoDB reporting output), 350
- serverStatus.globalLock.lockTime (MongoDB reporting output), 350
- serverStatus.globalLock.ratio (MongoDB reporting output), 350
- serverStatus.globalLock.totalTime (MongoDB reporting output), 350
- serverStatus.host (MongoDB reporting output), 347
- serverStatus.indexCounters (MongoDB reporting output), 352
- serverStatus.indexCounters.accesses (MongoDB reporting output), 353
- serverStatus.indexCounters.hits (MongoDB reporting output), 353
- serverStatus.indexCounters.misses (MongoDB reporting output), 353
- serverStatus.indexCounters.missRatio (MongoDB reporting output), 353
- serverStatus.indexCounters.resets (MongoDB reporting output), 353
- serverStatus.localTime (MongoDB reporting output), 348
- serverStatus.locks (MongoDB reporting output), 348
- serverStatus.locks.. (MongoDB reporting output), 348
- serverStatus.locks...timeAcquiringMicros (MongoDB reporting output), 348
- serverStatus.locks...timeAcquiringMicros.R (MongoDB reporting output), 348
- serverStatus.locks...timeAcquiringMicros.W (MongoDB reporting output), 348
- serverStatus.locks...timeLockedMicros (MongoDB reporting output), 348
- serverStatus.locks...timeLockedMicros.R (MongoDB reporting output), 348
- serverStatus.locks...timeLockedMicros.r (MongoDB reporting output), 348
- serverStatus.locks...timeLockedMicros.W (MongoDB reporting output), 348
- serverStatus.locks...timeLockedMicros.w (MongoDB reporting output), 348
- serverStatus.locks.<database> (MongoDB reporting output), 349
- serverStatus.locks.<database>.timeAcquiringMicros (MongoDB reporting output), 350
- serverStatus.locks.<database>.timeAcquiringMicros.r (MongoDB reporting output), 350
- serverStatus.locks.<database>.timeAcquiringMicros.w (MongoDB reporting output), 350
- serverStatus.locks.<database>.timeLockedMicros (MongoDB reporting output), 349
- serverStatus.locks.<database>.timeLockedMicros.r (MongoDB reporting output), 349
- serverStatus.locks.<database>.timeLockedMicros.w (MongoDB reporting output), 349
- serverStatus.locks.admin (MongoDB reporting output), 348
- serverStatus.locks.admin.timeAcquiringMicros (MongoDB reporting output), 349
- serverStatus.locks.admin.timeAcquiringMicros.r (MongoDB reporting output), 349

<code>serverStatus.locks.admin.timeAcquiringMicros.w</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.document.returned</code> (MongoDB reporting output), 360
<code>serverStatus.locks.admin.timeLockedMicros</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.document.updated</code> (MongoDB reporting output), 360
<code>serverStatus.locks.admin.timeLockedMicros.r</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.getLastError</code> (MongoDB reporting output), 360
<code>serverStatus.locks.admin.timeLockedMicros.w</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.getLastError.wtime</code> (MongoDB reporting output), 360
<code>serverStatus.locks.local</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.getLastError.wtime.num</code> (MongoDB reporting output), 360
<code>serverStatus.locks.local.timeAcquiringMicros</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.getLastError.wtime.totalMillis</code> (MongoDB reporting output), 360
<code>serverStatus.locks.local.timeAcquiringMicros.r</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.getLastError.wtimeouts</code> (MongoDB reporting output), 361
<code>serverStatus.locks.local.timeAcquiringMicros.w</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.operation</code> (MongoDB reporting output), 361
<code>serverStatus.locks.local.timeLockedMicros</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.operation.fastmod</code> (MongoDB reporting output), 361
<code>serverStatus.locks.local.timeLockedMicros.r</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.operation.idhack</code> (MongoDB reporting output), 361
<code>serverStatus.locks.local.timeLockedMicros.w</code> (MongoDB reporting output), 349	<code>serverStatus.metrics.operation.scanAndOrder</code> (MongoDB reporting output), 361
<code>serverStatus.mem</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.queryExecutor</code> (MongoDB reporting output), 361
<code>serverStatus.mem.bits</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.queryExecutor.scanned</code> (MongoDB reporting output), 361
<code>serverStatus.mem.mapped</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.record</code> (MongoDB reporting output), 361
<code>serverStatus.mem.mappedWithJournal</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.record.moves</code> (MongoDB reporting output), 361
<code>serverStatus.mem.resident</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.repl</code> (MongoDB reporting output), 361
<code>serverStatus.mem.supported</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.repl.apply</code> (MongoDB reporting output), 361
<code>serverStatus.mem.virtual</code> (MongoDB reporting output), 351	<code>serverStatus.metrics.repl.apply.batches</code> (MongoDB reporting output), 361
<code>serverStatus.metrics</code> (MongoDB reporting output), 360	<code>serverStatus.metrics.repl.apply.batches.num</code> (MongoDB reporting output), 361
<code>serverStatus.metrics.cursor</code> (MongoDB reporting output), 363	<code>serverStatus.metrics.repl.apply.batches.totalMillis</code> (MongoDB reporting output), 361
<code>serverStatus.metrics.cursor.open</code> (MongoDB reporting output), 363	<code>serverStatus.metrics.repl.apply.ops</code> (MongoDB reporting output), 361
<code>serverStatus.metrics.cursor.open.noTimeout</code> (MongoDB reporting output), 363	<code>serverStatus.metrics.repl.buffer</code> (MongoDB reporting output), 361
<code>serverStatus.metrics.cursor.open.pinned</code> (MongoDB reporting output), 364	<code>serverStatus.metrics.repl.buffer.count</code> (MongoDB reporting output), 362
<code>serverStatus.metrics.cursor.open.open.total</code> (MongoDB reporting output), 364	<code>serverStatus.metrics.repl.buffer.maxSizeBytes</code> (MongoDB reporting output), 362
<code>serverStatus.metrics.cursor.timedOut</code> (MongoDB reporting output), 363	<code>serverStatus.metrics.repl.buffer.sizeBytes</code> (MongoDB reporting output), 362
<code>serverStatus.metrics.document</code> (MongoDB reporting output), 360	<code>serverStatus.metrics.repl.network</code> (MongoDB reporting output), 362
<code>serverStatus.metrics.document.deleted</code> (MongoDB reporting output), 360	<code>serverStatus.metrics.repl.network.bytes</code> (MongoDB reporting output), 362
<code>serverStatus.metrics.document.inserted</code> (MongoDB reporting output), 360	

`serverStatus.metrics.repl.network.getmores` (MongoDB reporting output), 362
`serverStatus.metrics.repl.network.getmores.num` (MongoDB reporting output), 362
`serverStatus.metrics.repl.network.getmores.totalMillis` (MongoDB reporting output), 362
`serverStatus.metrics.repl.network.ops` (MongoDB reporting output), 362
`serverStatus.metrics.repl.network.readersCreated` (MongoDB reporting output), 362
`serverStatus.metrics.repl.oplog` (MongoDB reporting output), 362
`serverStatus.metrics.repl.oplog.insert` (MongoDB reporting output), 362
`serverStatus.metrics.repl.oplog.insert.num` (MongoDB reporting output), 362
`serverStatus.metrics.repl.oplog.insert.totalMillis` (MongoDB reporting output), 362
`serverStatus.metrics.repl.oplog.insertBytes` (MongoDB reporting output), 362
`serverStatus.metrics.repl.preload` (MongoDB reporting output), 362
`serverStatus.metrics.repl.preload.docs` (MongoDB reporting output), 363
`serverStatus.metrics.repl.preload.docs.num` (MongoDB reporting output), 363
`serverStatus.metrics.repl.preload.docs.totalMillis` (MongoDB reporting output), 363
`serverStatus.metrics.repl.preload.indexes` (MongoDB reporting output), 363
`serverStatus.metrics.repl.preload.indexes.num` (MongoDB reporting output), 363
`serverStatus.metrics.repl.preload.indexes.totalMillis` (MongoDB reporting output), 363
`serverStatus.metrics.ttl` (MongoDB reporting output), 363
`serverStatus.metrics.ttl.deletedDocuments` (MongoDB reporting output), 363
`serverStatus.metrics.ttl.passes` (MongoDB reporting output), 363
`serverStatus.network` (MongoDB reporting output), 354
`serverStatus.network.bytesIn` (MongoDB reporting output), 354
`serverStatus.network.bytesOut` (MongoDB reporting output), 354
`serverStatus.network.numRequests` (MongoDB reporting output), 354
`serverStatus.opcounters` (MongoDB reporting output), 356
`serverStatus.opcounters.command` (MongoDB reporting output), 356
`serverStatus.opcounters.delete` (MongoDB reporting output), 356
`serverStatus.opcounters.getmore` (MongoDB reporting output), 356
`serverStatus.opcounters.insert` (MongoDB reporting output), 356
`serverStatus.opcounters.query` (MongoDB reporting output), 356
`serverStatus.opcounters.update` (MongoDB reporting output), 356
`serverStatus.opcountersRepl` (MongoDB reporting output), 355
`serverStatus.opcountersRepl.command` (MongoDB reporting output), 356
`serverStatus.opcountersRepl.delete` (MongoDB reporting output), 356
`serverStatus.opcountersRepl.getmore` (MongoDB reporting output), 356
`serverStatus.opcountersRepl.insert` (MongoDB reporting output), 355
`serverStatus.opcountersRepl.query` (MongoDB reporting output), 355
`serverStatus.opcountersRepl.update` (MongoDB reporting output), 355
`serverStatus.process` (MongoDB reporting output), 347
`serverStatus.recordStats` (MongoDB reporting output), 358
`serverStatus.recordStats.<database>.accessesNotInMemory` (MongoDB reporting output), 359
`serverStatus.recordStats.<database>.pageFaultExceptionsThrown` (MongoDB reporting output), 359
`serverStatus.recordStats.accessesNotInMemory` (MongoDB reporting output), 359
`serverStatus.recordStats.admin.accessesNotInMemory` (MongoDB reporting output), 359
`serverStatus.recordStats.admin.pageFaultExceptionsThrown` (MongoDB reporting output), 359
`serverStatus.recordStats.local.accessesNotInMemory` (MongoDB reporting output), 359
`serverStatus.recordStats.local.pageFaultExceptionsThrown` (MongoDB reporting output), 359
`serverStatus.recordStats.pageFaultExceptionsThrown` (MongoDB reporting output), 359
`serverStatus.repl` (MongoDB reporting output), 355
`serverStatus.repl.hosts` (MongoDB reporting output), 355
`serverStatus.repl.ismaster` (MongoDB reporting output), 355
`serverStatus.repl.secondary` (MongoDB reporting output), 355
`serverStatus.repl.setName` (MongoDB reporting output), 355
`serverStatus.uptime` (MongoDB reporting output), 347
`serverStatus.uptimeEstimate` (MongoDB reporting output), 347
`serverStatus.version` (MongoDB reporting output), 347
`serverStatus.workingSet` (MongoDB reporting output), 359

- ul style="list-style-type: none; padding-left: 0;">
- serverStatus.workingSet.computationTimeMicros (MongoDB reporting output), 360
- serverStatus.workingSet.note (MongoDB reporting output), 359
- serverStatus.workingSet.overSeconds (MongoDB reporting output), 360
- serverStatus.workingSet.pagesInMemory (MongoDB reporting output), 359
- serverStatus.writeBacksQueued (MongoDB reporting output), 357
- set name, 616
- setParameter (database command), 318
- setShardVersion (database command), 291
- sh._adminCommand (shell method), 172
- sh._checkFullName (shell method), 172
- sh._checkMongos (shell method), 172
- sh._lastMigration (shell method), 172
- sh._lastMigration._id (MongoDB reporting output), 172
- sh._lastMigration.clientAddr (MongoDB reporting output), 173
- sh._lastMigration.details (MongoDB reporting output), 173
- sh._lastMigration.ns (MongoDB reporting output), 173
- sh._lastMigration.server (MongoDB reporting output), 172
- sh._lastMigration.time (MongoDB reporting output), 173
- sh._lastMigration.what (MongoDB reporting output), 173
- sh.addShard (shell method), 173
- sh.addShardTag (shell method), 174
- sh.addTagRange (shell method), 174
- sh.disableBalancing (shell method), 175
- sh.enableBalancing (shell method), 175
- sh.enableSharding (shell method), 175
- sh.getBalancerHost (shell method), 176
- sh.getBalancerState (shell method), 176
- sh.help (shell method), 176
- sh.isBalancerRunning (shell method), 177
- sh.moveChunk (shell method), 177
- sh.removeShardTag (shell method), 178
- sh.setBalancerState (shell method), 178
- sh.shardCollection (shell method), 178
- sh.splitAt (shell method), 179
- sh.splitFind (shell method), 179
- sh.startBalancer (shell method), 180
- sh.status (shell method), 180
- sh.status.databases._id (MongoDB reporting output), 182
- sh.status.databases.chunk-details (MongoDB reporting output), 182
- sh.status.databases.chunks (MongoDB reporting output), 182
- sh.status.databases.partitioned (MongoDB reporting output), 182
- sh.status.databases.primary (MongoDB reporting output), 182
- sh.status.databases.shard-key (MongoDB reporting output), 182
- sh.status.databases.tag (MongoDB reporting output), 182
- sh.status.sharding-version._id (MongoDB reporting output), 181
- sh.status.sharding-version.clusterId (MongoDB reporting output), 182
- sh.status.sharding-version.currentVersion (MongoDB reporting output), 181
- sh.status.sharding-version.minCompatibleVersion (MongoDB reporting output), 181
- sh.status.sharding-version.version (MongoDB reporting output), 181
- sh.status.shards._id (MongoDB reporting output), 182
- sh.status.shards.host (MongoDB reporting output), 182
- sh.status.shards.tags (MongoDB reporting output), 182
- sh.stopBalancer (shell method), 182
- sh.waitForBalancer (shell method), 183
- sh.waitForBalancerOff (shell method), 183
- sh.waitForDLock (shell method), 184
- sh.waitForPingChange (shell method), 184
- shard, 616
- shard key, 617
- Shard Key Index Type (MongoDB system limit), 607
- Shard Key is Immutable (MongoDB system limit), 607
- Shard Key Size (MongoDB system limit), 607
- Shard Key Value in a Document is Immutable (MongoDB system limit), 608
- shardCollection (database command), 291
- shardConnPoolStats (database command), 330
- shardConnPoolStats.createdByType (MongoDB reporting output), 330
- shardConnPoolStats.createdByType.master (MongoDB reporting output), 330
- shardConnPoolStats.createdByType.set (MongoDB reporting output), 331
- shardConnPoolStats.createdByType.sync (MongoDB reporting output), 331
- shardConnPoolStats.hosts (MongoDB reporting output), 330
- shardConnPoolStats.hosts.<host>.available (MongoDB reporting output), 330
- shardConnPoolStats.hosts.<host>.created (MongoDB reporting output), 330
- shardConnPoolStats.replicaSets (MongoDB reporting output), 330
- shardConnPoolStats.replicaSets.<name>.host (MongoDB reporting output), 330
- shardConnPoolStats.replicaSets.<name>.host[n].addr (MongoDB reporting output), 330
- shardConnPoolStats.replicaSets.<name>.host[n].hidden (MongoDB reporting output), 330
- shardConnPoolStats.replicaSets.<name>.host[n].ismaster (MongoDB reporting output), 330

- shardConnPoolStats.replicaSets.<name>.host[n].ok (MongoDB reporting output), 330
 - shardConnPoolStats.replicaSets.<name>.host[n].pingTimeMillis (MongoDB reporting output), 330
 - shardConnPoolStats.replicaSets.<name>.host[n].secondary (MongoDB reporting output), 330
 - shardConnPoolStats.replicaSets.<name>.host[n].tags (MongoDB reporting output), 330
 - shardConnPoolStats.threads (MongoDB reporting output), 331
 - shardConnPoolStats.threads.hosts (MongoDB reporting output), 331
 - shardConnPoolStats.threads.hosts.avail (MongoDB reporting output), 331
 - shardConnPoolStats.threads.hosts.created (MongoDB reporting output), 331
 - shardConnPoolStats.threads.hosts.host (MongoDB reporting output), 331
 - shardConnPoolStats.threads.seenNS (MongoDB reporting output), 331
 - shardConnPoolStats.totalAvailable (MongoDB reporting output), 331
 - shardConnPoolStats.totalCreated (MongoDB reporting output), 331
 - sharded cluster, 617
 - sharding, 617
 - config database, 593
 - Sharding Existing Collection Data Size (MongoDB system limit), 607
 - shardingState (database command), 292
 - shell helper, 617
 - shutdown (database command), 321
 - Single Document Modification Operations in Sharded Collections (MongoDB system limit), 607
 - single-master replication, 617
 - Size of Namespace File (MongoDB system limit), 604
 - slave, 617
 - sleep (database command), 370
 - Sorted Documents (MongoDB system limit), 608
 - Spherical Polygons must fit within a hemisphere. (MongoDB system limit), 608
 - split, 617
 - split (database command), 293
 - splitChunk (database command), 295
 - splitVector (database command), 296
 - SQL, 617
 - SSD, 617
 - stale, 617
 - standalone, 617
 - stopMongod (shell method), 185
 - stopMongoProgram (shell method), 185
 - stopMongoProgramByPid (shell method), 185
 - strict consistency, 617
 - sync, 617
 - syslog, 617
 - system
 - collections, 600
 - namespace, 600
 - system.indexes.key (MongoDB reporting output), 45
 - system.indexes.name (MongoDB reporting output), 45
 - system.indexes.ns (MongoDB reporting output), 45
 - system.indexes.v (MongoDB reporting output), 45
- ## T
- tag, 617
 - text (database command), 235
 - top (database command), 337
 - touch (database command), 321
 - TSV, 617
 - TTL, 618
- ## U
- unique index, 618
 - Unique Indexes in Sharded Collections (MongoDB system limit), 607
 - unordered query plan, 618
 - unsetSharding (database command), 293
 - update (database command), 222
 - update.n (MongoDB reporting output), 225
 - update.nModified (MongoDB reporting output), 225
 - update.ok (MongoDB reporting output), 225
 - update.upserted (MongoDB reporting output), 225
 - update.upserted._id (MongoDB reporting output), 225
 - update.upserted.index (MongoDB reporting output), 225
 - update.writeConcernError (MongoDB reporting output), 226
 - update.writeConcernError.code (MongoDB reporting output), 226
 - update.writeConcernError.errmsg (MongoDB reporting output), 226
 - update.writeErrors (MongoDB reporting output), 225
 - update.writeErrors.code (MongoDB reporting output), 225
 - update.writeErrors.errmsg (MongoDB reporting output), 225
 - update.writeErrors.index (MongoDB reporting output), 225
 - updateRole (database command), 260
 - updateUser (database command), 252
 - upsert, 618
 - usePowerOf2Sizes, 316
 - usePowerOf2Sizes (collection flag), 316
 - userInfo (database command), 257
 - UUID (shell method), 186
- ## V
- validate (database command), 335

validate.bytesWithHeaders (MongoDB reporting output), 337
 validate.bytesWithoutHeaders (MongoDB reporting output), 337
 validate.datasize (MongoDB reporting output), 336
 validate.deletedCount (MongoDB reporting output), 337
 validate.deletedSize (MongoDB reporting output), 337
 validate.errors (MongoDB reporting output), 337
 validate.extentCount (MongoDB reporting output), 335
 validate.extents (MongoDB reporting output), 335
 validate.extents.firstRecord (MongoDB reporting output), 336
 validate.extents.lastRecord (MongoDB reporting output), 336
 validate.extents.loc (MongoDB reporting output), 336
 validate.extents.nsdiag (MongoDB reporting output), 336
 validate.extents.size (MongoDB reporting output), 336
 validate.extents.xnext (MongoDB reporting output), 336
 validate.extents.xprev (MongoDB reporting output), 336
 validate.firstExtent (MongoDB reporting output), 335
 validate.firstExtentDetails (MongoDB reporting output), 336
 validate.firstExtentDetails.firstRecord (MongoDB reporting output), 337
 validate.firstExtentDetails.lastRecord (MongoDB reporting output), 337
 validate.firstExtentDetails.loc (MongoDB reporting output), 336
 validate.firstExtentDetails.nsdiag (MongoDB reporting output), 336
 validate.firstExtentDetails.size (MongoDB reporting output), 336
 validate.firstExtentDetails.xnext (MongoDB reporting output), 336
 validate.firstExtentDetails.xprev (MongoDB reporting output), 336
 validate.invalidObjects (MongoDB reporting output), 337
 validate.keysPerIndex (MongoDB reporting output), 337
 validate.lastExtent (MongoDB reporting output), 335
 validate.lastExtentSize (MongoDB reporting output), 336
 validate.nIndexes (MongoDB reporting output), 337
 validate.nrecords (MongoDB reporting output), 336
 validate.ns (MongoDB reporting output), 335
 validate.objectsFound (MongoDB reporting output), 337
 validate.ok (MongoDB reporting output), 337
 validate.padding (MongoDB reporting output), 336
 validate.valid (MongoDB reporting output), 337
 version (shell method), 195
 virtual memory, 618
 whatsmysuri (database command), 344
 working set, 618
 write concern, 618
 write lock, 618
 writebacklisten (database command), 367
 writeBacks, 618
 writeBacksQueued (database command), 366
 writeBacksQueued.hasOpsQueued (MongoDB reporting output), 366
 writeBacksQueued.queue (MongoDB reporting output), 366
 writeBacksQueued.queue.minutesSinceLastCall (MongoDB reporting output), 366
 writeBacksQueued.queue.n (MongoDB reporting output), 366
 writeBacksQueued.totalOpsQueued (MongoDB reporting output), 366
 WriteResult (shell method), 188
 WriteResult._id (MongoDB reporting output), 188
 WriteResult.hasWriteConcernError (shell method), 189
 WriteResult.hasWriteError (shell method), 189
 WriteResult.nInserted (MongoDB reporting output), 188
 WriteResult.nMatched (MongoDB reporting output), 188
 WriteResult.nModified (MongoDB reporting output), 188
 WriteResult.nRemoved (MongoDB reporting output), 188
 WriteResult.nUpserted (MongoDB reporting output), 188
 WriteResult.writeConcernError (MongoDB reporting output), 188
 WriteResult.writeConcernError.code (MongoDB reporting output), 188
 WriteResult.writeConcernError.errInfo (MongoDB reporting output), 188
 WriteResult.writeError (MongoDB reporting output), 188
 WriteResult.writeError.code (MongoDB reporting output), 188
 WriteResult.writeError.errmsg (MongoDB reporting output), 188

W

waitMongoProgramOnPort (shell method), 185
 waitProgram (shell method), 186
 WGS84, 618